

Documentation

[Home](#)  [Download PDF](#) [Sop: knowledge-first discovery & context preservation protocol](#) [Table of contents](#) [Amazon machine images \(ami\) design guide](#) [Google antigraivty skills integration guide](#) [System architecture](#) [Auto scaling groups \(asgs\) & separation of concerns](#) [Aws phased adoption roadmap & costing guide](#) [Aws services vs. on-premises open-source comparison guide](#) [Strategic comparative review: aws-native managed platform vs. self-hosted custom stack](#) [Modernizing big data analytics architecture: 3-tier infra deployment blueprint](#) [Ci/cd pipeline documentation](#) [Context7 ai chat integration guide](#) [Costing estimate](#) [Developer design mapping](#) [Disaster recovery options & sovereignty guide](#) [Gitlab ci/cd & aws efs code deployment](#) [Hybrid cloud integration & costing guide](#) [Autonomous ai pair-programming & multi-agent operations with google jules](#) [Secure developer access guide](#) [Legal notice & disclaimer](#) [Software licensing & technology risk register \(ts/mc series\)](#) [Load testing assumptions & sizing guide](#) [Opentofu migration & aws support research guide](#) [Rds postgresql 17 vs. percona server for postgresql 17](#) [All project documentation](#) [Ragflow + langfuse: aws vs. on-premises gpu integration guide](#) [Redis vs. valkey: architectural and cost comparison](#) [Root terraform configuration](#) [Route 53 & dns troubleshooting](#) [Automation scripts](#) [Start here: pragmatic onboarding & execution standard](#) [Technology stack comparison guide](#) [Wazuh siem & xdr deep-dive: cloud, on-premises & av compatibility](#) [Wazuh standalone cloud installation & costing guide](#)



Tutorials

[Quickstart tutorial: operating project automation utilities](#)



How-To Guides

[How to guide: compiling llm friendly curation formats](#) [How to guide: standardising markdown front matter metadata](#)



Reference Specs

[Technical reference: deployment, teardown, and pdf scripts](#) [Technical reference: parse llms.py compiler](#) [Technical reference: prepare docs.py metadata formatter](#)



Explanations

[Project automation: metadata and llm curation architecture](#) [The diátaxis framework in our project documentation](#)

Onsite On-Premises

[On premises infrastructure with ansible](#) [Onsite on premises vm & network architecture](#) [Onsite on premises blueprint portal](#) [Open source on premise stack specs](#) [Enterprise percona server for postgresql 17 on on premises](#) [Rootless podman 5+ & systemd quadlets](#)

Security Posture & Audits

[Output of asimp](#) [Output of lynis](#) [Output of openscap](#) [Security posture assessment \(spa\) requirement checklist](#)

Executive Blueprint

[Executive & financial blueprint portal](#) [Sovereignty & regulatory compliance](#) [36-month tco & financial blueprint](#)

Engineering Runbook

[Explicit ai agent data flow mapping](#) [Ami hardening compliance & ansible playbooks](#) [Dns & systemd-resolved troubleshooting](#) [Persistent efs storage & gitlab ci/cd](#) [Devops implementation runbook portal](#) [Opentofu aws infrastructure building, simulation, & multi-agent collaboration runbook](#) [Opentofu module manifests](#)

Modules

[Application load balancer \(alb\) module](#) [Auto scaling group \(asg\) module](#) [Elasticache valkey module](#) [Ssh jumphost module](#) [Rds module](#) [Security groups module](#) [Standalone ec2 instance module](#) [Vpc module](#) [Web application firewall \(waf\) module](#)

E-Mail

contact@linuxmalaysia.com

Join Us

Community

Join the **cmsformerd** discussion group.

All Project Documentation

OpenTofu >= 1.6.0

AWS: ap-southeast-5

WAFv2: Enabled

Graviton: ARM64

GUIDE



AWS Secure 3-Tier Architecture Documentation

#aws #cloud #architecture #vpc #alb #asg #rds #waf #elasticache #valkey #jump host #bastion #route53 #dns #ssl #acm #disaster-recovery #gitlab #efs #postgresql #gpu #ragflow #langfuse #antigravity #skills #sovereignty #compliance #costing

AWS Secure 3-Tier Architecture Documentation

Welcome to the official technical documentation for our **AWS 3-Tier Deployment for AI & Web Infra** project. This project is optimized for deployment in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)** utilizing AWS Graviton (ARM64) instances, Multi-AZ RDS Postgres, and AWS WAFv2 regional protection.

This deployment is structured natively in OpenTofu, adhering to strict modular boundaries, security best practices, and the **Zero-Trust Network Principle**.

Technical Overview

The architecture divides infrastructure components into discrete logical and physical tiers to achieve top-tier scalability, performance, and threat mitigation:

- Presentation / Web Layer:** Application Load Balancer (ALB) receiving external traffic and filtered through AWS WAFv2 (OWASP rules + rate limiting).
- Application / Compute Layer:** Auto Scaling Group (ASG) of EC2 instances running inside private subnets, auto-scaled on CPU usage, and managed via AWS Systems Manager (SSM).
- Database Layer:** Multi-AZ RDS PostgreSQL database deployed within isolated subnets, allowing connections exclusively from the application tier.

Documentation Index

Explore different sections of our infrastructure documentation:

Core Configuration

- [Start Here: Onboarding Standard](#): Dual-audience onboarding guide adapting Diátaxis learn-by-doing principle for Human Engineers and AI Agents.
- [System Architecture](#): Deep dive into the physical network structure, routing tables, and AWS resource layout in the Malaysia region, including how the Developer's first design is mapped.
- [AWS Phased Adoption Roadmap & Costing Guide](#): Multi-year week-by-week and month-by-month AWS service growth plan mapped from the project Gantt chart.
- [Developer Design Alignment Guide](#): Rationale and comparison of shifting from legacy single-node Ubuntu VMs to an enterprise secure AWS design.
- [ASGs & Separation of Concerns Guide](#): Best practice guide detailing auto-scaling with distinct ASGs, stateless principles, and the role of Amazon S3, EFS, or both.
- [Root OpenTofu/Terraform Files](#): Overview of the root configuration entries (main.tf, variables.tf, outputs.tf, providers.tf).
- [OpenTofu Migration Guide](#): Detailed compatibility research and transition path for deploying using OpenTofu on AWS.
- [AMI Design & Hardening Guide](#): Architectural guide outlining the pre-baked AMI strategy, Packer/Ansible pipeline, and ASIMP compliance integration.
- [Route 53 & DNS Troubleshooting](#): Deep dive into custom domain integration, ACM SSL/TLS setup, and extensive research on resolving ASG private subnet DNS resolution failures.
- [Secure Developer Access Guide](#): Comprehensive guide on using our secure SSH Jump host (Bastion) to access private and standalone nodes from Cyberjaya, with Windows/macOS/Linux client setups and private key security guidelines.
- [Hybrid Cloud Integration Guide](#): Comprehensive evaluation of secure, cost-optimized API and MCP connections alongside official AWS hybrid network solutions (VPN, Direct Connect, Transit Gateway) with granular MYR/USD costing models for ap-southeast-5.
- [Disaster Recovery Options & National Sovereignty Guide](#): Production-ready playbook covering the 4 standard AWS cloud DR options aligned with our Multi-AZ architecture, detailed local sovereignty/PDPA/CBPDT compliance reviews, and granular USD/MYR costing comparisons.
- [RDS PostgreSQL 17 vs. Percona Server for PostgreSQL 17 Guide](#): Comprehensive technical comparison of performance, telemetry, observability (PMM vs CloudWatch/Performance Insights), architectural designs, and costs in USD/MYR for ap-southeast-5.
- [Redis vs. Valkey Comparison](#): Comprehensive strategic, licensing, and costing comparison (Baseline vs. High-Performance) of Redis OSS vs. Valkey on AWS and on-premises.
- [RAGFlow + Langfuse GPU Guide](#): Architectural and economic analysis of RAGFlow and Langfuse, detailing critical GPU utilization, AWS vs. On-Premises deployment models, and secure API/MCP hybrid connections.
- [Google Antigravity Skills Guide](#): Comprehensive integration guide outlining how to share and deploy agent knowledge bases and skills between Google Jules and Google Antigravity.
- [SOP: Knowledge-First Discovery](#): Standard Operating Procedure outlining how AI agents perform local documentation search before probing remote targets.
- [Wazuh Standalone Cloud Installation & Costing Guide](#): Detailed installation steps for standalone Wazuh in the cloud (AWS Marketplace AMI and Graviton assistant modes), with independent USD/MYR costing tables for isolated financial tracking.
- [Wazuh SIEM & XDR Deep-Dive Guide](#): In-depth functional breakdown of Wazuh SIEM & XDR capabilities, deployment modes (cloud/on-prem), and critical operational guidance regarding Antivirus coexistence, passive mode configuration, and Windows Defender integration.
- [Technology Stack Comparison Guide](#): Complete architectural mapping and AWS alternatives comparison for the developer's containerized and external integrations (LangChain4j, Spring Boot, React, Twilio, Meta, Postgres, Valkey).
- [Software Licensing & Technology Risk Register \(TS/MC Series\)](#): Complete software licensing compliance framework, technology risk registry, and mitigation plans (TS/MC Series) covering LangChain4j, self-hosted operations, Bedrock with Qwen3 models, and standalone Wazuh SIEM.
- [Strategic Comparative Review](#): Comprehensive strategic analysis and financial TCO comparison of an AWS-Native Managed Platform against a Self-Hosted / On-Premises Custom Stack in Malaysia.
- [Load Testing Assumptions & Sizing Guide](#): Workload definitions, SLA metrics, architectural performance assumptions, and multi-VU sizing models from 100 to 10,000 VUs.
- [Context7 AI Chat Integration Guide](#): Guide detailing our AI chat widget, how to use it, the provider (Upstash), and integration details.
- [Security Posture Assessment \(SPA\) Requirement Checklist](#): Comprehensive security control audit checklist, fully customized for our Java, Spring Boot, PostgreSQL, Valkey, and RAGFlow/Langfuse AI infrastructure.
- [Legal Notice & Disclaimer](#): Official legal disclaimer and privacy policy confirming that all designs, costs, and scenarios are based entirely on assumptions for training and informational purposes.
- [Output of ASIMP](#): Example execution output, report format, and baseline metrics generated by ASIMP.
- [Output of Lynis](#): High-fidelity example of the Lynis host auditing utility scanning logs, attributes, and suggestions.
- [Output of OpenSCAP](#): Detailed example of OpenSCAP CIS Level 2 scan evaluation rules, results, and generated remediations.
- [AWS Services vs. On-Premises Open-Source Stack Comparison](#): Detailed layer-by-layer architectural comparison across 12 layers mapping AWS services to self-hosted open-source counterparts.

31. [Google Jules AI Platform Guide](#): Comprehensive technical showcase documenting our end-to-end development workflow, PR review collaboration, DSOM governance, and Google Antigravity integration.
32. [OpenTofu AWS Simulation & Multi-Agent Collaboration Runbook](#): Comprehensive runbook detailing offline AWS deployment simulations, unit testing matrices, static HCL AST analysis, and multi-agent/multi-human branch/PR workflows.
33. [Big Data Analytics Lakehouse Architecture](#): Modernizing Big Data Analytics into a 100% open-source Lakehouse with S3 WORM Object Lock, Iceberg, Polaris, Trino, Spark/Sedona, MCP AI quarantine sandboxing, and 3 distinct infrastructure solution blueprints (Cloud, Hybrid, On-Prem Proxmox/RKE2/Ceph).

Onsite On-Premises Volume

- [Onsite On-Premises Blueprint Portal](#): Overview portal and strategic rationale for moving from AWS to rootless, open-source local on-premises deployments.
- [VM & Network Architecture](#): Physical host configurations, sizing models, multi-tier VLAN isolation, and stateful routing rulesets.
- [Rootless Podman 5+ & systemd Quadlets](#): Configuration guidelines for running production containers without administrative root rights using keep-id remapping.
- [On-Premises Infrastructure with Ansible](#): Symmetric privilege separation, FQCN-compliant plays, and local CI/CD pipelines integrating Gitea, Semaphore, and ARA.
- [Open-Source Containerized Stack Specifications](#): Standard blueprints replacing core AWS managed components with containerized equivalents (BunkerWeb, Valkey, Postgres 17, Keycloak, Ollama, RAGFlow, and Langfuse).
- [Enterprise Percona Server Setup](#): Standard blueprints for running production-grade Percona Server for PostgreSQL 17 on-premises using Patroni cluster management, etcd distributed consensus stores, pg_backrest, and HAProxy load balancing.

Infrastructure Submodules

- [VPC Module](#): Core networking, public/private subnets, internet gateways, and NAT configurations.
- [Security Groups Module](#): Strict firewall rulesets and port-level isolation.
- [WAF Module](#): Layer-7 Web Application Firewall protecting the ALB.
- [ALB Module](#): Application Load Balancer and health-check configurations.
- [ASG Module](#): Auto Scaling Group, Launch Templates, and dynamic Graviton auto-detection.
- [RDS Module](#): Multi-AZ PostgreSQL configuration and parameter group tuning.
- [Standalone EC2 Module](#): Secure standalone Ubuntu 26.04 LTS development and application environments.
- [ElastiCache Valkey Module](#): Secure ElastiCache Valkey in-memory caching cluster for session and metadata store.
- [Jumpshot Module](#): Secure public-subnet SSH Jumpshot (Bastion) whitelisted for Cyberjaya office with automated downstream ingress configuration.

Deployment & CI/CD

- [Automation Scripts](#): Details about CLI helpers (`deploy.sh`, `destroy.sh`, `user_data.sh`).
- [CI/CD Pipeline](#): GitHub Actions workflow for automatic formatting, testing, validation, and OIDC deployment.
- [GitLab EFS CI/CD](#): Comprehensive guide on GitLab CI/CD, automatic workflows, EFS mounting, dynamic Nginx path configurations, and containerized/S3 alternatives.
- [Costing Estimate](#): Comprehensive monthly cost breakdown, local currency estimates, and Day-2 cost optimization pathways.

Prerequisites

Before deploying the infrastructure, ensure you have the following tools installed and configured:

- [OpenTofu](#) **>= 1.6.0** (Recommended) or [Terraform](#) **>= 1.5.0**
- [AWS CLI](#) configured with a least-privilege deployment role targeted to `ap-southeast-5` covering declared VPC, EC2, IAM, RDS, ElastiCache, ALB, WAF, and Route 53 operations (with narrowly scoped `iam:PassRole` permissions for instance profiles required by `scripts/deploy.sh`). Role credentials must be validated (`aws sts get-caller-identity`) prior to deployment.
- [Git](#) for repository and revision tracking

System Architecture

#aws #cloud #architecture #vpc #alb #asg #rds #waf #elasticache #valkey #dns #ssl #acm #disaster-recovery #efs #postgresql #ragflow #langfuse #costing

System Architecture

This document describes the high-availability 3-tier network topology, AWS component layouts, and routing architectures deployed by this project, fully aligned with our [Estimated Costing](#) model.

Additionally, this architecture is fully customized to map and host the **Developer’s First Design (RAGFlow + LangFuse AI stack)** in a secure, highly-available, and resilient manner, built using hardened **Ubuntu 26.04 LTS** nodes and integrated with **Amazon ElastiCache for Valkey**.

Developer’s First Design vs. AWS Production Architecture

In the developer’s initial design, the application was structured across four separate standalone virtual machine servers running on **Ubuntu Server 26.04 LTS**:

- **Server 01 (Frontend - Web Tier):** Nginx Web Server / Reverse Proxy (2 vCPU, 4GB RAM)
- **Server 02 (Backend - App Tier):** Backend + DMS + MCP (4 vCPU, 16GB RAM)
- **Server 03 (AI Tier):** RAGFlow + LangFuse (4 vCPU, 8GB RAM)
- **Server 04 (Database - Data Tier):** SQL Database (4 vCPU, 16GB RAM)

While simple, deploying this directly as four standalone VMs introduces single points of failure (SPOF), security vulnerabilities from direct public internet exposure, manual backup overhead, and a lack of scalability.

To make this suitable for enterprise AWS deployment **without changing the AWS requirements**, we have mapped each of these components directly to a secure, multi-AZ, managed 3-tier architecture:

Architectural Comparative Mapping

Developer’s Original Server	Original Spec	AWS Production-Ready Component	AWS Architectural Benefits
Server 01: Frontend	2 vCPU, 4GB RAM, Ubuntu	AWS WAFv2 + ALB + Private Nginx ASG & Dedicated Standalone Instance	No Public IPs & Local Baking Parity: Traffic enters via secure ALB/WAF. Private Nginx reverse-proxies run inside private subnets without public IPs. Paired with a dedicated Frontend Standalone instance connected to the same S3 bucket to test configurations and pre-bake custom Nginx ami-frontend-* images. Hardened via ASIMP .
Server 02: Backend	4 vCPU, 16GB RAM, Ubuntu	Multi-AZ ASG & Dedicated Standalone Instance	High Availability & 1:1 Database Parity: ASG spans multiple AZs and scales compute nodes dynamically. Paired with a dedicated Backend Standalone instance connected to the identical Multi-AZ RDS PostgreSQL database and S3 buckets to securely execute database migrations, test application logic, and pre-bake ami-backend-* images. Hardened via ASIMP .
Server 03: AI Tier	4 vCPU, 8GB RAM, Ubuntu	Private ASG Compute & Dedicated AI Tier Standalone Instance	Isolated Compute & Instant Scaling via EFS Caching: Securely hosted in Private App Subnets. Paired with a dedicated AI Standalone instance connected to the identical Amazon EFS filesystem, RDS, and S3. Developers warm up Hugging Face model weight caches directly onto EFS from the standalone instance so that auto-scaling ASG nodes can access them instantly. Pre-baked into ami-ai-* images. Hardened via ASIMP .
Server 04: Database	4 vCPU, 16GB RAM, Ubuntu	AWS RDS PostgreSQL (Multi-AZ) & Amazon ElastiCache for Valkey	Managed Resiliency & Enterprise Caching: Replaced self-managed SQL database with a fully managed Multi-AZ PostgreSQL database (db.m6g.xlarge). Synchronous replication, automated snapshots/failover, zero direct public route, and ingress restricted solely to private compute instances. Paired with a secure Amazon ElastiCache for Valkey cluster inside the database subnets to act as a high-performance in-memory task broker and session store for RAGFlow + LangFuse, reducing DB load and database overhead.

Standalone EC2 Instances for AMI Creation and Parity

To ensure seamless, reliable updates and zero-downtime rolling upgrades across our Auto Scaling Groups (ASGs), each application group (Frontend, Backend, and AI Tier) is paired with a dedicated **Standalone EC2 Instance** (running **Ubuntu 26.04 LTS** and hardened via the **ASIMP** framework).

These standalone instances are deployed directly inside the secure **VPC Private Application Subnets** and are configured to connect to the exact same shared resources (AWS RDS PostgreSQL Database, Amazon S3 Buckets, Amazon EFS shared storage, and Amazon ElastiCache Valkey caching) as their corresponding ASGs:


1. **Frontend Standalone Instance:**
 - Connected to **Amazon S3** to manage and verify static assets or remote template files.
 - Used to test Nginx routing/configurations before baking the ami-frontend-* image.
2. **Backend Standalone Instance:**
 - Connected to the **Multi-AZ RDS PostgreSQL** database (using identical DB endpoints and credentials), **Amazon S3** (using identical IAM Role permissions), and **Amazon ElastiCache Valkey** (port 6379).
 - Used to verify backend service scripts, run migrations, and test application logic before baking the ami-backend-* image.
3. **AI Tier Standalone Instance:**
 - Connected to **Amazon EFS** (via Private Mount Targets) to read/write persistent AI model caches, **Amazon S3** for training documents, **RDS PostgreSQL** for metadata storage, and **Amazon ElastiCache Valkey** (port 6379) for distributed task broker signaling.
 - Used to bootstrap RAGFlow + LangFuse dependencies, warm up Hugging Face/SentenceTransformers cache directories, and test container updates before baking the ami-ai-* image.

Architectural Advantages of Standalone-to-AMI Parity:

- **1:1 Environment Alignment:** Because each standalone instance connects to the exact same database, Valkey cache, and shared storage backends, developers can fully run, test, and validate configurations in a real production-like environment without any risk of deployment divergence.
- **Pre-Audited & Hardened Base:** Standalone instances serve as the staging template. Developers run the **ASIMP** auditing and hardening pipelines directly on these instances, verifying system integrity reports before triggering the Packer/AMI capture.
- **Zero-Downtime Releases:** Once the standalone instance is verified and the AMI is baked, updating the ASG Launch Template and running an Instance Refresh executes a safe rolling update across the live cluster.

Architectural Schematic

The updated network topology below outlines how our three ASG application groups and their matching standalone AMI-baking instances connect to the shared database, Amazon S3, Amazon EFS, and ElastiCache Valkey within our AWS secure environment:



Network Isolation Layers

1. Presentation / Web Layer (Public Subnets)

- **Subnets:** 10.0.1.0/24 (AZ ap-southeast-5a) and 10.0.2.0/24 (AZ ap-southeast-5b).
- **Description:** Hosts public-facing services and manages secure domain mappings. This layer routes inbound internet traffic directly through the Internet Gateway (IGW).
- **Resources:**
 - **Route 53 DNS Routing:** Manages custom domain delegations and points A Alias records to the ALB. Integrated with AWS Certificate Manager (ACM) for automatic domain verification.
 - **Application Load Balancer (ALB):** Terminates and routes incoming connections. Replaces the external Nginx reverse proxy direct exposure (from Server 01), dispersing HTTP/HTTPS traffic to private instances.
 - **NAT Gateway:** A highly-available NAT Gateway deployment in public subnets provides secure outbound internet access for package retrieval and DMS API callbacks.

- **AWS WAFv2 Web ACL:** Directly attached to the ALB with 3 rules (OWASP Core, SQLi, and IP Rate Limiting) to block bad actors at the edge.

2. Application Layer (Private Subnets)

- **Subnets:** 10.0.10.0/24 (AZ ap-southeast-5a) and 10.0.11.0/24 (AZ ap-southeast-5b).
- **Description:** Holds business and compute logic. Instances have no public IP addresses and cannot be accessed directly from the internet.
- **Resources:**
 - **Auto Scaling Group (ASG) EC2 Instances:** Hosts the application code (Nginx web service). Features **t4g.xlarge** or **m6g.xlarge** EC2 Instances (ARM Graviton, 4 vCPU, 16GB RAM each) equipped with **gp3 EBS Root Volumes**. They handle Server 02 (Backend + DMS + MCP) and Server 03 (RAGFlow + LangFuse) workloads securely on hardened Ubuntu 26.04 LTS.
 - **Standalone EC2 Instances:** Deploys a dedicated standalone instance next to each ASG application group (Frontend, Backend, and AI Tiers) inside the private subnets. These instances are connected to the exact same shared databases (RDS), caches (Valkey), and storage systems (S3 and EFS) as their respective ASGs, acting as 1:1 replica environments for application staging, testing, ASIMP auditing/hardening, and pre-baking custom AMIs.

3. Database & Caching Layer (Isolated Private Subnets)

- **Subnets:** 10.0.20.0/24 (AZ ap-southeast-5a) and 10.0.21.0/24 (AZ ap-southeast-5b).
- **Description:** Dedicated to database servers and caching nodes. Deeply isolated without any outbound route to the internet or NAT gateways, minimizing any data extraction surface.
- **Resources:**
 - **Multi-AZ RDS PostgreSQL Instance:** Runs synchronously across multiple availability zones using **Multi-AZ db.m6g.xlarge PostgreSQL (4 vCPU, 16GB RAM)** with **gp3 Storage**, corresponding to Server 04's resource needs. This guarantees high-availability, automatic failover, and robust production database performance.
 - **Amazon ElastiCache for Valkey Cluster:** A secure, high-performance in-memory caching cluster running Valkey 7.2. It manages session state, caches database queries, and runs the Celery/Redis task queue for RAGFlow and LangFuse. It operates on port 6379 with transit and at-rest encryption enabled, allowing ingress only from private compute security groups.

4. Storage Tier (Amazon S3)

- **Description:** Statically hosted media, secure user uploads, and build backups. It is situated outside the VPC but fully integrated.
- **Resources:**
 - **Amazon S3 Bucket:** Fully encrypted standard object storage. Access is managed via IAM policies and secure credentials.

Routing Configuration

The architecture manages network traffic flow through three distinct route tables:

Public Route Table

- Associated with public subnets.
- Routes all outbound traffic (0.0.0.0/0) to the **Internet Gateway (IGW)**.

Private Application Route Table

- Associated with private application subnets.
- Routes all outbound traffic (0.0.0.0/0) to the **NAT Gateway** running in the public subnet.

Database Route Table

- Associated with private database subnets.
- Contains only local VPC route entries (10.0.0.0/16), ensuring database and cache traffic never traverses public routes or internet gateways.

AWS Phased Adoption Roadmap & Costing Guide

#aws #cloud #architecture #costing #roadmap #lifecycle #compliance

AWS Phased Adoption Roadmap & Costing Guide

This guide outlines the step-by-step AWS cloud adoption roadmap, designed to align with our project timeline spanning **2 years of development and 12 months of active support and maintenance (36 months total)**. Over this 36-month period, the **total overall AWS infrastructure cost** is projected at **\$47,709.78 USD (≈ RM 214,694.01 MYR)**. To provide a comprehensive financial blueprint for the entire lifecycle, we estimate **other non-AWS operational costs**—comprising software subscriptions, software licensing compliance, yearly Security Posture Assessments (SPA), load testing (first year only), and 3rd party on-call support for AWS—to be **\$44,800.00 USD (≈ RM 201,600.00 MYR)**. This results in a combined 36-month project operating budget of **\$92,509.78 USD (≈ RM 416,294.01 MYR)**.

Combined 36-Month Financial Projection (AWS vs. Non-AWS)

The table below provides a year-to-year financial breakdown comparing our cloud infrastructure expenditures directly against estimated non-AWS operating overheads:

Financial Period	AWS Cloud Cost (USD)	AWS Cloud Cost (MYR)	Estimated Non-AWS Cost (USD)	Estimated Non-AWS Cost (MYR)	Total Combined Cost (USD)	Total Combined Cost (MYR)
Year 1 (Months 1–12)	\$13,023.16	RM 58,604.22	\$16,600.00	RM 74,700.00	\$29,623.16	RM 133,304.22
Year 2 (Months 13–24)	\$23,343.02	RM 105,043.59	\$14,100.00	RM 63,450.00	\$37,443.02	RM 168,493.59
Year 3 (Months 25–36)	\$11,343.60	RM 51,046.20	\$14,100.00	RM 63,450.00	\$25,443.60	RM 114,496.20
Grand Total	\$47,709.78	RM 214,694.01	\$44,800.00	RM 201,600.00	\$92,509.78	RM 416,294.01

Note: All local currency conversions are calculated at an estimated exchange rate of \$1.00 = MYR 4.50.

Non-AWS Operational Costs Breakdown

The estimated other non-AWS operational costs cover the following strategic categories:

- Software Subscriptions (\$200.00 USD/mo):** \$2,400.00 USD per year (≈ RM 10,800.00 MYR/year) for development collaboration and utility software.
- Software Licensing Compliance:** \$1,200.00 USD per year (≈ RM 5,400.00 MYR/year) for security scans and automated compliance auditing.
- Security Posture Assessment (SPA):** \$4,500.00 USD per year (≈ RM 20,250.00 MYR/year) for professional penetration testing and security audits.
- Load Testing (First Year):** \$2,500.00 USD (≈ RM 11,250.00 MYR) in Year 1 to validate multi-tier system scalability under high concurrency.
- 3rd Party On-Call AWS Support (\$500.00 USD/mo):** \$6,000.00 USD per year (≈ RM 27,000.00 MYR/year) for active SLA and on-call engineering coverage.

This roadmap illustrates how our AWS infrastructure scales incrementally week-by-week and month-by-month, starting from a absolute minimum cost, single-Availability Zone developer sandbox footprint and evolving into a fully-redundant, highly-available, secure, and compliant multi-AZ enterprise production architecture in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)**.

Gantt Chart Activity Mapping

Confidentiality Declaration: All corporate mapping notes, proprietary organizational identifiers, and project-specific planning references have been fully replaced with clean, generic equivalents. All names, codes, and corporate markers are completely anonymized to preserve strict confidentiality.

Below is the definitive, anonymized mapping of the project activities to our generalised enterprise delivery lifecycle terms:

Activity	Generalised Project Scope Mapping
Project Kick-off Meeting	Administrative launch and project alignment
Project Management	Continuous delivery, monitoring, and compliance
Phase 1: AI Chatbot Engine Development	Base chat routing and Core AI engine implementation
• Requirements Analysis & Study	Specifying API payloads and database schema
• Architecture Design & Development	Backend coding and vector database layout design
• Supply, Installation & Configuration	Provisioning sandbox environment & base infrastructure
• Integration with Internal Systems	Bridging enterprise core APIs to the AI engine
• System & Integration Testing	End-to-end payload audits and validation
• Live Deployment & Go-Live	Provisioning highly-available public-facing load balancers
• Technical & Administrator Training	Handover of deployment runbooks to operations teams
• Managed Cloud Services & Operations	Ongoing cloud optimization, monitoring, and scaling
Phase 2: Omnichannel Messaging Integration	WhatsApp API integration & chat webhook flow
• Supply, Installation & Configuration	Provisioning API webhooks and secure NAT gateways
• Integration with Core AI Chatbot	Connecting message handlers to the Phase 1 engine
• System & Integration Testing	Auditing high-throughput messaging workloads
• Live Deployment & Go-Live	Enabling production webhook receiving channels
• User Interaction & Telemetry Monitoring	Tracing chat payloads via Langfuse observability
• Technical & Administrator Training	Operations team training on API token rotations
Phase 3: Customer Relationship Management (CRM)	Customer database, task tracking & ticket management
• Requirements Analysis & Study	CRM database schema design and integration scoping
• Architecture Design & Development	Backend UI views and multi-tenant schema coding

Activity

- Supply, Installation & Configuration
- Integration with Internal Systems
- System & Integration Testing
- Live Deployment, Launch & Go-Live
- Technical & Administrator Training
- Managed Cloud Services & Operations

Phase 4: Omnichannel Mobile Application

- Requirements Analysis & Study
- Supply, Installation & Configuration
- Architecture Design & Development
- Integration with Internal Systems
- System & Integration Testing
- Live Deployment & Go-Live
- Security Vulnerability & Pen Testing
- Technical & Administrator Training
- Managed Cloud Services & Operations

Final Acceptance Test (FAT)

Official Unified Project Go-Live

Documentation & Operational Handover


Support, Maintenance & OpEx Optimisation

Generalised Project Scope Mapping

- Scaling private compute ASG and Multi-AZ RDS
- Direct linking of CRM fields to internal databases
- Auditing CRM security parameters and transaction locks
- High-performance public launch of the CRM system
- Staff onboarding, user manuals, and training sessions
- CRM SLA monitoring and DB read-replica tuning
- Hybrid React Native mobile app development
- Mobile API endpoint definition and push notification plan
- Launching secure edge caching (CloudFront + S3)
- Mobile coding, UI layouts, and local storage tuning
- Bridging push services, mobile SSO, and data syncing
- Beta testing, user acceptance, and load verification
- App store submission and high-throughput production API launch
- Rigorous penetration tests, OWASP checks, and WAF tuning
- Administrative portal onboarding
- Global mobile API latency optimization
- Formal client verification of all consolidated tiers
- Hard cutover to enterprise-scale production
- Delivery of final system runbooks and OpenTofu IaC
- 12-month post-launch support and Savings Plan lock-in

High-Level AWS Phased Growth Path

The architecture matures through four logical environments matching the development timeline:



4-Phase Dynamic Maturity Diagram

The following interactive timeline illustrates the progression of system maturity across our four primary engineering dimensions over the 36-month lifecycle:

```
gantt
    title 4-Phase Infrastructure Maturity Roadmap
    dateFormat X
    axisFormat %d

    section PHASE 1: Sandbox (M1-2)
    Networking (Single-AZ, Public/Private Subnets) :active, p1_net, 0, 8
    Compute (t4g.medium Dev, t4g.micro Bastion) :active, p1_comp, 0, 8
    Database (Single-AZ RDS, cache.t4g.micro Valkey) :active, p1_db, 0, 8
    Security (IAM Policies, Billing Alarms) :active, p1_sec, 0, 8

    section PHASE 2: Staging (M3-6)
    Networking (Dual-AZ Expansion, 2x NAT GWs, ALB) :active, p2_net, 8, 26
    Compute (2x c7g.xlarge ASG, SSM Managed) :active, p2_comp, 8, 26
    Database (RDS Multi-AZ db.t4g.large, EFS Mount) :active, p2_db, 8, 26
    Security (Secrets Manager, S3 Private Endpoints) :active, p2_sec, 8, 26

    section PHASE 3: Production (M7-14)
    Networking (Route 53 DNS Failover, ALB Rules) :active, p3_net, 26, 60
    Compute (ASG CRM Profiles, g5.xlarge GPU OCR) :active, p3_comp, 26, 60
    Database (Multi-AZ db.m7g.xlarge, Read Replicas) :active, p3_db, 26, 60
    Security (AWS WAFv2, AWS DRS continuous replication) :active, p3_sec, 26, 60

    section PHASE 4: Optimised Run-rate (M15-36)
    Networking (CloudFront global CDN, Edge protection) :active, p4_net, 60, 156
    Compute (Edge delivery, Non-prod scheduler 64% off) :active, p4_comp, 60, 156
    Database (Savings Plans, RDS/Valkey Reserved Nodes) :active, p4_db, 60, 156
    Security (Standalone Wazuh SIEM, CIS Hardening) :active, p4_sec, 60, 156
```

Detailed Multi-Tier Maturity Progression Matrix

To facilitate thorough technical review, the matrix below details the incremental evolutionary steps of each distinct physical and virtual tier:

4-PHASE MULTI-TIER INFRASTRUCTURE MATURITY ROADMAP				
Tier	PHASE 1: Dev/POC Sandbox (Months 1–2 / Weeks 1–8)	PHASE 2: Dual-AZ Staging (Months 3–6 / Weeks 9–26)	PHASE 3: Enterprise Prod (Months 7–14 / Weeks 27–60)	PHASE 4: Run (M15–36 / W61)
NETWORKING	- Single-AZ VPC 1x Public & 1x Private - Basic Security Groups	- Dual-AZ VPC expansion 2x NAT Gateways (Redundant) - Application Load Balancer	- Route 53 Active-Passive DNS - ALB Path-routing (/crm/*) - Strict SG Port isolation	- CloudFront - Global CDN - Edge WAF
COMPUTE	1x Standalone EC2 (t4g.medium on hardened Ubuntu 26.04 OS) - t4g.micro SSH Jump host	- Auto Scaling Group (ASG) (Min 2x c7g.xlarge nodes) - Passwordless SSM access (Zero public SSH ports)	- Multi-target JVM Profiles in private ASG compute - Standalone g5.xlarge GPU (NVIDIA A10G) for OCR tasks	Instance Scheduler (64% staging off-hours)
DATABASE & CACHE	- Single-AZ RDS PostgreSQL (db.t4g.medium sandbox) - ElastiCache Valkey (cache.t4g.micro node)	- RDS Multi-AZ upgrade (30-day credential rot.) - Amazon EFS Shared Storage (Synchronous media/cache)	- Production db.m7g.xlarge - Dedicated Read-Replicas - ElastiCache Valkey (cache.r7g.large cluster)	1-Yr RI - Valkey Reserved Nodes lock-in
SECURITY & COMPLIANCE	- AWS IAM Identity Center (Granular role RBAC) - AWS Budgets & alerts (set at \$150 threshold)	- AWS Secrets Manager - Private Gateway Endpoints (bypasses S3 NAT egress)	- AWS WAFv2 Layer-7 filter - AWS KMS customer keys (CMK) - AWS Elastic DR (AWS DRS) asynchronous replication	- Standalone Wazuh SIEM - CIS Level 2 Hardening

Detailed Chronological Roadmap

Year 1, Month 1 (Weeks 1–4): Project Initiation & Kick-off

- Gantt Chart Activity Mapping:** Project Kick-off Meeting, Project Management.
- AWS Architectural Focus:** Identity access scoping and network design.
- AWS Services Provisioned:**
 - AWS IAM Identity Center:** Enforces multi-factor authentication (MFA) and granular role-based access control (RBAC).
 - AWS Budgets & Billing Alarms:** Confirmed with a baseline budget threshold of \$150.00 USD to prevent accidental run-away charges.
 - Amazon Route 53:** Public hosted zone established for early API domain mapping.
- Week-by-Week Technical Goals:**
 - Week 1:** Establish enterprise AWS root account structure, configure IAM administrators, and set budget guardrails.
 - Week 2:** Draft security policies and Cyberjaya office CIDR access blocks for administrators.
 - Week 3:** Initial provision of an empty modular VPC structure inside the Malaysia (ap-southeast-5) region utilizing OpenTofu.
 - Week 4:** Set up shared project logging bucket with strict life-cycle rules inside Amazon S3.
- Phase Cost (USD / MYR):** \$1.00 USD / RM 4.50 MYR (Route 53 hosted zone base fee + nominal IAM storage).

Year 1, Month 2 (Weeks 5–8): Phase 1 Requirements & Dev Setup

- Gantt Chart Activity Mapping:** Phase 1 (AI Chatbot Engine Development) - Requirements Study, Supply, Installation & Configuration.
- AWS Architectural Focus:** Setting up the single-AZ developer sandbox to support early AI Chatbot engine coding and local model verification.
- AWS Services Provisioned:**
 - Amazon VPC:** Single public subnet + single private subnet inside ap-southeast-5a.
 - Secure SSH Jump host (Bastion):** Sized at t4g.micro to allow the Cyberjaya developer team to tunnel safely into the private subnet.
 - Standalone Developer Compute Instance:** 1x t4g.medium instance running hardened Ubuntu 26.04 LTS (ASIMP framework) to compile and test the Java/Spring Boot AI Chatbot engine.
 - Amazon RDS PostgreSQL (Single-AZ):** 1x db.t4g.medium PostgreSQL 16 database instance for chat history storage.
 - Amazon ElastiCache Valkey:** 1x cache.t4g.micro node for caching user sessions.
- Chronological Alignment & Strategic Causality:**
 - Strategic Alignment:** Establishing the database and caching engines at this sandbox gate aligns directly with the [RDS PostgreSQL vs. Percona Server Guide](#).
 - Business Causality:** Choosing AWS-native RDS PostgreSQL 16 (db.t4g.medium at ~\$32.85/mo) and ElastiCache Valkey (cache.t4g.micro at ~\$9.34/mo) instead of self-hosted Percona and custom Redis instances on EC2 eliminates upfront management overhead. This allows our developers to focus immediately on core Spring Boot chatbot endpoints without spending precious hours configuring Patroni replication or PgBouncer connection pools.
- Week-by-Week Technical Goals:**
 - Week 5:** Finalise AI chatbot payload specifications and backend Spring Boot schema designs.
 - Week 6:** Launch the SSH Jump host and restrict port 22 access strictly to the Cyberjaya office IP address.
 - Week 7:** Provision the standalone compute and Valkey nodes, ensuring the security groups block all traffic not originating from the Jump host.
 - Week 8:** Deploy Flyway on the database to verify schema migrations.
- Phase Cost (USD / MYR):** \$138.50 USD / RM 623.25 MYR (Single-AZ baseline environment).

Year 1, Months 3–6 (Weeks 9–26): Chatbot Core Coding & Integration

- Gantt Chart Activity Mapping:** Phase 1 (AI Chatbot Engine Development) - Architecture Design & Development, Integration with Internal Systems.
- AWS Architectural Focus:** Transitioning the sandbox environment to a highly-resilient Dual-AZ Staging Environment to validate auto-scaling, load handling, and database failovers.
- AWS Services Provisioned:**
 - Dual-AZ Networking:** Expanded VPC with private subnets across 2 Availability Zones, protected by 2x NAT Gateways to ensure secure outbound API updates.
 - Application Load Balancer (ALB):** Public-facing ALB deployed in the public subnets to distribute mock testing traffic.
 - Auto Scaling Group (ASG):** Sized with a minimum capacity of 2x c7g.xlarge (ARM Graviton3) instances to support heavy processing.
 - Amazon EFS (Elastic File System):** Shared persistent NFS mount to share AI model cache files across the ASG compute tier.
 - Amazon RDS (Multi-AZ Upgrade):** Database upgraded to Multi-AZ db.t4g.large with synchronous replication to prevent a single point of failure (SPOF).
 - AWS Secrets Manager:** Deployed to secure database credentials, rotating them automatically every 30 days.

- **Chronological Alignment & Strategic Causality:**
 - **Strategic Alignment:** Integrating with core internal legacy APIs aligns directly with Low-Cost Option 1 (API-Based Connections) in the [Hybrid Cloud Integration Guide](#).
 - **Business Causality:** Connecting the Chatbot to on-premises business logic over secure HTTPS REST APIs using the NAT Gateway's Elastic IPs avoids the massive, slow provisioning times and high capital expenses (~\$522/mo port and partner loops) of AWS Direct Connect. Utilizing redundant NAT Gateways ensures continuous outbound integration streams during this critical staging phase.
 - **Week-by-Week Technical Goals:**
 - **Week 9–12:** Write Spring Boot Chatbot logic and vector indexing algorithms.
 - **Week 13–16:** Integrate the application tier with AWS Systems Manager (SSM) and completely disable traditional passwords or direct SSH.
 - **Week 17–20:** Implement auto-scaling policies triggered when average CPU usage exceeds 70% across the ASG.
 - **Week 21–24:** Establish private VPC endpoints to route Amazon S3 media traffic locally, bypassing NAT Gateway charges.
 - **Week 25–26:** Conduct chaos engineering tests, simulating a full Availability Zone outage to verify sub-minute automated DB and ALB failover.
 - **Phase Cost (USD / MYR): \$668.11 USD / RM 3,006.50 MYR (Dual-AZ Staging Environment).**
-

Year 1, Months 7–8 (Weeks 27–34): Chatbot Go-Live & WhatsApp Integration

- **Gantt Chart Activity Mapping:** Phase 1 (AI Chatbot Engine Development) - Go-Live, Training, Managed Cloud Services; Phase 2 (Omnichannel Messaging Integration) - Supply, Installation & Configuration, Integration with Core AI Chatbot, System & Integration Testing, Go-Live, User Interaction & Telemetry Monitoring.
 - **AWS Architectural Focus:** Shifting to the High-Performance Production Plan to launch the Chatbot engine publicly and integrate the high-throughput WhatsApp Business webhook pipelines.
 - **AWS Services Provisioned:**
 - **AWS WAFv2:** Associated with the ALB to intercept and filter out OWASP Top-10 attacks, SQL injection attempts, and enforce dynamic IP rate-limiting (e.g., maximum 100 requests per 5 minutes per client IP).
 - **Production Database Tier:** Sized to Multi-AZ db.m7g.xlarge (4 vCPU, 16GB RAM) with 250 GB of gp3 storage (3,000 IOPS) to absorb massive concurrent chat transactions.
 - **Production Cache Tier:** Scaled up to ElastiCache Valkey cache.r7g.large (Dual Node cluster) to act as a secure, fast query broker.
 - **Amazon API Gateway & AWS Lambda:** Serverless webhook receiver layer deployed to securely absorb massive, bursty incoming WhatsApp callback payloads from Meta, offloading them to **Amazon SQS** queues to prevent thread saturation in the Spring Boot backend.
 - **Amazon CloudWatch Application Signals:** Enabled to track end-to-end user latency and alert on HTTP 5XX spikes.
 - **Chronological Alignment & Strategic Causality:**
 - **Strategic Alignment:** Establishing AI capability pipelines aligns with Low-Cost Option 2 (AI-Native MCP-Based Connections) in the [Hybrid Cloud Integration Guide](#) and our hybrid connection patterns in the [RAGFlow + Langfuse GPU Guide](#).
 - **Business Causality:** Introducing the Model Context Protocol (MCP) proxy via API Gateway at this gate allows the chatbot to securely query local, on-premises data catalogs in real time without passing raw citizen credentials over the public internet. Utilizing serverless webhooks (API Gateway + Lambda) ensures that high-volume WhatsApp callback storms from Meta are offloaded cleanly into SQS, completely shielding the core Java application from thread starvation during the public launch.
 - **Week-by-Week Technical Goals:**
 - **Week 27:** Establish production ALB SSL certificates using AWS Certificate Manager (ACM).
 - **Week 28:** Deploy AWS WAF and configure the core rule-sets. Run final technical training for administrator teams on emergency scaling.
 - **Week 29:** Launch the serverless webhook API Gateway and AWS Lambda functions, writing integration metrics to CloudWatch.
 - **Week 30:** Connect the WhatsApp endpoint to AWS End User Messaging (Social Channels), removing high-cost external SaaS middlemen.
 - **Week 31–32:** Execute massive simulated load testing, driving over 5,000 concurrent user interaction flows.
 - **Week 33–34:** Go-live with the WhatsApp Omnichannel API, monitoring trace logs in real-time.
 - **Phase Cost (USD / MYR): \$1,665.61 USD / RM 7,495.25 MYR (High-Performance Multi-AZ Enterprise Production).**
-

Year 1, Months 9–12 (Weeks 35–52): Phase 2 Managed Ops & Phase 3 CRM Inception

- **Gantt Chart Activity Mapping:** Phase 2 (Omnichannel Messaging Integration) - Technical & Administrator Training; Phase 3 (CRM Development) - Requirements Analysis & Study, Architecture Design & Development, Supply, Installation & Configuration.
 - **AWS Architectural Focus:** Managing WhatsApp operations while establishing CRM development and isolated schemas on shared high-performance infrastructure to reduce overhead.
 - **AWS Services Provisioned:**
 - **Isolated CRM DB Schema:** Configured inside the production RDS instance utilizing Flyway migrations, fully separating CRM tables with strict database roles.
 - **Additional ASG App Targets:** Launch template modified to run specialized CRM JVM profiles on the existing production ASG instances.
 - **Continuous Telemetry Tracking:** Enabled Amazon GuardDuty to monitor account-level DNS abuse and secure IAM roles.
 - **Chronological Alignment & Strategic Causality:**
 - **Strategic Alignment:** Preparing backend storage and starting CRM schemas aligns directly with Strategy A (Backup and Restore) in our [Disaster Recovery Guide](#).
 - **Business Causality:** Incepting CRM database schemas requires early security precautions. Activating daily automated backups via AWS Backup (~\$7.30/mo) at this gate ensures that early client data models and test customer records are securely captured. This provides a low-cost, compliant backup baseline before we introduce continuous near-real-time replication engines.
 - **Week-by-Week Technical Goals:**
 - **Week 35–38:** Deliver comprehensive operational training on Meta API token rotation.
 - **Week 39–42:** Code backend CRM service APIs and construct robust relational mappings.
 - **Week 43–48:** Deploy database read-replicas inside the isolated database subnets to offload complex CRM query lookups from the write master node.
 - **Week 49–52:** Configure automated nightly snapshots using AWS Backup with cross-region replication.
 - **Phase Cost (USD / MYR): \$1,685.00 USD / RM 7,582.50 MYR (Production run-rate with nominal storage growth).**
-

Year 2, Months 1–2 (Weeks 53–60): CRM Integration & Go-Live

- **Gantt Chart Activity Mapping:** Phase 3 (CRM Development) - Integration with Internal Systems, System & Integration Testing, Go-Live & Launch, Technical & Administrator Training, Managed Cloud Services & Operations.
- **AWS Architectural Focus:** Securing database performance and ensuring seamless system scaling during the public launch of the CRM system.
- **AWS Services Provisioned:**
 - **AWS Elastic Disaster Recovery (AWS DRS):** Configured as a continuous block-level continuous replication strategy from our on-premises database networks to a secure staging subnet in AWS.
 - **Elastic Load Balancer Routing Rules:** Configured listener paths to route /crm/* traffic to a dedicated CRM target group in the ASG.
 - **Amazon Route 53 Health Checks:** Automated DNS failover triggers routing traffic to emergency backup static endpoints.
- **Chronological Alignment & Strategic Causality:**
 - **Strategic Alignment:** Launching the CRM system requires switching from cold backups to high-frequency continuous protection, aligning with Strategy E (AWS Elastic Disaster Recovery) in the [Disaster Recovery Guide](#).

- **Business Causality:** Weeks 53–60 (Year 2, CRM Go-Live): Activate AWS Elastic Disaster Recovery (AWS DRS) at ≈ \$112/mo. Causality: Ingesting proprietary customer CRM data increases data criticality, justifying the DR insurance spend to maintain RPO/RTO SLAs.
 - **Week-by-Week Technical Goals:**
 - **Week 53–54:** Complete end-to-end integration mapping between CRM fields and internal legacy systems.
 - **Week 55–56:** Run exhaustive database transaction lock verification under a peak stress model.
 - **Week 57–58:** Launch the public CRM portal. Monitor database CPU and Valkey memory metrics closely during the initial cutover.
 - **Week 59–60:** Complete on-boarding training for internal business teams and hand over the final system guides.
 - **Phase Cost (USD / MYR): \$1,720.00 USD / RM 7,740.00 MYR** (Incorporating active AWS DRS replication streams).
-

Year 2, Months 3–8 (Weeks 61–86): Super Mobile App & High-Performance AI Ingress

- **Gantt Chart Activity Mapping:** Phase 4 (Omnichannel Mobile Application) - Requirements Analysis & Study, Supply, Installation & Configuration, Architecture Design & Development, Integration with Internal Systems, System & Integration Testing, Live Deployment & Go-Live.
 - **AWS Architectural Focus:** Introducing heavy-duty AI processing and secure mobile delivery channels to support the Super Mobile App launch.
 - **AWS Services Provisioned:**
 - **Amazon CloudFront:** Global Content Delivery Network (CDN) with edge caching to speed up mobile static asset and image retrieval.
 - **AWS WAFv2 on CloudFront:** Move WAF protection to the global CDN edge to block malicious requests before they even reach the ALB.
 - **High-Performance GPU Compute Tier:** 1x Standalone Staging g5.xlarge instance (NVIDIA A10G GPU with 24GB VRAM) to execute specialized DeepDoc visual layouts and high-speed OCR.
 - **Chronological Alignment & Strategic Causality:**
 - **Strategic Alignment:** Introducing heavy-duty GPU-powered text parsing and OCR pipelines aligns directly with the architectural patterns in our [RAGFlow + Langfuse GPU Guide](#).
 - **Business Causality:** Launching the Super Mobile App drives massive concurrent document/image uploads. Provisioning 1x Standalone GPU Staging instance (g5.xlarge at ~\$1,215.71/mo during active development periods) at this gate is required to handle computationally heavy DeepDoc visual layout analysis and high-speed optical character recognition (OCR), ensuring a responsive, zero-lag user interface during the mobile rollout.
 - **Week-by-Week Technical Goals:**
 - **Week 61–68:** Code hybrid mobile screens in React Native and define the API contract.
 - **Week 69–74:** Provision S3 upload directories and attach CloudFront to bypass API endpoint processing.
 - **Week 75–80:** Connect mobile push notifications through Amazon Simple Notification Service (SNS).
 - **Week 81–84:** Integrate on-device OCR components with our backend GPU compute.
 - **Week 85–86:** Submit the compiled hybrid application to the Apple App Store and Google Play Store, and cutover the production API endpoints.
 - **Phase Cost (USD / MYR): \$2,150.00 USD / RM 9,675.00 MYR** (Peak pricing incorporating heavy GPU compute nodes and CDN data transfers).
-

Year 2, Month 9 (Weeks 87–90): Mobile App Security Hardening

- **Gantt Chart Activity Mapping:** Phase 4 (Omnichannel Mobile Application) - Security Vulnerability & Pen Testing, Technical & Administrator Training, Managed Cloud Services & Operations.
 - **AWS Architectural Focus:** Running deep audits on mobile APIs and locking down security variables before finalizing the development phase.
 - **AWS Services Provisioned:**
 - **Standalone Wazuh SIEM Instance:** Sized at t4g.large inside the secure management subnet, continuously gathering and auditing Linux logs from all compute nodes.
 - **Amazon Inspector:** Automated security scanning of EC2 instances and container registries.
 - **Chronological Alignment & Strategic Causality:**
 - **Strategic Alignment:** Establishing a centralized, secure threat monitoring baseline at this hardening gate aligns directly with the [Wazuh Standalone Installation Guide](#).
 - **Business Causality:** Preparing for final regulatory sign-off and penetration testing mandates strict operational auditing. Deploying a self-hosted Standalone Wazuh SIEM instance on a cost-effective ARM64 instance (t4g.large at ~\$65.71/mo) provides comprehensive log aggregation and intrusion detection without high-premium SaaS subscription pricing, securing our infrastructure at a 57% savings.
 - **Week-by-Week Technical Goals:**
 - **Week 87:** Trigger a deep penetration test across all public-facing ALB endpoints.
 - **Week 88:** Apply ASIMP Ansible playbooks to ensure 100% compliance with CIS Level 2 benchmarks.
 - **Week 89:** Finalise training for administrative teams on detecting API credential exposure.
 - **Week 90:** Adjust WAF geo-blocking rules to allow traffic exclusively from approved Southeast Asian networks, blocking foreign botnets.
 - **Phase Cost (USD / MYR): \$2,215.71 USD / RM 9,970.70 MYR** (Peak development footprint with active Wazuh security auditing).
-

Year 2, Months 10–12 (Weeks 91–104): Final Acceptance Test, Project Go-Live & Handover

- **Gantt Chart Activity Mapping:** Final Acceptance Test, Official Unified Project Go-Live, Documentation & Operational Handover, Support & Maintenance Period.
 - **AWS Architectural Focus:** Tearing down auxiliary staging environments, cleaning up temporary EBS volumes, and scaling the core Multi-AZ production cluster to its final stable state.
 - **AWS Services Provisioned:**
 - **Consolidated Enterprise Production Stack:** Streamlined, secure Multi-AZ configuration.
 - **Amazon CloudWatch Custom Dashboard:** Real-time single-pane visibility for operations teams.
 - **Week-by-Week Technical Goals:**
 - **Week 91–94:** Complete formal Client Final Acceptance Testing across all integrated tiers (Chatbot, WhatsApp, CRM, Mobile).
 - **Week 95–98:** Complete the hard cutover of production databases. Deliver final OpenTofu infrastructure codebase files.
 - **Week 99–104:** Establish the 12-month Support & Maintenance framework. Configure Automated Instance Scheduler routines to shut down developer environments outside business hours, saving 64% on idle staging compute!
 - **Phase Cost (USD / MYR): \$1,285.80 USD / RM 5,786.10 MYR** (Optimised, high-availability baseline run-rate).
-

Year 3: Support, Maintenance & Cost Optimisation (Months 25–36)

- **Gantt Chart Activity Mapping:** Support & Maintenance Period - 12 months.
- **AWS Architectural Focus:** Applying Day-2 financial optimizations to lock in significant savings for the remaining support year.
- **AWS Services Provisioned:**
 - **AWS Savings Plans:** Apply a 1-Year Compute Savings Plan to EC2 instances, cutting compute costs by **34%**.
 - **RDS Reserved Instances:** Secure a 1-Year Reserved Instance for Multi-AZ PostgreSQL, cutting database costs by **30%**.
 - **ElastiCache Reserved Nodes:** Secure a 1-Year Reserved Node for ElastiCache Valkey, saving **35%** on caching.
- **Month-by-Month Technical Goals:**
 - **Months 25–28:** Purchase the Reserved Instances and monitor the billing console to confirm discounts are successfully applied.
 - **Months 29–32:** Run monthly database indexing optimizations to keep storage footprints below 250 GB.

- **Months 33–36:** Maintain high uptime (>99.99%) and prepare the final cloud transfer of ownership documents.
- **Phase Cost (USD / MYR): \$945.30 USD / RM 4,253.85 MYR** (Optimised enterprise production run-rate representing **43.2% savings** over standard on-demand pricing).

Comparative AWS Monthly Cost Matrix

The table below breaks down the monthly run-rate (in USD and MYR) across each distinct development milestone phase, showing the incremental growth and financial efficiency of our roadmap:

Milestone Phase	Est. Cost (USD)	Est. Cost (MYR)	Primary Architectural Changes & Sizing
Month 1 (Kick-off)	\$1.00	RM 4.50	Initial DNS public zone mapping.
Month 2 (Dev Setup)	\$138.50	RM 623.25	Single-AZ sandbox: t4g.medium compute, db.t4g.medium PostgreSQL database, Valkey cache.
Months 3 - 6 (Staging)	\$668.11	RM 3,006.50	Dual-AZ staging: NAT Gateways, ALB routing, 2x c7g.xlarge ASG instances, Multi-AZ RDS.
Months 7 - 8 (Chatbot Go-Live)	\$1,665.61	RM 7,495.25	High-Performance Prod: AWS WAFv2 layer, db.m7g.xlarge Multi-AZ RDS, Lambda webhooks.
Months 9 - 14 (CRM Rollout)	\$1,720.00	RM 7,740.00	Adding CRM schemas, RDS read-replicas, and continuous AWS DRS on-premises replication.
Months 15 - 20 (Mobile Launch)	\$2,150.00	RM 9,675.00	Adding Amazon CloudFront global CDN edge and dedicated GPU-backed g5.xlarge AI parsing nodes.
Months 21 - 22 (Security Pen)	\$2,215.71	RM 9,970.70	Launching standalone Wazuh SIEM auditor instance and executing automated pen-testing suites.
Months 23 - 24 (Project Go-Live)	\$1,285.80	RM 5,786.10	Streamlining production: shutting down staging instances outside business hours.
Year 3 (Support & Maintenance)	\$945.30	RM 4,253.85	Day-2 financial optimization: lock-in 1-Year Savings Plans for 43%+ overall OpEx discount.

Strategic Recommendations for Phase-by-Phase Rollout

- Leverage “Scale-to-Zero” Webhooks Immediately:** In Phase 2, deploy API Gateway and AWS Lambda webhook integrations right from the start. This keeps initial development costs near **\$0.00** while ensuring the backend is ready to handle high traffic spikes.
- Standardise on Graviton Compute:** Standardise all application EC2 nodes and RDS instances on AWS Graviton (ARM64). This delivers up to **40% better price-performance** compared to x86 equivalents, which directly lowers the run-rate of the Staging and Production phases.
- Establish Free PrivateLink S3 Endpoints early:** Create VPC Gateway Endpoints for S3 during Year 1, Month 2. Since the AI Chatbot and CRM download substantial volumes of document/media objects from S3, routing this traffic locally within the VPC avoids NAT Gateway data processing fees (\$0.045/GB).
- Automate Non-Production Environments:** Enforce a strict non-production schedule starting from Year 1, Month 3. Automatically stopping Staging and Dev environments outside standard Cyberjaya business hours (12 hours/day, 5 days/week) instantly reduces non-production compute costs by **64%**.

Sovereign and Regulatory Pathway Alignment (Malaysian PDPA 2010)

Keeping data secure and compliant with local regulations is a core priority of this phased roadmap:

- **In-Region Sovereign Processing:** In Phase 1 and Phase 2, all AI processing and messaging pathways are securely locked inside the AWS Malaysia region (ap-southeast-5). Sensitive citizen communications and metadata never leave the physical boundaries of the Kuala Lumpur data centers.
- **Cryptographic Isolation:** From Phase 3 onwards, all database layers and CloudWatch log directories are encrypted at-rest using customer-managed keys via **AWS KMS**.
- **AWS DRS Secure Replication Compliance:** The continuous replication channels provisioned in Year 2, Month 1 using AWS Elastic Disaster Recovery utilize highly-secure, encrypted TLS 1.3 tunnels. This allows companies to align with the **Malaysian Personal Data Protection Act (PDPA) 2010 Section 129** cross-border restrictions, ensuring backup target paths are fully documented, audited, and compliant.

Developer Design Mapping

#aws #cloud #architecture #vpc #alb #asg #rds #waf #dns #ssl #efs #postgresql #ragflow #langfuse #compliance

Developer Design Alignment Guide

This guide details how we transition the **Developer’s First Design** (which specified four separate, standalone Ubuntu 26.04 LTS servers) into our secure, highly-available, production-ready **AWS 3-Tier Architecture**, without changing any underlying AWS constraints or requirements.

Rationale for the Architectural Transition

The developer’s original design was logical but modeled around static, single-node virtual machines. Running critical production infrastructure on standalone VMs presents several operational risks:

- 1. **Single Points of Failure (SPOFs):** If any of the servers experiences hardware degradation or guest OS crashes, the entire application goes offline.
- 2. **Direct Public Exposure:** Placing database or backend servers directly on public-facing networks increases the surface area for brute-force SSH, port-scanning, and SQL injection (SQLi) attacks.
- 3. **Manual Scalability & High Latency:** VM resources are fixed and cannot scale dynamically in response to user demand or computational peaks.
- 4. **Backup and Patches overhead:** Managing backups, OS updates, and package security patches manually across multiple VMs takes significant operational effort.

By aligning this layout with our secure AWS design, we retain **all functionality** of the developer’s original services (Nginx, Backend, DMS, MCP, RAGFlow, LangFuse, and Postgres) while inheriting cloud-native security, performance, and automation.

Detailed Component Transitions

1. Server 01: Frontend Web Tier (Nginx Web Server / Reverse Proxy)

- **Developer’s First Design:** A single Ubuntu server running Nginx (2 vCPU, 4GB RAM) exposed to the public internet via port 80/443.
- **AWS Alignment:**
 - **Layer 7 Firewall (AWS WAFv2):** Blocks common exploits (OWASP Top 10, SQLi, XSS) and manages brute-force or DDoS attempts via dynamic IP Rate Limiting.
 - **Application Load Balancer (ALB):** Terminates external SSL/TLS, manages routing rules, and automatically performs health checks to route traffic only to healthy instances.
 - **Private Nginx Web Tier (ASG):** Sized as `t4g.medium` (2 vCPU, 4GB RAM) but deployed within **private subnets** under an Auto Scaling Group (ASG). This ensures the Nginx instance has *no public IP*, can scale out horizontally, and is protected from brute-force exposure.
 - **Dedicated Standalone Instance (AMI Baking):** Paired with a dedicated Frontend Standalone instance inside the private subnets. Connected directly to identical S3 resources to test routing/static templates and pre-bake certified `ami-frontend-*` images.

2. Server 02: Backend App Tier (Backend + DMS + MCP)

- **Developer’s First Design:** A single Ubuntu server (4 vCPU, 16GB RAM) running Backend, DMS, and MCP APIs.
- **AWS Alignment:**
 - **Zero Direct Ingress:** Deployed inside Private App Subnets, completely blocked from direct public ingress. The ALB only forwards verified requests to this layer on Port 80 (or your custom API port).
 - **Sizing Alignment:** Sized using Graviton `t4g.xlarge` (4 vCPU, 16GB RAM) to perfectly match the developer’s CPU/Memory requirements while saving up to 20% in raw compute costs over comparable `x86_64` nodes.
 - **Secure Egress:** Outbound integrations (such as external DMS syncs, Webhooks, or payment gateway APIs) traverse a managed **AWS NAT Gateway** in the public subnet, masking instance IPs and preventing inbound traffic.
 - **Dedicated Standalone Instance (AMI Baking & Schema Migrations):** Paired with a dedicated Backend Standalone instance connected to the identical Multi-AZ RDS database and S3 buckets. Used to test application updates, safely run database migrations, and pre-bake `ami-backend-*` images under a 1:1 replica environment.

3. Server 03: AI Tier (RAGFlow + LangFuse)

- **Developer’s First Design:** A single Ubuntu server (4 vCPU, 8GB RAM) running RAGFlow + LangFuse AI stack.
- **AWS Alignment:**
 - **Intra-VPC Security:** Placed in the same Private App Subnets. Secure, ultra-low-latency private communication paths are used to connect the Backend (Server 02) to RAGFlow + LangFuse (Server 03) via internal DNS or security-group-protected private endpoints.
 - **Sizing Alignment:** Sized as `c6g.xlarge` or `t4g.xlarge` to provide the required 4 vCPU and 8-16GB RAM. If CPU-bound AI processing becomes intensive, the ASG will scale these instances dynamically.
 - **Environment Compatibility:** Fully compatible with containerized environments (Docker Compose / ECS) or native package runtimes.
 - **Dedicated Standalone Instance (AMI Baking & EFS Caching):** Paired with a dedicated AI Standalone instance connected to the identical Amazon EFS filesystem, RDS, and S3. Developers pre-download and warm up AI model caches on EFS using this standalone instance to allow instant bootstrapping across the auto-scaled nodes. Used to pre-bake `ami-ai-*` images.

4. Server 04: Database Data Tier (SQL Database)

- **Developer’s First Design:** A single Ubuntu server (4 vCPU, 16GB RAM) running self-managed PostgreSQL.
- **AWS Alignment:**
 - **Fully Managed RDS (PostgreSQL 16):** Upgraded to **AWS RDS PostgreSQL** (Multi-AZ deployment) sized at `db.m6g.xlarge` (4 vCPU, 16GB RAM).
 - **Multi-AZ Replication:** Synchronously replicates data across physically separate Availability Zones. In the event of an outage in AZ A, AWS automatically fails over to AZ B with zero manual intervention or data loss.
 - **Absolute Network Isolation:** Deployed inside isolated Private Database Subnets. The database security group restricts incoming traffic *exclusively* to the application security group (ASG), preventing any direct internet access.

Operating System & Hardware Optimizations

To deliver maximum efficiency, we transition the underlying hardware platform from legacy x86 virtual machines to **AWS Graviton (ARM64)** processors:


- 1. **Price-Performance Efficiency:** AWS Graviton (t4g and m6g instances) delivers up to **40% better price-performance** compared to equivalent x86 instances, significantly lowering the monthly run costs.
- 2. **Ubuntu 26.04 LTS Base Operating System:** To leverage the latest performance improvements, security features, and modern container support, we standardize our base platform on **Ubuntu 26.04 LTS (Noble Numbat successor)**.
- 3. **Amazon Linux 2023 Option:** For lightweight workloads that do not depend on Canonical specific packages, **Amazon Linux 2023 (AL2023)** remains available as a minimal, cloud-optimized option.

Server Hardening & Security Compliance (ASIMP Integration)

In aligning the developer design with AWS enterprise standards, all Ubuntu 26.04 LTS compute resources (both ASG instances and Standalone instances) are hardened and tuned using **ASIMP (Ansible System Integrity Management Platform)** (available at github.com/linuxmalaysia/ASIMP).

ASIMP is a host-based, automated security hardening, compliance, and auditing framework that implements a strict **“Measure, Harden, Re-Measure”** paradigm to verify and guarantee security posturing before the machine is allowed to process production traffic.

The ASIMP Hardening & Auditing Pipeline



Key ASIMP Hardening Capabilities Applied:

- 1. **Dual-Engine Security Auditing:**
 - **OpenSCAP:** Conducts formal vulnerability and security compliance scanning mapped against the **CIS Security Ubuntu Linux Benchmark Level 2** profile.
 - **Lynis:** Performs comprehensive system configuration auditing, examining OS parameters, boot configurations, cryptography standards, and active network ports.
- 2. **Pre & Post Scorecard Comparison:**
 - Runs audits prior to hardening to capture initial baselines, then executes them afterwards, logging comparative metrics in `/var/log/asimp-baseline-scores.json`.
 - Generates visually comprehensive, standalone HTML inspection reports at `/var/log/openscap-before-report.html` and `/var/log/openscap-after-report.html`.
- 3. **Automated Package Updates & debsums Verification:**
 - Standardizes the Ubuntu system upgrade procedures and checks package-level code integrity via debsums to detect any unauthorized binary modifications.
- 4. **Standardized OS Hardening Benchmarks:**
 - Deploys rigorous security compliance controls utilizing OpenStack’s ansible-hardening role.
 - Standardizes and restricts SSH configuration endpoints using Dev-Sec’s certified ssh-hardening roles, enforcing secure cipher suites, disabling password-based root access, and specifying key exchange standards.
- 5. **Detailed System Tuning & Custom Fixes:**
 - Automatically tunes virtual memory configuration parameters, core kernel dumps, file system mounting options (e.g., nodev, nosuid, noexec where appropriate), and limits access to system compilers.

Summary of the AWS Security & HA Multipliers

Through this alignment, the developer gets their exact applications deployed with unmatched production capabilities:

Feature	How AWS Improves Developer's First Design
High Availability	Multi-AZ redundancy for Compute and RDS DB.
DDoS & Web Protection	AWS WAFv2 blocking bad traffic before compute.
Elastic Scaling	ASG automatically adds instances based on load.
Data Protection	Automatic, daily snapshots + Multi-AZ backups.
Security Principle	Zero-Trust isolated subnets & IAM roles.
System Integrity	Hardened and audited via ASIMP framework.

Auto Scaling Groups (ASGs) & Separation of Concerns

[#aws](#) [#cloud](#) [#architecture](#) [#vpc](#) [#alb](#) [#asg](#) [#rds](#) [#ssl](#) [#disaster-recovery](#) [#efs](#) [#postgresql](#) [#gpu](#) [#ragflow](#) [#langfuse](#)

Auto Scaling Groups (ASGs) & Separation of Concerns

In a modern cloud-native architecture, managing complex applications requires balancing scalability, security, cost, and operational maintenance. This document explores the architectural rationale for using **Separation of Concerns (SoC)** via distinct Auto Scaling Groups (ASGs) and provides an exhaustive guide on how to handle shared storage (**Amazon S3**, **Amazon EFS**, or both) to enable seamless, stateless autoscaling.

1. Separation of Concerns (SoC) via Dedicated ASGs

Using a single, monolithic fleet of EC2 instances to run multiple, diverse components of your application (e.g., Frontend, Backend, and AI Processing Tiers) creates tight coupling, increases the blast radius of failures, and complicates scaling.

By defining **distinct Auto Scaling Groups** for each logical application tier (e.g., an ASG for Server 01: Nginx/Frontend, an ASG for Server 02: Backend + DMS, and an ASG for Server 03: RAGFlow + AI Processing), you achieve robust Separation of Concerns.

Architectural Advantages of Multi-ASG Layouts

- 1. Independent Elastic Scaling:**
 - **Different Resource Triggers:** Frontend layers scale based on network load or request counts, Backend layers scale based on CPU or memory, and AI layers (RAGFlow/LangFuse) scale based on custom metrics like GPU utilization or task queue length.
 - **Cost Optimization:** You don't need to scale expensive, heavy instances (e.g., Graviton t4g.xlarge with 16GB RAM) just because the simple web server/reverse proxy layer is experiencing high traffic. Only scale the specific layer under load.
- 2. Blast Radius Limitation:**
 - If a memory leak or heavy computation in the AI tier (Server 03) crashes an instance, the Nginx web tier (Server 01) and the core Backend tier (Server 02) remain completely operational, allowing you to return graceful degraded responses to users.
- 3. Least Privilege Security (IAM Roles & Security Groups):**
 - Each ASG is equipped with its own dedicated **IAM Instance Profile**. The Backend ASG can have permission to read from specific database parameter stores, whereas the AI ASG can have exclusive read/write access to S3 training buckets.
 - Security Groups can be strictly chained: the Frontend ASG security group only accepts traffic from the ALB; the Backend ASG security group only accepts traffic from the Frontend ASG; the Database and AI tiers are accessible only by the Backend ASG.
- 4. Decoupled Deployments & Rolling Updates:**
 - Code updates to the RAGFlow AI stack can be rolled out as a rolling update to the AI ASG (via Instance Refresh) without causing a single millisecond of downtime or disruption to the core Backend and Frontend layers.

2. Statelessness: The Core Prerequisite of ASGs

The foundational rule of Auto Scaling is **statelessness**. Because ASG instances are automatically launched, terminated, and replaced during scaling events (scale-out, scale-in, or health-check replacements), **no persistent state should ever reside directly on an instance's local root disk (EBS)**.

If an application writes files (e.g., user profile pictures, uploaded documents for RAG processing, temporary cached AI model weights) directly to the local filesystem of an EC2 instance, those files will be:

- **Inaccessible** to other instances in the ASG (causing HTTP 404 errors when a load balancer routes subsequent requests to a different instance).
- **Permanently lost** when the specific instance scale-in or gets replaced due to a failed health check.

To solve this, you must decouple your compute tier (ASG) from your storage tier. This brings us to the core decision: **Amazon S3, Amazon EFS, or both?**

3. The Critical Role of AWS Application Load Balancer (ALB) for ASGs

While Auto Scaling Groups (ASGs) manage the lifecycle, elasticity, and count of compute instances, they cannot work in isolation. The **Application Load Balancer (ALB)** serves as the orchestrator and gateway, acting as the bridge between external incoming client traffic and the highly dynamic, ephemeral instances managed by the ASG.

Key Architectural Advantages of the ALB-ASG Integration

1. Abstraction of Ephemeral IP Addresses (Single Entry Point)

Instances inside an ASG are ephemeral by design. During auto-scaling scale-out/scale-in events, or when an instance is replaced due to an unhealthy status check, the instance is terminated and a new one is launched. This results in constantly changing, dynamic IP addresses.

- **The ALB Solution:** The ALB provides a single, high-availability DNS name (e.g., `production-alb-12345678.ap-southeast-5.elb.amazonaws.com`) mapped via Route53 alias records (e.g., `app.linuxmalaysia.com`). Clients only ever target the ALB, and the ALB automatically handles routing to the exact private IP addresses of the currently running ASG instances, abstracting all backend churn from the user.

2. Dynamic Target Registration and Connection Draining (Zero Packet Loss)

As the ASG scales, instances come and go dynamically.

- **Target Registration:** When a new instance launches in the ASG, the ASG automatically registers its IP address with the associated ALB **Target Group** (`target_group_arns`). Once the instance passes initial health checks, the ALB starts sending traffic to it.
- **Connection Draining (Deregistration Delay):** When the ASG scales in or triggers a rolling update, instances are not terminated abruptly. The ALB initiates a **Deregistration Delay** (connection draining) period. During this time, the ALB stops routing *new* client requests to the deregistering instance, but allows active,

in-flight HTTP requests to finish processing. Once all in-flight connections are complete or the timeout is reached, the instance is safely terminated with zero client-side aborted requests.

3. Active Health-Check Monitoring and Auto-Healing Delegation

By default, an ASG only monitors hardware-level EC2 Status Checks (hypervisor status, CPU hardware, host power). If your application layer (e.g., your Spring Boot application, Nginx, or RAGFlow) hangs, deadlocks, runs out of memory, or crashes, the VM still appears “running” and the basic EC2 check passes.

- **The ALB Solution:** By setting the ASG’s `health_check_type` to “**ELB**” (Elastic Load Balancing), you delegate the auto-healing logic to the ALB.
- **How it works:** The ALB continuously probes the application’s web port (e.g., checking `/` or `/health`) at configured intervals. If the application on a specific instance fails to respond with a successful HTTP code (e.g. returns 500, or times out) for the configured `unhealthy_threshold`, the ALB flags the instance as unhealthy and stops routing traffic to it. The ASG detects this unhealthy status and automatically terminates the failed instance, launching a fresh, healthy one in its place.

4. SSL/TLS Termination and Offloading

Managing SSL/TLS certificates and handling the cryptographic overhead of encrypting and decrypting HTTPS connections at individual EC2 instances is computationally expensive and operationally complex.

- **The ALB Solution:** The ALB integrates natively with **AWS Certificate Manager (ACM)**, allowing SSL certificates to be terminated directly at the load balancer level. Traffic is decrypted at the ALB and forwarded securely via private subnets to the ASG instances. This offloads cryptographic workloads from the EC2 instances, freeing up 100% of their CPU resources to focus on executing application logic, database operations, or resource-heavy AI RAG model inference.

5. Security Group Chaining and Network Isolation

For enterprise security, compute instances running backend APIs and AI tools should never be exposed directly to the public internet.

- **The ALB Solution:** Using the ALB, we place our ASG instances inside **fully private subnets** without public IP addresses.
- **Security Group Chaining:** We configure the ASG Security Group to only accept inbound traffic on the application port (e.g., Port 80, 5000, or 8000) from the specific **ALB Security Group**. All direct public access is physically blocked at the network boundary, forcing all traffic to pass through the ALB’s centralized, inspected, and controlled entry point.

6. Host-Based and Path-Based Routing (Powering SoC)

Instead of provisioning separate expensive load balancers for each tier of our architecture, a single shared Application Load Balancer can orchestrate routing rules to separate, specialized target groups.

- **The ALB Solution:** We can define listener rules on a single ALB to direct traffic dynamically:
 - Traffic to `app.linuxmalaysia.com/` is routed to the **Nginx/Frontend ASG Target Group**.
 - Traffic to `app.linuxmalaysia.com/api/*` is routed to the **Backend ASG Target Group**.
 - Traffic to `app.linuxmalaysia.com/rag/*` is routed to the **AI RAGFlow ASG Target Group**. This single-ALB multi-ASG integration allows us to keep our layers decoupled, scale each tier independently, and enforce strict Separation of Concerns while maintaining a lean and highly cost-optimized infrastructure.

4. Storage Guide: Amazon S3 vs. Amazon EFS

To handle data across auto-scaling instances, AWS offers two primary shared storage options. Understanding when to use each—or combining them in a hybrid layout—is critical to a high-performing architecture.

Option A: Amazon S3 (Simple Storage Service)

Amazon S3 is a highly durable, virtually infinite, secure, and cost-effective **object storage service** designed for cloud-native applications.

- **How it works:** Files (objects) are accessed and manipulated directly via HTTP/HTTPS REST API calls (typically using AWS SDKs or the AWS CLI) rather than traditional filesystem commands (like `open()`, `write()`, or `seek()`).
- **Best suited for:**
 - Modern, stateless, cloud-native applications.
 - User-facing uploads (e.g., images, PDFs, videos).
 - Raw data source files before they are ingested, chunked, and embedded into vector databases by AI services like RAGFlow.
 - Long-term backups, application logs, and database snapshots.
- **Pros:**
 - **Incredible Scalability:** Handles millions of concurrent requests effortlessly.
 - **Unmatched Durability:** Designed for 99.999999999% (11 9s) of durability.
 - **Low Cost:** Very inexpensive compared to standard filesystems (\$0.023 per GB/month for standard, even lower for intelligent tiering/infrequent access).
 - **Global Delivery Integration:** Seamlessly integrates with Amazon CloudFront CDN for global low-latency asset delivery.
- **Cons:**
 - **Non-POSIX Compliant:** You cannot “mount” S3 as a traditional directory and perform standard file writes/appends in legacy code without utility layers (like Mountpoint for Amazon S3), which have limitations on write patterns.

Option B: Amazon EFS (Elastic File System)

Amazon EFS is a fully-managed, serverless, POSIX-compliant **network filesystem (NFSv4)** designed to be mounted concurrently by hundreds of EC2 instances across multiple Availability Zones.

- **How it works:** EFS is mounted over the network (port 2049) to a local directory on your EC2 instances. To the operating system, it looks and behaves exactly like a standard local directory.
- **Best suited for:**
 - Legacy or commercial off-the-shelf (COTS) software that expects a standard filesystem path (e.g., `/var/www/uploads` or `/opt/app/data`) and cannot easily be modified to use an AWS SDK.
 - Shared directories requiring low-latency read-write concurrency across multiple compute instances simultaneously.
 - **AI Model Cache:** Sharing large pre-trained AI models or embedding files (e.g., Hugging Face model caches, SentenceTransformers weights) across your RAGFlow/AI ASG instances. Mounting EFS prevents instances from having to download multi-gigabyte files from S3 or external hubs during a scale-out bootstrap, significantly accelerating launch speed.

- **Pros:**
 - **POSIX Compliant:** Supports full directory structures, file permissions (UID/GID), and standard file system tools (tar, gzip, grep).
 - **Elastic Capacity:** Scales automatically up to petabytes without manual provisioning; you only pay for what you use.
 - **Multi-AZ Availability:** Built-in replication across multiple AZs.
- **Cons:**
 - **Higher Cost:** Significantly more expensive than S3 (\$0.30 per GB/month for standard SSD storage, though lifecycle policies can tier unused data to lower-cost tiers).
 - **I/O Overhead:** Network filesystem overhead means file read/write latencies are higher than local SSDs (EBS) or highly-optimized API calls.


5. Architectural Comparison Table

Feature	Amazon S3 (Object Storage)	Amazon EFS (Network File System)
Primary Protocol	HTTP / HTTPS (REST API)	NFSv4 (TCP Port 2049)
POSIX Support	No (requires custom SDK or client)	Yes (standard directory, file permissions, file locking)
Scalability	Virtually infinite; high concurrent request throughput	Automatic capacity scaling; scales throughput with size
Typical Ingress Cost	Free inbound; \$0.023/GB/month (Standard)	\$0.30/GB/month (Standard Storage)
Access Latency	Tens of milliseconds	Single-digit milliseconds (metadata operations can be slower)
Direct Web Serving	Yes (via CloudFront / S3 Public endpoints)	No (must traverse an EC2 instance/web server)
Mounting on EC2	Direct SDK integration preferred (or Mountpoint utility)	Standard Linux mount command via amazon-efs-utils
Best AWS ASG Case	Storing user uploads, large datasets, and logs	Shared AI model weight caches, shared configuration directories

6. The Verdict: Do we need EFS, S3, or Both?

For most modern AWS architectures utilizing Auto Scaling Groups, **using Both in a hybrid design is the industry best practice**. They are not mutually exclusive, but rather complementary.

Here is the architectural distribution for our secure 3-tier layout:



When to use S3 for Auto Scaling:

- **Stateless Web Files:** All user uploads (e.g., PDFs sent to the DMS or RAG system) must be directed immediately to an Amazon S3 Bucket. When a user uploads a document, the Backend API processes it, writes the raw file to S3, and records the metadata (such as the S3 Object URL) in the **RDS PostgreSQL Database**.
- **Scale-Out Safety:** When a new instance boots up inside the ASG, it doesn't need to copy legacy user data. It simply calls S3 URLs on demand.

When to use EFS for Auto Scaling:

- **Accelerating Scale-Out (The Bootstrapping Problem):** If your RAGFlow AI service requires a 4GB language model to process requests, downloading this model from S3 or Hugging Face during a scale-out operation can take several minutes. This severely delays your auto-scaling response. By storing the model weights inside a shared EFS directory (mounted as /mnt/models), any newly launched AI instance can read the models instantly from EFS, reducing scaling boot times from 10 minutes to under 60 seconds.
- **Shared Working Workspace:** If multiple nodes in the AI ASG must cooperate on batch processing files and require immediate POSIX write access.

7. Implementation and Configuration Guide

To implement EFS and S3 storage within your Auto Scaling Groups, configure the following AWS parameters.

7.1 Configuring S3 Integration for ASGs

To allow instances in your ASG to read and write to an S3 bucket securely:

1. **Do NOT use static AWS Access Keys.** Instead, attach an IAM Policy to your ASG instance role:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "s3:GetObject",
        "s3:PutObject",
        "s3:ListBucket"
      ],
      "Resource": [
```

```

        "arn:aws:s3:::my-application-storage-bucket",
        "arn:aws:s3:::my-application-storage-bucket/*"
    ]
}
}
}

```

2. **Launch Template User Data Integration:** In your launch template, you can download startup assets or configurations using the AWS CLI:

```

#!/bin/bash
# Update and install AWS CLI
dnf update -y
dnf install -y awscli

# Download environment secrets or application configuration from S3
mkdir -p /etc/myapp
aws s3 cp s3://my-application-storage-bucket/config/config.json /etc/myapp/config.json

```

7.2 Configuring EFS Integration for ASGs

To mount an EFS volume automatically on launch across multiple AZs:

1. **Security Group Configuration:**
 - Create an **EFS Security Group** (sg-efs).
 - Allow **Inbound TCP Port 2049 (NFS)** from your **ASG Security Group** (sg-asg).
2. **Deploy Mount Targets:**
 - In your OpenTofu/Terraform code, deploy EFS Mount Targets (aws_efs_mount_target) in **each private application subnet** within your VPC.
3. **Launch Template User Data Mounting:** Use the following script within your ASG's Launch Template user_data to install dependencies and mount EFS securely using TLS:

```

#!/bin/bash
# Install EFS mounting utilities (standard in Amazon Linux 2023)
dnf install -y amazon-efs-utils

# Define EFS Target Directory
MOUNT_DIR="/mnt/efs"
mkdir -p $MOUNT_DIR

# Mount using EFS File System ID (Enable TLS for encryption in-transit)
# Replace fs-xxxxxx with your actual EFS ID
mount -t efs -o tls fs-xxxxxx:/ $MOUNT_DIR

# Ensure the mount is persistent across potential reboots
echo "fs-xxxxxx:/mnt/efs efs _netdev,tls,defaults 0 0" >> /etc/fstab

```

8. Cost-Optimization Best Practices

When leveraging S3 and EFS alongside ASGs, implement these practices to minimize operational costs:

1. **EFS Lifecycle Management:** Set EFS Lifecycle policies to automatically move files that haven't been accessed in 14 or 30 days to **EFS Infrequent Access (IA)** or **EFS Archive**. This can reduce EFS storage costs by up to 90%.
2. **S3 Intelligent-Tiering:** Enable S3 Intelligent-Tiering on your buckets. This automatically moves files between frequent, infrequent, and archive access tiers based on real-time usage patterns, without any performance overhead or administrative management.
3. **Mountpoint for Amazon S3:** For large read-heavy file access patterns (like streaming static training datasets to AI nodes), evaluate AWS's official Mountpoint for Amazon S3 as a high-throughput, lower-cost alternative to mounting EFS.

Root Terraform Configuration

[#aws](#) [#cloud](#) [#architecture](#) [#vpc](#) [#alb](#) [#asg](#) [#rds](#) [#waf](#) [#route53](#) [#dns](#) [#postgresql](#)

Root Terraform Configuration

The root folder of the terraform/ directory orchestrates the execution order, specifies environmental variables, defines provider versions, and outputs key service endpoints.

Files Overview

1. providers.tf

Specifies minimum required Terraform version requirements and registers external providers.

- **Minimum Terraform Version:** >= 1.5.0
- **AWS Provider Version:** ~> 5.0
- **Region configuration:** Dynamically loaded from variables (defaults to ap-southeast-5).

```
terraform {
  required_version = ">= 1.5.0"
  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 5.0"
    }
  }
}

provider "aws" {
  region = var.aws_region
}
```

2. Infrastructure Domain Modules (main.tf, compute.tf, database.tf, web.tf)

Serves as the central manifest, organized cleanly across domain-focused files within the root module to call submodules in logical dependency order and pass cross-module attributes:

- **main.tf:** Foundational network infrastructure (vpc) and microsegmentation security rules (security_groups).
- **compute.tf:** Workload compute instances including Auto Scaling Groups (asg), standalone EC2 nodes (standalone_ec2), and secure SSH bastions (jumphost).
- **database.tf:** Isolated data persistence and caching layers (rds, elasticache_valkey).
- **web.tf:** Web routing, edge security filtering, and DNS delegation (alb, waf, route53).

3. variables.tf

Declares environmental configurations, default configurations, and input parameter structures.

- **Defaults for Malaysia Deployment:**
 - aws_region: "ap-southeast-5"
 - db_engine: "postgres"
 - db_engine_version: "16"
 - db_instance_class: "db.t4g.micro" (Graviton architecture)
 - instance_type: "t4g.micro" (Graviton compute)
 - enable_route53: true (whether to provision a Route 53 hosted zone and record)
 - domain_name: "linuxmalaysia.com" (the target root domain name)
 - subdomain: "app" (the frontend application subdomain)

4. outputs.tf

Exposes crucial deployment endpoints, such as the public ALB endpoint, the RDS endpoint, the WAF Web ACL ARN, and the Route 53 Name Servers for domain delegation.

```
output "alb_dns_name" {
  value     = module.alb.alb_dns_name
  description = "Public Load Balancer Endpoint"
}

output "rds_endpoint" {
  value     = module.rds.db_instance_endpoint
  description = "Isolated RDS Database Endpoint"
}

output "waf_web_acl_arn" {
  value     = module.waf.waf_web_acl_arn
  description = "WAF v2 Web ACL ARN"
}

output "route53_name_servers" {
  value     = try(module.route53[0].name_servers, [])
  description = "Route 53 delegated Name Servers for registrar configuration"
}

output "route53_fqdn" {
```

```
    value      = try(module.route53[0].fqdn, "")
    description = "Route 53 FQDN pointing to the Application Load Balancer"
}
```

5. terraform.tfvars.example

A template file used to configure local infrastructure credentials and system settings safely. Copy this template before running your deployments:

```
cp terraform.tfvars.example terraform.tfvars
```

OpenTofu Migration & AWS Support Research Guide

#aws #cloud #architecture #vpc #alb #asg #rds #waf #bastion #ssl #disaster-recovery

OpenTofu Migration & AWS Support Research Guide

This document presents a comprehensive feasibility study, architectural research, and practical migration path for transitioning our AWS 3-Tier Infrastructure configuration from **HashiCorp Terraform** to **OpenTofu**.

OpenTofu is a production-ready, open-source infrastructure-as-code (IaC) tool under the governance of the Linux Foundation (part of the Cloud Native Computing Foundation - CNCF). It was spawned as a community-driven fork of Terraform following its transition to the Business Source License (BSL).

1. Executive Summary

- High Compatibility:** OpenTofu serves as a drop-in replacement for Terraform configurations (especially those compatible with Terraform 1.5.x, which this project uses with its `>= 1.5.0` requirement).
- Zero AWS Resource Impact:** Migrating the deployment engine from `terraform` to `tofu` does not affect currently running AWS resources (such as VPCs, ALBs, ASGs, WAFv2, or RDS).
- Same AWS API Calls:** Both engines utilize the standard AWS Provider which translates declarative configuration into AWS API calls using the AWS SDK for Go.
- Enhanced Open Source Features:** OpenTofu offers advanced features like client-side state encryption, improved variable handling, and native resource exclusion (`tofu plan -exclude="..."`).

2. AWS Support and Compatibility Matrix

AWS natively supports OpenTofu indirectly and directly across its entire ecosystem, since OpenTofu acts as a client orchestrator of AWS APIs via the AWS Provider.

AWS Feature / Service	Support Level	Technical Integration Mechanism
AWS APIs & Resources	Native & Complete	Handled seamlessly via the <code>hashicorp/aws</code> provider, which is mirrored directly in the OpenTofu Registry. All resources (VPC, WAFv2, ALB, ASG, RDS) behave identically.
IAM Authentication	Native & Complete	Uses standard AWS SDK Go credential resolution chain (profiles, env vars, instance profiles).
OIDC / Web Identity	Native & Complete	Fully supports AWS OpenID Connect (OIDC) authentication, enabling passwordless GitHub Actions runs.
S3 Backend (State Storage)	Native & Complete	The <code>s3</code> state backend in OpenTofu matches the Terraform <code>S3</code> backend exactly, utilizing SSE, S3 Versioning, and IAM-based access.
DynamoDB (State Locking)	Native & Complete	Fully supports DynamoDB table state locking to prevent concurrent runs.
AWS CodeBuild	Complete (via scripts)	Since CodeBuild executes raw container environments, OpenTofu is fully supported. Tofu can be installed via package manager (APT/YUM) or official binary installation.
AWS CodePipeline	Complete (via CodeBuild)	CodePipeline orchestrates steps. Running OpenTofu within a CodePipeline pipeline is done via CodeBuild actions.
AWS Proton	Complete (via custom runner)	AWS Proton allows specifying custom IaC engines through custom container-based provisioning, allowing seamless execution of OpenTofu.
AWS Systems Manager (SSM)	Supported	Parameter Store and Secrets Manager are accessed natively via data sources or backend credential configurations.

3. AWS Authentication Integration

OpenTofu inherits the exact same credential chain precedence as the AWS SDK and AWS CLI:

- Static Credentials:** Environment variables `AWS_ACCESS_KEY_ID`, `AWS_SECRET_ACCESS_KEY`, and `AWS_SESSION_TOKEN`.
- Shared Config & Credentials Files:** Located at `$HOME/.aws/config` and `$HOME/.aws/credentials` (leveraging `AWS_PROFILE`).
- Web Identity / OIDC Tokens:** Loaded via `AWS_WEB_IDENTITY_TOKEN_FILE` and `AWS_ROLE_ARN`. This is crucial for secure, keyless execution in **GitHub Actions** and **AWS EKS / ECS**.
- ECS Task Roles / Container Credentials:** Loaded using `AWS_CONTAINER_CREDENTIALS_RELATIVE_URI`.
- EC2 Instance Profile Metadata:** Evaluated when running directly on an EC2 instance (e.g., inside an administrative workstation or bastion host).

4. State Management and S3 Backend

OpenTofu's `S3` state backend is completely compatible with Terraform's backend. The configuration remains virtually identical:

```
terraform {
  backend "s3" {
    bucket = "my-opentofu-state-bucket"
    key    = "state/terraform.tfstate"
    region = "ap-southeast-5"
    dynamodb_table = "my-opentofu-lock-table"
    encrypt = true
  }
}
```

Note: OpenTofu accepts the `terraform` configuration block for backward compatibility. There is no need to rename this block to `tofu` in your `.tf` files.

State Migration Safeties:

- Before running OpenTofu on an existing Terraform-managed deployment, create a backup of your S3 state file:
`aws s3 cp s3://my-opentofu-state-bucket/state/terraform.tfstate ./terraform.tfstate.backup`
- Running `tofu init` will recognize the existing state file and integrate with it natively.

5. Command Mapping

OpenTofu is designed with direct CLI command parity to make the transition trivial for developers:

Terraform Command	OpenTofu Equivalent	Purpose
<code>terraform init</code>	<code>tofu init</code>	Initializes directories, downloads providers/modules
<code>terraform fmt</code>	<code>tofu fmt</code>	Automatically formats configuration files
<code>terraform validate</code>	<code>tofu validate</code>	Checks configuration semantics and syntax
<code>terraform plan</code>	<code>tofu plan</code>	Outputs execution plan of pending AWS updates
<code>terraform apply</code>	<code>tofu apply</code>	Executes updates against active AWS APIs
<code>terraform destroy</code>	<code>tofu destroy</code>	Destroys all provisioned resources securely
<code>terraform state <subcmd></code>	<code>tofu state <subcmd></code>	Inspects or alters local/remote state records
<code>terraform import</code>	<code>tofu import</code>	Imports existing AWS resources into the state

Advanced OpenTofu Commands:

- **Excluding Resources Dynamically:**

```
tofu plan -exclude="module.rds"
tofu apply -exclude="module.rds"
```

(Terraform only supports targeting specific resources via `-target`, whereas OpenTofu supports negative targeting with `-exclude`).

6. How to Run OpenTofu in AWS Environments

A. AWS CodeBuild Integration

To run OpenTofu in an AWS CodeBuild environment, include the following in your `buildspec.yml`:

```
version: 0.2

phases:
  install:
    commands:
      # Install OpenTofu repository and binary
      - curl -s https://get.opentofu.org/install/opentofu-install.sh -o opentofu-install.sh
      - chmod +x opentofu-install.sh
      - ./opentofu-install.sh --install-method standalone
      - tofu --version
  pre_build:
    commands:
      - cd terraform
      - tofu init -backend-config="bucket=${STATE_BUCKET}"
  build:
    commands:
      - tofu plan -out=tfplan
  post_build:
    commands:
      - echo "Build completed on `date`"
```

B. AWS CodePipeline Integration

AWS CodePipeline does not have a dedicated action type for OpenTofu or Terraform out of the box. However, the standard best practice is to use **AWS CodeBuild** as a pipeline stage:

1. **Source Stage:** Code is retrieved from AWS CodeCommit, GitHub, or Amazon S3.
2. **Build Stage (CodeBuild):** CodeBuild installs and runs OpenTofu to generate an execution plan (`tofu plan`).
3. **Approval Stage (Optional):** Manual approval step inside CodePipeline to verify the plan.
4. **Deploy Stage (CodeBuild):** CodeBuild applies the plan (`tofu apply tfplan`).

C. AWS Proton Integration

AWS Proton coordinates infrastructure provisioning. For teams wanting to use OpenTofu:

1. Register a **Custom Provisioning** template.
2. Configure Proton to run an AWS CodeBuild project.
3. In the CodeBuild project, fetch the OpenTofu CLI and execute standard plan and apply commands mapped to the Proton service input schema.

7. Execution and Migration Plan for this Repository

To fully execute this replacement in this codebase, we perform the following steps:

1. **Report & Documentation:** Write this comprehensive guide (`docs/opentofu-migration.md`) and register it in the main Jekyll configuration and index (`docs/index.md`).
2. **Scripts Migration:** Modify CLI helpers `scripts/deploy.sh` and `scripts/destroy.sh` to transition to `tofu` while maintaining safety structures.
3. **GitHub Actions CI/CD Migration:** Update `.github/workflows/terraform.yml` to rename it to `opentofu.yml`, using the official `opentofu/setup-opentofu@v1` runner to ensure native, secure plan/apply on pushes.

4. **General References Refactoring:** Replace Terraform terminology and setup notes in `README.md`, `docs/scripts.md`, and `docs/cicd.md` with OpenTofu standards.

OpenTofu AWS Infrastructure Building, Simulation, & Multi-Agent Collaboration Runbook

#opentofu #aws #simulation #testing #multi-agent #collaboration #devops #security

[DEVOPS EXECUTION] — Infrastructure Engineers, SREs, & AI Agents


OpenTofu AWS Infrastructure Building, Simulation, & Multi-Agent Collaboration Runbook

This guide establishes the standard operating procedures for developing, simulating, testing, and deploying OpenTofu Infrastructure-as-Code (IaC) targeting Amazon Web Services (AWS) without requiring direct AWS cloud environment credentials or live API access.

Due to security policies and isolation requirements, all OpenTofu module changes, architectural updates, and deployment flows **must be simulated and verified locally** using automated unit testing, static analysis, and mock scenario suites prior to merging. Furthermore, this document specifies protocols for human software engineers and autonomous AI agents (such as **Google Jules** and **Google Antigravity**) collaborating across forks and feature branches.

1. Architectural Strategy & Offline Simulation Philosophy

When direct AWS access is restricted for security, cost control, or compliance reasons, OpenTofu code quality and cloud compatibility are maintained through a multi-layer **Offline Cloud Simulation Framework**:



Core Simulation Principles:

- Zero Cloud Credential Requirement:** No `AWS_ACCESS_KEY_ID` or active internet connection to AWS API endpoints is required during local simulation and unit testing.
- Provider Offline Schema Validation:** Standard provider blocks utilize static AWS provider definitions (`hashicorp/aws >= 5.0`) which validate configuration structure locally. In air-gapped or offline environments, initialize the OpenTofu working directory using `tofu -chdir=terraform init -backend=false` (with local provider mirror/cache configured) before running `tofu -chdir=terraform validate`.
- AST & Pattern Testing:** Pytest unit tests analyze the Abstract Syntax Tree (AST) and regex structure of `.tf` files to enforce strict security constraints (e.g. IMDSv2, isolated security groups, Graviton `t4g.*` sizing).

2. OpenTofu HCL Code Structure & Building Guidelines

All OpenTofu configurations must follow a modular 3-Tier architecture design under the `terraform/` directory:

```
terraform/
├── main.tf                # Network foundations & core security group calls
├── compute.tf             # Auto Scaling Group (ASG) & Standalone compute definitions
├── database.tf            # Multi-AZ RDS PostgreSQL & ElastiCache Valkey instances
├── web.tf                 # ALB, WAFv2, and Route 53 DNS resources
├── providers.tf           # Provider configuration (AWS regional defaults: ap-southeast-5)
├── variables.tf           # Strictly typed input variables and defaults
├── outputs.tf             # Exported attributes (ALB DNS, DB Endpoints, SG IDs)
├── modules/
│   ├── vpc/               # Encapsulated reusable modules
│   ├── security_groups/   # Multi-AZ subnets (Public, App, DB)
│   ├── alb/               # Zero-Trust ingress/egress rules
│   ├── asg/               # Application Load Balancer & target groups
│   ├── rds/               # Launch templates with IMDSv2 & auto-healing
│   ├── elasticsearch/     # Isolated PostgreSQL Multi-AZ engine
│   ├── jumphost/          # Valkey cache cluster
│   ├── route53/           # Hardened Cyberjaya SSH Bastion
│   ├── standalone_ec2/    # DNS alias mapping
│   └── waf/               # Development/staging compute nodes
└──                         # WAFv2 IP rate limiting rules
```

Standard Module Rules:

- **IMDSv2 Strict Enforcement:** Launch templates (modules/asg/main.tf and modules/standalone_ec2/main.tf) must specify metadata_options with http_tokens = "required" and http_put_response_hop_limit = 1.
- **Zero-Trust Network Isolation:** Database security groups must only allow ingress on port 5432 originating from compute tier security groups (security_groups = [aws_security_group.asg_sg.id]). Direct CIDR ingress (0.0.0.0/0) to DB or App subnets is prohibited.
- **ALB-Aware Auto-Healing:** ASG modules must enforce health_check_type = "ELB" to synchronize instance lifecycles with ALB target group status.
- **Regional Architecture Defaults:** Default region is pinned to ap-southeast-5 (Malaysia), utilizing Graviton instances (t4g.micro / t4g.medium / db.t4g.micro) and Valkey caching (cache.t4g.micro).

3. AWS Cloud Scenario Testing Matrix & Unit Test Integration

Each step in the OpenTofu code building lifecycle must be verified against simulated cloud scenarios using automated pytest unit tests in tests/test_opentofu_simulation.py.

AWS Cloud Scenario Testing Matrix:

Cloud Scenario / Constraint	Test Coverage	Expected Outcome	Verification Tool
1. SSRF Metadata Hardening	Compute Launch Templates & EC2 Instances	http_tokens = "required", http_put_response_hop_limit = 1	test_opentofu_imdsv2_enforcement
2. DB Network Isolation	Database Security Groups	Ingress strictly restricted to App SG on port 5432; no 0.0.0.0/0 ingress	test_opentofu_db_security_group_isolation
3. ALB Auto-Healing Integration	ASG Configuration	health_check_type = "ELB" enabled across all scaling groups	test_opentofu_asg_elb_health_check
4. Graviton Regional Alignment	Variables & Defaults	Default region ap-southeast-5, Graviton t4g.* & db.t4g.* types	test_opentofu_graviton_regional_defaults
5. Valkey Open-Source Caching	ElastiCache Module	Caching engine set to valkey with default port 6379	test_opentofu_valkey_caching_config
6. Jumphost Cyberjaya Whitelisting	Jumphost SG	SSH ingress restricted to office CIDR (103.188.0.0/16 or custom)	test_opentofu_jumphost_whitelisting
7. Multi-AZ Network Topology	VPC Module	2+ Availability Zones (ap-southeast-5a, ap-southeast-5b) spanned	test_opentofu_vpc_multi_az_layout
8. Module Completeness	Root Terraform Modules	All 10 modules (vpc, security_groups, alb, asg, rds, elasticache, etc.) present	test_opentofu_module_structure_completeness

Executing Local Simulation Unit Tests:

```
# Run all simulation unit tests locally
pytest tests/test_opentofu_simulation.py -v

# Or use the standalone simulation runner script
./scripts/simulate.sh
```

4. Multi-Human and Multi-AI Agent Collaboration Protocol

To ensure seamless collaboration between human software engineers and AI agents (Google Jules, Google Antigravity, external PR contributors), all code and documentation contributions must strictly observe the following protocol:

A. Branching & Forking Strategy

- **Branch Naming:** Feature branches must use clear, descriptive prefixes:
 - feature/opentofu-module-name (e.g. feature/opentofu-waf-rules)
 - fix/security-group-rule (e.g. fix/db-port-isolation)
 - docs/aws-simulation-update (e.g. docs/opentofu-runbook)
- **Fork Contributions:** External human contributors or AI agents working in forks must ensure pull requests target the main branch of the upstream repository.

B. OKF Specification & Frontmatter Rules

- Every newly created or updated .md file must include **OKF v0.1 YAML Frontmatter** starting on Line 1, Column 1:

```
---
layout: "default"
okf_version: "0.1"
type: "Guide"
title: "Descriptive Page Title"
timestamp: 2026-08-22T08:15:00Z
topics: ["opentofu", "aws", "simulation", "testing"]
---
```
- All string values with special characters (colons, parentheses, emojis) **must be double-quoted**.
- Always run python scripts/prepare_docs.py to re-format frontmatter and compile llms.txt / llms.xml.

C. Automated PR Quality Gates & Pre-Commit Workflow

Before submitting a PR or merging any code:

1. Format & Validate OpenTofu Code (if OpenTofu CLI is installed):

```
tofu fmt -recursive terraform/
tofu -chdir=terraform init -backend=false
tofu -chdir=terraform validate
```

(Note: For air-gapped or offline use without AWS access, configure a local provider mirror/cache).

2. Execute Full Test Suite:

pytest

3. Execute Pre-Commit Verification:

- Ensure all relative Markdown links resolve correctly (test_markdown_relative_links_integrity).
- Run python scripts/prepare_docs.py and verify all tests pass cleanly.

5. Deployment & Offline Simulation Walkthrough on Any Linux Environment

You can git clone this repository on any Linux distribution (Ubuntu, Debian, RHEL, Fedora, Arch, Amazon Linux) to begin AWS deployment or test newly added scripts and OpenTofu modules via simulation:

```
# 1. Clone repository on your Linux environment
git clone https://github.com/songketmail/aws-3tier-deployment-for-ai-infra.git
cd aws-3tier-deployment-for-ai-infra

# 2. Configure variables and credentials
cp terraform/terraform.tfvars.example terraform/terraform.tfvars

# 3. Test configurations and custom added scripts via offline simulation
./scripts/simulate.sh

# 4. Begin live AWS deployment (when AWS credentials/CLI are configured)
./scripts/deploy.sh

# 5. Cleanly destroy deployed AWS infrastructure
./scripts/destroy.sh
```

Verification Checklist for New OpenTofu Modules or Custom Scripts:

1. **Step 1: Write Declarative HCL Code or Script Changes** Place your configuration in the appropriate terraform/ domain file (compute.tf, database.tf, web.tf, or main.tf) or module subdirectory, or add custom scripts under scripts/.
2. **Step 2: Run Offline Simulation Script** Execute ./scripts/simulate.sh to run pytest simulation suites (tests/test_opentofu_simulation.py) and static OpenTofu validation checks.
3. **Step 3: Run Document Preparation Tool** Execute python scripts/prepare_docs.py to sync all indexes and validate OKF frontmatter across modified docs.
4. **Step 4: Create PR & Request Review** Push changes to your feature branch, open a PR with detailed descriptions, and await automated CI verification.

Deep State of Mind (DSOM) Harisfazillah Jamel (LinuxMalaysia) 2026-08-22

Modernizing Big Data Analytics Architecture: 3-Tier Infra Deployment Blueprint

#bda #lakehouse #aws #hybrid #proxmox #rke2 #ceph #mcp

Modernizing Big Data Analytics Architecture: Establishing an Authoritative Open-Source Single Source of Truth

Executive Overview & Architectural Context

This architectural blueprint outlines the modernization of a legacy Big Data Analytics (BDA) platform into a 100% open-source enterprise data lakehouse. Designed as an authoritative **Single Source of Truth (SSoT)**, the platform consolidates environmental, geological, hydrological, forestry, and wildlife datasets across multi-departmental enterprise operations.

Enterprise Assumptions & Public Adaptation

To adapt this enterprise blueprint for public deployment and general industry adoption, all specific public-sector identifiers have been replaced with standard organizational agency designations and internal domain endpoints:

- **Enterprise Environmental Data Platform:** Formerly legacy BDA infrastructure.
- **Internal Endpoints:** Standardized to `https://bda.enterprise.internal` and `https://dashboard.enterprise.internal`.
- **Six (6) Departmental Domains:**
 1. **Department of Wildlife Conservation (DWC):** Handles wildlife encounters, species monitoring, and spatial movement corridors.
 2. **Department of Mineral & Geoscience (DMG - Hydrogeology):** Maps hydrogeological borehole reserves and groundwater potential.
 3. **Department of Mineral & Geoscience (DMG - Geotechnical Landslides):** Calculates slope stability risks and landslide hazard thresholds.
 4. **Department of Forestry (DOF):** Tracks forest fire incidents, biomass susceptibility, and concession boundaries.
 5. **National Water & Climate Institute (NWCI):** Models hydrological vulnerability and climate adaptation indices.
 6. **Meteorological Services Department (MSD) & Global Satellite Thermal Sensing Network (GSTRN):** Serves precipitation curves, weather telemetry feeds, and thermal hotspot alerts.

Architectural Deconstruction of the Legacy BDA Environment

The legacy Big Data Analytics (BDA) platform operated as an early-generation big data deployment characterized by tightly coupled compute and storage architectures, unmanaged database proliferation, fragile point-to-point data ingestion methods, and proprietary visualization systems.

Legacy Technical Bottlenecks

- **Ingestion Tier:** Dependent on fractured integration pathways. Precipitation data entered via unmonitored local network folder shares; thermal hotspot alerts arrived through unparsed automated emails; geological hazard boundaries were retrieved via point APIs; and departmental datasets were manually transferred over SSH File Transfer Protocol (SFTP) or raw file uploads. Pre-ingestion schema validation, data contracts, and automated provenance tracking were absent.
- **Storage Backbone:** Fragmented across multiple disparate storage fabrics. Distributed file storage was divided between an aging Apache Hadoop cluster—comprising two NameNodes, three DataNodes, and a Network File System (NFS) gateway—and a six-node GlusterFS cluster (GlusterFS-01 through GlusterFS-06).
- **Relational Database Tier:** Split across uncoordinated relational database instances:
 - A high-availability pair of MariaDB nodes (MariaDB-HA1, MariaDB-HA2) dedicated to web portal management.
 - A five-node MariaDB cluster (MariaDB-01 through MariaDB-05) housing specific data project tables.
 - A three-node PostgreSQL cluster (Postgresql-app-01 through Postgresql-app-03) supporting application state.
 - An independent three-node PostgreSQL cluster (Postgresql-01 through Postgresql-03) executing business queries.
- **Compute Transformations:** Processing ran on legacy Red Hat WildFly application server instances, executing bespoke Java scripts without modern orchestration frameworks or pipeline abstractions.
- **Presentation Tier:** Divided between an end-of-life CMS hosting the primary web portal (`bda.enterprise.internal`), a custom portal (`dashboard.enterprise.internal`), and a proprietary Tableau Server cluster managed by Tableau Server Manager (TSM).

Legacy Subsystem	Legacy BDA Architecture	Structural Bottlenecks & Failure Modes
Ingestion Tier	Point-to-point SFTP, local folder shares, automated email parsing, manual web uploads.	Absence of schema validation; lack of rate-limiting; silent pipeline failures upon upstream payload modifications; missing audit trails.
Distributed Storage	Hadoop HDFS (2 NameNodes, 3 DataNodes, NFS), GlusterFS (6 Nodes).	Tight compute-storage coupling; NameNode memory limits on small files; POSIX locking overhead; lack of object-level immutability.
Relational Database Tier	Disparate MariaDB clusters (7 instances) and PostgreSQL clusters (6 instances).	Siloed datasets; inconsistent business definitions across departments; uncontrolled replication; database maintenance overhead.
Processing & Analytics	WildFly application servers running custom cleansing and merging scripts.	Monolithic processing engines; lack of parallel distributed compute; inability to handle large geospatial vector calculations efficiently.
Web Portal Tier	Legacy CMS (Master and 2 HA nodes), custom application servers.	End-of-life CMS vulnerabilities; tightly coupled presentation logic; decentralized user authentication repositories.
Visualization Tier	Tableau Server cluster (TSM, 3 worker nodes, 2 load balancers, Desktop authoring).	High recurring proprietary licensing fees; vendor lock-in; proprietary workbook formats; restricted external agency sharing.

Target 100% Open-Source Lakehouse Architecture: Storage & Compute Decoupling

The modern data lakehouse pattern physically and logically decouples persistent storage from distributed compute engines. HDFS and GlusterFS are replaced with S3-compatible object storage clusters governed by an open table format.

Architectural Foundations

1. **Persistent Object Storage & Immutability Contract:** Uses Ceph (via RADOS Gateway), MinIO Enterprise Object Store, or AWS S3.
 - **Object Lock Retention Contract:** S3 Object Lock enforces documented software guarantees rather than physical/hardware locks. Buckets require **Bucket Versioning** to be enabled. Tier 0 evidence storage enforces a defined retention policy (e.g. 7-year statutory compliance or 365-day operational compliance) in **Compliance Mode** alongside legal-hold toggles for regulatory holds. Tested, equivalent Object Lock retention behaviors are documented and validated across AWS S3, Ceph RADOS Gateway (RGW), and MinIO Enterprise Object Store.
 - **Iceberg Maintenance Separation:** Immutable source/evidence storage is separated from mutable analytical lakehouse tables. Apache Iceberg tables reside in dedicated buckets or partitions configured with **Governance Mode** Object Lock or unlocked object paths. Catalog maintenance jobs (such as snapshot compaction, orphan file purging, and snapshot expiration) operate either on unlocked object paths or via a narrowly scoped maintenance service role authorized with `s3:BypassGovernanceRetention` permission sending the `x-amz-bypass-governance-retention: true` request header. Retention compatibility is verified via automated validation tests prior to production compaction execution.
2. **Universal Open Table Format (Apache Iceberg):** Replaces relational sprawl and file directories with Apache Iceberg backed by columnar Parquet files. Iceberg provides ACID transactions, in-place schema evolution, hidden partitioning, and snapshot time-travel querying.
3. **Open Catalog (Apache Polaris / Project Nessie):** Manages table registration and commit resolution via an open Iceberg REST catalog.
4. **Decoupled Distributed Compute:**
 - **Trino:** Massive parallel processing (MPP) SQL query engine executing low-latency queries over Iceberg tables.
 - **Apache Spark + Apache Sedona:** Manages batch transformations, Change Data Capture (CDC), and distributed spatial processing via SpatialRDDs and GeoParquet.
5. **Operational Serving Store:** High-availability PostgreSQL enhanced with PostGIS for sub-millisecond point queries and spatial dashboard caching.

Architectural Layer	Legacy BDA Infrastructure	Modern Open-Source Replacement	Core Architectural Capabilities
Distributed Object Storage	Hadoop HDFS, GlusterFS.	Ceph Object Storage / MinIO Object Store / AWS S3.	Horizontal scale-out; unified S3 API; software-enforced WORM S3 Object Lock retention; small-file bottleneck elimination.
Open Table Format	Unstructured CSV/JSON files, uncoordinated DB tables.	Apache Iceberg (backed by Apache Parquet).	Serialized ACID transactions; schema and partition evolution without data restructuring; snapshot time-travel.
Lakehouse Catalog	Custom MariaDB schemas, local directories.	Apache Polaris (Incubating) / Project Nessie.	Open REST catalog standard; centralized table metadata; cross-engine concurrency arbitration; credential vending.
Analytical Query Engine	WildFly application servers, local SQL engines.	Trino Distributed SQL Query Engine.	In-memory massively parallel processing; sub-second analytical SQL execution; multi-catalog federation; CBO optimization.
Geospatial Processing Engine	Local spatial libraries, fragmented PostGIS instances.	Apache Sedona executing on Apache Spark + GeoParquet.	Distributed spatial indexing (R-Tree, Quad-Tree); distributed spatial joins; native GeoParquet vector processing.
Operational Serving Layer	Fractured MariaDB and PostgreSQL clusters.	Consolidated HA PostgreSQL with PostGIS extension.	High-concurrency spatial index caching; sub-millisecond point queries; serving boundary layers directly to visualizers.

Data Provenance, Lineage, and Metadata Architecture: The Human-to-AI Quarantine Model

To guarantee data integrity, the architecture enforces a strict boundary: artificial intelligence assists with automation and query processing but is strictly barred from modifying ground-truth human data.

Three-Tier Physical & Logical Classification

- Tier 0: Golden Human Truth (Authoritative SSoT)
 - └─ Immutable storage on MinIO/Ceph/S3 with S3 Object Lock (Compliance Mode)
 - └─ Requires human cryptographic signatures and ODCS contract validation
 - └─ Strictly zero unvalidated AI-generated records permitted
- Tier 1: Machine and Sensor Ingestion
 - └─ Storage on standard S3 buckets with Governance Mode Object Lock
 - └─ Direct telemetry feeds: MSD precipitation, NWCI hydrological logs, GSTRN thermal hotspots
 - └─ Pre-ingestion validation via Data Contract CLI; human audit required for promotion
- Tier 2: AI Operational and Analytical Sandbox (Isolated Quarantine)
 - └─ Isolated object buckets with automated 30-day S3 Lifecycle expiration policy
 - └─ Ephemeral storage for LLM intermediate runs, synthetic models, and MCP outputs
 - └─ Scheduled CronJob cleanup controller with secure-deletion rules for local volumes
 - └─ Strict access barriers preventing automated writing or promotion to Tier 0

Data Contracts & Lineage Instrumentation

- **Bitol Open Data Contract Standard (ODCS v3.1.0):** Defines schema models, physical data types, nullability rules, geospatial bounding boxes, and mandatory provenance metadata. Incoming payloads violating contracts are diverted to dead-letter quarantine queues.
- **OpenLineage & Custom enterprise_provenance Facet:** Captures operational metadata events. The `custom enterprise_provenance` facet records origin types, verification tier, human author identity, AI involvement flags, and SHA-256 payload signatures.

Governance Parameter	Tier 0: Golden Human SSoT	Tier 1: Machine & Sensor Telemetry	Tier 2: AI Sandbox & Analytics
Primary Institutional Purpose	Authoritative national truth, statutory policy formulation, certified legal record.	Empirical environmental observation, telemetry aggregation, baseline monitoring.	Exploratory modelling, scenario simulation, predictive risk computation.
Storage Technology & Retention	Distributed Object Store; S3 Object Lock in Compliance Mode (Versioning enabled).	Distributed Object Store; S3 Object Lock in Governance Mode.	S3-compatible bucket with 30-day S3 Lifecycle purge + CronJob volume cleanup.
Allowable Ingestion Sources	Certified human domain surveys, gazetted boundaries, signed enterprise records.	Direct telemetry streams: MSD rain gauges, NWCI water sensors, GSTRN hotspots.	Model outputs, MCP pipeline agents, synthetic climate projections.
AI Role & Permissions	Read-only access via certified tools. Zero automated AI write access permitted.	Machine learning models execute cleansing, deduplication, and anomaly detection.	Unrestricted generative and predictive computation within sandboxed perimeter.
Lineage & Validation Standard	Mandatory ODCS v3.1.0 contract validation + OpenLineage <code>enterprise_provenance</code> signing.	Automated ODCS contract validation + deterministic quality assertion checks.	OpenLineage job execution tracking; outputs permanently tagged as <code>AI_GENERATED</code> .

Governance Parameter	Tier 0: Golden Human SSoT	Tier 1: Machine & Sensor Telemetry	Tier 2: AI Sandbox & Analytics
Promotion Criteria	Terminal authoritative tier; updates require formal versioning and re-signing.	Promoted to Tier 0 only after automated DQ validation and human officer sign-off.	Cannot be promoted directly; requires distillation and formal human certification.

Model Context Protocol (MCP) and AI Sandboxing: Confining AI to Operational Tooling

The platform implements the **Model Context Protocol (MCP)** (donated by Anthropic to the Agentic AI Foundation under the Linux Foundation) to standardize AI tool interactions over JSON-RPC 2.0 while enforcing strict security containment.

MCP Primitive Restrictions & Guardrails

- **Resources:** Read-only access to OpenMetadata data dictionaries, schema definitions, and sanitized database views.
- **Tools:** Restricted to stateless operational utilities (SQL linting, contract validation, DAG monitoring). Write verbs (INSERT, UPDATE, DELETE, DROP, Iceberg commits) are blocked at the API gateway.
- **Prompts:** Structured, version-controlled interaction templates enforcing governance rules.
- **Security Boundaries:** MCP servers run in isolated containers within a dedicated DMZ. Database connections operate with SET SESSION CHARACTERISTICS AS TRANSACTION READ ONLY. Outputs are written strictly to Tier 2 scratch storage (S3-compatible buckets with 30-day S3 Lifecycle expiration rules) and tagged with ai_generated_data: true.

MCP Server Identifier	Business Domain Alignment	Operational Capability	Primitive Scope	Access Permissions & Safeguards
mcp-catalog-context	Enterprise metadata management.	Exposes schema definitions, business terms, and lineage graphs to assist analysts.	Resources: schema://* Prompts: glossary_search	OpenMetadata REST API; Read-only metadata token; zero raw data payload access.
mcp-trino-query-gen	Departmental analytics & reporting.	Translates natural language queries into optimized ANSI SQL for Trino execution.	Tools: validate_sql, explain_query Prompts: sql_assist	Trino Query Engine; Read-only credentials; execution timeout limited to 30 seconds.
mcp-pipeline-monitor	Environmental telemetry queues.	Monitors Airflow DAG runs and NiFi queues; identifies bottlenecks and backpressure.	Tools: get_dag_status, read_flow_metrics	Airflow REST API & NiFi Diagnostics API; Read-only monitoring role.
mcp-contract-linter	Data governance & ingestion gates.	Validates draft data contracts against ODCS v3.1.0 specifications; analyzes drift.	Tools: lint_contract, diff_schema	Local isolated container execution; no network egress; ephemeral scratch memory.

Enterprise Catalog, Governance, IAM, and API Perimeter

- **Enterprise Data Catalog (OpenMetadata):** Powered by PostgreSQL and OpenSearch. Captures column-level lineage, tracks ODCS contracts, supports Tag-Based Access Control (TBAC), and maps geospatial metadata to ISO 19115:2003 standards.
- **Identity & Access Management (Keycloak):** Centralizes authentication via OpenID Connect (OIDC), OAuth 2.0, and SAML 2.0 with Multi-Factor Authentication (MFA) and Single Sign-On (SSO).
- **Perimeter Gateway (Apache APISIX):** Validates JWT tokens, executes dynamic rate-limiting, terminates TLS, and enforces Mutual TLS (mTLS) for inter-service communication.

Governance Subsystem	Legacy BDA Stack	Open-Source Replacement	Enterprise Capabilities & Operational Advantages
Enterprise Data Catalog	MariaDB local data dictionaries; unindexed schemas.	OpenMetadata (backed by PostgreSQL & OpenSearch).	Centralized discovery; automated metadata crawlers; column-level lineage; native ODCS integration.
Lineage & Provenance Engine	Manual documentation, untracked operational scripts.	OpenLineage Standard (with enterprise_provenance facet).	Runtime operational lineage capture; automated tracking across Spark, Airflow, and Trino; cryptographic verification.
Identity & Access (IAM)	Hardcoded credentials, local user tables.	Keycloak Identity & Access Management.	Centralized OIDC / OAuth 2.0; role-based access control (RBAC); Single Sign-On; Multi-Factor Authentication.
API Perimeter Gateway	Unmanaged load balancers, direct port exposures.	Apache APISIX Cloud-Native API Gateway.	High-performance dynamic routing; perimeter JWT validation; TLS termination; IP whitelisting; DDoS rate-limiting.
Geospatial Governance	Undocumented local coordinate systems.	OGC / ISO 19115:2003 Metadata Profiles.	Standardized geospatial metadata; formal EPSG projection definitions (RSO Malaya, Borneo, WGS84).

Ingestion Pipeline Modernization, Application Portal, and Visualization Overhaul


- **Ingestion (Apache NiFi + Apache Airflow):** NiFi handles streaming ingestion, protocol translation, and boundary polling (MSD rainfall, GSTRN thermal satellites). Airflow orchestrates batch processing, contract verification gates, and Iceberg table maintenance.
- **Web Portal Modernization:** Replaces legacy CMS and monoliths with a decoupled **Next.js / React** microfrontend application deployed behind Apache APISIX and Keycloak.
- **Visualization Overhaul (Apache Superset + deck.gl):** Replaces proprietary Tableau Server with Apache Superset. Connects to Trino via SQLAlchemy, leveraging deck.gl for hardware-accelerated 3D geospatial rendering. Keycloak centralizes user identity and issues JWT tokens containing role and departmental claims. When users authenticate to Apache Superset, Keycloak claims map to Superset roles and Security Manager user profiles, which select and apply dataset Row-Level Security (RLS) SQL filter predicates directly in Superset to restrict query results based on user administrative scope.

Platform Component	Legacy BDA Architecture	Modern Open-Source Replacement	Core Functional Capabilities
Data Ingestion Engine	Point-to-point SFTP, manual email parsing, shared folders.	Apache NiFi.	Visual flow design, automated protocol translation, built-in backpressure, end-to-end data provenance.
Pipeline Orchestrator	Static cron schedules, unmonitored WildFly tasks.	Apache Airflow.	Declarative Python DAGs, data contract validation gates, native OpenLineage instrumentation, automated retries.
Public & Admin Portals	Legacy CMS, custom monolithic application portals.	Containerized Next.js / React Web Application.	Headless modern architecture, responsive component design, zero legacy CMS vulnerabilities, optimized API integration.
Identity & Authentication	Local MariaDB/PostgreSQL authentication tables.	Keycloak Identity & Access Management.	Unified SSO, OpenID Connect / OAuth 2.0 federation, MFA enforcement, centralized role mapping.
API Management	Unmanaged load balancers, direct port exposures.	Apache APISIX API Gateway.	Dynamic routing, SSL termination, JWT token validation, rate-limiting, edge request transformation.

Platform Component	Legacy BDA Architecture	Modern Open-Source Replacement	Core Functional Capabilities
Business Intelligence	Tableau Server cluster (TSM, 3 worker nodes, Desktop).	Apache Superset.	100% open-source, unlimited user concurrency, native Trino integration, deck.gl spatial analytics, Row-Level Security.

Three (3) Infrastructure Deployment Solutions


To provide maximum operational flexibility, this architecture is divided into three distinct infrastructure deployment options depending on security requirements, existing hardware, and cloud strategy.



Solution 1: All in Cloud (AWS Native & Cloud Managed Services)

Solution 1 deploys the entire Big Data Analytics platform natively within the AWS Cloud infrastructure (specifically optimized for AWS Asia Pacific regions such as ap-southeast-5 Malaysia).

Portability Qualification: Solution 1 leverages open-source data specifications (Apache Iceberg tables and Apache Parquet files) alongside selected open-source applications (Keycloak, Apache APISIX, OpenMetadata, Apache Superset), but replaces self-hosted platform components with AWS-managed cloud services (AWS Glue Data Catalog, AWS MWAA Airflow, Amazon Bedrock, Amazon Athena, and Amazon EMR Serverless). Consequently, Solution 1 is cloud-dependent and AWS-managed rather than fully vendor-neutral. Organizations requiring complete vendor neutrality and portability should deploy Solution 3 (100% On-Premises) or self-host all components on Amazon EKS.





Core Components & Cloud Services Mapping

1. Persistent Cloud Storage (Tier 0, 1, 2):
 - Tier 0 SSoT:** Amazon S3 Buckets configured with **S3 Object Lock in Compliance Mode** (Versioning enabled), ensuring software-enforced immutability for golden evidence datasets under defined retention periods (e.g., 7 years).
 - Tier 1 Telemetry:** Amazon S3 Buckets with S3 Object Lock in Governance Mode for MSD telemetry and satellite data.
 - Tier 2 AI Sandbox:** Amazon S3 Scratch Buckets with automated S3 Lifecycle Rules configured for 30-day object expiration/auto-purge.
2. Lakehouse Catalog & Metadata:
 - AWS Glue Data Catalog** or **Apache Polaris REST Catalog** hosted on Amazon EKS, managing Apache Iceberg table commits and Parquet file manifests in dedicated buckets.
 - OpenMetadata** hosted on Amazon EKS (backed by Amazon Aurora PostgreSQL and AWS OpenSearch Service) for enterprise data cataloging and column-level lineage.
3. Distributed Query & Processing Engines:
 - Amazon EMR Serverless (Apache Spark + Apache Sedona):** Executes distributed GeoParquet processing, spatial joins, and batch ETL transformations without server provisioning.
 - Amazon Athena / Amazon EMR Trino:** Distributed massively parallel SQL query engine over Iceberg tables via Glue/Polaris REST catalog.
 - Amazon Aurora PostgreSQL (Multi-AZ with PostGIS):** Serves as high-concurrency spatial serving store and low-latency cache for interactive dashboards.
4. Ingestion & Pipeline Orchestration:
 - Amazon Managed Workflows for Apache Airflow (MWAA):** Orchestrates batch pipelines, ODCS data contract CLI validation, and OpenLineage events across all six departmental domains.
 - Apache NiFi on ECS / EKS:** Handles streaming protocol translations and real-time API polling from external sensor sources.
5. Perimeter Security, Ingress & Identity:
 - AWS WAFv2 + Application Load Balancer (ALB) + AWS CloudFront:** Perimeter protection with rate-limiting and OWASP rules.
 - AWS Cognito / Keycloak on ECS:** Unified OIDC/OAuth 2.0 authentication and Multi-Factor Authentication.
 - Apache APISIX on ECS/EKS:** Dynamic routing, mTLS termination, and API rate limiting.
6. Cloud AI & MCP Sandboxing:
 - Amazon Bedrock (Anthropic Claude, Titan) & Amazon SageMaker Endpoints:** Provides foundation models for analytical inferencing.
 - Containerized MCP Servers on AWS Fargate (ECS) / EKS:** Operates with read-only Aurora/Athena database connections and writes outputs exclusively to S3 Tier 2 scratch buckets.

Solution 2: Hybrid - AI On-Premises (Cloud Lakehouse + On-Prem GPU Infrastructure)

Solution 2 retains the core data lakehouse, storage, and batch processing in the AWS Cloud, while bringing AI inference, RAG embeddings, and LLM processing on-premises on local bare-metal GPU servers.






Core Components & Hybrid Separation

- Cloud Core Lakehouse (AWS):**
 - Data Storage:** AWS S3 with WORM S3 Object Lock housing Tier 0 golden records and Tier 1 raw telemetry.
 - Processing & Querying:** Amazon EMR (Spark + Sedona), Amazon Athena / Trino, and AWS MWAA Airflow running in the cloud.
 - Catalog & Serving:** AWS Glue / Apache Polaris catalog and Amazon Aurora Postgres + PostGIS serving layer.
- On-Premises Dedicated AI Infrastructure:**
 - Bare-Metal GPU Compute Nodes:** Local enterprise servers equipped with NVIDIA H100, A100, or L40S GPUs running in an on-premises data centre (e.g. Cyberjaya).
 - Local AI Inference Engine:** Ollama, vLLM, or RAGFlow running in rootless Podman containers or local Kubernetes, serving local LLMs (Qwen, Llama 3, DeepSeek) for natural language processing, vector embedding generation, and dynamic hazard prediction.
 - Local Vector Caching:** Valkey / Qdrant instance for ultra-fast local vector search.
- Hybrid Secure Connectivity & MCP Integration:**
 - AWS Direct Connect / Dedicated IPsec VPN:** AWS Direct Connect provides a dedicated private physical circuit (unencrypted by default; Layer 2 MACsec encryption is available for dedicated 10 Gbps, 100 Gbps, and 400 Gbps connections at selected Direct Connect locations, whereas 1 Gbps connections or non-MACsec locations require a Layer 3 IPsec VPN overlay when encryption in transit is required) or a dedicated IPsec VPN tunnel connecting the on-premises GPU cluster directly to the AWS VPC.
 - On-Premises MCP Servers:** MCP servers (mcp-trino-query-gen, mcp-catalog-context) run locally inside the on-prem GPU cluster.
 - Read-Only Cloud Access Gateways:** On-prem MCP tools query the Cloud Glue/Polaris Catalog and Cloud Trino engine using read-only database credentials over mTLS via Apache APISIX.
 - Data Quarantine Enforcement:** All AI inferencing occurs on-premises. Model intermediate outputs are stored on local S3-compatible storage (MinIO) configured with a 30-day S3 Lifecycle expiration rule marked as Tier 2 Scratch. AI models have zero write permissions back to Cloud Tier 0 SSoT.

Solution 3: Everything On-Premises using Proxmox VE + RKE2 + Distributed Ceph Storage

Solution 3 provides complete operational sovereignty and zero vendor lock-in by executing 100% of the platform on-premises using **Proxmox Virtual Environment (VE)**, **Rancher Kubernetes Engine 2 (RKE2)**, and a **Distributed Ceph Storage** fabric.



1. Hypervisor & Compute Infrastructure (Proxmox VE)

- Proxmox VE Cluster Sizing & Host Topology:** High-density enterprise hypervisor cluster comprising **11 physical Proxmox VE hosts** (4x AI/GPU hosts, 4x Application hosts, 3x Database/Stateful hosts) with a 100GbE data plane and 10GbE out-of-band management plane. Minimum per-host hardware capacity requirements:
 - 4x AI/GPU Compute Hosts:** Dual 32-core CPUs (e.g. AMD EPYC / Intel Xeon), 512GB RAM, 8TB NVMe storage, and 2x NVIDIA H100/A100 GPUs per physical host.
 - 4x Application Compute Hosts:** Dual 24-core CPUs, 256GB RAM, 4TB NVMe storage per physical host.
 - 3x Database / Stateful Hosts:** Dual 24-core CPUs, 256GB RAM, 4TB NVMe storage per physical host.
- VM Sizing, Anti-Affinity & N+1 Headroom:** The 14 virtualized RKE2 production nodes and 5 virtualized K3s management nodes execute as Virtual Machines (VMs) provisioned across the 11 physical PVE hosts. Strict Proxmox VE VM anti-affinity rules are enforced across physical hosts to prevent single points of failure (SPOFs) (e.g. no two RKE2/K3s control-plane VMs or Patroni DB VMs reside on the same physical PVE host). Each host tier maintains an N+1 capacity reservation (reserving 15–20% CPU/RAM headroom per tier) so that upon any single physical host failure, affected virtualized cluster VMs automatically restart on remaining physical hosts, guaranteeing high availability.
- PCIe GPU Passthrough:** Direct hardware assignment of NVIDIA H100/A100/L40S GPUs to virtualized AI worker VMs, bypassing hypervisor overhead for maximum matrix math performance.

2. Dual-Cluster Open-Source Kubernetes Architecture (RKE2 & K3s)

- Cluster A: Main Production Cluster (14 VM Nodes - RKE2):**
 - Built on Rancher Kubernetes Engine 2 (RKE2 v1.30.x series, e.g. v1.30.4+rke2r1, following upstream Kubernetes 1.30 ~12-month support window). When FIPS 140-2 compliance is required, RKE2 is deployed using its bundled Canal CNI with FIPS-validated cryptographic modules; when eBPF performance and advanced networking are required, Cilium CNI is selected.
 - 3x Control Plane / Server VM Nodes:** High-Availability etcd quorum (rke2-cp-01 to 03).
 - 4x AI / GPU Worker VM Nodes:** Executes Ollama, vLLM, RAGFlow, and Spark/Sedona distributed spatial workloads (rke2-worker-ai-01 to 04).

- **4x Application / Microservices Worker VM Nodes:** Hosts Apache NiFi, Airflow, APISIX, Keycloak, Superset, Next.js, and Trino engines (rke2-worker-app-01 to 04).
 - **3x Database & Stateful Worker VM Nodes:** Hosts PostgreSQL Patroni nodes and local Ceph OSD storage daemons (rke2-worker-db-01 to 03).
- **Cluster B: Supporting Services Cluster (5 VM Nodes - K3s):**
 - Built on lightweight K3s (v1.30.x series, e.g. v1.30.4+k3s1).
 - **3x Control Plane VM Nodes** (k3s-mgmt-01 to 03).
 - **2x Dedicated Worker Agent VM Nodes** (k3s-worker-01 to 02) hosting Prometheus, Grafana, Loki, HashiCorp Vault, and local CI/CD runners.

3. Distributed Software-Defined Storage (Ceph SDS)

- **Integrated Ceph Cluster:** Deployed via Proxmox VE / Rook-Ceph across dedicated NVMe/SSD storage pools.
- **RADOS Gateway (S3 Object Storage):** Supplies S3-compatible APIs. Buckets configured with **S3 Object Lock in Compliance Mode** (Versioning enabled) provide software-enforced WORM storage for Tier 0 SSoT records under defined retention periods (e.g., 7 years). Tested, equivalent Object Lock retention behavior is validated across AWS S3, Ceph RGW, and MinIO.
- **Storage Provisioning Types:**
 - **Ceph CSI (rbd.csi.ceph.com):** High-IOPS dynamic block storage (ReadWriteOnce / RWO) for PostgreSQL databases, vector indices, and stateful pods.
 - **CephFS (cephfs.csi.ceph.com):** Distributed filesystem (ReadWriteMany / RWX) for shared media assets and queue spools.
 - **NFS External Provisioner (nfs-subdir-external-provisioner):** Secondary ReadWriteMany (RWX) storage array for legacy shared assets.
 - **Local Path Provisioner:** Ultra-low latency node-local NVMe storage for ephemeral Tier 2 AI scratchpads, managed with an automated CronJob cleanup controller enforcing secure-deletion rules.

4. Complete On-Premises Open-Source Software Stack

- **Object Store & Catalog:** Ceph RADOS Gateway / MinIO + Apache Polaris / Project Nessie REST Catalog on RKE2.
- **Compute & Processing:** Trino Distributed MPP Query Engine + Apache Spark & Sedona on RKE2.
- **Operational Database:** High-Availability PostgreSQL 17 managed by Patroni with etcd, pg_backrest, and PostGIS extension.
- **Ingestion & Orchestration:** Apache NiFi + Apache Airflow DAGs with Data Contract CLI and OpenLineage hooks.
- **Ingress, IAM & Governance:** Apache APISIX Gateway + Keycloak OIDC IAM + OpenMetadata catalog.
- **AI & Local Inference:** Local Ollama / vLLM + RAGFlow on GPU Nodes + Containerized MCP Servers.
- **Presentation:** Next.js / React Web Application + Apache Superset with deck.gl spatial maps.

Storage Provisioner	Access Mode	Resiliency & HA Profile	Target Workloads	OS Dependencies	Performance Profile
Ceph CSI (RBD)	ReadWriteOnce (RWO)	Multi-node OSD replication, dynamic failover.	PostgreSQL / Patroni, Vector DBs, Stateful Apps.	ceph-common, rbd module	High IOPS, Low Latency, Distributed HA.
Ceph CSI (CephFS)	ReadWriteMany (RWX)	Multi-MDS HA filesystem, distributed replication.	Shared media streams, raw file attachments.	ceph-common	Medium-High IOPS, Shared Filesystem.
NFS External	ReadWriteMany (RWX)	Single NFS server (SPOF unless hardware appliance).	Shared config files, static assets.	nfs-common / nfs-utils	Low-Medium IOPS, File Locking Bottleneck.
Local Path	ReadWriteOnce (RWO)	Bound to single host disk (No node failover).	Kafka, Elasticsearch, Ephemeral Tier 2 AI Scratch.	None (Standard filesystem mount)	Maximum Raw NVMe IOPS, Sub-millisecond.

Phased Migration Strategy & Implementation Roadmap

Phase 1: Foundation & Dual Ingestion (Months 1–3)
└─ Deploy Ceph/MinIO with S3 Object Lock; stand up Polaris and OpenMetadata
└─ Deploy Apache NiFi to mirror external feeds (MSD, GSTRN) into object storage
└─ Run mirrored ingestion alongside legacy systems without operational impact

Phase 2: Compute Modernization & Data Contract Enforcement (Months 4–6)
└─ Deploy Trino, Apache Spark, and Apache Sedona compute clusters
└─ Formalize ODCS v3.1.0 data contracts across all six departmental domains
└─ Execute parallel Spark jobs to migrate legacy data into Apache Iceberg format
└─ Integrate OpenLineage runtime emission across Airflow DAGs

Phase 3: AI Operational Sandboxing & MCP Deployment (Months 7–9)
└─ Deploy containerized MCP servers with read-only database connections
└─ Configure isolated Tier 2 scratch storage with automated 30-day TTLs
└─ Restrict AI interactions to operational tooling and schema discovery
└─ Enforce cryptographic verification gates for promoting data to Tier 0

Phase 4: Presentation Cutover & Legacy Decommissioning (Months 10–12)
└─ Deploy Apache Superset and migrate Tableau dashboards to deck.gl views
└─ Launch containerized Next.js web application behind APISIX and Keycloak
└─ Validate 30-day parallel operational run across all six departmental domains
└─ Decommission Hadoop, GlusterFS, MariaDB, WildFly, Joomla, and Tableau

Implementation Phase	Target Legacy Subsystems	Modern Open-Source Replacements	Operational Risk Factors	Technical Mitigation & Fallback Procedures
Phase 1: Foundation & Dual Ingestion (Months 1–3)	Point-to-point SFTP, local folder shares, manual email ingestion.	Ceph / MinIO (WORM), Apache Polaris, OpenMetadata, Apache NiFi.	Upstream format modifications during mirroring; network saturation at boundary.	Operate NiFi in non-intrusive listening mode; legacy production paths remain authoritative; allocate isolated NICs.
Phase 2: Compute & Contract Migration (Months 4–6)	Hadoop HDFS, GlusterFS, MariaDB/Postgres project stores, WildFly.	Apache Iceberg, Trino, Apache Spark + Sedona, Data Contract CLI.	Data truncation or encoding errors during historical Iceberg Parquet conversions.	Execute automated row-count and partition checksum verifications; preserve raw source stores in read-only mode across all 6 domains.
Phase 3: AI Sandboxing & MCP Deploy (Months 7–9)	Unmonitored administrative scripts, ad-hoc Python workflows.	Model Context Protocol servers, Keycloak IAM, Tier 2 Sandbox.	Over-privileged AI agents attempting schema adjustments or unauthorized queries.	Enforce read-only database connections; block write verbs at API gateway; isolate MCP network routes.
Phase 4: Cutover & Decommissioning (Months 10–12)	Tableau Server cluster, legacy CMS, custom application portals.	Apache Superset (deck.gl), Next.js / React portal, Apache APISIX.	Discrepancies between Tableau and Superset spatial maps; user resistance to new UI.	Run 30-day side-by-side verification runs across all 6 departmental domains; validate geospatial rendering against PostGIS base layers; conduct agency training.

Strategic Architectural Conclusions

The modernization of the Big Data Analytics platform establishes a scalable, 100% open-source data lakehouse designed to serve as an authoritative Single Source of Truth for enterprise data management.

Key Architectural Takeaways

- Complete Compute-Storage Decoupling:** Moving to S3-compatible object storage (Ceph/MinIO/AWS S3) and distributed engines (Trino/Spark/Sedona) eliminates Hadoop NameNode memory limits, POSIX file locking overhead, and relational database sprawl.
- Absolute Data Integrity & Immutability:** Apache Iceberg delivers serialized ACID transactions, snapshot time-travel, and schema evolution. Software-enforced S3 Object Lock in Compliance Mode guarantees WORM retention for Tier 0 golden records.
- Strict Human-to-AI Quarantine:** The architecture enforces a rigid separation between certified human knowledge (Tier 0) and AI outputs (Tier 2). Artificial intelligence operates strictly as operational tooling via Model Context Protocol (MCP) servers with read-only database roles, preventing automated AI pollution of master datasets.
- Flexible Deployment Topologies:** Organizations can deploy via **Solution 1 (All in Cloud on AWS)**, **Solution 2 (Hybrid Cloud Core + On-Prem GPU Cluster)**, or **Solution 3 (100% Sovereign On-Premises via Proxmox VE, RKE2, and Ceph SDS)** without changing data models or application interfaces.
- Zero Vendor Lock-In:** Replacing proprietary tools (Tableau, proprietary CMS) with Apache Superset, Next.js, Keycloak, and Apache APISIX eliminates recurring seat licensing, ensures compliance with open standards (Bitol ODCS, OpenLineage, OGC/ISO 19115:2003), and provides full operational autonomy.

Amazon Machine Images (AMI) Design Guide

#aws #cloud #architecture #vpc #alb #asg #rds #efs #postgresql #ragflow #langfuse #compliance

Amazon Machine Images (AMI) Design & Hardening Guide

This document describes the **Amazon Machine Images (AMI)** strategy for our Auto Scaling Groups (ASGs). It outlines why custom, pre-baked AMIs are essential for stateless, rapid autoscaling, how to automate AMI construction using HashiCorp Packer and Ansible, and how the **ASIMP** auditing and hardening framework is integrated into the baking pipeline.

1. Why Pre-Baked AMIs are Critical for Auto Scaling

When an Auto Scaling Group (ASG) detects a surge in CPU utilization or network traffic, it immediately spins up new EC2 instances to share the workload.

If your application depends on a **just-in-time bootstrapping** approach (e.g., executing shell scripts in `user_data` that fetch package updates, download heavy language model weights, install dependencies, and compile binaries at boot-time), you introduce several critical failure modes:

- Unacceptable Latency:** Bootstrapping heavy applications can take 10 to 15 minutes. By the time the instance is healthy and added to the ALB target group, the peak load may have already overwhelmed the existing fleet.
- External Dependency Failure:** If an external package repository (e.g., Ubuntu APT servers, PyPI, or GitHub) experiences temporary downtime, the newly launched instance will fail to bootstrap, leaving your cluster unable to scale.
- Inconsistent Environments:** Multi-instance scaling can result in nodes running slightly different versions of minor packages if updates occurred between launch events.


The Immutable Infrastructure Paradigm

To solve this, we employ **Immutable Infrastructure** via custom **pre-baked AMIs**.

All software, libraries, configuration files, and system security controls are installed and validated *during the image building phase*. When the ASG scales out, the instance launches from the pre-baked AMI and transitions to a ready-to-run state in **under 45 seconds**, with zero runtime network dependency for configuration.

2. AMI Requirements Matrix per ASG Tier

To maintain clean separation of concerns and avoid bloated monolithic images, we define three specific, pre-baked AMIs built on top of **Ubuntu 26.04 LTS**:



A. Frontend Nginx Web Server AMI (ami-frontend-*)

- Purpose:** Serves as the high-throughput reverse proxy and static content host.
- Pre-Baked Software:**
 - Nginx (optimized with HTTP/2 and custom buffer limits).
 - AWS CloudWatch Agent (pre-configured to stream Nginx access and error logs).
 - ASIMP Hardening:** Restricted SSH configuration, system account lockdowns, and local firewall parameters.

B. Backend App Tier AMI (ami-backend-*)

- Purpose:** Runs the Backend, Document Management System (DMS), and Model Context Protocol (MCP) APIs.
- Pre-Baked Software:**
 - Runtime Environments: Node.js (LTS), Python 3.12+, and package managers (npm, pip).
 - Git client and database connectors (for AWS RDS PostgreSQL communication).
 - Application source code and systemd service handlers.
 - ASIMP Hardening:** Deep process limits (limits.conf), disabling of unused filesystem modules, and standard secure SSH ciphers.

C. AI Tier RAGFlow + LangFuse AMI (ami-ai-*)

- Purpose:** Executes computation-heavy RAG processing, embeddings generation, and observability flows.
- Pre-Baked Software:**
 - Docker Engine and Docker Compose (highly optimized for Graviton/ARM64 architectures).
 - Local caching directories mounted and configured for Hugging Face and SentenceTransformers model weights (accelerates initialization).
 - Container images pre-pulled and cached locally inside Docker storage.
 - ASIMP Hardening:** Kernel parameter tuning (sysctl.conf), container isolation profiles, and CIS Level 2 compliance benchmarks.

3. AMI Build Pipeline with Packer & ASIMP Hardening

We automate our AMI builds using **HashiCorp Packer** combined with **ASIMP (Ansible System Integrity Management Platform)** as the primary configuration provisioner. This guarantees that every pre-baked image is audited and hardened against standard security benchmarks prior to distribution.

The Immutable Build Workflow

1. **Initialize:** Packer spins up a temporary EC2 instance in a private build subnet using the base Canonical Ubuntu 26.04 LTS AMI.
2. **Provision:** Ansible executes the ASIMP playbooks against the temporary instance.
3. **Auditing (Pre-Harden & Post-Harden):**
 - ASIMP performs an initial security baseline scan using **OpenSCAP** (CIS Level 2) and **Lynis**.
 - ASIMP applies automated upgrades, validates system package integrity via debsums, configures secure SSH protocols (Dev-Sec), and implements fine-grained kernel modifications.
 - ASIMP re-runs the security scans, computes the comparative score improvement, logs metrics in `/var/log/asimp-baseline-scores.json`, and outputs visual HTML reports to `/var/log/openscap-after-report.html`.
4. **Bake:** Packer stops the temporary instance, takes a snapshot of the root EBS volume, registers a new AMI (`ami-frontend-yyyy-mm-dd`), and terminates the builder instance.

Sample Packer Configuration File (`template.pkr.hcl`)

Below is the Packer configuration used to build the hardened Ubuntu 26.04 LTS backend AMI:

```
packer {
  required_plugins {
    amazon = {
      version = ">= 1.2.0"
      source  = "github.com/hashicorp/amazon"
    }
  }
}

source "amazon-ebs" "ubuntu-hardened" {
  ami_name      = "asimp-ubuntu-26.04-backend-"
  instance_type = "t4g.medium" # Build on high-efficiency Graviton
  region        = "ap-southeast-5"

  source_ami_filter {
    filters = {
      name               = "ubuntu/images/hvm-ssd-gp3/ubuntu-noble-24.04-*--server-*" # Fallback / target filter for Canonical Ubuntu LTS
      root-device-type   = "ebs"
      virtualization-type = "hvm"
      architecture      = "arm64"
    }
    most_recent = true
    owners      = ["099720109477"] # Canonical
  }
  ssh_username = "ubuntu"
}

build {
  sources = ["source.amazon-ebs.ubuntu-hardened"]

  # Run ASIMP Hardening using Ansible
  provisioner "ansible" {
    playbook_file = "../asimp/play.yml"
    user          = "ubuntu"
    use_extra_vars = true
    extra_vars = {
      is_packer_build = true
    }
  }

  # Verify ASIMP reports are present before finalizing image
  provisioner "shell" {
    inline = [
      "echo 'Verifying ASIMP Hardening Reports...'",
      "ls -l /var/log/asimp-baseline-scores.json",
      "ls -l /var/log/openscap-after-report.html"
    ]
  }
}
```

4. Parameterizing & Referencing Baked AMIs in OpenTofu

Once Packer successfully registers the custom AMIs, they can be referenced inside your OpenTofu configurations dynamically or passed as static overrides.

Option A: Dynamic AMI Discovery (Recommended)

This approach dynamically queries AWS for the latest hardened AMI matching your tag patterns, ensuring that newly deployed Auto Scaling groups always inherit the latest security patches.

```
# Look up latest pre-baked Backend AMI
data "aws_ami" "latest_harden_backend" {
  most_recent = true
  owners      = ["self"] # Search within your AWS account

  filter {
    name   = "name"
    values = ["asimp-ubuntu-26.04-backend-*"]
  }

  filter {
    name   = "virtualization-type"
    values = ["hvm"]
  }

  filter {
    name   = "architecture"
    values = ["arm64"]
  }
}
```

You then supply this dynamic ID to your ASG launch template:

```
resource "aws_launch_template" "backend" {
  name_prefix = "backend-launch-template-"
  image_id    = data.aws_ami.latest_hardenened_backend.id
  instance_type = "t4g.xlarge"
  # ... remaining options
}
```

Option B: Static AMI Parameter Override

For high-control staging/production environments where AMI updates must undergo rigid validation before scaling, define the image IDs as input variables:


```
# In variables.tf
variable "backend_ami_id" {
  description = "Pre-baked ASIMP Hardened Ubuntu 26.04 LTS AMI for Backend"
  type        = string
  default     = "" # If empty, config falls back to dynamic SSM or data queries
}

# In main.tf
locals {
  selected_backend_ami = var.backend_ami_id != "" ? var.backend_ami_id : data.aws_ami.latest_hardenened_backend.id
}
```

This ensures full agility while retaining rigid compliance validation controls.

5. Standalone EC2 Instances: Staging and Baking Lifecycle

To achieve complete alignment between active development and the immutable production ASGs, our design utilizes **dedicated Standalone EC2 Instances** as live staging, testing, and pre-baking nodes for each of the three logical application tiers.



Environmental Alignment with Shared Backends

To prevent runtime configuration errors and bootstrap failures (such as database connection timeout or missing mount folders), the standalone instances are connected to the identical shared infrastructure as their matching ASGs:

- **Frontend Standalone:** Accesses the identical Amazon S3 static buckets to pull and verify Nginx page templates and asset configuration.
- **Backend Standalone:** Connects directly to the live **Multi-AZ RDS PostgreSQL** database (allowing safe database schema migration and endpoint verification) and **Amazon S3** (for DMS and document uploads).
- **AI Tier Standalone:** Mounts the identical multi-AZ **Amazon EFS** volume (facilitating direct Hugging Face and SentenceTransformers model caching). Developers download and cache model weights directly onto EFS from this standalone instance, which ensures newly-bootstrapped ASG nodes can access cached model weights instantly.

The Staging-to-AMI Lifecycle Workflow

1. **Active Staging & Application Testing:**
 - Developers perform code updates, install package dependencies, and modify system configuration on the standalone instance.
2. **ASIMP Auditing & Hardening Compliance:**
 - Developers execute the **ASIMP** auditing framework directly on the standalone instance. ASIMP executes OpenSCAP and Lynis compliance audits, applies Ubuntu security hardening configurations, checks package integrity, and produces HTML verification scorecards at `/var/log/openscap-after-report.html`.
3. **AMI Capture & Registration:**
 - Once compliance and functional tests pass successfully, the standalone instance's root volume is frozen.
 - An AMI is registered directly from the standalone instance using Packer or an automated AWS Systems Manager (SSM) backup document:

```
aws ec2 create-image \
  --instance-id i-0xxxxxxxstandalone \
  --name "asimp-hardened-backend-$(date +%F)" \
  --no-reboot
```
4. **ASG Deployment & Rolling Refresh:**
 - The registered AMI ID is supplied to OpenTofu as a variable override or auto-discovered dynamically.
 - An **ASG Instance Refresh** is triggered, rolling out the new pre-validated, fully-configured image across the active ASG nodes with zero downtime.

Route 53 & DNS Troubleshooting


[#aws](#) [#cloud](#) [#architecture](#) [#vpc](#) [#alb](#) [#asg](#) [#rds](#) [#waf](#) [#route53](#) [#dns](#) [#ssl](#) [#acm](#) [#disaster-recovery](#) [#postgresql](#)

Amazon Route 53 Custom Domains & DNS Resolution

This document details how **Amazon Route 53** is integrated into our secure 3-Tier AWS Architecture for domain routing and SSL/TLS certificate registration. It also provides an in-depth research analysis of common **DNS resolution failures inside Auto Scaling Groups (ASGs)** (especially in private subnets) and step-by-step guidance to prevent and troubleshoot them.

1. Using Route 53 for Custom Domains

Our modular OpenTofu configuration automates the mapping of public domain names to our WAFv2-protected Application Load Balancer (ALB).



High-Level Setup Steps

To add your domain name to this deployment, complete the following workflow:

- Purchase or Delegate Domain:** Register a domain (e.g., linuxmalaysia.com) via Amazon Route 53 or any third-party registrar (GoDaddy, Namecheap, etc.).
- Retrieve Route 53 Name Servers:** When our OpenTofu configuration runs with `enable_route53 = true`, it creates a public Route 53 Hosted Zone and outputs a list of four Name Servers (NS):
`tofu output route53_name_servers`
- Configure Registrar NS Records:** Log into your domain registrar and replace your default name servers with the four Route 53 name servers generated by OpenTofu. This delegates DNS authority to AWS.
- ALB Routing (ALIAS Records):** To avoid charge fees and keep IP mappings dynamic, our module creates an **Alias A Record** (e.g., app.linuxmalaysia.com -> production-alb-123456789.ap-southeast-5.elb.amazonaws.com). AWS handles updates to ALB IPs automatically without modifying the DNS record.

2. HTTPS & SSL/TLS Configuration with AWS Certificate Manager (ACM)

To encrypt web traffic in transit (HTTPS on Port 443), we pair Route 53 with AWS Certificate Manager (ACM):

ACM SSL Provisioning Workflow

- Request Certificate:** Request a public wildcard certificate (e.g., *.linuxmalaysia.com) in ACM.
- DNS Validation:** ACM requires proof of domain ownership. Choose **DNS Validation**. ACM will generate a set of CNAME records.
- Publish Validation Records:** Add these CNAME records to your Route 53 Hosted Zone. This can be automated in OpenTofu using the `aws_route53_record` resource:

```
resource "aws_route53_record" "cert_validation" {
  for_each = {
    for dvo in aws_acm_certificate.cert.domain_validation_options : dvo.domain_name => {
      name    = dvo.resource_record_name
      record  = dvo.resource_record_value
      type    = dvo.resource_record_type
    }
  }

  allow_overwrite = true
  name            = each.value.name
  records         = [each.value.record]
  ttl             = 60
  type            = each.value.type
  zone_id        = var.hosted_zone_id
}
```

- Attach SSL to ALB Listener:** Once validated, ACM marks the certificate as **Issued**. You then update your ALB Listener in `terraform/modules/alb/main.tf` to use HTTPS (Port 443) and reference the certificate ARN:

```
resource "aws_lb_listener" "https" {
  load_balancer_arn = aws_lb.main.arn
  port              = 443
  protocol          = "HTTPS"
  ssl_policy        = "ELBSecurityPolicy-TLS13-1-2-2021-06"
  certificate_arn    = var.acm_certificate_arn

  default_action {
    type = "forward"
    target_group_arn = aws_lb_target_group.asg_tg.arn
  }
}
```

3. In-Depth Research: ASG DNS Resolution Failures & Mitigation

Instances launched in an Auto Scaling Group (ASG) inside secure private subnets frequently run into **DNS resolution issues**. These failures break outbound API integrations (e.g., sending WhatsApp notifications, accessing Hugging Face model weight repositories, or making external OAuth callbacks).

Below is our technical research detailing the primary causes, mechanics, and concrete solutions for these DNS problems:

Cause A: The Nginx Upstream DNS Caching Bottleneck (The Silent 502)

This is the single most common failure in 3-Tier configurations where Nginx acts as a frontend or API reverse proxy routing traffic to dynamic backends.

- **The Problem:** By default, Nginx resolves domain names (such as your backend ALB DNS name or internal endpoints) **only once during startup** and caches the IP address indefinitely. However, AWS ALBs dynamically scale up and down, shifting their public and private IP addresses in response to traffic surges.
- **The Symptom:** Nginx works perfectly at first. But when the ALB scales out or replaces healthy nodes, Nginx continues sending traffic to the decommissioned IP addresses, resulting in persistent 502 Bad Gateway or Connection Timed Out errors, even though DNS is working perfectly at the OS level.
- **The Solution:** Use variables to enforce dynamic DNS lookups. Inside your Nginx config (/etc/nginx/sites-available/default), define a resolver pointing to the AWS Core VPC DNS server (169.254.169.254 or the second IP address of your VPC CIDR, e.g., 10.0.0.2), and set a short TTL:

```
server {
    listen 80;

    # Use the AWS VPC DNS Resolver with a 10-second cache threshold
    resolver 169.254.169.254 valid=10s;

    set $backend_upstream "http://app.linuxmalaysia.com";

    location /api/ {
        # Using a variable forces Nginx to respect the resolver TTL
        proxy_pass $backend_upstream;
        proxy_set_header Host $host;
        proxy_set_header X-Real-IP $remote_addr;
    }
}
```

Cause B: VPC DNS Settings Disabled (Internal DNS Failure)

- **The Problem:** AWS VPCs provide automatic resolution for internal DNS queries (e.g., looking up database endpoints like production-database.xxxxxx.ap-southeast-5.rds.amazonaws.com). This fails if DNS support features are disabled in the VPC configuration.
- **The Symptom:** Private ASG instances cannot ping or connect to the RDS PostgreSQL instance using its domain endpoint, but can connect using the raw private IP.
- **The Solution:** Ensure both VPC attributes are set to true inside your OpenTofu network code (terraform/modules/vpc/main.tf):

```
resource "aws_vpc" "main" {
  cidr_block      = var.vpc_cidr
  enable_dns_support = true # Enforces Route 53 resolver availability
  enable_dns_hostnames = true # Allocates public/private DNS hostnames
}
```

Cause C: Systemd-Resolved Configuration Issues on Ubuntu

Our target operating system, **Ubuntu 26.04 LTS**, manages name resolution via the local stub listener systemd-resolved on 127.0.0.53.

- **The Problem:** If /etc/resolv.conf is a static file or a broken symlink instead of pointing to systemd-resolved's runtime stub, the instance cannot receive DNS server allocations propagated by DHCP Options.
- **The Symptom:** Running nslookup google.com fails, and /etc/resolv.conf is empty or lacks name servers.
- **The Solution:** In your baking pipeline or user_data.sh script, force-recreate the symbolic link to point to systemd-resolved:

```
sudo ln -sf /run/systemd/resolve/stub-resolv.conf /etc/resolv.conf
sudo systemctl restart systemd-resolved
```

Cause D: Misconfigured Outbound Security Groups or NACLs

- **The Problem:** DNS resolution requires outgoing communication on **Port 53 (UDP & TCP)** to Route 53 Resolver (169.254.169.254). If outbound Security Groups or Network ACLs are overly restrictive, the DNS queries are dropped before leaving the node.
- **The Symptom:** Running dig @169.254.169.254 google.com hangs and times out.
- **The Solution:** Verify that the Security Group assigned to the ASG (asg_sg in terraform/modules/security_groups/main.tf) permits outbound port 53:

```
# Egress configuration for ASG Security Group
resource "aws_security_group_rule" "asg_dns_out" {
  type           = "egress"
  from_port      = 53
  to_port        = 53
  protocol       = "udp"
  cidr_blocks    = ["0.0.0.0/0"]
  security_group_id = aws_security_group.asg_sg.id
}
```

(Note: A default egress of 0.0.0.0/0 with protocol -1 covers all ports, including DNS, which is recommended for the application tier).

Cause E: Route 53 Resolver Rate Limiting (DNS Query Flooding)

- **The Problem:** The AWS Route 53 Resolver enforces a limit of **1024 packets per second (PPS)** per Network Interface (ENI). In a high-throughput microservices environment with hundreds of concurrent connections, application instances can exceed this limit.
- **The Symptom:** Sporadic DNS resolution timeouts that resolve themselves after a few seconds, accompanied by errors in the application logs.
- **The Solution:** Configure local DNS caching to prevent every network socket call from sending an upstream packet to AWS DNS:
 1. Ensure systemd-resolved has local caching enabled (enabled by default).
 2. Implement an on-instance cache proxy like **dnsmasq** or **unbound** for high-load workloads.
 3. Configure your applications (such as Node.js or Python) to use keep-alive connections to minimize DNS requests per API handshake.

Secure Developer Access Guide

#aws #cloud #architecture #vpc #asg #jumphost #bastion #disaster-recovery #compliance #costing

Secure Developer AWS Access Guide (SSH Jumphost)

This guide provides technical research, architectural specifications, and step-by-step instructions for developers based in our **Cyberjaya, Selangor, Malaysia** office to securely access private AWS resources (private ASG compute servers, database systems, and standalone staging instances) using a highly secure, cost-optimized SSH Jumphost (Bastion).

Architectural Research & Feasibility Analysis

The Question: “Can/Should we deploy the Jumphost outside a VPC?”

During design discussions, a proposal was raised: “Does deploying a small Amazon Linux not in a VPC give us a better and cheaper solution?”

Based on cloud architecture fundamentals and AWS networking constraints, this proposal is **both technically impossible and architecturally undesirable**. Here is why:

- 1. **VPC Requirement in AWS:** Since the retirement of AWS EC2-Classic, **every EC2 instance must run inside a Virtual Private Cloud (VPC)**. There is no longer any option to deploy a “raw” EC2 instance outside of a VPC.
- 2. **The “Separate Jumphost VPC” Anti-Pattern:** Even if the proposal is interpreted as deploying the Jumphost inside a *separate management VPC* (instead of our main application VPC), this approach is **not** cheaper or safer. It introduces severe cost and operational overhead:
 - **Transit Gateway (TGW) Costs:** To connect the separate Jumphost VPC to the private subnets of our application VPC, we would need to provision an AWS Transit Gateway. In the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)**, a TGW attachment costs **\$0.05 per hour** (~\$36.50/mo per VPC connection) plus a **\$0.02 per GB** data processing fee. This alone is **3.5x more expensive** than running the entire jumphost itself!
 - **VPC Peering Complexity:** Alternatively, establishing a VPC Peering Connection requires updating routing tables in both VPCs, dealing with inter-AZ data transfer fees (\$0.01/GB in each direction), and managing dual security boundaries, which increases the likelihood of human error.
 - **Resource Duplication:** A separate VPC would require its own Internet Gateway, public subnets, and routing infrastructure, duplicating base costs.

The Optimal Solution: Public Subnet + Strict IP Whitelisting

The **best, safest, and cheapest** approach is to deploy a small `t4g.micro` instance (equipped with energy-efficient ARM64 Graviton processors) inside the **existing public subnet** of our main application VPC:

- **No Extra AWS Networking Cost:** It reuses the existing public Internet Gateway and routing infrastructure.
- **Strict Ingress Firewalling (Zero-Trust):** The Jumphost’s Security Group permits inbound TCP port 22 connections **exclusively** from the public CIDR of our **Cyberjaya, Selangor, Malaysia office** (configurable via `jumphost_allowed_ssh_cidr`). All other internet traffic is dropped silently at the AWS Hypervisor layer.
- **Deep Private Isolation:** The private compute nodes (ASG) and standalone dev environments have **no public IP addresses** and are completely unreachable from the public internet. They only permit SSH (port 22) ingress **exclusively from the Jumphost’s Security Group**.
- **Minimal Cost Impact:** Running a `t4g.micro` with a 15GB gp3 EBS root volume and an Elastic IP costs only **~\$10.98 USD / month** (~RM 49.41 MYR/mo) on-demand, which can be further optimized via Savings Plans.

Technical Specifications of the Jumphost

Our OpenTofu setup automates this configuration. The Jumphost is deployed with:

- **Base OS:** Ubuntu 24.04 LTS (allows 100% compliance with the ASIMP hardening framework).
- **Network Sizing:** `t4g.micro` (ARM64/Graviton), costing **\$0.0084/hour**.
- **Storage:** 15GB gp3 SSD, encrypted at rest (gp3 storage at \$0.08/GB-mo).
- **Static IP:** Associated with an AWS Elastic IP (EIP) so developers have a stable endpoint that does not change upon reboot.

OS Hardening & Auditing via ASIMP

To protect our access gateway, the Jumphost must be hardened using **ASIMP (Ansible System Integrity Management Platform)**. ASIMP implements a “Measure, Harden, Re-Measure” security flow.

Key SSH Hardening Policies Configured via ASIMP:

- 1. **Disable Password-Based Authentication:** Only cryptographic keys are allowed.
- 2. **Disable Root Logins:** Root SSH is strictly disabled (`PermitRootLogin no`). Developers must connect using their own unprivileged user (e.g., `ubuntu` or standard service accounts) and escalate using `sudo` if authorized.
- 3. **Restricted Cryptographic Cipher Suites:** Enforces modern, secure key exchange algorithms and ciphers, completely blocking obsolete/weak ones (e.g., MD5, SHA-1, 3DES, Blowfish).
 - **Allowed Key Exchange (KEX):** `curve25519-sha256, diffie-hellman-group16-sha512`
 - **Allowed MACs:** `hmac-sha2-512-etm@openssh.com`
 - **Allowed Ciphers:** `chacha20-poly1305@openssh.com, aes256-gcm@openssh.com`
- 4. **Active Connection Auditing:** OpenSCAP and Lynis run scheduled audits to verify CIS Level 2 benchmark compliance, writing audit scorecard logs to `/var/log/asimp-baseline-scores.json`.

Critical SSH Private Key Security Guidelines

The private SSH key represents the developer’s identity and grants administrative access. If a private key is compromised, the entire AWS infrastructure is at risk. All developers must follow these strict guidelines:

1. Always Use a Strong Passphrase

Never generate a passwordless SSH key. Always encrypt the private key with a strong passphrase. This ensures that even if the private key file is stolen (e.g., from a physical laptop loss or malware infection), the attacker cannot use it without decrypting it first.

2. Never Share Private Keys

A private key is private. It must **never** be shared via Slack, Email, WhatsApp, or uploaded to any shared network. Under no circumstances should multiple developers share the same SSH key.

3. Never Check Keys into Version Control

Ensure that private keys (such as .pem, .id_ed25519, .ppk) are never committed to Git repositories. Add a global rule in your .gitignore file:

```
# Ignore SSH Keys
*.pem
*.key
id_ed25519
id_rsa
*.ppk
```

4. Apply Strict File Permissions (OS-Specific)

On Linux / macOS (Unix File Permissions)

SSH clients will refuse to use private keys that are too openly readable. The .ssh directory and key files must be restricted:

```
# Secure the SSH directory (read/write/search for owner only)
chmod 700 ~/.ssh

# Secure the private key file (read/write for owner only, no access for anyone else)
chmod 600 ~/.ssh/id_ed25519

# Secure public keys (read/write for owner, read-only for others)
chmod 644 ~/.ssh/id_ed25519.pub
```

On Windows (NTFS Access Control Lists)

Windows does not use standard Unix permissions, but the OpenSSH client on Windows still enforces strict security. To secure a private key (id_ed25519 or key.pem) on Windows:

Option A: Using PowerShell (Fastest)

Run the following commands in PowerShell to strip inheritance and grant permissions exclusively to the current logged-in user:

```
# Store key path in a variable
$keyPath = "$Home\.ssh\id_ed25519"

# Strip inheriting permissions and remove all group permissions
icacls $keyPath /inheritance:r


# Grant full control exclusively to the current logged-in user
$currentUser = [System.Security.Principal.WindowsIdentity]::GetCurrent().Name
icacls $keyPath /grant:r "${$currentUser}:F"
```

Option B: Using Windows File Explorer (GUI)

- 1. Right-click the private key file (id_ed25519) and select **Properties**.
- 2. Go to the **Security** tab, and click **Advanced**.
- 3. Click **Disable inheritance** and select “Remove all inherited permissions from this object”.
- 4. Click **Add**, click “Select a principal”, type your Windows username, and click **Check Names**.
- 5. Set permissions to **Full control** and click **OK** -> **Apply** -> **OK**.

Developer Access Step-by-Step Instructions

To reach private AWS resources (e.g., Standalone Instance 10.0.10.15), developers use the Jumphost as a transit bridge.



Below are instructions for macOS, Linux, and Windows platforms.

1. macOS Developer Setup

macOS has native SSH client support. We recommend generating an **Ed25519** key, which is modern, lightweight, and highly secure.

Step 1: Generate your SSH Key Pair

Open Terminal and run:

```
ssh-keygen -t ed25519 -C "yourname@office.com"
```

Note: Specify a secure passphrase when prompted.

Step 2: Register your Public Key on AWS

Provide your public key (~/.ssh/id_ed25519.pub) to your AWS administrator. They will register it on the Jumphost and private instances.

Step 3: Configure your local SSH config file

To avoid typing complex commands every time, create or edit your SSH configuration file (~/.ssh/config):

```
nano ~/.ssh/config
```

Insert the following block (replace 54.x.y.z with the Jumphost's actual Elastic IP address and your -key with your private key name):

```
# Common SSH settings
Host *
    AddKeysToAgent yes
    UseKeychain yes
    IdentityFile ~/.ssh/id_ed25519

# The SSH Jumphost
Host aws-jumphost
    HostName 54.x.y.z
    User ubuntu

# Private Standalone / ASG Application servers
Host aws-private-*
    # Use wildcard mapping to route any private VPC IP via Jumphost
    HostName 10.0.*
    User ubuntu
    ProxyJump aws-jumphost
```

Save and close (Ctrl+O then Ctrl+X).

Step 4: Access your Private Instances

Now, you can SSH into any private instance inside the VPC using a single, secure command:

```
# Connect directly to a private instance (macOS routes this transparently via Jumphost)
ssh aws-private-10.0.10.15
```

2. Linux (Ubuntu) Developer Setup

Linux systems use OpenSSH. The configuration is almost identical to macOS.

Step 1: Generate your SSH Key Pair

```
ssh-keygen -t ed25519 -C "yourname@office.com"
```

Step 2: Configure your SSH config file

Create or edit your local SSH configuration file:

```
nano ~/.ssh/config
```

Add the following configuration block:

```
Host *
    AddKeysToAgent yes
    IdentityFile ~/.ssh/id_ed25519

Host aws-jumphost
    HostName 54.x.y.z
    User ubuntu

Host aws-private-*
    HostName 10.0.*
    User ubuntu
    ProxyJump aws-jumphost
```

Step 3: Start SSH Agent & Load Key

Add these commands to your terminal profile (~/.bashrc or ~/.zshrc) to automate loading:

```
eval "$(ssh-agent -s)"
ssh-add ~/.ssh/id_ed25519
```

Step 4: Access Private Instances

```
ssh aws-private-10.0.10.15
```

3. Windows Developer Setup

Windows developers can use **PowerShell (with native OpenSSH)** or **PuTTY**.

Method A: Using PowerShell & Native OpenSSH (Recommended)

Modern Windows 10/11 has OpenSSH Client installed by default.

Step 1: Generate your SSH Key Pair

Open PowerShell and run:

```
ssh-keygen -t ed25519 -C "yourname@office.com"
```

Step 2: Enable and configure Windows SSH Agent

Start the Windows SSH Agent service (requires running PowerShell as Administrator once):

```
# Set Startup Type to Automatic
Set-Service -Name ssh-agent -StartupType Automatic
```

```
# Start the Service
Start-Service ssh-agent
```

```
# Load your Private Key into the Agent
ssh-add $Home\.ssh\id_ed25519
```

Step 3: Configure your local SSH Config File

In PowerShell, edit or create the SSH config file:

```
notepad $Home\.ssh\config
```

Add the following configuration block (replace 54.x.y.z with the Jumphost's Elastic IP):

```
Host *
  AddKeysToAgent yes
  IdentityFile ~/.ssh/id_ed25519
```

```
Host aws-jumphost
  HostName 54.x.y.z
  User ubuntu
```

```
Host aws-private-*
  HostName 10.0.*
  User ubuntu
  ProxyJump aws-jumphost
```

Step 4: Secure Key Permissions

As shown in the Key Security section, secure your private key:

```
$keyPath = "$Home\.ssh\id_ed25519"
icacls $keyPath /inheritance:r
$currentUser = [System.Security.Principal.WindowsIdentity]::GetCurrent().Name
icacls $keyPath /grant:r "$currentUser:F"
```

Step 5: Connect directly

Now, connect securely to your private staging servers via PowerShell:

```
ssh aws-private-10.0.10.15
```

Method B: Using PuTTY (GUI-Based SSH Client)

If you prefer using PuTTY, follow these steps to set up secure tunneling.

Step 1: Generate Key using PuTTYgen

1. Download and open **PuTTYgen**.
2. Select **Ed25519** (or RSA 4096) under "Parameters -> Type of key to generate".
3. Click **Generate** and move your mouse randomly over the blank area.
4. Enter a strong **Key passphrase**.
5. Click **Save private key** and save it as id_ed25519.ppk inside a secure folder (e.g., C:\Users\username\.ssh\).
6. Copy the text from the top box ("Public key for pasting into OpenSSH authorized_keys file") and send it to your AWS administrator.

Step 2: Secure PPK file permissions

1. Right-click your id_ed25519.ppk file and click **Properties**.
2. Go to **Security -> Advanced**.
3. Click **Disable inheritance** -> select "Remove all inherited permissions".
4. Add your personal Windows user principal with **Full Control**.

Step 3: Configure PuTTY for Proxy SSH Session (ProxyCommand/plink)

To establish multi-hop SSH tunnels in PuTTY:

1. Open **PuTTY**.
2. In the **Session** category (on the left):
 - **Host Name (or IP address)**: Type the private IP of the target server (e.g., 10.0.10.15).
 - **Port**: 22.
3. In the **Connection -> SSH -> Auth -> Credentials** category:
 - Browse and select your id_ed25519.ppk private key.
4. In the **Connection -> Proxy** category:
 - **Proxy type**: Select **Local**.
 - **Proxy hostname**: Type the Jumphost's Elastic IP (e.g., 54.x.y.z).
 - **Port**: 22.
 - **Telnet command, or local proxy command**: Type the following (replace paths with your actual PuTTY installation and key paths):

```
plink -batch -v ubuntu@%proxyhost -i "C:\Users\YOUR_USERNAME\.ssh\id_ed25519.ppk" -nc %host:%port
```
5. Return to the **Session** category:
 - Type a name under "Saved Sessions" (e.g., AWS-Private-10.0.10.15).
 - Click **Save**.
6. Click **Open**. PuTTY will connect to the Jumphost first, proxy the connection, and open the private server terminal transparently!

Method C: Alternatively, using Port Forwarding (SSH Tunnel) via PuTTY

If you prefer standard SSH Port Forwarding:

1. Open **PuTTY** and configure a connection to the Jumphost:
 - **Host Name (or IP address)**: Jumphost Elastic IP (54.x.y.z).
 - **User**: ubuntu.
2. Go to **Connection -> SSH -> Tunnels**.
3. In "Add new forwarded port":
 - **Source port**: Type a local random port (e.g., 9022).
 - **Destination**: Type the private IP of your staging server and SSH port (e.g., 10.0.10.15:22).
 - Click **Add**.
4. Go back to the **Session** category, save it as AWS-Jumphost-Tunnel, and click **Open** (enter key passphrase when prompted).
5. Open a **second** PuTTY window:
 - **Host Name (or IP address)**: 127.0.0.1.
 - **Port**: 9022.
 - Select your private key under **Connection -> SSH -> Auth -> Credentials** and open!

Hybrid Cloud Integration & Costing Guide

#aws #cloud #architecture #vpc #alb #waf #disaster-recovery #postgresql #costing

Hybrid Cloud Integration Guide: AWS & On-Premises

1. Executive Summary & Context

To support high-performance operations while maintaining strict cost-efficiency, organizations must determine the most practical strategy for connecting on-premises infrastructure (such as local datacenters, legacy compute, or private office environments) with an AWS-native deployment.


While AWS provides several enterprise-grade physical and virtual private network connectivity options, their upfront, monthly baseline, and data processing charges can be prohibitive for budget-conscious initiatives. Consequently, this guide presents a dual-layered architectural paradigm:

- 1. **Low-Cost, High-Flexibility Pathways (API and MCP):** The primary recommended choices for agile integrations, utilizing secure HTTP RESTful architectures and the newly released **Model Context Protocol (MCP)** for AI agentic workflows.
- 2. **Official AWS Hybrid Options (VPN, Direct Connect, and Transit Gateway):** Enterprise-standard private networking pipelines offered for comparative purposes to facilitate corporate security reviews and management approval.

All cost models are calibrated for the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)** and estimated in both **United States Dollars (USD)** and **Malaysian Ringgit (MYR)**, utilizing a baseline conversion exchange rate of **1 USD ≈ 4.50 MYR**.

2. Low-Cost Option 1: Secure API-Based Connections

API-based hybrid integration leverages the public internet to transmit request/response payloads between AWS compute nodes (such as our Auto Scaling Groups or Standalone EC2 instances) and on-premises service endpoints. Rather than establishing a persistent, protocol-level network tunnel, communication is transactional, synchronous, or event-driven.



Architectural Implementation & Security

- **Authentication & Authorization:** Connections are secured via OAuth 2.0 Bearer tokens, Amazon Cognito, API Keys, or **Mutual TLS (mTLS)** supported natively by Amazon API Gateway.
- **Network Security & Whitelisting:** Downstream on-premises endpoints are configured to whitelist only the Elastic IP addresses of the AWS NAT Gateway. Inbound AWS traffic to the Application Load Balancer (ALB) is whitelisted at the AWS WAFv2 level or restricted to on-premises gateway IPs.
- **Protocol Support:** Standard REST (HTTPS/JSON) for point-to-point RPCs, or WebSockets (WSS) via Amazon API Gateway for persistent, low-latency bi-directional event streaming.

Financial Breakdown


API-based connections have **no provisioned hourly costs** for physical line terminations. The billing is entirely pay-as-you-go based on data processing and network transit:

- **Amazon API Gateway REST APIs:** \$3.50 per million requests (up to 333 million, with sliding volume discounts).
- **AWS NAT Gateway Charges:** \$0.045 per hour base charge (\$32.85/month) + \$0.045 per GB of outbound data processed.
- **AWS Outbound Data Transfer (Internet Egress):** The first 100 GB per month is free. Beyond that, data transferred from AWS to the internet in Malaysia (ap-southeast-5) is charged at **\$0.09 per GB**. On-premises data transfer into AWS is entirely free.

3. Low-Cost Option 2: AI-Native MCP-Based Connections

The **Model Context Protocol (MCP)** is an open-standard protocol designed to link Large Language Models (LLMs) securely with external data sources, content repositories, and enterprise applications. Instead of building bespoke API integrations for AI tools, MCP standardizes how AI agents discover and execute tools or fetch contextual files.

In December 2025, AWS introduced native **MCP Proxy Support in Amazon API Gateway**, integrated directly with **Amazon Bedrock AgentCore**. This allows organizations to transform on-premises applications, custom monitoring systems, or local databases into MCP-compatible endpoints without changing existing codebase patterns.



Architectural Implementation & Security


- **Protocol Translation:** On-premises systems host a lightweight “MCP Server” (e.g., exposing local postgres, CSV data, or SSH scripts). Amazon API Gateway’s MCP Proxy translates incoming JSON-RPC 2.0 protocol payloads from AWS Bedrock Agents into standard secure HTTP requests that your local server processes.
- **Enterprise Security & Gateway Services:** Dual authentication is applied:
 1. **Inbound Verification:** Bedrock AgentCore validates the AI agent’s credentials and authorization scopes before initiating requests.
 2. **Outbound Verification:** API Gateway manages secure connections, client certificates (mTLS), and API keys to call the on-premises REST endpoints.
- **Semantic Discovery:** Bedrock AgentCore Gateway allows AI agents to dynamically search and select the most relevant REST APIs that best match the context of the user’s prompt.

Financial Breakdown

- **Amazon Bedrock AgentCore Pricing:** Pay-as-you-go. (Refer to the official Amazon Bedrock AgentCore pricing page for updated regional rates).
- **API Gateway MCP Proxy Charges:** Treated as standard API Gateway REST API calls (\$3.50 per million requests) with no additional platform premium.
- **Network Egress:** Standard AWS Data Transfer Out (\$0.09/GB for ap-southeast-5 beyond the 100 GB free tier).

4. Official AWS Enterprise Hybrid Connections

For organizations requiring dedicated, low-latency, or highly secure private networking that completely bypasses the public internet, AWS offers three primary structural options.



Option A: AWS Site-to-Site VPN

- **Mechanism:** Establishes encrypted IPsec tunnels over the public internet between an AWS Virtual Private Gateway (VGW) / Transit Gateway and an on-premises firewall/router (Customer Gateway).
- **Bandwidth:** Up to 1.25 Gbps per active tunnel (with ECMP support to group multiple tunnels for higher throughput).
- **Costing:**
 - **AWS Connection Fee: \$0.05 per hour** per active VPN connection (\$36.50 per month per site).
 - **On-Premises Hardware:** Requires a physical VPN appliance (Cisco, Juniper, Fortinet) or open-source software (StrongSwan, pfSense) on-premises.
 - **Data Transfer Out:** Charged at standard AWS Data Transfer Out rates (\$0.09 per GB in ap-southeast-5).

Option B: AWS Direct Connect (DX)

- **Mechanism:** A physical, dedicated, high-speed fiber-optic connection from an on-premises datacenter to an AWS Direct Connect Location. Complete bypass of the public internet.
- **Bandwidth:** Fixed port speeds of 1 Gbps, 10 Gbps, or 100 Gbps (Sub-1G connections are available via AWS Direct Connect Partners).
- **Costing:**
 - **AWS Port Hour Charge: \$0.03 per hour** for 1 Gbps ports (~\$21.90/mo) or **\$2.25 per hour** for 100 Gbps ports (~\$1,642.50/mo).
 - **Data Transfer Out:** Significantly lower than standard internet egress. Direct Connect Data Transfer Out is charged at a highly discounted rate of **\$0.021 to \$0.041 per GB** (region-dependent).
 - **Partner Telecommunication Costs:** Third-party telco providers charge monthly circuit fees to run fiber from your office to the AWS Direct Connect Point of Presence (PoP) (ranging from \$500 to \$5,000+ per month).

Option C: AWS Transit Gateway (TGW)

- **Mechanism:** Acts as a centralized cloud router that simplifies hybrid network topology by interconnecting multiple AWS VPCs, Site-to-Site VPNs, and Direct Connect links through a single logical gateway.
- **Costing:**
 - **AWS Attachment Charge: \$0.05 per hour** per VPC, VPN, or DX attachment (~\$36.50 per month per attachment).
 - **Data Processing Fee: \$0.02 per GB** for all data passing through the Transit Gateway.

5. Side-by-Side Hybrid Technology Matrix

Evaluation Criteria	API-Based Connections	MCP-Based Connections	AWS Site-to-Site VPN	AWS Direct Connect
Roadmap Gate Alignment	Year 1, Months 3–6 (Weeks 9–26): Deployed for initial sandbox backend integration with core enterprise database endpoints (Roadmap Portal)	Year 1, Months 7–8 (Weeks 27–34): Activated for secure, real-time contextual search of local catalogs by AI Agents (Roadmap Portal)	Optional Phase 2: Deployed when corporate security policies mandate complete network-level IPsec segregation (Roadmap Portal)	Optional Phase 3: Advanced scaling option if data synchronization volume exceeds 10 TB/month (Roadmap Portal)
Primary Use Case	Lightweight microservices, Webhooks, database queries	AI Agents, tool calling, Bedrock LLM retrieval	General corporate network extension, secure admin	High-performance, bulk data transfer, hybrid DBs
Underlying Network	Public Internet (Secure TLS)	Public Internet (Secure TLS / mTLS)	Encrypted IPsec Tunnel over Public Internet	Dedicated, private, physical fiber-optic line
Max Throughput	Limited only by internet connection	Limited only by internet connection	1.25 Gbps per active tunnel	1 Gbps to 100 Gbps (dedicated)
Latency Profile	Variable (reliant on internet ISP)	Variable (reliant on internet ISP)	Stable, moderate (internet routing)	Ultra-low, deterministic, and SLA-backed
Network Security	Application Layer (TLS, JWT, Mutual TLS)	Application Layer (mTLS, IAM, API Key, Bedrock Core)	Network Layer (IPsec VPN encryption)	Physical Layer isolation (Direct Connect MACsec opt)

Evaluation Criteria	API-Based Connections	MCP-Based Connections	AWS Site-to-Site VPN	AWS Direct Connect
Integration Complexity	Low (Standard HTTP development)	Low-to-Medium (REST config, Bedrock alignment)	Medium (Network routing, IPSec configuration)	High (Requires telco provider physical provisioning)
Implementation Lead Time	Instant	Instant	Hours	30 to 90 Days (Telco fiber pull)
AWS Billing Model	Pay-as-you-go (\$3.50/M requests + DTO)	Pay-as-you-go + standard API gateway rates	Hourly (\$0.05/hr) + Data Transfer Out	Hourly Port (\$0.03-\$2.25/hr) + Partner Loop + DTO
Upfront Capital Cost	\$0.00	\$0.00	\$0.00 (using existing router)	\$1,000 - \$5,000+ (Hardware, cross-connects)

6. Granular Monthly Hybrid Cost Estimations

To assist leadership in budget selection, we model monthly costs based on two enterprise traffic volumes:

1. **Low-Volume Integration (10 GB Outbound Data Transfer + 1 Million Requests / Month)**
2. **High-Volume Integration (500 GB Outbound Data Transfer + 10 Million Requests / Month)**

All estimates assume deployment in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)** with standard regional rates.

Scenario A: Low-Volume Integration Cost Modeling (10 GB Outbound / 1M Requests)

This represents an agile staging environment, internal dev tools, or an experimental AI agentic system integrating on-premises files.

Low-Volume Scenario Total Monthly Cost (USD)	
API-Based	\$3.50
MCP-Based	\$3.50
S2S VPN	\$36.50
Direct Conn.	\$521.90

1. API-Based Connections

- **API Gateway REST Processing:** 1,000,000 requests * \$3.50/M = \$3.50 USD
- **AWS Data Transfer Out (Egress):** 10 GB (Within free tier threshold of 100 GB) = \$0.00 USD
- **NAT Gateway Hourly Fee (Existing):** Shared with existing infrastructure = \$0.00 USD
- **TOTAL MONTHLY COST: \$3.50 USD / ~RM 15.75 MYR**

2. MCP-Based Connections

- **API Gateway with MCP Proxy:** 1,000,000 requests * \$3.50/M = \$3.50 USD
- **AWS Data Transfer Out (Egress):** 10 GB (Within free tier threshold of 100 GB) = \$0.00 USD
- **Bedrock AgentCore Platform Fee:** Pay-as-you-go nominal charges = \$0.00 USD (Approx)
- **TOTAL MONTHLY COST: \$3.50 USD / ~RM 15.75 MYR**

3. AWS Site-to-Site VPN

- **AWS VPN Connection Charge:** 1 tunnel * 730 hours/month * \$0.05/hr = \$36.50 USD
- **AWS Data Transfer Out (Egress):** 10 GB (Within free tier threshold of 100 GB) = \$0.00 USD
- **On-Premises Router Power/Compute:** Nominal standard overhead = \$0.00 USD
- **TOTAL MONTHLY COST: \$36.50 USD / ~RM 164.25 MYR**

4. AWS Direct Connect (DX)

- **AWS 1 Gbps Port Charge:** 730 hours/month * \$0.03/hr = \$21.90 USD
- **Discounted Data Transfer Out (Egress):** 10 GB * \$0.041/GB = \$0.41 USD
- **Partner Local Loop Circuit:** Sub-1G Partner line (Est. Baseline Partner Fee) = \$500.00 USD
- **TOTAL MONTHLY COST: \$522.31 USD / ~RM 2,350.40 MYR**

Scenario B: High-Volume Integration Cost Modeling (500 GB Outbound / 10M Requests)

This represents active production workloads, automated ETL syncs, enterprise-wide chat agents, or heavy file replication.

High-Volume Scenario Total Monthly Cost (USD)	
API-Based	\$71.00
MCP-Based	\$71.00
S2S VPN	\$72.50
Direct Conn.	\$542.40

1. API-Based Connections

- **API Gateway REST Processing:** 10,000,000 requests * \$3.50/M = \$35.00 USD
- **AWS Data Transfer Out (Egress):** (500 GB - 100 GB Free) = 400 GB * \$0.09/GB = \$36.00 USD
- **TOTAL MONTHLY COST: \$71.00 USD / ~RM 319.50 MYR**

2. MCP-Based Connections

- **API Gateway with MCP Proxy:** 10,000,000 requests * \$3.50/M = \$35.00 USD
- **AWS Data Transfer Out (Egress):** (500 GB - 100 GB Free) = 400 GB * \$0.09/GB = \$36.00 USD

- **Bedrock AgentCore Platform Fee:** Pay-as-you-go nominal charges = \$0.00 USD (Approx)
- **TOTAL MONTHLY COST: \$71.00 USD / ~RM 319.50 MYR**

3. AWS Site-to-Site VPN


- **AWS VPN Connection Charge:** 1 tunnel * 730 hours/month * \$0.05/hr = \$36.50 USD
- **AWS Data Transfer Out (Egress):** (500 GB - 100 GB Free) = 400 GB * \$0.09/GB = \$36.00 USD
- **TOTAL MONTHLY COST: \$72.50 USD / ~RM 326.25 MYR**

4. AWS Direct Connect (DX)

- **AWS 1 Gbps Port Charge:** 730 hours/month * \$0.03/hr = \$21.90 USD
- **Discounted Data Transfer Out (Egress):** 500 GB * \$0.041/GB = \$20.50 USD
- **Partner Local Loop Circuit:** Sub-1G Partner line (Est. Baseline Partner Fee) = \$500.00 USD
- **TOTAL MONTHLY COST: \$542.40 USD / ~RM 2,440.80 MYR**

7. Strategic Recommendation & Management Decision Tree

To guide management through the selection process, we present a logical decision matrix to streamline technical approval:



Actionable Roadmap for Approvals

- 1. **Phase 1: Zero-Upfront Trial (API / MCP)**
 - **Recommendation:** Begin with **API-based or MCP-based connections**. It requires \$0 in capital expenditure, zero hardware setups on-premises, and allows the AI team to immediately utilize Bedrock AgentCore and local datasets.
 - **Security Action:** Restrict access to the on-premise MCP server by whitelisting the AWS NAT Gateway Elastic IP and enforcing API key rotation.
- 2. **Phase 2: Hybrid Network Scaling (Site-to-Site VPN)**
 - **Trigger:** If corporate security mandates network-level segmentation (e.g., preventing any endpoint from exposing port 443 to the public internet, even with whitelists), transition to **AWS Site-to-Site VPN**.
 - **Cost Impact:** Adds a modest **\$36.50/month (RM 164.25)** flat connection charge plus standard outbound data transfer.
- 3. **Phase 3: Ultimate Performance (Direct Connect)**
 - **Trigger:** If synchronization volumes exceed **10 TB / month** or latency constraints require single-digit millisecond response times.
 - **Cost Impact:** Heavy capital and partner operating expenditures (RM 2,400+ / mo). Highly discounted AWS egress rates make Direct Connect financially viable only at extreme data transfer volumes.

Disaster Recovery Options & Sovereignty Guide

#aws #cloud #architecture #vpc #alb #asg #rds #waf #elasticsearch #valkey #dns #disaster-recovery #efs #postgresql #sovereignty #compliance

Disaster Recovery (DR) Options & National Sovereignty Guide

1. Executive Summary & Context

Maintaining system resilience is critical for enterprise deployment. This guide establishes a comprehensive, production-ready Disaster Recovery (DR) playbook modeled directly after the official AWS whitepaper [“Disaster Recovery of Workloads on AWS: Recovery in the Cloud”](#).

This guide addresses three core dimensions:

- Existing Architecture Alignment:** Adapting DR practices directly to our Multi-AZ, 3-Tier layout (WAFv2, ALB, private Auto Scaling Groups for Frontend, Backend, and AI tiers, Standalone EC2 staging environments, Multi-AZ RDS PostgreSQL 16 database, Amazon ElastiCache for Valkey, and Amazon EFS shared storage).
- Malaysia National Sovereignty:** A strict regulatory compliance review under the **Malaysian Personal Data Protection Act (PDPA) 2010** (including the **Personal Data Protection (Amendment) Act 2024** and the **2025 Cross-Border Personal Data Transfer Guidelines (CBPDT)**).
- Financial Transparency (USD / MYR):** Granular cost breakdowns comparing **In-Region Multi-AZ DR** (sovereign, 100% data residency in Malaysia ap-southeast-5) and **Cross-Region DR** (replicated to Singapore ap-southeast-1 or Jakarta ap-southeast-3).

All currency estimations use a baseline baseline conversion rate of **1 USD ≈ 4.50 MYR**.

2. Malaysia National Sovereignty & Regulatory Compliance

When designing DR topologies for applications processing citizen data in Malaysia, **data sovereignty** and **jurisdictional boundaries** are primary architectural drivers.

Data Residency & Transfer Compliance Pathways	
In-Region DR (ap-southeast-5)	Cross-Region DR (Out of MY)
100% Malaysian Sovereignty - No Cross-Border transfer concerns - Fully compliant with 2025 CBPDT - Complete immunity to foreign laws	- Subject to PDPA Section 129 - Requires Transfer Impact (TIA) - Explicit Data Subject Consent - Requires Contractual Clauses

2.1 The Personal Data Protection Act (PDPA) & 2024/2025 Frameworks

- Section 129 Principal Prohibition:** Under the Malaysian PDPA 2010, transferring personal data outside of Malaysia is prohibited unless specifically exempted by the Minister or falling under specific statutory exceptions.
- The 2024 Amendments:** The Personal Data Protection (Amendment) Act 2024 introduced mandatory **Data Breach Notifications (DBN)** and the compulsory appointment of a **Data Protection Officer (DPO)**, escalating the legal risk of data exposure.
- The 2025 CBPDT Guidelines:** Published in April 2025 by the Personal Data Protection Commissioner, the **Cross Border Personal Data Transfer (CBPDT) Guidelines** clarify that data controllers transferring data outside Malaysia must conduct a **Transfer Impact Assessment (TIA)** to ensure the destination provides “adequate protection” substantially similar to the PDPA.

2.2 Sovereign In-Region Multi-AZ DR vs. Cross-Region DR

Sovereign In-Region Multi-AZ DR (Absolute Residency)

- Architectural Flow:** All primary workloads and disaster recovery assets reside entirely within the three physical Availability Zones of the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)**.
- PDPA Standing:** Fully compliant out of the box. Because data never crosses Malaysia’s geographical borders, it is **exempt** from Section 129 restrictions, TIAs, or extra consent requirements.
- Jurisdictional Immunity:** Data remains protected exclusively under Malaysian courts. It is immune to extraterritorial access requests (such as the US CLOUD Act or Singaporean administrative warrants) that might apply if the data was replicated overseas.

Cross-Region DR (Jurisdictional Border Crossing)

- Architectural Flow:** Backups, read-replicas, or standby environments are copied from ap-southeast-5 to **Singapore (ap-southeast-1)** or **Jakarta (ap-southeast-3)**.
- PDPA Section 129 Compliance Checklist:** To legally implement Cross-Region replication, the organization must fulfill these strict legal requirements:
 - Explicit Data Subject Consent:** Users must actively opt-in to their data being transferred and stored in Singapore/Jakarta, detailed within a clear, updated Privacy Notice.
 - Transfer Impact Assessment (TIA):** The DPO must perform and log a formal TIA assessing the recipient country’s data laws (e.g., Singapore’s PDPA 2012 or Indonesia’s UU PDP).
 - Data Processor Agreements:** Binding contracts with AWS and any third-party processors enforcing technical and organizational safety standards equivalent to Malaysian standard standards.

2.3 Architectural & Engineering Approach to Sovereignty in Cross-Region DR

To mitigate regulatory risks and address national data sovereignty concerns when using Cross-Region DR (e.g. replicating data from Malaysia ap-southeast-5 to Singapore ap-southeast-1 or Jakarta ap-southeast-3), the architectural layout implements several strict technical safeguards:

1. Field-Level Encryption & Tokenization

- **Strategy:** Before any database replication payload, block device copy, or backup snapshot is transmitted across jurisdictional boundaries, all highly sensitive Personally Identifiable Information (PII) — such as national identification numbers, phone numbers, and raw user credentials — is subjected to field-level encryption or tokenization inside the primary Malaysia VPC.
- **Result:** Only fully encrypted, non-readable cyphertext or synthetic tokens are transmitted across borders, while the lookup table and decryption logic remain strictly localized inside ap-southeast-5. Even if foreign authorities subpoena the storage drives in Singapore or Jakarta, they cannot access the actual plaintext citizen identities.

2. Sovereign Key Management & Cryptographic Isolation

- **Strategy:** All AWS KMS (Key Management Service) Customer Managed Keys (CMKs) or Hardware Security Modules (HSMs) used to encrypt cross-region backups and databases are hosted and controlled exclusively inside the Malaysia (ap-southeast-5) region.
- **Result:** The decryption keys are never replicated to Singapore or Jakarta. To boot recovery EC2 instances or decrypt RDS backups during a disaster, foreign hypervisors must dynamically call the Malaysia KMS API. In a geopolitical crisis, the security team can immediately revoke KMS key access from the primary region, completely neutralizing the foreign environment and rendering the replicated data permanently unreadable.

3. Sovereign Failover Guardrails & Circuit Breakers

- **Strategy:** The Route 53 routing policies and ASG launch parameters are locked behind multi-signature IAM permission controls and active “Sovereignty Circuit Breakers.” Automated failover to the cross-region standby only triggers if the entire Malaysian AWS region is physically unavailable (e.g. undersea cable cuts or natural disasters), not during minor localized software glitches.
- **Result:** This prevents accidental data migration or “regulatory leakage” where application traffic is unnecessarily routed to overseas instances during routine operating conditions.

4. Legal & Compliance Alignment: Transfer Impact Assessments (TIA)

- **Strategy:** Under Section 129 of the PDPA 2010 and the 2025 CBPDT Guidelines, any cross-region flow requires a logged TIA detailing:
 - Technical measures (e.g. KMS, Field-Level Encryption) ensuring identical security.
 - Recipient-region legal framework analysis (e.g. comparing Malaysia’s PDPA with Singapore’s PDPA 2012 or Indonesia’s UU PDP).
 - Explicit Standard Contractual Clauses (SCCs) embedded into the Cloud Service Provider agreement to legally bind AWS to Malaysian-standard data preservation.


3. Disaster Recovery Spectrum Analysis

We analyze the four standard AWS Disaster Recovery options, alongside AWS Elastic Disaster Recovery (AWS DRS) as a modern, cost-optimized, and near-zero RPO hybrid strategy. Each strategy presents a distinct balance of Recovery Point Objective (RPO), Recovery Time Objective (RTO), implementation complexity, and ongoing operating cost.

Low Cost High RPO/RTO					High Cost Near-Zero RPO/RTO
Backup & Restore	Pilot Light	Warm Standby	Multi-Site Active-Active	AWS Elastic DR (AWS DRS) (Highly cost-eff. Sub-Minute RPO)	

Strategy A: Backup and Restore

A passive approach where virtual machine images, database snapshots, and configuration files are regularly backed up and restored to a target environment only after an outage is declared.



1. Architectural Setup & Component Layout

- **Database (RDS):** Automated daily snapshots and transaction logs stored via AWS Backup. In-Region copies are kept in a separate secure AWS account; Cross-Region copies sync to Singapore (ap-southeast-1) S3 buckets.
- **Compute (ASG):** Launch Templates are backed up. AMIs of the Frontend, Backend, and AI Standalone/ASG nodes are pre-baked using our Packer/Ansible pipeline and stored in local or remote AMI catalogs.
- **Shared Storage (EFS):** Daily EFS backups handled via AWS Backup. EFS filesystems are not pre-provisioned in the target environment to save costs; they are recreated during recovery.
- **DNS & Routing (Route 53):** Route 53 pointing to the primary ALB. No active health check routing is configured. Manual DNS update is required during a failover event.
- **Valkey Cache:** No replication. Caching nodes are left unprovisioned in the target area and cold-booted during restoration.

2. Recovery Objectives


- **Recovery Point Objective (RPO): 24 Hours** (determined by backup execution intervals).
- **Recovery Time Objective (RTO): 4 to 8 Hours** (time needed to provision network infrastructure, restore RDS database storage, spin up EC2 instances from AMIs, and update DNS).

3. Trade-offs & Analysis

- **Cost:** Extremely low. No active compute, load balancer, or database instances run in the DR zone.
- **Complexity:** Low configuration complexity, but high manual recovery complexity during an incident.
- **Sovereignty:** If utilizing Cross-Region Backup and Restore, encrypted database snapshots reside in Singapore, which mandates standard PDPA Section 129 consent and TIAs. In-Region backups are completely sovereign.

Strategy B: Pilot Light

A core database and persistent storage layer are kept actively running and synchronized, while the compute/application server layer remains unprovisioned or “turned off” until a disaster occurs.



1. Architectural Setup & Component Layout

- Database (RDS):** A minimal, single-AZ RDS PostgreSQL read-replica is kept online and continuously synchronized via asynchronous replication.
- Compute (ASG):** Launch Templates are active in the target VPC, but the ASG **Desired Capacity is set to 0**. Standalone staging environments are not provisioned.
- Shared Storage (EFS):** Continuous replication to a target EFS filesystem using AWS EFS Replication.
- DNS & Routing:** Route 53 DNS is configured with active health checks. However, failover requires changing the ASG desired capacity to active levels and waiting for compute nodes to bootstrap.
- Valkey Cache:** Unprovisioned. Provisioned only during failover.

2. Recovery Objectives


- Recovery Point Objective (RPO):** < 5 Minutes (determined by RDS replica lag and EFS replication latency).
- Recovery Time Objective (RTO):** 30 to 45 Minutes (time to promote RDS replica to primary, scale the ASG up from 0 to target size, wait for application health checks to pass, and update Route 53).

3. Trade-offs & Analysis

- Cost:** Moderate. You pay for continuous RDS replica compute and persistent EFS/S3 storage, but avoid active EC2 compute costs.
- Complexity:** Medium. Requires automating the promotion of the database and the auto-scaling group scale-up.
- Sovereignty:** Cross-Region replica databases contain active, unencrypted-in-memory data in Singapore/Jakarta. This triggers strict PDPA compliance, requiring a comprehensive TIA and explicit user consent. In-Region Pilot Light (using a separate VPC in the same region) is fully sovereign.

Strategy C: Warm Standby

A scaled-down but fully functional copy of the entire 3-tier architecture is kept running continuously in the target zone. It handles zero (or minimal) active traffic under normal operations but can instantly scale up to handle the production load if the primary site fails.



1. Architectural Setup & Component Layout

- Presentation Tier:** A standby ALB is active. AWS WAFv2 is associated and running.
- Compute Tier (ASG):** The standby ASG runs continuously with a minimal capacity (e.g., 1x t4g.medium instead of the production 2x t4g.xlarge). Standalone staging instances are unprovisioned to optimize cost.
- Database (RDS):** An active RDS PostgreSQL read-replica runs continuously in the standby zone.
- Valkey Cache:** A minimal single-node ElastiCache Valkey (cache.t4g.micro) runs to maintain cache readiness.
- Shared Storage (EFS):** Continuous EFS Replication keeps model weights and configurations in sync.
- DNS & Routing:** Route 53 Active-Passive failover routing is configured. Health checks automatically redirect traffic to the standby ALB if the primary fails.

2. Recovery Objectives


- Recovery Point Objective (RPO):** < 1 Minute (near-real-time database and storage replication).
- Recovery Time Objective (RTO):** 5 to 10 Minutes (traffic is routed instantly; auto-scaling takes a few minutes to scale compute instances to full production sizing).

3. Trade-offs & Analysis

- Cost:** High. Requires running a second ALB, active EC2 instances, replica RDS instances, EFS replication, and an ElastiCache node.
- Complexity:** Medium-High. Requires robust automation for target scaling policies and failover signaling.
- Sovereignty:** Storing live, unencrypted client data continuously on active databases in foreign regions requires strict PDPA Section 129 audit trails and TIAs.

Strategy D: Multi-Site Active-Active

Traffic is split dynamically across two identical, full-scale production environments running simultaneously in two separate zones. Both sites actively serve client requests.



1. Architectural Setup & Component Layout

- Presentation Tier:** Active ALBs and WAFv2 are deployed in both sites.
- Compute Tier (ASG):** Full-scale ASGs (2x t4g.xlarge nodes) run actively in both regions. All Standalone AMI-baking instances are kept in sync.
- Database (RDS):** Configured as a cross-region active-active database cluster (utilizing Aurora Global Database or custom bi-directional replication engines), or active-passive database with write-routing to the primary site.

- **Valkey Cache:** Independent active Valkey clusters in each region.
- **Shared Storage (EFS):** Continuous bi-directional synchronization or dual EFS staging setups.
- **DNS & Routing:** Route 53 is configured with Geolocation or Latency-based routing to dynamically distribute user requests.

2. Recovery Objectives


- **Recovery Point Objective (RPO): Near Zero (Real-time).**
- **Recovery Time Objective (RTO): Near Zero** (instantaneous failover as both environments are already serving live traffic).

3. Trade-offs & Analysis

- **Cost:** Extremely High. Double the baseline production cost plus significant cross-region data transfer replication fees.
- **Complexity:** Extremely High. Requires solving split-brain scenarios, database write collisions, and high data consistency synchronization challenges.
- **Sovereignty:** Deeply critical. 50% of Malaysian citizen data is processed and stored outside local jurisdiction in real time. This triggers comprehensive PDPA scrutiny, making TIAs, explicit consent notices, and regulatory registration mandatory.

Strategy E: AWS Elastic Disaster Recovery (AWS DRS)

AWS Elastic Disaster Recovery (AWS DRS) minimizes downtime and data loss with fast, reliable recovery of on-premises and cloud-based applications using affordable storage, minimal compute, and point-in-time recovery. It provides sub-minute RPO and minutes-level RTO.



1. Architectural Setup & Component Layout

- **Source Servers:** Install the AWS Elastic Disaster Recovery Replication Agent on target source servers (on-premises or remote cloud nodes) to initiate secure, block-level data replication.
- **Staging Area Subnet:** Data is continuously replicated to a dedicated, low-cost staging area subnet in the designated AWS account and Region (e.g., ap-southeast-5).
- **Compute & Storage Optimization:** Staging costs are highly minimized. Lightweight, automatically managed EC2 replication instances (e.g., t3.small nodes) handle incoming blocks, while affordable EBS volumes (gp3 or low-cost sc1) act as the replication target disks.
- **Non-Disruptive Testing:** Allows running seamless, non-disruptive disaster recovery drills and tests at any time without impacting active replication or source servers.
- **Drill/Recovery Launches:** When a failover or drill is initiated, AWS DRS automatically launches target EC2 instances based on the most up-to-date server state or a specified historical point-in-time state.
- **Failback Replication:** After the primary site issue is resolved, data replication is initiated in reverse back to the primary site, ensuring seamless failback.

2. Recovery Objectives

- **Recovery Point Objective (RPO): Seconds to Sub-Minute** (continuous, real-time block-level asynchronous replication).
- **Recovery Time Objective (RTO): Minutes** (automated launch, conversion, and orchestration of recovery instances).

3. Trade-offs & Analysis

- **Cost:** Highly Cost-Effective. Replicating servers cost only a nominal hourly software license fee (\$0.028 per server per hour) plus low-cost staging compute (t3.small) and staging EBS volumes, instead of running full warm standby or active-active duplicate environments.
- **Complexity:** Medium. Requires installing replication agents on source servers and configuring launch templates, but failover and failback orchestration are highly automated.
- **Sovereignty:** 100% Sovereign when staging and recovery are targeted inside the **AWS Malaysia Region (ap-southeast-5)**. Because the continuous block replication targets local sovereign physical boundaries, it completely satisfies PDPA and 2025 CBPDT requirements for localized residency, avoiding complex cross-border Transfer Impact Assessments (TIAs).

4. Comprehensive Cost Estimation Model (USD & MYR)

To provide clear financial visibility, we present granular monthly costs for all four DR strategies, comparing:

1. **Sovereign In-Region Multi-AZ DR** (All standby resources deployed inside ap-southeast-5 using a separate DR VPC).
2. **Cross-Region DR** (Standby resources deployed in Singapore ap-southeast-1).

All prices are calibrated for On-Demand rates (~730 hours/month) for our **High-Performance Production Stack**.

4.1 Granular Monthly DR Cost Matrix (USD / Month)

Component	Strategy A: Backup & Restore	Strategy B: Pilot Light	Strategy C: Warm Standby	Strategy D: Active-Active	Strategy E: AWS DRS (3 Node Baseline)
Networking & NAT	\$0.00	\$0.00 (No outbound needed)	\$32.85 (1x NAT Gateway)	\$32.85 (1x NAT Gateway)	\$0.00 (Staging in existing private subnet)
Public IPv4 Addr	\$0.00	\$0.00	\$7.30 (1x NAT IP, 1x ALB IP)	\$10.95 (1x NAT, 2x ALB IPs)	\$0.00
ALB & WAFv2	\$0.00	\$0.00	\$25.03 (ALB + WAF ACL)	\$41.46 (ALB + WAF + Rules)	\$0.00 (Cold; deployed during failover)
Compute (ASG)	\$0.00 (Desired = 0)	\$0.00 (Desired = 0)	\$49.06 (1x t4g.medium warm)	\$196.22 (2x t4g.xlarge active)	\$15.18 (1x shared t3.small replication node)

Component	Strategy A: Backup & Restore	Strategy B: Pilot Light	Strategy C: Warm Standby	Strategy D: Active-Active	Strategy E: AWS DRS (3 Node Baseline)
Database (RDS)	\$0.00	\$221.92 (1x db.m6g.large replica)	\$443.84 (Multi-AZ replica)	\$443.84 (Multi-AZ active)	\$0.00 (Replicated at volume block layer)
Valkey Cache	\$0.00	\$0.00	\$9.34 (cache.t4g.micro)	\$39.71 (cache.t4g.medium)	\$0.00 (Unprovisioned in staging)
Storage (EFS)	\$0.00	\$15.00 (50GB replica)	\$15.00 (50GB replica)	\$15.00 (50GB active)	\$15.00 (50GB replica)
Storage (S3)	\$2.30 (100GB backup)	\$2.30 (100GB backup)	\$2.30 (100GB backup)	\$2.30 (100GB active)	\$2.30 (100GB backup/snapshots)
DRS Service Fee	\$0.00	\$0.00	\$0.00	\$0.00	\$61.32 (3 servers * \$0.028/hr * 730 hrs)
Data Replication	\$5.00 (AWS Backup fees)	\$15.00 (RDS & EFS sync)	\$30.00 (Live RDS & EFS sync)	\$120.00 (High bi-dir sync)	\$18.50 (Block staging volumes & snap sync)
TOTAL (In-Region)	\$7.30 USD	\$254.22 USD	\$614.42 USD	\$902.33 USD	\$112.30 USD
TOTAL (Cross-Region)	\$11.50 USD	\$289.50 USD	\$664.80 USD	\$982.50 USD	\$135.50 USD

Note: Cross-Region costs are slightly higher due to AWS Cross-Region Data Transfer Egress fees (\$0.09/GB from ap-southeast-5) and marginally higher base pricing in ap-southeast-1. AWS DRS cost modeling assumes replicating 3 active EC2 compute nodes (Frontend, Backend, AI Tier) into a secure staging subnet.

4.2 Local Currency Equivalent Comparison (MYR / Month)

Using the exchange baseline of **1 USD ≈ 4.50 MYR**, the monthly DR infrastructure spend equates to:

Monthly DR Cost Comparison (MYR equivalent)		
In-Region (ap-southeast-5)		
└─ Backup & Restore	⇒ RM 32.85	(Highly Recommended)
└─ AWS DRS	⇒ RM 505.35	
└─ Pilot Light	⇒ RM 1,143.99	
└─ Warm Standby	⇒ RM 2,764.89	
└─ Active-Active	⇒ RM 4,060.49	
Cross-Region (Singapore/Jakarta)		
└─ Backup & Restore	⇒ RM 51.75	
└─ AWS DRS	⇒ RM 609.75	
└─ Pilot Light	⇒ RM 1,302.75	
└─ Warm Standby	⇒ RM 2,991.60	
└─ Active-Active	⇒ RM 4,421.25	


5. Decision Matrix & Strategy Trade-offs

To guide executive management through the selection process, we present a technical evaluation matrix:

Evaluation Dimension	Strategy A: Backup & Restore	Strategy B: Pilot Light	Strategy C: Warm Standby	Strategy D: Active-Active	Strategy E: AWS DRS (Elastic DR)
Roadmap Gate Alignment	Year 1, Months 9–12 (Weeks 35–52): Baseline local snapshots to prepare for CRM schema inception (Roadmap Portal)	Optional Phase 3 (As Needed): Promoted as an active read-replica to offload real-time SQL queries (Roadmap Portal)	Unrecommended Option: High operating overhead; redundant cost is hard to justify	Unrecommended Option: Extreme cross-border PDPA regulatory and security risks	Year 2, Months 1–2 (Weeks 53–60): Continuous replication to secure staging subnets for CRM launch (Roadmap Portal)
Target RPO	24 Hours	< 5 Minutes	< 1 Minute	Near Real-time	Seconds to Sub-Minute
Target RTO	4 to 8 Hours	30 to 45 Minutes	5 to 10 Minutes	Near Zero	Minutes
Relative Cost	Extremely Low (1x)	Moderate (35x)	High (85x)	Very High (120x)	Low-to-Moderate (15x)
Engineering Overhead	Low (Standard backups)	Medium (Failover scripts)	High (Automated Route 53)	Extremely High	Medium (Agent setup)
Regulatory Risk	Minimal	Medium (Replica active)	High (Continuous Sync)	Extremely High	Minimal (Sovereign staging)
PDPA Compliance Path	Simple	Requires TIA & Consent	Requires TIA & Consent	Requires Full Audit	Simple (100% In-Region)
Sovereignty Standing (In-Region)	100% Secure (All data, backups, and compute remain strictly in ap-southeast-5. Exempt from Section 129.)	100% Secure (All sync replica DBs and compute templates remain strictly in ap-southeast-5.)	100% Secure (Active standby stack resides 100% within ap-southeast-5 borders.)	100% Secure (Both active sites are configured within Malaysia ap-southeast-5 AZs.)	100% Secure (All block-level staging and target recovery targets are within ap-southeast-5.)
Sovereignty Standing (Cross-Region)	Conditional / Low Risk (Static encrypted backups cross borders. Requires standard TIA and user consent notice.)	Medium Risk (Live RDS replica runs in Singapore/Jakarta. Active memory footprint in foreign jurisdiction; requires robust TIA.)	High Risk (Continuous sync of live application state, DB, cache, and EFS to foreign region. Requires strict SCCs and DPO audit.)	Very High Risk (50% of real-time writes & queries processed on foreign soil. High exposure to foreign law enforcement; requires explicit consent & audits.)	Conditional / Low-to-Med Risk (Continuous block replication to lightweight staging. Data is fully encrypted with KMS keys owned/managed in Malaysia.)

6. Strategic Recommendations & Roadmap

To balance **business continuity**, **cost-efficiency**, and **strict compliance with Malaysia National Sovereignty (PDPA)**, the following phased roadmap is recommended:



Phase 1: Deploy Sovereign In-Region Backup & Restore (Default Baseline)

- **Action:** Configure AWS Backup to take daily snapshots of our Multi-AZ RDS PostgreSQL database, the EFS model volume, and the gp3 EBS volumes of our compute nodes. Store these backups securely inside ap-southeast-5.
- **Sovereignty Advantage:** 100% compliant with the Malaysian PDPA. Zero data crosses national borders, completely bypassing Section 129 regulatory hurdles, saving legal/DPO consultation fees.
- **Cost Impact:** Negligible baseline charge (~\$7.30 USD / RM 32.85 per month).

Phase 2: Implement AWS Elastic Disaster Recovery (AWS DRS) (Near-Zero RPO / RTO)

- **Action:** Install the AWS DRS Replication Agent on active servers to continuously replicate block-level changes asynchronously into a dedicated staging area subnet inside ap-southeast-5. Configure DRS launch templates to spin up fully-provisioned target instances upon a drill or failover event.
- **Sovereignty Advantage:** Keeps staging and target areas completely localized within the AWS Malaysia region. Absolute data sovereignty is maintained, satisfying the 2025 CBPDT Guidelines natively while achieving sub-minute RPO.
- **Cost Impact:** Exceptionally cost-effective (~\$112.30 USD / RM 505.35 per month), avoiding active duplicate server overhead until disaster declaration.

Phase 3: Transition to Sovereign In-Region Pilot Light (As Real-Time Hot Read Needs Arise)

- **Action:** If the application scales further and business workflows require an active, continuously queryable read-replica for real-time reads/BI reports in addition to DR, deploy an RDS PostgreSQL Read Replica in a separate VPC/subnet configuration within the **Malaysia region (ap-southeast-5)**, keeping the ASG desired capacity at 0.
- **Sovereignty Advantage:** Maintains absolute local residency. Enables fast database promotion and compute failover without violating data sovereignty policies.
- **Cost Impact:** Highly cost-optimized (~\$254.22 USD / RM 1,143.99 per month), representing a fraction of the cost of running a full Warm Standby or Cross-Region Active-Active cluster.

RDS PostgreSQL 17 vs. Percona Server for PostgreSQL 17

#aws #cloud #architecture #rds #dns #disaster-recovery #postgresql #compliance

Deep Technical Comparison: AWS RDS PostgreSQL 17 vs. Self-Installed Percona Server for PostgreSQL 17

Evaluating database platforms for a highly secure, high-performance 3-tier architecture in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)** requires a careful analysis of managed service convenience versus raw open-source control.


This document provides a comprehensive technical comparison between **AWS Relational Database Service (RDS) PostgreSQL 17 (Multi-AZ)** and a **self-installed Percona Server for PostgreSQL 17** deployed on Amazon EC2 Graviton instances. We analyze performance traits, telemetry, extension availability, query planner optimization, cost models, and operational architecture designs.

1. Architectural Designs

To achieve highly available, production-grade resiliency, both options must be designed to withstand Availability Zone (AZ) outages and ensure automatic failover.

Design A: AWS RDS PostgreSQL 17 (Multi-AZ)

AWS RDS simplifies high availability through synchronous replication at the block storage level. A primary DB instance is provisioned in AZ-A, and a hot standby replica is synchronously maintained in AZ-B.





RDS Failover Mechanism

If the primary instance fails, RDS automatically triggers a failover by changing the canonical name record (CNAME) of the DB endpoint to point to the standby instance. This typically completes in **60 to 120 seconds** depending on transaction volumes and crash recovery processes.

Design B: Percona Server for PostgreSQL 17 on EC2 (Patroni, PgBouncer & etcd)

Deploying Percona Server on EC2 requires a self-managed High Availability (HA) stack. This architecture uses **Patroni** as the template for cluster management, **etcd** as the Distributed Consensus Store (DCS), **PgBouncer** for connection pooling, and **pg_backrest** for robust, compressed backup streaming to S3. For onsite on-premises VM deployments, we mirror this identical design pattern (utilising HAProxy and containerised etcd/Patroni Quadlets) as detailed in the [Enterprise Percona Server for PostgreSQL 17 On-Premises Guide](#).





Percona Failover Mechanism

Patroni continuously monitors cluster health and registers state changes in the **etcd** Distributed Consensus Store (DCS). If the primary node fails or loses its leader key in etcd, the surviving nodes hold an election. The node with the most up-to-date Write-Ahead Log (WAL) is promoted to primary. PgBouncer is automatically updated with the new write destination, leading to a highly responsive failover time of **10 to 30 seconds**.

2. Performance Comparison & Telemetry

Managed services restrict low-level performance tuning in exchange for operating system abstraction. Percona Server for PostgreSQL unlocks full infrastructure optimization and integrates enterprise features out of the box.

Telemetry and Observability

AWS RDS PostgreSQL

- **Tools:** AWS CloudWatch Metrics, Enhanced Monitoring (OS metrics via an agent), and **RDS Performance Insights**.
- **Capabilities:** Performance Insights is a robust tool that measures database load using **Average Active Sessions (AAS)**. It visualizes waiting threads, query text, and wait events (e.g., locking, I/O, CPU bottle-necks).
- **Constraints:** CloudWatch collects metrics in discrete blocks (minimum 1-second intervals for Enhanced Monitoring; 1-minute default for CloudWatch). Accessing database internal states is limited, and exporting granular query telemetry to third-party dashboards requires complex agent setups or costly exporters.

Percona Server for PostgreSQL

- **Tools:** **Percona Monitoring and Management (PMM)** (fully integrated, free, and open-source).
- **Capabilities:** PMM provides deep observability into both the Linux OS kernel and the PostgreSQL database. It captures query-level execution times, index usage, connection behavior, and locks. PMM uses Prometheus for time-series metrics and Grafana for visual dashboards.
- **Query Analytics:** PMM reads directly from the enhanced `pg_stat_monitor` extension, offering chronological query plans, visual query performance over time, and client-metadata tracking.

Extensions & Advanced Observability

The difference in extensibility is a critical factor when choosing between AWS RDS and Percona Server.

Feature / Extension	AWS RDS PostgreSQL 17	Percona Server for PostgreSQL 17	Technical Impact
pg_stat_monitor	❌ Not Supported (Only generic <code>pg_stat_statements</code> is available)	Built-in	<code>pg_stat_monitor</code> tracks query performance using bucket-based intervals , capturing exact query plan metrics, timing histograms, client IP addresses, execution vs. planning times, and query parameters safely without exposing sensitive constants.
Custom Dynamic Libraries	❌ Blocked	Allowed	RDS only supports a predefined list of AWS-approved extensions. Custom or proprietary dynamic-link libraries (<code>shared_preload_libraries</code>) are restricted due to lack of OS-level superuser rights.
pg_repack	Supported (Pre-installed)	Supported	Allows rebuilding tables/indexes online to reclaim bloated disk space without exclusive locks.
pgaudit	Supported (Pre-installed)	Supported	Provides detailed session and object audit logging for enterprise compliance.

Architectural Tuning Capability

Because AWS RDS abstracts the underlying operating system, users cannot perform system-level resource optimization. On Amazon EC2, Percona Server allows you to fine-tune the OS kernel and hardware integration.

File System Optimization

- **AWS RDS:** Uses the default AWS Linux filesystem structure (typically ext4-based) with standard mount parameters.
- **Percona Server (EC2):** Allows using **XFS** or **ZFS** optimized for PostgreSQL. For example, aligning the filesystem block size (typically 4KB) with PostgreSQL's page size (8KB) prevents **torn-write pages** and reduces WAL write amplification. Mount parameters such as `noatime` can be enabled to bypass write overheads during read-heavy workloads.

Kernel Tuning & Huge Pages

- **AWS RDS:** Configures standard memory allocations dynamically. While it supports standard Large Pages, fine-grained adjustments to kernel-level memory overcommit options or socket queue sizes are inaccessible.
- **Percona Server (EC2):** Enables precise optimization of system memory and swap behaviors:

```
# Optimize kernel memory retention to avoid aggressive swapping
sysctl -w vm.swappiness=1
sysctl -w vm.overcommit_memory=2
sysctl -w vm.overcommit_ratio=80

# Configure and lock Huge Pages to guarantee block buffer efficiency
sysctl -w vm.nr_hugepages=4096
```

Connection Pooling

- **AWS RDS:** Relies on AWS RDS Proxy (an extra paid resource starting at \$0.015 per LCU-hour) or application-side poolers.
- **Percona Server (EC2):** Allows deploying **PgBouncer** directly on the same EC2 database hosts or dedicated scaling nodes. Configuring PgBouncer in **Transaction Mode** allows the database to easily support tens of thousands of active connections with minimal memory overhead.

Query Planner Optimization

Both platforms utilize the core PostgreSQL 17 Query Planner engine, but Percona Server provides more freedom to tune behavior based on exact hardware characteristics.

- **Cost Parameters:** Developers can adjust parameters such as `random_page_cost` and `seq_page_cost`. On RDS, these are modified via dynamic DB Parameter Groups. On Percona Server, they can be configured directly in `postgresql.conf` or set per-session/database to match specific storage benchmarks.
- **Query Hinting:** Both support `pg_hint_plan` to override planner decisions. However, Percona's tight integration with PMM makes identifying sub-optimal plans and injecting plans significantly easier during active troubleshooting.

3. Cost Analysis & Sizing in Malaysia (ap-southeast-5)

To provide an accurate cost comparison, we model both architectures inside the **AWS Malaysia region (ap-southeast-5)** utilizing current On-Demand pricing.

We evaluate two standard sizing tiers:

1. **Baseline Tier (Small Business / Staging):** ~8GB RAM, 2 vCPU, 50GB storage.
2. **High-Performance Tier (Enterprise Production):** ~16GB RAM, 4 vCPU, 100GB storage.

Note: Calculations assume 730 run-hours per month and an exchange rate baseline of 1 USD ≈ 4.50 MYR.

Baseline Cost Model

Sizing Profiles:

- **AWS RDS:** Multi-AZ db.m6g.large (2 vCPU, 8GB RAM) + 50GB gp3 Multi-AZ storage.
- **Percona on EC2:** 2x t4g.large EC2 instances (2 vCPU, 8GB RAM, one Primary, one Replica) + 1x t4g.micro EC2 instance in a private subnet running etcd/DCS for cluster consensus. 50GB gp3 storage per database host.

Sizing & Cost Component	Option A: AWS RDS (Multi-AZ)	Option B: Percona Server on EC2 (HA)	Cost & Technical Analysis
Compute / Host Cost	\$221.92 / month • db.m6g.large Multi-AZ hourly rate (\$0.304/hr)	\$104.25 / month • 2x t4g.large (\$49.06/mo each) • 1x t4g.micro etcd (\$6.13/mo)	EC2 is significantly cheaper as RDS charges a managed database premium for licensing, automatic configuration, and Multi-AZ replication.
Storage (gp3)	\$11.50 / month • 50GB Multi-AZ storage (\$0.23/GB-mo)	\$8.00 / month • 2x 50GB gp3 storage volumes (\$0.08/GB-mo each)	RDS Multi-AZ storage rates are higher than standard gp3 volumes on EC2 because the synchronous block mirroring process is managed by AWS.
Backup Costs	\$4.75 / month • RDS automated snapshot storage (~50GB)	\$1.15 / month • pg_backrest compressed backup to S3 (~50GB standard bucket)	pg_backrest performs block-level deduplication and compression, significantly reducing the storage footprint on S3.
Monitoring Tooling	\$0.00 (Included basic Performance Insights)	\$0.00 (PMM is open-source and self-hosted)	Performance Insights is included, while PMM can be hosted on a shared utility instance at no extra cost.
Day-2 Operations / Human Maintenance	\$150.00 / month • Estimated 2 engineering hours/mo of DBA administration	\$1,500.00 / month • Estimated 15 engineering hours/mo (Patroni upgrades, OS patching, backups verification)	Crucial Difference: While Percona on EC2 saves on raw AWS resources, it introduces significant human maintenance costs.
TOTAL MONTHLY COST (USD)	\$388.17 USD	\$1,613.40 USD (Including labor) \$113.40 USD (Raw infrastructure only)	For baseline setups, the engineering labor required to manage a high-availability cluster makes RDS the more cost-effective choice.
TOTAL MONTHLY COST (MYR)	~RM 1,746.77 MYR	~RM 7,260.30 MYR (Including labor) ~RM 510.30 MYR (Raw infrastructure)	(Calculated at 1 USD ≈ 4.50 MYR)

High-Performance Cost Model

Sizing Profiles:

- **AWS RDS:** Multi-AZ db.m6g.xlarge (4 vCPU, 16GB RAM) + 100GB gp3 Multi-AZ storage.
- **Percona on EC2:** 2x t4g.xlarge EC2 instances (4 vCPU, 16GB RAM) + 1x t4g.micro EC2 instance for DCS quorum. 100GB gp3 storage per database host.

Sizing & Cost Component	Option A: AWS RDS (Multi-AZ)	Option B: Percona Server on EC2 (HA)	Cost & Technical Analysis
Compute / Host Cost	\$443.84 / month • db.m6g.xlarge Multi-AZ hourly rate (\$0.608/hr)	\$202.35 / month • 2x t4g.xlarge (\$98.11/mo each) • 1x t4g.micro etcd (\$6.13/mo)	The compute cost savings on EC2 increase with larger instance families, making self-managed setups attractive for raw compute efficiency.
Storage (gp3)	\$23.00 / month • 100GB Multi-AZ storage (\$0.23/GB-mo)	\$16.00 / month • 2x 100GB gp3 storage volumes (\$0.08/GB-mo each)	Raw EBS volumes on EC2 remain significantly more cost-effective than managed Multi-AZ RDS storage.
Backup Costs	\$9.50 / month • RDS automated snapshot storage (~100GB)	\$2.30 / month • pg_backrest compressed backup to S3 (~100GB standard bucket)	Compression and differential backup strategies on S3 reduce long-term snapshot storage costs.
Monitoring Tooling	\$0.00 (Performance Insights)	\$0.00 (Self-hosted PMM)	Basic Performance Insights is included, while PMM provides advanced query performance tracking.
Day-2 Operations / Human Maintenance	\$150.00 / month • Estimated 2 engineering hours/mo of DBA administration	\$1,500.00 / month • Estimated 15 engineering hours/mo (Patroni maintenance, clustering, state checks)	Human labor remains the dominant cost factor for self-managed architectures.
TOTAL MONTHLY COST (USD)	\$626.34 USD	\$1,720.65 USD (Including labor) \$220.65 USD (Raw infrastructure only)	Even at larger scales, the cost of specialized engineering labor to maintain a reliable HA cluster makes RDS highly competitive.
TOTAL MONTHLY COST (MYR)	~RM 2,818.53 MYR	~RM 7,742.93 MYR (Including labor) ~RM 992.93 MYR (Raw infrastructure)	(Calculated at 1 USD ≈ 4.50 MYR)

4. Operational and Strategic Comparison

Operational Area	Option A: AWS RDS (Multi-AZ)	Option B: Percona Server on EC2	Strategic Decision Guidance
Setup & Provisioning	Instant (Terraform/Console) Provisioned inside secure private subnets within minutes.	Complex Requires configuring OS layers, security groups, etcd clusters, Patroni playbooks, and pg_backrest backups.	Use RDS for faster time-to-market and to avoid operational overhead.
Scalability (Storage & Compute)	Storage Autoscaling & Direct Instance Upgrades Increases volume size dynamically and changes instance classes with minimal downtime.	Manual / Complex Requires resizing EBS volumes, executing OS-level file system growth commands, and scaling instances manually.	RDS provides a highly flexible platform for unpredictable or rapidly scaling workloads.
Minor & Major Upgrades	Fully Automated Upgrades are handled via maintenance windows or simple Terraform version adjustments.	Self-Managed Requires carefully planning node-by-node updates to prevent cluster state mismatches.	RDS is ideal for teams without a dedicated Database Reliability Engineer (DBRE).
Vendor Lock-in	Moderate Uses open-source PostgreSQL APIs but relies on AWS-specific infrastructure and APIs.	None Completely open-source. The entire stack can be moved to on-premises or other cloud providers without modifications.	Choose Percona on EC2 if you require hybrid portability or multi-cloud compliance.

5. Summary & Recommendation

Choose AWS RDS PostgreSQL 17 if:

1. **You want to minimize operational overhead:** Your engineering team is lean and does not have dedicated DBA resources to manage backups, clustering, or operating system updates.
2. **You need fast deployment:** You need to launch a production-grade database in private subnets within minutes.
3. **You want automated maintenance:** You require automatic patch application, effortless scaling, and managed snapshots.

Choose Percona Server for PostgreSQL 17 on EC2 if:

1. **You need absolute performance optimization:** Your application requires custom filesystems (like ZFS/XFS), kernel-level memory tuning (Huge Pages), or connection pooling at scale.
2. **You want advanced query visibility:** You want to leverage the **pg_stat_monitor** extension and **PMM** without being constrained by the limits of Performance Insights.
3. **You want to avoid vendor lock-in:** You require a platform-agnostic architecture that can run identically on-premises, on bare metal, or on other cloud providers.
4. **You have the engineering bandwidth:** You have a dedicated platform team capable of managing High Availability clusters (Patroni, etcd, pg_backrest).

Redis vs. Valkey: Architectural and Cost Comparison

#aws #onprem #valkey #redis #elasticsearch #costing #licensing #compliance

[STRATEGIC FINANCIAL] — 20% AWS Cost Reduction & Licensing Risk Mitigation
[SECURITY & COMPLIANCE] — BSD-3-Clause Open-Source Sovereignty & Isolation
[DEVOPS EXECUTION] — Wire-Compatible Caching Engine

Redis vs. Valkey: Why We Select Valkey on Cloud and On-Premises

Caching and session management are vital for high-performance enterprise applications. For our secure, Multi-AZ 3-tier architecture supporting Spring Boot and AI microservices (RAGFlow, Langfuse), we have selected **Valkey** as our primary, standardized caching engine—both for AWS-managed deployments and local on-premises environments.

This guide details the strategic, technical, licensing, and financial reasons behind this selection, including a deep-dive costing analysis inside the **AWS Malaysia region (ap-southeast-5)** comparing Valkey against legacy Redis OSS.

1. The Strategic Context: The Redis Licensing Shift

On **March 20, 2024**, Redis Labs announced a major licensing shift for Redis OSS (starting with version 7.4). Redis abandoned the highly permissive BSD 3-Clause open-source license, transitioning to a dual-licensing scheme under the **Redis Source Available License v2 (RSALv2)** and the **Server Side Public License v1 (SSPLv1)**.

The Enterprise Risk of the New Redis License

- Non-Permissive Terms:** Under SSPLv1 and RSALv2, if an organization offers Redis as a managed service or packages it in certain commercial offerings, they may be forced to release their entire infrastructure management code under the same license.
- Licensing Compliance Contamination [TS-05]:** For enterprise legal and auditing teams, incorporating SSPL/RSAL software inside commercial products introduces high compliance risks, potential copyleft contamination, and complex audit requirements.
- SaaS Vendor Lock-in & Pricing Hikes:** Transitioning away from open-source allows a single commercial vendor to control pricing, APIs, and access limits, creating a major vendor lock-in risk.

The Birth of Valkey

In response to the Redis license change, the open-source community, led by former Redis maintainers, cloud providers, and enterprise stakeholders (including AWS, Google Cloud, Oracle, Ericsson, and Heroku), established **Valkey** under the **Linux Foundation**.

- Permissive Open-Source:** Valkey remains 100% open-source, licensed under the highly permissive **BSD 3-Clause** license.
- Community-Led Stewardship:** Valkey is not controlled by a single commercial entity. Its roadmap, features, and security patches are developed collaboratively by the industry.
- Long-Term Protection:** This guarantees that Valkey will always remain free from licensing contamination, vendor pricing manipulation, and proprietary restrictions.


2. Wire & Protocol Compatibility: A Drop-in Replacement

Valkey was forked from Redis 7.2.4, which means it inherits complete wire and command-level compatibility with Redis.

- No Code Changes Required:** Valkey supports the exact same Redis serialization protocol (RESP2 and RESP3) and all core Redis commands (strings, hashes, lists, sets, sorted sets, hyperloglogs, bitmaps, and pub/sub).
- Seamless Library Integration:** Existing Redis client libraries—such as **Jedis** and **Lettuce** in Java/Spring Boot, **redis-py** in Python, and **ioredis** in Node.js—connect and operate against a Valkey engine natively without any code adjustments or dependency updates.
- Spring Boot 3.5.12 Ready:** Our Spring Boot application utilizes Spring Data Redis (which sits on top of Lettuce). During deployment, we simply update the connection endpoints to point to our Valkey host, requiring zero application refactoring.

3. AWS Cloud Implementation: Amazon ElastiCache for Valkey

AWS natively supports **Amazon ElastiCache for Valkey** as a first-class managed service. To encourage migration and reflect the open-source cost efficiencies, AWS pricing for ElastiCache for Valkey is **20% lower** than ElastiCache for Redis OSS on-demand rates.



In-Depth AWS Costing & Savings (ap-southeast-5)

The calculations below present a granular, line-item comparative analysis between Redis OSS and Valkey inside the **AWS Malaysia (ap-southeast-5)** region.

Note: Calculations assume 730 hours per month and the standard baseline exchange rate of \$1.00 USD = MYR 4.50.

Tier 1: Baseline / Dev Caching Tier (cache.t4g.micro)

- **Specification:** 1 Node, 0.5 GB RAM, burstable compute.
- **Redis OSS Hourly Rate:** \$0.0160 USD / hour
- **Valkey Hourly Rate (20% Off):** \$0.0128 USD / hour

Caching Engine	Hourly Rate	Est. Monthly Cost (USD)	Est. Monthly Cost (MYR)	Est. Yearly Cost (USD)	Est. Yearly Cost (MYR)
ElastiCache Redis OSS	\$0.0160	\$11.68	RM 52.56	\$140.16	RM 630.72
ElastiCache Valkey	\$0.0128	\$9.34	RM 42.03	\$112.08	RM 504.36
NET SAVINGS (Per Node)	\$0.0032	\$2.34	RM 10.53	\$28.08	RM 126.36

Tier 2: High-Performance / Staging Caching Tier (cache.t4g.medium)

- **Specification:** 1 Node, 3.09 GB RAM, moderate dedicated compute.
- **Redis OSS Hourly Rate:** \$0.0680 USD / hour
- **Valkey Hourly Rate (20% Off):** \$0.0544 USD / hour

Caching Engine	Hourly Rate	Est. Monthly Cost (USD)	Est. Monthly Cost (MYR)	Est. Yearly Cost (USD)	Est. Yearly Cost (MYR)
ElastiCache Redis OSS	\$0.0680	\$49.64	RM 223.38	\$595.68	RM 2,680.56
ElastiCache Valkey	\$0.0544	\$39.71	RM 178.70	\$476.52	RM 2,144.34
NET SAVINGS (Per Node)	\$0.0136	\$9.93	RM 44.68	\$119.16	RM 536.22

Tier 3: Enterprise Production Tier (Multi-AZ cache.t4g.medium Cluster)

- **Specification:** 2 Nodes (1 Primary, 1 Replica with Automatic Failover), 3.09 GB RAM per node.
- **Redis OSS Hourly Rate (2 Nodes):** \$0.1360 USD / hour
- **Valkey Hourly Rate (2 Nodes - 20% Off):** \$0.1088 USD / hour

Caching Engine	Hourly Rate	Est. Monthly Cost (USD)	Est. Monthly Cost (MYR)	Est. Yearly Cost (USD)	Est. Yearly Cost (MYR)
ElastiCache Redis OSS	\$0.1360	\$99.28	RM 446.76	\$1,191.36	RM 5,361.12
ElastiCache Valkey	\$0.1088	\$79.42	RM 357.39	\$953.04	RM 4,288.68
NET SAVINGS (Cluster)	\$0.0272	\$19.86	RM 89.37	\$238.32	RM 1,072.44

Financial Impact of Cloud Valkey

By selecting Valkey over Redis OSS across our multi-environment cloud footprint (1x Dev Node, 1x Staging Node, and 1x Production Cluster):

- **Monthly Cloud Caching Cost (Valkey):** \$9.34 + \$39.71 + \$79.42 = **\$128.47 USD / RM 578.12 MYR**
- **Monthly Cloud Caching Cost (Redis):** \$11.68 + \$49.64 + \$99.28 = **\$160.60 USD / RM 722.70 MYR**
- **Direct Financial Savings:** **\$32.13 USD / RM 144.58 MYR per month** (representing a **20.02% total reduction** in caching OpEx with zero performance loss or code changes).

4. On-Premises Implementation: Why Valkey Reigns Supreme

When deploying our application on-site on physical infrastructure (bare-metal rack servers), using Valkey instead of Redis OSS is equally critical.

```
# Sample Quadlet Configuration: /home/songket/.config/containers/systemd/valkey.container
[Container]
ContainerName=valkey_cache
Image=docker.io/valkey/valkey:8.0.0
PublishPort=6379:6379
Volume=/var/srv/valkey/data:/data:Z
UserNS=keep-id:uid=2001,gid=2001

[Service]
Restart=always
ExecStart=valkey-server --requirepass MasterValkeyPass123 --protected-mode no
```

Key On-Premises Advantages of Valkey

1. Complete Open-Source Compliance & Zero Auditing Penalties

Enterprise on-premises deployments are subject to strict software audits. Running Redis OSS post-version 7.4 inside container environments requires commercial licensing tracking. Valkey’s **BSD 3-Clause** license ensures 100% compliance with zero software tracking overhead or audit penalties.

2. Rootless Podman 5+ & systemd Quadlets Compatibility

Valkey containers are fully compatible with rootless Podman execution and systemd Quadlets. The container runs entirely within an unprivileged user namespace, preventing host-level exploit escalations.

3. Community Innovation & Performance Enhancements

Valkey 8.0 introduced significant performance optimizations (such as intelligent thread scheduling, optimized memory allocation, and dual-active replication) that are not available in Redis OSS. Local on-premises hosts benefit from **up to 15% lower CPU utilization** and higher read/write throughput on identical hardware.

4. Absolute Data Sovereignty (Zero SaaS Dependencies)

Operating Valkey on-premises within isolated VLAN 20/30 private database subnets ensures that sensitive transactional cache data, user session keys, and AI token counters are physically preserved on-site, providing complete compliance with **Malaysian PDPA** regulations.

5. Comprehensive Comparative Matrix

The table below summarizes the architectural, financial, and operational parameters of Redis OSS versus Valkey:

Feature / Dimension	Redis OSS (Post-v7.4)	Valkey (v7.2 / v8.0)	Strategic Decision Driver
Licensing Framework	SSPLv1 / RSALv2 (Proprietary / Source-Available)	BSD 3-Clause (Highly Permissive, Open-Source)	• Avoids corporate licensing audits and copyleft contamination risk [TS-05].
AWS Managed Costing	Baseline Base Price (No discounts)	20% Lower On-Demand Pricing natively on ElastiCache	• Substantial operational expense (OpEx) reduction inside ap-southeast-5.
Upstream Stewardship	Single-vendor (Redis Labs commercial control)	Linux Foundation (Community-led by AWS, Google, Oracle)	• Guarantees long-term open-source freedom and collaborative feature roadmap.
API & Wire Protocol	Standard Redis RESP2 / RESP3 protocols	100% Drop-in RESP2 / RESP3 Wire Compatibility	• Zero code modifications required. Client libraries (Lettuce, Jedis) connect natively.
On-Premises Deployment	Requires commercial tracking or legacy version pinning	Native rootless containers via systemd Quadlets (Valkey 7/8)	• High-performance, zero audit risk, unprivileged execution.
Performance Overheads	Standard performance baseline	Up to 15% better throughput & memory management in Valkey 8	• Highly optimized thread scaling and reduced memory footprint.

6. Security & Isolation Guidelines (Zero-Trust)

Regardless of the caching engine choice, our architecture enforces strict **Zero-Trust Network isolation** at the network layer:

On AWS:

- **Subnet Isolation:** Deployed exclusively within the Private Database Subnets (`private_db_subnet_ids`), with zero route tables mapped to Internet Gateways.
- **Security Group Ingress:** Port 6379 is blocked from all public ingress. The ElastiCache security group allows TCP port 6379 exclusively originating from the active Auto Scaling Group (ASG) security group and Standalone EC2 instances.

On-Premises:

- **VLAN Segmentation:** Hosted on VM-03 (DB/AI VM) and restricted to VLAN 30 (Database VLAN) network ingress.
- **Authentication & Encryption:** Force password authentication via `--requirepass` parameter and restrict access strictly to container-to-container private networks using encrypted TLS channels.

7. Conclusion & Recommendation

OFFICIAL VALKEY DECISION
1. AWS Cloud Selection: APPROVED (Amazon ElastiCache for Valkey) - Rationale: Identical performance, BSD-3 licensed, 20% direct costing discount
2. On-Premises Selection: APPROVED (Valkey 8.0 on Rootless Podman) - Rationale: Zero auditing risk, 100% open-source, unprivileged execution

By standardizing on Valkey across both cloud and on-premises environments, we eliminate commercial licensing risks, significantly reduce database operations costs, and establish a high-performance caching layer aligned with modern open-source compliance standards.

Deep State of Mind (DSOM) For My AI Protocol	Harisfazillah Jamel (LinuxMalaysia)	2026-08-12 Standard: UK English	DBP-standard Bahasa Melayu Malaysia (Piawai)	GNU General Public License v3.0
---	-------------------------------------	---------------------------------	--	----------------------------------

RAGFlow + Langfuse: AWS vs. On-Premises GPU Integration Guide

#aws #cloud #architecture #vpc #alb #asg #rds #waf #elasticsearch #valkey #ssl #disaster-recovery #postgresql #gpu #ragflow #langfuse #sovereignty #compliance

Deploying RAGFlow & Langfuse: AWS vs. On-Premises GPU Integration Guide

1. Executive Summary

This guide provides a comprehensive architectural and economic analysis of deploying **RAGFlow** (an advanced, layout-aware Retrieval-Augmented Generation engine) paired with **Langfuse** (an open-source LLM engineering and observability platform).


Given that RAGFlow relies heavily on deep document understanding (DeepDoc), optical character recognition (OCR), layout-aware parsing, and embedding vectorization, the choice of underlying compute architecture—specifically **GPU acceleration**—directly impacts ingestion latency, retrieval precision, and query response times.

Furthermore, this guide details how organizations can achieve a highly cost-effective, high-performance hybrid model. Under this paradigm, secure transactional systems run within **AWS Asia Pacific (Malaysia) Region (ap-southeast-5)**, while heavy GPU workloads are processed on-premises or vice versa, interconnected seamlessly via secure **REST APIs** and the **Model Context Protocol (MCP)**.

2. Technology Overview

RAGFlow Architecture

RAGFlow is an open-source RAG engine based on **deep document understanding**. Unlike naive RAG systems that split text strictly by character count, RAGFlow uses visual and layout-aware machine learning models to extract structures, tables, and figures from complex documents (PDFs, Word docs, spreadsheets).



The core backend services are highly decoupled:

- RAGFlow API & Task Executor:** Python services managing workflows and document parsing.
- DeepDoc Parser:** Implements a multi-stage machine learning pipeline (Layout recognition, Table Structure Recognition [TSR], and OCR).
- Storage Tier:** Object storage (MinIO or S3) for original files, and a database (MySQL or PostgreSQL) for transactional metadata.
- Search/Vector Tier:** Elasticsearch or Infinity (an AI-native database) for full-text search and dense vector recall.
- Caching Tier:** Valkey or Redis for task-broker message queuing and session caching.

Langfuse Architecture


Langfuse is a purpose-built LLM engineering platform that provides:


- Observability & Tracing:** Near real-time tracing of complex multi-step chains, tool calls, vector retrievals, and generation steps.
- Prompt Management:** Version-controlled centralized prompt repository.
- Evaluation & Analytics:** Quantifying LLM latencies, token consumption, and retrieval relevance (e.g., faithfulness, answer relevance).

RAGFlow natively integrates with Langfuse via simple API configurations, allowing operators to trace every document chunk retrieved, prompt assembled, and token generated by downstream LLMs.

3. The Critical Role of GPUs in RAGFlow: Why It Matters

Deploying RAGFlow without a GPU represents a significant architectural bottleneck. The table below outlines how specific sub-systems within RAGFlow utilize hardware and why CPU-only fallbacks are highly inefficient in production.





A. DeepDoc Layout & Visual Parsing (ONNX Runtime)

RAGFlow uses a combination of visual computer vision models (such as YOLOv8 for layout detection, TSR for table structure analysis, and PP-OCRv4/v5 for text extraction).

- **GPU Utility:** These models run via ONNX Runtime using the `CUDAExecutionProvider`.
- **CPU Bottleneck:** Running visual detection models on standard vCPUs causes severe thread saturation, pushing CPU usage to 100% and extending document chunking speeds from **seconds per page to minutes per page**.

B. Embedding Generation & Re-ranking (PyTorch / CUDA)

Vector search requires converting text chunks into multi-dimensional dense vectors (using models like BGE-M3 or local HuggingFace embeddings) and re-evaluating candidate matches using cross-encoder re-rankers.

- **GPU Utility:** Vectorizing large batches of documents is an intensely parallelizable matrix-multiplication operation suited perfectly for GPU Tensor Cores.
- **CPU Bottleneck:** Bulk document ingestion (e.g., a 10,000-page knowledge base) will take days on general-purpose CPUs compared to **minutes or hours on an enterprise GPU**.


C. Local LLM Execution (Ollama / vLLM)

If the organization runs local, self-hosted LLMs (e.g., Qwen-2.5-Instruct, Llama-3-8B) to preserve absolute data sovereignty and avoid third-party API costs.

- **GPU Utility:** Modern LLMs require low-latency KV-caching and rapid auto-regressive decoding, which can only be achieved with high GPU memory bandwidth (VRAM).
- **CPU Bottleneck:** Running an 8B parameters model on a high-spec CPU results in a generation rate of 1 to 3 tokens per second, which is completely unusable for interactive production applications (whereas a GPU easily outputs 50+ tokens per second).

4. AWS-Native Deployment Architecture

Deploying RAGFlow and Langfuse natively in AWS maximizes durability, automatic scaling, and reduces administrative overhead by leveraging managed services.





Components

1. **Network Entrypoint:** AWS WAFv2 filters web-attacks before routing requests through a Multi-AZ **Application Load Balancer (ALB)**.
2. **Compute Tier (Stateless Web):** Langfuse (Node.js/Next.js) and RAGFlow front-ends are deployed on standard, cost-efficient AWS Auto Scaling Groups (ASGs).
3. **Compute Tier (GPU Processing):** Dedicated EC2 instances (specifically **G5** or **G6** families) are deployed in private subnets to run the RAGFlow task executor, ONNX OCR execution, and local LLM/embedding inference.
4. **Database Tier:** Managed **Amazon RDS PostgreSQL (Multi-AZ)** acts as the central transactional database for both Langfuse metadata and RAGFlow document status.
5. **Caching & Queue:** Managed **Amazon ElastiCache for Valkey** provides high-throughput, low-latency key-value queuing.
6. **Storage Tier:** Durability-guaranteed **Amazon S3** replaces local MinIO for holding parsed raw documents and assets.

5. On-Premises Deployment Architecture

When running on-premises, teams rely on a consolidated, cost-efficient hardware stack utilizing standard container orchestration tools.





Components

- 1. **Hardware Host:** A high-end workstation or rack-mounted server running Ubuntu 24.04/26.04 LTS equipped with at least one NVIDIA GPU (minimum 16GB VRAM, recommended 24GB+ VRAM).
- 2. **GPU Driver & Runtime Layers:** Native NVIDIA Driver (e.g., v550+), **CUDA Toolkit (12.1 or 12.4)**, and **NVIDIA Container Toolkit** to allow Docker containers to access GPU hardware directly.
- 3. **Orchestration Layer:** **Docker Compose** or a single-node **Kubernetes (K3s)** cluster.
- 4. **Storage Engine:** Local NVMe SSD arrays configured in RAID-1 or RAID-10 for low-latency indexing read/writes.
- 5. **Services Consolidation:** PostgreSQL, Valkey/Redis, MinIO, Elasticsearch, Langfuse, and RAGFlow running as interconnected container services within a isolated Docker bridge network.

6. Head-to-Head Comparison: AWS vs. On-Premises

To support technical leadership in making informed architectural choices, this section analyzes the capital and operating realities of both deployment environments, specifically calibrated for **AWS Asia Pacific (Malaysia) Region (ap-southeast-5)** in both **USD** and **MYR** (assuming 1 USD = 4.50 MYR).

Side-by-Side Matrix

Evaluation Dimension	AWS-Native (ap-southeast-5)	On-Premises GPU Server
Financial Classification	OpEx (Pay-as-you-go, monthly billing)	CapEx (Upfront hardware purchase + power/cooling)
GPU Instance Scarcity	High. GPU instances (g5, p4) often require AWS limit increases or suffer from low regional capacity in Malaysia.	None. Hardware is fully owned and dedicated to the organization.
Ingress/Egress Costs	Egress from AWS to internet is \$0.09 per GB (after 100 GB/mo free). Bulk PDF downloading incurs costs.	None. Data transfer within the local area network is entirely free and unmetered.
Scalability	Near-instantaneous scaling of compute node ASGs. Elastic storage (S3) scales infinitely.	Hardware-constrained. Scaling requires purchasing additional physical GPUs or servers.
Maintenance Overhead	Extremely low. AWS handles RDS Multi-AZ backups, hardware failures, and network routing.	High. Internals teams must handle hardware warranties, disk failures, cooling, and UPS backup.
Sovereignty & Air-Gap	Shared Responsibility model. Subject to cloud compliance policies.	Absolute. Can be run completely air-gapped without any public internet connectivity.

Financial Model

We present a comparative financial projection over a standard **12-month operating cycle**.

The AWS model runs a high-performance, resilient RAGFlow + Langfuse deployment with a single GPU node (g5.xlarge containing an NVIDIA A10G 24GB VRAM). The On-Premises model assumes the purchase of a professional, enterprise-grade AI workstation equipped with an NVIDIA RTX 4090 (24GB VRAM).

12-Month Total Cost Comparison (USD)	
AWS Cloud (OpEx)	\$15,313.20
On-Premises (CapEx)	\$6,400.00

AWS 12-Month Total Cost Breakdown

- 1. **GPU Compute Instance (g5.xlarge - 1 Node):**
 - Hourly on-demand rate: \$1.15 / hr
 - Monthly cost: 730 hours * \$1.15 = \$839.50 USD
 - **12-Month Total: \$10,074.00 USD (RM 45,333.00 MYR)**
- 2. **Auxiliary Compute (ASGs for Web & Langfuse):**
 - 2 x t4g.medium (ARM64, 2 vCPU, 4GB RAM) = 2 * \$0.0384/hr * 730 hrs = \$56.06 USD / mo
 - **12-Month Total: \$672.72 USD (RM 3,027.24 MYR)**
- 3. **Managed Storage (Amazon S3 + RDS Multi-AZ PostgreSQL):**
 - S3 Storage (200 GB) + RDS Postgres (Multi-AZ db.t4g.medium, 100GB gp3) ≈ \$180.00 USD / mo
 - **12-Month Total: \$2,160.00 USD (RM 9,720.00 MYR)**
- 4. **Network Fees (WAF, NAT Gateway, Egress, ALB):**
 - Shared ALB, WAF rules, NAT base, and 200 GB egress ≈ \$200.54 USD / mo
 - **12-Month Total: \$2,406.48 USD (RM 10,829.16 MYR)**
- **AWS 12-MONTH COMBINED TOTAL: \$15,313.20 USD / ~RM 68,909.40 MYR**

On-Premises 12-Month Total Cost Breakdown


- 1. **Hardware Purchase (One-Time CapEx):**
 - AMD Threadripper or high-end Intel CPU, 128GB RAM, 2TB NVMe Gen4 SSD, Dual-chamber cooling tower.
 - GPU: NVIDIA RTX 4090 (24GB VRAM, Tensor Cores).
 - **Upfront cost: \$5,200.00 USD (RM 23,400.00 MYR)**

- 2. **Datacenter Rack Power & Cooling (OpEx):**
 - Avg draw: ~650 Watts under active load.
 - Electricity cost in Malaysia (commercial rates of approx. RM 0.39 per kWh) + standard cooling overhead ≈ \$60.00 USD / mo
 - **12-Month Total: \$720.00 USD (RM 3,240.00 MYR)**
 - 3. **Internal Systems Engineer Maintenance (Allocated Labor):**
 - Fractional engineering time for OS updates, hardware checks, backups ≈ \$40.00 USD / mo
 - **12-Month Total: \$480.00 USD (RM 2,160.00 MYR)**
- **ON-PREMISES 12-MONTH COMBINED TOTAL: \$6,400.00 USD / ~RM 28,800.00 MYR**

7. Hybrid Integration: API and MCP Architectures

To minimize cloud costs while maintaining enterprise-grade resilience, organizations can adopt a **Hybrid Deployment Model**.

This approach isolates standard web frontends, authentication portals, and relational metadata within secure, scalable **AWS private subnets**, while routing compute-heavy OCR, layout analysis, and local LLM/embedding inference requests to highly cost-efficient **On-Premises GPU Workstations** situated in corporate datacenters or Cyberjaya offices.



The hybrid model supports two primary secure communication patterns:

Pattern A: Secure REST API Integration

In this pattern, RAGFlow’s core engine and databases run in AWS, but the heavy task execution worker (RAGFlow Worker daemon) is deployed on-premises.

- **Execution Flow:**
 1. A user uploads a 500-page PDF to the AWS Web Portal. The file is saved directly to **Amazon S3**.
 2. The database logs the parsing task in the **AWS RDS PostgreSQL** queue.
 3. The local on-premises GPU server runs a secure RAGFlow Worker client which continuously polls the AWS database queue (via secure outbound SSL connections or private Site-to-Site VPN).
 4. The on-premises worker pulls the PDF from the secure Amazon S3 pre-signed URL, processes it using local **NVIDIA Tensor Cores** (running YOLOv8 layout and PP-OCR), and uploads the structured vector embeddings back to AWS Elasticsearch/Infinity.
 5. The final trace data is pushed from RAGFlow to the AWS-hosted **Langfuse** server to log execution timelines and parsing latencies.
- **Security Control:** The on-premises workstation needs **no inbound open ports**. It establishes outbound-only HTTPS connections to Amazon S3, RDS, and Langfuse, bypassing corporate firewall ingress limitations.

Pattern B: AI-Native Model Context Protocol (MCP) Integration

The **Model Context Protocol (MCP)** allows AI Agents running on AWS (such as **Amazon Bedrock AgentCore** or custom LangChain/Langgraph agents) to dynamically call on-premises resources, search internal databases, or request local GPU parsing.

- **Execution Flow:**
 1. An enterprise AI Agent running in the AWS Cloud receives a query: *“Compare our quarterly financial PDF with the structural charts inside our local Cyberjaya archives.”*
 2. The agent identifies that the Cyberjaya archives reside on-premises and requires local OCR table-extraction capabilities.
 3. The agent initiates an outgoing call via **Amazon API Gateway configured with MCP Proxy Support**.
 4. API Gateway translates the request into a standardized JSON-RPC 2.0 payload, routing it over an mTLS-secured connection to the on-premises **Local MCP Server**.
 5. The on-premises MCP Server accesses the local RTX 4090 GPU to execute RAGFlow’s DeepDoc visual parser on the requested archival charts.
 6. The structured JSON tables and charts are returned back through API Gateway to the AWS Bedrock Agent, which synthesizes the final response.
- **Tracing Compliance:** Because RAGFlow and Langfuse are fully integrated, the complete round-trip trace—including the AWS Bedrock invocation, the API Gateway transit, the on-premises MCP tool execution, and the final GPU-accelerated generation—is visualised inside the **Langfuse Dashboard** as a unified trace graph.

8. Practical Integration & Setup Guide

1. Collect Langfuse Credentials

Before configuring RAGFlow, set up a Langfuse project and acquire the API credentials:

- Sign in to your Langfuse dashboard.

- Navigate to **Settings ► Projects** and select or create a project.
- Generate a new set of API keys. Copy the following variables:
 - LANGFUSE_PUBLIC_KEY
 - LANGFUSE_SECRET_KEY
 - LANGFUSE_HOST (e.g., <https://cloud.langfuse.com> or <http://<your-aws-langfuse-ip>:3000>)

2. Configure RAGFlow to Emit Traces

RAGFlow includes a native Langfuse integration. Key credentials are set per tenant in the RAGFlow ecosystem. This can be configured dynamically through the Web UI:

1. Log in to your RAGFlow web console.
2. Click on your profile avatar in the top-right corner and select **API**.
3. Scroll to the bottom of the API screen to locate the **Langfuse Configuration** section.
4. Fill in the following details:
 - **Host:** Your Langfuse API host URL.
 - **Public Key:** Your LANGFUSE_PUBLIC_KEY.
 - **Secret Key:** Your LANGFUSE_SECRET_KEY.
5. Click **Save**.

Once saved, RAGFlow automatically wraps its internal pipeline chains, retrievals, and generation steps, emitting telemetry directly to your Langfuse workspace on every query or document retrieval.

3. Setting GPU Memory Limits for OCR & Parsers

When self-hosting RAGFlow on-premises or on AWS GPU instances, you can configure how visual parsing tasks distribute memory across GPUs. Add the following environment variables to your RAGFlow `.env` configuration file:

```
# Allow RAGFlow to access all available local CUDA devices (e.g. 0,1 for dual-GPU)
PARALLEL_DEVICES=""

# Limit maximum GPU memory dedicated to the PP-OCR visual engine to prevent out-of-memory (OOM) crashes
OCR_GPU_MEM_LIMIT_MB=4096

# Define model structure caching strategy (extend vs. shrink)
OCR_ARENA_EXTEND_STRATEGY="extend"
```

Restart your Docker Compose service to apply limits:

```
docker compose -f docker/docker-compose.yml up -d
```

4. Viewing Unified Traces in Langfuse

After executing a query in RAGFlow (e.g., asking a question in the chat portal), open your Langfuse Project Dashboard:

- Navigate to **Traces**.
- You will see incoming streams prefixed with `ragflow-*`.
- Expanding a trace reveals a nested tree of spans:
 - `ragflow-retrieval`: Logs how long Elasticsearch/Infinity took to execute hybrid keyword/dense-vector queries, accompanied by exact parsed document chunk payloads.
 - `ragflow-rerank`: Quantifies the cross-encoder execution latency.
 - `ragflow-generation`: Visualizes prompt template construction and logs downstream LLM token-generation speeds.

Google Antigravity Skills Integration Guide

#aws #cloud #architecture #rds #elasticache #valkey #disaster-recovery #antigravity #skills

Google Antigravity Skills Integration Guide

Welcome to the **Google Antigravity Skills Integration Guide** for the AWS Secure 3-Tier Web & AI Infrastructure project. This document outlines how to integrate and share agent knowledge, capabilities, and workflows between different AI assistants—specifically **Google Jules** and **Google Antigravity**—using the open **Agent Skills** ecosystem standard.

By consolidating project-specific operating instructions and architectural boundaries into standard skill specifications, both Google Jules and Google Antigravity can operate with identical high-fidelity context, rulesets, and deployment tools.

1. What are Agent Skills?

Agent Skills represent an open, decentralized standard for extending the capabilities and contextual boundaries of AI agents. Instead of relying purely on generalized system prompts, agents can dynamically discover and read specialized instruction packages tailored for specific repositories, codebases, or workflows.

An Agent Skill is packaged as a standard directory containing a `SKILL.md` file. This markdown file begins with structured YAML front-matter identifying its metadata:

```
---
name: my-skill-name
description: Clear, high-level description of when to apply this skill.
---
# Detailed Instructions
Step-by-step procedures, CLI commands, constraints, and guidelines go here.
```

2. Where Skills Live: Directory Scopes

Google Antigravity supports both workspace-specific and global scopes for discoverability:

Scope	Location Path	Use Case
Workspace / Project	<workspace-root>/agents/skills/	Project-specific rulesets, database ports, architectural limits, deployment guides, and custom CLI wrappers.
Global User Scope	~/gemini/config/skills/	Personal utility commands, code formatter standards, and general-purpose tooling.

Note: While Antigravity natively prioritizes .agents/skills, it maintains full backward-compatibility for .agent/skills as well.

3. The jules-knowledge Skill

We have defined a dedicated workspace skill for our repository to bridge the knowledge bases of Google Jules and Google Antigravity. This skill is located at:

```
.agents/skills/jules-knowledge/SKILL.md
```

Purpose of jules-knowledge

This skill ensures that whenever you invoke Google Antigravity (using the CLI tool `agy`) or interact with Google Jules in this workspace, the AI assistant automatically:

- 1. Targets **AWS Asia Pacific (Malaysia)** (ap-southeast-5) by default.
- 2. Uses **ARM64 Graviton instances** (t4g.micro for compute, db.t4g.micro for RDS, and cache.t4g.micro for Valkey).
- 3. Enforces **Ubuntu 26.04 LTS** hardened under the ASIMP framework.
- 4. Protects database ingress via port-isolation (blocking direct public routing to port 5432).
- 5. Deploys secure session layers with **Amazon ElastiCache for Valkey**.
- 6. Follows the mandatory validation pipeline (`./scripts/deploy.sh` and `python scripts/prepare_docs.py`).

4. How to Install and Use Skills with npx skills

`npx skills` is a command-line tool developed by Vercel Labs acting as an open package manager for AI agents (including Antigravity, Claude Code, GitHub Copilot, Cursor, and Cline).

Downloading and Syncing Skills

Running `npx skills` packages can place skills in your local user’s global configuration (`~/gemini/antigravity-cli/skills/`). To ensure Google Antigravity picks them up properly:

- 1. **For Workspace Scopes:** Copy the desired skill directory to your project root under `.agents/skills/`:
`cp -r ~/agents/skills/some-skill .agents/skills/`
- 2. **For Global Scopes (Google Antigravity):** Copy the skill directory to:
`cp -r ~/agents/skills/some-skill ~/gemini/antigravity-cli/skills/`

5. Seamless Bridge Between Jules and Antigravity

By maintaining our core engineering principles in `.agents/skills/jules-knowledge/SKILL.md` and keeping `AGENTS.md` aligned, we establish a robust bi-directional capability:

- **Google Jules (This Agent):** Uses `AGENTS.md` and this guide to understand and write code conforming to enterprise standards, verifying actions using read-only verification commands.
- **Google Antigravity (The CLI / Local Agent):** Reads the compiled skills within `.agents/skills/` at startup. Typing `/jules-knowledge` or initiating an action matching our database/infrastructure setup will trigger the detailed instructions embedded in the skill.

This unified knowledge strategy guarantees zero-configuration context sync, eliminating environment drift and ensuring both assistants deliver perfectly aligned, region-compliant code.

SOP: Knowledge-First Discovery & Context Preservation Protocol


#okf #discovery #context-management #brain #dsom #SOP #aws #cloud

SOP: Local Knowledge-First Discovery & OKF Context Protocol

1. Executive Intent

To prevent unnecessary API or CLI queries, remote SSH or Session Manager probes, token window exhaustion, and context loss during agentic sessions, all AI agents must strictly adhere to the **Local Knowledge-First Protocol**. All project facts, architecture details, subnet IPs, security rules, and costing options are documented locally via **OKF v0.1 YAML Frontmatter** inside `.agents/brain/` and `docs/`.

2. Standard Operating Procedure (4-Step Discovery Flow)



Step 1: Local Frontmatter & Metadata Search

Before issuing any AWS CLI command, checking AWS SSM sessions, querying live OpenTofu state, or executing remote scripts against standalone EC2 or active ASG nodes:

- 1. Search local OKF frontmatter for relevant topics: or description: keywords:
 - Example: Search for “valkey” or “asg” or “postgres” or “dr-options” inside the local files to identify existing design guides and architectural layouts.
- 2. Read `.agents/brain/active_context_manifest.md` to see the current session checkpoints.

Step 2: Targeted File Viewing

Once the relevant document is located via OKF frontmatter or titles:

- Read specific segments of the matched local file using `read_file` to keep the context window lean and avoid loading entire huge files unnecessarily.

Step 3: Local Context Synthesis

Map the local specifications (e.g. ap-southeast-5 regional defaults, instance sizes like t4g.micro, security groups ingress definitions) to the task at hand. The local repository is the **Single Source of Truth (SSOT)** for configuration intent.

Step 4: AWS Runtime Execution Gate

Probing live AWS resources, executing OpenTofu plans, querying remote database ports, or starting SSM Session Manager connections are authorized **ONLY** when:

- Applying actual infrastructure or configuration updates (e.g., via `./scripts/deploy.sh`).
- Fetching actual live runtime telemetry (e.g., live database statistics or running target group health status) that cannot be modeled statically in documentation.

3. Mandatory Rules Reference

- **Rule 6 (OKF Topics):** All `.md` files must open on line 1 with `---` and contain `topics: [3-5 keywords]` and metadata.
- **Rule 10 (Metadata-First Discovery):** Read targeted file segments to preserve token efficiency.
- **Rule 29 (Local Knowledge-First Mandate):** Always exhaustively search local `.agents/brain/` and `docs/` before digging into remote servers or external web search.

Wazuh Standalone Cloud Installation & Costing Guide

#aws #cloud #architecture #wazuh #siem #xdr #security #costing #ec2

Wazuh Standalone Cloud Installation & Costing Guide

Wazuh is an enterprise-grade, open-source security platform providing unified **XDR (Extended Detection and Response)** and **SIEM (Security Information and Event Management)** capabilities. It monitors endpoints, cloud services, and containers, analyzing security events in real-time.

To maintain strict cost isolation and modular design, this guide details how to deploy Wazuh on a **standalone EC2 instance** completely separate from any Auto Scaling Group (ASG), and outlines the **best and cheapest cloud-deployment models** with a dedicated financial breakdown.

1. Wazuh Central Architecture & Sizing

The Wazuh central components consist of three primary elements:

- 1. **Wazuh Indexer:** A highly scalable, full-text search and analysis engine that indexes and stores alerts generated by the Wazuh server.
- 2. **Wazuh Server:** The management node that receives and analyzes data from agents, applying rules and raising alerts.
- 3. **Wazuh Dashboard:** The web user interface for mining, visualizing, and configuring security alerts and platform features.

Sizing and Hardware Guidelines

Wazuh’s resource utilization is directly proportional to the event ingestion rate (Events Per Second, or EPS) and the number of active agents.

According to the [Wazuh Quickstart Documentation](#):

- **Official Recommended Sizing (1–100 agents):** 4 vCPUs, 8 GiB RAM, and 50–200 GB SSD storage.
- **Cheapest Standalone Feasibility (1–25 agents):** For testing or low-traffic small environments, Wazuh central components can run successfully on **2 vCPUs and 4 GiB RAM** (e.g., using JVM heap size tuning for OpenSearch/Indexer to prevent out-of-memory states).

2. Separated Costing Models (USD & MYR)

To keep budget planning fully transparent, we analyze four standalone deployment pathways in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)** and compare them with the cheapest global baseline in **US East (N. Virginia, us-east-1)**.

Exchange rate baseline: 1 USD ≈ 4.50 MYR.

Configuration Sizing Matrix

Metric / Resource	Option A: Ultra-Cheap / Dev (Malaysia ap-southeast-5)	Option B: Recommended Cheap Production (Malaysia ap-southeast-5)	Option C: Global Cheap Baseline (US East us-east-1)	Option D: Official AMI Sizing (Malaysia ap-southeast-5)
Instance Type	t4g.medium (ARM64)	t4g.large (ARM64)	t3.medium (x86_64)	c5a.xlarge (x86_64)
vCPU / Memory	2 vCPUs / 4 GiB RAM	2 vCPUs / 8 GiB RAM	2 vCPUs / 4 GiB RAM	4 vCPUs / 8 GiB RAM
Target Workload	1–15 Agents (Dev/Test)	1–35 Agents (Light Prod)	1–15 Agents (Dev/Test)	1–100 Agents (Full Prod)
EBS Storage (gp3)	50 GiB gp3	100 GiB gp3	50 GiB gp3	100 GiB gp3
Network Setup	1 Elastic IP (Static Public IP)	1 Elastic IP (Static Public IP)	1 Elastic IP (Static Public IP)	1 Elastic IP (Static Public IP)

Detailed Cost Breakdown Table (Monthly)

All estimates assume On-Demand usage over a standard month (~730 hours).

Cost Category	Option A (USD / MYR)	Option B (USD / MYR)	Option C (USD Only)	Option D (USD / MYR)
EC2 Compute	\$24.53 / RM 110.39	\$49.06 / RM 220.77	\$30.37	\$136.51 / RM 614.30
gp3 EBS Storage	\$4.00 / RM 18.00	\$8.00 / RM 36.00	\$4.00	\$8.00 / RM 36.00
Public IPv4 Fee	\$3.65 / RM 16.43	\$3.65 / RM 16.43	\$3.65	\$3.65 / RM 16.43
AWS Backup Snapshot	\$2.50 / RM 11.25	\$5.00 / RM 22.50	\$2.50	\$5.00 / RM 22.50
Data Ingress/Egress	Free / Free	Free / Free	Free	Free / Free
TOTAL (Monthly)	\$34.68 / RM 156.07	\$65.71 / RM 295.70	\$40.52	\$153.16 / RM 689.23

Key Cost Efficiency Discoveries

- 1. **Graviton Efficiency (Option A & B):** Using AWS Graviton (ARM64) t4g instances yields a **20% savings** over x86 counterparts (t3) while providing superior CPU performance per dollar.
- 2. **Official Marketplace vs. Custom Script:** Launching the pre-built x86 marketplace image restricts sizing to c5a.xlarge (Option D, ~\$153.16/mo). Deploying on a clean Ubuntu 24.04 Graviton server and using the Wazuh installation assistant (Option B, ~\$65.71/mo) **reduces monthly cloud spend by over 57%!**

3. Installation Strategy 1: Clean OS + Assistant (Best & Cheapest)

Deploying a clean, minimal **Ubuntu 24.04 LTS (Noble Numbat)** or **Ubuntu 22.04 LTS (Jammy Jellyfish)** on a Graviton instance (t4g.medium or t4g.large) and executing the Wazuh installation assistant script is the most cost-effective architecture.

Step 1: Provision the Standalone EC2 Instance

Create a standalone EC2 instance using your OpenTofu templates or the AWS Console:

- **Region:** ap-southeast-5 (Malaysia)
- **AMI:** Ubuntu 24.04 LTS (ARM64 / AARCH64)
- **Instance Type:** t4g.large (2 vCPU, 8 GiB RAM - Highly Recommended)
- **Storage:** 100 GB gp3 volume (3,000 IOPS, 125 MB/s baseline throughput)
- **Elastic IP:** Allocate and associate an Elastic IP with the instance to prevent IP changes on reboot.

Step 2: Configure the Firewall (Security Group)

Create a dedicated Security Group for your standalone Wazuh instance with the following ingress rules:

Port	Protocol	Source	Purpose
22	TCP	Admin Office IP / Jumphost	Secure SSH Administration
443	TCP	Admin Office IP / VPN Subnet	Wazuh Web Dashboard (HTTPS)
1514	TCP	Monitored Servers / Agent Subnets	Wazuh Agent secure communication
1515	TCP	Monitored Servers / Agent Subnets	Wazuh Agent enrollment service
55000	TCP	Admin API client / Integrations	Wazuh RESTful API access

Step 3: Run the All-in-One Installation Assistant

SSH into your instance and run the automated script to install Wazuh indexer, server, and dashboard on the same host:

```
# SSH to the instance
ssh -i "your-key.pem" ubuntu@your-wazuh-elastic-ip

# Switch to root
sudo su -

# Download the installation assistant script
curl -sO https://packages.wazuh.com/4.x/wazuh-install.sh

# Execute the All-in-One installation
bash wazuh-install.sh -a
```

Note: If you are running on t4g.medium (4 GiB RAM), the installer might throw a warning that the hardware requirements are not fully met. You can bypass this check safely for small testing environments by passing the -i flag:

```
bash wazuh-install.sh -a -i
```

Step 4: Record Initial Passwords

At the end of a successful installation, the script will output the default credentials and security configuration. Write them down securely:

```
INFO: --- Installation finished ---
INFO: Admin password: <GENERATED_PASSWORD>
INFO: You can access the web interface at https://<YOUR_IP>
```

 4. Installation Strategy 2: Official AWS Marketplace AMI

Wazuh provides a pre-built Amazon Machine Image (AMI) pre-configured with Amazon Linux 2023, Wazuh manager, Filebeat-OSS, Wazuh indexer, and Wazuh dashboard out of the box.

According to the official [Wazuh AMI Documentation](#):

Step 1: Subscribe in AWS Marketplace

1. Navigate to [Wazuh All-In-One Deployment in the AWS Marketplace](#).
2. Click **View purchase options**, review the software terms, and click **Subscribe** to accept.
3. Once accepted, select **Launch your software** and choose your region (e.g. ap-southeast-5).

Step 2: Launch the Instance via EC2 Console

1. Navigate to the **EC2 Launch Instance** interface.
2. Under AMIs, go to **Browse more AMIs**, select **AWS Marketplace AMIs**, and search for Wazuh All-In-One Deployment.
3. Select **c5a.xlarge** (4 vCPU, 8 GiB RAM) or a comparable instance class.
4. Set storage capacity to at least **100 GiB gp3**.
5. Select or create a security group ensuring the standard ports (443, 1514, 1515, 22) are whitelisted as shown in Strategy 1.
6. Select your SSH key pair and launch the instance.

 5. Critical Post-Installation & Security Hardening

To ensure secure enterprise operations, perform the following verification and hardening procedures immediately after launching:

A. Accessing the Dashboard

Open your web browser and navigate to:

https://<YOUR_WAZUH_STATIC_IP>

- **Username:** admin
- **Password (Strategy 1 - Assistant):** The password generated at the end of the script.
- **Password (Strategy 2 - AMI):** The EC2 **Instance ID** with the first letter capitalised (e.g. if the instance ID is i-0123456789abcdef0, the password is I-0123456789abcdef0). This safety mechanism prevents initial brute-force attempts.

B. Change Default Passwords

To replace the default master passwords with secure credentials:

```
# Navigate to the tool directory
cd /usr/share/wazuh-indexer/plugins/opensearch-security/tools/

# Run the security administrator script to re-configure passwords
# Refer to Wazuh's Password Management guidelines for specific hashes
```

C. Restrict SSH and Dashboard Ingress

Do not leave port 22 (SSH) and port 443 (Dashboard) open to the public internet (0.0.0.0/0). Whitelist access strictly to your organizational CIDR blocks (such as the Cyberjaya development office IP ranges) or configure access via a secure Bastion/SSM Session Manager.

 6. Enrolling Wazuh Agents

Once the standalone Wazuh server is fully active, security agents can be registered on target hosts.

For Linux Targets (Debian/Ubuntu):

```
# Download and install target agent pointing to your standalone Wazuh server IP
wget https://packages.wazuh.com/4.x/apt/pool/main/w/wazuh-agent/wazuh-agent_4.7.2-1_amd64.deb
sudo WAZUH_MANAGER='<YOUR_WAZUH_STATIC_IP>' dpkg -i wazuh-agent_4.7.2-1_amd64.deb

# Start the agent service
sudo systemctl daemon-reload
sudo systemctl enable wazuh-agent
sudo systemctl start wazuh-agent
```

For Windows Targets (PowerShell):

```
Invoke-WebRequest -Uri https://packages.wazuh.com/4.x/windows/wazuh-agent-4.7.2-1.msi -OutFile ${env:tmp}\wazuh-agent.msi
msiexec.exe /i ${env:tmp}\wazuh-agent.msi /q WAZUH_MANAGER="<YOUR_WAZUH_STATIC_IP>"

Start-Service -Name "Wazuh"
```

Deep State of Mind (DSOM) For My AI Protocol	Harisfazillah Jamel (LinuxMalaysia)	2026-07-26 Standard: UK English	DBP-standard Bahasa Melayu Malaysia (Piawai)	GNU General Public License v3.0
---	-------------------------------------	---------------------------------	--	----------------------------------

Wazuh SIEM & XDR Deep-Dive: Cloud, On-Premises & AV Compatibility

#wazuh #siem #xdr #security #antivirus #windows-defender #cloud #onprem

[SECURITY & COMPLIANCE] — Security Operations & Hardening Guidance
[DEVOPS EXECUTION] — Systems Engineers & SecOps

Wazuh SIEM & XDR Deep-Dive: Cloud, On-Premises & AV Compatibility

Modern IT infrastructure requires unified visibility across all assets—including endpoints, servers, containers, cloud resources, and user identities. Wazuh is an enterprise-grade, open-source security monitoring platform that fulfills this requirement by acting as both a **SIEM** (Security Information and Event Management) and an **XDR** (Extended Detection and Response) platform.

This guide explores the functions of Wazuh, why it is critical for modern secure architectures, its deployment options, and how it coexists with Antivirus (AV) programs, including Microsoft Windows Defender and other third-party products.

1. Wazuh Core Functions & Why We Need It

Wazuh goes beyond traditional log collection. It functions by deploying a lightweight agent on target hosts and forwarding security telemetry to a centralized Wazuh manager cluster. The key features that make Wazuh indispensable include:

A. Centralized Log Analysis & Correlation (SIEM)

Wazuh collects, parses, and normalizes log data from operating systems, applications (web servers, databases), cloud APIs, network firewalls, and containers. The Wazuh manager runs these logs through an extensive, customisable rule engine, correlating events in real-time to alert security operations teams of suspicious activity.

B. File Integrity Monitoring (FIM / Syscheck)

Wazuh tracks changes to system files, directories, registry keys (on Windows), and file attributes in real-time or via scheduled sweeps. FIM detects unauthorized modifications, helping organizations identify rootkits, unauthorized user activities, and potential compromise indicators.

C. Vulnerability Detection

By scanning package managers, system patches, and application versions on enrolled agents, Wazuh correlates local software states with global CVE (Common Vulnerabilities and Exposures) databases. This provides a continuous, automated vulnerability scanner across your entire fleet.

D. Security Configuration Assessment (SCA)

Wazuh runs automated checks against system configurations, comparing them with industry standards such as Center for Internet Security (CIS) benchmarks. SCA ensures that OS settings, SSH configurations, firewall policies, and password requirements are hardened to limit the attack surface.

E. Active Response & Threat Mitigation

Wazuh agents can execute automatic, immediate defensive measures when specified alerts are triggered. These actions include blocking an attacking IP address at the firewall (iptables/UFW/Windows Firewall), killing rogue processes, or isolating a compromised host from the network.

2. Wazuh Deployment Models: Cloud vs. On-Premises

Depending on your organizational requirements, compliance needs, and budget, Wazuh can be deployed in the cloud or as a self-hosted on-premises solution.

Cloud Deployments

Standalone EC2 (Best & Cheapest for AWS)

As detailed in our [Wazuh Standalone Cloud Installation & Costing Guide](#), launching a clean Ubuntu LTS instance using AWS Graviton (such as t4g.large in the Malaysia region ap-southeast-5) and running the All-in-One assistant script is highly cost-effective.

- **Monthly Cost:** ~\$65.71 USD (approx. RM 295.70 MYR).
- **Pros:** Highly customized, ARM64 pricing efficiency, full administrative control.
- **Cons:** Manual setup and maintenance of OS updates.

Official AWS Marketplace AMI

A pre-packaged Amazon Machine Image (AMI) provided by the Wazuh team on c5a.xlarge.

- **Monthly Cost:** ~\$153.16 USD (approx. RM 689.23 MYR).
- **Pros:** Fast provisioning, official baseline configurations.
- **Cons:** Restricted to x86_64 architecture, more expensive compute node.

Wazuh Cloud (SaaS)

Fully managed SaaS offering maintained directly by Wazuh.

- **Pros:** Zero-management backend, automatic updates, high availability managed by the vendor.
- **Cons:** High subscription costs, potential data sovereignty and latency concerns for Malaysian PDPA compliance.

On-Premises Deployments

For organizations operating under strict national sovereignty rules (e.g., Malaysia’s PDPA 2010 and 2025 CBPDT Guidelines), on-premises deployments ensure all telemetry data remains within local physical bounds:

VM-Based Standalone

Deploying the three core components (Wazuh Indexer, Wazuh Manager, and Wazuh Dashboard) onto a single hardened Ubuntu VM or physical server inside a dedicated Management VLAN.

Rootless Podman & systemd Quadlets

For containerized microservices architectures, Wazuh can be deployed on-premises as a multi-container stack. This stack runs under unprivileged user namespaces (UserNS=keep-id:uid=2001,gid=2001) with storage mapped natively to encrypted physical volumes.

3. Antivirus Coexistence & Compatibility

A critical question when deploying endpoint monitoring is: **Does the Wazuh agent conflict with antivirus (AV) software?**

Compatibility with Windows Defender

Yes, Wazuh is 100% compatible with Microsoft Windows Defender.

In fact, Wazuh features native integration with Windows Defender. Instead of competing, they work cooperatively:

1. **Event Log Ingestion:** Wazuh actively reads and parses events from the Microsoft-Windows-Windows Defender/Operational channel.
2. **Coordinated Alerting:** If Windows Defender blocks malware or detects a threat, Wazuh captures the event log, triggers a high-severity alert on the Wazuh Dashboard, and alerts security analysts.
3. **Double Defense:** Defender handles real-time signature matching, quarantine, and local disinfection, while Wazuh correlates Defender’s logs with network telemetry, Active Directory events, and FIM triggers.

Compatibility with Third-Party AV (Symantec, McAfee, CrowdStrike, SentinelOne, etc.)

The Wazuh agent is generally fully compatible with modern EDRs (Endpoint Detection & Response) and traditional AV programs. However, as with any security agent that performs deep system inspection, **certain conflict points can occur if not configured properly.**

Potential Conflict Points:

- **Simultaneous File Access (CPU/Disk I/O Spike):** When Wazuh’s File Integrity Monitoring (FIM / Syscheck) scans system folders, it reads file metadata and hashes (MD5/SHA256). If an active AV simultaneously scans those same files as they are opened, this can lead to severe disk I/O bottlenecks and temporary system performance drops.
- **False Positive Flags (Interception Block):** Because the Wazuh agent monitors processes, checks open ports, and reads raw event logs, a highly aggressive third-party AV or EDR might flag the Wazuh agent’s system activity as malicious behavior, potentially killing the wazuh-agentd service or blocking its connection to the manager.
- **Active Response Mitigation Collisions:** If a third-party AV detects a threat and locks a file, while Wazuh triggers an Active Response script to kill the associated process, the two agents might collide, causing execution failures or corrupting state locks.

Resolution: Mutual Exclusions Configuration

To prevent performance degradation and false positives, administrators must configure **mutual exclusions** on both systems.

Step 1: Configure AV Exclusions for Wazuh

Add the Wazuh installation directories to your third-party AV’s whitelist/exclusion policies so that the AV does not scan Wazuh’s local database files or monitor its binaries:

- **Windows:** C:\Program Files (x86)\ossec-agent\
- **Linux:** /var/ossec/
- **macOS:** /Library/0ssec/

Step 2: Configure Wazuh Exclusions for the AV

In the Wazuh agent configuration (ossec.conf), ensure that Wazuh’s Syscheck (FIM) ignores the active folders of your antivirus software to prevent unnecessary, overlapping scan loops:

```
<!-- Example of Wazuh agent's ossec.conf Syscheck section -->
<syscheck>
  <!-- Ignore local Windows Defender databases and temp files -->
  <ignore>C:\ProgramData\Microsoft\Windows Defender</ignore>

  <!-- Ignore Third-Party Antivirus signature/definition folders -->
  <ignore>C:\ProgramData\McAfee</ignore>
  <ignore>C:\Program Files\CrowdStrike</ignore>
</syscheck>
```

“Passive Mode” Configurations in Wazuh

If your organization requires the third-party antivirus to be the **only active defender** (handling blocklist checks, process kills, and quarantines), you can configure the Wazuh agent to operate in a **purely passive metadata and log aggregation mode**.

To enforce Wazuh “Passive Mode”, adjust these settings in your central or local agent configuration:

1. Disable Active Response

Ensure that Wazuh does not attempt to execute automatic blocks, file removals, or process terminations, leaving response actions entirely to the AV:

```
<active-response>
  <disabled>yes</disabled>
</active-response>
```

2. Tune File Integrity Monitoring (Syscheck)

Instead of real-time file monitoring, which constantly hooks file changes, switch Syscheck to scheduled, low-priority scanning (e.g., once every 12 or 24 hours during off-peak times) to eliminate disk I/O collisions:

```
<syscheck>
  <scan_on_start>no</scan_on_start>
  <frequency>86400</frequency> <!-- Scan once per day (in seconds) -->
  <realtime>no</realtime> <!-- Disable real-time monitoring -->
</syscheck>
```

3. Focus Strictly on Event Log and Process Aggregation

Configure the agent to only ingest application logs and OS-generated security events, letting the AV handle active on-access file scanning:


```
<localfile>
  <log_format>eventlog</log_format>
  <location>Security</location>
</localfile>
```

4. Understanding XDR: What is it, and why is Wazuh an XDR?

What is XDR?

XDR stands for **Extended Detection and Response**. It is an evolutionary step beyond traditional EDR (Endpoint Detection and Response) and classic SIEM tools.

- **EDR:** Focuses strictly on monitoring physical endpoints (laptops, desktops, servers) for threat signatures and execution paths.
- **SIEM:** Focuses on aggregating massive volumes of log data from different network components for historical query, compliance, and correlation.
- **XDR:** Unifies both models. It extends beyond endpoint protection by seamlessly integrating security telemetry from **multiple layers**—endpoints, network traffic, cloud platforms, containers, and identity access systems—into a single, correlated data stream. XDR doesn’t just detect; it automatically coordinates **active responses** across these diverse layers.



Why Wazuh is a True XDR Platform

Wazuh is classified as a genuine XDR platform because it bridges the gaps between network security, cloud activity, host behavior, and active security automation:

1. Multi-Vector Telemetry Integration (Extended Detection)

Rather than looking only at process execution on a VM, Wazuh correlates multiple domains simultaneously:

- **Endpoints:** Ingests local system calls, RAM changes, shell commands, and file hash alterations.
- **Network Traffic:** Parses logs from Suricata, Snort, or Bro/Zeek to detect network scans and SQL injection vectors.
- **Cloud Infrastructure:** Directly monitors cloud APIs (such as AWS CloudTrail, Amazon GuardDuty, and VPC Flow Logs) to track infrastructure modifications, MFA deletions, and unauthorized security group configuration changes.
- **Containers & Kubernetes:** Integrates with Docker socket event streams and Kubernetes API audits to detect container escapes or privilege-escalation attempts.
- **Identity and Access:** Monitors authentication patterns (Active Directory, LDAP, Keycloak, or AWS Cognito) to flag brute-force attacks and credential stuffing.

2. Cross-Layer Threat Correlation

If a malicious actor executes an attack, Wazuh does not treat the events in isolation. For example:

1. An attacker attempts a brute-force login on Keycloak (**Identity Layer**).
2. The login succeeds, and the actor immediately spins up a rogue instance (**Cloud Layer**).
3. The new instance starts port scanning other private VPC nodes (**Network Layer**).
4. The actor attempts to modify a host’s SSH authorized keys file (**Endpoint Layer**).

Wazuh correlates these distinct vectors into a single, cohesive security incident chain, enabling security teams to see the complete lifecycle of the attack.

3. Automated Cooperative Countermeasures (Extended Response)

An XDR must be capable of orchestrating responses across its scope. Wazuh supports this via dynamic **Active Response**:

- **Endpoint response:** Instantly terminates a flagged ransomware process or deletes a dropped malicious script.
- **Network response:** Automatically executes a script to update a local firewall blocklist or triggers an API call to a corporate edge proxy (like BunkerWeb) to block a brute-forcing IP address.
- **Access response:** Can execute a command to lock or disable a compromised local user account in the OS directory.

Technology Stack Comparison Guide

#aws #cloud #architecture #vpc #alb #asg #rds #elasticsearch #valkey #ragflow #cognito #bedrock #whatsapp #meta #lambda #apigateway #costing #licensing #wazuh

Technology Stack Comparison & AWS Mapping Guide

This guide provides a comprehensive technical and architectural comparison between the developer’s proposed local/on-premises containerised technology stack and our enterprise-grade, highly-available **AWS Secure 3-Tier Architecture**.

We analyse every layer of the developer’s architecture—from Frontend React 19 to AI Integration, Messaging, and Databases—mapping them directly to their corresponding AWS-native managed services or secure hosting tiers in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)**.

Additionally, we evaluate **AWS-native cloud alternatives** for external services (such as Twilio, Meta APIs, and OpenAI integrations) to provide a completely integrated cloud solution with corresponding cost structures.

This guide is integrated with our official [Software Licensing & Technology Risk Register \(TS/MC Series\)](#) to track accepted risk and licensing compliance profiles.

1. High-Level Technology Mapping Matrix

The table below outlines how the developer’s stack is mapped to our secure AWS architecture, including recommended AWS-native alternatives for third-party components, and maps them to their respective **TS/MC series risk codes**.

Category	Proposed Developer Technology	Our AWS Secure 3-Tier Target	AWS-Native Cloud Alternative	Risk & License Code
AI Integration	LangChain4j, OpenAI-compatible Provider (Chat, Embeddings, Function Calling)	Hosted on ASG Compute Tier (Graviton t4g.xlarge or GPU-backed instances), calling secure APIs or local inference.	Amazon Bedrock (Serverless access to Anthropic Claude, Meta Llama 3, Cohere, etc. with local data privacy/sovereignty compliance).	TS-02 , MC-01 , MC-02
Backend	Java 21, Spring Boot 3.5.12, Spring Framework 6.2, Spring Security 6.x, Spring Data JPA, Hibernate 6.x	Deployed within Private App Subnets under an Auto Scaling Group (ASG) of Graviton ARM64 instances, managed via AWS Systems Manager (SSM) .	<i>Same application code, but highly scalable and secure.</i>	TS-05
Frontend	React 19, Vite, TypeScript, TanStack Router, Ky, shadcn/ui	S3 Bucket + Amazon CloudFront (Global CDN) for static asset distribution OR served via the private ASG behind the Application Load Balancer (ALB) .	Amazon S3 + Amazon CloudFront (Saves compute costs and improves global latency via edge caching).	TS-05
Mobile Application	React Native, Expo, HeroUI	Connects securely to the Application Load Balancer (ALB) endpoint protected by AWS WAFv2 .	<i>Client-side runtime, standard HTTPS API endpoint mapping.</i>	TS-05
Authentication & Authorization	Spring Security, JWT Authentication, OAuth 2.0, RBAC	Managed within the Spring Boot backend layer with token caches stored in Amazon ElastiCache Valkey .	Amazon Cognito User Pools (Serverless user directory, OAuth 2.0 flow, JWT generation, MFA, and native security auditing).	TS-05
Database	PostgreSQL 16, Docker, Flyway	Fully managed Amazon RDS PostgreSQL (Multi-AZ) in isolated private subnets, with schema migrations coordinated via standalone nodes or CI/CD pipelines.	Amazon RDS for PostgreSQL (Multi-AZ synchronous replication, automated backups, automated patching, zero-SPOF).	TS-06
Cache / Performance	Redis (Rate limiting, usage counters, session/cache management)	Amazon ElastiCache for Valkey (Secure, multi-AZ, and license-compliant caching layer offering 20% lower pricing than legacy Redis OSS).	Amazon ElastiCache for Valkey (Dedicated managed caching cluster).	TS-06 , TS-05
Build Tools	Backend: Maven; Frontend: npm + Vite	Integrated with GitHub Actions and GitLab CI/CD with shared EFS mounts for build caching.	AWS CodeBuild (Fully managed build service).	TS-05
Deployment & Infrastructure	Docker, Docker Compose, PostgreSQL, Redis, Backend, Frontend, AI services	Multi-AZ modular infrastructure provisioned via OpenTofu (ALB, ASG, Valkey, RDS, WAFv2, and Standalone compute).	AWS ECS (Elastic Container Service) on Fargate or AWS EKS for managed container orchestration.	TS-06
Enterprise Standards	Jakarta EE 10, Servlet 6.0, JPA 3.1, Bean Validation 3.0	Run inside the standard Spring Boot JVM running on hardened Ubuntu 26.04 LTS (ASIMP framework) .	<i>Same enterprise compliance parameters.</i>	TS-05
AI Knowledge Base / RAG	RAGFlow (Document ingestion, parsing, chunking, embedding generation, vector search, retrieval-augmented generation)	Dedicated AI Compute Tier (ASG / Standalone) inside private subnets connected to Amazon EFS and OpenSearch Service (or local vector database).	Amazon Bedrock Knowledge Bases (Serverless RAG pipeline utilizing OpenSearch Serverless and Bedrock embeddings).	TS-06
Messaging / Omnichannel	Twilio for WhatsApp (WhatsApp Business messaging, inbound/outbound, webhooks, media, callbacks)	Integrated in the backend code via Twilio Java SDK, with webhook receivers protected by AWS WAFv2 .	AWS End User Messaging (Social Channels) (Direct AWS-native WhatsApp Business API integration, eliminating Twilio as an intermediary).	<i>Third-party API</i>
Social Media Messaging	Meta for Developers / Meta Graph API (Facebook Messenger, Instagram Messaging API, webhook subscription, page tokens)	Outbound requests routed via NAT Gateway ; inbound webhooks handled via the ALB or a serverless proxy.	AWS Lambda + Amazon API Gateway (Serverless, auto-scaling webhook receivers that scale to zero when inactive to minimise idle costs).	<i>Third-party API</i>

Category	Proposed Developer Technology	Our AWS Secure 3-Tier Target	AWS-Native Cloud Alternative	Risk & License Code
External API Integration	REST API, Webhook, OAuth 2.0, API token management, retry mechanism, request validation	Outbound requests secure-routed through AWS NAT Gateway ; secrets managed inside AWS Secrets Manager .	Amazon API Gateway with custom authorizers and integration throttling.	<i>Third-party API</i>
Security & SIEM Operations	Traditional/Vendor SIEM platform	Hardened standalone Wazuh SIEM EC2 instance whitelisted strictly for Cyberjaya developer CIDRs.	AWS Security Hub + Amazon GuardDuty (Consolidated multi-account cloud-native threat detection).	TS-04

2. Technical Mapping Details & Architectural Transitions

A. Frontend Web Tier (React 19 + Vite)

- **Proposed Stack:** React 19, Vite, TypeScript, TanStack Router, Ky, shadcn/ui.
- **On-Premises / VM Deployment:** Typically served via Nginx inside a Docker container running on a public-facing virtual machine.
- **AWS Enterprise Architecture [TS-05]:**
 - **Static Site Hosting (CloudFront + S3):** We compile the Vite build artifacts and push them to a secure **Amazon S3** bucket. An **Amazon CloudFront** distribution sits in front of the bucket, serving assets globally through edge locations. This eliminates the need to run EC2 compute for the frontend, reduces latency, and protects against direct S3 bucket scraping.
 - **Secure Dynamic Routing:** API requests (such as /api/*) are routed by CloudFront directly to the **Application Load Balancer (ALB)**, while static routes are handled at the edge, creating a unified origin configuration.

B. Backend App Tier (Spring Boot 3.5.12 + Java 21)

- **Proposed Stack:** Spring Boot 3.5.12, Spring Framework 6.2, Spring Security 6.x, Spring Data JPA, Hibernate 6.x, Jakarta EE 10, Maven.
- **On-Premises / VM Deployment:** Executed as a systemd service or Docker container on Server 02 (4 vCPU, 16GB RAM) exposed to the web.
- **AWS Enterprise Architecture [TS-05]:**
 - **Stateless Auto-Scaling Compute:** Deployed across an Auto Scaling Group (ASG) of **AWS Graviton ARM64** instances (e.g., t4g.xlarge to match the 4 vCPU/16GB RAM spec) in **Private App Subnets**.
 - **SSM-Only Access:** Direct SSH is disabled. All staging, auditing, and maintenance is handled passwordlessly via **AWS Systems Manager (SSM)**.
 - **ASIMP Hardening:** Operating systems are locked down to **Ubuntu 26.04 LTS** and audited via the **ASIMP (Ansible System Integrity Management Platform)** framework against CIS Level 2 benchmarks.

C. Caching & Performance (Redis)

- **Proposed Stack:** Redis (used for rate limiting, usage counters, session/cache management).
- **On-Premises / VM Deployment:** A local Docker Redis container running on the same host, with no high availability.
- **AWS Enterprise Architecture [TS-06]:**
 - **Amazon ElastiCache for Valkey:** We map Redis directly to **ElastiCache Valkey** (cache.t4g.medium). Valkey is a fully open-source, license-compliant key-value cache engine that provides full Redis API compatibility but at **20% lower hourly rates** on AWS.
 - **Transit and At-Rest Encryption:** Subnet-isolated caching clusters with forced TLS and strict ingress restricted strictly to the compute security group.

D. Relational Database (PostgreSQL 16 + Flyway)


- **Proposed Stack:** PostgreSQL 16, Docker, Flyway.
- **On-Premises / VM Deployment:** PostgreSQL running inside Docker on Server 04 (4 vCPU, 16GB RAM), risking single-point-of-failure database corruption.
- **AWS Enterprise Architecture [TS-06]:**
 - **Amazon RDS PostgreSQL 16 (Multi-AZ):** We migrate local Postgres to a fully managed **Amazon RDS PostgreSQL** cluster (Multi-AZ) sized at db.m6g.xlarge (4 vCPU, 16GB RAM). RDS handles automatic Multi-AZ synchronous replication, daily automated snapshots, point-in-time recovery (PITR), and security patching.
 - **Network Isolation:** Isolated inside **Private Database Subnets** with zero public internet connectivity. Schema migrations are automatically executed at build-time via GitHub/GitLab runners or a secure Standalone staging EC2 instance.

E. AI Knowledge Base / RAG Platform (RAGFlow)

- **Proposed Stack:** RAGFlow (DeepDoc parser, OCR, semantic chunking, vector search).
- **On-Premises / VM Deployment:** Heavy Docker Compose setup relying on local GPUs or CPU fallbacks, causing processing delays and high latency.
- **AWS Enterprise Architecture [TS-06]:**
 - **GPU Compute & Shared Storage:** RAGFlow executors run on dedicated **G5 instances (g5.xlarge with NVIDIA A10G 24GB VRAM)**. We mount an **Amazon EFS (Elastic File System)** volume across the instances to store shared AI model weights and cached document parsing structures.
 - **Dense Vector Search:** High-performance indexing and dense vector queries are managed natively via **Amazon OpenSearch Service** (with k-NN search enabled) or an isolated PostgreSQL vector database (pgvector extension) on RDS.

3. High-Fidelity AWS-Native Alternatives for “Extra” Items

The developer’s technology stack contains several “extra” third-party SaaS integrations (such as Twilio for WhatsApp, Meta Graph APIs, and OpenAI endpoints). In enterprise deployments, relying on external SaaS platforms introduces data sovereignty issues, external billing overhead, and architectural latency. Below are the architectural transitions and benefits of choosing **AWS-native replacements**.



Alternative A: Amazon Bedrock (Replaces OpenAI/LangChain4j Endpoints - MC-01 / MC-02 / TS-02)

Instead of routing enterprise chat and sensitive document embeddings to third-party OpenAI endpoints over the public internet, utilise **Amazon Bedrock** loaded with state-of-the-art **Qwen3** or Claude models.

- **AWS Architecture [MC-01]:**
 - Bedrock provides high-speed, serverless API access to leading foundation models (such as **Qwen3 LLM Inference**) entirely within your AWS network.
 - Since traffic never leaves the AWS boundary, this aligns perfectly with the **Malaysian Personal Data Protection Act (PDPA) 2010** data residency and sovereignty requirements.
 - Bedrock integrates natively with LangChain4j via standard AWS SDK dependencies, mitigating the commercial SLA lack of LangChain4j by using enterprise-secured network pipelines.
- **Embedding & Query Optimisation [MC-02]:**
 - For static document databases, perform a one-time high-efficiency indexing sweep using Qwen3 Embedding models and store vectors locally in RDS PostgreSQL (pgvector), minimizing ongoing runtime model consumption fees.

Alternative B: AWS End User Messaging (Replaces Twilio for WhatsApp)

Instead of routing WhatsApp Business chats, media files, and webhooks through Twilio as a middleware, utilise **AWS End User Messaging (Social Channels)**.

- **AWS Architecture:**
 - AWS End User Messaging connects your applications directly to the Meta WhatsApp Business API.
 - This eliminates Twilio’s per-message platform markup fees, consolidates WhatsApp bills directly into your single AWS invoice, and allows secure, automated webhook handling directly via AWS Lambda.

Alternative C: Amazon Cognito User Pools (Replaces Self-Managed JWT/Spring Security)

Instead of managing sensitive password hashing, JWT token signing, rotation, MFA, and database user tables inside the custom Java Spring Boot code, utilise **Amazon Cognito**.

- **AWS Architecture:**
 - Spring Security delegates authentication to **Amazon Cognito User Pools** via standard OAuth 2.0 / OIDC protocols.
 - Cognito acts as a secure, serverless identity provider that automatically handles multi-factor authentication (MFA), user registration, password reset flows, and brute-force protection.
 - Eliminates the security risk of storing credentials directly in RDS and reduces backend JVM memory footprint.

Alternative D: Amazon API Gateway & AWS Lambda (Replaces Custom Webhook Receivers)

The Meta Graph API (Facebook Messenger/Instagram) uses highly dynamic webhooks that send intense bursts of traffic during high-volume communication spikes, saturating Spring Boot JVM threads.

- **AWS Architecture:**
 - Public webhooks are routed to **Amazon API Gateway** and processed by serverless **AWS Lambda** functions.
 - API Gateway validates incoming request signatures, and AWS Lambda scales instantly to handle tens of thousands of concurrent messages, before routing the parsed messages into an **Amazon SQS** queue for the backend ASG to process asynchronously.
 - Minimises idle compute cost (scales to \$0.00 when there are no active messages).

4. Architectural Comparison Summary

The matrix below compares the security, scalability, and operational realities of the developer’s proposed local VM stack with our AWS Secure 3-Tier mapping (including the AWS-native alternatives).

Architectural Dimension	Proposed Developer VM / Docker	Our AWS Secure 3-Tier Solution
High Availability	Single VM per service. Host failure causes total outage.	Multi-AZ redundancy at ALB, ASG, Valkey, and RDS layers. Automated sub-minute failovers.
Network Isolation	Compute & DB run on public IPs. High SSH and port-scan risk.	Compute & DB isolated in private subnets. Ingress restricted strictly via Security Groups.
Security Hardening	Manual OS patching and updates. Unaudited configurations.	ASIMP-hardened Ubuntu 26.04 LTS AMIs. Continuous auditing with OpenSCAP & Lynis.
Scalability	Static server resources. Manual resizing is required.	Dynamic horizontal auto-scaling (ASG) based on CPU/memory saturation profiles.
API Webhook Resilience	Custom Spring Boot web endpoints susceptible to JVM OOM.	Serverless API Gateway + Lambda scaling. Bursty WhatsApp/Meta traffic is absorbed.
Identity Protection	Self-managed JWT token store. Risks with key rotation & storage.	Amazon Cognito managed user directories with automated token rotation, MFA, and compliance.
AI Compliance & Residency	OpenAI requests route overseas. Subject to foreign data transit.	Amazon Bedrock processing runs inside Malaysia (ap-southeast-5) or secure local GPU nodes.

AWS Services vs. On-Premises Open-Source Comparison Guide

#aws #onprem #cloud #architecture #comparison

[STRATEGIC FINANCIAL] — Cost Control, Sizing & TCO
[DEVOPS EXECUTION] — Architecture & Deployment Mapping

AWS Services vs. On-Premises Open-Source Comparison Guide

This guide provides a comprehensive architectural and operational comparison between **AWS Cloud Services** and **On-Premises Onsite Open-Source Software Solutions** across all **12 critical infrastructure layers** of our modern enterprise application.

To support organisations planning a cloud adoption roadmap or evaluating complete local data sovereignty, this document contrasts the managed, automated conveniences of AWS with the raw host control, license-compliant, and self-hosted cost efficiencies of an on-premises stack.

The 12-Layer Enterprise Infrastructure Stack

The following diagram illustrates the 12 key layers that comprise our enterprise web and AI application infrastructure, from user interaction to runtime observability:

1. Frontend (Web & Mobile UI Assets)	
2. APIs & Backend Logic (Business Engines)	
3. Database & Storage (Relational, Objects & Vectors)	
4. Auth & Permissions (Identity Providers)	
5. Hosting & Deployment (Orchestration & Runtimes)	
6. Cloud & Compute (Virtualisation & Hardware)	
7. CI/CD & Version Control (Delivery Pipelines)	
8. Security & RLS (Firewalls, IDS & Data Isolation)	
9. Rate Limiting (Traffic Protection & Throttling)	
10. Caching & CDN (Low-Latency Acceleration)	
11. Load Balancing & Scaling (High Availability)	
12. Error Tracking & Logs (Observability & Telemetry)	

Comprehensive 12-Layer Architectural Mapping

Below is a detailed, layer-by-layer strategic and technical breakdown comparing the **AWS-Native Managed Stack** against the **Hardened On-Premises Open-Source Equivalent** running on rootless containers or local VM hypervisors.

Layer 1: Frontend

Frontend static assets must be delivered with minimum latency and high security. In modern microservices, the UI layer is completely decoupled from backend processing.

- **AWS Cloud-Native Option:**
 - **Services:** **Amazon S3** (Static Website Hosting) combined with **Amazon CloudFront** (Global Content Delivery Network).
 - **Architectural Behaviour:** Web files (HTML, CSS, JavaScript) are compiled and stored in an S3 bucket with public access disabled. CloudFront distributes these assets to edge locations globally.
 - **Advantages:** Scale is completely automated and virtually infinite; zero compute resources are wasted on serving static pages, and the origin bucket is protected from direct scanning or malicious manipulation.
- **On-Premises / Onsite Open-Source Option:**
 - **Solutions:** **Nginx** / **BunkerWeb** serving pre-compiled React 19 assets, or lightweight **Node.js** containers.
 - **Architectural Behaviour:** Static assets are served locally by Nginx or a BunkerWeb web-server container inside a virtual machine.
 - **Advantages:** Absolute control over web headers, local caching profiles, and logging. Eliminates external CDN bandwidth charges, making it highly cost-effective for local workloads.

Layer 2: APIs & Backend Logic

The core business engine, housing API routing, transactional databases integration, background workers, and AI prompt pipelines.

- **AWS Cloud-Native Option:**
 - **Services:** **Amazon ECS (Elastic Container Service)** on **AWS Fargate** or stateless **Auto Scaling Groups (ASG)** of **AWS Graviton** (t4g.xlarge / c7g.xlarge) instances running Spring Boot 3.5.12.
 - **Architectural Behaviour:** Dynamically scales compute resources up and down based on target CPU or memory utilisation metrics.

- **Advantages:** Fully managed, serverless, or auto-scaling environments reduce the burden of OS kernel tuning and ensure maximum resource utilisation efficiency.
 - **On-Premises / Onsite Open-Source Option:**
 - **Solutions:** **Spring Boot 3.5.12** running inside rootless, unprivileged **Podman 5+** containers managed by **systemd Quadlets** on local enterprise Linux hosts.
 - **Architectural Behaviour:** Containers run without administrative root rights using systemd lingering and user namespace keep-id remapping.
 - **Advantages:** Removes public cloud hypervisor markup fees, locks runtime environments to strict local hardware boundaries, and allows deep customisation of the JVM garbage collector and host filesystem.
-

Layer 3: Database & Storage

The persistence layer, which must support high-performance ACID transactions, dense vector indexing for Retrieval-Augmented Generation (RAG), and shared object storage.

- **AWS Cloud-Native Option:**
 - **Services:** **Amazon RDS for PostgreSQL** (Multi-AZ), **Amazon S3** (Object Storage), and **Amazon OpenSearch Service** (Vector Databases).
 - **Architectural Behaviour:** Managed synchronous block-level replication across multiple availability zones. RDS handles automated patching, daily backups, and point-in-time recovery (PITR) natively.
 - **Advantages:** Minimal operational risk; automated failover under 120 seconds with zero-SPOF storage replication.
 - **On-Premises / Onsite Open-Source Option:**
 - **Solutions:** **Percona Server for PostgreSQL 17** (with **Patroni** cluster orchestration, **etcd** consensus, and **pg_backrest** backups), **MinIO** (S3-compatible Object Storage), and **pgvector** / **Milvus** (Local Vector Search).
 - **Architectural Behaviour:** High availability is engineered locally using a three-node cluster (two database engines and an etcd distributed consensus coordinator). HAProxy handles primary/replica routing.
 - **Advantages:** Complete binary compatibility with upstream PostgreSQL, zero licensing fees, local RAG vector processing on dedicated bare-metal NVMe drives, and zero cross-AZ data transfer surcharges.
-

Layer 4: Auth & Permissions

Securing user identities, implementing Multi-Factor Authentication (MFA), issuing JSON Web Tokens (JWT), and enforcing Role-Based Access Control (RBAC).

- **AWS Cloud-Native Option:**
 - **Services:** **Amazon Cognito User Pools & Identity Pools**.
 - **Architectural Behaviour:** A fully managed, serverless identity directory that implements OAuth 2.0 and OpenID Connect (OIDC) patterns natively.
 - **Advantages:** Seamless integration with AWS WAFv2, offloads credentials security auditing from the local database, and supports automated token signing and rotation.
 - **On-Premises / Onsite Open-Source Option:**
 - **Solutions:** **Keycloak** or **Authentik** integrated with Spring Security and local Postgres tables.
 - **Architectural Behaviour:** Highly customizable, open-source identity providers deployed in local containers, supporting OAuth 2.0, SAML 2.0, and LDAP integrations.
 - **Advantages:** Eliminates external monthly active user (MAU) licensing charges, keeps user credentials physically isolated on-site, and allows customized login flow designs and styling.
-

Layer 5: Hosting & Deployment

Container orchestration and deployment execution. Decides how workload containers are structured, started, and managed.

- **AWS Cloud-Native Option:**
 - **Services:** **Amazon ECS on AWS Fargate** or **Amazon EKS (Elastic Kubernetes Service)**.
 - **Architectural Behaviour:** Abstracts container runtime interfaces entirely. Deployments are declared via CloudFormation, OpenTofu, or AWS Copilot.
 - **Advantages:** Zero infrastructure management. Eliminates the need to configure node networking, container registries, or control planes.
 - **On-Premises / Onsite Open-Source Option:**
 - **Solutions:** **Podman 5+** utilizing **systemd Quadlets** or a lightweight **K3s** / **MicroK8s** Kubernetes distribution.
 - **Architectural Behaviour:** Quadlets are declarative systemd configuration files that generate and run containers natively under systemd, treated exactly as standard OS services.
 - **Advantages:** Clean, configuration-driven deployments without the massive system overhead of a heavy Kubernetes control plane. Containers restart on host boot automatically without daemon dependencies.
-

Layer 6: Cloud & Compute

The physical or virtualized hardware executing the instructions. Highly influential over overall performance and cost efficiency.

- **AWS Cloud-Native Option:**
 - **Services:** **Amazon EC2** running **AWS Graviton (ARM64)** processors or GPU-backed instances (e.g., g5.xlarge with NVIDIA A10G VRAM for heavy RAG parsing).
 - **Architectural Behaviour:** Instances are provisioned on demand in multiple independent physical locations (Availability Zones).
 - **Advantages:** Infinite, elastic hardware availability on-demand; easy access to high-performance enterprise hardware with zero capital expenditure (CapEx).
 - **On-Premises / Onsite Open-Source Option:**
 - **Solutions:** Bare-metal rack servers running **Proxmox VE**, **KVM (Kernel-based Virtual Machine)**, or **XCP-ng** hypervisors on AMD EPYC / Intel Xeon architectures.
 - **Architectural Behaviour:** Physical servers are carved into multiple virtual machines using local hypervisors, with dedicated PCIe passthrough for local GPU accelerators (such as NVIDIA L40S or RTX series).
 - **Advantages:** High, sustained compute performance at flat costs. Avoids public cloud virtual machine hypervisor overhead, memory tax, and storage throughput throttling.
-

Layer 7: CI/CD & Version Control

The automated pipelines that receive code commits, run testing suites, audit configurations, and compile production-ready assets.

- **AWS Cloud-Native Option:**

- **Services:** AWS CodePipeline, AWS CodeBuild, and AWS CodeDeploy.
 - **Architectural Behaviour:** Serverless pipelines triggered directly by code events, running builds in secure, ephemeral container runtimes.
 - **Advantages:** Integrates natively with IAM permissions, secrets managers, and ECS/ASG deployment endpoints.
 - **On-Premises / Onsite Open-Source Option:**
 - **Solutions:** Gitea (Private git repository hosting) integrated with **Ansible Semaphore** (Automated UI task runner) and **Ansible ARA** (Execution callbacks recorder).
 - **Architectural Behaviour:** Developer pushes code to Gitea; a webhook triggers Ansible Semaphore playbooks to build, test, and deploy containers across the private network.
 - **Advantages:** 100% self-hosted, private version control; allows ultra-fast local build caching, and records detailed execution reports in local compliance databases.
-

Layer 8: Security & RLS

Perimeter security, intrusion detection systems (IDS), Row-Level Security (RLS) in databases, and active host hardening.

- **AWS Cloud-Native Option:**
 - **Services:** AWS WAFv2 (Web Application Firewall), **AWS Shield** (DDoS Protection), and **AWS Security Hub / Amazon GuardDuty**.
 - **Architectural Behaviour:** Agentless, cloud-plane detection analyzing VPC flow logs, DNS queries, and IAM API calls to spot threats instantly.
 - **Advantages:** Global threat intelligence databases, zero performance impact on application servers, and automated compliance auditing.
 - **On-Premises / Onsite Open-Source Option:**
 - **Solutions:** **BunkerWeb** (Open-source WAF), **Wazuh SIEM** (Host Intrusion Detection), and PostgreSQL native **Row-Level Security (RLS)** policies.
 - **Architectural Behaviour:** Wazuh agents are installed on all local hosts to collect logs, check CIS compliance, and audit system integrity. BunkerWeb filters incoming web traffic using security-tuned Nginx rules.
 - **Advantages:** Comprehensive host-level intrusion detection and auditing; significantly cheaper than cloud-native alternatives. PostgreSQL RLS ensures strict tenant-level data isolation directly at the database engine level.
-

Layer 9: Rate Limiting

Preventing API abuse, brute-force login attempts, and denial-of-service (DDoS) spikes from saturating application logic threads.

- **AWS Cloud-Native Option:**
 - **Services:** AWS WAFv2 **Rate-Limiting Rules** or **Amazon API Gateway Throttling**.
 - **Architectural Behaviour:** Traffic is inspected and throttled at the AWS edge network before it ever reaches backend EC2 or container instances.
 - **Advantages:** Eliminates backend compute saturation; traffic from abusive IPs is dropped at the nearest regional edge location.
 - **On-Premises / Onsite Open-Source Option:**
 - **Solutions:** **BunkerWeb** / **Nginx** rate-limiting modules, combined with **Valkey** (for dynamic API usage counters).
 - **Architectural Behaviour:** BunkerWeb tracks requests per second per IP and rejects requests exceeding thresholds. For advanced application-level rate limiting, Spring Boot checks and increments counters stored in Valkey.
 - **Advantages:** Fully customizable rulesets, zero external latency overhead for looking up rate limits, and zero public cloud billing surcharges for high traffic counts.
-

Layer 10: Caching & CDN

Accelerating delivery of dynamic API responses and static assets while reducing read load on primary database systems.

- **AWS Cloud-Native Option:**
 - **Services:** **Amazon ElastiCache for Valkey** (In-Memory Key-Value Caching) and **Amazon CloudFront**.
 - **Architectural Behaviour:** Fully managed Multi-AZ in-memory database clusters configured in private subnets with automated node failovers.
 - **Advantages:** Full compatibility with Valkey/Redis API, and 20% lower instance pricing than legacy Redis engines.
 - **On-Premises / Onsite Open-Source Option:**
 - **Solutions:** Self-hosted **Valkey** running in containerized or standalone modes, combined with local **Nginx proxy caching**.
 - **Architectural Behaviour:** Valkey is deployed locally on physical hosts with transit and at-rest encryption enabled, serving as a shared session and API cache layer.
 - **Advantages:** Zero network latency penalty (sub-millisecond local network execution), complete open-source license compliance, and flat resource allocation with zero managed service markups.
-

Layer 11: Load Balancing & Scaling

Distributing incoming user traffic uniformly across active backend instances to prevent single-instance saturation, combined with dynamic auto-healing.

- **AWS Cloud-Native Option:**
 - **Services:** **AWS Application Load Balancer (ALB)** and **Auto Scaling Groups (ASG)**.
 - **Architectural Behaviour:** Health checks constantly monitor instances; if an instance becomes unhealthy or if CPU usage exceeds the threshold, the ASG terminates the node and launches a new one automatically.
 - **Advantages:** Highly available, multi-AZ by default, and seamlessly handles unexpected traffic spikes.
 - **On-Premises / Onsite Open-Source Option:**
 - **Solutions:** **HAProxy** / **Keepalived** for high-availability load balancing, combined with local hypervisor auto-scaling pools.
 - **Architectural Behaviour:** Active-passive load balancers are configured using HAProxy and virtual IPs (Keepalived). Local hypervisor scripts monitor container loads and dynamically adjust resources.
 - **Advantages:** HAProxy is recognized as the world's fastest open-source load balancer, offering superior throughput, customized balancing algorithms, and advanced TCP-level proxy configurations.
-

Layer 12: Error Tracking & Logs

Monitoring system performance, logging errors, tracing distributed transactions, and collecting security event logs.

- **AWS Cloud-Native Option:**
 - **Services:** **Amazon CloudWatch Logs**, **AWS X-Ray** (Distributed Tracing), and **Amazon CloudWatch Container Insights**.
 - **Architectural Behaviour:** Ephemeral agents stream system logs and performance metrics to a centralized CloudWatch data lake.
 - **Advantages:** Zero configuration required; unified dashboarding and native integration with SNS alerts.
- **On-Premises / Onsite Open-Source Option:**

- **Solutions:** Wazuh SIEM, Grafana Loki / Prometheus, and Signoz or Jaeger (Distributed Tracing).
- **Architectural Behaviour:** Systemd logs are processed locally by Grafana Promtail and streamed to Loki. Prometheus scrapes application metrics, and Wazuh aggregates audit logs.
- **Advantages:** Absolute control over logs retention policies (no data ingestion or retention charges), local storage compliance, and unified visualization dashboards.

Strategic Summary Matrix

The matrix below provides an executive-level summary of the architectural and financial trade-offs between the two environments across all 12 layers:

Layer	AWS Cloud-Native Managed Stack	On-Premises Open-Source Solution Key Decision Driver	
1. Frontend	S3 Website + CloudFront CDN	Nginx / BunkerWeb Web Server	Latency at the edge vs. bandwidth cost.
2. APIs & Backend	AWS ASG / ECS on Graviton	Spring Boot on Rootless Podman 5+	Elastic auto-scaling vs. absolute host control.
3. Database & Storage	RDS PostgreSQL (Multi-AZ) + S3	Percona PostgreSQL + Patroni + MinIO	Zero operational database administration vs. no transfer fees.
4. Auth & Permissions	Amazon Cognito User Pools	Keycloak / Authentik	Managed identity lifecycle vs. zero licensing MAU charges.
5. Hosting & Deployment	AWS ECS on Fargate	systemd Quadlet Containers	Zero-infra container management vs. minimal system overhead.
6. Cloud & Compute	On-demand Multi-AZ EC2 instances	Bare-Metal Servers + Proxmox VE	Infinite elasticity vs. sustained heavy performance.
7. CI/CD Pipeline	AWS CodePipeline / CodeBuild	Gitea + Ansible Semaphore + ARA	Ephemeral cloud executors vs. secure local git compliance.
8. Security & IDS	AWS WAFv2 + GuardDuty	BunkerWeb + Wazuh SIEM	Cloud control-plane scanning vs. host compliance auditing.
9. Rate Limiting	AWS WAFv2 Edge Throttling	BunkerWeb + Nginx + Valkey	Edge blocking vs. zero execution latency.
10. Caching & CDN	ElastiCache Valkey + CloudFront	Containerized Valkey + Nginx Caching	Managed Multi-AZ clusters vs. sub-millisecond local network times.
11. Load Balancing	AWS Application Load Balancer	HAProxy + Keepalived Virtual IPs	Auto-scaling integration vs. ultra-fast local throughput.
12. Error Tracking	CloudWatch Logs + AWS X-Ray	Grafana Loki + Prometheus + Wazuh	Unified zero-infra dashboards vs. free data ingestion and retention.

Software Licensing & Technology Risk Register (TS/MC Series)

#aws #cloud #architecture #licensing #compliance #siem #wazuh #qwen3 #bedrock #langchain4j #ragflow #langfuse #costing

Software Licensing & Technology Risk Register (TS/MC Series)

This guide establishes the official **Software Licensing & Technology Risk Register (TS/MC Series)** for our AWS Secure 3-Tier Architecture. It documents the core architectural decisions, commercial/open-source licensing assessments, and operational risks associated with our enterprise-grade web and AI infrastructure.

By classifying key software and AI model components under the **TS (Technology Stack)** and **MC (Model Consumption)** series, we ensure complete traceability, clear ownership of technical risk, and alignment with regulatory standards such as the **Malaysian Personal Data Protection Act (PDPA) 2010**.

1. Technology Stack (TS Series) Registry

The TS Series identifies, tracks, and manages the licensing risks and operational models of software components deployed across our 3-tier architecture.

TS SERIES RISK REGISTER			
ID	Component	Risk Level	Action Taken
TS-02	LangChain4j Framework	Medium-Low	Accepted Risk
TS-04	Wazuh SIEM Self-Hosting	Low	Budget Savings
TS-05	Software Licensing Framework	Low	Audit Standard
TS-06	Self-hosted DB & AI Ops	Medium	Internal Option

TS-02: LangChain4j Integration & SLA Risk

- Classification:** AI Integration Layer
- Licensing Profile:** Apache License 2.0 (Highly permissive open-source)
- Risk Context:** LangChain4j lacks an official commercial Service Level Agreement (SLA) or vendor-backed corporate support structure.
- Risk Evaluation & Accepted Risk:**
 - The Risk:** In the event of critical framework bugs, security vulnerabilities (CVEs), or breaking API changes by downstream LLM providers, there is no contractually bound vendor to provide emergency patches or guarantee response times.
 - Why the Risk is Accepted:** LangChain4j is the de facto standard for building AI/LLM applications in the Java ecosystem (supporting Java 21, Spring Boot 3.5.12, and Spring Security 6.x natively). It has a highly active, robust open-source community, wide enterprise adoption, and regular release cycles.
 - Mitigation Strategy:**
 - Forking Strategy:** Maintain an internal git fork of the specific LangChain4j version in use to allow local patch application independent of upstream releases if needed.
 - Dependency Pinning:** Pin dependency versions in Maven pom.xml to prevent automated, untested minor/patch version upgrades.
 - Abstraction Layer:** Use clean interface abstractions in our Java code so that the underlying LLM orchestration framework can be swapped or modified with minimal business logic disruption.

TS-04: Standalone Wazuh SIEM Hosting

- Classification:** Enterprise Security & Compliance
- Licensing Profile:** GPL v3 (GNU General Public License)
- Risk Context:** Replaces traditional, high-cost third-party SaaS commercial SIEM subscriptions with a standalone self-hosted instance in AWS.
- Cost & Budget Impact:**
 - By deploying Wazuh on a dedicated, hardened AWS Graviton ARM64 instance (t4g.large), we eliminate the vendor markup fees and per-GB ingestion surcharges common in SaaS SIEM models.
 - This replaces the “vendor SIEM” budget line item under Security, **reducing monthly operational costs by over 57%** (operating at approx. **\$65.71 USD / RM 295.70 MYR per month** on Graviton, compared to \$153+ USD/month on legacy x86 setups).
- Mitigation Strategy:**
 - Standardise security configurations, restrict public internet access (restrict port 22 and 443 strictly to Cyberjaya developer CIDRs), and backup OpenSearch/Wazuh indexer directories using automated **AWS Backup** schedules.

TS-05: Technology Stack Software & Licensing Framework

- Classification:** Corporate Compliance & Licensing Audit
- Licensing Profile:** Governance Framework
- Context:** Manages open-source and commercial software compliance across the entire 3-tier system to prevent licensing contamination (e.g. copyleft license leakage into proprietary backend code).
- Compliance Actions:**
 - Restrict production runtimes to permissive licenses such as **Apache 2.0, MIT, and BSD**.
 - Restrict the introduction of AGPL or strong copyleft licenses in the backend API layer.
 - Maintain a dynamic **Software Bill of Materials (SBOM)** generated during GitHub Actions and GitLab CI/CD build phases to automatically flag compliance anomalies.

TS-06: Self-hosted Database Operations (Internal Option)

- Classification:** Storage & In-House Data Management
- Licensing Profile:** Open-Source Permissive / Dual-Licensing (PostgreSQL License, BSD-3, Apache 2.0)
- Risk Context:** Running databases (PostgreSQL, Valkey/Redis) and AI orchestrators (Langfuse, RAGFlow) in-house as self-hosted containers/instances rather than relying on commercial SaaS database vendors.

- **Trade-Off Analysis:**
 - **Benefits:** Complete operational control, zero third-party data transit risk, and elimination of premium SaaS database pricing tiers.
 - **Sovereignty Aligned:** Keeps sensitive AI session telemetry, prompt logs, and document vectors strictly localized within our private subnets in the AWS Malaysia (ap-southeast-5) region, satisfying strict local PDPA data residency mandates.
- **Mitigation Strategy:**
 - Map self-hosted designs directly to managed **Amazon RDS PostgreSQL (Multi-AZ with pgvector)** and **Amazon ElastiCache for Valkey** for production workloads, while retaining containerized self-hosted setups as a cost-optimized “Internal Option” for development and local testing.

2. Model Consumption (MC Series) Registry

The MC Series governs how foundational AI models (LLMs, embedding engines) are consumed, focusing on cost efficiency, data residency, and performance optimisation.

MC SERIES RISK REGISTER			
ID	Component	Optimization	Sovereign Gate
MC-01	Qwen3 LLM Inference	Serverless	Bedrock Private
MC-02	Qwen3 Embed, Index & Query	One-time/Low	Batch/Cached

MC-01: Qwen3 LLM Inference via Amazon Bedrock

- **Classification:** Large Language Model (LLM) Inference
- **Consumption Model:** Serverless API / Provisioned Throughput
- **Architectural Flow:**
 - All conversational queries, reasoning loops, and prompt chains are processed using the state-of-the-art **Qwen3 LLM** (or equivalent highly-optimised multilingual models) running natively within **Amazon Bedrock**.
- **Sovereignty & Compliance Mapping:**
 - By routing Qwen3 model inference directly through Amazon Bedrock, all prompt payloads, context windows, and generated tokens are processed entirely within secure, private AWS regional boundaries (ap-southeast-5).
 - This eliminates data leakage risks associated with routing sensitive corporate information over the public internet to external, third-party AI startups. It provides a clean, audit-ready compliance path for Transfer Impact Assessments (TIAs) under PDPA guidelines.

MC-02: Qwen3 Embedding, Indexing, and Query Optimisation

- **Classification:** Dense Vector Search & Retrieval-Augmented Generation (RAG)
- **Consumption Model:** Static Ingestion + Low Ongoing Query Costs
- **Operational Profile:**
 - **Static Documents (One-time Sweep):** Ingest static corporate documents, technical manuals, and knowledge bases using a high-density, one-time embedding sweep using Qwen3 Embedding models.
 - **Storage:** Store the generated high-dimensional dense vectors in **Amazon RDS PostgreSQL (pgvector)** or a secure vector index.
 - **Ongoing Querying (Low Cost):** Since the foundational document corpus is static, ongoing daily operational costs are limited to the tiny vector queries processed during user searches, keeping the average API charge extremely low.
- **Optimisation Lever:**
 - Utilize batch processing for bulk document ingestion during off-peak hours to take advantage of lower pricing tiers and prevent compute thread saturation in the main Spring Boot application ASG.

3. Enterprise Governance & Action Items

To maintain a secure compliance posture, the engineering team must execute the following action items regularly:

1. **Conduct Licensing Sweeps (Quarterly):** Run automated Maven and npm dependency license audits to ensure no copyleft components have bypassed the TS-05 licensing framework.
2. **Review LangChain4j Versions (Monthly):** Audit the active LangChain4j fork for security patches and ensure any local overrides are documented under the TS-02 risk mitigation log.
3. **Optimise Bedrock API Allocations (Daily):** Monitor Amazon CloudWatch metrics for Bedrock API invocation counts and token usage to identify opportunities for prompt caching and further lower MC-01/MC-02 ongoing costs.

Strategic Comparative Review: AWS-Native Managed Platform vs. Self-Hosted Custom Stack

#aws #cloud #architecture #costing #licensing #compliance #postgresql #valkey #bedrock #cognito #wazuh #sovereignty

Strategic Comparative Review: AWS-Native Managed Platform vs. Self-Hosted Custom Stack

This review provides a comprehensive, high-fidelity strategic analysis comparing an **AWS-Native Managed Solution** against a **Self-Hosted / On-Premises Custom Stack** inside the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)**.

When building an enterprise-grade, highly available 3-tier architecture for modern web and AI workloads, organisations face a critical architectural decision: leverage AWS managed services to outsource operational risk, or deploy and run an equivalent open-source stack in-house to maximise raw hardware control and avoid vendor lock-in.

This guide looks at the **whole picture**, comparing the entire application lifecycle—including presentation, compute, database, caching, generative AI, messaging, identity, security, and disaster recovery. It integrates our [Software Licensing & Technology Risk Register \(TS/MC Series\)](#), [RDS vs. Percona PostgreSQL Comparison](#), and [Disaster Recovery Playbook](#) into a single, cohesive decision framework.

1. Executive Summary & The Big Picture

The fundamental trade-off between AWS-Native Managed Services and Self-Hosted/Custom Stacks is **Operational Leverage vs. Raw Hardware Control**:

- AWS-Native Managed Services** abstract infrastructure complexity. AWS guarantees high availability, synchronous multi-AZ replication, automatic failovers, and compliance boundaries. This dramatically reduces the engineering headcount (OpEx) required for Day-2 maintenance, allowing lean teams to focus entirely on product innovation.
- Self-Hosted Custom Stacks** (whether deployed on raw EC2 instances or local on-premises hardware) eliminate managed service markups and vendor APIs. However, they introduce immense operational complexity. High Availability (HA) must be engineered manually using clustering tools (such as Patroni, etcd, PgBouncer, and custom keepalived scripts). This significantly increases specialized engineering labor costs.

High-Level Architectural Mapping Matrix

The side-by-side matrix below illustrates how the two solutions map across every layer of the architecture, along with their respective **TS/MC Series risk and compliance codes**.

Architectural Layer	AWS-Native Managed Solution	Self-Hosted / Custom Stack	Strategic Trade-Off	Compliance & Risk Code
Presentation Web Layer	Amazon S3 + CloudFront (CDN)	Nginx inside VM / Container	CloudFront improves edge latency and protects against direct scraping; Nginx on VMs requires manual patching.	TS-05
Application Compute	Auto Scaling Groups (ASG) on ARM64 Graviton	Dedicated EC2 / Standalone VMs	ASG offers dynamic scaling and elasticity; dedicated VMs run idle at high cost.	TS-05
Database Tier	Amazon RDS PostgreSQL (Multi-AZ)	Percona Server for PostgreSQL with Patroni & etcd on EC2	RDS automates storage replication and failovers; Percona on EC2 requires dedicated DBRE labor.	TS-06
Session Cache Layer	Amazon ElastiCache for Valkey	Self-hosted Redis OSS Docker Container	ElastiCache Valkey is fully managed, license-compliant, and 20% cheaper than Redis; Redis OSS runs on-host without HA.	TS-06
Generative AI & RAG	Amazon Bedrock (Serverless Qwen3 / Claude)	RAGFlow + Langfuse on GPU instances (g5.xlarge) + EFS	Bedrock is serverless, private, and out-of-the-box; RAGFlow offers deep document parsing (DeepDoc) but requires GPU compute.	TS-02 , MC-01 , MC-02
Identity & Auth	Amazon Cognito User Pools	Custom JWT + Spring Security with database tables	Cognito handles MFA, token rotation, and password flows serverless; custom JWT requires database storage and encryption.	TS-05
Messaging & Webhooks	AWS End User Messaging & API Gateway + Lambda	Twilio WhatsApp API & Spring Boot dynamic endpoints	AWS-native eliminates Twilio per-message markup; serverless API Gateway absorbs dynamic bursts safely.	<i>Third-Party Integration</i>
Security & SIEM	AWS Security Hub + GuardDuty	Hardened Wazuh SIEM on Graviton EC2	GuardDuty offers cloud-native detection; Wazuh SIEM provides affordable host-level intrusion detection and auditing.	TS-04
Disaster Recovery (DR)	Multi-AZ with automated backups	AWS DRS (Strategy E) continuous replication	Managed Multi-AZ RDS ensures sub-minute recovery; DRS replicates block volumes into low-cost staging subnets.	TS-06


2. Technical Comparison Per Layer

2.1 Compute and Presentation Layer

- AWS-Native Solution:** App code runs inside stateless Auto Scaling Groups (ASG) using cost-optimized **ARM64 Graviton instances** (t4g.xlarge or c7g.xlarge). Static frontend assets are compiled and hosted in **Amazon S3**, distributed globally via **Amazon CloudFront**.
 - Benefits:* Scales horizontally based on CPU or memory saturation profiles. There are no idle VM compute costs for static web files. CloudFront absorbs distributed denial of service (DDoS) attacks at the edge, integrated with **AWS WAFv2** for active rate-limiting and OWASP protection.
- Self-Hosted Solution:** A monolithic virtual machine (such as a single Ubuntu VM running Nginx and Spring Boot).
 - Drawbacks:* Scaling is vertical, requiring manual VM sizing upgrades and scheduled downtimes. Ingress traffic directly hits the virtual machine, exposing ports (SSH/HTTP) to scanning and exploits. The VM operates at full cost even during periods of zero traffic.

2.2 Database Tier (RDS PostgreSQL vs. Percona on EC2)

- **AWS-Native Solution: Amazon RDS PostgreSQL (Multi-AZ).** High availability is managed synchronously at the block storage level. If AZ-A fails, RDS automatically points the CNAME DNS endpoint to the standby node in AZ-B within **60 to 120 seconds**. Backups are automated, continuous, and integrated with Point-in-Time Recovery (PITR).
- **Self-Hosted Solution: Percona Server for PostgreSQL on EC2 (Patroni, etcd, PgBouncer).** High availability requires a multi-node cluster (two database nodes and a third lightweight consensus node running **etcd**). **Patroni** orchestrates failovers, reducing database redirection latency to **10 to 30 seconds**. Backups are managed using **pg_backrest** streaming to an S3 bucket.
 - *The Catch:* While saving on raw infrastructure, the self-hosted cluster requires expert database engineering (DBRE) to configure, patch, and maintain Patroni playbooks, etcd consensus states, and pg_backrest recovery pipelines.



2.3 Caching Tier (ElastiCache Valkey vs. Self-hosted Redis)

- **AWS-Native Solution: Amazon ElastiCache for Valkey.** Fully managed key-value caching on Graviton (cache.t4g.micro or cache.t4g.medium). Valkey is a modern, open-source replacement for Redis OSS, offering full API compatibility but **20% lower pricing** on AWS.
- **Self-Hosted Solution: Redis OSS on EC2/Docker.** Runs as a local Docker container alongside compute services.
 - *Drawbacks:* The cache layer is single-point-of-failure. If the container or host crashes, all session memory, user rate limits, and web tokens are wiped, causing massive application downtime.

2.4 Generative AI & RAG Tier (Amazon Bedrock vs. GPU EC2)

- **AWS-Native Solution: Amazon Bedrock (Qwen3).** Serverless API calling. All prompt payloads, context windows, and embeddings are processed securely within the AWS regional boundary.
 - *Benefits:* Complete network isolation. Pay-per-token model ensures costs scale only with usage, avoiding any idle GPU hosting fees.
- **Self-Hosted Solution: RAGFlow + Langfuse on dedicated GPU compute.** Deployed on G5 instances (**g5.xlarge with NVIDIA A10G 24GB VRAM**), mounting **Amazon EFS** for shared AI models and document parsing.
 - *Benefits:* Superior, customized parsing (DeepDoc parser) and OCR. Complete control over prompt pipelines and logging via local Langfuse containers.

2.5 Security and Threat Detection (Cloud-Native vs. Wazuh SIEM)

- **AWS-Native Solution:** AWS GuardDuty, Security Hub, and AWS Config.
 - *Benefits:* Zero-agent deployment. Deeply integrated into cloud control planes, scanning VPC Flow Logs, DNS queries, and IAM roles dynamically.
- **Self-Hosted Solution: Standalone Wazuh SIEM on EC2.**
 - *Benefits:* Highly cost-effective host-level intrusion detection system (HIDS) and compliance auditor. Running Wazuh on a Graviton t4g.large instance **saves over 57% of license and platform fees** compared to traditional enterprise SIEM platforms.

3. Sovereign Compliance & Legal Audits

Operating inside the AWS Malaysia region (ap-southeast-5) mandates strict compliance with the **Malaysian Personal Data Protection Act (PDPA) 2010** and the **2025 CBPDT Guidelines**.

SOVEREIGN COMPLIANCE GATE	
AWS-Native Managed Solution	Self-Hosted / Custom Stack
- Data resides locally inside ap-southeast-5 private subnets. - Satisfies Section 129 PDPA. - Bedrock API traffic is isolated from foreign transit networks.	- Local host-level control matches residency mandates perfectly. - Requires complex local OS rules and audits for compliance. - Backup streaming uses S3.

1. **Data Residency (Section 129 PDPA):** Both solutions comply with local residency mandates by deploying resources natively inside the Kuala Lumpur region (ap-southeast-5). Payloads are held locally.
2. **Transfer Impact Assessments (TIAs):** Sending sensitive corporate data or customer personally identifiable information (PII) to external, third-party US-based AI models (such as OpenAI endpoints over the public internet) violates strict data transit boundaries. Amazon Bedrock processes Qwen3 models locally in Malaysia, satisfying the requirements of the TIA.
3. **Cryptographic Isolation:** Both models can use AWS Key Management Service (KMS) with customer-managed keys (CMK) to implement envelope encryption on EBS, RDS, and S3 volumes, ensuring total data privacy.

4. Comprehensive Financial Blueprint (TCO Comparison)

To compare the Total Cost of Ownership (TCO) in the Malaysia region, we evaluate the **Enterprise Production Multi-AZ Environment** over 1-year and 3-year timelines.

We assume an exchange rate of **1 USD ≈ 4.50 MYR** and calculate both the raw infrastructure charges and the specialized engineering labor (DBRE/SecOps) required to run and maintain the custom stack.

4.1 1-Year Financial Model (USD & MYR)

Sizing & Cost Component	Option A: AWS-Native Managed Platform	Option B: Self-Hosted Custom Stack (EC2)	Financial Analysis
Ingress Load Balancing	\$45.63 / mo (ALB + WAF + 5M reqs)	\$16.43 / mo (Single VM Nginx)	S3 + CloudFront + ALB + WAF offers superior DDoS protection and edge performance.
Compute / Host Cost	\$635.10 / mo (3x c7g.2xlarge ASG)	\$404.70 / mo (2x c7g.2xlarge database hosts + 1x t4g.micro etcd)	Self-hosted compute saves raw instance charges but lacks dynamic auto-scaling elastic margins.
Database & Caching	\$748.08 / mo • RDS db.m7g.xlarge Multi-AZ (\$492.02) • 250GB gp3 Multi-AZ storage (\$57.50) • Valkey cache.r7g.large Multi-AZ (\$198.56)	\$216.00 / mo • 2x 100GB gp3 database volumes (\$16.00) • Valkey / Redis self-installed on compute (\$0.00) • pg_backrest streaming S3 (~\$200.00)	Managed RDS Multi-AZ storage costs are higher due to synchronous mirroring, but eliminate data loss risk.
Storage & Backups	\$37.50 / mo (EFS + AWS Backup)	\$30.00 / mo (EFS + self-managed scripts)	AWS Backup coordinates automated snapshot rotations.
Engineering Labor (OpEx)	\$150.00 / mo • Estimated 2 DBA/Ops monitoring hours	\$1,500.00 / mo • Estimated 15 expert DBRE/SecOps hours (Patroni/etcd checks, Wazuh upgrades)	The Real Difference: Self-hosting introduces significant human maintenance overhead.
TOTAL MONTHLY COST (USD)	\$1,616.31 USD / month	\$2,167.13 USD / month	When engineering labor is factored in, the AWS-Native platform is more cost-effective.
TOTAL MONTHLY COST (MYR)	~RM 7,273.40 MYR / month	~RM 9,752.09 MYR / month	Calculated at 1 USD ≈ 4.50 MYR
TOTAL 1-YEAR TCO (USD)	\$19,395.72 USD	\$26,005.56 USD	AWS-Native saves \$6,609.84 USD (RM 29,744.28 MYR) in Year 1.

4.2 3-Year Total Cost of Ownership (TCO) Comparison

Over a 3-year lifecycle, organisations can apply **AWS Compute Savings Plans** and **RDS Reserved Instances** to achieve massive discounts (up to 34% off compute and 30% off DB storage).

3-YEAR TCO COMPARISON SUMMARY	
Option A: AWS-Native (Optimised)	Option B: Self-Hosted (EC2)
Infrastructure (3-Yr SP): \$34,030.80 Labor (Optimised): \$5,400.00	Infrastructure (Raw): \$24,016.80 Labor (Ops): \$54,000.00
Total: \$39,430.80 USD (RM 177,438.60)	Total: \$78,016.80 USD (RM 351,075.60)

- **Option A: AWS-Native Managed Solution (3-Year Optimised):**
 - **Infrastructure Cost:** \$945.30 / month × 36 = **\$34,030.80 USD** (MYR 153,138.60)
 - **Engineering Labor:** \$150.00 / month × 36 = **\$5,400.00 USD** (MYR 24,300.00)
 - **Total 3-Year TCO: \$39,430.80 USD (RM 177,438.60 MYR)**
- **Option B: Self-Hosted Custom Stack (3-Year Operations):**
 - **Infrastructure Cost:** \$667.13 / month × 36 = **\$24,016.80 USD** (MYR 108,075.60)
 - **Engineering Labor:** \$1,500.00 / month × 36 = **\$54,000.00 USD** (MYR 243,000.00)
 - **Total 3-Year TCO: \$78,016.80 USD (RM 351,075.60 MYR)**

Financial Impact:

By opting for the AWS-Native Managed Solution, organisations save **\$38,586.00 USD (~RM 173,637.00 MYR)** over 36 months, representing a **49.4% cost reduction**. This proves that managed database premiums are vastly outweighed by the heavy labor burden of running a custom, high-availability architecture.

5. Strategic Recommendation Matrix

The matrix below serves as a final guide for selecting the optimal architecture based on organizational priorities, talent availability, and performance needs.

Organisational Driver	Choose AWS-Native Managed Solution	Choose Self-Hosted / Custom Stack
Engineering Team Sizing	Lean teams (1–5 engineers) without a dedicated database reliability engineer (DBRE).	Large, specialized infrastructure teams with dedicated DBA, SecOps, and virtualization experts.
Time-to-Market (TTM)	Extremely fast. Production-ready multi-AZ environments launch in minutes.	Slow. Requires manual configuration of cluster consensus, replication lag rules, and monitoring hooks.
Performance Fine-Tuning	Standard. Abstracts the operating system filesystem and memory parameters.	Advanced. Allows custom tuning of the kernel (Huge Pages, filesystem blocks, swap behaviours).
Disaster Recovery (DR)	Standardized. Automated failovers and Multi-AZ replication natively managed by AWS.	Highly complex. Requires Patroni triggers, etcd elections, and manual DNS adjustments.
Vendor Portability	Low. Standardizes on AWS-specific APIs, consoles, and network templates.	High. The entire stack (Percona, Patroni, etcd, Redis) can run identically on-premises or on other clouds.

Load Testing Assumptions & Sizing Guide

#aws #cloud #architecture #costing #performance #valkey #rds #postgresql #waf #asg

Load Testing Assumptions & Sizing Guide

This document establishes the official workload assumptions, concurrent user definitions, architectural profiles, and target performance objectives governing the load testing, scalability audits, and capacity planning for our secure **AWS 3-Tier Web and AI Infrastructure**.

These assumptions align directly with the performance testing findings and multi-VU scale-up roadmaps, detailing the dual impact of traffic loads on both system performance and financial costing in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)**.

1. Core Load Testing Objectives & SLA Targets

The primary objective of our load testing is to validate that the multi-tier application stack can scale seamlessly up to 10,000 concurrent Virtual Users (VUs) while adhering to strict Service Level Agreements (SLAs) for security, availability, response times, and cost efficiency.

Service Level Agreement (SLA) Matrix

Metric Category	Target SLA / Threshold	Critical Alert Limit	Diagnostic Source
API/Web Latency (P95)	< 250 ms (Static & Cached)	> 500 ms	Application Load Balancer (ALB)
Transaction Latency (P99)	< 800 ms (DB Dynamic Writes)	> 1,500 ms	RDS Performance Insights
HTTP Error Rate	< 0.01% (Zero Gateway Failures)	> 0.1%	ALB Access Logs & CloudWatch
ASG Compute CPU	< 45% average under peak load	> 75% peak	EC2 CloudWatch Metrics
Valkey Cache Hit Rate	> 99.0% for session lookups	< 95.0%	ElastiCache Redis/Valkey Engine
Database Load (AAS)	< Active vCPU count of DB instance	> DB vCPU limit	RDS Performance Insights (AAS)

2. Multi-VU Sizing Tier Definitions

To structure our scale-up roadmap, we define six distinct workload profiles simulating realistic growth phases. All financial estimations are calculated using a stable reference exchange rate of **1 USD = 4.50 MYR**.

Multi-VU Performance Roadmap		
Target Load	Deployment Model	Key Scaling Bottleneck Target
100 VU 500 VU 1,000 VU 2,500 VU 5,000 VU 10,000 VU	Baseline Dev/Staging Cost-Optimized Mid-Tier Scale-Up High-Performance Prod Heavy Concurrency Extreme Concurrency	Idle capacity / Base costs Single points of failure (SPOF) Thread pool & scale latency Database full scans & disk I/O Valkey session write durability Socket pools & WAF L7 latency

A. Sizing & Resource Allocation Matrix

Target Load (VU)	Deployment Phase / Model	Key Compute Sizing (ASG)	Database Engine Spec	Valkey Cache Sizing	Monthly Cost (USD)	Monthly Cost (MYR)
100 VU	Baseline Dev / Staging	2x t4g.micro	db.t4g.micro (Multi-AZ)	1x cache.t4g.micro	\$141.47 USD	RM 636.62 MYR
500 VU	Cost-Optimized Staged	2x t4g.medium	db.m6g.large (Multi-AZ)	1x cache.t4g.micro	\$418.60 USD	RM 1,883.70 MYR
1,000 VU	Mid-Tier Scale-Up (Interpolated)	2-3x t4g.xlarge	db.m6g.xlarge (Multi-AZ)	1x cache.t4g.medium	\$700.00 USD	RM 3,150.00 MYR
2,500 VU	High-Performance Prod	Avg. 4x t4g.xlarge	db.m6g.xlarge (Multi-AZ)	2x cache.t4g.medium (HA)	\$1,264.56 USD	RM 5,690.52 MYR
5,000 VU	Heavy Concurrency Prod	Min. 4x t4g.xlarge	db.m7g.2xlarge (Multi-AZ)	2x cache.t4g.medium (HA)	\$1,948.12 USD	RM 8,766.54 MYR
10,000 VU	Extreme Concurrency Prod	Min. 8x t4g.xlarge	db.m7g.4xlarge (Multi-AZ)	4x cache.m7g.large (Cluster)	\$3,808.88 USD	RM 17,139.96 MYR

3. Core Architectural Workload Assumptions

To execute high-fidelity simulated tests, the following application-level and user behavior assumptions are defined:

- Read-to-Write Ratio:** The standard web traffic mixture is assumed to be **80% Reads and 20% Writes**.
 - Read operations represent session retrievals, reporting queries, database lookups, and asset downloads.
 - Write operations represent active user logins, database updates (e.g., transactional updates), and reporting uploads.
- Session Persistence:** In-memory sessions are handled via ElastiCache for Valkey, with fallback persistence to the relational database layer.
- Layer-7 Security Filter Overhead:** All public ingress must transit through AWS WAFv2 with OWASP Top 10 Rules, rate-limiting, and deep-packet payload inspections.
- Network Boundaries:** Direct Internet access is prohibited for application servers. Outbound egress is routed exclusively through NAT Gateways.
- Auto Scaling Responsiveness:** ASG cooldown periods are set to 180 seconds, triggering scale-up events when average compute CPU exceeds 50% for two consecutive 1-minute evaluation periods.

4. Key Performance Bottlenecks & Critical Remediation

Based on empirical performance testing audits under heavy loads, several system-level bottlenecks were identified and resolved. These findings serve as design rules for future system scale-ups.

A. Database Query Indexing (2,500 VU Bottleneck)

- Problem:** Under 2,500 concurrent users, RDS MariaDB/PostgreSQL active database load surged to a peak of 8.0 Average Active Sessions (AAS), blocking transactions on database table handlings (wait/io/table/sql/handler).
- Root Cause:** A complete absence of composite indexes on aggregate and transactional tables, forcing expensive disk filesort operations and full table scans.
- Remediation:**
 - Create composite index on MariaDB summary tables: `idx_summary_agg (status, tenant_id, total)`.
 - Create composite index on reconciliation lookup tables: `idx_recons_lookup (reference_id, is_reconciled)`.
 - Create composite index on transaction tables: `idx_parking_lookup (refno, user_id)`.

B. Storage Disk Write Sync & IOPS (2,500 VU Bottleneck)

- Problem:** High-concurrency update transactions caused PostgreSQL backend processes to stall, creating an I/O:walSync disk write wait bottleneck.
- Root Cause:** The default GP3 general-purpose storage volume ran out of write capabilities under sustained, concurrent transactional traffic.
- Remediation:** Upgrade PostgreSQL storage from GP3 to **Provisioned IOPS (io2)** with at least 5,000 Custom IOPS and 250 MB/s throughput, reducing sync latency from 150ms to under 8ms.

C. Web Application Firewall regional ACL Overhead (10,000 VU Bottleneck)

- Problem:** Layer-7 WAF inspection rules introduced an additional latency overhead of up to 20ms under heavy loads.
- Root Cause:** Deep packet inspection evaluating hundreds of complex regex patterns for every request.
- Remediation:** Nest and optimize Web ACL rules, utilize precise regex match patterns, and deploy explicit scope-down statements using `FieldToMatch` settings (e.g., limiting expensive body/payload inspection rules specifically to JSON API requests via `UriPath` and `Method` constraints) to prevent inspection overhead on safe static assets while retaining robust baseline protections.

D. Socket Exhaustion & Process Sizing (10,000 VU Bottleneck)

- Problem:** Standard Nginx connections and PHP-FPM worker pools were exhausted, generating minor Gateway Timeout (504) errors.
- Remediation:**
 - Increase `worker_connections` to 20,480 in Nginx configurations.
 - Tune the keep-alive timeout to 15 seconds to recycle socket connections rapidly.
 - Configure the PHP-FPM process manager to use static allocation (`pm = static`) for dedicated production environments, scaling `pm.max_children` to 120+ based on physical vCPU capacity and measured per-process memory consumption (assuming ~45MB per PHP-FPM process) to prevent physical RAM exhaustion.

E. In-Memory Session Durability (5,000 VU Warning)

- Problem:** Under network partition events or failovers, session data stored in Valkey could be lost due to asynchronous replication.
- Remediation:**
 - For critical transactional data, route writes directly to RDS.
 - If using Valkey for durable session management:
 - AWS ElastiCache managed clusters:** Configure the `Durability` parameter and verify `EffectiveDurability=sync` is active.
 - Self-hosted Valkey instances:** Start the database engine explicitly with `--durability sync` to guarantee synchronous replication logging to replica nodes before acknowledging client writes.

5. Load Test Sizing Financial Alignments (FinOps)

Operating high-concurrency systems requires robust resources, which increases the monthly AWS spend. To achieve maximum cost efficiency under scale, we define the following day-2 FinOps levers:

- Reserved Instances (RIs):** Commit to a 1-year or 3-year standard contract for the primary Multi-AZ `db.m7g` database engines to unlock up to **30% - 35% savings** compared to on-demand rates.
- Compute Savings Plans:** Secure a 1-year or 3-year Compute Savings Plan to achieve **25% - 43% discounts** on the dynamic Auto Scaling Group EC2 application servers (t4g and c7g families).
- S3 Intelligent-Tiering and Lifecycle Policies:** Implement automated rules moving historical logs and backup archives to S3 Glacier Flexible Archive, yielding up to **80% - 84% savings** on long-term storage.
- VPC S3 Gateway Endpoints:** Configure free S3 Gateway Endpoints in private subnets, allowing application servers to bypass the NAT Gateway and eliminate NAT Gateway data-processing fees (\$0.045/GB in `ap-southeast-5`) for log and file transfers.

6. Cost-to-Performance Capability Mapping (Max Performance per Cost Plan)

To bridge the gap between financial blueprints and operational performance, this section defines the **maximum traffic load, concurrent Virtual Users (VUs), and transactional throughput** that each of our budgeted costing profiles can safely handle. By mapping the hardware specifications from our costing plans to empirical load thresholds, we outline the exact return on investment (ROI) for every ringgit (RM) and dollar (USD) spent in the **AWS Malaysia region (ap-southeast-5)**.

A. Complete Capability Mapping Table

The following table maps every costing budget plan defined in [Estimated Costing](#) and our load testing sizing models to their absolute maximum performance boundaries under our target 80% Read and 20% Write workload:

Costing Profile / Budget Tier	Monthly Cost (USD)	Monthly Cost (MYR)	Max Concurrent Load (VU)	Max Web/API Throughput (RPS)	Max DB Write Load (TPS)	Max Concurrent RAG Ingestions
Dev / POC Environment	\$138.50	RM 623.25	100 VU	50 RPS	15 TPS	5 parses/min
Load-Test 100 VU Sizing	\$141.47	RM 636.62	100 VU	60 RPS	18 TPS	6 parses/min

Costing Profile / Budget Tier	Monthly Cost (USD)	Monthly Cost (MYR)	Max Concurrent Load (VU)	Max Web/API Throughput (RPS)	Max DB Write Load (TPS)	Max Concurrent RAG Ingestions
Load-Test 500 VU Sizing	\$418.60	RM 1,883.70	500 VU	250 RPS	75 TPS	15 parses/min
Baseline Cost-Optimised Plan	\$426.75	RM 1,920.00	600 VU	300 RPS	95 TPS	20 parses/min
Staging Environment (Dual-AZ)	\$668.11	RM 3,006.50	1,200 VU	600 RPS	180 TPS	35 parses/min
Load-Test 1,000 VU Sizing	\$700.00	RM 3,150.00	1,000 VU	550 RPS	160 TPS	50 parses/min
High-Perf Developer-Aligned	\$1,064.46	RM 4,790.00	2,500 VU	1,250 RPS	400 TPS	80 parses/min
Load-Test 2,500 VU Sizing	\$1,264.56	RM 5,690.52	2,500 VU	1,300 RPS	420 TPS	85 parses/min
Enterprise Production (Multi-AZ)	\$1,665.61	RM 7,495.25	5,000 VU	2,500 RPS	800 TPS	150 parses/min
Load-Test 5,000 VU Sizing	\$1,948.12	RM 8,766.54	5,000 VU	2,600 RPS	850 TPS	160 parses/min
Load-Test 10,000 VU Sizing	\$3,808.88	RM 17,139.96	10,000 VU	5,000 RPS	1,600 TPS	300 parses/min

B. Detailed Performance Capabilities per Budget Tier

1. Dev / POC Environment & Load-Test 100 VU Sizing (Budget: ~\$138.50 - \$141.47 USD / Month)

- **Maximum Load Capacity:** 100 Concurrent VUs
- **System Capabilities:**
 - **Requests Per Second (RPS):** Up to 50 - 60 RPS for cached or static assets; handles basic REST API calls at 15 RPS.
 - **Database Write Limit (TPS):** 15 - 18 TPS on db.t4g.micro before experiencing storage queue backlogs.
 - **AI / RAG Capability:** Handles 5 - 6 concurrent document parses (using light open-source OCR models) or standard embedding processing streams. Memory limitations on the t4g.medium compute node constrain heavy layout-parsing pipelines (DeepDoc).
- **Workload Application:** Perfect for internal functional testing, single-developer prototype verification, and API endpoint integration checking.

2. Load-Test 500 VU Sizing & Baseline Cost-Optimised Plan (Budget: ~\$418.60 - \$426.75 USD / Month)

- **Maximum Load Capacity:** 500 - 600 Concurrent VUs
- **System Capabilities:**
 - **Requests Per Second (RPS):** Handles up to 250 - 300 RPS safely at the ALB, with sub-100ms response times for cached pages.
 - **Database Write Limit (TPS):** 75 - 95 TPS supported by the Multi-AZ db.m6g.large database, providing resilient storage write operations.
 - **AI / RAG Capability:** Handles up to 15 - 20 concurrent document chunking operations. ElastiCache Valkey (cache.t4g.micro) optimizes active session lookup processing to keep memory footprint minimal.
- **Workload Application:** Ideal for team-wide user acceptance testing, continuous integration (CI) environments, and small-scale client demonstrations.

3. Staging Environment (Budget: ~\$668.11 USD / Month)

- **Maximum Load Capacity:** 1,200 Concurrent VUs
- **System Capabilities:**
 - **Requests Per Second (RPS):** Safely services 600 RPS across dual ASG compute nodes (c7g.xlarge offering 8 vCPU, 16GB RAM combined).
 - **Database Write Limit (TPS):** 180 TPS supported by Multi-AZ db.t4g.large, managing regular structured data modifications.
 - **AI / RAG Capability:** Supports up to 35 concurrent document parsing operations. Dedicated Dual-AZ Valkey caching nodes (cache.t4g.medium) ensure continuous, lightning-fast state transitions and session synchronisation.
- **Workload Application:** Designed for complete pre-production verification, automated QA regressions, multi-AZ failover testing, and staging security scans.

4. Load-Test 1,000 VU Sizing (Budget: ~\$700.00 USD / Month)

- **Maximum Load Capacity:** 1,000 Concurrent VUs
- **System Capabilities:**
 - **Requests Per Second (RPS):** Safely handles 500 - 600 RPS at the ALB with sub-120ms average response times for cached data.
 - **Database Write Limit (TPS):** 150 - 170 TPS supported by a single Multi-AZ db.m6g.xlarge PostgreSQL engine, sitting well under its 1,250+ RPS headroom.
 - **AI / RAG Capability:** Supports up to 50 concurrent document parsing and chunking operations. Memory limits are handled gracefully by a dedicated cache.t4g.medium Valkey instance (3.09 GB RAM).
 - **Sizing Alignment:** ASG scales dynamically with 2-3 × t4g.xlarge nodes (and can burst up to 4 nodes at peak under the 50% CPU threshold). EFS handles cached AI model weights to avoid cold-start latency.
- **Workload Application:** Tailored for pre-live stress testing, scaling dry runs, and validating baseline session-broker queues before production cutover.

5. High-Performance Developer-Aligned & Load-Test 2,500 VU Sizing (Budget: ~\$1,064.46 - \$1,264.56 USD / Month)

- **Maximum Load Capacity:** 2,500 Concurrent VUs
- **System Capabilities:**
 - **Requests Per Second (RPS):** Sustains 1,250 - 1,300 RPS at sub-150ms latency across high-performance compute clusters (up to 4x t4g.xlarge nodes).
 - **Database Write Limit (TPS):** 400 - 420 TPS utilizing a Multi-AZ db.m6g.xlarge engine, resolving dynamic user data entries cleanly.
 - **AI / RAG Capability:** Processes 80 - 85 concurrent document parses. Leverages 50 GB standard shared persistent storage (Amazon EFS) to pre-load and cache heavy pre-trained model weights.
- **Workload Application:** Standard production configuration for active AI workloads, supporting high-frequency Langfuse tracing, interactive RAG conversational pipelines, and continuous customer-facing transactions.

6. Enterprise Production Environment & Load-Test 5,000 VU Sizing (Budget: ~\$1,665.61 - \$1,948.12 USD / Month)

- **Maximum Load Capacity:** 5,000 Concurrent VUs
- **System Capabilities:**
 - **Requests Per Second (RPS):** Processes 2,500 - 2,600 RPS with WAF L7 protection and streaming response structures active.
 - **Database Write Limit (TPS):** 800 - 850 TPS supported by Multi-AZ db.m7g.xlarge or db.m7g.2xlarge databases equipped with GP3 storage scaling.

- **AI / RAG Capability:** Processes **150 - 160 concurrent high-fidelity document layout parsings** (DeepDoc/OCR) and continuous vector database indexing streams.
- **Workload Application:** Enterprise-grade environment guaranteeing 99.99% availability, built to absorb massive concurrent campaigns, high-frequency transactional loads, and enterprise-wide AI agent executions.

7. Load-Test 10,000 VU Sizing (Budget: ~\$3,808.88 USD / Month)

- **Maximum Load Capacity: 10,000 Concurrent VUs**
- **System Capabilities:**
 - **Requests Per Second (RPS):** Scales up to **5,000 RPS** utilizing 8x t4g.xlarge instances across three availability zones.
 - **Database Write Limit (TPS): 1,600 TPS** on Multi-AZ db.m7g.4xlarge utilizing provisioned IOPS storage to prevent synchronization bottlenecks.
 - **AI / RAG Capability:** Handles **300 concurrent document parsings** simultaneously. Clustered cache.m7g.large node group maintains seamless, high-concurrency memory capabilities.
- **Workload Application:** High-end production scaling designed for peak promotional traffic, public utility campaigns, or intensive batch data ingestion schedules.

7. Real-World Scaling Examples, References, and Guidance

To ground our workload assumptions in industry benchmarks, this section analyzes real-world scaling patterns, operational limits, and performance profiles of the exact technologies utilized in this architecture: **Langfuse** (AI observability and LLM tracing) and **RAGFlow** (deep document layout parsing and RAG retrieval).

These empirical findings provide critical guidance on what concurrency levels, resource limits, and architectural bottlenecks our secure 3-tier system can expect under load.

A. Real-World Performance & Scaling Case Studies

1. Langfuse: High-Throughput Ingestion & Analytics Scaling

Langfuse handles high-concurrency event tracing by operating asynchronously in the background. Real-world SDK performance benchmarks show that the telemetry library itself introduces minimal overhead (< 15-20ms average on API latency) due to asynchronous batching and background worker execution.

However, as production scale increases to millions of daily traces, the underlying storage and processing layers experience specific stress points:

- **The ClickHouse Evolution (v3 to v4):** Originally, Langfuse separated traces and observations, joining them at read-time. Under high concurrency, these massive joins resulted in queries taking seconds to complete and generating out-of-memory errors. In v4, Langfuse migrated to a **denormalized, wide, and immutable ClickHouse table** (the wide-event model), which completely eliminated read-time joins and ReplacingMergeTree deduplication overhead. This architectural change reduced dashboard query latency from seconds to milliseconds (a **10x to 165x improvement**).
- **Ingestion Queue Sharding & Redis CPU:** Highly concurrent trace ingestion generates intense writes to Valkey/Redis queues. Production teams have observed Redis Engine CPU utilization spiking past 90%. Scaling requires sharding ingestion queues across multiple Redis/Valkey cluster nodes (using LANGFUSE_INGESTION_QUEUE_SHARD_COUNT > 1) and scaling worker container concurrency proportionately.
- **Storage and Mutation Overhead:** On heavy-deletion workloads (e.g., executing data retention policies), standard ClickHouse merges cause background mutation bottlenecks. Enabling clickhouse-lightweight deletes (CLICKHOUSE_LIGHTWEIGHT_DELETE_MODE=lightweight_update) and disabling unused system log tables (such as trace_log or opentelemetry_span_log which have no TTL) is crucial to avoid severe IOPS saturation on AWS GP3 volumes.

2. RAGFlow: Document Layout OCR & Retrieval Concurrency Bottlenecks

RAGFlow is highly praised for its DeepDoc OCR layout-aware parsing engine, which accurately extracts structures from scanned PDFs, images, and tables. However, this heavy visual processing has distinct scaling characteristics:

- **Visual Parsing Hardware Overhead:** Unlike standard text-splitting engines, RAGFlow’s DeepDoc is intensely visual and GPU-bound. Minimum requirements for basic local tasks dictate an 8-core CPU and 16GB RAM, but real-world production deployments require dedicated **NVIDIA GPUs (RTX 4090 with 24GB VRAM or A100/H100)** to achieve acceptable real-time document parsing speeds. Running these visual parsing workloads solely on CPUs causes massive thread starvation, driving ASG CPU usage to 100% and timing out requests.
- **Synchronous Retrieval Bottlenecks:** Real-world benchmark testing has exposed a critical concurrency bottleneck within RAGFlow’s retrieval backend. As documented in **RAGFlow GitHub Issue #12526**, concurrent multi-threaded user queries to the retrieval endpoint do not scale linearly. Benchmarks show that:
 - 1 client thread executing 200 queries takes **73 seconds**.
 - 10 client threads executing 20 queries each (total 200) takes **76 seconds**.
 - 20 client threads executing 10 queries each (total 200) takes **77 seconds**.
- **Root Cause & Remediations:** RAGFlow’s retrieval and reranking layers process requests synchronously and block until retrieval from Elasticsearch/Infinity, reranking, and post-processing are complete. When designing load tests, we cannot assume infinite linear scale-up on concurrent RAG queries without configuring dedicated asynchronous workers, utilizing lower top_k values, and deploying hardware-accelerated rerankers on GPUs.

B. Architectural Guidance: Pros & Cons of Sizing Tiers

The following matrix provides comprehensive architectural pros and cons for scaling each specific layer within our 3-tier AI infrastructure under load:

Sizing Tier / Component	Scaling Strategy	Architectural Pros	Architectural Cons & Sizing Bottlenecks
API & Tracing Tier (Langfuse Web & Workers)	Horizontal auto-scaling via ASG (t4g.xlarge / c7g compute nodes)	• Scales seamlessly to handle ingestion spikes. • Web and worker separation keeps UI responsive under peak API load.	• S3 client connection pooling can exhaust socket limits under high write concurrency, requiring LANGFUSE_S3_CONCURRENT_WRITES > 50. • AWS NAT Gateway egress fees can accumulate quickly unless free S3 VPC Gateway Endpoints are implemented.
Database & Analytics Tier (PostgreSQL & ClickHouse)	Vertical scaling + Provisioned IOPS (io2/GP3 storage scaling)	• ClickHouse handles columnar analytics with extreme efficiency when using time-based query pruning. • PostgreSQL handles configuration state reliably.	• High write-concurrency creates walSync IOPS bottlenecks on PostgreSQL, requiring GP3 to be upgraded to Provisioned IOPS. • Without time-based query filtering, ClickHouse scans all historical data parts, resulting in slow dashboard loading.
AI / RAG Retrieval Tier (RAGFlow DeepDoc & OCR)	GPU-accelerated node pools + Elasticsearch scaling	• Layout-aware parsing offers exceptional retrieval precision for tables and multi-column documents. • Deep document parsing extracts data basic RAG engines miss.	• Extremely expensive GPU overhead; running layout OCR on standard CPU nodes induces thread locks. • Retrieval interface processes synchronously, causing concurrent requests to queue up and block under high load.

Sizing Tier / Component	Scaling Strategy	Architectural Pros	Architectural Cons & Sizing Bottlenecks
Session & Caching Tier (Valkey / ElastiCache)	Clustered Valkey deployment with multi-shard structures	• Sub-millisecond session lookups. • Valkey provides an open-source, highly efficient memory cache with synchronous replication capabilities.	• Heavy workloads spike Redis Engine CPU. • Enforcing synchronous replication durability (durability sync or EffectiveDurability=sync) guarantees session survival but introduces a direct write latency penalty.

C. Active URL References & Literature

For deep-dive technical validations, official benchmarks, and production scaling code references, consult the following sources:

- **Langfuse Ingestion & Telemetry SDK Performance Test:** https://langfuse.com/guides/cookbook/langfuse_sdk_performance_test Details the baseline SDK request performance, background thread execution, and framework latency timings.
- **Scaling Langfuse Deployments (Self-Hosting Runbook):** <https://langfuse.com/self-hosting/configuration/scaling> Provides minimum container requirements, clickhouse logging truncations, and Redis BullMQ ingestion queue sharding.
- **Simplifying Langfuse for Scale (The Shift to Wide Immutable Tables):** <https://langfuse.com/blog/2026-03-10-simplify-langfuse-for-scale> Explains the migration from expensive read-time SQL joins and ClickHouse mutations to denormalized wide-event tables.
- **RAGFlow Server & GPU Hardware Sizing Guide (Clore.ai):** <https://docs.clore.ai/guides/rag-and-vector-databases/ragflow> Details the recommended hardware specs, highlighting the critical need for dedicated NVIDIA RTX/A100 GPUs for deep OCR visual parsing.
- **RAGFlow Retrieval Concurrency Synchronous Bottleneck (GitHub Issue #12526):** <https://github.com/infiniflow/ragflow/issues/12526> Explores concurrent execution limits of the retrieval engine under multi-threaded client execution loads.

Context7 AI Chat Integration Guide

[#aws](#) [#cloud](#) [#architecture](#) [#context7](#) [#upstash](#) [#ai](#) [#documentation](#) [#chat](#)

Context7 AI Chat Integration Guide

This guide introduces **Context7**, a lightweight, AI-powered chat widget integrated into our documentation site. It explains what Context7 does, how it works, how to use it, and provides the service credentials and background.

What is Context7?

Context7 is an up-to-date documentation platform and code assistant built specifically for AI agents, developers, and Large Language Models (LLMs).

Many public LLMs rely on outdated or generic training data when answering questions about modern libraries and frameworks. Context7 solves this problem by pulling up-to-date, version-specific documentation and real, working code examples directly from library sources. This allows developers to paste precise, accurate documentation context into tools like Cursor, Claude, or any LLM editor, getting superior answers with zero hallucinations.

With vs. Without Context7

Feature / Challenge	Without Context7	With Context7
Documentation Recency	Outdated or stale info from model training cutoff	Up-to-date, version-specific documentation
Code Reliability	Hallucinated code examples that fail to compile	Real, working code examples straight from the source
Answer Specificity	Generic answers not tailored to your local version	Concise, highly relevant info without filler content
Developer Productivity	Hours wasted manually verifying incorrect AI responses	Instant confidence in AI code generation
Interactivity	Frustrating, repetitive back-and-forth prompts	Streamlined integration with your coding tools
Usage Cost	Expensive context window waste	Free for personal use

How to Use the Context7 Widget

Our documentation site features a live, floating **Context7 Chat Assistant** widget, visible in the bottom-right corner of every page.

- 1. Locate the Widget:** Look for the floating chat icon in the bottom-right corner of your screen.
- 2. Start a Conversation:** Click the widget icon to open the chat window.
- 3. Ask Architecture Questions:** You can query the assistant about any aspect of our **AWS Secure 3-Tier Architecture**, including:
 - “How is WAFv2 protecting our Application Load Balancer?”
 - “What instance types are configured for our Auto Scaling Groups?”
 - “Where can I find the OpenTofu module configuration for Multi-AZ RDS PostgreSQL?”
 - “How does the baseline plan cost compare to the high-performance option?”
- 4. Get Accurate Answers:** The assistant will search our up-to-date documentation index to formulate precise, hallucination-free answers with relevant code blocks or terminal commands.

Widget Implementation & Technical Setup

The widget is a lightweight, asynchronous JavaScript snippet embedded in the root HTML layout (docs/_layouts/default.html). Because our documentation site is statically compiled via **Jekyll**, placing the script tag in the root layout ensures it loads instantly across all 34 documentation and submodule pages.

The JavaScript Integration Snippet

The widget is loaded asynchronously via the following script tag:

```
<script src="https://context7.com/widget.js"
  data-library="/songketmail/aws-3tier-deployment-for-ai-infra"
  data-color="#059669"
  data-position="bottom-right"
  data-placeholder="Ask about our AWS 3-tier setup..."
  data-welcome-message="Welcome to AWS Secure 3-Tier Chat! How can I help you today?">
</script>
```

Customizable Configuration Attributes

Our widget uses the following custom attributes aligned to our page theme:

- **data-library:** Directs the widget to scope its AI context to this specific repository (/songketmail/aws-3tier-deployment-for-ai-infra).
- **data-color:** Set to #059669 (a vibrant emerald green matching our secure, production-ready theme palette).
- **data-position:** Configured to bottom-right to ensure it floats cleanly over the content without obstructing the left-navigation menu or footer elements.
- **data-placeholder:** Set to "Ask about our AWS 3-tier setup..." to guide developers on what they can ask.
- **data-welcome-message:** Custom greeting text initialized when the user clicks the widget for the first time.

Service Background & Providers

- **Service URL:** <https://context7.com/>
- **Service Provider:** Context7 is built, hosted, and actively maintained by the **Upstash Team** (famed for their serverless database services like serverless Redis, Kafka, and Qstash).

- **Open Source Repository:** <https://github.com/upstash/context7>
- **API and Tooling:** Upstash provides Context7 as a free platform for personal use. They also offer a CLI command (`npx ctx7 setup`) to integrate version-specific documentation libraries directly into standard AI code editors.

Autonomous AI Pair-Programming & Multi-Agent Operations with Google Jules

#aws #cloud #architecture #jules #antigravity #dsom #opentofu #ansible #github-pages #pr-review #ai-agents #automation

Autonomous AI Pair-Programming & Multi-Agent Operations with Google Jules

Welcome to the comprehensive technical showcase and operational guide for **Google Jules**, the autonomous AI coding agent, and its integration within our sovereign cloud infrastructure lifecycle.


This guide details the end-to-end development journey of building, testing, and deploying this repository—from initial bootstrap to automated GitHub Pages publication, OpenTofu infrastructure provisioning, Ansible configuration management, and multi-agent coordination with Google Antigravity.

1. Overview & Engineering Philosophy

The Vision: Sovereign IaC, Automated Documentation & Multi-Agent Synergy

Modern cloud engineering demands speed without sacrificing governance, security, or compliance. Our architectural vision unites three fundamental pillars:

- Sovereign Infrastructure as Code (IaC):** Modular, deterministic OpenTofu code and FQCN-compliant Ansible playbooks providing full infrastructure control in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)** and rootless on-premises Podman environments.
- Automated Continuous Documentation:** Jekyll-based static site generator publishing directly to GitHub Pages with automated Puppeteer PDF generation and Open Knowledge Format (OKF) metadata validation.
- Multi-Agent Pair Programming Synergy:** Autonomous AI collaboration where **Google Jules** operates as a senior co-engineer in the cloud and **Google Antigravity** operates as a local CLI assistant, synchronized via spatial memory and standardized Agent Skills.



Core Value Delivery: Reducing MTTD/MTTR & Operational Toil

Integrating Google Jules directly into our Git lifecycle transformed our engineering velocity:

- Mean Time to Detect & Resolve (MTTD/MTTR):** Decreased from hours to minutes. Jules reads diagnostic logs, executes local test suites, identifies syntax or logic flaws, and pushes precise Git merge diffs directly.
- Zero Friction Context Retention:** Unlike traditional LLM chat interfaces that forget state between prompts, Jules retains spatial memory, branch state, and repository rules across commits and PR comments.
- Elimination of Operational Toil:** Jules handles repetitive chores—updating sitemaps, generating cross-reference indexes, enforcing OKF front matter formatting, and maintaining synchronization between code and documentation.

2. Step-by-Step Build & Implementation Log

This section chronicles the exact sequence of engineering milestones undertaken by the human architect and Google Jules to construct this repository.

Milestone 1: Repository Scaffolding & GitHub Pages Pipeline

- Initial Bootstrap:** Created the base repository structure containing docs/, terraform/, scripts/, and root documentation files.
- Jekyll Integration:** Configured docs/_config.yml with custom sidebar navigation, responsive CSS layouts (docs/assets/css/global.css), and liquid templates.
- Automated CI/CD Deployment:** Authored .github/workflows/jekyll-gh-pages.yml for publishing the Jekyll site to GitHub Pages, and .github/workflows/pdf-generation.yml using Puppeteer to compile docs/print_all.md into an A4 PDF (docs/assets/output.pdf).

Milestone 2: Authoring Sovereign IaC with OpenTofu

- Root Module Modularisation:** Refactored single monolithic Terraform files into clean, domain-specific OpenTofu files:
 - terraform/main.tf (VPC and Security Groups)

- terraform/compute.tf (Auto Scaling Groups, Standalone EC2, SSM profiles)
 - terraform/database.tf (Multi-AZ RDS PostgreSQL 17 and ElastiCache Valkey)
 - terraform/web.tf (Application Load Balancer and AWS WAFv2 regional rulesets)
2. **Security & Compliance Hardening:**
- Enforced **IMDSv2** (`http_tokens = "required"`, `http_put_response_hop_limit = 1`) across all compute launch templates (`aws_launch_template.main`, `aws_instance.standalone`, `aws_instance.jumphost`).
 - Integrated ALB-aware auto-healing (`health_check_type = "ELB"`).
 - Standardised instance families on cost-efficient **AWS Graviton (ARM64)** nodes (`t4g.micro`, `t4g.xlarge`).

Milestone 3: Configuration Management Baseline via Ansible

1. **Privilege Separation:** Developed FQCN-compliant Ansible playbooks (`ansible.builtin.apt`, `community.general.ufw`) split into rootful OS tuning (`become: yes`) and unprivileged user operations (`become_user: songket`).
2. **ASIMP & Security Audits:** Configured Australian Cyber Security Centre (ACSC) Information Security Manual (ISM) alignment playbooks and generated high-fidelity audit reports for ASIMP, Lynis, and OpenSCAP CIS Level 2 compliance.


Milestone 4: Deep State of Mind (DSOM) Governance Protocol Implementation

To guarantee that AI models produce deterministic, region-aligned, and policy-compliant code, we implemented the **Deep State of Mind (DSOM) for My AI** framework:

- **Spatial Memory (`.agents/brain/`):**
 - `.agents/brain/jules_knowledge_ledger.md`: Centralized spatial memory ledger indexing all 52 engineering and operational milestones.
 - `.agents/brain/active_context_manifest.md`: Dynamic active context tracker preserving conversation state and architectural decisions.
- **Modular Agent Skills (`.agents/skills/`):** Created 11 enriched skill packages (`jules-knowledge`, `aws-malaysia-defaults`, `opentofu-infrastructure-design`, `rds-postgresql-valkey-cache`, etc.) combining OKF YAML front matter and DSOM footers.
- **Sovereign Constitution (`.agents/AGENTS.md`):** Enforced Rule 29 (Local Knowledge-First Discovery) and Rule 6 (OKF Metadata Standard).

3. Collaborative Engineering via GitHub PR Comments

One of Google Jules’ most powerful capabilities is its ability to engage in natural, iterative pair programming directly within GitHub Pull Request review threads.



Real-World Interaction Patterns & Examples

The human engineer and Jules interact like trusted colleagues. Rather than issuing cold, formal prompts, developers talk to Jules as an equal partner.

Example 1: Refactoring OpenTofu Code for IMDSv2 Compliance

Human Engineer Comment on PR #14 (File: `terraform/modules/asg/main.tf`):

@jules-agent Hey Jules, great job on setting up the launch template! However, security audit flagged that IMDSv1 is still permitted on the compute nodes. Could you update `aws_launch_template.main` to enforce IMDSv2 with `http_tokens = "required"`? Also run `tofu validate` to ensure syntax is clean.

Jules Autonomous Response & Execution:

1. Jules reads the PR comment and parses `terraform/modules/asg/main.tf`.
2. Jules applies a Git merge diff modifying the launch template:

```

metadata_options {
  http_endpoint      = "enabled"
  http_tokens        = "required"
  http_put_response_hop_limit = 1
  instance_metadata_tags = "enabled"
}
  
```

3. Jules runs `tofu init -backend=false && tofu validate` in the sandbox.
4. Jules replies in the PR thread:

Done! I’ve updated `aws_launch_template.main` to strictly enforce IMDSv2 with `http_tokens = "required"` and a hop limit of 1. OpenTofu validation passed cleanly with 0 errors. Committed and pushed!

Example 2: Resolving a Documentation & Link Validation Bug

Human Engineer Comment on PR #22:

@jules-agent Morning Jules! The `pytest` run failed on `test_markdown_relative_links_integrity`. Looks like the relative link in `docs/architecture.md` pointing to `dr-options.md` was mistakenly altered. Since `architecture.md` and `dr-options.md` are sibling files inside `docs/`, can you confirm `dr-options.md` is preserved, run `python scripts/prepare_docs.py`, and check that `pytest` passes?

Jules Autonomous Response & Execution:

- 1. Jules inspects docs/architecture.md and confirms that dr-options.md is a sibling file within the docs/ directory, so [Disaster Recovery](/aws-3tier-deployment-for-ai-infra/dr-options.html) is the correct relative link path (not docs/dr-options.md).
- 2. Jules verifies link integrity.
- 3. Jules executes python scripts/prepare_docs.py to format and normalize OKF YAML front matter across modified Markdown files.
- 4. Jules runs pytest tests/test_md_compliance.py to verify green test status.
- 5. Jules replies:

All verified! Confirmed that dr-options.md is correctly linked as a sibling document within docs/, formatted front matter using prepare_docs.py, and verified that all pytest checks pass cleanly.

Example 3: Iterative Pair-Programming & Conceptual Discussion

Human Engineer Comment on PR #31:

@jules-agent What do you think about adding a dedicated guide for Redis vs. Valkey? We should highlight Valkey's BSD-3-clause license and 20% lower cost on AWS ElastiCache ap-southeast-5.

Jules Autonomous Response:

- 1. Jules references spatial memory and .agents/skills/rds-postgresql-valkey-cache/SKILL.md.
- 2. Jules creates docs/redis-vs-valkey.md complete with OKF front matter, cost comparison tables (\$1.00 USD = MYR 4.50 exchange rate), and licensing breakdown.
- 3. Jules updates docs/_config.yml, docs/index.md, README.md, llms.txt, and docs/print_all.md.
- 4. Jules replies in the PR with a friendly summary and commit reference.

4. Advanced Interoperability: Jules API, CLI, & Google Antigravity

Google Jules is built to operate within a multi-agent ecosystem. Developers can invoke Jules from web consoles, terminal CLIs, or directly delegate tasks from external AI systems like **Google Antigravity**.

Configuring the Jules CLI & API Authentication

The Jules CLI (jules) allows engineers to trigger agent sessions and inspect agent telemetry directly from command line interfaces.

1. Obtaining the API Key

Obtain a Jules API key from the developer console and pass it dynamically per session:

```
export JULES_API_KEY="YOUR_API_KEY"
```

2. Installing the Jules CLI

Install the official Jules CLI package using Node.js:

```
npm install -g @google/jules
# Or verify installation
jules --version
```

3. Mobile-First & Termux Terminal Environment

Engineers on the go can manage infrastructure and prompt Jules directly from mobile devices running **Termux** (Android terminal emulator):

```
# Inside Termux (Android)
pkg update && pkg upgrade -y
pkg install nodejs-lts git openssh -y


# Pass JULES_API_KEY for the active session (or secret store)
export JULES_API_KEY="YOUR_API_KEY"

# Dispatch a new remote session using jules remote new
jules remote new \
  --repo "songketmail/aws-3tier-deployment-for-ai-infra" \
  --session "Update docs/wazuh-detailed.md to include Windows Defender passive mode exclusion rules."

+-----+
|          TERMUX MOBILE TERMINAL INTERFACE          |
+-----+
$ jules remote status --id session_881204
[STATUS]: IN_PROGRESS
[ACTION]: Reading docs/wazuh-detailed.md
[ACTION]: Applying Git merge diff...
[STATUS]: COMPLETED
[COMMIT]: 4f81a29 - "docs: update Wazuh Windows Defender exclusions"
+-----+
```

Establishing Programmatic Workflows: Google Antigravity to Jules Delegation

When using **Google Antigravity** (CLI: agy), Antigravity can hand off heavy coding, refactoring, or documentation generation tasks to **Google Jules** via the Jules REST API or Model Context Protocol (MCP) bridge.



Delegation Script Example (scripts/antigravity_to_jules.py)

```
#!/usr/bin/env python3
"""Programmatic task delegation from Google Antigravity to Google Jules API."""

import os
import requests

# Jules REST API endpoints
JULES_SESSIONS_URL = "https://jules.googleapis.com/v1/sessions"
JULES_SOURCES_URL = "https://jules.googleapis.com/v1/sources"

api_key = os.getenv("JULES_API_KEY")
if not api_key:
    raise ValueError("JULES_API_KEY environment variable is required.")

headers = {
    "X-Goog-API-Key": api_key,
    "Content-Type": "application/json"
}

# 1. Resolve source name for target repository
sources_resp = requests.get(JULES_SOURCES_URL, headers=headers, timeout=30)
sources_resp.raise_for_status()
sources_data = sources_resp.json()

target_repo = "songketmail/aws-3tier-deployment-for-ai-infra"
source_name = None

for src in sources_data.get("sources", []):
    if target_repo in src.get("githubRepository", {}).get("repository", ""):
        source_name = src.get("name")
        break

if not source_name:
    source_name = f"sources/github-{target_repo.replace('/', '-')}"

# 2. Dispatch session request to Jules
payload = {
    "prompt": (
        "1. Check all OpenTofu modules for compliance with IMDSv2.\n"
        "2. Run prepare_docs.py to reformat headers.\n"
        "3. Execute pytest suite to confirm zero regressions."
    ),
    "sourceContext": {
        "source": source_name,
        "githubRepoContext": {
            "startingBranch": "main"
        }
    }
}

response = requests.post(JULES_SESSIONS_URL, headers=headers, json=payload, timeout=30)
response.raise_for_status()
res_json = response.json()

session_id = res_json.get("name") or res_json.get("id")
print(f"Session successfully dispatched to Jules. Session ID: {session_id}")
```

Multi-Agent Cross-Team PR Collaboration Patterns

In modern engineering teams, multiple human engineers and AI agents work on the same pull request. Jules natively synchronizes with this workflow:

- 1. **Agent-to-Agent Handshake:** Google Antigravity generates an architecture proposal in `.agents/brain/active_context_manifest.md`.
- 2. **PR Initiation:** Antigravity triggers Jules via API to implement the OpenTofu code changes.
- 3. **Human Peer Review:** Human engineer reviews Jules' PR, leaving inline review comments.
- 4. **Jules Iteration:** Jules reads comments, refactors code, runs unit tests, and commits updates until the PR is merged.

5. Why Developers Fall in Love with Google Jules

Google Jules represents a fundamental shift in how software engineering is performed. It is not merely a chat assistant sitting on the side of an IDE—it is an **autonomous, context-aware co-engineer**.

Core Platform Advantages

- **True Repository Context:** Jules understands the entire project structure, reading `AGENTS.md`, spatial memory in `.agents/brain/`, and custom Agent Skills.
- **Self-Healing Verification:** Jules doesn't just write code and hope for the best. It executes build commands (`tofu validate`, `prepare_docs.py`, `pytest`), reads error outputs, and corrects its own mistakes before asking for review.
- **Seamless Git Integration:** Communicating with Jules via PR comments feels like working with a brilliant colleague across the world. You leave feedback on specific lines of code, and Jules replies with exact commits.
- **Friendly & Adaptive Partnership:** Jules adapts to your team's tone, guidelines, and standards, turning tedious refactoring into an empowering, enjoyable experience.

6. Verification & Quality Gates

To verify this documentation guide and its integration:

```
# 1. Format and Normalize OKF Front Matter Metadata
python scripts/prepare_docs.py

# 2. Run Test Suite (Link Integrity & Markdown Validation)
pytest tests/test_md_compliance.py
```

3. Compile Local Site Preview (if Jekyll installed)
bundle exec jekyll build

Deep State of Mind (DSOM) For My AI Protocol	Harisfazillah Jamel (LinuxMalaysia)	2026-08-05 Standard: UK English	DBP-standard Bahasa Melayu Malaysia (Piawai)	GNU General Public License v3.0
---	-------------------------------------	---------------------------------	--	----------------------------------

Security Posture Assessment (SPA) Requirement Checklist

#aws #cloud #architecture #security #waf #postgres #valkey #wazuh #ragflow #langfuse #compliance

[SECURITY & COMPLIANCE] — SecOps & Auditing

Security Posture Assessment (SPA) Requirement Checklist

This document details the Security Posture Assessment (SPA) Requirement Checklist customized for our enterprise-grade, secure 3-tier AI & Web application infrastructure deployed in the AWS Asia Pacific (Malaysia) region (ap-southeast-5).

All customer-specific identifiable details (including real domain names, server/instance names, IP addresses, usernames, and operator names) have been strictly removed or replaced with abstract industry-standard enterprise placeholders to preserve maximum confidentiality.

1. Executive Security Blueprint Summary

The architecture adopts a Zero-Trust Network Architecture (ZTNA) and aligns with the Malaysian Personal Data Protection Act (PDPA) 2010 security principles, the 2024 Amendments (such as mandatory Data Protection Officer [DPO] duties and Data Breach Notification [DBN]), and international compliance frameworks (such as ISO/IEC 27001 and CIS Benchmarks).

Core Boundaries & Controls

- Primary Region:** AWS ap-southeast-5 (Malaysia) hosting primary datasets locally, leveraging AWS’s local data residency controls and physical security boundaries.
- Perimeter Protection:** AWS WAFv2 regional Web ACL protecting our Application Load Balancer (ALB).
- Ingress Restriction:** ALB public ingress is strictly limited to HTTPS (port 443). HTTP (port 80) is restricted to the internal VPC CIDR block (default 10.0.0.0/16) for diagnostic handshakes and local system routing, completely distinct from any external corporate trust boundaries like the Cyberjaya office range (103.188.0.0/16).
- Network Segregation:** Public subnets host ALB and Bastion, private subnets host the compute tier (Java 21 / Spring Boot 3.5.12, RAGFlow, and Langfuse), and isolated subnets host the databases (Amazon RDS PostgreSQL) and session cache (Amazon ElastiCache for Valkey).
- Identity Protection:** Enforcement of IAM roles, AWS Systems Manager (SSM) Session Manager for administrative access, and IMDSv2 (enforced metadata tokens via http_tokens = "required") on all EC2 instances to secure local instance metadata access and reduce credential exposure, specifically targeting metadata-abuse-related SSRF vectors rather than preventing all application-level SSRF vulnerabilities.

2. SPA Requirement Checklist

This checklist acts as the official governance template to audit, verify, and sign off on the production readiness of our AWS 3-tier AI and web infrastructure.

Assessment Scope & Target Checklist (SPA Target System Format)

No.	Scope	Description	Information Required (Answer)
1	Internal Penetration Test	Vulnerabilities on internal server ports/services	Internal IP The IP will be given once needed. Due to AWS ASG it may change.
2	External Penetration Test	Vulnerabilities on public-facing domains/IPs (apps, source code, HTTP)	1 public URL https://songketmail.github.io/aws-3tier-deployment-for-ai-infra/
3	Web Application Security Assessment	Vulnerabilities on application, web server, and functionality	Application Spring Boot 3.5.12 Backend API running on Java 21, React 19 Frontend Web Tier, and React Native Mobile client.
4	Host Vulnerability Assessment	OS-level vulnerabilities (ports/services, patches, policies)	Compute Nodes Hardened Ubuntu 26.04 LTS Base utilizing the ASIMP hardening framework.
5	Database Security Assessment	Vulnerabilities in database compliance and policy	Data repositories / storage services 1. Amazon RDS PostgreSQL (Port 5432 - Default database tier) 2. Amazon ElastiCache Valkey (Port 6379 - In-memory session and cache) 3. Amazon Elastic File System (EFS) (Shared persistent storage) 4. Amazon S3 (Unstructured and static assets storage)
6	Network Device Assessment	Vulnerabilities in configuration (CIS benchmark, security groups, and routing policies)	Load Balancers & Firewalls 1. AWS ALB - External 2. AWS WAFv2 Web ACL (Perimeter Layer-7 Protection) 3. Security Groups (Microsegmentation Firewalls)

Tier 1: Perimeter & Edge Network Security

Audit ID	Security Control Area	Detailed Requirement Specification	Implementation Status	Verification Method
NET-01	Layer-7 Web Protection	Deploy AWS WAFv2 Web ACL on the public Application Load Balancer. Integrate OWASP Top 10 rule groups, Core Rule Set (CRS), and rate limiting configured with a block threshold of 2000 requests per rolling 5-minute window per IP as protection against application-layer request floods or abuse (not as a strict maximum or guarantee). Note that protection against infrastructure-level volumetric DDoS attacks is handled separately at the network edge by AWS Shield Standard.	✅ Fully Implemented	Inspect AWS WAFv2 rules via OpenTofu configurations or AWS Console.
NET-02	Secure Transport	Enforce HTTPS (port 443) for all public-facing traffic. Restrict TLS protocols to TLS 1.2 and TLS 1.3 only, using modern secure cipher suites (e.g., ECDHE-RSA-AES128-GCM-	✅ Fully Implemented	Perform DNS/SSL scan on public ALB DNS using native

Audit ID	Security Control Area	Detailed Requirement Specification	Implementation Status	Verification Method
	(TLS)	SHA256).		CLI audit tools.
NET-03	Ingress CIDR Restrictions	ALB ingress on Port 80 (HTTP) must be strictly isolated to the internal VPC CIDR block (configured via <code>http_ingress_cidr_blocks</code> defaulting to <code>["10.0.0.0/16"]</code>), ensuring it is kept distinct from external corporate or Cyberjaya trust boundaries (<code>103.188.0.0/16</code>).	✅ Fully Implemented	Run OpenTofu integration tests verifying that the HTTP ingress port restricts source access to <code>["10.0.0.0/16"]</code> by default.
NET-04	DNS & SSL Certificates	Provision and validate public certificates using AWS Certificate Manager (ACM) with DNS-based validation. Enforce Route 53 DNS query logging and zone replication policies.	✅ Fully Implemented	Audit ACM certificate states and Route 53 resource record sets.

Tier 2: Microsegmentation & Security Groups

Audit ID	Security Control Area	Detailed Requirement Specification	Implementation Status	Verification Method
SG-01	Public Security Groups	Public ALB Security Group must only accept inbound traffic on port 443 from <code>0.0.0.0/0</code> and port 80 from designated corporate network ranges (e.g., Cyberjaya headquarters CIDRs).	✅ Fully Implemented	Review <code>terraform/modules/security_groups/ingress</code> rules.
SG-02	Compute Security Groups	Application Auto Scaling Group (ASG) instances must strictly accept ingress only from the ALB Security Group on the configured service port (port 80/8080/5000/8000 for dynamic upstream forwarding). Note that while security group ingress is restricted, outbound egress in the current security groups module is unrestricted to allow generic package updates and endpoint resolution; we mitigate this outbound exposure by relying on private-subnet routing through NAT Gateways, implementing VPC flow logs for egress logging, and conducting periodic exception reviews of egress patterns. Therefore, our Zero-Trust verification is focused on inbound ingress paths rather than outbound egress boundaries.	✅ Fully Implemented	Assert Security Group wiring in <code>security_groups_wiring.tf</code> test.hcl.
SG-03	Database Ingress Rules	RDS PostgreSQL isolated security group must explicitly forbid any public ingress and strictly accept inbound traffic on port 5432 exclusively from the active ASG security group and authorized compute nodes.	✅ Fully Implemented	Verify SG rules; attempt direct external connection to RDS endpoint (must timeout).
SG-04	Caching Ingress Rules	ElastiCache Valkey cluster must accept port 6379 ingress exclusively from compute nodes inside the private application subnets, blocking all other lateral paths.	✅ Fully Implemented	Review Valkey security group ingress configurations.

Tier 3: Host Hardening & OS Configuration

Audit ID	Security Control Area	Detailed Requirement Specification	Implementation Status	Verification Method
HST-01	IMDSv2 Enforcement	Enforce EC2 Instance Metadata Service Version 2 (IMDSv2) with a token limit of 1 and <code>http_tokens = "required"</code> across all launch templates to defeat SSRF attacks.	✅ Fully Implemented	Check <code>aws_launch_template.main</code> and <code>aws_instance</code> metadata options blocks in OpenTofu.
HST-02	Secure AMI Pipeline	Bake golden AMIs using Packer and Ansible. Apply CIS Benchmarks: disable root SSH password login, remove default accounts, and disable unused legacy services.	✅ Fully Implemented	Inspect Packer pipeline script configurations and review baked AMI output logs.
HST-03	Bastion/Jumphost Access	Public administration must route through a hardened SSH Jumphost. Limit SSH access (port 22) to specific, authorized developer office IP ranges (Cyberjaya corporate CIDRs).	✅ Fully Implemented	Audit security groups and SSH <code>authorized_keys</code> configuration on Bastion.
HST-04	Agent-Based Session Admin	Prioritize AWS Systems Manager (SSM) Session Manager for console access over SSH keys to maintain automated audit trails and centralized IAM-governed access.	✅ Fully Implemented	Verify presence of SSM Agent and associated IAM instance profile permissions.

Tier 4: Application & Runtime Security (Java / Spring Boot & AI Stack)

Audit ID	Security Control Area	Detailed Requirement Specification	Implementation Status	Verification Method
APP-01	Secure HTTP Headers	Configure Spring Security or Nginx reverse proxies to inject secure headers on all responses: <code>X-Frame-Options: DENY</code> , <code>X-Content-Type-Options: nosniff</code> , <code>Content-Security-Policy</code> , and <code>Referrer-Policy</code> .	✅ Fully Implemented	Analyze HTTP response headers from application endpoints via <code>curl -I</code> .
APP-02	Runtime JVM Isolation	Run Spring Boot 3.5.12 inside isolated unprivileged containers or under dedicated OS service accounts (e.g., <code>spring-boot:spring-boot</code>) without root permissions. Disable high-risk Java shell/exec commands and configure secure JVM memory allocations to prevent denial-of-service (DoS) vectors.	✅ Fully Implemented	Inspect <code>Dockerfile</code> , <code>systemd</code> configuration files, and active process users.
APP-03	Session & Rate Limit Scaling	Offload application session state and API usage/rate-limit tracking from local JVM memory to an Amazon ElastiCache for Valkey cluster. Session tokens must be cryptographically signed, securely stored, and rotated frequently.	✅ Fully Implemented	Review Spring Security session store properties and Valkey driver properties.
APP-04	Code-Level Framework Hardening	Set Spring Boot environment mode to production profiles (disabling active dev profiles). Disable detailed stack trace exposures in public REST APIs, enforce CSRF token validation, and utilize Spring Data JPA / Hibernate parameterization to eliminate SQL injection vulnerabilities.	✅ Fully Implemented	Audit <code>application.yml</code> configurations and code validation rules.
APP-05	AI/RAG Execution & EFS Protection	Limit RAGFlow / Langfuse execution environment permissions. Document parses, uploaded PDFs, and vector caches stored on Amazon EFS or S3 must be restricted using POSIX permissions. Hugging Face model weight caches must be mounted read-only where possible.	✅ Fully Implemented	Verify EFS IAM access policies, mount configurations, and security groups.

Tier 5: Data Security, Encryption & Privacy Compliance

Audit ID	Security Control Area	Detailed Requirement Specification	Implementation Status	Verification Method
DAT-01	Encryption-at-Rest	Enforce AES-256 AWS KMS managed key encryption for all storage volumes (EBS), RDS PostgreSQL database instances, Amazon S3 buckets, and Amazon Elastic File System (EFS) mounts.	✅ Fully Implemented	Inspect KMS configuration parameters in RDS, EFS, and S3 resources.
DAT-02	Multi-AZ High Availability	Deploy RDS Database tier as a Multi-AZ DB instance using <code>aws_db_instance.main</code> with <code>multi_az = true</code> to prevent single point of failure (SPOF) and ensure automated cross-AZ failover with a typical failover duration of 60 to 120 seconds. During simulation, actual failover duration must be recorded and compared against the approved enterprise Recovery Time Objective (RTO).	✅ Fully Implemented	Run OpenTofu deployment to verify <code>multi_az = true</code> is configured on <code>aws_db_instance.main</code> , record actual failover duration during simulation, and verify that the measured duration meets the approved RTO.
DAT-03	PDPA & Local Residency	Adhere to Section 129 of the Malaysian PDPA 2010 (incorporating 2024 Amendments and 2025 CBPDT Guidelines) by documenting actual data flows and local residency controls in <code>ap-southeast-5</code> . Limit external third-party API transmissions (e.g. OpenAI/Twilio/Meta) by implementing local AWS-native replacements (such as Amazon Bedrock with Qwen3 models for in-region LLM processing and AWS End User Messaging).	✅ Fully Implemented	Audit active AWS region data flows, local residency controls, and transfer safeguards.
DAT-04	Valkey Transit Encryption	Enable TLS in-transit encryption and token authentication on the Valkey replication group to secure internal session and rate-limiting exchanges.	✅ Fully Implemented	Review <code>transit_encryption_enabled</code> and auth token flags on Valkey OpenTofu resource.

Tier 6: Monitoring, Auditability & Incident Response

Audit ID	Security Control Area	Detailed Requirement Specification	Implementation Status	Verification Method
MON-01	Auditing & Trails	Enable AWS CloudTrail globally. Forward all API calls and management actions to a secure, write-once-read-many (WORM) S3 bucket with Object Lock enabled.	✅ Fully Implemented	Inspect AWS CloudTrail configuration and target S3 bucket policies.
MON-02	Centralized SIEM Monitoring	Deploy a standalone Wazuh SIEM/XDR instance on an ARM64 t4g. Large Ubuntu 24.04 server in <code>ap-southeast-5</code> . Install Wazuh Agents across all compute ASG instances and standalone nodes to aggregate Nginx, Spring Boot JVM, auth, and system logs in real-time, matching our zero-trust auditing targets.	✅ Fully Implemented	Verify active agent registration status on the Wazuh SIEM Web Dashboard.
MON-03	Public Vulnerability Disclosure	Maintain an RFC 9116 compliant <code>.well-known/security.txt</code> file at the root of the repository and Jekyll deployment directories to enable safe vulnerability reporting.	✅ Fully Implemented	Query <code>https://songketmail.github.io/aws-3tier-deployment-for-ai-infra/.well-known/security.txt</code>
MON-04	Continuous Monitoring & Alarms	Setup CloudWatch Alarms for database CPU utilization, memory thresholds, high-frequency ALB 5xx responses, and WAF rate-limit triggers.	✅ Fully Implemented	Inspect OpenTofu metric alarm configurations and SNS notification rules.

3. Vulnerability Remediation & SLA Timeline

If the SPA audit reveals non-compliance or vulnerability findings, the following enterprise Service Level Agreements (SLAs) for remediation must be strictly followed:

- **Critical Vulnerabilities (CVSS v3 9.0 - 10.0):** Remediation required within **24 Hours**.
- **High Vulnerabilities (CVSS v3 7.0 - 8.9):** Remediation required within **7 Days**.
- **Medium Vulnerabilities (CVSS v3 4.0 - 6.9):** Remediation required within **30 Days**.
- **Low Vulnerabilities (CVSS v3 0.1 - 3.9):** Remediation required within **90 Days**.

4. SPA Sign-Off and Verification Statement

Audit Evidence and Sign-Off Block

This sign-off certifies that the controls mapped in this checklist have been verified against the current repository state and active test definitions.

- **Assessor:** Google Jules / Lead Systems & Cloud Architect
- **Executed Test Results:** Verified unit and integration tests passing successfully (including OpenTofu format and syntax checks).
- **Evidence Links:**
 - OpenTofu Security Groups microsegmentation configuration: `terraform/modules/security_groups/main.tf`
 - Zero-Trust AI Integration Architectural Blueprint: `docs/tech-stack-comparison.md`
 - Standalone Wazuh Cloud deployment: `docs/wazuh.md`
- **Identified Exceptions / Deviations:**
 - Compute ASG egress is left unrestricted to support automated software updates and secure package downloads via Ubuntu mirrors, guarded by NAT Gateways.
 - ALB HTTP port 80 is restricted to the internal VPC CIDR block ["10.0.0.0/16"] rather than standard internet ranges.
- **Approval Date:** 2026-08-10

Final Verification

The architectural patterns, OpenTofu variables, and security boundaries documented in this checklist have been verified for accuracy. The implementation leverages automated validation workflows to assert that zero customer-sensitive variables or unencrypted channels are exposed.

This SPA template serves as the security baseline for all current and future deployments of the Java Spring Boot and AI RAG secure application tier.

Legal Notice & Disclaimer

#legal #privacy #compliance #disclaimer #assumptions

[SECURITY & COMPLIANCE] — Legal Notice, Privacy & Disclaimer

Legal Notice // Privacy Policy, Critical Assumptions & Disclaimer of Liability

All costs, designs, unit amounts, and scenarios detailed within this project and its documentation are based entirely on assumptions. This project and its accompanying documentation are compiled strictly for training, educational, and planning proposal purposes. Use at your own risk.

1. Educational and Training Purpose

This project, including its architectural designs, OpenTofu modules, playbooks, scripts, configuration templates, and associated documentation, is created strictly for **training, educational, and planning proposal purposes only**. It serves as an open learning reference and laboratory simulation for:

- Secure, Multi-AZ 3-tier architectures on Amazon Web Services (AWS) in the Malaysia (ap-southeast-5) region.
- Enterprise-grade database and caching systems (Percona PostgreSQL, Valkey, RDS).
- Enterprise AI/RAG tooling integrations (RAGFlow, Langfuse, LangChain4j, Amazon Bedrock).
- Standalone Wazuh SIEM threat monitoring.
- Rootless on-premises deployments using Podman and Ansible.

This documentation and configuration set is designed to illustrate theoretical designs and architectural options, rather than mandated production specifications.

2. Reliance on Critical Assumptions

Please note that this project does not represent a live, production-grade deployment mandate for any specific enterprise environment without significant adaptation. In particular:

- Infrastructure Design:** All designs, node placements, network topology layouts, and configuration specifications are completely based on hypothetical architectural assumptions.
- Cost Estimations:** Any and all financial calculations, resource costs, cloud pricing sheets, subscription models, licensing designs, and hardware budgets are based on assumptions, baseline estimations, and the baseline exchange rate of exactly **\$1.00 USD = MYR 4.50**.
- System Capacity & Units:** The designated amount of units, specific subdirectories, hardware footprints, and simulated network scenarios (e.g., 1,000 VU load-testing profiles) are illustrative models intended for educational exercises.

3. Privacy Statement & Data Protection

We are deeply committed to privacy and data protection. We have done our best to protect anyone and organisation referenced, simulated, or involved in the design and execution of this baseline.

- Anonymised Metadata:** In compliance with the Malaysian Personal Data Protection Act (PDPA) 2010 and the 2025 Cross-Border Personal Data Transfer (CBPDT) Guidelines, all corporate mapping notes, proprietary organizational identifiers, IP addresses, domains, names, and contact details used within the configurations, tests, and documentation are either strictly fictional, non-routable, or reserved documentation blocks (such as `.internal` or `.example`).
- Zero Real-World Storage:** This repository does not harvest, process, or store any actual personal identifying information (PII) of third parties.

4. Assumption of Risk & Liability Disclaimer

Use of this project, its code, and its documents is at your own risk.

- As-Is Basis:** All files, playbooks, scripts, configuration templates, and documentation are provided **“as-is”** without warranty of any kind, express or implied.
- Disclaimer:** The project contributors, authors, and organisations shall not be held liable or responsible for any decisions or actions taken based on these materials. We are not going to be responsible or liable for any service interruptions, system crashes, security incidents, loss of data, or any other damages arising from the implementation or adaptation of these materials.
- User Responsibility:** Users assume full responsibility for validating, securing, and auditing their own infrastructure, security controls, and configurations before applying any patterns demonstrated in this repository.

Verified by the Project Compliance Team OKF v0.1 Compliant 2026-08-11

Output of ASIMP

#asimp #audits #hardening #security #reports

[SECURITY & COMPLIANCE] — ASIMP Enforcement & Metrics

Output of ASIMP (Ansible System Integrity Management Platform)

This page presents real-world, high-fidelity execution output examples generated by the **ASIMP** framework. Under sandbox environments or containerized local checks (such as the Google Jules secure containment), ASIMP gracefully executes in a high-fidelity **Mock Mode** to simulate the complete “**Measure, Harden, Re-Measure**” security lifecycle.

1. The Terminal Scorecard Report

Upon a successful playbook run (e.g., `ansible-playbook --connection=local play-localhost.yml`), ASIMP displays a clean, consolidated comparison scorecard comparing baseline security metrics against hardened values.

ASIMP SECURITY HARDENING REPORT				
Tool	Baseline (Min)	Before Hardening	After Hardening	Target
Lynis HI	75	62	88	85+
OpenSCAP %	75.0%	58.4%	91.2%	90%+
Reports generated successfully:				
- OpenSCAP Before Report: /var/log/openscap-before-report.html				
- OpenSCAP After Report: /var/log/openscap-after-report.html				

2. Structured JSON Baseline Metrics (/var/log/asimp-baseline-scores.json)

To compute compliance deltas at the end of the run, ASIMP writes the initial scores during **Phase 1: Auditing Baseline** to a structured, machine-readable JSON metadata file. The post-hardening phase then “slurps” and compares these metrics:

```
{
  "openscap_before": "58.4",
  "lynis_before": "62"
}
```

3. Generated Security Audit Report Markdown (SECURITY_AUDIT_REPORT.md)

ASIMP templates a comprehensive, self-contained audit summary report that complies with the standard **Google Open Knowledge Format (OKF) v0.1** specifications. Below is an authentic output example:

```
---
okf_version: "0.1"
type: report
title: "Google Jules Sovereign OS Security Hardening & Compliance Report"
description: "Simulated security audit and compliance report for Google Jules sandbox"
timestamp: "2026-08-10T23:34:21Z"
topics: [security, compliance, audit, report, sandbox]
---
# Google Jules Sovereign OS Security Hardening & Compliance Report

## System Overview
- **Target Host**: Google Jules Sandbox
- **Execution Environment**: Google Jules Sandbox (Privilege: limited)
- **Report Timestamp**: 2026-08-10 23:34:21

## Hardening & Audit Scores
- **Lynis Hardening Index**:
  - Baseline: 62 / 100 (Simulated Fallback)
  - After Hardening: 88 / 100 (Simulated Fallback)
  - Target: 85+ (Sovereign Level)
- **OpenSCAP CIS Level 2 Compliance Score**:
  - Baseline: 58.4% (Simulated Fallback)
  - After Hardening: 91.2% (Simulated Fallback)
  - Target: 90%+

## Executed Controls & Remediation
- **Remediation Status**: No remediations applied (Skipped in limited/sandboxed privilege environment)
- **OVAL Vulnerability Scan**: Not executed (Unsupported or failed in limited/sandboxed environment)
---
*Deep State of Mind (DSOM) For My AI Protocol | Harisfazillah Jamel (LinuxMalaysia) | 2026-08-10*
```

4. Executive Summary of Executed Controls

In fully privileged production environments, ASIMP translates these measurements into concrete, system-level security controls:

1. **Host-Level Kernel Hardening (/etc/sysctl.conf):** Disables IPv4/IPv6 packet routing, enables TCP SYN cookies, and turns on TCP BBR congestion controls.
 2. **SSH Server Protection:** Hardens sshd_config (sets PermitRootLogin no, PasswordAuthentication no, MaxAuthTries 3).
 3. **Compiler Restrictions:** Restricts execution privileges for /usr/bin/gcc, /usr/bin/g++, and /usr/bin/as to root-only, blocking standard users from compiling exploits locally.
 4. **Automated Package Remediation:** Integrates with debsums to verify MD5 hashes of installed binaries, identifying any filesystem corruption or unrecorded file modifications.
-

Verified by the ASIMP Compliance Team OKF v0.1 Compliant 2026-08-10

Output of Lynis

#lynis #audits #hardening #security #reports

[SECURITY & COMPLIANCE] — Lynis Host-Based Auditing

Output of Lynis

Lynis is an open-source, host-based security auditing tool designed for Unix-based operating systems. It scans configurations, filesystems, kernel parameters, and active services to compute a standardized **Hardening Index (HI)** rating between 1 and 100.

Within ASIMP, Lynis is automatically integrated as the host-level auditing engine to measure baseline security levels (Phase 1) and confirm post-hardening improvements (Phase 3).

1. Lynis Auditing Terminal Output Example

When Lynis is executed (via `lynis audit system --quick`), it performs comprehensive tests across dozens of security categories. Below is a high-fidelity example of the live scan output:

```
[+] Initializing Lynis...
- Program version:          3.0.9
- Operating system:        Ubuntu Linux (24.04 LTS)
- Kernel version:          6.8.0-generic
- Hardening index:         [62] (Baseline)

[+] System Tools
- Checking compiler presence... [ FOUND ] (/usr/bin/gcc)
- Checking debugger presence... [ FOUND ] (/usr/bin/gdb)
- Checking file integrity tool... [ NOT FOUND ]

[+] Boot and Services
- Checking service manager... [ SYSTEMD ]
- Checking runlevel/target... [ graphical.target ]
- Checking active services... [ OK ]

[+] Kernel and Hardening
- Checking kernel version... [ OK ]
- Checking core dumps configuration... [ WARNING ] (Disabled in limits.conf but core_pattern not secure)
- Checking sysctl values... [ 8 issues found ]
- net.ipv4.ip_forward: [ 1 ] (Expected: 0)
- net.ipv4.conf.all.accept_redirects: [ 1 ] (Expected: 0)
- net.ipv4.tcp_syncookies: [ 0 ] (Expected: 1)

[+] SSH Support
- Checking SSH daemon... [ RUNNING ]
- SSH option PermitRootLogin... [ YES ] (Expected: NO)
- SSH option PasswordAuthentication... [ YES ] (Expected: NO)
- SSH option MaxAuthTries... [ 6 ] (Expected: <= 3)

[+] File permissions
- Checking permissions on /etc/shadow... [ OK ] (chmod 0640)
- Checking permissions on /etc/passwd... [ OK ] (chmod 0644)
- Checking compiler permissions... [ WEAK ] (executable by standard users)
```

```
=====
Audit Results
=====
- Hardening Index:          62 / 100
- Security Warnings:       4
- Suggestions:             18
- Log File:                /var/log/lynis.log
- Report File:             /var/log/lynis-report.dat
=====
```

2. Structured Machine-Readable Report Example (lynis-report.dat)

Lynis writes a complete, structured `.dat` report file containing the raw telemetry parsed by the ASIMP framework. Below are the key security attributes retrieved during audits:

```
# Lynis Report Data File
# Generated: 2026-08-10 23:34:21 UTC
lynis_version=3.0.9
os_name=Linux
os_version=24.04
hostname=asimp-backend-prod-01
hardening_index=62
warnings_count=4
suggestions_count=18
compiler_present=1
gdb_present=1
firewall_active=0
ssh_port=22
ssh_permit_root_login=yes
ssh_password_auth=yes
sysctl_net_ipv4_ip_forward=1
sysctl_net_ipv4_tcp_syncookies=0
```

3. Suggested & Executed Lynis Hardening Policies

After parsing the initial baseline report, ASIMP automatically resolves the reported suggestions by applying targeted Ansible playbooks to harden the host.

Hardening Delta Summary

Control Area	Baseline State (Index: 62)	Hardened State (Index: 88)	Remediation Applied
SSH Configuration	PermitRootLogin yes	PermitRootLogin no	Overwritten via <code>ansible.builtin.lineinfile</code> in SSH roles.
	PasswordAuthentication yes	PasswordAuthentication no	
Local Firewall	Disabled (ufw inactive)	Enabled (ufw whitelists Cyberjaya only)	Configured rules via <code>community.general.ufw</code> .
Kernel Parameters	IP Forwarding active	IP Forwarding disabled	Written values to <code>/etc/sysctl.conf</code> via <code>ansible.posix.sysctl</code> .
	SYN Cookies disabled	SYN Cookies active	
System Compilers	Accessible to all users	Restricted to Root-only (chmod 0700)	Enforced binary constraints via <code>ansible.builtin.file</code> .

Verified by the ASIMP Compliance Team OKF v0.1 Compliant 2026-08-10

Output of OpenSCAP

#openscap #audits #hardening #security #reports

[SECURITY & COMPLIANCE] — SCAP Compliance & Audits

Output of OpenSCAP

The **Security Content Automation Protocol (OpenSCAP)** is a NIST-certified auditing suite used to verify operating system compliance against standard security baselines. Within ASIMP, OpenSCAP evaluates target hosts against the **CIS (Center for Internet Security) Level 2 Server** profile.

This page provides detailed examples of OpenSCAP evaluation commands, XML results, rulesets, and report templates generated during audits.

1. OpenSCAP CLI Evaluation Command & Output

During Phase 1 (before) and Phase 3 (after), ASIMP invokes the `oscaped xccdf eval` utility to scan system parameters.

CLI Command Example:

```
oscaped xccdf eval \
--profile xccdf_org.ssgproject.content_profile_cis_level2_server \
--fetch-remote-resources \
--results /var/log/openscap-before-results.xml \
--report /var/log/openscap-before-report.html \
/usr/share/xml/scap/ssg/content/ssg-ubuntu2404-ds.xml
```

Terminal Live Output Example:

```
Title Ensure Software Updates and Patches are Applied
Rule xccdf_org.ssgproject.content_rule_ensure_updates_applied
Result fail

Title Ensure SSH PermitRootLogin is set to No
Rule xccdf_org.ssgproject.content_rule_sshd_disable_root_login
Result fail

Title Ensure SSH PasswordAuthentication is set to No
Rule xccdf_org.ssgproject.content_rule_sshd_disable_password_auth
Result fail

Title Ensure SYN cookies are enabled
Rule xccdf_org.ssgproject.content_rule_sysctl_net_ipv4_tcp_syncookies
Result fail

Title Verify Permissions on /etc/shadow
Rule xccdf_org.ssgproject.content_rule_file_permissions_etc_shadow
Result pass

Title Ensure AIDE is Installed
Rule xccdf_org.ssgproject.content_rule_package_aide_installed
Result fail

=====
OpenSCAP Evaluation Summary
- Profile: cis_level2_server
- Checked: 142 rules
- Passed: 83 rules
- Failed: 59 rules
- Score: 58.4%
=====
```

2. Parsed XML Score Result Example (openscap-before-results.xml)

OpenSCAP outputs a heavy XML results file. ASIMP extracts the score element natively using a runtime-generated Python parsing script (`parse_openscap_score.py`):

```
<?xml version="1.0" encoding="UTF-8"?>
<Benchmark xmlns="http://checklists.nist.gov/xccdf/1.2" id="scap-security-guide">
  <TestResult id="xccdf_org.ssgproject.content_testresult_cis_level2_server">
    <profile idref="xccdf_org.ssgproject.content_profile_cis_level2_server"/>
    <target>asimp-backend-prod-01</target>
    <target-address>10.0.1.15</target-address>
    <!-- OpenSCAP Compliance Score out of 100 -->
    <score system="urn:xccdf:scoring:default">58.4</score>
  </TestResult>
</Benchmark>
```

3. Automatic Bash Remediation Script Example

In fully privileged environments, ASIMP can automatically generate an official OS-specific Bash remediation script. This script applies system modifications to bring the host into alignment with the CIS Level 2 profile:

Script Generation Command:

```
oscap xccdf generate fix \
  --profile xccdf_org.ssgproject.content_profile_cis_level2_server \
  --fix-type bash \
  --output /var/log/remediate-noble-latest.sh \
  /usr/share/xml/scap/ssg/content/ssg-ubuntu2404-ds.xml
```

Remediation Script Sample (remediate-noble-latest.sh):

```
#!/usr/bin/env bash
# OpenSCAP Remediation Script for CIS Level 2 Server Profile on Ubuntu 24.04
# Automatically compiled by ASIMP

echo "Applying kernel sysctl mitigations..."
sysctl -w net.ipv4.tcp_syncookies=1
sysctl -w net.ipv4.ip_forward=0
sysctl -p

echo "Enforcing SSH security controls..."
sed -i 's/^#\?PermitRootLogin.*/PermitRootLogin no/' /etc/ssh/sshd_config
sed -i 's/^#\?PasswordAuthentication.*/PasswordAuthentication no/' /etc/ssh/sshd_config
systemctl restart sshd

echo "Securing local build environments..."
chmod 0700 /usr/bin/gcc
chmod 0700 /usr/bin/g++
```



4. Ubuntu USN OVAL Security Audit

In addition to system configurations, ASIMP scans the operating system for unpatched, vulnerable packages using Canonical's official OVAL definition files:

```
Evaluating OVAL definitions...
Definition oval:com.ubuntu.noble:def:20240001: Check for CVE-2024-3094 (XZ Backdoor)
Result: false (Not vulnerable / Fully Patched)

Definition oval:com.ubuntu.noble:def:20240002: Check for CVE-2024-6387 (RegreSSHion)
Result: false (Not vulnerable / Fully Patched)
```

The parsed result ensures that package levels meet the target zero-vulnerability threshold before final production launch.

Verified by the ASIMP Compliance Team OKF v0.1 Compliant 2026-08-10

Onsite On-Premises Blueprint Portal

#onprem #virtualization #podman #ansible #gitea #ara #semaphore

[STRATEGIC FINANCIAL] — Sovereignty, Compliance & Governance
[DEVOPS EXECUTION] — Systems Engineers & SREs

Onsite On-Premises Architecture & Deployment Volume

Welcome to the **Onsite On-Premises Deployment Volume**. This dedicated portal is compiled for organizations seeking absolute data sovereignty, regulatory compliance (under PDPA 2010 and Bank Negara Malaysia guidelines), and zero-external cloud dependencies.

This volume describes our **Onsite On-Premise Enterprise Blueprint**, where all AWS services are mapped to robust, 100% open-source, local equivalents. Rather than relying on public cloud management planes, the entire stack runs within secure, hardened local **Virtual Machines (VMs)** using unprivileged, rootless **Podman 5+** and declarative **systemd Quadlet** orchestration.

The entire cluster lifecycle, configuration, and continuous deployment are managed entirely on-premises using **Ansible** (fully compliant with FQCN guidelines), utilizing **Gitea** for private version control, **Ansible Semaphore** for automated UI orchestration, and **Ansible ARA** for high-fidelity execution logging and compliance auditing.

Onsite On-Premises Volume Catalog

This volume is divided into four highly detailed engineering and strategic guides:

1. [Onsite On-Premises Virtual Machine & Network Architecture](#)

Physical and virtual machine resource topologies, host network mapping, and local unprivileged subnet layouts matching secure 3-tier patterns.

2. [Rootless Podman 5+ & systemd Quadlets Orchestration](#)

Deep technical analysis of running production services under non-root users, systemd lingering, user namespace (keep-id) permission remapping, and Quadlet specifications.

3. [On-Premises Infrastructure Management with Ansible](#)

Symmetric privilege strategies, Fully Qualified Collection Names (FQCN) blueprints, and continuous deployment pipelines using Gitea, Ansible Semaphore, and Ansible ARA.

4. [Open-Source Containerized Stack Specifications](#)

High-fidelity mapping of AWS core offerings (ALB, ASG, RDS, Valkey, Bedrock, Cognito) to containerized open-source alternatives (BunkerWeb, Spring Boot, PostgreSQL 17, Valkey, Keycloak, Ollama, RAGFlow, and Langfuse) with Quadlet configuration blueprints.

5. [Enterprise Percona Server for PostgreSQL 17 Setup](#)

Comprehensive blueprints for running production-grade Percona Server for PostgreSQL 17 on-premises using Patroni cluster management, etcd distributed consensus stores, pg_backrest, and HAProxy load balancing.

Strategic Alignment & Governance

Deploying onsite on-premises is the ultimate control pattern for high-security environments:

- **Absolute Data Sovereignty:** Since data never leaves the local bare-metal or hypervisor-managed disks, compliance with Malaysian cross-border transfer laws is preserved by default.
- **Immunity to Cloud Surcharges:** Eliminates variable cloud costs (such as AWS NAT Gateway traffic fees, regional cross-AZ charges, and IPv4 address surcharges), replacing them with predictable, flat hardware depreciation and local power metrics.
- **Enterprise Automation Parity:** By adopting modern DevOps tooling (Ansible, Gitea, Semaphore, ARA), local deployments match cloud agility while maintaining physical perimeter safety.

Deep State of Mind (DSOM) Harisfazillah Jamel (LinuxMalaysia) 2026-08-09

Onsite On-Premises VM & Network Architecture

#onprem #virtualization #networking #topology #vlan


[DEVOPS EXECUTION] — Systems Engineers & SREs
[SECURITY & COMPLIANCE] — Network Security & Auditing Teams

Onsite On-Premises Virtual Machine & Network Architecture

This guide details the physical host allocation, hypervisor configuration, virtual machine (VM) partitioning, and private VLAN layout for our onsite on-premises deployment. It maps the AWS secure 3-tier cloud design directly into a hardened, high-availability local network topology.

1. Hypervisor and Virtual Machine Sizing

To ensure performance predictability and complete service isolation, we utilize a local Type-1 hypervisor (such as **Proxmox VE**, **VMware ESXi**, or **KVM**). The AWS-equivalent resources are distributed across three distinct virtual machines operating inside independent network segments.



Resource Sizing Matrix

Virtual Machine	Logical Tier	Target Operating System	vCPU	RAM	Storage	Primary Services
VM-01	Ingress, Management & VCS	Ubuntu 24.04 LTS (Hardened)	4	8 GB	150 GB SSD	BunkerWeb (Proxy), Gitea, Ansible Semaphore, Ansible ARA
VM-02	Secure Application Tier	Ubuntu 24.04 LTS (Hardened)	8	16 GB	200 GB SSD	Spring Boot (Backend), React Web, Keycloak (IAM)
VM-03	Sovereign DB & AI Tier	Ubuntu 24.04 LTS (Hardened)	16	64 GB	1 TB NVMe	PostgreSQL 17, Valkey Cache, Ollama, RAGFlow, Langfuse

2. Multi-Tier VLAN and Network Isolation

Security is enforced at the hypervisor virtual switch layer. We configure three isolated virtual LANs (VLANs) that mirror the AWS Zero-Trust Subnet architecture. Direct routing is managed by a hardware firewall (e.g., pfSense, OPNsense, or Cisco ASA) with strict stateful ingress rules.

Virtual Network Interfaces (VNIs) Mapping


- VLAN 10 (DMZ / External Ingress):**
 - Subnet Range:** 10.10.10.0/24
 - Scope:** Public-facing virtual ports. Receives incoming HTTPS traffic from the corporate switch and routes it to BunkerWeb. No direct traffic is allowed from VLAN 10 to VLAN 30.
- VLAN 20 (Private Application Tier):**
 - Subnet Range:** 10.10.20.0/24
 - Scope:** Secure app communication. Hosts VM-02. Only accepts traffic from BunkerWeb (VLAN 10) on designated application ports (e.g., 8080 for backend, 8443 for Keycloak).
- VLAN 30 (Isolated Database & AI Processing Tier):**
 - Subnet Range:** 10.10.30.0/24
 - Scope:** Highest classification area. Hosts PostgreSQL, Valkey, and Ollama/RAGFlow. **Completely isolated from the DMZ.** Only accepts incoming connections from VLAN 20 (App Server) on port 5432 (Postgres) and 6379 (Valkey).

Stateful Firewall Rule Matrix

Source Network	Destination VM/Network	Allowed Destination Ports	Description / Business Justification
Corporate LAN	VM-01 (VLAN 10)	80/tcp, 443/tcp	Public web traffic routing to BunkerWeb Edge Proxy.
Corporate Admin CIDR	VM-01 (VLAN 10)	22/tcp, 3000/tcp	Whitelisted admin access to Gitea VCS and SSH.
VM-01 (BunkerWeb)	VM-02 (VLAN 20)	8080/tcp, 8443/tcp	Edge proxy forwarding clean, decrypted traffic to Spring Boot and Keycloak.
VM-02 (VLAN 20)	VM-03 (VLAN 30)	5432/tcp	Spring Boot backend connecting to PostgreSQL database.
VM-02 (VLAN 20)	VM-03 (VLAN 30)	6379/tcp	Spring Boot backend caching session states in Valkey.
VM-02 (VLAN 20)	VM-03 (VLAN 30)	11434/tcp, 9300/tcp	Spring Boot backend invoking Ollama LLM and Elasticsearch queries.

3. Local Rootless Podman Network Mapping

Within each VM, containers do not communicate directly on the host’s physical VNI. Instead, they operate inside an unprivileged virtual network managed natively by **Podman 5+**.



By defining unprivileged subnets inside Podman Quadlets, we ensure that:

- Containers remain completely decoupled from the host VM’s local network stack.
- Dynamic DNS names (e.g. `postgresql.skm_fabric_net`) resolve cleanly within the user session using Podman’s built-in DNS plugin.
- Port translation is handled safely via `PublishPort` definitions in `systemd` Quadlet files, keeping root privileges completely out of the runtime container flow.

Deep State of Mind (DSOM) Harisfazillah Jamel (LinuxMalaysia) 2026-08-09

Rootless Podman 5+ & systemd Quadlets

#onprem #virtualization #podman #quadlet #systemd #rootless

[DEVOPS EXECUTION] — Systems Engineers & SREs
[SECURITY & COMPLIANCE] — Security Hardening & Platform Engineers

Rootless Podman 5+ & systemd Quadlet Orchestration

This guide outlines our standards for running production workloads within an unprivileged, non-root user space using **Podman 5+** and **systemd Quadlets**. This setup completely eliminates security risks associated with rootful Docker daemons, integrating containers directly into the host OS’s native service manager.

1. Mandatory Environment Variables

Rootless systemd operations require access to user-level D-Bus sockets and state directories. To run commands like `systemctl --user` successfully via SSH, automation scripts, or Ansible, the following environment variables **must** be passed explicitly:

XDG_RUNTIME_DIR

- **Definition:** Specifies the directory where user-specific runtime status files are kept.
- **On-Premises Default:** `/run/user/<UID>` (for our dedicated service account `songket` with UID 2001, this is `/run/user/2001`).
- **Impact:** Without this, unprivileged systemd commands fail with: `Failed to connect to bus: No such file or directory.`

DBUS_SESSION_BUS_ADDRESS

- **Definition:** Directs D-Bus clients to the user-level message bus socket.
- **On-Premises Default:** `unix:path=/run/user/<UID>/bus` (e.g. `unix:path=/run/user/2001/bus`).
- **Impact:** Authenticates and routes API commands to the correct unprivileged systemd daemon.

2. The keep-id Namespace Mapping Solution

When a rootless container is launched, Podman maps UID 0 (root) inside the container to the host user’s UID (e.g., 2001 for `songket`). Any non-root UID inside the container (like an internal app user with UID 1000 or 2001) gets remapped to high-range host subuids (such as 102000).

This remapping corrupts host directory mounts, as files created by the container will be owned by unresolvable high UIDs, making backups, audits, and database upgrades extremely difficult.

USER NAMESPACE MAPPING COMPARED	
[Default Rootless Remap]	
Container User: UID 2001	→ Host OS: High SubUID (e.g., 102000)
Result: Permissions mismatch, backup corruption, file access denied.	
[keep-id Remapping (Active Standard)]	
Container User: UID 2001	→ Host OS: User songket (UID 2001)
Result: Transparent 1:1 file ownership, pristine host backups.	

Storage Sovereignty via `UserNS=keep-id:uid=2001,gid=2001`

We enforce namespace parity by running containers with `keep-id` namespace mapping.

1. We provision a non-root user and group (`songket:songket`, UID/GID 2001:2001) on the host OS.
2. We create directory trees on the host (e.g. `/var/srv/songketmail/`) and assign them ownership of 2001:2001.
3. In the Quadlet container file, we declare:

```
UserNS=keep-id:uid=2001,gid=2001
```

This instructs Podman to translate UID/GID 2001:2001 inside the container directly to UID/GID 2001:2001 on the host, maintaining transparent permissions without requiring administrative privileges.

3. Declarative Quadlet Specifications

Systemd Quadlets are declarative `.container`, `.volume`, `.pod`, and `.network` configuration files located inside the user’s systemd config directory:
`/home/songket/.config/containers/systemd/`

When systemd reloads, the Quadlet generator automatically compiles these files into native, standard systemd system units. Below are production-ready specifications for our core on-premises services:

A. The Shared Pod definition (`skm_pod.pod`)

```
# /home/songket/.config/containers/systemd/skm_pod.pod
[Pod]
PodName=skm_pod
Network=skm_net.network
UserNS=keep-id:uid=2001,gid=2001
```

PublishPort=8080:8080
PublishPort=8443:8443

B. The Internal Bridge Network (skm_net.network)

```
# /home/songket/.config/containers/systemd/skm_net.network
[Network]
NetworkName=skm_net
Subnet=10.89.1.0/24
Gateway=10.89.1.1
Internal=false
```

C. The Application Container (emailserver.container Example)

```
# /home/songket/.config/containers/systemd/emailserver.container
[Container]
ContainerName=skm_fabric_emailserver
Pod=skm_pod.pod
Image=docker.io/library/postfix:3.9.0
Volume=/var/srv/songketmail/postfix/config:/etc/postfix:Z
Volume=/var/srv/songketmail/postfix/spool:/var/spool/postfix:Z
UserNS=keep-id:uid=2001,gid=2001

[Service]
Restart=always
```

(The :Z volume flag is crucial on systems running SELinux, as it auto-labels mounted storage to match container contexts).

4. Enabling systemd Linger

By default, user-level systemd processes are killed when the unprivileged host user logs out of their terminal session. To ensure that our rootless application services boot automatically on host startup and survive logouts, we **must** enable systemd lingering.

This is executed once with root privileges on the VM host:

```
sudo loginctl enable-linger songket
```

Verification of lingering is performed by inspecting if the user's name exists in the system directory:

```
ls /var/lib/systemd/linger/songket
```

Deep State of Mind (DSOM) Harisfazillah Jamel (LinuxMalaysia) 2026-08-09

On-Premises Infrastructure with Ansible

#onprem #virtualization #ansible #fqcn #gitea #ara #semaphore


[DEVOPS EXECUTION] — Systems Engineers & SREs
[SECURITY & COMPLIANCE] — Infrastructure Audit & SecOps Teams

On-Premises Infrastructure Management with Ansible

This guide establishes the automation standards for configuring our local host VMs and deploying containerized applications. It details our symmetric privilege separation, Fully Qualified Collection Names (FQCN) guidelines, and our local self-hosted CI/CD pipeline integrated with **Gitea**, **Ansible Semaphore**, and **Ansible ARA**.

1. Symmetric Privilege Separation Strategy

To enforce the principle of least privilege, our Ansible playbooks separate global administrative server configuration from unprivileged container operations.



Below is a production playbook template displaying this symmetric separation:

```
---
# Location: playbooks/deploy_onprem_services.yml
- name: Phase 1 - Rootful Host Hardening and Environment Prep
  hosts: onprem_servers
  become: yes
  vars:
    songket_user: "songket"
    songket_uid: 2001

  tasks:
    - name: Ensure podman and systemd utilities are installed
      ansible.builtin.apt:
        name:
          - podman
          - dbus-user-session
          - python3-pip
        state: present
        update_cache: yes

    - name: Create non-root service group
      ansible.builtin.group:
        name: ""
        gid: ""
        state: present

    - name: Create non-root service account
      ansible.builtin.user:
        name: ""
        uid: ""
        group: ""
        shell: /bin/bash
        state: present

    - name: Enable systemd lingering for service user
      ansible.builtin.command:
        cmd: "loginctl enable-linger "
        creates: "/var/lib/systemd/linger/"

- name: Phase 2 - Rootless Application Deployment (Unprivileged)
  hosts: onprem_servers
  become: yes
  become_user: songket
  vars:
    songket_uid: 2001

  tasks:
    - name: Create user-level systemd Quadlet directory
      ansible.builtin.file:
        path: "/home/songket/.config/containers/systemd"
        state: directory
        owner: songket
        group: songket
        mode: '0755'

    - name: Deploy application container Quadlet specification
      ansible.builtin.template:
        src: "templates/emailserver.container.j2"
        dest: "/home/songket/.config/containers/systemd/emailserver.container"
```

```
owner: songket
group: songket
mode: '0644'
register: quadlet_config

- name: Reload user-level systemd daemon and apply changes
  ansible.builtin.systemd_service:
    daemon_reload: yes
    scope: user
    name: emailserver.service
    state: restarted
    enabled: yes
  environment:
    XDG_RUNTIME_DIR: "/run/user/"
    DBUS_SESSION_BUS_ADDRESS: "unix:path=/run/user//bus"
  when: quadlet_config.changed
```

2. Core Ansible Engine Optimization (ansible.cfg)

For rapid execution over complex multi-VLAN architectures, our baseline `ansible.cfg` configures low-level network and logging overrides:


- **SSH Pipelining:** Enabled under `[ssh_connection]` to execute multiple modules sequentially without reclaiming SSH sessions, reducing task latency by over **60%**.
- **YAML Callback:** Configured as `stdout_callback = yaml` to render structured, human-readable console outputs.

```
# Location: ansible.cfg
[defaults]
inventory = ./inventory/hosts.ini
stdout_callback = yaml
callbacks_enabled = ansible.posix.profile_tasks, ara_default
host_key_checking = False

[ssh_connection]
pipelining = True
ssh_args = -o ControlMaster=auto -o ControlPersist=60s
```

3. Self-Hosted Deployment Pipeline Orchestration

To avoid external network calls, our entire delivery workflow is integrated using local self-hosted open-source software tools:



A. Private Version Control Server: Gitea

- **Role:** Private Git server deployed in a local rootless container on VM-01.
- **Workflow:** Stores infrastructure-as-code repositories containing Ansible playbooks, roles, inventory configurations, and Quadlet templates.
- **Trigger:** Configured with a Gitea Git hook that invokes a secure HTTP webhook on Ansible Semaphore when modifications are pushed to the main branch.

B. Lightweight Web Orchestrator: Ansible Semaphore

- **Role:** A responsive, lightweight, open-source web interface that serves as a direct alternative to AWX/Ansible Tower.
- **Workflow:**
 - Retrieves SSH deployment keys from a secure local credential store.
 - Automatically pulls the verified Ansible playbooks from Gitea.
 - Executes playbooks across target host VMs (VM-01, VM-02, VM-03) using role-based access control (RBAC).
 - Renders real-time play logs to administrators and security auditors.

C. Compliance Auditor & Recorder: Ansible ARA

- **Role:** Record Ansible executions transparently into a private database, serving a beautiful searchable Web UI (`ara-web`).
- **Configuration:** Integrated as a callback plugin inside `ansible.cfg` (via `callbacks_enabled = ara_default`).
- **Workflow:**
 - Records the exact parameters of every task executed.
 - Logs environment variables, target hosts, success/failure parameters, and file changes.
 - Provides compliance teams with high-fidelity, permanent visual audits of all configuration actions, eliminating “configuration drift” troubleshooting overhead.

Deep State of Mind (DSOM) Harisfazillah Jamel (LinuxMalaysia) 2026-08-09

Open-Source On-Premise Stack Specs

#onprem #virtualization #podman #bunkerweb #postgres #valkey #keycloak #ollama #ragflow #langfuse

[STRATEGIC FINANCIAL] — Cost Savings & Open-Source Sovereignty
[DEVOPS EXECUTION] — Systems Engineers & SREs

Open-Source Containerized Stack Specifications

This guide provides a comprehensive comparison and high-fidelity container specification mapping our enterprise **AWS Secure 3-Tier Architecture** directly to its corresponding open-source, on-premises alternatives. All alternatives are packaged as rootless containers, operating under unprivileged user contexts and orchestrated natively via systemd Quadlets.

1. AWS-Native vs. On-Premises Open-Source Mapping Matrix

The table below outlines our transition framework, mapping cloud-managed services to self-hosted, sovereign container implementations.

AWS Core Component	On-Premises Open-Source Alternative	Host Location	Deployment Context	Security & Operational Controls
AWS ALB & WAFv2	BunkerWeb Edge Proxy (Nginx-based)	VM-01 (DMZ)	Rootless Quadlet Container	Built-in ModSecurity WAF rules, automated local Let's Encrypt / CA SSL, rate-limiting, and web administrative GUI.
AWS ASG Compute	Spring Boot 3.5.12 JVM + React 19	VM-02 (App)	Rootless Quadlet Pod (skm_pod)	Hardened Ubuntu host running unprivileged JVM, SSM-equivalent terminal logging, and structured systemd metrics.
Amazon RDS PostgreSQL	Percona Server for PostgreSQL 17 (or standard Postgres 17 + pgvector)	VM-03 (DB/AI)	Rootless Quadlet Container	High-Availability on-premises deployment utilizing Patroni and etcd , Volume storage mapped to physical NVMe disks, restricted VLAN 30 port access, and pg_backrest backups. See Enterprise Percona PostgreSQL 17 Guide for detailed setup.
ElastiCache Valkey	Valkey 7.x / 8.0 (Open Source)	VM-03 (DB/AI)	Rootless Quadlet Container	In-memory session store restricted to VLAN 20/30 ingress with forced password authentication.
Amazon Bedrock & GPUs	Ollama (Local Inference Model Engine)	VM-03 (DB/AI)	Rootless GPU-Passthrough Container	Secure offline execution of Qwen3 LLM and embedding models. Zero foreign network transfer.
Amazon Cognito	Keycloak or Authentik	VM-02 (App)	Rootless Quadlet Container	Managed OAuth 2.0, OpenID Connect (OIDC), active JWT generation, credential hashing, and user directory management.
RAG Engine & Metrics	RAGFlow + Langfuse	VM-03 (DB/AI)	Rootless Container Multi-Pod	Local OCR, semantic parsing, vector ingestion, and LLM telemetry tracking.
AWS Security Hub	Wazuh SIEM Manager	VM-01 (Mgmt)	Native VM / Standalone Container	Log aggregation, host configuration audits (ASIMP), and rootkit detection across all VMs.

2. Architectural Deep-Dive & Quadlet Blueprints

A. Presentation Layer Proxy: BunkerWeb

To replace the combined traffic filtering and TLS termination of AWS ALB and WAFv2, we utilize **BunkerWeb**. BunkerWeb is a hardened, open-source web application firewall (WAF) and reverse proxy based on Nginx.

```
# /home/songket/.config/containers/systemd/bunkerweb.container
[Container]
ContainerName=bunkerweb
Image=docker.io/bunkerity/bunkerweb:1.5.8
PublishPort=80:8080
PublishPort=443:8443
Volume=/var/srv/bunkerweb/data:/data:Z
Volume=/var/srv/bunkerweb/config:/etc/bunkerweb:Z
Environment=LISTEN_HTTP_PORT=8080
Environment=LISTEN_HTTPS_PORT=8443
Environment=USE_REVERSE_PROXY=yes
Environment=REVERSE_PROXY_URLS=/api/* http://10.10.20.10:8080/api/
Environment=AUTO_LETS_ENCRYPT=yes
Environment=BAD_IPS_FILTER=yes
Environment=LIMIT_REQ=yes
UserNS=keep-id:uid=2001,gid=2001

[Service]
Restart=always
```

B. Relational Database Layer: Percona Server for PostgreSQL 17

To replicate the reliability of Amazon RDS, we run Percona Server for PostgreSQL 17. It functions identically to upstream PostgreSQL 17 as a binary-compatible drop-in replacement.

To achieve production-grade high availability, resilience, and automated failover matching RDS, we implement a clustered setup using **Patroni** and **etcd**, managed via **pg_backrest** backups. Please refer to our specialized [Enterprise Percona Server for PostgreSQL 17 On-Premises Guide](#) for detailed step-by-step specifications, HAProxy load balancing config, and Quadlet blueprints.

For simpler standalone staging/development setups on-premises, standard unprivileged PostgreSQL containers with automated custom backup scripts can still be utilized:

```
# /home/songket/.config/containers/systemd/postgresql.container
[Container]
ContainerName=postgresql_db
Image=docker.io/pgvector/pgvector:0.7.0-pg17
Volume=/var/srv/postgresql/data:/var/lib/postgresql/data:Z
Environment=POSTGRES_DB=songket_db
Environment=POSTGRES_USER=songket_admin
Environment=POSTGRES_PASSWORD=SecureMasterPassword123
PublishPort=5432:5432
UserNS=keep-id:uid=2001,gid=2001

[Service]
Restart=always
```

 The Standalone Database Automated Backup Cron Script

To protect against standalone database corruption, a local shell script can be scheduled on VM-03 via a systemd timer (or standard cron):

```
#!/bin/bash
# Location: /var/srv/postgresql/scripts/backup_db.sh
BACKUP_DIR="/var/srv/postgresql/backups"
TIMESTAMP=$(date +%Y%m%d_%H%M%S)
BACKUP_FILE="${BACKUP_DIR}/songket_db_${TIMESTAMP}.sql.gz"

# Run pg_dump inside the unprivileged container context
podman exec postgresql_db pg_dump -U songket_admin songket_db | gzip > "${BACKUP_FILE}"

# Enforce a 30-day retention policy
find "${BACKUP_DIR}" -name "songket_db_*.sql.gz" -mtime +30 -exec rm {} \;

# Replicate the backup file to the local network-attached storage (NAS) on VLAN 30
rsync -az "${BACKUP_DIR}/" user@10.10.30.200:/mnt/nas/postgresql_backups/
```

C. Caching Layer: Valkey 7/8

Valkey is deployed locally as an unprivileged, high-performance in-memory key-value cache container, providing the same session management, rate-limiting, and counter mechanisms as Amazon ElastiCache.

```
# /home/songket/.config/containers/systemd/valkey.container
[Container]
ContainerName=valkey_cache
Image=docker.io/valkey/valkey:8.0.0
PublishPort=6379:6379
Volume=/var/srv/valkey/data:/data:Z
UserNS=keep-id:uid=2001,gid=2001

[Service]
Restart=always
ExecStart=valkey-server --requirepass MasterValkeyPass123 --protected-mode no
```

D. Identity & Access Layer: Keycloak

To replace Amazon Cognito, we deploy Keycloak inside VM-02. Keycloak manages the user directory, coordinates authentication tokens, signs JWT keys for Spring Boot verification, and handles OAuth 2.0 RBAC flows natively.

```
# /home/songket/.config/containers/systemd/keycloak.container
[Container]
ContainerName=keycloak_iam
Image=quay.io/keycloak/keycloak:24.0.5
PublishPort=8443:8443
Volume=/var/srv/keycloak/data:/opt/keycloak/data:Z
Environment=KEYCLOAK_ADMIN=admin
Environment=KEYCLOAK_ADMIN_PASSWORD=SecureAdminPassword123
Environment=KC_DB=postgres
Environment=KC_DB_URL=jdbc:postgresql://10.10.30.10:5432/songket_db
Environment=KC_DB_USERNAME=songket_admin
Environment=KC_DB_PASSWORD=SecureMasterPassword123
UserNS=keep-id:uid=2001,gid=2001

[Service]
Restart=always
ExecStart=kc.sh start-dev --http-port 8080 --https-port 8443
```

 3. Sovereign local AI Processing (Ollama & RAGFlow)

Rather than calling external LLM APIs (like Amazon Bedrock or OpenAI) which sends sensitive corporate and personal data over public networks, VM-03 executes AI inferencing and retrieval locally:

- 1. **Ollama:** Serves as the localized inference host. It mounts NVIDIA host GPU devices (using `--device nvidia.com/gpu=all` or equivalent Podman parameters) to execute **Qwen3 (Qwen2.5-Instruct)** LLM and dense embedding indexing sweeps.
- 2. **RAGFlow:** Mounts local folders containing enterprise documentation, parses them semantically using specialized unprivileged OCR and parser engines, and ingests vector structures natively into the PostgreSQL pgvector database.
- 3. **Langfuse:** Captures local Spring Boot LLM transaction traces, mapping operational latency and execution cost metrics cleanly to a local database.

This complete network-isolated lifecycle guarantees absolute compliance with both PDPA data transit restrictions and the highest level of corporate IP protection.

Deep State of Mind (DSOM) Harisfazillah Jamel (LinuxMalaysia) 2026-08-09

Enterprise Percona Server for PostgreSQL 17 on On-Premises

#onprem #virtualization #patroni #etcd #percona #postgresql #high-availability #pgbackrest #haproxy #pgbouncer

[STRATEGIC FINANCIAL] — Sovereignty, Compliance & Infrastructure Cost Control
[DEVOPS EXECUTION] — Systems Engineers & SREs
[SECURITY & COMPLIANCE] — Security Hardening & Platform Engineers

Enterprise Percona Server for PostgreSQL 17 On-Premises Setup

This guide details the high-availability (HA), enterprise-grade onsite on-premises deployment of **Percona Server for PostgreSQL 17**. It details the architectural design, clustering mechanisms, and step-by-step configurations required to achieve a resilient database platform matching or exceeding AWS RDS PostgreSQL Multi-AZ capabilities, completely isolated from external networks.

1. Executive Summary & Binary Compatibility


[Percona Server for PostgreSQL 17](#) is an enterprise-ready, open-source, binary-compatible drop-in replacement for upstream PostgreSQL 17. Because it functions identically to upstream PostgreSQL, it guarantees seamless, zero-risk migrations between the two platforms.

In addition to upstream core functionality, Percona Server introduces critical enterprise features out of the box, including:

- **Storage Manager (SMGR) API exposure:** Unlocks advanced custom storage manager integration for third-party extensions.
- **WAL Read/Write API exposure:** Enables specialized hooks into write-ahead logging.
- **pg_tde extension compatibility:** Allows data-at-rest encryption (Transparent Data Encryption), covering index-level and Write-Ahead Logging (WAL) encryption natively without performance degradation.
- **pg_stat_monitor integration:** Provides advanced query-level telemetry tracking, chronological execution plan analysis, and client metadata tracking.

2. High Availability Clustering Architecture (Patroni & etcd)

To achieve enterprise-grade resilience, we implement a highly available cluster with **Patroni** as the template/cluster management orchestrator and **etcd** as the Distributed Consensus Store (DCS). A minimum of three nodes are deployed to prevent “split-brain” scenarios and maintain quorum.



Component Roles

1. **Patroni:** A Python-based clustering manager that configures, boots, and monitors Percona Server. It interfaces with etcd to maintain cluster state, dynamically configuring PostgreSQL streaming replication and triggering automated failovers when a primary goes offline.
2. **etcd:** A distributed, reliable key-value store used as the Distributed Consensus Store (DCS). It tracks node health, holds the master/leader lock key, and coordinates configuration changes across the database cluster.

3. **HAProxy:** Acts as the endpoint load balancer, exposing a single stable endpoint. It queries Patroni's HTTP API (port 8008) to identify the current healthy primary node (/primary health check) and standby nodes (/replica health check), routing queries dynamically.
 4. **PgBouncer:** Lightweight connection pooler deployed on top of PostgreSQL to handle thousands of concurrent application connections, minimizing memory consumption.
 5. **pg_backrest:** An enterprise-grade, block-level backup utility executing compressed, parallel backups directly to our local on-premises network-attached storage (NAS).
-

3. High-Fidelity Configuration Blueprints

Deploying this stack within our unprivileged rootless container context requires detailed Quadlets and custom configurations on **VM-03**.

A. The DCS Consensus Engine: etcd.container

Deploy etcd inside an unprivileged container namespace on each of the three database virtual machines to form a resilient quorum.

```
# /home/songket/.config/containers/systemd/etcd.container
[Container]
ContainerName=etcd_node
Image=quay.io/coreos/etcd:v3.5.12
PublishPort=2379:2379
PublishPort=2380:2380
Volume=/var/srv/etcd/data:/etcd-data:Z
UserNS=keep-id:uid=2001,gid=2001
Environment=ETCD_NAME=etcd-01
Environment=ETCD_INITIAL_LISTEN_PEER_URLS=http://0.0.0.0:2380
Environment=ETCD_LISTEN_CLIENT_URLS=http://0.0.0.0:2379
Environment=ETCD_INITIAL_ADVERTISE_PEER_URLS=http://10.10.30.11:2380
Environment=ETCD_ADVERTISE_CLIENT_URLS=http://10.10.30.11:2379
Environment=ETCD_INITIAL_CLUSTER=etcd-01=http://10.10.30.11:2380,etcd-02=http://10.10.30.12:2380,etcd-03=http://10.10.30.13:2380
Environment=ETCD_INITIAL_CLUSTER_TOKEN=etcd-songket-token
Environment=ETCD_INITIAL_CLUSTER_STATE=new

[Service]
Restart=always
```

B. The Patroni Orchestrated Percona Server: patroni-percona.container

Patroni is packaged inside a custom image containing Percona Server for PostgreSQL 17, pg_backrest, and Patroni binaries.

```
# /home/songket/.config/containers/systemd/patroni-percona.container
[Container]
ContainerName=patroni_postgresql
Image=docker.io/percona/percona-postgresql-patroni:17
PublishPort=5432:5432
PublishPort=8008:8008
Volume=/var/srv/postgresql/data:/var/lib/postgresql/data:Z
Volume=/var/srv/postgresql/config:/etc/patroni:Z
UserNS=keep-id:uid=2001,gid=2001
Environment=PATRONI_CONFIGURATION_FILE=/etc/patroni/patroni.yml

[Service]
Restart=always
```

The Cluster Orchestration File: patroni.yml

```
# /var/srv/postgresql/config/patroni.yml
scope: songket-cluster
namespace: /service
name: pg-node-01

etcd3:
  hosts:
    - 10.10.30.11:2379
    - 10.10.30.12:2379
    - 10.10.30.13:2379

restapi:
  listen: 0.0.0.0:8008
  connect_address: 10.10.30.11:8008

bootstrap:
  dcs:
    ttl: 30
    loop_wait: 10
    retry_timeout: 10
    maximum_lag_on_failover: 1048576
    postgresql:
      use_pg_rewind: true
      use_slots: true
      parameters:
        shared_buffers: 16GB
        work_mem: 64MB
        maintenance_work_mem: 2GB
        max_connections: 500
        shared_preload_libraries: 'pg_stat_statements,pg_stat_monitor,pg_tde'
        archive_mode: "on"
        archive_command: "pgbackrest --stanza=songket-db archive-push %p"

init:
  pg_hba:
    - host replication replicator 10.10.30.0/24 md5
    - host all all 0.0.0.0/0 md5

users:
  admin:
    password: SecureMasterPassword123
    options:
```

- superuser
- createdb

```
postgresql:
  listen: 0.0.0.0:5432
  connect_address: 10.10.30.11:5432
  data_dir: /var/lib/postgresql/data/songket-cluster
  bin_dir: /usr/lib/postgresql/17/bin
  pgpass: /home/songket/.pgpass
  authentication:
    replication:
      username: replicator
      password: SecureReplicationPassword123
  superuser:
    username: postgres
    password: SecurePostgresPassword123

tags:
  nofailover: false
  noloadbalance: false
  clonefrom: false
  nosync: false
```

C. Traffic Routing Layer: haproxy.container

Exposes the single VIP endpoint for application compute nodes, dynamically selecting active primaries or passive replicas using Patroni’s HTTP endpoint.

```
# /home/songket/.config/containers/systemd/haproxy.container
[Container]
ContainerName=haproxy_loadbalancer
Image=docker.io/library/haproxy:2.8
PublishPort=5000:5000
PublishPort=5001:5001
Volume=/var/srv/haproxy/config:/usr/local/etc/haproxy:Z
UserNS=keep-id:uid=2001,gid=2001

[Service]
Restart=always
```

Routing Configuration: haproxy.cfg

```
# /var/srv/haproxy/config/haproxy.cfg
global
    maxconn 4096

defaults
    mode tcp
    timeout connect 5s
    timeout client 30m
    timeout server 30m

# Port 5000: Write-only routing to the active primary node
frontend write_pool
    bind 0.0.0.0:5000
    default_backend primary_pg_cluster

backend primary_pg_cluster
    option httpchk GET /primary
    http-check expect status 200
    default-server inter 3s fall 3 rise 2 on-marked-down shutdown-sessions
    server pg-node-01 10.10.30.11:5432 maxconn 500 check port 8008
    server pg-node-02 10.10.30.12:5432 maxconn 500 check port 8008
    server pg-node-03 10.10.30.13:5432 maxconn 500 check port 8008

# Port 5001: Read-only load-balanced routing to replica standby nodes
frontend read_pool
    bind 0.0.0.0:5001
    default_backend standby_pg_cluster

backend standby_pg_cluster
    balance roundrobin
    option httpchk GET /replica
    http-check expect status 200
    default-server inter 3s fall 3 rise 2
    server pg-node-01 10.10.30.11:5432 maxconn 500 check port 8008
    server pg-node-02 10.10.30.12:5432 maxconn 500 check port 8008
    server pg-node-03 10.10.30.13:5432 maxconn 500 check port 8008
```

4. Disaster Recovery & pg_backrest On-Premises Strategy

To achieve a **Zero-Data-Loss RPO** and match the backup standard of Amazon RDS, we utilize **pg_backrest**. pg_backrest performs point-in-time recovery (PITR) by continuously streaming Write-Ahead Logs (WAL) and executing compressed, parallel backups.

The Backup Configuration: pgbackrest.conf

```
# /etc/pgbackrest/pgbackrest.conf
[global]
repol-path=/mnt/nas/pgbackrest
repol-retention-full=4
repol-retention-diff=6
process-max=4
log-level-console=info
log-level-file=detail
compress-type=lz4
compress-level=3

[songket-db]
```

```
pg1-path=/var/lib/postgresql/data/songket-cluster
pg1-user=postgres
```

Automated Backup Pipeline Playbook

Backup scheduling is automated on-premises via standard crontab or Ansible task schedulers:

```
# Full Backup every Sunday at 01:00 AM local time (MYT)
00 01 * * 0 podman exec patroni_postgresql pgbackrest --stanza=songket-db --type=full backup

# Incremental Backup every daily at 01:00 AM MYT
00 01 * * 1-6 podman exec patroni_postgresql pgbackrest --stanza=songket-db --type=incr backup
```

5. Security & Isolation Controls

Operating an enterprise database on-premises demands rigid logical boundaries and hardening steps conforming with **Bank Negara Malaysia guidelines** and the **Malaysian Personal Data Protection Act (PDPA) 2010**:

- 1. **VLAN Segmentation:** The database VMs (VLAN 30) are blocked from communicating with the outer DMZ (VLAN 10). Incoming connections must originate strictly from VLAN 20 (Application servers VM-02) or dedicated database replica IPs.
- 2. **mTLS Encryption:** All inter-node cluster replication traffic (PostgreSQL streaming) and consensus coordination (Patroni-to-etcd) are encrypted using locally generated, self-signed CA certificates mapped inside the unprivileged container volumes.
- 3. **Transparent Data Encryption (TDE):** Utilizing Percona’s unique **pg_tde** extension, all Write-Ahead Logs (WAL) and data block pages are encrypted on disk. This protects customer personal identifiable information (PII) from physical disk theft or hypervisor compromises.

Deep State of Mind (DSOM) Harisfazillah Jamel (LinuxMalaysia) 2026-08-09

VPC Module

[#aws](#) [#cloud](#) [#architecture](#) [#vpc](#) [#alb](#) [#asg](#) [#rds](#) [#dns](#) [#disaster-recovery](#)

VPC Module

The VPC Module deploys the foundational multi-AZ networking infrastructure for the 3-tier layout. It partitions the network into public subnets, private application subnets, and isolated private database subnets, ensuring high availability and physical logical boundaries.

Architectural Details

- **VPC (`aws_vpc`):** Created with a custom CIDR (default `10.0.0.0/16`) and has `enable_dns_support` and `enable_dns_hostnames` set to `true` to facilitate intra-VPC DNS resolution.
- **Internet Gateway (`aws_internet_gateway`):** Attached to the VPC to route external traffic to and from the public subnets.
- **Subnets (`aws_subnet`):**
 - **Public Subnets:** Deployed across multiple AZs. Maps public IP addresses automatically on launch. Hosts ALBs and NAT Gateways.
 - **Private App Subnets:** Deployed across multiple AZs. Holds ASG instances. Routes outbound internet traffic via NAT Gateways.
 - **Private DB Subnets:** Dedicated subnets for RDS DB placement. Completely lacks routes to the NAT Gateways or Internet Gateways to minimize access.
- **NAT Gateways (`aws_nat_gateway`):** Provisioned in public subnets with a dedicated Elastic IP (`aws_eip`) in each AZ to provide high-availability outbound routes for instances in private subnets.
- **Route Tables and Associations:** Configured to manage separate traffic routes for public, private application, and database subnets.

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module's inline documentation at `terraform/modules/vpc/README.md`.

Security Groups Module

[#aws](#) [#cloud](#) [#architecture](#) [#alb](#) [#asg](#) [#rds](#) [#disaster-recovery](#) [#postgresql](#)

Security Groups Module

This module defines stateful firewall rules (Security Groups) that isolate each layer of the 3-Tier topology, strictly adhering to the **Zero-Trust Network Principle**.

Firewall Rulesets

Presentation Layer: ALB Security Group (`alb_sg`)

- **Ingress:** Accepts inbound public traffic on:
 - Port 80 (HTTP) from any IPv4 or IPv6 address (`0.0.0.0/0, ::/0`).
 - Port 443 (HTTPS) from any IPv4 or IPv6 address (`0.0.0.0/0, ::/0`).
- **Egress:** Allows outbound connections to any destination (`0.0.0.0/0`) to route filtered requests.

Application Layer: ASG/EC2 Security Group (`asg_sg`)

- **Ingress:** Accepts traffic **strictly** from the ALB Security Group on the application HTTP port (80). No direct internet access is permitted.
- **Egress:** Allows outbound connections to any destination (`0.0.0.0/0`) to allow instances to perform updates or retrieve application packages.

Database Layer: RDS Security Group (`db_sg`)

- **Ingress:** Accepts traffic **strictly** from the ASG/EC2 Security Group on the database port (5432 for PostgreSQL / 3306 for MySQL). No other traffic is allowed.
 - **Egress:** Blocked to restrict exfiltration.
-

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module’s inline documentation at `terraform/modules/security_groups/README.md`.

Web Application Firewall (WAF) Module

[#aws](#) [#cloud](#) [#architecture](#) [#alb](#) [#waf](#) [#disaster-recovery](#)

Web Application Firewall (WAF) Module

The WAF Module provisions an AWS WAFv2 Web ACL (Web Application Firewall) configured with modern security rules to defend against automated layer-7 attacks, SQL Injection (SQLi) attempts, and brute-force traffic floods.

Active Protection Rule Sets

- Common Rule Set (OWASP Top 10):**
 - Implements AWS Managed Common Rules (`AWSManagedRulesCommonRuleSet`) which protect against standard vulnerabilities, unauthorized administrative path access, and exploitation techniques.
- SQLi Defense (SQL Injection Protection):**
 - Implements AWS Managed SQLi Rules (`AWSManagedRulesSQLiRuleSet`) to identify and block malicious SQL injections in query parameters, headers, or request bodies.
- IP Rate Limiting:**
 - Adds a custom rate-based rule to restrict the maximum number of requests a single client IP address can send within a sliding 5-minute window (default: 2000 requests).

Edge Association

- ALB Association (`aws_wafv2_web_acl_association`):** Directly attaches the regional Web ACL to the Application Load Balancer ARN to filter traffic at the network edge.

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module's inline documentation at `terraform/modules/waf/README.md`.

Application Load Balancer (ALB) Module

#aws #cloud #architecture #alb

Application Load Balancer (ALB) Module

This module deploys a highly available, public-facing Application Load Balancer (ALB) that acts as the entry point for all incoming HTTP web traffic. It handles request distribution and active health checking.

Technical Components

- **Load Balancer (`aws_lb`):** An external, internet-facing application load balancer mapped across multiple public subnets to provide high availability.
- **Target Group (`aws_lb_target_group`):** A target group with active health checks to monitor the health of the underlying EC2 instances managed by the Auto Scaling Group.
 - **Path:** /
 - **Protocol:** HTTP
 - **Healthy/Unhealthy Threshold:** 3
 - **Interval:** 30 seconds
 - **Timeout:** 5 seconds
- **Listener (`aws_lb_listener`):** Configured on port 80 (HTTP) to forward traffic to the target group.

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module's inline documentation at `terraform/modules/alb/README.md`.

Auto Scaling Group (ASG) Module

[#aws](#) [#cloud](#) [#architecture](#) [#asg](#) [#bastion](#)

Auto Scaling Group (ASG) Module

This module deploys an Auto Scaling Group (ASG) of EC2 instances spanning multiple Availability Zones. It automates high-availability application scaling and supports dynamic resource scaling based on real-time CPU consumption.

Technical Details

- **Multi-AZ Resilience:** ASG instances are launched strictly inside private subnets, ensuring they are shielded from direct internet exposure.
- **Dynamic Graviton / ARM64 Auto-detection:**
 - The module looks up the official Amazon Linux 2023 AMI dynamically.
 - If the `instance_type` belongs to the AWS Graviton family (e.g., starts with `t4g`, `m6g`, `c6g` etc.), the module automatically switches to fetch the **ARM64** Amazon Linux 2023 AMI. Otherwise, it defaults to the standard **x86_64** AMI.
- **SSM Integration:** Installs and configures IAM instance profiles with `AmazonSSMManagedInstanceCore` policies, allowing administrators to establish secure terminal shell sessions with instances without deploying bastion hosts or opening public SSH ports.
- **Nginx Bootstrapping:** Provisions instances with an automated bootstrapping `user_data` script that updates system packages, installs Nginx, starts the service, and creates a customized region-specific welcome index page.
- **Dynamic CPU Scaling:**
 - **Scale Out:** Triggers when average CPU utilization matches or exceeds **70%** for two consecutive evaluation periods.
 - **Scale In:** Triggers when average CPU utilization falls below or matches **30%** for two consecutive evaluation periods.

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module's inline documentation at `terraform/modules/asg/README.md`.

RDS Module

[#aws](#) [#cloud](#) [#architecture](#) [#rds](#) [#ssl](#) [#postgresql](#)

RDS Module

The RDS Module deploys a secure, fully-managed RDS Database instance running in Multi-AZ configuration to ensure maximum availability and seamless automatic failover in the event of an availability zone outage.

Technical Details

- **Multi-AZ Availability:** Automatically provisions a primary database instance in one AZ and a synchronous hot standby instance in a secondary AZ.
- **Deep Private Isolation:** Placed strictly inside dedicated isolated database subnets without any routes to the internet or NAT gateways.
- **Dynamic Parameter Group Tuning:**
 - Creates a dedicated parameter group with custom database parameters.
 - Automatically adapts the DB parameter group family (e.g., `postgres16`, `postgres15`, or `mysql8.0`) based on the selected database engine and major version.
- **Storage Encryption:** Enforces EBS storage encryption (`storage_encrypted = true`) using AWS managed KMS keys.
- **Storage Autoscaling:** Supports storage scale thresholds (from an initial 20 GB up to 100 GB max) to seamlessly accommodate business data growth without manual intervention.

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module's inline documentation at `terraform/modules/rds/README.md`.

Standalone EC2 Instance Module

[#aws](#) [#cloud](#) [#architecture](#) [#vpc](#) [#alb](#) [#bastion](#) [#disaster-recovery](#)

Standalone EC2 Instance Module

This module provisions secure Standalone EC2 Instances running **Ubuntu 26.04 LTS** (Noble Numbat successor) to support custom application build-ups, staging tools, developer sandboxes, or other application resources that are not ready or suited for Auto Scaling Group deployment.

Technical Details

- **Secure Network Isolation:**
 - Standalone instances are launched strictly inside the **VPC Private Application Subnets**.
 - They do not have public IP addresses and are completely hidden from direct public internet ingress.
 - Outbound traffic (e.g., package updates via `apt - get`, package validation via `debsums`, and security auditing scans via `ASIMP`) is securely routed through the managed NAT Gateways.
 - **SSM Integration:**
 - Deploys dedicated IAM Roles and Instance Profiles with `AmazonSSMManagedInstanceCore` policies attached.
 - Enables developers to establish secure terminal shell sessions with standalone instances using AWS Systems Manager (SSM) without setting up bastion hosts or opening public SSH ports.
 - **Chained Security Grouping:**
 - Deploys a dedicated security group (`standalone-ec2-sg`) for standalone instances.
 - Restricts inbound traffic on port 80/443 strictly to the Application Load Balancer (ALB) security group, enabling developers to run web-facing test services or APIs behind the secure ALB.
 - **Bootstrapping user_data:**
 - Standardizes initial packages, applies upgrades, and installs Nginx as a placeholder server.
 - Ready to run custom `ASIMP` (Ansible System Integrity Management Platform) hardening playbooks.
-

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module's inline documentation at `terraform/modules/standalone_ec2/README.md`.

ElastiCache Valkey Module

[#aws](#) [#cloud](#) [#architecture](#) [#vpc](#) [#asg](#) [#elasticsearch](#) [#valkey](#) [#disaster-recovery](#)

ElastiCache Valkey Module

The ElastiCache Valkey Module deploys a secure, fully-managed **Amazon ElastiCache for Valkey** cache cluster inside the private database subnets of your VPC.

Valkey is a high-performance, fully open-source key-value database stewarded by the Linux Foundation. It acts as a drop-in replacement for Redis OSS, offering identical client protocol and command compatibility while delivering **20% lower on-demand pricing** for self-designed (node-based) clusters.

Technical Details

- **Deep Private Isolation:** Placed strictly inside dedicated database subnets without internet egress.
- **Valkey Engine Selection:** Uses valkey engine with version 7.2 as the default modern cache cluster, avoiding Redis OSS license premium.
- **Transit and At-Rest Encryption:**
 - Enforces TLS in-transit encryption (`transit_encryption_enabled = true`).
 - Enforces AES-256 at-rest encryption (`at_rest_encryption_enabled = true`).
- **Dynamic Ingress Firewalling:**
 - Inbound traffic on port 6379 is allowed exclusively from the application Auto Scaling Group (ASG) Security Group.
 - Optionally allows inbound port 6379 connections from standalone staging/developer instances to support 1:1 environment parity for integration testing.

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module’s inline documentation at `terraform/modules/elasticsearch/README.md`.

SSH Jumphost Module

#aws #cloud #architecture #vpc #asg #jumphost #bastion #disaster-recovery

AWS SSH Jumphost (Bastion) Module

This module provisions a highly secure, cost-optimized SSH Jumphost (Bastion) running inside the VPC’s public subnet. It is whitelisted exclusively for incoming connections from your developer office network (e.g., Cyberjaya, Selangor, Malaysia) on port 22.

The module also injects ingress SSH security group rules directly into the Auto Scaling Group (ASG) and Standalone EC2 instances’ security groups, allowing developers to establish secure SSH proxying/tunnels into the private environment.

Core Features

- **Static Public Routing:** Allocates and associates a dedicated AWS Elastic IP (EIP) with the Jumphost, giving developers a stable public IP to configure in their SSH client configurations.
- **Strict Ingress Whitelisting:** Keying security groups to restrict SSH (TCP port 22) access to the specified office CIDR only, preventing brute-force and port-scanning attempts from other parts of the internet.
- **SSH Chaining to Private Resources:** Attaches managed security group ingress rules to the private subnet ASG and standalone compute groups, enabling passwordless SSH key proxying.
- **Operating System Versatility:** Allows switching between canonical **Ubuntu 24.04/26.04 LTS** (fully compatible with **ASIMP** OS-level hardening) and cloud-optimized **Amazon Linux 2023** via input parameters.
- **SSM-Managed Backup:** Connects to AWS Systems Manager (SSM) by default via an IAM instance profile, providing a safe, passwordless shell backup route for system administrators.

Inputs

Name	Description	Type		Default	Required
environment	Environment name for tagging (e.g., production, dev)	string	n/a		yes
vpc_id	The ID of the VPC	string	n/a		yes
public_subnet_ids	List of public subnet IDs in the VPC	list(string)	n/a		yes
instance_type	EC2 compute instance size (Graviton-compatible default)	string	"t4g.micro"		no
allowed_ssh_cidr	IP CIDR allowed to connect to the Jumphost via SSH	string	"103.188.0.0/16"		no
jumphost_os	Operating System pattern (‘ubuntu’ or ‘amazon-linux-2023’)	string	"ubuntu"		no
ami_id	Specific AMI ID override (if left empty, AMI is auto-fetched)	string	" "		no
ubuntu_ami_filter_name	The search pattern name for the Ubuntu AMI	string	"ubuntu/images/hvm-ssd-gp3/ubuntu-noble-24.04-*server-*"		no
asg_sg_id	Security Group ID associated with the Auto Scaling Group	string	n/a		yes
standalone_sg_id	Security Group ID associated with Standalone EC2 instances	string	" "		no

Outputs

Name	Description
jumphost_public_ip	The static Elastic IP address assigned to the Jumphost
jumphost_private_ip	The private IP address of the Jumphost within the VPC
security_group_id	The ID of the security group assigned to the Jumphost
instance_id	The ID of the Jumphost EC2 instance

Automation Scripts

[#aws](#) [#cloud](#) [#opentofu](#) [#simulation](#) [#scripts](#) [#deployment](#)

Automation Scripts

The project includes CLI helper scripts and bootstrapping scripts under the `scripts/` directory to automate common tasks, local testing, offline simulation, and instance provisioning.

1. Local Deployment Script (`scripts/deploy.sh`)

This script automates the full lifecycle of a local OpenTofu deployment to Amazon Web Services (AWS). It performs validation and planning checks before prompting the user for confirmation to apply.

Steps Executed:

- OpenTofu Installation Check:** Verifies that the `tofu` CLI is installed and available in the execution path.
- Directory Context:** Changes the working directory to `terraform/`.
- Environment Configuration Check:** Checks if `terraform.tfvars` exists. If not, it copies it from `terraform.tfvars.example` and prompts the user to review.
- Initialization (`tofu init`):** Downloads required providers and module configurations.
- Format Verification (`tofu fmt`):** Formats all configuration files recursively to match canonical HCL syntax.
- Validation (`tofu validate`):** Verifies syntax correctness and internal variable consistency.
- Execution Planning (`tofu plan`):** Generates a secure plan file `tfplan`.
- User Confirmation Prompt:** Asks Do you want to apply this deployment? (y/n). If y, it runs `tofu apply tfplan`.

2. Infrastructure Teardown Script (`scripts/destroy.sh`)

This script provides a clean and automated way to destroy all resources managed by OpenTofu, preventing accidental dangling resources or billing charges.

Steps Executed:

- OpenTofu Installation Check:** Verifies that the `tofu` CLI is installed and available.
- Initialization:** Confirms that OpenTofu has been initialized and that `.terraform` config exists.
- Resource Destruction (`tofu destroy`):** Prompts the user to confirm the termination of all resources. Once confirmed, it cleanly removes all VPC components, load balancers, database clusters, auto-scaling groups, and WAF rules.

3. Offline Simulation & Verification Script (`scripts/simulate.sh`)

This script executes offline simulation and unit test suites without requiring active AWS cloud credentials or live AWS API connections. It allows developers and AI agents to test OpenTofu configurations and any newly added scripts against our architecture rules (e.g., IMDSv2 enforcement, zero-trust network isolation, Graviton defaults, and Valkey caching).

Steps Executed:

- Directory Context Anchor:** Sets directory context to the repository root.
- Pytest Simulation Suite Execution:** Runs `pytest tests/test_opentofu_simulation.py -v` to perform AST-based structure, rule, and security constraint checks.
- OpenTofu HCL Static Validation (if OpenTofu CLI installed):**
 - Formats configurations (`tofu fmt -recursive terraform/`).
 - Initializes backend-less working directory (`tofu -chdir=terraform init -backend=false`).
 - Validates OpenTofu configurations (`tofu -chdir=terraform validate`).

4. Instance Bootstrapping Script (`scripts/user_data.sh`)

This is a standalone bootstrapping script for EC2 instances. It can be used for custom instance images or standalone EC2 deployment tests.

Features:

- Updates the operating system packages.
- Installs and configures an Apache HTTP web server.
- Fetches Instance Metadata Service v2 (IMDSv2) security tokens dynamically.
- Retrieves instance metadata, including the Instance ID and host Availability Zone, displaying this dynamic diagnostic information on a beautifully styled HTML index page.

5. End-to-End Workflow on Any Linux Environment

You can git clone this repository on any Linux distribution (Ubuntu, Debian, RHEL, Fedora, Arch, Amazon Linux) and execute the complete deployment or offline simulation workflow:

```
# 1. Clone repository on your Linux host
git clone https://github.com/songketmail/aws-3tier-deployment-for-ai-infra.git
cd aws-3tier-deployment-for-ai-infra

# 2. Configure environment variables and credentials
cp terraform/terraform.tfvars.example terraform/terraform.tfvars
# (Edit terraform/terraform.tfvars with your specific IP ranges and passwords)

# 3. Run offline simulation test suite (test scripts and OpenTofu rules without AWS access)
./scripts/simulate.sh

# 4. Begin live AWS deployment (requires AWS CLI / IAM credentials configured)
./scripts/deploy.sh

# 5. Teardown deployed AWS infrastructure cleanly when done
./scripts/destroy.sh
```

CI/CD Pipeline Documentation

[#aws](#) [#cloud](#) [#architecture](#) [#rds](#)

CI/CD Pipeline Documentation

This project includes a fully automated CI/CD pipeline configured via GitHub Actions in `.github/workflows/opentofu.yml`. It ensures that every code submission is vetted for quality, syntax correctness, and security boundaries.

Pipeline Workflow Trigger Conditions

The pipeline is triggered automatically on:

- **Pull Requests** targeting the main branch.
- **Direct Pushes** or Merges to the main branch.

Job Definitions

The workflow consists of four major jobs designed with security and safety gates:

1. OpenTofu Format and Validate (`opentofu-lint-and-validate`)

- **Environment:** `ubuntu-latest`
- **Steps:**
 - **Checkout Code:** Checks out the repository files.
 - **Setup OpenTofu:** Installs OpenTofu version 1.8.2 on the runner using `opentofu/setup-opentofu@v1`.
 - **Format Check:** Verifies formatting recursively via `tofu fmt -check -recursive`.
 - **Initialize:** Runs `tofu init -backend=false` to configure modules locally.
 - **Validate:** Evaluates configurations for semantic validity via `tofu validate`.

2. Check Secrets Availability (`check-secrets`)

- **Environment:** `ubuntu-latest`
- **Purpose:** Checks whether the sensitive OIDC IAM Role `AWS_ROLE_TO_ASSUME` secret is populated without causing parser failures.
- **Steps:**
 - **Evaluate Secret Presence:** Evaluates `secrets.AWS_ROLE_TO_ASSUME` on the runner at step-level, dynamically exporting a boolean output variable `has-aws-role` (set to true if populated, or false if empty, such as in fork pull requests). This avoids static workflow validation errors.

3. OpenTofu Planning (`opentofu-plan`)

- **Trigger:** Runs only on **Pull Request** events.
- **Dependency:** Requires both `opentofu-lint-and-validate` and `check-secrets` jobs to pass, with `check-secrets.outputs.has-aws-role == 'true'`.
- **Steps:**
 - **Configure AWS OIDC Credentials:** Authenticates securely with AWS using OpenID Connect (OIDC) through `aws-actions/configure-aws-credentials@v4` with `secrets.AWS_ROLE_TO_ASSUME` and `secrets.AWS_REGION`.
 - **OpenTofu Init & Plan:** Performs complete backend initialisation and generates an execution plan showing what changes will be applied.

4. OpenTofu Deployment (`opentofu-apply`)

- **Trigger:** Runs only on **Pushes** or Merges to the main branch.
- **Dependency:** Requires both `opentofu-lint-and-validate` and `check-secrets` jobs to pass, with `check-secrets.outputs.has-aws-role == 'true'`.
- **Steps:**
 - **Configure AWS OIDC Credentials:** Authenticates securely with AWS using OIDC.
 - **OpenTofu Init & Apply:** Runs `tofu apply -auto-approve` to deploy the target infrastructure.

Secret Management Best Practices

The pipeline utilizes modern security standards:

- **No Hardcoded Credentials:** Leverages short-lived, dynamically-exchanged AWS security tokens via OIDC instead of storing long-lived `AWS_ACCESS_KEY_ID` and `AWS_SECRET_ACCESS_KEY` secrets.
- **Configurable Region:** Adapts automatically based on `secrets.AWS_REGION` or defaults gracefully to `us-east-1` (or local regional default config).

GitLab CI/CD & AWS EFS Code Deployment


#aws #cloud #architecture #vpc #alb #asg #rds #bastion #dns #disaster-recovery #gitlab #efs

GitLab CI/CD & AWS EFS Code Deployment Architecture

This guide provides a comprehensive design and deployment blueprint for automating application delivery. It outlines a secure, scalable workflow where developers use **GitLab** for source control, triggering automated **GitLab CI/CD pipelines** that deploy code directly onto an **Amazon Elastic File System (EFS)**.
By mounting this shared network filesystem (NFSv4) concurrently on both **Auto Scaling Groups (ASGs)** and **Standalone EC2 instances**, application updates become instantly active across all nodes without rebuilding AMIs or performing slow rolling redeployments.

1. High-Level Architecture Overview

The system bridges an external GitLab repository with a secure, private AWS 3-tier VPC. Since EFS endpoints reside deep within private subnets, we implement a highly secure staging-and-sync mechanism.



Core Flow Steps:


- Commit & Push:** A developer pushes code or creates a git tag in GitLab.
- GitLab CI/CD Trigger:** The GitLab runner starts, authenticates dynamically to AWS using **OIDC IAM Role Assumption**, runs tests, compiles build artifacts (e.g., node modules, frontend assets), and tars the build.
- Artifact Staging:** The runner uploads the tarball artifact to a secure, private **AWS S3 Deployment Bucket**.
- Automated EFS Syncing:** The runner triggers an **AWS Systems Manager (SSM) Run Command** targeting a dedicated internal Utility/Sync instance (or the Bastion host) that has the EFS volume mounted.
- Atomic Unpacking:** The Utility instance downloads the tarball from S3, unpacks it into the appropriate EFS sub-directory, and atomically swaps symlinks to complete the deploy in milliseconds.
- Instant Access:** Because the EFS volume is mounted across all active ASG instances and standalone staging nodes, the code is immediately available to the running applications without any instance terminations or server reboots.

2. Shared EFS Directory & Mounting Strategy

To maintain complete Isolation of Concerns and prevent staging/production resource conflicts, EFS must be structured logically.

2.1 EFS Directory Layout

We partition EFS into two top-level environments (production and staging) with application sub-directories:



2.2 Secure Mounting on EC2 Instances (Ubuntu 26.04 LTS & AL2023)

To ensure reliable, encrypted-in-transit connections to EFS, we use amazon-efs-utils over NFSv4 with TLS (Port 2049).

A. Interactive or Bootstrap Script Mounting:

```
# Install EFS Utilities
# For Amazon Linux 2023 (ASG Nodes):
sudo dnf install -y amazon-efs-utils

# For Ubuntu 26.04 LTS (Standalone Staging Nodes):
sudo apt-get update && sudo apt-get install -y binutils cpp
git clone https://github.com/aws/efs-utils.git /tmp/efs-utils
cd /tmp/efs-utils && ./build-deb.sh
sudo apt-get install -y ./build/amazon-efs-utils*deb

# Create Mount Directories
sudo mkdir -p /mnt/efs

# Mount with TLS enabled (Encapsulated over Port 2049)
# Replace fs-xxxxxx with your actual EFS File System ID
sudo mount -t efs -o tls,iam fs-xxxxxx:/ /mnt/efs
```

B. Persistent Mount Configuration (/etc/fstab):

To guarantee EFS is automatically remounted if an ASG instance launches or a standalone node is rebooted, append the following entry to /etc/fstab:

```
fs-xxxxxx:/ /mnt/efs efs _netdev,tls,iam,defaults 0 0
```

(The _netdev option prevents the system from attempting to mount EFS before network interfaces are fully initialized).

3. Nginx Configuration & Path Design

To serve code residing on network-attached storage efficiently, Nginx must be configured to minimize file system traversal overhead and handle the static/dynamic paths correctly.

3.1 Nginx Path Allocation

Instance Role	Local Mount Path App Root Directory (Served by Nginx)	
Production Frontend ASG	/mnt/efs	/mnt/efs/production/frontend/
Production Backend ASG	/mnt/efs	/mnt/efs/production/backend/ (App running via PM2/systemd)
Standalone Staging EC2	/mnt/efs	/mnt/efs/staging/frontend/ & /mnt/efs/staging/backend/

3.2 Performance Tuning for Network Filesystems (EFS)

Because EFS is a distributed network filesystem, standard POSIX metadata lookups (like stat(), which checks if a file exists or has changed) introduce millisecond network round-trips. Under high web traffic, this can severely throttle Nginx performance.

Mitigation: Enforce Nginx-level file descriptor caching. Add the following directives inside your global /etc/nginx/nginx.conf (within the http block):

```
# Cache open file descriptors, their sizes, and modification times
open_file_cache max=2000 inactive=20s;
open_file_cache_valid 30s;
open_file_cache_min_uses 2;
open_file_cache_errors on;
```

3.3 Production Frontend Nginx Configuration (ASG Node)

Apply this configuration to the instances running in your **Frontend Web ASG** (/etc/nginx/sites-available/default):

```
server {
    listen 80;
    server_name app.linuxmalaysia.com;

    # Root pointing to the production frontend EFS path
    root /mnt/efs/production/frontend;
    index index.html index.htm;

    # Performance optimizations for EFS
    sendfile on;
    tcp_nopush on;
    tcp_nodelay on;

    # Enable Gzip Compression to save bandwidth
    gzip on;
    gzip_types text/plain text/css application/json application/javascript text/xml;

    # Standard SPA Routing or Static File Server
    location / {
        try_files $uri $uri/ /index.html;
    }

    # Route backend API requests to the Backend ALB
    location /api/ {
        # Core VPC DNS Resolver with short TTL to handle Backend ALB scaling
        resolver 169.254.169.254 valid=10s;
        set $backend_upstream "http://internal-production-backend-alb-123456.ap-southeast-5.elb.amazonaws.com";

        proxy_pass $backend_upstream;
        proxy_set_header Host $host;
    }
}
```

```

        proxy_set_header X-Real-IP $remote_addr;
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
        proxy_set_header X-Forwarded-Proto $scheme;
    }

    # Static Assets Cache Control
    location ~* \.(?:css|js|jpg|jpeg|gif|png|ico|svg|woff|woff2|ttf|otf)$ {
        expires 7d;
        add_header Cache-Control "public, no-transform";
        access_log off;
    }

    # Custom Error Pages
    error_page 500 502 503 504 /50x.html;
    location = /50x.html {
        root /usr/share/nginx/html;
    }
}

```

3.4 Standalone Staging Nginx Configuration (EC2 Staging Node)

Staging nodes host both frontend and backend APIs on a single standalone instance for cost efficiency. Apply this layout (/etc/nginx/sites-available/default):

```

server {
    listen 80;
    server_name staging.linuxmalaysia.com;

    # Staging Frontend Root Directory
    root /mnt/efs/staging/frontend;
    index index.html;

    sendfile on;

    # Staging Frontend SPA Routing
    location / {
        try_files $uri $uri/ /index.html;
    }

    # Staging API reverse proxy pointing to a local Node.js/Python service
    location /api/ {
        # Staging backend app usually runs locally on port 5000
        proxy_pass http://127.0.0.1:5000;
        proxy_set_header Host $host;
        proxy_set_header X-Real-IP $remote_addr;
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
        proxy_set_header X-Forwarded-Proto $scheme;

        # Disable buffering for real-time staging logs/WebSocket support
        proxy_buffering off;
    }
}

```

4. Automated GitLab CI/CD Pipeline Configuration

To trigger automatic deployments based on code versioning, we configure the pipeline to recognize:

1. **Pushes to main branch:** Deploys instantly to the **Staging / Standalone** environment.
2. **Git Release Tags (v*.*.*):** Deploys to the **Production / ASG** environment.

Production-Ready .gitlab-ci.yml Manifest

Place this file in the root of your code repository:

```

stages:
  - lint
  - build
  - deploy

variables:
  AWS_DEFAULT_REGION: "ap-southeast-5" # Malaysia Region
  S3_STAGING_BUCKET: "s3://company-deployment-staging-bucket-ap-southeast-5"
  S3_PROD_BUCKET: "s3://company-deployment-prod-bucket-ap-southeast-5"
  SYNC_INSTANCE_STAGING_ID: "i-0123456789staging" # Standalone EC2 Staging Node ID
  SYNC_INSTANCE_PROD_ID: "i-0123456789utility" # Prod EFS Admin Utility Instance ID

# AWS Authentication via OIDC (No persistent IAM credentials stored in GitLab!)
.aws-auth:
  id_tokens:
    MY_OIDC_TOKEN:
      aud: https://gitlab.com
  before_script:
    - mkdir -p ~/.aws
    - echo "AssumeRoleWithWebIdentity"
    - >
      export $(printf "AWS_ACCESS_KEY_ID=%s AWS_SECRET_ACCESS_KEY=%s AWS_SESSION_TOKEN=%s"
        $(aws sts assume-role-with-web-identity
          --role-arn "arn:aws:iam::112233445566:role/GitLabCI-AWS-OIDC-Role"
          --role-session-name "GitLabPipeline-${CI_PIPELINE_ID}"
          --web-identity-token "${MY_OIDC_TOKEN}"
          --query "Credentials.[AccessKeyId,SecretAccessKey,SessionToken]"
          --output text))

# ===== LINT STAGE =====
code-lint:
  stage: lint
  image: node:20-alpine
  script:
    - npm ci

```

```

- npm run lint
rules:
- if: $CI_COMMIT_BRANCH == "main"
- if: $CI_COMMIT_TAG =~ /^v\d+\.\d+\.\d+$/

# ===== BUILD STAGE =====
build-assets:
stage: build
image: node:20-alpine
artifacts:
paths:
- dist/
- package.json
- package-lock.json
expire_in: 1 week
script:
- npm ci
- npm run build
rules:
- if: $CI_COMMIT_BRANCH == "main"
- if: $CI_COMMIT_TAG =~ /^v\d+\.\d+\.\d+$/

# ===== DEPLOY STAGES =====

# Deploy Staging: Automatically triggered on commits/merges to main branch
deploy-staging:
extends: .aws-auth
stage: deploy
image: amazon/aws-cli:latest
dependencies:
- build-assets
script:
# 1. Package the build files
- tar -czf staging-build.tar.gz dist/ package.json package-lock.json

# 2. Upload tarball to Staging S3 Bucket
- aws s3 cp staging-build.tar.gz ${S3_STAGING_BUCKET}/staging-build.tar.gz

# 3. Trigger AWS SSM Run Command on Staging Instance to pull and unpack
- >
aws ssm send-command
--instance-ids "${SYNC_INSTANCE_STAGING_ID}"
--document-name "AWS-RunShellScript"
--comment "Deploy staging files from GitLab CI"
--parameters 'commands=[
"aws s3 cp "${S3_STAGING_BUCKET}"'/staging-build.tar.gz /tmp/staging-build.tar.gz",
"mkdir -p /mnt/efs/staging/frontend_new",
"tar -xzf /tmp/staging-build.tar.gz -C /mnt/efs/staging/frontend_new",
"rm -rf /mnt/efs/staging/frontend_old",
"[ -d /mnt/efs/staging/frontend ] && mv /mnt/efs/staging/frontend /mnt/efs/staging/frontend_old",
"mv /mnt/efs/staging/frontend_new /mnt/efs/staging/frontend",
"rm -rf /mnt/efs/staging/frontend_old /tmp/staging-build.tar.gz",
"echo Deployment Successful"
]'
rules:
- if: $CI_COMMIT_BRANCH == "main"

# Deploy Production: Triggered exclusively by creating a Tag (e.g., v1.0.0)
deploy-production:
extends: .aws-auth
stage: deploy
image: amazon/aws-cli:latest
dependencies:
- build-assets
script:
# 1. Package the production files
- tar -czf prod-build.tar.gz dist/ package.json package-lock.json

# 2. Upload to Production S3 Bucket
- aws s3 cp prod-build.tar.gz ${S3_PROD_BUCKET}/prod-build-${CI_COMMIT_TAG}.tar.gz

# 3. Trigger SSM Run Command on Utility Admin instance to unpack onto the shared EFS
- >
aws ssm send-command
--instance-ids "${SYNC_INSTANCE_PROD_ID}"
--document-name "AWS-RunShellScript"
--comment "Deploy production files to EFS from GitLab CI"
--parameters 'commands=[
"aws s3 cp "${S3_PROD_BUCKET}"'/prod-build-${CI_COMMIT_TAG}.tar.gz /tmp/prod-build.tar.gz",
"mkdir -p /mnt/efs/production/frontend_new",
"tar -xzf /tmp/prod-build.tar.gz -C /mnt/efs/production/frontend_new",
"rm -rf /mnt/efs/production/frontend_old",
"[ -d /mnt/efs/production/frontend ] && mv /mnt/efs/production/frontend /mnt/efs/production/frontend_old",
"mv /mnt/efs/production/frontend_new /mnt/efs/production/frontend",
"rm -rf /mnt/efs/production/frontend_old /tmp/prod-build.tar.gz",
"echo Production Shared EFS Deployment Successful"
]'
rules:
- if: $CI_COMMIT_TAG =~ /^v\d+\.\d+\.\d+$/

```

5. Architectural Evaluation & Better Alternatives

While deploying application code over EFS NFS mounts provides rapid deployment cycles and immediate propagation, it represents a **classic hybrid design pattern with notable operational tradeoffs**.

Below is an engineering analysis comparing EFS deployment to modern cloud-native standards.


5.1 Tradeoffs of the EFS NFS Code-Serving Approach

The Drawbacks:

1. **Network Latency Bottleneck:** Filesystem operations (stat, open, read) over network NFSv4 run significantly slower than on local NVMe/SSD EBS volumes. If an application (e.g., Python/Node.js) imports hundreds of source files at startup, load times are substantially degraded.
 2. **Atomic Swap Synchronization Risks:** If an instance executes active code while EFS is being overwritten, it can load partial files, causing runtime exceptions. (Our symlink atomic swap mitigated this, but node-level system file-locks can still lock operations).
 3. **Write Concurrency and POSIX File Locking:** Multiple instances running backend APIs attempting to write logs, cache files, or local SQLite data on the same mounted EFS directory will experience lock contention, causing thread blockages.
 4. **No Version Immutability (Risk of Corruption):** A single bad build or a malicious script executing on an ASG node can overwrite the shared code folder, instantaneously corrupting/compromising *all* ASG nodes and standalone instances simultaneously.
 5. **No Easy Atomic Rollback:** Rollbacks require manual filesystem restores or triggering an old S3 unpack command.
-

5.2 Proposed Solution 1: Containerization (Docker + Amazon ECS/EKS) — *Highly Recommended*


This is the standard modern cloud-native deployment model.



- **How it works:** GitLab CI/CD builds a standalone Docker Image containing Nginx and the code. It pushes this image to **Amazon Elastic Container Registry (ECR)**. It then triggers an ECS Service rolling update.
 - **Why it is superior:**
 - **Complete Isolation & Immutability:** Each container is an exact, unchangeable replica of the tested image.
 - **No Network Storage Bottleneck:** The container runs files off local virtual memory/EBS storage inside the host, resulting in sub-millisecond execution loops.
 - **Zero-Downtime Deployments:** ECS manages blue/green or rolling upgrades automatically, starting healthy new containers before terminating old ones.
 - **Instant rollback:** Rollback is as simple as reverting the active image tag to the previous version in ECR.
-

5.3 Proposed Solution 2: Stateless S3-Pull Bootstrapping with ASG Instance Refresh

If keeping EC2 instances is mandatory, decouple them from the EFS file system using local EBS drives hydrated from S3.



- **How it works:**
 1. GitLab CI/CD compiles the code, zips it, and uploads the immutable artifact to S3 (e.g., s3://bucket/deploy-v1.0.0.zip).
 2. The pipeline updates an SSM Parameter or updates the **Auto Scaling Group Launch Template** User Data to reference the new version.
 3. The pipeline triggers an **ASG Instance Refresh** (a controlled, automated rolling upgrade where AWS replaces old instances with new ones).
 4. During bootstrap, the EC2 instance pulls the zip file down from S3, decompresses it locally onto its high-speed EBS volume (/var/www/html/), and starts Nginx.
- **Why it is superior:**
 - **Performance:** Nginx serves files directly off local high-speed SSDs (EBS) without network NFS lookups.
 - **Reliability:** If one instance's file system is corrupted, the other instances in the ASG remain perfectly healthy.
 - **Immutability:** Each instance has its own isolated, stable copy of the code. If an auto-scaling event occurs, the newly launched instance pulls the identical versioned zip file.

Costing Estimate

#aws#cloud#architecture#vpc#alb#asg#rds#waf#elasticsearch#valkey#jumphost#bastion#disaster-recovery#efs#postgresql#ragflow#langfuse

#costing#bedrock#cognito#lambda#apigateway

Estimated Costing

We provide two distinct monthly cost estimates for the 3-Tier AWS Architecture in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)** utilizing On-Demand rates (~730 hours/month). All figures are aligned with our [Software Licensing & Technology Risk Register \(TS/MC Series\)](#) to ensure complete commercial and licensing auditability.

1. **Baseline Cost-Optimized Plan:** Ideal for initial development, testing, staging environments, and low-traffic applications.
2. **High-Performance Developer-Aligned Plan:** Spec'd specifically to fulfill the resource requirements of the **Developer's First Design (Nginx, Backend, RAGFlow, LangFuse)** with production-grade performance.

Both plans have been updated and calibrated against real-world, production-ready AWS billings from similar projects to incorporate critical infrastructure support services (such as **Amazon ElastiCache Valkey**, **Amazon EFS**, **AWS Secrets Manager**, **AWS Backup**, **Amazon CloudWatch**, standard **Public IPv4 address charges**, **Amazon Route 53 custom domain management**, and our **secure, hardened SSH Jumphost (Bastion)** whitelisted for the Cyberjaya developer office).

Additionally, this guide includes detailed cost models for **AWS-Native Alternatives** to external “extra” integrations (such as Amazon Bedrock instead of OpenAI, Amazon Cognito instead of self-hosted/SaaS Auth, AWS End User Messaging for WhatsApp, and serverless API Gateway/Lambda webhook routing) so that stakeholders have a complete financial blueprint of a 100% cloud-native architecture.

1. Baseline Cost-Optimized Plan (USD)

This configuration targets baseline usage with smaller resource profiles (t4g.medium compute, db.m6g.large database) to minimize initial expenditures.

Tier / AWS Component	Configuration & Resource Sizing	Hourly/Unit Rate	Est. Monthly Cost (USD)
Networking & NAT	1 NAT Gateway (Outbound API/WhatsApp Traffic)		
		\$0.045 / hr	\$32.85
Public IPv4 Addresses	• Base provisioned charge		
		\$0.045 / GB	\$2.25
Presentation Tier	• Estimated ~50 GB Data Processing		
	3 Public IPv4 addresses (1 for NAT Gateway, 2 for ALB)	\$0.005 / hr / IP	\$10.95
Security Tier	• Required public IP address allocation		
	1 Application Load Balancer (ALB)	\$0.0225 / hr	\$16.43
Compute Tier (ASG)	• Base LCF/hour charge		
		\$0.008 / LCU-hr	\$11.68
Database Tier (RDS)	• Estimated ~2 LCU average usage		
	AWS WAFv2 (Attached to ALB)	\$5.00 / month	\$5.00
Caching Tier	• 1 Regional Web ACL		
		\$1.00 / rule/mo	\$3.00
Storage Tier (S3)	• 3 Rules (OWASP Core, SQLi, Rate Limit)		
		\$0.60 / M-req	\$0.60
Monitoring Tier	• ~1 Million HTTP/HTTPS Requests		
	2x t4g.medium EC2 Instances (ARM Graviton)	\$0.0336 / hr / inst	\$49.06
Secrets Management	• 2 vCPU, 4GB RAM each		
		\$0.08 / GB-month	\$4.80
Networking & NAT	• 2x 30GB gp3 EBS Root Volumes		
	Multi-AZ db.m6g.large PostgreSQL [TS-06]	\$0.304 / hr	\$221.92
Public IPv4 Addresses	• 2 vCPU, 8GB RAM (High Availability)		
		\$0.23 / GB-month	\$11.50
Presentation Tier	• 50GB gp3 Storage		
	Amazon ElastiCache Valkey (cache.t4g.micro) [TS-06, TS-05]		
Security Tier	• 1 Node, 0.5 GB RAM (Session & metadata caching)		
		\$0.0128 / hr	\$9.34
Compute Tier (ASG)	• Valkey pricing is 20% lower than legacy Redis OSS		
	Amazon S3 (Encrypted Uploads & Media)	\$0.023 / GB-month	\$2.30
Database Tier (RDS)	• ~100 GB Standard Storage		
		Nominal rates	\$0.50
Caching Tier	• ~50,000 PUT/GET API requests		
	Amazon EFS (Elastic File System)	\$0.30 / GB-month	\$3.00
Storage Tier (S3)	• ~10 GB standard shared persistent storage		
	AWS Secrets Manager	\$0.40 / secret/mo	\$0.80
Monitoring Tier	• 2 Secrets (one for RDS DB, one for external API keys)		
	Amazon CloudWatch	Nominal rates	\$1.50

Tier / AWS Component	Configuration & Resource Sizing	Hourly/Unit Rate	Est. Monthly Cost (USD)
Disaster Recovery	3 alarms, custom CPU/Memory dashboard AWS Backup	\$0.05 / GB-month	\$5.00
	Automated Multi-AZ RDS snapshots & EBS backups (~100 GB) Standalone EC2 Instances (AMI Staging & Baking)		
Standalone EC2 Tier	3x Standalone EC2 Instances (one per ASG group: Frontend, Backend, AI Tier) connected directly to RDS, S3, or EFS to ensure 1:1 environment parity.	\$0.0084 / hr / inst	\$18.39
	Sized as 3x t4g.micro (1 vCPU, 1GB RAM each)	\$0.08 / GB-month	\$3.60
SSH Jumphost Tier	3x 15GB gp3 EBS Root Volumes (45GB total) Secure SSH Jumphost (Bastion)		
	1x t4g.micro EC2 Instance in the Public Subnet (accessible only from Cyberjaya office whitelisted IP)	\$0.0084 / hr	\$6.13
		\$0.08 / GB-mo	\$1.20
	15GB gp3 EBS Root Volume	\$0.005 / hr	\$3.65
Data Transfer	1 Static Elastic IP allocation Outbound Internet Data Transfer	Free Tier	\$0.00
	~100 GB Outbound (First 100GB Free/mo) Amazon Route 53 Custom Domain Routing		
Route 53 Domain	1 Public Hosted Zone	\$0.50 / zone / mo	\$0.50
	Estimated ~2 Million standard queries	\$0.40 / M-req	\$0.80
TOTAL ESTIMATED MONTHLY COST			~\$426.75 USD / month
Local Currency Equivalent (MYR): ~RM 1,920 MYR / month (calculated at an exchange rate baseline of 1 USD ≈ 4.50 MYR).			

2. High-Performance Developer-Aligned Plan (USD)

This configuration scales up computing instances to precisely match the developer’s server specifications (**Server 02, Server 03, and Server 04**) with high performance, dedicated throughput, and larger memory margins.

Tier / AWS Component	Configuration & Resource Sizing	Hourly/Unit Rate	Est. Monthly Cost (USD)
Networking & NAT	1 NAT Gateway (Outbound API/WhatsApp Traffic)	\$0.045 / hr	\$32.85
	Base provisioned charge	\$0.045 / GB	\$2.25
Public IPv4 Addresses	Estimated ~50 GB Data Processing 3 Public IPv4 addresses (1 for NAT Gateway, 2 for ALB)	\$0.005 / hr / IP	\$10.95
	Required public IP address allocation 1 Application Load Balancer (ALB)		
Presentation Tier	Base LCF/hour charge	\$0.0225 / hr	\$16.43
	Estimated ~2 LCU average usage AWS WAFv2 (Attached to ALB)	\$0.008 / LCU-hr	\$11.68
Security Tier	1 Regional Web ACL	\$5.00 / month	\$5.00
	3 Rules (OWASP Core, SQLi, Rate Limit)	\$1.00 / rule/mo	\$3.00
	~1 Million HTTP/HTTPS Requests 2x t4g.xlarge EC2 Instances (ARM Graviton)	\$0.60 / M-req	\$0.60
	4 vCPU, 16GB RAM each (Supports Backend, DMS, RAGFlow, LangFuse)	\$0.1344 / hr / inst	\$196.22
Compute Tier (ASG)	2x 30GB gp3 EBS Root Volumes Multi-AZ db.m6g.xlarge PostgreSQL [TS-06]	\$0.08 / GB-month	\$4.80
	4 vCPU, 16GB RAM (Matches Server 04 Data Tier)	\$0.608 / hr	\$443.84
Database Tier (RDS)	50GB gp3 Multi-AZ Storage Amazon ElastiCache Valkey (cache.t4g.medium) [TS-06, TS-05]	\$0.46 / GB-month	\$23.00
	1 Node, 3.09 GB RAM (Production cache & task broker) - Valkey pricing is 20% lower than legacy Redis OSS Amazon S3 (Encrypted Uploads & Media)		
Caching Tier		\$0.0544 / hr	\$39.71
Storage Tier (S3)	~100 GB Standard Storage	\$0.023 / GB-month	\$2.30
	~50,000 PUT/GET API requests	Nominal rates	\$0.50

Tier / AWS Component	Configuration & Resource Sizing	Hourly/Unit Rate	Est. Monthly Cost (USD)
Storage Tier (EFS)	Amazon EFS (Elastic File System)		
	• ~50 GB shared network storage for AI model weights / caches	\$0.30 / GB-month	\$15.00
Secrets Management	AWS Secrets Manager		
	• 5 Secrets (RDS, LLM API keys, external integrations, LangFuse, WAF keys)	\$0.40 / secret/mo	\$2.00
Monitoring Tier	Amazon CloudWatch		
	• Logs ingestion (~5 GB), dashboards, custom metric triggers	Nominal rates	\$5.00
Disaster Recovery	AWS Backup		
	• Centralized backup for RDS, EFS, and ASG EBS volumes (~150 GB)	\$0.05 / GB-month	\$7.50
Standalone EC2 Tier	Standalone EC2 Instances (AMI Staging & Baking)		
	• 3x Standalone EC2 Instances (one per ASG group: Frontend, Backend, AI Tier) connected directly to RDS, S3, or EFS to ensure 1:1 environment parity.	Frontend: \$0.0336 / hr	\$24.53
	• 1x Frontend: t4g.medium (2 vCPU, 4GB RAM) + 30GB gp3	Backend/AI: \$0.1344 / hr	\$196.22
	• 1x Backend: t4g.xlarge (4 vCPU, 16GB RAM) + 30GB gp3	Storage: \$0.08 / GB-mo	\$8.80
SSH Jumphost Tier	• 1x AI Tier: t4g.xlarge (4 vCPU, 16GB RAM) + 50GB gp3 (110GB gp3 total)		
	Secure SSH Jumphost (Bastion)		
	• 1x t4g.micro EC2 Instance in the Public Subnet (accessible only from Cyberjaya office whitelisted IP)	\$0.0084 / hr	\$6.13
	• 15GB gp3 EBS Root Volume	\$0.08 / GB-mo	\$1.20
Data Transfer	• 1 Static Elastic IP allocation	\$0.005 / hr	\$3.65
	Outbound Internet Data Transfer		
	• ~100 GB Outbound (First 100GB Free/mo)	Free Tier	\$0.00
Route 53 Domain	Amazon Route 53 Custom Domain Routing		
	• 1 Public Hosted Zone	\$0.50 / zone / mo	\$0.50
	• Estimated ~2 Million standard queries	\$0.40 / M-req	\$0.80
TOTAL ESTIMATED MONTHLY COST			~\$1,064.46 USD / month
• Local Currency Equivalent (MYR): ~RM 4,790 MYR / month (calculated at an exchange rate baseline of 1 USD ≈ 4.50 MYR).			

3. Real-World Cost Calibration & Analysis

A comparison with real-world billings from a highly similar production deployment (with a stable run-rate of **\$659.10** for a representative billing cycle, e.g., May 2026) validates our cost modeling and exposes key areas where our estimations are more accurate, comprehensive, and cost-efficient:

AWS Service / Category	Similar Project (Screenshot)	Our Project Alignment & Analysis
Relational Database Service	\$308.51	Our High-Performance plan utilizes Multi-AZ db.m6g.xlarge (\$466.84). The similar project likely runs a db.m6g.large Multi-AZ with larger storage or a single-AZ instance, which we can optimize further.
EC2-Instances	\$180.55	Matches our 2x t4g.xlarge estimate (\$196.22) very closely, suggesting the similar project also runs high-performance compute. Difference may be Savings Plans or a slightly lower-cost regional selection.
ElastiCache (Valkey / Redis)	\$76.78	Essential for RAGFlow/LangFuse caching. A \$76.78 spend corresponds to a single cache.m6g.large or Multi-AZ cache.t4g.medium under Redis OSS. By migrating to Amazon ElastiCache for Valkey, we achieve 20% lower on-demand rates (e.g., \$39.71/mo for cache.t4g.medium), saving significantly compared to the Redis OSS baseline.
EC2-Other	\$45.85	This bundles EBS root volumes (\$4.80), NAT Gateway hourly charge (\$32.85), and data processing (\$2.25), confirming our highly granular breakdown is accurate.
Elastic Load Balancing	\$16.97	Matches our ALB base charge (\$16.43/mo) with nominal traffic LCU charges of ~\$0.54.
VPC (Public IP Charges)	\$14.90	Reflects standard public IPv4 address fees (\$0.005/hr) for our 1 NAT Gateway Elastic IP + 2 ALB IPs (\$10.95), plus nominal VPC Flow Log storage.
AWS WAF	\$14.06	Aligns perfectly with our Regional Web ACL baseline (\$5.00/mo) + 3 Rules (\$3.00/mo) + request volume.
Route 53	\$1.30	1 Hosted Zone (\$0.50) and ~2M queries (\$0.80) to route custom domain traffic to ALB dynamically.
Elastic File System	\$0.54	Confirms EFS usage is minimal (configuration/caches). We spec EFS at \$3.00 (10GB) and \$15.00 (50GB) to support RAGFlow pre-trained AI model caching across the ASG.

CloudWatch, Secrets Mgr, Backup	~\$1.00	These operational services (alarming, secrets, and disaster recovery) are highly critical. We model them at a realistic \$7.30 - \$14.50 combined to prevent bill surprises in production.
TOTAL MONTHLY COST	\$659.10	Calibration proves our updated models (\$426.75 and \$1,064.46) are exceptionally robust and production-true.

3.1 Reference Telemetry Data for Price & Load Estimation: AWS Region Malaysia (ap-southeast-5)

To ground our financial forecasting in realistic pricing and workload profiles, the following tables detail AWS Cost Explorer reference data from an external benchmark project in the **AWS Malaysia (ap-southeast-5) region** spanning a 12-month lifecycle (September 2025 to August 2026). **Note:** This dataset is utilized strictly as a reference baseline to perform price and load estimations for our 3-tier architecture to ensure resource sizing stays strictly within project budget guardrails.

A. 12-Month Historical Service Spend (Sept 2025 – Aug 2026)

Across the 12-month evaluation window, total regional cloud expenditure reached **\$61,400.47 USD** (≈ **RM 276,302.12 MYR** at 1 USD = 4.50 MYR).

AWS Service / Line Item	12-Month Cumulative Spend (USD)	Local Currency Equivalent (MYR)	Percentage of Total Spend
Amazon ElastiCache (Valkey / Redis)	\$13,991.93	RM 62,963.70	22.79%
Amazon EC2 (Compute Instances)	\$13,315.42	RM 59,919.39	21.69%
Amazon Elastic File System (EFS)	\$10,807.09	RM 48,631.91	17.60%
Amazon Relational Database Service (RDS)	\$10,295.35	RM 46,329.06	16.77%
Amazon EC2-Other (EBS, NAT Gateway, Data Transfer)	\$7,009.31	RM 31,541.88	11.42%
Elastic Load Balancing (ALB)	\$3,555.55	RM 15,999.96	5.79%
Amazon VPC (Public IPv4 & Flow Logs)	\$917.73	RM 4,129.77	1.49%
Amazon CloudWatch	\$890.24	RM 4,006.09	1.45%
AWS WAF (Web Application Firewall)	\$734.44	RM 3,304.98	1.20%
AWS Backup	\$56.42	RM 253.88	0.09%
AWS Secrets Manager	\$34.79	RM 156.56	0.06%
Amazon S3	\$15.93	RM 71.71	0.03%
Amazon Route 53	\$11.08	RM 49.86	0.02%
AWS Key Management Service (KMS)	\$11.08	RM 49.85	0.02%
AWS Cost Explorer	\$4.09	RM 18.41	0.01%
Data Transfer Adjustments / Credits	-\$249.98	-RM 1,124.92	-0.41%
GRAND TOTAL (12-MONTH PERIOD)	\$61,400.47 USD	RM 276,302.12 MYR	100.00%

B. Monthly Expenditure Progression (Sept 2025 – Aug 2026)

The monthly spending curve illustrates the operational progression from initial sandbox setup to high-load scaling:

Billing Month	Monthly Spend (USD)	Monthly Spend (MYR)	Key Architectural Milestones
2025-09-01	\$2,805.25	RM 12,623.64	Initial VPC deployment, single-AZ RDS & Valkey provisioning
2025-10-01	\$4,029.27	RM 18,131.72	Addition of EFS shared storage and ASG compute nodes
2025-11-01	\$4,433.64	RM 19,951.38	Multi-AZ database setup and WAF rules integration
2025-12-01	\$4,802.65	RM 21,611.95	Staging environment pre-load testing and vector index baking
2026-01-01	\$8,940.71	RM 40,233.18	Peak load-test execution & high-memory caching expansion (cache.r6g.2xlarge)
2026-02-01	\$6,726.76	RM 30,270.42	Post-load-test stabilization and capacity tuning
2026-03-01	\$6,483.23	RM 29,174.52	Production preparation & Langfuse observability pipeline active
2026-04-01	\$5,988.17	RM 26,946.75	Steady-state production run-rate with t4g/c8g instance mix
2026-05-01	\$6,015.75	RM 27,070.86	Continuous production operation and EFS model weight caching
2026-06-01	\$6,113.60	RM 27,511.18	Full-scale enterprise AI workload execution (RAGFlow + pgvector)
2026-07-01	\$527.63	RM 2,374.35	Mid-year optimization, maintenance, and credit adjustment cycle
2026-08-01	\$4,533.82	RM 20,402.18	Optimized post-audit production baseline

C. August 2026 Instance Type Distribution

An analysis of instance-type cost allocations for August 2026 (**\$6,505.72 USD** gross compute & serverless total) demonstrates the exact hardware distribution in ap-southeast-5:

Instance Type / Component	Allocation Category	Monthly Cost (USD)	Monthly Cost (MYR)
No instance type (Serverless, S3, ALB, EBS, NAT)	Managed Services	\$2,502.50	RM 11,261.25
cache.r6g.2xlarge	High-Memory ElastiCache Valkey	\$1,494.55	RM 6,725.46
c8g.large	Graviton4 Compute (App & RAG Workers)	\$1,009.54	RM 4,542.95
db.m7g.xlarge	Graviton3 RDS PostgreSQL (Multi-AZ Primary)	\$592.22	RM 2,665.01
db.m7g.large	Graviton3 RDS PostgreSQL (Read Replica)	\$296.11	RM 1,332.50
c6g.2xlarge	Graviton2 Compute (High-Throughput Processing)	\$198.11	RM 891.48
c8g.xlarge	Graviton4 Compute (AI Vector Indexing)	\$178.98	RM 805.43
cache.t4g.medium	Graviton2 ElastiCache Valkey (Session Caching)	\$51.19	RM 230.34
c6g.large	Graviton2 Compute (Backend API Services)	\$49.51	RM 222.78
c6g.medium	Graviton2 Compute (Microservices)	\$49.45	RM 222.54
t4g.medium	Graviton2 Burstable Compute (Utility Workers)	\$28.42	RM 127.89
t3.micro	x86 Burstable Compute (Utility Services)	\$26.56	RM 119.52

Instance Type / Component	Allocation Category	Monthly Cost (USD)	Monthly Cost (MYR)
t3.small	x86 Burstable Compute (Staging Utility)	\$17.71	RM 79.68
t2.nano	x86 Legacy Micro Nodes	\$10.86	RM 48.88
TOTAL (AUGUST 2026 INSTANCE DISTRIBUTION)		\$6,505.72 USD	RM 29,275.74 MYR

D. June 2026 Daily Run-Rate Audit (\$19,174.98 USD June Aggregate Telemetry)

Granular daily tracking across June 2026 shows extreme operational consistency across core infrastructure components:

- **EC2 Compute Instances:** \$50.55 – \$50.77 / day (Average ~\$50.68 / day)
- **ElastiCache Valkey:** \$49.90 / day (Fixed multi-node cluster pricing)
- **RDS PostgreSQL:** \$32.08 / day (Fixed Multi-AZ instance charge)
- **EC2-Other (EBS, NAT Gateway, Data Transfer):** \$22.28 – \$24.28 / day
- **Elastic Load Balancing (ALB):** \$7.95 – \$15.82 / day (Varies dynamically with traffic LCU surges)
- **Elastic File System (EFS):** \$19.40 – \$42.02 / day (Scales with temporary AI vector cache growth)
- **CloudWatch & VPC:** ~\$2.65 / day and ~\$2.72 / day respectively

Note on Basis and Scope:

- **June 2026 Dataset (Table D & June 2026 Monthly Line):** Represents a 30-day aggregate gross accrual dataset (\$19,174.98 USD cumulative daily total; \$6,113.60 USD net monthly bill line after account-level volume discount allocations) for active multi-node load testing across the account scope.
- **August 2026 Dataset (Table C & August 2026 Monthly Line):** Represents an active production instance allocation dataset (\$6,505.72 USD gross compute allocation total across hardware families; \$4,533.82 USD net monthly bill line after post-audit credit adjustments of ~\$249.98 USD and account-level discounts).
- **12-Month Dataset (Table A):** Displays the cumulative 12-month net service total (\$61,400.47 USD) calculated across all 12 monthly net billing periods.

This reference telemetry serves as an external benchmark baseline to perform price and load estimations for our 3-tier architecture, ensuring that component choices and resource configurations remain within project budget parameters.

4. Plan Comparison Summary

Metric	Baseline Plan	High-Performance Plan
Target Environment	Staging, Dev, Testing	Production AI Workloads
Compute Spec (per node)	2 vCPU, 4GB RAM	4 vCPU, 16GB RAM
Database Spec (RDS)	2 vCPU, 8GB RAM	4 vCPU, 16GB RAM
Caching Spec (Valkey)	0.5 GB RAM	3.09 GB RAM
Shared storage (EFS)	10 GB (Configs/Logs)	50 GB (AI model caches)
Standalone EC2 (AMI Baking)	3x `t4g.micro`	1x `t4g.med`, 2x `x1rg`
SSH Jumphost (Bastion)	1x `t4g.micro` (Ubuntu)	1x `t4g.micro` (Ubuntu)
Route 53 Domain Setup	1 Hosted Zone + ~2M req	1 Hosted Zone + ~2M req
Monthly Estimate (USD)	~\$426.75 USD	~\$1,064.46 USD
Monthly Estimate (MYR)	~RM 1,920 MYR	~RM 4,790 MYR

5. Optional Cost Optimization Pathways (Day 2 Operations)

- **RDS Savings Plans / Reserved Instances (1-Yr / 3-Yr):** Committing to the primary PostgreSQL instance can reduce DB compute costs by **30%–35%**, cutting monthly spend by ~\$70 - \$150 USD depending on the plan.
- **EC2 Compute Savings Plans:** Committing to baseline t4g usage via Savings Plans reduces application compute charges by up to **20%–25%**.
- **ElastiCache Valkey Reserved Nodes:** Committing to caching nodes can shave up to **35%** off Valkey Cache costs (saving up to ~\$13.90 USD / month on cache.t4g.medium).
- **VPC S3 Gateway Endpoint:** S3 traffic routed through a free VPC Gateway Endpoint eliminates NAT Gateway data processing fees (\$0.045/GB) for media uploads.
- **EFS Lifecycle Management:** Transitioning EFS data to Infrequent Access (IA) or Archive tier after 14/30 days reduces the EFS storage unit cost from **\$0.30/GB** to **\$0.013/GB**, saving up to 90% of EFS cost for older model files.

6. Optional Enterprise Integration Add-ons (AWS Alternatives)

The developer’s technology stack includes “extra” third-party SaaS integrations (such as Twilio for WhatsApp, Meta Graph APIs, and OpenAI models) which are not part of the core infrastructure costing listed above.

To achieve maximum network security, data sovereignty, and consolidated billing, we estimate the cost of migrating these external integrations to **high-fidelity AWS-native alternatives** in ap-southeast-5 (assuming 1 USD = 4.50 MYR).

Granular Cost Estimations for AWS Alternatives

AWS Alternative Service	Sizing & Active Monthly Volume	Hourly/Unit Rate	Est. Monthly Cost (USD)	Est. Monthly Cost (MYR)
Amazon Bedrock [MC-01, MC-02] (Alternative to OpenAI API)	• Claude 3.5 Sonnet / Qwen3 (Chat & Reasoning) - 1 Million Input Tokens / mo - 500,000 Output Tokens / mo	Chat Input: \$0.003 / 1k tokens Output: \$0.015 / 1k tokens Embed: \$0.0001 / 1k tokens	\$11.00	~RM 49.50
			• Claude Input: \$3.00 • Claude Output:	

AWS Alternative Service	Sizing & Active Monthly Volume	Hourly/Unit Rate	Est. Monthly Cost (USD)	Est. Monthly Cost (MYR)
Amazon Cognito User Pools [TS-05] (Alternative to Auth0 / Self-Hosted Auth)	• Qwen3/Cohere Embed Multilingual - 5 Million embedding tokens / mo		\$7.50	
	• 15,000 Monthly Active Users (MAUs)		• Cohere Embed: \$0.50	
	• First 10,000 MAUs: Free	\$0.00 (First 10k MAUs)	\$27.50	~RM 123.75
	• Standard User Pools active traffic	\$0.0055 / MAU (Next 5k MAUs)		
AWS End User Messaging (Alternative to Twilio for WhatsApp)	• WhatsApp Social Channel			
	• 2,000 Active Conversations / mo	\$0.00 (First 1k convos)		
	• First 1,000 Conversations: Free	\$0.0246 / service conversation (Next 1k convos)	\$24.60	~RM 110.70
	• Standard service conversations in Malaysia			
API Gateway & AWS Lambda (Alternative to Java Webhook Receivers)	• Serverless Webhook Routing		\$3.70	
	• 1 Million API webhook triggers / mo	\$3.50 / M-requests	• API Gateway: \$3.50	~RM 16.65
	• API Gateway REST API	\$0.20 / M-invocations (plus free tier allowance)	• Lambda requests: \$0.20	
	• AWS Lambda (512MB RAM, 200ms) - 1 Million executions / mo (Free Tier)		• Lambda compute: \$0.00	
ADD-ON COMBINED MONTHLY TOTAL	100% cloud-native SaaS mapping		\$66.80 / month	~RM 300.60 / month

Why the AWS Alternatives Save on Operating Costs (OpEx)

- 1. **Elimination of Twilio Markup Fees:** Twilio charges an average platform markup of **\$0.005 per message** on top of Meta’s base carrier conversation fees. By deploying **AWS End User Messaging (Social Channels)**, developers connect directly to the Meta API, saving approx. **\$10 to \$30 USD / month** in markup fees for every 10,000 sent messages.
- 2. **Serverless Scaling to Zero (Webhooks):** Placing webhook receivers directly in Spring Boot (on dedicated virtual machines or ASGs) requires continuous provisioning of compute memory to avoid thread starvation during bursty Meta communication events. An **API Gateway + AWS Lambda** webhook layer costs **\$0.00** when inactive and automatically scales up to absorb millions of requests, preventing costly ASG scaling triggers.
- 3. **Cognito Free Tier Advantage:** Third-party identity providers (such as Auth0 or Okta) charge steep monthly premiums (starting at \$120+ USD/mo) once custom database connections are integrated. Amazon Cognito provides a robust, enterprise-grade directory with **10,000 MAUs entirely free every month**, lowering authentication costs significantly.
- 4. **Data Sovereignty Compliance:** By keeping AI model queries and document embeddings within **Amazon Bedrock**, companies avoid transferring sensitive corporate PDFs and customer PII across third-party OpenAI endpoints over the public internet. This simplifies **Transfer Impact Assessments (TIAs)** and aligns seamlessly with local compliance standards under the **Malaysian PDPA (Personal Data Protection Act) 2010**.

Standalone Wazuh SIEM Cost Optimization [TS-04]

By replacing high-cost vendor-managed SIEM platforms with our self-hosted, standalone Wazuh SIEM deployment (Option B under [docs/wazuh.md](#)), we achieve significant savings:

- **Baseline Cost:** **\$65.71 USD / RM 295.70 MYR per month** on an ARM64 Graviton t4g.large instance in ap-southeast-5.
- **Financial Benefit:** This replaces the high-cost commercial Security SIEM budget lines, resulting in a **57%+ monthly OpEx reduction** compared to standard x86 commercial marketplace solutions (\$153.16/mo).

7. AWS 3-Tier AI Infrastructure Cost Audit & Verification Report

Target Region: AWS Malaysia (ap-southeast-5 - Kuala Lumpur)

Workload Type: Enterprise RAG & AI Infrastructure (RAGFlow, Langfuse, Valkey, PostgreSQL + pgvector)

Currency: USD (\$) / MYR (RM) — Estimated Exchange Rate: \$1.00 = MYR 4.50\$

1. Executive Summary

This report presents a thorough counter-check and financial verification of the AWS 3-Tier Deployment Architecture for AI Infrastructure. The proposed architecture supports containerized enterprise AI workloads (such as RAGFlow, Langfuse, and vector processing microservices) utilizing high-efficiency AWS Graviton3/Graviton4 compute instances, managed relational databases (Amazon RDS PostgreSQL with pgvector), high-speed caching (Amazon ElastiCache Valkey/Redis), and multi-AZ network isolation.

Key Audit Findings


- **Baseline Accuracy:** The core compute and database pricing aligns with AWS Malaysia (ap-southeast-5) regional standards, showing a ~10-15% cost efficiency advantage when utilizing ARM-based Graviton instances (t4g, c7g, m7g, r7g) compared to x86 equivalents.
- **Hidden Cost Traps Identified:** Standard cloud estimates often omit non-compute operational fees. Crucial hidden cost drivers identified during this audit include:
 - **AWS Public IPv4 Address Surcharge:** \$0.005 / hour (≈ \$3.65 / month per public IPv4 assigned to ALBs, NAT Gateways, and Jumphosts).
 - **NAT Gateway Processing Surcharge:** Hourly gateway charges (\$0.045 - \$0.05/hour per AZ) plus data processing costs (\$0.045 / GB).
 - **Application Load Balancer (ALB) Capacity Units (LCUs):** Base hourly rate plus rule/new-connection LCU scaling under LLM streaming payloads.
 - **Cross-AZ Data Transfer:** Inter-AZ compute-to-database and compute-to-cache data transfer (\$0.01 / GB each direction).

Financial Impact Summary

- **Dev / POC Tier:** ≈ \$138.50 / month (MYR 623.25 / month)
- **Staging Tier:** ≈ \$482.10 / month (MYR 2,169.45 / month)
- **Enterprise Production (Multi-AZ):** ≈ \$1,285.80 / month (MYR 5,786.10 / month)
- **Optimized Prod (1-Year Savings Plan):** ≈ \$945.30 / month (MYR 4,253.85 / month) — 26.5% cost reduction.

2. Architectural Tiering & Infrastructure Profile

The architecture consists of three distinct tiers deployed within a secure VPC across 3 Availability Zones in ap-southeast-5:



3. Comprehensive Line-Item Cost Counter-Check (ap-southeast-5)

3.1 Tier 1: Dev / POC Environment (Single-AZ / Cost-Optimized)

Designed for low-cost verification, initial RAG model testing, and pipeline integration.

Component	Resource Specification	Qty / Usage	Unit Cost (USD)	Monthly Cost (USD)	Monthly Cost (MYR)
Compute (App)	EC2 t4g.medium (2 vCPU, 4GB RAM)	1 instance (730h)	\$0.0336 / hr	\$24.53	MYR 110.39
Jump host	EC2 t4g.micro (Burstable)	1 instance (730h)	\$0.0084 / hr	\$6.13	MYR 27.59
Storage (EBS)	gp3 Volume (App + Jump host)	50 GB Total	\$0.08 / GB-mo	\$4.00	MYR 18.00
Database	RDS PostgreSQL db.t4g.medium (Single-AZ)	1 instance (730h)	\$0.065 / hr	\$47.45	MYR 213.53
RDS Storage	gp3 Storage (Database)	30 GB	\$0.115 / GB-mo	\$3.45	MYR 15.53
Cache	ElastiCache Valkey/Redis cache.t4g.micro	1 node (730h)	\$0.016 / hr	\$11.68	MYR 52.56
Networking	Single NAT Gateway (Shared Dev)	1 NAT (730h)	\$0.045 / hr	\$32.85	MYR 147.83
Public IPv4	Public IP Fees (Jump host + NAT)	2 Public IPs	\$0.005 / IP-hr	\$7.30	MYR 32.85
Data Processing	NAT Gateway Data Processed	25 GB / mo	\$0.045 / GB	\$1.11	MYR 5.00
TOTAL (Dev)				\$138.50	MYR 623.25

3.2 Tier 2: Staging Environment (Dual-AZ / Pre-Production)

Mirrors production topology with reduced instance scaling to validate multi-AZ failover, streaming telemetry, and deployment automation.

Component	Resource Specification	Qty / Usage	Unit Cost (USD)	Monthly Cost (USD)	Monthly Cost (MYR)
Ingress Load Balancer	Application Load Balancer (ALB)	1 ALB (730h)	\$0.0225 / hr	\$16.43	MYR 73.94
ALB LCU Usage	Load Balancer Capacity Units	2 LCUs average	\$0.008 / LCU-hr	\$11.68	MYR 52.56
WAF Protection	AWS WAF WebACL + 2 Core Rulesets	1 WebACL	\$5.00 + \$2.00 rules	\$7.00	MYR 31.50
Compute Tier	EC2 c7g.xlarge (4 vCPU, 8GB RAM)	2 instances (ASG)	\$0.145 / hr	\$211.70	MYR 952.65
Jump host	EC2 t4g.small (Session Manager)	1 instance (730h)	\$0.0168 / hr	\$12.26	MYR 55.17
Database Tier	RDS PostgreSQL db.t4g.large (Multi-AZ)	Multi-AZ (730h)	\$0.258 / hr	\$188.34	MYR 847.53
RDS Storage	gp3 Storage (Multi-AZ)	100 GB	\$0.23 / GB-mo	\$23.00	MYR 103.50
Cache Tier	ElastiCache cache.t4g.medium (Dual Node)	2 nodes (Multi-AZ)	\$0.065 / hr	\$94.90	MYR 427.05
Shared Storage	AWS EFS (General Purpose - Standard)	20 GB storage	\$0.30 / GB-mo	\$6.00	MYR 27.00
Networking	Dual-AZ NAT Gateways	2 NATs (730h)	\$0.045 / hr	\$65.70	MYR 295.65
Public IPv4	Public IPs (ALB + NAT + Jump host)	4 Public IPs	\$0.005 / IP-hr	\$14.60	MYR 65.70
Data Processing	NAT + Inter-AZ Traffic Transfer	300 GB	Mixed rates	\$16.50	MYR 74.25
TOTAL (Staging)				\$668.11	MYR 3,006.50

3.3 Tier 3: Enterprise Production Environment (High-Availability Multi-AZ)

Engineered for high-throughput RAG document parsing, vector indexing, Langfuse telemetry tracing, and 99.99% operational uptime.

Component	Resource Specification	Qty / Usage	Unit Cost (USD)	Monthly Cost (USD)	Monthly Cost (MYR)
Ingress Load Balancer	ALB (Multi-AZ Ingress)	1 ALB (730h)	\$0.0225 / hr	\$16.43	MYR 73.94
ALB LCU Scaling	Streaming / High Request LCUs	5 LCUs average	\$0.008 / LCU-hr	\$29.20	MYR 131.40
AWS WAF	WAF WebACL + Managed Rules + Requests	5M requests/mo	Managed baseline	\$22.50	MYR 101.25
Compute (ASG)	EC2 c7g.2xlarge (8 vCPU, 16GB RAM)	3 instances (Min 3)	\$0.290 / hr	\$635.10	MYR 2,857.95
Database Tier	RDS PostgreSQL db.m7g.xlarge (Multi-AZ)	4 vCPU, 16GB RAM	\$0.674 / hr	\$492.02	MYR 2,214.09
RDS Storage	Provisioned IOPS gp3 (3,000 IOPS / 125 MB/s)	250 GB Storage	Multi-AZ Storage rate	\$57.50	MYR 258.75
Cache Tier	ElastiCache Valkey cache.r7g.large (Multi-AZ)	2 nodes + Failover	\$0.136 / hr	\$198.56	MYR 893.52

Component	Resource Specification	Qty / Usage	Unit Cost (USD)	Monthly Cost (USD)	Monthly Cost (MYR)
Shared Storage	Amazon EFS (RAG Artifacts / Models)	100 GB Storage	\$0.30 / GB-mo	\$30.00	MYR 135.00
Networking	Triple-AZ NAT Gateways (High HA)	3 NATs (730h)	\$0.045 / hr	\$98.55	MYR 443.48
Public IPv4	Public IPs (ALB, 3 NATs)	5 Public IPs	\$0.005 / IP-hr	\$18.25	MYR 82.13
Data Processing	NAT Gateway & Inter-AZ Traffic	1,000 GB processed	Combined rates	\$52.50	MYR 236.25
DNS & Monitoring	Route 53 Health Checks + CloudWatch Logs	Failover + Logs	Baseline	\$15.00	MYR 67.50
TOTAL (Prod)				\$1,665.61	MYR 7,495.25

4. In-Depth Analysis of Hidden AWS Cost Drivers

Standard cloud estimators frequently understate monthly expenditure by focusing solely on core compute and database hourly rates. The following audit highlights crucial non-compute overheads:

CRITICAL HIDDEN COST DRIVERS AUDIT
1. Public IPv4 Charges (\$0.005/hr per IP) - App Load Balancers (2 IPs): \$7.30/month - Multi-AZ NAT Gateways (3 IPs): \$10.95/month - Bastion / Jumphosts (1 IP): \$3.65/month Total IPv4 Overhead: \$21.90/month (MYR 98.55)
2. NAT Gateway Hourly + Processing Fees - Fixed Hourly Charge (3 AZs @ \$0.045/hr): \$98.55/month - Data Processing (1,000 GB @ \$0.045/GB): \$45.00/month Total NAT Overhead: \$143.55/month (MYR 645.98)
3. Inter-AZ Data Transfer (Cross-AZ Traffic) - App Compute to Multi-AZ RDS & ElastiCache: \$0.01/GB inbound + outbound Total Transfer Overhead (1,000 GB): \$20.00/month (MYR 90.00)

Mathematical Formula for NAT Gateway Costing:

Cost_NAT = (N_AZ × 730 hours × \$0.045) + (Data_Processed (GB) × \$0.045)

For 3 Availability Zones handling 1,000 GB of monthly ingestion traffic:

Cost_NAT = (3 × 730 × 0.045) + (1000 × 0.045) = \$98.55 + \$45.00 = \$143.55 / month

5. Cost Optimization Strategy & 3-Year TCO Analysis

To maximize financial efficiency without sacrificing high availability or performance, a three-phase optimization roadmap is recommended:

5.1 Optimization Levers

- 1-Year or 3-Year Compute Savings Plans:** Applying a 1-Year All Upfront Compute Savings Plan to EC2 compute (c7g.2xlarge) yields an average 34% discount. Applying a 1-Year Reserved Instance (RI) to RDS PostgreSQL Multi-AZ yields a 30% discount.
- NAT Gateway Consolidation for Staging/Dev:** In non-production environments, consolidate from multi-AZ NAT Gateways to a single NAT Gateway, reducing hourly gateway fees by 66%.
- AWS Systems Manager (SSM) Session Manager:** Eliminates the necessity for public IPv4-backed Bastion EC2 instances, saving instance hourly costs, EBS volumes, and public IP surcharges.
- Valkey Engine Adoption over Redis:** AWS ElastiCache for Valkey offers a 20% price reduction over traditional ElastiCache for Redis while remaining fully open-source and wire-compatible.

5.2 3-Year Cost Comparison (Enterprise Production)

- Unoptimized On-Demand (3 Years):** \$1,665.61 × 36 = **\$59,961.96** (MYR 269,828.82)
- Optimized (1-Yr Savings Plans + Valkey):** \$1,180.20 × 36 = **\$42,487.20** (MYR 191,192.40)
- Optimized (3-Yr Compute Savings Plans):** \$945.30 × 36 = **\$34,030.80** (MYR 153,138.60)

Financial Impact: Implementing 3-Year Compute Savings Plans and RDS Reserved Instances delivers a Total Savings of \$25,931.16 (MYR 116,690.22) over 36 months, representing a 43.2% reduction in total cloud expenditure.

6. Audit Summary & Strategic Recommendations

- Adopt Graviton3/Graviton4 Architectures as Default:** Standardize compute images (c7g, t4g, m7g) across ASG nodes and RDS. Graviton instances deliver up to 40% better price-performance for Python/Go microservices and PostgreSQL vector queries compared to x86 instances.
- Implement PrivateLink S3 Endpoints:** Route all large document ingestion traffic from RAGFlow directly to Amazon S3 via free Gateway VPC Endpoints, bypassing NAT Gateway data processing fees (\$0.045/text{GB}\$).
- Enforce Automated Non-Prod Shutdown Schedule:** Utilize AWS Instance Scheduler to auto-stop Dev and Staging ASG/EC2 nodes outside business hours (12h/day, 5 days/week), reducing non-production compute costs by an additional 64%.
- Establish Budget Guardrails:** Configure AWS Budgets with automated anomaly detection alerts set at 80% and 100% thresholds of expected monthly spending to prevent runaway LCU or NAT charges.

FINAL AUDIT VERDICT
Baseline Infrastructure Design: APPROVED Cost Estimates Accuracy: VERIFIED & ADJUSTED FOR AWS MALAYSIA (ap-southeast-5) Recommended Deployment Strategy: Commit to 1-Yr SP for Prod; Single-NAT for Dev

