

All Project Documentation

OpenTofu >= 1.6.0 AWS: ap-southeast-5 WAFv2: Enabled Graviton: ARM64

GUIDE

AWS Secure 3-Tier Architecture Documentation



AWS Secure 3-Tier Architecture Documentation

Welcome to the official technical documentation for our **AWS 3-Tier Deployment for AI & Web Infra** project. This project is optimized for deployment in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)** utilizing AWS Graviton (ARM64) instances, Multi-AZ RDS Postgres, and AWS WAFv2 regional protection.

This deployment is structured natively in OpenTofu, adhering to strict modular boundaries, security best practices, and the **Zero-Trust Network Principle**.

Technical Overview

The architecture divides infrastructure components into discrete logical and physical tiers to achieve top-tier scalability, performance, and threat mitigation:

- **Presentation / Web Layer:** Application Load Balancer (ALB) receiving external traffic and filtered through AWS WAFv2 (OWASP rules + rate limiting).
 - **Application / Compute Layer:** Auto Scaling Group (ASG) of EC2 instances running inside private subnets, auto-scaled on CPU usage, and managed via AWS Systems Manager (SSM).
 - **Database Layer:** Multi-AZ RDS PostgreSQL database deployed within isolated subnets, allowing connections exclusively from the application tier.
-

Documentation Index

Explore different sections of our infrastructure documentation:

Core Configuration

- . [System Architecture](#): Deep dive into the physical network structure, routing tables, and AWS resource layout in the Malaysia region, including how the Developer's first design is mapped.
- . [AWS Phased Adoption Roadmap & Costing Guide](#): Multi-year week-by-week and month-by-month AWS service growth plan mapped from the project Gantt chart.
- . [Developer Design Alignment Guide](#): Rationale and comparison of shifting from legacy single-node Ubuntu VMs to an enterprise secure AWS design.
- . [ASGs & Separation of Concerns Guide](#): Best practice guide detailing auto-scaling with distinct ASGs, stateless principles, and the role of Amazon S3, EFS, or both.
- . [Root OpenTofu/Terraform Files](#): Overview of the root configuration entries (main.tf, variables.tf, outputs.tf, providers.tf).
- . [OpenTofu Migration Guide](#): Detailed compatibility research and transition path for deploying using OpenTofu on AWS.
- . [AMI Design & Hardening Guide](#): Architectural guide outlining the pre-baked AMI strategy, Packer/Ansible pipeline, and ASIMP

compliance integration.

- . **Route 53 & DNS Troubleshooting**: Deep dive into custom domain integration, ACM SSL/TLS setup, and extensive research on resolving ASG private subnet DNS resolution failures.
- . **Secure Developer Access Guide**: Comprehensive guide on using our secure SSH Jump host (Bastion) to access private and standalone nodes from Cyberjaya, with Windows/macOS/Linux client setups and private key security guidelines.
- . **Hybrid Cloud Integration Guide**: Comprehensive evaluation of secure, cost-optimized API and MCP connections alongside official AWS hybrid network solutions (VPN, Direct Connect, Transit Gateway) with granular MYR/USD costing models for ap-southeast-5.
- . **Disaster Recovery Options & National Sovereignty Guide**: Production-ready playbook covering the 4 standard AWS cloud DR options aligned with our Multi-AZ architecture, detailed local sovereignty/PDPA/CBPDT compliance reviews, and granular USD/MYR costing comparisons.
- . **RDS PostgreSQL 17 vs. Percona Server for PostgreSQL 17 Guide**: Comprehensive technical comparison of performance, telemetry, observability (PMM vs CloudWatch/Performance Insights), architectural designs, and costs in USD/MYR for ap-southeast-5.
- . **RAGFlow + Langfuse GPU Guide**: Architectural and economic analysis of RAGFlow and Langfuse, detailing critical GPU utilization, AWS vs. On-Premises deployment models, and secure API/MCP hybrid connections.
- . **Google Antigravity Skills Guide**: Comprehensive integration guide outlining how to share and deploy agent knowledge bases and skills between Google Jules and Google Antigravity.
- . **SOP: Knowledge-First Discovery**: Standard Operating Procedure outlining how AI agents perform local documentation search before probing remote targets.
- . **Wazuh Standalone Cloud Installation & Costing Guide**: Detailed installation steps for standalone Wazuh in the cloud (AWS Marketplace AMI and Graviton assistant modes), with independent USD/MYR costing tables for isolated financial tracking.

- . **Technology Stack Comparison Guide**: Complete architectural mapping and AWS alternatives comparison for the developer's containerized and external integrations (LangChain4j, Spring Boot, React, Twilio, Meta, Postgres, Valkey).
- . **Software Licensing & Technology Risk Register (TS/MC Series)**: Complete software licensing compliance framework, technology risk registry, and mitigation plans (TS/MC Series) covering LangChain4j, self-hosted operations, Bedrock with Qwen3 models, and standalone Wazuh SIEM.
- . **Strategic Comparative Review**: Comprehensive strategic analysis and financial TCO comparison of an AWS-Native Managed Platform against a Self-Hosted / On-Premises Custom Stack in Malaysia.
- . **Load Testing Assumptions & Sizing Guide**: Workload definitions, SLA metrics, architectural performance assumptions, and multi-VU sizing models from 100 to 10,000 VUs.

Infrastructure Submodules

- **VPC Module**: Core networking, public/private subnets, internet gateways, and NAT configurations.
- **Security Groups Module**: Strict firewall rulesets and port-level isolation.
- **WAF Module**: Layer-7 Web Application Firewall protecting the ALB.
- **ALB Module**: Application Load Balancer and health-check configurations.
- **ASG Module**: Auto Scaling Group, Launch Templates, and dynamic Graviton auto-detection.
- **RDS Module**: Multi-AZ PostgreSQL configuration and parameter group tuning.
- **Standalone EC2 Module**: Secure standalone Ubuntu 26.04 LTS development and application environments.
- **ElastiCache Valkey Module**: Secure ElastiCache Valkey in-memory caching cluster for session and metadata store.
- **Jumphost Module**: Secure public-subnet SSH Jumphost (Bastion) whitelisted for Cyberjaya office with automated downstream

ingress configuration.

Deployment & CI/CD

- [Automation Scripts](#): Details about CLI helpers (`deploy.sh` , `destroy.sh` , `user_data.sh`).
 - [CI/CD Pipeline](#): GitHub Actions workflow for automatic formatting, testing, validation, and OIDC deployment.
 - [GitLab EFS CI/CD](#): Comprehensive guide on GitLab CI/CD, automatic workflows, EFS mounting, dynamic Nginx path configurations, and containerized/S3 alternatives.
 - [Costing Estimate](#): Comprehensive monthly cost breakdown, local currency estimates, and Day-2 cost optimization pathways.
-

Prerequisites

Before deploying the infrastructure, ensure you have the following tools installed and configured:

- [OpenTofu](#) **>= 1.6.0** (Recommended) or **Terraform >= 1.5.0**
- [AWS CLI](#) configured with admin-level credentials for `ap-southeast-5`
- **Git** for repository and revision tracking

System Architecture

#aws #cloud #architecture #vpc #alb #asg #rds #waf #elasticache #valkey
#dns #ssl #acm #disaster-recovery #efs #postgresql #ragflow #langfuse
#costing

System Architecture

This document describes the high-availability 3-tier network topology, AWS component layouts, and routing architectures deployed by this project, fully aligned with our [Estimated Costing](#) model.

Additionally, this architecture is fully customized to map and host the **Developer's First Design (RAGFlow + LangFuse AI stack)** in a secure, highly-available, and resilient manner, built using hardened **Ubuntu 26.04 LTS** nodes and integrated with **Amazon ElastiCache for Valkey**.

Developer's First Design vs. AWS Production Architecture

In the developer's initial design, the application was structured across four separate standalone virtual machine servers running on **Ubuntu Server 26.04 LTS**:

- **Server 01 (Frontend - Web Tier):** Nginx Web Server / Reverse Proxy (2 vCPU, 4GB RAM)
- **Server 02 (Backend - App Tier):** Backend + DMS + MCP (4 vCPU, 16GB RAM)
- **Server 03 (AI Tier):** RAGFlow + LangFuse (4 vCPU, 8GB RAM)
- **Server 04 (Database - Data Tier):** SQL Database (4 vCPU, 16GB RAM)

While simple, deploying this directly as four standalone VMs introduces single points of failure (SPOF), security vulnerabilities from direct public internet exposure, manual backup overhead, and a lack of scalability.

To make this suitable for enterprise AWS deployment **without changing the AWS requirements**, we have mapped each of these components directly to a secure, multi-AZ, managed 3-tier architecture:

Architectural Comparative Mapping

Developer's Original Server	Original Spec	AWS Production-Ready Component	AWS Architectural Benefits
Server 01: Frontend	2 vCPU, 4GB RAM, Ubuntu	AWS WAFv2 + ALB + Private Nginx ASG & Dedicated Standalone Instance	No Public IPs & Local Baking Parity: Traffic enters via secure ALB/WAF. Private Nginx reverse-proxies run inside private subnets without public IPs. Paired with a dedicated Frontend Standalone instance connected to the same S3 bucket to test configurations and pre-bake custom Nginx ami-frontend-* images. Hardened via ASIMP .
Server 02: Backend	4 vCPU, 16GB RAM, Ubuntu	Multi-AZ ASG & Dedicated Standalone Instance	High Availability & 1:1 Database Parity: ASG spans multiple AZs and scales compute nodes dynamically. Paired with a dedicated Backend Standalone instance connected to the identical Multi-AZ RDS PostgreSQL database and S3 buckets to securely execute database migrations, test application logic, and pre-bake ami-backend-* images. Hardened via ASIMP .

Developer's Original Server	Original Spec	AWS Production-Ready Component	AWS Architectural Benefits
Server 03: AI Tier	4 vCPU, 8GB RAM, Ubuntu	Private ASG Compute & Dedicated AI Tier Standalone Instance	Isolated Compute & Instant Scaling via EFS Caching: Securely hosted in Private App Subnets. Paired with a dedicated AI Standalone instance connected to the identical Amazon EFS filesystem, RDS, and S3. Developers warm up Hugging Face model weight caches directly onto EFS from the standalone instance so that auto-scaling ASG nodes can access them instantly. Pre-baked into ami-ai-* images. Hardened via ASIMP .

Developer's Original Server	Original Spec	AWS Production-Ready Component	AWS Architectural Benefits
Server 04: Database	4 vCPU, 16GB RAM, Ubuntu	AWS RDS PostgreSQL (Multi-AZ) & Amazon ElastiCache for Valkey	Managed Resiliency & Enterprise Caching: Replaced self-managed SQL database with a fully managed Multi-AZ PostgreSQL database (db.m6g.xlarge). Synchronous replication, automated snapshots/failover, zero direct public route, and ingress restricted solely to private compute instances. Paired with a secure Amazon ElastiCache for Valkey cluster inside the database subnets to act as a high-performance in-memory task broker and session store for RAGFlow + LangFuse, reducing DB load and database overhead.

Standalone EC2 Instances for AMI Creation and Parity

To ensure seamless, reliable updates and zero-downtime rolling upgrades across our Auto Scaling Groups (ASGs), each application group (Frontend, Backend, and AI Tier) is paired with a dedicated **Standalone EC2 Instance** (running **Ubuntu 26.04 LTS** and hardened via the **ASIMP** framework).

These standalone instances are deployed directly inside the secure **VPC Private Application Subnets** and are configured to connect to the exact same shared resources (AWS RDS PostgreSQL Database, Amazon S3 Buckets, Amazon EFS shared storage, and Amazon ElastiCache Valkey caching) as their corresponding ASGs:

. **Frontend Standalone Instance:**

- Connected to **Amazon S3** to manage and verify static assets or remote template files.
- Used to test Nginx routing/configurations before baking the `ami-frontend-*` image.

. **Backend Standalone Instance:**

- Connected to the **Multi-AZ RDS PostgreSQL** database (using identical DB endpoints and credentials), **Amazon S3** (using identical IAM Role permissions), and **Amazon ElastiCache Valkey** (port 6379).
- Used to verify backend service scripts, run migrations, and test application logic before baking the `ami-backend-*` image.

. **AI Tier Standalone Instance:**

- Connected to **Amazon EFS** (via Private Mount Targets) to read/write persistent AI model caches, **Amazon S3** for training documents, **RDS PostgreSQL** for metadata storage, and **Amazon ElastiCache Valkey** (port 6379) for distributed task broker signaling.
- Used to bootstrap RAGFlow + LangFuse dependencies, warm up Hugging Face/SentenceTransformers cache directories, and test container updates before baking the `ami-ai-*` image.

Architectural Advantages of Standalone-to-AMI Parity:


- **1:1 Environment Alignment:** Because each standalone instance connects to the exact same database, Valkey cache, and shared storage backends, developers can fully run, test, and validate configurations in a real production-like environment without any risk of deployment divergence.
- **Pre-Audited & Hardened Base:** Standalone instances serve as the staging template. Developers run the **ASIMP** auditing and hardening


pipelines directly on these instances, verifying system integrity reports before triggering the Packer/AMI capture.

- **Zero-Downtime Releases:** Once the standalone instance is verified and the AMI is baked, updating the ASG Launch Template and running an Instance Refresh executes a safe rolling update across the live cluster.

Architectural Schematic

The updated network topology below outlines how our three ASG application groups and their matching standalone AMI-baking instances connect to the shared database, Amazon S3, Amazon EFS, and ElastiCache Valkey within our AWS secure environment:





Network Isolation Layers

1. Presentation / Web Layer (Public Subnets)

- **Subnets:** 10.0.1.0/24 (AZ ap-southeast-5a) and 10.0.2.0/24 (AZ ap-southeast-5b).

- **Description:** Hosts public-facing services and manages secure domain mappings. This layer routes inbound internet traffic directly through the Internet Gateway (IGW).
- **Resources:**
 - **Route 53 DNS Routing:** Manages custom domain delegations and points A Alias records to the ALB. Integrated with AWS Certificate Manager (ACM) for automatic domain verification.
 - **Application Load Balancer (ALB):** Terminates and routes incoming connections. Replaces the external Nginx reverse proxy direct exposure (from Server 01), dispersing HTTP/HTTPS traffic to private instances.
 - **NAT Gateway:** A highly-available NAT Gateway deployment in public subnets provides secure outbound internet access for package retrieval and DMS API callbacks.
 - **AWS WAFv2 Web ACL:** Directly attached to the ALB with 3 rules (OWASP Core, SQLi, and IP Rate Limiting) to block bad actors at the edge.

2. Application Layer (Private Subnets)

- **Subnets:** 10.0.10.0/24 (AZ ap-southeast-5a) and 10.0.11.0/24 (AZ ap-southeast-5b).
- **Description:** Holds business and compute logic. Instances have no public IP addresses and cannot be accessed directly from the internet.
- **Resources:**
 - **Auto Scaling Group (ASG) EC2 Instances:** Hosts the application code (Nginx web service). Features **t4g.xlarge** or **m6g.xlarge** EC2 Instances (ARM Graviton, 4 vCPU, 16GB RAM each) equipped with **gp3 EBS Root Volumes**. They handle Server 02 (Backend + DMS + MCP) and Server 03 (RAGFlow + LangFuse) workloads securely on hardened Ubuntu 26.04 LTS.
 - **Standalone EC2 Instances:** Deploys a dedicated standalone instance next to each ASG application group (Frontend, Backend, and AI Tiers) inside the private subnets. These instances are connected to the exact same shared databases (RDS), caches

(Valkey), and storage systems (S3 and EFS) as their respective ASGs, acting as 1:1 replica environments for application staging, testing, ASIMP auditing/hardening, and pre-baking custom AMIs.

3. Database & Caching Layer (Isolated Private Subnets)

- **Subnets:** 10.0.20.0/24 (AZ ap-southeast-5a) and 10.0.21.0/24 (AZ ap-southeast-5b).
- **Description:** Dedicated to database servers and caching nodes. Deeply isolated without any outbound route to the internet or NAT gateways, minimizing any data extraction surface.
- **Resources:**
 - **Multi-AZ RDS PostgreSQL Instance:** Runs synchronously across multiple availability zones using **Multi-AZ db.m6g.xlarge PostgreSQL (4 vCPU, 16GB RAM)** with **gp3 Storage**, corresponding to Server 04's resource needs. This guarantees high-availability, automatic failover, and robust production database performance.
 - **Amazon ElastiCache for Valkey Cluster:** A secure, high-performance in-memory caching cluster running Valkey 7.2. It manages session state, caches database queries, and runs the Celery/Redis task queue for RAGFlow and LangFuse. It operates on port 6379 with transit and at-rest encryption enabled, allowing ingress only from private compute security groups.

4. Storage Tier (Amazon S3)

- **Description:** Statically hosted media, secure user uploads, and build backups. It is situated outside the VPC but fully integrated.
- **Resources:**
 - **Amazon S3 Bucket:** Fully encrypted standard object storage. Access is managed via IAM policies and secure credentials.

Routing Configuration

The architecture manages network traffic flow through three distinct route tables:

Public Route Table

- Associated with public subnets.
- Routes all outbound traffic (0.0.0.0/0) to the **Internet Gateway (IGW)**.

Private Application Route Table

- Associated with private application subnets.
- Routes all outbound traffic (0.0.0.0/0) to the **NAT Gateway** running in the public subnet.

Database Route Table

- Associated with private database subnets.
- Contains only local VPC route entries (10.0.0.0/16), ensuring database and cache traffic never traverses public routes or internet gateways.

AWS Phased Adoption Roadmap & Costing Guide

#aws

#cloud

#architecture

#costing

#roadmap

#lifecycle

#compliance

AWS Phased Adoption Roadmap & Costing Guide

This guide outlines the step-by-step AWS cloud adoption roadmap, designed to align with our project timeline spanning **2 years of development and 12 months of active support and maintenance (36 months total)**. Over this 36-month period, the **total overall AWS infrastructure cost** is projected at **\$47,709.78 USD** ($\approx$ **RM 212,308.88 MYR**). To provide a comprehensive financial blueprint for the entire lifecycle, we estimate **other non-AWS operational costs**—comprising software subscriptions, software licensing compliance, yearly Security Posture Assessments (SPA), load testing (first year only), and 3rd party on-call support for AWS—to be **\$44,800.00 USD** ($\approx$ **RM 199,360.00 MYR**). This results in a combined 36-month project operating budget of **\$92,509.78 USD** ($\approx$ **RM 411,668.88 MYR**).

Combined 36-Month Financial Projection (AWS vs. Non-AWS)

The table below provides a year-to-year financial breakdown comparing our cloud infrastructure expenditures directly against estimated non-AWS operating overheads:

Financial Period	AWS Cloud Cost (USD)	AWS Cloud Cost (MYR)	Estimated Non-AWS Cost (USD)	Estimated Non-AWS Cost (MYR)	Total Combine Cost (U
Year 1 (Months 1-12)	\$13,023.16	RM 57,953.18	\$16,600.00	RM 73,870.00	\$29,623
Year 2 (Months 13-24)	\$23,343.02	RM 103,876.62	\$14,100.00	RM 62,745.00	\$37,443
Year 3 (Months 25-36)	\$11,343.60	RM 50,479.08	\$14,100.00	RM 62,745.00	\$25,443
Grand Total	\$47,709.78	RM 212,308.88	\$44,800.00	RM 199,360.00	\$92,509

Note: All local currency conversions are calculated at a baseline exchange rate of 1 USD ≈ 4.45 MYR.

Non-AWS Operational Costs Breakdown

The estimated other non-AWS operational costs cover the following strategic categories:

- . **Software Subscriptions (\$200.00 USD/mo):** \$2,400.00 USD per year (≈ RM 10,680.00 MYR/year) for development collaboration and utility software.
- . **Software Licensing Compliance:** \$1,200.00 USD per year (≈ RM 5,340.00 MYR/year) for security scans and automated compliance auditing.
- . **Security Posture Assessment (SPA):** \$4,500.00 USD per year (≈ RM 20,025.00 MYR/year) for professional penetration testing and security audits.
- . **Load Testing (First Year):** \$2,500.00 USD (≈ RM 11,125.00 MYR) in Year 1 to validate multi-tier system scalability under high

concurrency.

- . **3rd Party On-Call AWS Support (\$500.00 USD/mo):** \$6,000.00 USD per year ($\approx$ RM 26,700.00 MYR/year) for active SLA and on-call engineering coverage.

This roadmap illustrates how our AWS infrastructure scales incrementally week-by-week and month-by-month, starting from a absolute minimum cost, single-Availability Zone developer sandbox footprint and evolving into a fully-redundant, highly-available, secure, and compliant multi-AZ enterprise production architecture in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)**.

Gantt Chart Activity Mapping

Confidentiality Declaration: All corporate mapping notes, proprietary organizational identifiers, and project-specific planning references have been fully replaced with clean, generic equivalents. All names, codes, and corporate markers are completely anonymized to preserve strict confidentiality.

Below is the definitive, anonymized mapping of the project activities to our generalised enterprise delivery lifecycle terms:

Activity	Generalised Project Scope Mapping
Project Kick-off Meeting	Administrative launch and project alignment
Project Management	Continuous delivery, monitoring, and compliance
Phase 1: AI Chatbot Engine Development	Base chat routing and Core AI engine implementation
• Requirements Analysis & Study	Specifying API payloads and database schema
• Architecture Design & Development	Backend coding and vector database layout design
• Supply, Installation & Configuration	Provisioning sandbox environment & base infrastructure
• Integration with Internal Systems	Bridging enterprise core APIs to the AI engine
• System & Integration Testing	End-to-end payload audits and validation
• Live Deployment & Go-Live	Provisioning highly-available public-facing load balancers
• Technical & Administrator Training	Handover of deployment runbooks to operations teams
• Managed Cloud Services & Operations	Ongoing cloud optimization, monitoring, and scaling
Phase 2: Omnichannel Messaging Integration	WhatsApp API integration & chat webhook flow
• Supply, Installation & Configuration	Provisioning API webhooks and secure NAT gateways


Activity	Generalised Project Scope Mapping
• Integration with Core AI Chatbot	Connecting message handlers to the Phase 1 engine
• System & Integration Testing	Auditing high-throughput messaging workloads
• Live Deployment & Go-Live	Enabling production webhook receiving channels
• User Interaction & Telemetry Monitoring	Tracing chat payloads via Langfuse observability
• Technical & Administrator Training	Operations team training on API token rotations
Phase 3: Customer Relationship Management (CRM)	Customer database, task tracking & ticket management
• Requirements Analysis & Study	CRM database schema design and integration scoping
• Architecture Design & Development	Backend UI views and multi-tenant schema coding
• Supply, Installation & Configuration	Scaling private compute ASG and Multi-AZ RDS
• Integration with Internal Systems	Direct linking of CRM fields to internal databases
• System & Integration Testing	Auditing CRM security parameters and transaction locks
• Live Deployment, Launch & Go-Live	High-performance public launch of the CRM system

Activity	Generalised Project Scope Mapping
• Technical & Administrator Training	Staff onboarding, user manuals, and training sessions
• Managed Cloud Services & Operations	CRM SLA monitoring and DB read-replica tuning
Phase 4: Omnichannel Mobile Application	Hybrid React Native mobile app development
• Requirements Analysis & Study	Mobile API endpoint definition and push notification plan
• Supply, Installation & Configuration	Launching secure edge caching (CloudFront + S3)
• Architecture Design & Development	Mobile coding, UI layouts, and local storage tuning
• Integration with Internal Systems	Bridging push services, mobile SSO, and data syncing
• System & Integration Testing	Beta testing, user acceptance, and load verification
• Live Deployment & Go-Live	App store submission and high-throughput production API launch
• Security Vulnerability & Pen Testing	Rigorous penetration tests, OWASP checks, and WAF tuning
• Technical & Administrator Training	Administrative portal onboarding
• Managed Cloud Services & Operations	Global mobile API latency optimization
Final Acceptance Test (FAT)	Formal client verification of all consolidated tiers

Activity	Generalised Project Scope Mapping
Official Unified Project Go-Live	Hard cutover to enterprise-scale production
Documentation & Operational Handover	Delivery of final system runbooks and OpenTofu IaC
Support, Maintenance & OpEx Optimisation	12-month post-launch support and Savings Plan lock-in

High-Level AWS Phased Growth Path

The architecture matures through four logical environments matching the development timeline:



Detailed Chronological Roadmap

Year 1, Month 1 (Weeks 1–4): Project Initiation & Kick-off

- **Gantt Chart Activity Mapping:** Project Kick-off Meeting, Project Management.
 - **AWS Architectural Focus:** Identity access scoping and network design.
 - **AWS Services Provisioned:**
 - **AWS IAM Identity Center:** Enforces multi-factor authentication (MFA) and granular role-based access control (RBAC).
 - **AWS Budgets & Billing Alarms:** Confirmed with a baseline budget threshold of \$150.00 USD to prevent accidental run-away charges.
 - **Amazon Route 53:** Public hosted zone established for early API domain mapping.
 - **Week-by-Week Technical Goals:**
 - **Week 1:** Establish enterprise AWS root account structure, configure IAM administrators, and set budget guardrails.
 - **Week 2:** Draft security policies and Cyberjaya office CIDR access blocks for administrators.
 - **Week 3:** Initial provision of an empty modular VPC structure inside the Malaysia (ap-southeast-5) region utilizing OpenTofu.
 - **Week 4:** Set up shared project logging bucket with strict life-cycle rules inside Amazon S3.
 - **Phase Cost (USD / MYR): \$1.00 USD / RM 4.45 MYR** (Route 53 hosted zone base fee + nominal IAM storage).
-

Year 1, Month 2 (Weeks 5–8): Phase 1 Requirements & Dev Setup

- **Gantt Chart Activity Mapping:** Phase 1 (AI Chatbot Engine Development) - Requirements Study, Supply, Installation & Configuration.
- **AWS Architectural Focus:** Setting up the single-AZ developer sandbox to support early AI Chatbot engine coding and local model verification.
- **AWS Services Provisioned:**

- **Amazon VPC:** Single public subnet + single private subnet inside ap-southeast-5a .
 - **Secure SSH Jumphost (Bastion):** Sized at t4g.micro to allow the Cyberjaya developer team to tunnel safely into the private subnet.
 - **Standalone Developer Compute Instance:** 1x t4g.medium instance running hardened Ubuntu 26.04 LTS (ASIMP framework) to compile and test the Java/Spring Boot AI Chatbot engine.
 - **Amazon RDS PostgreSQL (Single-AZ):** 1x db.t4g.medium PostgreSQL 16 database instance for chat history storage.
 - **Amazon ElastiCache Valkey:** 1x cache.t4g.micro node for caching user sessions.
 - **Week-by-Week Technical Goals:**
 - **Week 5:** Finalise AI chatbot payload specifications and backend Spring Boot schema designs.
 - **Week 6:** Launch the SSH Jumphost and restrict port 22 access strictly to the Cyberjaya office IP address.
 - **Week 7:** Provision the standalone compute and Valkey nodes, ensuring the security groups block all traffic not originating from the Jumphost.
 - **Week 8:** Deploy Flyway on the database to verify schema migrations.
 - **Phase Cost (USD / MYR): \$138.50 USD / RM 616.33 MYR** (Single-AZ baseline environment).
-

Year 1, Months 3–6 (Weeks 9–26): Chatbot Core Coding & Integration

- **Gantt Chart Activity Mapping:** Phase 1 (AI Chatbot Engine Development) - Architecture Design & Development, Integration with Internal Systems.
- **AWS Architectural Focus:** Transitioning the sandbox environment to a highly-resilient Dual-AZ Staging Environment to validate auto-scaling, load handling, and database failovers.

- **AWS Services Provisioned:**

- **Dual-AZ Networking:** Expanded VPC with private subnets across 2 Availability Zones, protected by **2x NAT Gateways** to ensure secure outbound API updates.
- **Application Load Balancer (ALB):** Public-facing ALB deployed in the public subnets to distribute mock testing traffic.
- **Auto Scaling Group (ASG):** Sized with a minimum capacity of 2x c7g.xlarge (ARM Graviton3) instances to support heavy processing.
- **Amazon EFS (Elastic File System):** Shared persistent NFS mount to share AI model cache files across the ASG compute tier.
- **Amazon RDS (Multi-AZ Upgrade):** Database upgraded to Multi-AZ db.t4g.large with synchronous replication to prevent a single point of failure (SPOF).
- **AWS Secrets Manager:** Deployed to secure database credentials, rotating them automatically every 30 days.

- **Week-by-Week Technical Goals:**

- **Week 9–12:** Write Spring Boot Chatbot logic and vector indexing algorithms.
- **Week 13–16:** Integrate the application tier with AWS Systems Manager (SSM) and completely disable traditional passwords or direct SSH.
- **Week 17–20:** Implement auto-scaling policies triggered when average CPU usage exceeds 70% across the ASG.
- **Week 21–24:** Establish private VPC endpoints to route Amazon S3 media traffic locally, bypassing NAT Gateway charges.
- **Week 25–26:** Conduct chaos engineering tests, simulating a full Availability Zone outage to verify sub-minute automated DB and ALB failover.

- **Phase Cost (USD / MYR): \$668.11 USD / RM 2,973.11 MYR** (Dual-AZ Staging Environment).

Year 1, Months 7–8 (Weeks 27–34): Chatbot Go-Live & WhatsApp Integration

- **Gantt Chart Activity Mapping:** Phase 1 (AI Chatbot Engine Development) - Go-Live, Training, Managed Cloud Services; Phase 2 (Omnichannel Messaging Integration) - Supply, Installation & Configuration, Integration with Core AI Chatbot, System & Integration Testing, Go-Live, User Interaction & Telemetry Monitoring.
- **AWS Architectural Focus:** Shifting to the High-Performance Production Plan to launch the Chatbot engine publicly and integrate the high-throughput WhatsApp Business webhook pipelines.
- **AWS Services Provisioned:**
 - **AWS WAFv2:** Associated with the ALB to intercept and filter out OWASP Top-10 attacks, SQL injection attempts, and enforce dynamic IP rate-limiting (e.g., maximum 100 requests per 5 minutes per client IP).
 - **Production Database Tier:** Sized to Multi-AZ db.m7g.xlarge (4 vCPU, 16GB RAM) with 250 GB of gp3 storage (3,000 IOPS) to absorb massive concurrent chat transactions.
 - **Production Cache Tier:** Scaled up to ElastiCache Valkey cache.r7g.large (Dual Node cluster) to act as a secure, fast query broker.
 - **Amazon API Gateway & AWS Lambda:** Serverless webhook receiver layer deployed to securely absorb massive, bursty incoming WhatsApp callback payloads from Meta, offloading them to **Amazon SQS** queues to prevent thread saturation in the Spring Boot backend.
 - **Amazon CloudWatch Application Signals:** Enabled to track end-to-end user latency and alert on HTTP 5XX spikes.
- **Week-by-Week Technical Goals:**
 - **Week 27:** Establish production ALB SSL certificates using AWS Certificate Manager (ACM).
 - **Week 28:** Deploy AWS WAF and configure the core rule-sets. Run final technical training for administrator teams on emergency scaling.
 - **Week 29:** Launch the serverless webhook API Gateway and AWS Lambda functions, writing integration metrics to CloudWatch.

- **Week 30:** Connect the WhatsApp endpoint to AWS End User Messaging (Social Channels), removing high-cost external SaaS middlemen.
 - **Week 31-32:** Execute massive simulated load testing, driving over 5,000 concurrent user interaction flows.
 - **Week 33-34:** Go-live with the WhatsApp Omnichannel API, monitoring trace logs in real-time.
 - **Phase Cost (USD / MYR): \$1,665.61 USD / RM 7,411.98 MYR** (High-Performance Multi-AZ Enterprise Production).
-

Year 1, Months 9-12 (Weeks 35-52): Phase 2 Managed Ops & Phase 3 CRM Inception

- **Gantt Chart Activity Mapping:** Phase 2 (Omnichannel Messaging Integration) - Technical & Administrator Training; Phase 3 (CRM Development) - Requirements Analysis & Study, Architecture Design & Development, Supply, Installation & Configuration.
- **AWS Architectural Focus:** Managing WhatsApp operations while establishing CRM development and isolated schemas on shared high-performance infrastructure to reduce overhead.
- **AWS Services Provisioned:**
 - **Isolated CRM DB Schema:** Configured inside the production RDS instance utilizing Flyway migrations, fully separating CRM tables with strict database roles.
 - **Additional ASG App Targets:** Launch template modified to run specialized CRM JVM profiles on the existing production ASG instances.
 - **Continuous Telemetry Tracking:** Enabled Amazon GuardDuty to monitor account-level DNS abuse and secure IAM roles.
- **Week-by-Week Technical Goals:**
 - **Week 35-38:** Deliver comprehensive operational training on Meta API token rotation.
 - **Week 39-42:** Code backend CRM service APIs and construct robust relational mappings.

- **Week 43–48:** Deploy database read-replicas inside the isolated database subnets to offload complex CRM query lookups from the write master node.
 - **Week 49–52:** Configure automated nightly snapshots using AWS Backup with cross-region replication.
 - **Phase Cost (USD / MYR): \$1,685.00 USD / RM 7,498.25 MYR**
(Production run-rate with nominal storage growth).
-

Year 2, Months 1–2 (Weeks 53–60): CRM Integration & Go-Live

- **Gantt Chart Activity Mapping:** Phase 3 (CRM Development) - Integration with Internal Systems, System & Integration Testing, Go-Live & Launch, Technical & Administrator Training, Managed Cloud Services & Operations.
- **AWS Architectural Focus:** Securing database performance and ensuring seamless system scaling during the public launch of the CRM system.
- **AWS Services Provisioned:**
 - **AWS Elastic Disaster Recovery (AWS DRS):** Configured as a continuous block-level continuous replication strategy from our on-premises database networks to a secure staging subnet in AWS.
 - **Elastic Load Balancer Routing Rules:** Configured listener paths to route /crm/* traffic to a dedicated CRM target group in the ASG.
 - **Amazon Route 53 Health Checks:** Automated DNS failover triggers routing traffic to emergency backup static endpoints.
- **Week-by-Week Technical Goals:**
 - **Week 53–54:** Complete end-to-end integration mapping between CRM fields and internal legacy systems.
 - **Week 55–56:** Run exhaustive database transaction lock verification under a peak stress model.
 - **Week 57–58:** Launch the public CRM portal. Monitor database CPU and Valkey memory metrics closely during the initial cutover.

- **Week 59–60:** Complete on-boarding training for internal business teams and hand over the final system guides.
 - **Phase Cost (USD / MYR): \$1,720.00 USD / RM 7,654.00 MYR**
(Incorporating active AWS DRS replication streams).
-

Year 2, Months 3–8 (Weeks 61–86): Super Mobile App & High-Performance AI Ingress

- **Gantt Chart Activity Mapping:** Phase 4 (Omnichannel Mobile Application) - Requirements Analysis & Study, Supply, Installation & Configuration, Architecture Design & Development, Integration with Internal Systems, System & Integration Testing, Live Deployment & Go-Live.
- **AWS Architectural Focus:** Introducing heavy-duty AI processing and secure mobile delivery channels to support the Super Mobile App launch.
- **AWS Services Provisioned:**
 - **Amazon CloudFront:** Global Content Delivery Network (CDN) with edge caching to speed up mobile static asset and image retrieval.
 - **AWS WAFv2 on CloudFront:** Move WAF protection to the global CDN edge to block malicious requests before they even reach the ALB.
 - **High-Performance GPU Compute Tier:** 1x Standalone Staging g5.xlarge instance (NVIDIA A10G GPU with 24GB VRAM) to execute specialized DeepDoc visual layouts and high-speed OCR.
- **Week-by-Week Technical Goals:**
 - **Week 61–68:** Code hybrid mobile screens in React Native and define the API contract.
 - **Week 69–74:** Provision S3 upload directories and attach CloudFront to bypass API endpoint processing.
 - **Week 75–80:** Connect mobile push notifications through Amazon Simple Notification Service (SNS).
 - **Week 81–84:** Integrate on-device OCR components with our backend GPU compute.

- **Week 85–86:** Submit the compiled hybrid application to the Apple App Store and Google Play Store, and cutover the production API endpoints.
 - **Phase Cost (USD / MYR): \$2,150.00 USD / RM 9,567.50 MYR** (Peak pricing incorporating heavy GPU compute nodes and CDN data transfers).
-

Year 2, Month 9 (Weeks 87–90): Mobile App Security Hardening

- **Gantt Chart Activity Mapping:** Phase 4 (Omnichannel Mobile Application) - Security Vulnerability & Pen Testing, Technical & Administrator Training, Managed Cloud Services & Operations.
 - **AWS Architectural Focus:** Running deep audits on mobile APIs and locking down security variables before finalizing the development phase.
 - **AWS Services Provisioned:**
 - **Standalone Wazuh SIEM Instance:** Sized at t4g.large inside the secure management subnet, continuously gathering and auditing Linux logs from all compute nodes.
 - **Amazon Inspector:** Automated security scanning of EC2 instances and container registries.
 - **Week-by-Week Technical Goals:**
 - **Week 87:** Trigger a deep penetration test across all public-facing ALB endpoints.
 - **Week 88:** Apply ASIMP Ansible playbooks to ensure 100% compliance with CIS Level 2 benchmarks.
 - **Week 89:** Finalise training for administrative teams on detecting API credential exposure.
 - **Week 90:** Adjust WAF geo-blocking rules to allow traffic exclusively from approved Southeast Asian networks, blocking foreign botnets.
 - **Phase Cost (USD / MYR): \$2,215.71 USD / RM 9,860.00 MYR** (Peak development footprint with active Wazuh security auditing).
-

Year 2, Months 10–12 (Weeks 91–104): Final Acceptance Test, Project Go-Live & Handover

- **Gantt Chart Activity Mapping:** Final Acceptance Test, Official Unified Project Go-Live, Documentation & Operational Handover, Support & Maintenance Period.
 - **AWS Architectural Focus:** Tearing down auxiliary staging environments, cleaning up temporary EBS volumes, and scaling the core Multi-AZ production cluster to its final stable state.
 - **AWS Services Provisioned:**
 - **Consolidated Enterprise Production Stack:** Streamlined, secure Multi-AZ configuration.
 - **Amazon CloudWatch Custom Dashboard:** Real-time single-pane visibility for operations teams.
 - **Week-by-Week Technical Goals:**
 - **Week 91–94:** Complete formal Client Final Acceptance Testing across all integrated tiers (Chatbot, WhatsApp, CRM, Mobile).
 - **Week 95–98:** Complete the hard cutover of production databases. Deliver final OpenTofu infrastructure codebase files.
 - **Week 99–104:** Establish the 12-month Support & Maintenance framework. Configure Automated Instance Scheduler routines to shut down developer environments outside business hours, saving 64% on idle staging compute!
 - **Phase Cost (USD / MYR): \$1,285.80 USD / RM 5,721.81 MYR**
(Optimised, high-availability baseline run-rate).
-

Year 3: Support, Maintenance & Cost Optimisation (Months 25–36)

- **Gantt Chart Activity Mapping:** Support & Maintenance Period - 12 months.
- **AWS Architectural Focus:** Applying Day-2 financial optimizations to lock in significant savings for the remaining support year.
- **AWS Services Provisioned:**

- **AWS Savings Plans:** Apply a 1-Year Compute Savings Plan to EC2 instances, cutting compute costs by **34%**.
 - **RDS Reserved Instances:** Secure a 1-Year Reserved Instance for Multi-AZ PostgreSQL, cutting database costs by **30%**.
 - **ElastiCache Reserved Nodes:** Secure a 1-Year Reserved Node for ElastiCache Valkey, saving **35%** on caching.
 - **Month-by-Month Technical Goals:**
 - **Months 25–28:** Purchase the Reserved Instances and monitor the billing console to confirm discounts are successfully applied.
 - **Months 29–32:** Run monthly database indexing optimizations to keep storage footprints below 250 GB.
 - **Months 33–36:** Maintain high uptime (>99.99%) and prepare the final cloud transfer of ownership documents.
 - **Phase Cost (USD / MYR): \$945.30 USD / RM 4,206.59 MYR** (Optimised enterprise production run-rate representing **43.2% savings** over standard on-demand pricing).
-

Comparative AWS Monthly Cost Matrix

The table below breaks down the monthly run-rate (in USD and MYR) across each distinct development milestone phase, showing the incremental growth and financial efficiency of our roadmap:

Milestone Phase	Est. Cost (USD)	Est. Cost (MYR)	Primary Deliverables
Month 1 (Kick-off)	\$1.00	RM 4.45	Initial Project Setup
Month 2 (Dev Setup)	\$138.50	RM 616.33	System Architecture & Design
Months 3 - 6 (Staging)	\$668.11	RM 2,973.11	Development & Integration
Months 7 - 8 (Chatbot Go-Live)	\$1,665.61	RM 7,411.98	Initial Deployment & Testing
Months 9 - 14 (CRM Rollout)	\$1,720.00	RM 7,654.00	Advanced CRM Integration
Months 15 - 20 (Mobile Launch)	\$2,150.00	RM 9,567.50	Advanced CRM Integration & Mobile App Development
Months 21 - 22 (Security Pen)	\$2,215.71	RM 9,860.00	Security Audit & Penetration Testing
Months 23 - 24 (Project Go-Live)	\$1,285.80	RM 5,721.81	System Integration & Deployment
Year 3 (Support & Maintenance)	\$945.30	RM 4,206.59	Deployment & Ongoing Support

Strategic Recommendations for Phase-by-Phase Rollout

- Leverage “Scale-to-Zero” Webhooks Immediately:** In Phase 2, deploy API Gateway and AWS Lambda webhook integrations right from the

start. This keeps initial development costs near **\$0.00** while ensuring the backend is ready to handle high traffic spikes.

- . **Standardise on Graviton Compute:** Standardise all application EC2 nodes and RDS instances on AWS Graviton (ARM64). This delivers up to **40% better price-performance** compared to x86 equivalents, which directly lowers the run-rate of the Staging and Production phases.
- . **Establish Free PrivateLink S3 Endpoints early:** Create VPC Gateway Endpoints for S3 during Year 1, Month 2. Since the AI Chatbot and CRM download substantial volumes of document/media objects from S3, routing this traffic locally within the VPC avoids NAT Gateway data processing fees (\$0.045/GB).
- . **Automate Non-Production Environments:** Enforce a strict non-production schedule starting from Year 1, Month 3. Automatically stopping Staging and Dev environments outside standard Cyberjaya business hours (12 hours/day, 5 days/week) instantly reduces non-production compute costs by **64%**.

Sovereign and Regulatory Pathway Alignment (Malaysian PDPA 2010)

Keeping data secure and compliant with local regulations is a core priority of this phased roadmap:

- **In-Region Sovereign Processing:** In Phase 1 and Phase 2, all AI processing and messaging pathways are securely locked inside the AWS Malaysia region (ap-southeast-5). Sensitive citizen communications and metadata never leave the physical boundaries of the Kuala Lumpur data centers.
- **Cryptographic Isolation:** From Phase 3 onwards, all database layers and CloudWatch log directories are encrypted at-rest using customer-managed keys via **AWS KMS**.
- **AWS DRS Secure Replication Compliance:** The continuous replication channels provisioned in Year 2, Month 1 using AWS Elastic Disaster Recovery utilize highly-secure, encrypted TLS 1.3

tunnels. This allows companies to align with the **Malaysian Personal Data Protection Act (PDPA) 2010 Section 129** cross-border restrictions, ensuring backup target paths are fully documented, audited, and compliant.

Developer Design Mapping

#aws

#cloud

#architecture

#vpc

#alb

#asg

#rds

#waf

#dns

#ssl

#efs

#postgresql

#ragflow

#langfuse

#compliance

Developer Design Alignment Guide

This guide details how we transition the **Developer's First Design** (which specified four separate, standalone Ubuntu 26.04 LTS servers) into our secure, highly-available, production-ready **AWS 3-Tier Architecture**, without changing any underlying AWS constraints or requirements.

Rationale for the Architectural Transition

The developer's original design was logical but modeled around static, single-node virtual machines. Running critical production infrastructure on standalone VMs presents several operational risks:

- . **Single Points of Failure (SPOFs):** If any of the servers experiences hardware degradation or guest OS crashes, the entire application goes offline.
- . **Direct Public Exposure:** Placing database or backend servers directly on public-facing networks increases the surface area for brute-force SSH, port-scanning, and SQL injection (SQLi) attacks.
- . **Manual Scalability & High Latency:** VM resources are fixed and cannot scale dynamically in response to user demand or computational peaks.
- . **Backup and Patches overhead:** Managing backups, OS updates, and package security patches manually across multiple VMs takes significant operational effort.

By aligning this layout with our secure AWS design, we retain **all functionality** of the developer's original services (Nginx, Backend, DMS, MCP, RAGFlow, LangFuse, and Postgres) while inheriting cloud-native security, performance, and automation.

Detailed Component Transitions

1. Server 01: Frontend Web Tier (Nginx Web Server / Reverse Proxy)

- **Developer's First Design:** A single Ubuntu server running Nginx (2 vCPU, 4GB RAM) exposed to the public internet via port 80/443.
- **AWS Alignment:**
 - **Layer 7 Firewall (AWS WAFv2):** Blocks common exploits (OWASP Top 10, SQLi, XSS) and manages brute-force or DDoS attempts via dynamic IP Rate Limiting.
 - **Application Load Balancer (ALB):** Terminates external SSL/TLS, manages routing rules, and automatically performs health checks to route traffic only to healthy instances.
 - **Private Nginx Web Tier (ASG):** Sized as t4g.medium (2 vCPU, 4GB RAM) but deployed within **private subnets** under an Auto Scaling Group (ASG). This ensures the Nginx instance has *no public IP*, can scale out horizontally, and is protected from brute-force exposure.
 - **Dedicated Standalone Instance (AMI Baking):** Paired with a dedicated Frontend Standalone instance inside the private subnets. Connected directly to identical S3 resources to test routing/static templates and pre-bake certified ami-frontend-* images.

2. Server 02: Backend App Tier (Backend + DMS + MCP)

- **Developer's First Design:** A single Ubuntu server (4 vCPU, 16GB RAM) running Backend, DMS, and MCP APIs.
- **AWS Alignment:**

- **Zero Direct Ingress:** Deployed inside Private App Subnets, completely blocked from direct public ingress. The ALB only forwards verified requests to this layer on Port 80 (or your custom API port).
- **Sizing Alignment:** Sized using Graviton t4g.xlarge (4 vCPU, 16GB RAM) to perfectly match the developer's CPU/Memory requirements while saving up to 20% in raw compute costs over comparable x86_64 nodes.
- **Secure Egress:** Outbound integrations (such as external DMS syncs, Webhooks, or payment gateway APIs) traverse a managed **AWS NAT Gateway** in the public subnet, masking instance IPs and preventing inbound traffic.
- **Dedicated Standalone Instance (AMI Baking & Schema Migrations):** Paired with a dedicated Backend Standalone instance connected to the identical Multi-AZ RDS database and S3 buckets. Used to test application updates, safely run database migrations, and pre-bake ami-backend-* images under a 1:1 replica environment.

3. Server 03: AI Tier (RAGFlow + LangFuse)

- **Developer's First Design:** A single Ubuntu server (4 vCPU, 8GB RAM) running RAGFlow + LangFuse AI stack.
- **AWS Alignment:**
 - **Intra-VPC Security:** Placed in the same Private App Subnets. Secure, ultra-low-latency private communication paths are used to connect the Backend (Server 02) to RAGFlow + LangFuse (Server 03) via internal DNS or security-group-protected private endpoints.
 - **Sizing Alignment:** Sized as c6g.xlarge or t4g.xlarge to provide the required 4 vCPU and 8-16GB RAM. If CPU-bound AI processing becomes intensive, the ASG will scale these instances dynamically.
 - **Environment Compatibility:** Fully compatible with containerized environments (Docker Compose / ECS) or native package runtimes.
 - **Dedicated Standalone Instance (AMI Baking & EFS Caching):** Paired with a dedicated AI Standalone instance connected to the

identical Amazon EFS filesystem, RDS, and S3. Developers pre-download and warm up AI model caches on EFS using this standalone instance to allow instant bootstrapping across the auto-scaled nodes. Used to pre-bake `ami-ai-*` images.

4. Server 04: Database Data Tier (SQL Database)

- **Developer's First Design:** A single Ubuntu server (4 vCPU, 16GB RAM) running self-managed PostgreSQL.
 - **AWS Alignment:**
 - **Fully Managed RDS (PostgreSQL 16):** Upgraded to **AWS RDS PostgreSQL** (Multi-AZ deployment) sized at `db.m6g.xlarge` (4 vCPU, 16GB RAM).
 - **Multi-AZ Replication:** Synchronously replicates data across physically separate Availability Zones. In the event of an outage in AZ A, AWS automatically fails over to AZ B with zero manual intervention or data loss.
 - **Absolute Network Isolation:** Deployed inside isolated Private Database Subnets. The database security group restricts incoming traffic *exclusively* to the application security group (ASG), preventing any direct internet access.
-

Operating System & Hardware Optimizations

To deliver maximum efficiency, we transition the underlying hardware platform from legacy x86 virtual machines to **AWS Graviton (ARM64)** processors:

- . **Price-Performance Efficiency:** AWS Graviton (`t4g` and `m6g` instances) delivers up to **40% better price-performance** compared to equivalent x86 instances, significantly lowering the monthly run costs.
- . **Ubuntu 26.04 LTS Base Operating System:** To leverage the latest performance improvements, security features, and modern container support, we standardize our base platform on **Ubuntu 26.04 LTS (Noble Numbat successor)**.


- . **Amazon Linux 2023 Option:** For lightweight workloads that do not depend on Canonical specific packages, **Amazon Linux 2023 (AL2023)** remains available as a minimal, cloud-optimized option.
-

Server Hardening & Security Compliance (ASIMP Integration)

In aligning the developer design with AWS enterprise standards, all Ubuntu 26.04 LTS compute resources (both ASG instances and Standalone instances) are hardened and tuned using **ASIMP (Ansible System Integrity Management Platform)** (available at github.com/linuxmalaysia/ASIMP).

ASIMP is a host-based, automated security hardening, compliance, and auditing framework that implements a strict **“Measure, Harden, Re-Measure”** paradigm to verify and guarantee security posturing before the machine is allowed to process production traffic.

The ASIMP Hardening & Auditing Pipeline



Key ASIMP Hardening Capabilities Applied:

- . **Dual-Engine Security Auditing:**

- **OpenSCAP:** Conducts formal vulnerability and security compliance scanning mapped against the **CIS Security Ubuntu Linux Benchmark Level 2** profile.
 - **Lynis:** Performs comprehensive system configuration auditing, examining OS parameters, boot configurations, cryptography standards, and active network ports.
 - . **Pre & Post Scorecard Comparison:**
 - Runs audits prior to hardening to capture initial baselines, then executes them afterwards, logging comparative metrics in `/var/log/asimp-baseline-scores.json`.
 - Generates visually comprehensive, standalone HTML inspection reports at `/var/log/openscap-before-report.html` and `/var/log/openscap-after-report.html`.
 - . **Automated Package Updates & debsums Verification:**
 - Standardizes the Ubuntu system upgrade procedures and checks package-level code integrity via debsums to detect any unauthorized binary modifications.
 - . **Standardized OS Hardening Benchmarks:**
 - Deploys rigorous security compliance controls utilizing OpenStack's `ansible-hardening` role.
 - Standardizes and restricts SSH configuration endpoints using Dev-Sec's certified `ssh-hardening` roles, enforcing secure cipher suites, disabling password-based root access, and specifying key exchange standards.
 - . **Detailed System Tuning & Custom Fixes:**
 - Automatically tunes virtual memory configuration parameters, core kernel dumps, file system mounting options (e.g., `nodev`, `nosuid`, `noexec` where appropriate), and limits access to system compilers.
-

Summary of the AWS Security & HA Multipliers

Through this alignment, the developer gets their exact applications deployed with unmatched production capabilities:

Feature	How AWS Improves Developer's First Design
High Availability	Multi-AZ redundancy for Compute and RDS DB.
DDoS & Web Protection	AWS WAFv2 blocking bad traffic before compute
Elastic Scaling	ASG automatically adds instances based on load
Data Protection	Automatic, daily snapshots + Multi-AZ backups
Security Principle	Zero-Trust isolated subnets & IAM roles.
System Integrity	Hardened and audited via ASIMP framework.

Auto Scaling Groups (ASGs) & Separation of Concerns

#aws #cloud #architecture #vpc #alb #asg #rds #ssl #disaster-recovery
#efs #postgresql #gpu #ragflow #langfuse

Auto Scaling Groups (ASGs) & Separation of Concerns

In a modern cloud-native architecture, managing complex applications requires balancing scalability, security, cost, and operational maintenance. This document explores the architectural rationale for using **Separation of Concerns (SoC)** via distinct Auto Scaling Groups (ASGs) and provides an exhaustive guide on how to handle shared storage (**Amazon S3**, **Amazon EFS**, or both) to enable seamless, stateless autoscaling.

1. Separation of Concerns (SoC) via Dedicated ASGs

Using a single, monolithic fleet of EC2 instances to run multiple, diverse components of your application (e.g., Frontend, Backend, and AI Processing Tiers) creates tight coupling, increases the blast radius of failures, and complicates scaling.

By defining **distinct Auto Scaling Groups** for each logical application tier (e.g., an ASG for Server 01: Nginx/Frontend, an ASG for Server 02: Backend + DMS, and an ASG for Server 03: RAGFlow + AI Processing), you achieve robust Separation of Concerns.

Architectural Advantages of Multi-ASG Layouts

. **Independent Elastic Scaling:**

- **Different Resource Triggers:** Frontend layers scale based on network load or request counts, Backend layers scale based on CPU or memory, and AI layers (RAGFlow/LangFuse) scale based on custom metrics like GPU utilization or task queue length.
- **Cost Optimization:** You don't need to scale expensive, heavy instances (e.g., Graviton t4g.xlarge with 16GB RAM) just because the simple web server/reverse proxy layer is experiencing high traffic. Only scale the specific layer under load.

. **Blast Radius Limitation:**

- If a memory leak or heavy computation in the AI tier (Server 03) crashes an instance, the Nginx web tier (Server 01) and the core Backend tier (Server 02) remain completely operational, allowing you to return graceful degraded responses to users.

. **Least Privilege Security (IAM Roles & Security Groups):**

- Each ASG is equipped with its own dedicated **IAM Instance Profile**. The Backend ASG can have permission to read from specific database parameter stores, whereas the AI ASG can have exclusive read/write access to S3 training buckets.
- Security Groups can be strictly chained: the Frontend ASG security group only accepts traffic from the ALB; the Backend ASG security group only accepts traffic from the Frontend ASG; the Database and AI tiers are accessible only by the Backend ASG.

. **Decoupled Deployments & Rolling Updates:**

- Code updates to the RAGFlow AI stack can be rolled out as a rolling update to the AI ASG (via Instance Refresh) without causing a single millisecond of downtime or disruption to the core Backend and Frontend layers.

2. **Statelessness: The Core Prerequisite of ASGs**

The foundational rule of Auto Scaling is **statelessness**. Because ASG instances are automatically launched, terminated, and replaced during scaling events (scale-out, scale-in, or health-check replacements), **no persistent state should ever reside directly on an instance's local root disk (EBS)**.

If an application writes files (e.g., user profile pictures, uploaded documents for RAG processing, temporary cached AI model weights) directly to the local filesystem of an EC2 instance, those files will be:

- **Inaccessible** to other instances in the ASG (causing HTTP 404 errors when a load balancer routes subsequent requests to a different instance).
- **Permanently lost** when the specific instance scale-in or gets replaced due to a failed health check.

To solve this, you must decouple your compute tier (ASG) from your storage tier. This brings us to the core decision: **Amazon S3, Amazon EFS, or both?**

3. Storage Guide: Amazon S3 vs. Amazon EFS

To handle data across auto-scaling instances, AWS offers two primary shared storage options. Understanding when to use each—or combining them in a hybrid layout—is critical to a high-performing architecture.

Option A: Amazon S3 (Simple Storage Service)

Amazon S3 is a highly durable, virtually infinite, secure, and cost-effective **object storage service** designed for cloud-native applications.

- **How it works:** Files (objects) are accessed and manipulated directly via HTTP/HTTPS REST API calls (typically using AWS SDKs or the AWS CLI) rather than traditional filesystem commands (like `open()`, `write()`, or `seek()`).

- **Best suited for:**
 - Modern, stateless, cloud-native applications.
 - User-facing uploads (e.g., images, PDFs, videos).
 - Raw data source files before they are ingested, chunked, and embedded into vector databases by AI services like RAGFlow.
 - Long-term backups, application logs, and database snapshots.
- **Pros:**
 - **Incredible Scalability:** Handles millions of concurrent requests effortlessly.
 - **Unmatched Durability:** Designed for 99.999999999% (11 9s) of durability.
 - **Low Cost:** Very inexpensive compared to standard filesystems (\$0.023 per GB/month for standard, even lower for intelligent tiering/infrequent access).
 - **Global Delivery Integration:** Seamlessly integrates with Amazon CloudFront CDN for global low-latency asset delivery.
- **Cons:**
 - **Non-POSIX Compliant:** You cannot “mount” S3 as a traditional directory and perform standard file writes/appends in legacy code without utility layers (like Mountpoint for Amazon S3), which have limitations on write patterns.

Option B: Amazon EFS (Elastic File System)

Amazon EFS is a fully-managed, serverless, POSIX-compliant **network filesystem (NFSv4)** designed to be mounted concurrently by hundreds of EC2 instances across multiple Availability Zones.

- **How it works:** EFS is mounted over the network (port 2049) to a local directory on your EC2 instances. To the operating system, it looks and behaves exactly like a standard local directory.
- **Best suited for:**
 - Legacy or commercial off-the-shelf (COTS) software that expects a standard filesystem path (e.g., /var/www/uploads or /opt/app/data) and cannot easily be modified to use an AWS SDK.

- Shared directories requiring low-latency read-write concurrency across multiple compute instances simultaneously.
 - **AI Model Cache:** Sharing large pre-trained AI models or embedding files (e.g., Hugging Face model caches, SentenceTransformers weights) across your RAGFlow/AI ASG instances. Mounting EFS prevents instances from having to download multi-gigabyte files from S3 or external hubs during a scale-out bootstrap, significantly accelerating launch speed.
 - **Pros:**
 - **POSIX Compliant:** Supports full directory structures, file locking, permissions (UID/GID), and standard file system tools (tar, gzip, grep).
 - **Elastic Capacity:** Scales automatically up to petabytes without manual provisioning; you only pay for what you use.
 - **Multi-AZ Availability:** Built-in replication across multiple AZs.
 - **Cons:**
 - **Higher Cost:** Significantly more expensive than S3 (\$0.30 per GB/month for standard SSD storage, though lifecycle policies can tier unused data to lower-cost tiers).
 - **I/O Overhead:** Network filesystem overhead means file read/write latencies are higher than local SSDs (EBS) or highly-optimized API calls.
-


4. Architectural Comparison Table

Feature	Amazon S3 (Object Storage)	Amazon EFS (Network File System)
Primary Protocol	HTTP / HTTPS (REST API)	NFSv4 (TCP Port 2049)
POSIX Support	No (requires custom SDK or client)	Yes (standard directory, file permissions, file locking)
Scalability	Virtually infinite; high concurrent request throughput	Automatic capacity scaling; scales throughput with size
Typical Ingress Cost	Free inbound; \$0.023/GB/month (Standard)	\$0.30/GB/month (Standard Storage)
Access Latency	Tens of milliseconds	Single-digit milliseconds (metadata operations can be slower)
Direct Web Serving	Yes (via CloudFront / S3 Public endpoints)	No (must traverse an EC2 instance/web server)
Mounting on EC2	Direct SDK integration preferred (or Mountpoint utility)	Standard Linux mount command via amazon-efs-utils
Best AWS ASG Case	Storing user uploads, large datasets, and logs	Shared AI model weight caches, shared configuration directories

5. The Verdict: Do we need EFS, S3, or Both?

For most modern AWS architectures utilizing Auto Scaling Groups, **using Both in a hybrid design is the industry best practice.** They are not mutually exclusive, but rather complementary.

Here is the architectural distribution for our secure 3-tier layout:



When to use S3 for Auto Scaling:

- **Stateless Web Files:** All user uploads (e.g., PDFs sent to the DMS or RAG system) must be directed immediately to an Amazon S3 Bucket. When a user uploads a document, the Backend API processes it, writes the raw file to S3, and records the metadata (such as the S3 Object URL) in the **RDS PostgreSQL Database**.
- **Scale-Out Safety:** When a new instance boots up inside the ASG, it doesn't need to copy legacy user data. It simply calls S3 URLs on demand.

When to use EFS for Auto Scaling:

- **Accelerating Scale-Out (The Bootstrapping Problem):** If your RAGFlow AI service requires a 4GB language model to process requests, downloading this model from S3 or Hugging Face during a scale-out operation can take several minutes. This severely delays your auto-scaling response. By storing the model weights inside a shared EFS directory (mounted as `/mnt/models`), any newly launched AI instance can read the models instantly from EFS, reducing scaling boot times from 10 minutes to under 60 seconds.
 - **Shared Working Workspace:** If multiple nodes in the AI ASG must cooperate on batch processing files and require immediate POSIX write access.
-

6. Implementation and Configuration Guide

To implement EFS and S3 storage within your Auto Scaling Groups, configure the following AWS parameters.

6.1 Configuring S3 Integration for ASGs

To allow instances in your ASG to read and write to an S3 bucket securely:

- . **Do NOT use static AWS Access Keys.** Instead, attach an IAM Policy to your ASG instance role:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "s3:GetObject",
        "s3:PutObject",
        "s3:ListBucket"
      ],
      "Resource": [
        "arn:aws:s3:::my-application-storage-bucket",
        "arn:aws:s3:::my-application-storage-bucket/*"
      ]
    }
  ]
}
```

- . **Launch Template User Data Integration:** In your launch template, you can download startup assets or configurations using the AWS CLI:

```
#!/bin/bash
# Update and install AWS CLI
dnf update -y
dnf install -y awscli

# Download environment secrets or application configuration from S3
mkdir -p /etc/myapp
aws s3 cp s3://my-application-storage-bucket/config/config.json
/etc/myapp/config.json
```

6.2 Configuring EFS Integration for ASGs

To mount an EFS volume automatically on launch across multiple AZs:

- . **Security Group Configuration:**

- Create an **EFS Security Group** (sg-efs).
- Allow **Inbound TCP Port 2049 (NFS)** from your **ASG Security Group** (sg-asg).
- . **Deploy Mount Targets:**
 - In your OpenTofu/Terraform code, deploy EFS Mount Targets (aws_efs_mount_target) in **each private application subnet** within your VPC.
- . **Launch Template User Data Mounting:** Use the following script within your ASG's Launch Template user_data to install dependencies and mount EFS securely using TLS:

```
#!/bin/bash
# Install EFS mounting utilities (standard in Amazon Linux 2023)
dnf install -y amazon-efs-utils

# Define EFS Target Directory
MOUNT_DIR="/mnt/efs"
mkdir -p $MOUNT_DIR

# Mount using EFS File System ID (Enable TLS for encryption in-transit)
# Replace fs-xxxxxx with your actual EFS ID
mount -t efs -o tls fs-xxxxxx:/ $MOUNT_DIR

# Ensure the mount is persistent across potential reboots
echo "fs-xxxxxx:/ /mnt/efs efs _netdev,tls,defaults 0 0" >>
/etc/fstab
```

7. Cost-Optimization Best Practices

When leveraging S3 and EFS alongside ASGs, implement these practices to minimize operational costs:

- . **EFS Lifecycle Management:** Set EFS Lifecycle policies to automatically move files that haven't been accessed in 14 or 30

days to **EFS Infrequent Access (IA)** or **EFS Archive**. This can reduce EFS storage costs by up to 90%.

- . **S3 Intelligent-Tiering**: Enable S3 Intelligent-Tiering on your buckets. This automatically moves files between frequent, infrequent, and archive access tiers based on real-time usage patterns, without any performance overhead or administrative management.
- . **Mountpoint for Amazon S3**: For large read-heavy file access patterns (like streaming static training datasets to AI nodes), evaluate AWS's official Mountpoint for Amazon S3 as a high-throughput, lower-cost alternative to mounting EFS.

Root Terraform Configuration

#aws

#cloud

#architecture

#vpc

#alb

#asg

#rds

#waf

#route53

#dns

#postgresql

Root Terraform Configuration

The root folder of the terraform/ directory orchestrates the execution order, specifies environmental variables, defines provider versions, and outputs key service endpoints.

Files Overview

1. providers.tf

Specifies minimum required Terraform version requirements and registers external providers.

- **Minimum Terraform Version:** `>= 1.5.0`
- **AWS Provider Version:** `~> 5.0`
- **Region configuration:** Dynamically loaded from variables (defaults to `ap-southeast-5`).

```

terraform {
  required_version = ">= 1.5.0"
  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 5.0"
    }
  }
}

provider "aws" {
  region = var.aws_region
}

```

2. main.tf

Serves as the central manifest, calling modules in sequential dependency orders (Networking -> Security Groups -> ALB -> WAF -> ASG -> RDS -> Route 53) and passing cross-module attributes.

- **Example:** VPC subnet IDs are fed directly into ALB subnets, ASG subnets, and RDS subnets.
- **Example:** Security Group IDs are automatically mapped to protect dependencies.
- **Example:** Route 53 setup uses the ALB DNS name and Canonical Zone ID to route custom domain aliases.

3. variables.tf

Declares environmental configurations, default configurations, and input parameter structures.

- **Defaults for Malaysia Deployment:**
 - `aws_region` : "ap-southeast-5"
 - `db_engine` : "postgres"
 - `db_engine_version` : "16"
 - `db_instance_class` : "db.t4g.micro" (Graviton architecture)

- `instance_type`: "t4g.micro" (Graviton compute)
- `enable_route53`: true (whether to provision a Route 53 hosted zone and record)
- `domain_name`: "linuxmalaysia.com" (the target root domain name)
- `subdomain`: "app" (the frontend application subdomain)

4. outputs.tf

Exposes crucial deployment endpoints, such as the public ALB endpoint, the RDS endpoint, the WAF Web ACL ARN, and the Route 53 Name Servers for domain delegation.

```

output "alb_dns_name" {
  value      = module.alb.alb_dns_name
  description = "Public Load Balancer Endpoint"
}

output "rds_endpoint" {
  value      = module.rds.db_instance_endpoint
  description = "Isolated RDS Database Endpoint"
}

output "waf_web_acl_arn" {
  value      = module.waf.waf_web_acl_arn
  description = "WAF v2 Web ACL ARN"
}

output "route53_name_servers" {
  value      = try(module.route53[0].name_servers, [])
  description = "Route 53 delegated Name Servers for registrar
configuration"
}

output "route53_fqdn" {
  value      = try(module.route53[0].fqdn, "")
  description = "Route 53 FQDN pointing to the Application Load
Balancer"
}

```

5. terraform.tfvars.example

A template file used to configure local infrastructure credentials and system settings safely. Copy this template before running your deployments:

```
cp terraform.tfvars.example terraform.tfvars
```

OpenTofu Migration & AWS Support Research Guide

#aws

#cloud

#architecture

#vpc

#alb

#asg

#rds

#waf

#bastion

#ssl

#disaster-recovery

OpenTofu Migration & AWS Support Research Guide

This document presents a comprehensive feasibility study, architectural research, and practical migration path for transitioning our AWS 3-Tier Infrastructure configuration from **HashiCorp Terraform** to **OpenTofu**.

OpenTofu is a production-ready, open-source infrastructure-as-code (IaC) tool under the governance of the Linux Foundation (part of the Cloud Native Computing Foundation - CNCF). It was spawned as a community-driven fork of Terraform following its transition to the Business Source License (BSL).

1. Executive Summary

- **High Compatibility:** OpenTofu serves as a drop-in replacement for Terraform configurations (especially those compatible with Terraform 1.5.x, which this project uses with its `>= 1.5.0` requirement).
- **Zero AWS Resource Impact:** Migrating the deployment engine from terraform to tofu does not affect currently running AWS resources (such as VPCs, ALBs, ASGs, WAFv2, or RDS).
- **Same AWS API Calls:** Both engines utilize the standard AWS Provider which translates declarative configuration into AWS API calls using the AWS SDK for Go.

- **Enhanced Open Source Features:** OpenTofu offers advanced features like client-side state encryption, improved variable handling, and native resource exclusion (`tofu plan -exclude="..."`).
-

2. AWS Support and Compatibility Matrix

AWS natively supports OpenTofu indirectly and directly across its entire ecosystem, since OpenTofu acts as a client orchestrator of AWS APIs via the AWS Provider.

AWS Feature / Service	Support Level	Technical Integration Mechanism
AWS APIs & Resources	Native & Complete	Handled seamlessly via the hashicorp/aws provider, which is mirrored directly in the OpenTofu Registry. All resources (VPC, WAFv2, ALB, ASG, RDS) behave identically.
IAM Authentication	Native & Complete	Uses standard AWS SDK Go credential resolution chain (profiles, env vars, instance profiles).
OIDC / Web Identity	Native & Complete	Fully supports AWS OpenID Connect (OIDC) authentication, enabling passwordless GitHub Actions runs.
S3 Backend (State Storage)	Native & Complete	The s3 state backend in OpenTofu matches the Terraform S3 backend exactly, utilizing SSE, S3 Versioning, and IAM-based access.
DynamoDB (State Locking)	Native & Complete	Fully supports DynamoDB table state locking to prevent concurrent runs.
AWS CodeBuild	Complete (via scripts)	Since CodeBuild executes raw container environments, OpenTofu is fully supported. Tofu can be installed via package manager (APT/YUM) or official binary installation.
AWS CodePipeline	Complete (via CodeBuild)	CodePipeline orchestrates steps. Running OpenTofu within a CodePipeline pipeline is done via CodeBuild actions.

AWS Feature / Service	Support Level	Technical Integration Mechanism
AWS Proton	Complete (via custom runner)	AWS Proton allows specifying custom IaC engines through custom container-based provisioning, allowing seamless execution of OpenTofu.
AWS Systems Manager (SSM)	Supported	Parameter Store and Secrets Manager are accessed natively via data sources or backend credential configurations.

3. AWS Authentication Integration

OpenTofu inherits the exact same credential chain precedence as the AWS SDK and AWS CLI:

- . **Static Credentials:** Environment variables `AWS_ACCESS_KEY_ID`, `AWS_SECRET_ACCESS_KEY`, and `AWS_SESSION_TOKEN`.
- . **Shared Config & Credentials Files:** Located at `$HOME/.aws/config` and `$HOME/.aws/credentials` (leveraging `AWS_PROFILE`).
- . **Web Identity / OIDC Tokens:** Loaded via `AWS_WEB_IDENTITY_TOKEN_FILE` and `AWS_ROLE_ARN`. This is crucial for secure, keyless execution in **GitHub Actions** and **AWS EKS / ECS**.
- . **ECS Task Roles / Container Credentials:** Loaded using `AWS_CONTAINER_CREDENTIALS_RELATIVE_URI`.
- . **EC2 Instance Profile Metadata:** Evaluated when running directly on an EC2 instance (e.g., inside an administrative workstation or bastion host).

4. State Management and S3 Backend

OpenTofu's S3 state backend is completely compatible with Terraform's backend. The configuration remains virtually identical:

```
terraform {  
  backend "s3" {  
    bucket      = "my-opentofu-state-bucket"  
    key         = "state/terraform.tfstate"  
    region      = "ap-southeast-5"  
    dynamodb_table = "my-opentofu-lock-table"  
    encrypt     = true  
  }  
}
```

Note: OpenTofu accepts the terraform configuration block for backward compatibility. There is no need to rename this block to tofu in your .tf files.

State Migration Safeties:

- Before running OpenTofu on an existing Terraform-managed deployment, create a backup of your S3 state file:

```
aws s3 cp s3://my-opentofu-state-bucket/state/terraform.tfstate  
./terraform.tfstate.backup
```

- Running `tofu init` will recognize the existing state file and integrate with it natively.

5. Command Mapping

OpenTofu is designed with direct CLI command parity to make the transition trivial for developers:

Terraform Command	OpenTofu Equivalent	Purpose
terraform init	tofu init	Initializes directories, downloads providers/modules
terraform fmt	tofu fmt	Automatically formats configuration files
terraform validate	tofu validate	Checks configuration semantics and syntax
terraform plan	tofu plan	Outputs execution plan of pending AWS updates
terraform apply	tofu apply	Executes updates against active AWS APIs
terraform destroy	tofu destroy	Destroys all provisioned resources securely
terraform state <subcmd>	tofu state <subcmd>	Inspects or alters local/remote state records
terraform import	tofu import	Imports existing AWS resources into the state

Advanced OpenTofu Commands:

- **Excluding Resources Dynamically:**

```
tofu plan -exclude="module.rds"
tofu apply -exclude="module.rds"
```

(Terraform only supports targeting specific resources via -target, whereas OpenTofu supports negative targeting with -exclude).

6. How to Run OpenTofu in AWS Environments

A. AWS CodeBuild Integration

To run OpenTofu in an AWS CodeBuild environment, include the following in your `buildspec.yml`:

```
version: 0.2

phases:
  install:
    commands:
      # Install OpenTofu repository and binary
      - curl -s https://get.opentofu.org/install/opentofu-install.sh -o opentofu-install.sh
      - chmod +x opentofu-install.sh
      - ./opentofu-install.sh --install-method standalone
      - tofu --version
  pre_build:
    commands:
      - cd terraform
      - tofu init -backend-config="bucket=${STATE_BUCKET}"
  build:
    commands:
      - tofu plan -out=tfplan
  post_build:
    commands:
      - echo "Build completed on `date`"
```

B. AWS CodePipeline Integration

AWS CodePipeline does not have a dedicated action type for OpenTofu or Terraform out of the box. However, the standard best practice is to use **AWS CodeBuild** as a pipeline stage:

- . **Source Stage:** Code is retrieved from AWS CodeCommit, GitHub, or Amazon S3.

- . **Build Stage (CodeBuild):** CodeBuild installs and runs OpenTofu to generate an execution plan (`tofu plan`).
- . **Approval Stage (Optional):** Manual approval step inside CodePipeline to verify the plan.
- . **Deploy Stage (CodeBuild):** CodeBuild applies the plan (`tofu apply tfplan`).

C. AWS Proton Integration

AWS Proton coordinates infrastructure provisioning. For teams wanting to use OpenTofu:

- . Register a **Custom Provisioning** template.
- . Configure Proton to run an AWS CodeBuild project.
- . In the CodeBuild project, fetch the OpenTofu CLI and execute standard plan and apply commands mapped to the Proton service input schema.

7. Execution and Migration Plan for this Repository

To fully execute this replacement in this codebase, we perform the following steps:

- . **Report & Documentation:** Write this comprehensive guide (`docs/opentofu-migration.md`) and register it in the main Jekyll configuration and index (`docs/index.md`).
- . **Scripts Migration:** Modify CLI helpers `scripts/deploy.sh` and `scripts/destroy.sh` to transition to `tofu` while maintaining safety structures.
- . **GitHub Actions CI/CD Migration:** Update `.github/workflows/terraform.yml` to rename it to `opentofu.yml`, using the official `opentofu/setup-opentofu@v1` runner to ensure native, secure plan/apply on pushes.

- . **General References Refactoring:** Replace Terraform terminology and setup notes in README.md, docs/scripts.md, and docs/cicd.md with OpenTofu standards.

Amazon Machine Images (AMI)

Design Guide

#aws

#cloud

#architecture

#vpc

#alb

#asg

#rds

#efs

#postgresql

#ragflow

#langfuse

#compliance

Amazon Machine Images (AMI) Design & Hardening Guide

This document describes the **Amazon Machine Images (AMI)** strategy for our Auto Scaling Groups (ASGs). It outlines why custom, pre-baked AMIs are essential for stateless, rapid autoscaling, how to automate AMI construction using HashiCorp Packer and Ansible, and how the **ASIMP** auditing and hardening framework is integrated into the baking pipeline.

1. Why Pre-Baked AMIs are Critical for Auto Scaling

When an Auto Scaling Group (ASG) detects a surge in CPU utilization or network traffic, it immediately spins up new EC2 instances to share the workload.

If your application depends on a **just-in-time bootstrapping** approach (e.g., executing shell scripts in `user_data` that fetch package updates, download heavy language model weights, install dependencies, and compile binaries at boot-time), you introduce several critical failure modes:

- **Unacceptable Latency:** Bootstrapping heavy applications can take 10 to 15 minutes. By the time the instance is healthy and added

to the ALB target group, the peak load may have already overwhelmed the existing fleet.

- . **External Dependency Failure:** If an external package repository (e.g., Ubuntu APT servers, PyPI, or GitHub) experiences temporary downtime, the newly launched instance will fail to bootstrap, leaving your cluster unable to scale.
- . **Inconsistent Environments:** Multi-instance scaling can result in nodes running slightly different versions of minor packages if updates occurred between launch events.


The Immutable Infrastructure Paradigm

To solve this, we employ **Immutable Infrastructure** via custom **pre-baked AMIs**.

All software, libraries, configuration files, and system security controls are installed and validated *during the image building phase*. When the ASG scales out, the instance launches from the pre-baked AMI and transitions to a ready-to-run state in **under 45 seconds**, with zero runtime network dependency for configuration.

2. AMI Requirements Matrix per ASG Tier

To maintain clean separation of concerns and avoid bloated monolithic images, we define three specific, pre-baked AMIs built on top of **Ubuntu 26.04 LTS**:



A. Frontend Nginx Web Server AMI (`ami-frontend-*`)

- **Purpose:** Serves as the high-throughput reverse proxy and static content host.
- **Pre-Baked Software:**
 - Nginx (optimized with HTTP/2 and custom buffer limits).
 - AWS CloudWatch Agent (pre-configured to stream Nginx access and error logs).
 - **ASIMP Hardening:** Restricted SSH configuration, system account lockdowns, and local firewall parameters.

B. Backend App Tier AMI (`ami-backend-*`)

- **Purpose:** Runs the Backend, Document Management System (DMS), and Model Context Protocol (MCP) APIs.
- **Pre-Baked Software:**
 - Runtime Environments: Node.js (LTS), Python 3.12+, and package managers (`npm` , `pip`).
 - Git client and database connectors (for AWS RDS PostgreSQL communication).
 - Application source code and systemd service handlers.
 - **ASIMP Hardening:** Deep process limits (`limits.conf`), disabling of unused filesystem modules, and standard secure SSH ciphers.

C. AI Tier RAGFlow + LangFuse AMI (`ami-ai-*`)

- **Purpose:** Executes computation-heavy RAG processing, embeddings generation, and observability flows.
- **Pre-Baked Software:**
 - Docker Engine and Docker Compose (highly optimized for Graviton/ARM64 architectures).
 - Local caching directories mounted and configured for Hugging Face and SentenceTransformers model weights (accelerates initialization).
 - Container images pre-pulled and cached locally inside Docker storage.

- **ASIMP Hardening:** Kernel parameter tuning (`sysctl.conf`), container isolation profiles, and CIS Level 2 compliance benchmarks.
-

3. AMI Build Pipeline with Packer & ASIMP Hardening

We automate our AMI builds using **HashiCorp Packer** combined with **ASIMP (Ansible System Integrity Management Platform)** as the primary configuration provisioner. This guarantees that every pre-baked image is audited and hardened against standard security benchmarks prior to distribution.

The Immutable Build Workflow

- . **Initialize:** Packer spins up a temporary EC2 instance in a private build subnet using the base Canonical Ubuntu 26.04 LTS AMI.
- . **Provision:** Ansible executes the ASIMP playbooks against the temporary instance.
- . **Auditing (Pre-Harden & Post-Harden):**
 - ASIMP performs an initial security baseline scan using **OpenSCAP** (CIS Level 2) and **Lynis**.
 - ASIMP applies automated upgrades, validates system package integrity via `debsums`, configures secure SSH protocols (Dev-Sec), and implements fine-grained kernel modifications.
 - ASIMP re-runs the security scans, computes the comparative score improvement, logs metrics in `/var/log/asimp-baseline-scores.json`, and outputs visual HTML reports to `/var/log/openscap-after-report.html`.
- . **Bake:** Packer stops the temporary instance, takes a snapshot of the root EBS volume, registers a new AMI (`ami-frontend-yyyy-mm-dd`), and terminates the builder instance.

Sample Packer Configuration File (`template.pkr.hcl`)

Below is the Packer configuration used to build the hardened Ubuntu 26.04 LTS backend AMI:

```

packer {
  required_plugins {
    amazon = {
      version = ">= 1.2.0"
      source  = "github.com/hashicorp/amazon"
    }
  }
}

source "amazon-ebs" "ubuntu-hardened" {
  ami_name      = "asimp-ubuntu-26.04-backend-"
  instance_type = "t4g.medium" # Build on high-efficiency Graviton
  region        = "ap-southeast-5"

  source_ami_filter {
    filters = {
      name           = "ubuntu/images/hvm-ssd-gp3/ubuntu-noble-
24.04-*--server-*" # Fallback / target filter for Canonical Ubuntu LTS
      root-device-type = "ebs"
      virtualization-type = "hvm"
      architecture    = "arm64"
    }
    most_recent = true
    owners      = ["099720109477"] # Canonical
  }
  ssh_username = "ubuntu"
}

build {
  sources = ["source.amazon-ebs.ubuntu-hardened"]

  # Run ASIMP Hardening using Ansible
  provisioner "ansible" {
    playbook_file = "../asimp/play.yml"
    user          = "ubuntu"
    use_extra_vars = true
    extra_vars = {
      is_packer_build = true
    }
  }
}

```

```
    }  
  }  
  
  # Verify ASIMP reports are present before finalizing image  
  provisioner "shell" {  
    inline = [  
      "echo 'Verifying ASIMP Hardening Reports...'",  
      "ls -l /var/log/asimp-baseline-scores.json",  
      "ls -l /var/log/openscap-after-report.html"  
    ]  
  }  
}
```

4. Parameterizing & Referencing Baked AMIs in OpenTofu

Once Packer successfully registers the custom AMIs, they can be referenced inside your OpenTofu configurations dynamically or passed as static overrides.

Option A: Dynamic AMI Discovery (Recommended)

This approach dynamically queries AWS for the latest hardened AMI matching your tag patterns, ensuring that newly deployed Auto Scaling groups always inherit the latest security patches.

```

# Look up latest pre-baked Backend AMI
data "aws_ami" "latest_hardenened_backend" {
  most_recent = true
  owners      = ["self"] # Search within your AWS account

  filter {
    name   = "name"
    values = ["asimp-ubuntu-26.04-backend-*"]
  }

  filter {
    name   = "virtualization-type"
    values = ["hvm"]
  }

  filter {
    name   = "architecture"
    values = ["arm64"]
  }
}

```

You then supply this dynamic ID to your ASG launch template:

```

resource "aws_launch_template" "backend" {
  name_prefix   = "backend-launch-template-"
  image_id      = data.aws_ami.latest_hardenened_backend.id
  instance_type = "t4g.xlarge"
  # ... remaining options
}

```

Option B: Static AMI Parameter Override

For high-control staging/production environments where AMI updates must undergo rigid validation before scaling, define the image IDs as input variables:

```

# In variables.tf
variable "backend_ami_id" {
    description = "Pre-baked ASIMP Hardened Ubuntu 26.04 LTS AMI for Backend"
    type        = string
    default     = "" # If empty, config falls back to dynamic SSM or data queries
}


# In main.tf
locals {
    selected_backend_ami = var.backend_ami_id != "" ? var.backend_ami_id
    : data.aws_ami.latest_hardened_backend.id
}

```

This ensures full agility while retaining rigid compliance validation controls.

5. Standalone EC2 Instances: Staging and Baking Lifecycle

To achieve complete alignment between active development and the immutable production ASGs, our design utilizes **dedicated Standalone EC2 Instances** as live staging, testing, and pre-baking nodes for each of the three logical application tiers.



Environmental Alignment with Shared Backends

To prevent runtime configuration errors and bootstrap failures (such as database connection timeout or missing mount folders), the standalone instances are connected to the identical shared infrastructure as their matching ASGs:

- **Frontend Standalone:** Accesses the identical Amazon S3 static buckets to pull and verify Nginx page templates and asset configuration.
- **Backend Standalone:** Connects directly to the live **Multi-AZ RDS PostgreSQL** database (allowing safe database schema migration and endpoint verification) and **Amazon S3** (for DMS and document uploads).
- **AI Tier Standalone:** Mounts the identical multi-AZ **Amazon EFS** volume (facilitating direct Hugging Face and SentenceTransformers model caching). Developers download and cache model weights directly onto EFS from this standalone instance, which ensures newly-bootstrapped ASG nodes can access cached model weights instantly.

The Staging-to-AMI Lifecycle Workflow

- . **Active Staging & Application Testing:**
 - Developers perform code updates, install package dependencies, and modify system configuration on the standalone instance.
- . **ASIMP Auditing & Hardening Compliance:**
 - Developers execute the **ASIMP** auditing framework directly on the standalone instance. ASIMP executes OpenSCAP and Lynis compliance audits, applies Ubuntu security hardening configurations, checks package integrity, and produces HTML verification scorecards at `/var/log/openscap-after-report.html`.
- . **AMI Capture & Registration:**
 - Once compliance and functional tests pass successfully, the standalone instance's root volume is frozen.
 - An AMI is registered directly from the standalone instance using Packer or an automated AWS Systems Manager (SSM) backup document:

```
aws ec2 create-image \  
  --instance-id i-0xxxxxxstandalone \  
  --name "asimp-hardened-backend-$(date +%F)" \  
  --no-reboot
```

. **ASG Deployment & Rolling Refresh:**

- The registered AMI ID is supplied to OpenTofu as a variable override or auto-discovered dynamically.
- An **ASG Instance Refresh** is triggered, rolling out the new pre-validated, fully-configured image across the active ASG nodes with zero downtime.

Route 53 & DNS

Troubleshooting

#aws

#cloud

#architecture

#vpc

#alb

#asg

#rds

#waf

#route53

#dns

#ssl

#acm

#disaster-recovery

#postgresql

Amazon Route 53 Custom Domains & DNS Resolution

This document details how **Amazon Route 53** is integrated into our secure 3-Tier AWS Architecture for domain routing and SSL/TLS certificate registration. It also provides an in-depth research analysis of common **DNS resolution failures inside Auto Scaling Groups (ASGs)** (especially in private subnets) and step-by-step guidance to prevent and troubleshoot them.

1. Using Route 53 for Custom Domains

Our modular OpenTofu configuration automates the mapping of public domain names to our WAFv2-protected Application Load Balancer (ALB).

```
[ Web Client ]
```

```
|
```

```
▼ (HTTPS: app.linuxmalaysia.com)
```

```
[ Amazon Route 53 ] → (Resolves to ALB IP Addresses via A ALIAS Record)
```

```
|
```

```
▼
```

```
[ Application Load Balancer (ALB) ]
```

High-Level Setup Steps

To add your domain name to this deployment, complete the following workflow:

- . **Purchase or Delegate Domain:** Register a domain (e.g., `linuxmalaysia.com`) via Amazon Route 53 or any third-party registrar (GoDaddy, Namecheap, etc.).
- . **Retrieve Route 53 Name Servers:** When our OpenTofu configuration runs with `enable_route53 = true`, it creates a public Route 53 Hosted Zone and outputs a list of four Name Servers (NS):

```
tofu output route53_name_servers
```

- . **Configure Registrar NS Records:** Log into your domain registrar and replace your default name servers with the four Route 53 name servers generated by OpenTofu. This delegates DNS authority to AWS.
- . **ALB Routing (ALIAS Records):** To avoid charge fees and keep IP mappings dynamic, our module creates an **Alias A Record** (e.g., `app.linuxmalaysia.com -> production-alb-123456789.ap-southeast-5.elb.amazonaws.com`). AWS handles updates to ALB IPs automatically without modifying the DNS record.

2. HTTPS & SSL/TLS Configuration with AWS Certificate Manager (ACM)

To encrypt web traffic in transit (HTTPS on Port 443), we pair Route 53 with AWS Certificate Manager (ACM):

ACM SSL Provisioning Workflow

- . **Request Certificate:** Request a public wildcard certificate (e.g., `*.linuxmalaysia.com`) in ACM.
- . **DNS Validation:** ACM requires proof of domain ownership. Choose **DNS Validation**. ACM will generate a set of CNAME records.

- . **Publish Validation Records:** Add these CNAME records to your Route 53 Hosted Zone. This can be automated in OpenTofu using the `aws_route53_record` resource:

```
resource "aws_route53_record" "cert_validation" {
  for_each = {
    for dvo in aws_acm_certificate.cert.domain_validation_options :
    dvo.domain_name => {
      name    = dvo.resource_record_name
      record  = dvo.resource_record_value
      type    = dvo.resource_record_type
    }
  }

  allow_overwrite = true
  name            = each.value.name
  records         = [each.value.record]
  ttl             = 60
  type            = each.value.type
  zone_id         = var.hosted_zone_id
}
```

- . **Attach SSL to ALB Listener:** Once validated, ACM marks the certificate as **Issued**. You then update your ALB Listener in `terraform/modules/alb/main.tf` to use HTTPS (Port 443) and reference the certificate ARN:

```
resource "aws_lb_listener" "https" {
  load_balancer_arn = aws_lb.main.arn
  port              = 443
  protocol          = "HTTPS"
  ssl_policy        = "ELBSecurityPolicy-TLS13-1-2-2021-06"
  certificate_arn    = var.acm_certificate_arn

  default_action {
    type             = "forward"
    target_group_arn = aws_lb_target_group.asg_tg.arn
  }
}
```

3. In-Depth Research: ASG DNS Resolution Failures & Mitigation

Instances launched in an Auto Scaling Group (ASG) inside secure private subnets frequently run into **DNS resolution issues**. These failures break outbound API integrations (e.g., sending WhatsApp notifications, accessing Hugging Face model weight repositories, or making external OAuth callbacks).

Below is our technical research detailing the primary causes, mechanics, and concrete solutions for these DNS problems:

Cause A: The Nginx Upstream DNS Caching Bottleneck (The Silent 502)

This is the single most common failure in 3-Tier configurations where Nginx acts as a frontend or API reverse proxy routing traffic to dynamic backends.

- **The Problem:** By default, Nginx resolves domain names (such as your backend ALB DNS name or internal endpoints) **only once during startup** and caches the IP address indefinitely. However, AWS ALBs

dynamically scale up and down, shifting their public and private IP addresses in response to traffic surges.

- **The Symptom:** Nginx works perfectly at first. But when the ALB scales out or replaces healthy nodes, Nginx continues sending traffic to the decommissioned IP addresses, resulting in persistent 502 Bad Gateway or Connection Timed Out errors, even though DNS is working perfectly at the OS level.
- **The Solution:** Use variables to enforce dynamic DNS lookups. Inside your Nginx config (/etc/nginx/sites-available/default), define a resolver pointing to the AWS Core VPC DNS server (169.254.169.254 or the second IP address of your VPC CIDR, e.g., 10.0.0.2), and set a short TTL:

```
server {
    listen 80;

    # Use the AWS VPC DNS Resolver with a 10-second cache threshold
    resolver 169.254.169.254 valid=10s;

    set $backend_upstream "http://app.linuxmalaysia.com";

    location /api/ {
        # Using a variable forces Nginx to respect the resolver TTL
        proxy_pass $backend_upstream;
        proxy_set_header Host $host;
        proxy_set_header X-Real-IP $remote_addr;
    }
}
```

Cause B: VPC DNS Settings Disabled (Internal DNS Failure)

- **The Problem:** AWS VPCs provide automatic resolution for internal DNS queries (e.g., looking up database endpoints like production-database.xxxxxx.ap-southeast-5.rds.amazonaws.com). This fails if DNS support features are disabled in the VPC configuration.
- **The Symptom:** Private ASG instances cannot ping or connect to the RDS PostgreSQL instance using its domain endpoint, but can

connect using the raw private IP.

- **The Solution:** Ensure both VPC attributes are set to true inside your OpenTofu network code (terraform/modules/vpc/main.tf):

```
resource "aws_vpc" "main" {
  cidr_block      = var.vpc_cidr
  enable_dns_support = true # Enforces Route 53 resolver
  availability
  enable_dns_hostnames = true # Allocates public/private DNS
  hostnames
}
```

Cause C: Systemd-Resolved Configuration Issues on Ubuntu

Our target operating system, **Ubuntu 26.04 LTS**, manages name resolution via the local stub listener systemd-resolved on 127.0.0.53 .

- **The Problem:** If /etc/resolv.conf is a static file or a broken symlink instead of pointing to systemd-resolved's runtime stub, the instance cannot receive DNS server allocations propagated by DHCP Options.
- **The Symptom:** Running nslookup google.com fails, and /etc/resolv.conf is empty or lacks name servers.
- **The Solution:** In your baking pipeline or user_data.sh script, force-recreate the symbolic link to point to systemd-resolved:

```
sudo ln -sf /run/systemd/resolve/stub-resolv.conf /etc/resolv.conf
sudo systemctl restart systemd-resolved
```

Cause D: Misconfigured Outbound Security Groups or NACLs

- **The Problem:** DNS resolution requires outgoing communication on **Port 53 (UDP & TCP)** to Route 53 Resolver (169.254.169.254). If outbound Security Groups or Network ACLs are overly restrictive, the DNS queries are dropped before leaving the node.

- **The Symptom:** Running `dig @169.254.169.254 google.com` hangs and times out.
- **The Solution:** Verify that the Security Group assigned to the ASG (`asg_sg` in `terraform/modules/security_groups/main.tf`) permits outbound port 53:

```
# Egress configuration for ASG Security Group
resource "aws_security_group_rule" "asg_dns_out" {
  type           = "egress"
  from_port      = 53
  to_port        = 53
  protocol       = "udp"
  cidr_blocks    = ["0.0.0.0/0"]
  security_group_id = aws_security_group.asg_sg.id
}
```

(Note: A default egress of 0.0.0.0/0 with protocol -1 covers all ports, including DNS, which is recommended for the application tier).

Cause E: Route 53 Resolver Rate Limiting (DNS Query Flooding)

- **The Problem:** The AWS Route 53 Resolver enforces a limit of **1024 packets per second** (PPS) per Network Interface (ENI). In a high-throughput microservices environment with hundreds of concurrent connections, application instances can exceed this limit.
- **The Symptom:** Sporadic DNS resolution timeouts that resolve themselves after a few seconds, accompanied by errors in the application logs.
- **The Solution:** Configure local DNS caching to prevent every network socket call from sending an upstream packet to AWS DNS:
 1. Ensure `systemd-resolved` has local caching enabled (enabled by default).
 2. Implement an on-instance cache proxy like **dnsmasq** or **unbound** for high-load workloads.

3. Configure your applications (such as Node.js or Python) to use keep-alive connections to minimize DNS requests per API handshake.

Secure Developer Access Guide

#aws

#cloud

#architecture

#vpc

#asg

#jumphost

#bastion

#disaster-recovery

#compliance

#costing

Secure Developer AWS Access Guide (SSH Jumphost)

This guide provides technical research, architectural specifications, and step-by-step instructions for developers based in our **Cyberjaya, Selangor, Malaysia** office to securely access private AWS resources (private ASG compute servers, database systems, and standalone staging instances) using a highly secure, cost-optimized SSH Jumphost (Bastion).

Architectural Research & Feasibility Analysis

The Question: “Can/Should we deploy the Jumphost outside a VPC?”

During design discussions, a proposal was raised: **“Does deploying a small Amazon Linux not in a VPC give us a better and cheaper solution?”**

Based on cloud architecture fundamentals and AWS networking constraints, this proposal is **both technically impossible and architecturally undesirable**. Here is why:

- . **VPC Requirement in AWS:** Since the retirement of AWS EC2-Classic, **every EC2 instance must run inside a Virtual Private Cloud (VPC)**. There is no longer any option to deploy a “raw” EC2 instance outside of a VPC.

- **The “Separate Jumphost VPC” Anti-Pattern:** Even if the proposal is interpreted as deploying the Jumphost inside a *separate management VPC* (instead of our main application VPC), this approach is **not** cheaper or safer. It introduces severe cost and operational overhead:
 - **Transit Gateway (TGW) Costs:** To connect the separate Jumphost VPC to the private subnets of our application VPC, we would need to provision an AWS Transit Gateway. In the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)**, a TGW attachment costs **\$0.05 per hour** (~\$36.50/mo per VPC connection) plus a **\$0.02 per GB** data processing fee. This alone is **3.5x more expensive** than running the entire jumphost itself!
 - **VPC Peering Complexity:** Alternatively, establishing a VPC Peering Connection requires updating routing tables in both VPCs, dealing with inter-AZ data transfer fees (\$0.01/GB in each direction), and managing dual security boundaries, which increases the likelihood of human error.
 - **Resource Duplication:** A separate VPC would require its own Internet Gateway, public subnets, and routing infrastructure, duplicating base costs.

The Optimal Solution: Public Subnet + Strict IP Whitelisting

The **best, safest, and cheapest** approach is to deploy a small t4g.micro instance (equipped with energy-efficient ARM64 Graviton processors) inside the **existing public subnet** of our main application VPC:

- **No Extra AWS Networking Cost:** It reuses the existing public Internet Gateway and routing infrastructure.
- **Strict Ingress Firewalling (Zero-Trust):** The Jumphost’s Security Group permits inbound TCP port 22 connections **exclusively** from the public CIDR of our **Cyberjaya, Selangor, Malaysia office** (configurable via `jumphost_allowed_ssh_cidr`). All other internet traffic is dropped silently at the AWS Hypervisor layer.
- **Deep Private Isolation:** The private compute nodes (ASG) and standalone dev environments have **no public IP addresses** and are

completely unreachable from the public internet. They only permit SSH (port 22) ingress **exclusively from the Jumphost's Security Group**.

- **Minimal Cost Impact:** Running a t4g.micro with a 15GB gp3 EBS root volume and an Elastic IP costs only **~\$10.98 USD / month** (~RM 49.41 MYR/mo) on-demand, which can be further optimized via Savings Plans.

Technical Specifications of the Jumphost

Our OpenTofu setup automates this configuration. The Jumphost is deployed with:

- **Base OS:** Ubuntu 24.04 LTS (allows 100% compliance with the ASIMP hardening framework).
- **Network Sizing:** t4g.micro (ARM64/Graviton), costing **\$0.0084/hour**.
- **Storage:** 15GB gp3 SSD, encrypted at rest (gp3 storage at \$0.08/GB-mo).
- **Static IP:** Associated with an AWS Elastic IP (EIP) so developers have a stable endpoint that does not change upon reboot.

OS Hardening & Auditing via ASIMP

To protect our access gateway, the Jumphost must be hardened using **ASIMP (Ansible System Integrity Management Platform)**. ASIMP implements a "Measure, Harden, Re-Measure" security flow.

Key SSH Hardening Policies Configured via ASIMP:

- . **Disable Password-Based Authentication:** Only cryptographic keys are allowed.
- . **Disable Root Logins:** Root SSH is strictly disabled (PermitRootLogin no). Developers must connect using their own unprivileged user (e.g., ubuntu or standard service accounts) and escalate using sudo if authorized.

- . **Restricted Cryptographic Cipher Suites:** Enforces modern, secure key exchange algorithms and ciphers, completely blocking obsolete/weak ones (e.g., MD5, SHA-1, 3DES, Blowfish).
 - **Allowed Key Exchange (KEX):** curve25519-sha256 , diffie-hellman-group16-sha512
 - **Allowed MACs:** hmac-sha2-512-etm@openssh.com
 - **Allowed Ciphers:** chacha20-poly1305@openssh.com , aes256-gcm@openssh.com
 - . **Active Connection Auditing:** OpenSCAP and Lynis run scheduled audits to verify CIS Level 2 benchmark compliance, writing audit scorecard logs to /var/log/asimp-baseline-scores.json .
-

Critical SSH Private Key Security Guidelines

The private SSH key represents the developer's identity and grants administrative access. If a private key is compromised, the entire AWS infrastructure is at risk. All developers must follow these strict guidelines:

1. Always Use a Strong Passphrase

Never generate a passwordless SSH key. Always encrypt the private key with a strong passphrase. This ensures that even if the private key file is stolen (e.g., from a physical laptop loss or malware infection), the attacker cannot use it without decrypting it first.

2. Never Share Private Keys

A private key is private. It must **never** be shared via Slack, Email, WhatsApp, or uploaded to any shared network. Under no circumstances should multiple developers share the same SSH key.

3. Never Check Keys into Version Control

Ensure that private keys (such as .pem , .id_ed25519 , .ppk) are never committed to Git repositories. Add a global rule in your .gitignore file:

```
# Ignore SSH Keys
*.pem
*.key
id_ed25519
id_rsa
*.ppk
```

4. Apply Strict File Permissions (OS-Specific)

On Linux / macOS (Unix File Permissions)

SSH clients will refuse to use private keys that are too openly readable. The `.ssh` directory and key files must be restricted:

```
# Secure the SSH directory (read/write/search for owner only)
chmod 700 ~/.ssh

# Secure the private key file (read/write for owner only, no access for
# anyone else)
chmod 600 ~/.ssh/id_ed25519

# Secure public keys (read/write for owner, read-only for others)
chmod 644 ~/.ssh/id_ed25519.pub
```

On Windows (NTFS Access Control Lists)

Windows does not use standard Unix permissions, but the OpenSSH client on Windows still enforces strict security. To secure a private key (`id_ed25519` or `key.pem`) on Windows:

Option A: Using PowerShell (Fastest)

Run the following commands in PowerShell to strip inheritance and grant permissions exclusively to the current logged-in user:

```
# Store key path in a variable
$keyPath = "$Home\.ssh\id_ed25519"

# Strip inheriting permissions and remove all group permissions
icacls $keyPath /inheritance:r

# Grant full control exclusively to the current logged-in user
$currentUser =
[System.Security.Principal.WindowsIdentity]::GetCurrent().Name
icacls $keyPath /grant:r "${currentUser}:F"
```


Option B: Using Windows File Explorer (GUI)

- . Right-click the private key file (id_ed25519) and select **Properties**.
- . Go to the **Security** tab, and click **Advanced**.
- . Click **Disable inheritance** and select "Remove all inherited permissions from this object".
- . Click **Add**, click "Select a principal", type your Windows username, and click **Check Names**.
- . Set permissions to **Full control** and click **OK** -> **Apply** -> **OK**.

Developer Access Step-by-Step Instructions

To reach private AWS resources (e.g., Standalone Instance 10.0.10.15), developers use the Jumphost as a transit bridge.

```
[ Cyberjaya Office ]
(Public IP)
| (Port 22 Inbound)
```



Below are instructions for macOS, Linux, and Windows platforms.

1. macOS Developer Setup

macOS has native SSH client support. We recommend generating an **Ed25519** key, which is modern, lightweight, and highly secure.

Step 1: Generate your SSH Key Pair

Open Terminal and run:

```
ssh-keygen -t ed25519 -C "yourname@office.com"
```

Note: Specify a secure passphrase when prompted.

Step 2: Register your Public Key on AWS

Provide your public key (`~/.ssh/id_ed25519.pub`) to your AWS administrator. They will register it on the Jumphost and private instances.

Step 3: Configure your local SSH config file

To avoid typing complex commands every time, create or edit your SSH configuration file (`~/.ssh/config`):

```
nano ~/.ssh/config
```

Insert the following block (replace `54.x.y.z` with the Jumphost's actual Elastic IP address and `your-key` with your private key name):

```
# Common SSH settings
Host *
    AddKeysToAgent yes
    UseKeychain yes
    IdentityFile ~/.ssh/id_ed25519

# The SSH Jumphost
Host aws-jumphost
    HostName 54.x.y.z
    User ubuntu

# Private Standalone / ASG Application servers
Host aws-private-*
    # Use wildcard mapping to route any private VPC IP via Jumphost
    HostName 10.0.*
    User ubuntu
    ProxyJump aws-jumphost
```

Save and close (Ctrl+O then Ctrl+X).

Step 4: Access your Private Instances

Now, you can SSH into any private instance inside the VPC using a single, secure command:


```
# Connect directly to a private instance (macOS routes this
transparently via Jumphost)
ssh aws-private-10.0.10.15
```

2. Linux (Ubuntu) Developer Setup

Linux systems use OpenSSH. The configuration is almost identical to macOS.

Step 1: Generate your SSH Key Pair

```
ssh-keygen -t ed25519 -C "yourname@office.com"
```

Step 2: Configure your SSH config file

Create or edit your local SSH configuration file:

```
nano ~/.ssh/config
```

Add the following configuration block:

```
Host *
  AddKeysToAgent yes
  IdentityFile ~/.ssh/id_ed25519
```

```
Host aws-jumphost
  HostName 54.x.y.z
  User ubuntu
```

```
Host aws-private-*
  HostName 10.0.*
  User ubuntu
  ProxyJump aws-jumphost
```

Step 3: Start SSH Agent & Load Key

Add these commands to your terminal profile (~/.bashrc or ~/.zshrc) to automate loading:

```
eval "$(ssh-agent -s)"  
ssh-add ~/.ssh/id_ed25519
```

Step 4: Access Private Instances

```
ssh aws-private-10.0.10.15
```

3. Windows Developer Setup

Windows developers can use **PowerShell (with native OpenSSH)** or **PuTTY**.

Method A: Using PowerShell & Native OpenSSH (Recommended)

Modern Windows 10/11 has OpenSSH Client installed by default.

Step 1: Generate your SSH Key Pair

Open PowerShell and run:

```
ssh-keygen -t ed25519 -C "yourname@office.com"
```

Step 2: Enable and configure Windows SSH Agent

Start the Windows SSH Agent service (requires running PowerShell as Administrator once):

```
# Set Startup Type to Automatic
Set-Service -Name ssh-agent -StartupType Automatic

# Start the Service
Start-Service ssh-agent

# Load your Private Key into the Agent
ssh-add $Home\.ssh\id_ed25519
```

Step 3: Configure your local SSH Config File

In PowerShell, edit or create the SSH config file:

```
notepad $Home\.ssh\config
```

Add the following configuration block (replace 54.x.y.z with the Jumphost's Elastic IP):

```
Host *
  AddKeysToAgent yes
  IdentityFile ~/.ssh/id_ed25519

Host aws-jumphost
  HostName 54.x.y.z
  User ubuntu

Host aws-private-*
  HostName 10.0.*
  User ubuntu
  ProxyJump aws-jumphost
```

Step 4: Secure Key Permissions

As shown in the Key Security section, secure your private key:

```
$keyPath = "$Home\.ssh\id_ed25519"
icaccls $keyPath /inheritance:r
$currentUser =
[System.Security.Principal.WindowsIdentity]::GetCurrent().Name
icaccls $keyPath /grant:r "${currentUser}:F"
```

Step 5: Connect directly

Now, connect securely to your private staging servers via PowerShell:

```
ssh aws-private-10.0.10.15
```

Method B: Using PuTTY (GUI-Based SSH Client)

If you prefer using PuTTY, follow these steps to set up secure tunneling.

Step 1: Generate Key using PuTTYgen

- . Download and open **PuTTYgen**.
- . Select **Ed25519** (or RSA 4096) under "Parameters -> Type of key to generate".
- . Click **Generate** and move your mouse randomly over the blank area.
- . Enter a strong **Key passphrase**.
- . Click **Save private key** and save it as `id_ed25519.ppk` inside a secure folder (e.g., `C:\Users\username\.ssh\`).
- . Copy the text from the top box ("Public key for pasting into OpenSSH authorized_keys file") and send it to your AWS administrator.

Step 2: Secure PPK file permissions

- . Right-click your `id_ed25519.ppk` file and click **Properties**.
- . Go to **Security -> Advanced**.

- . Click **Disable inheritance** -> select "Remove all inherited permissions".
- . Add your personal Windows user principal with **Full Control**.

Step 3: Configure PuTTY for Proxy SSH Session (ProxyCommand/plink)

To establish multi-hop SSH tunnels in PuTTY:

- . Open **PuTTY**.
- . In the **Session** category (on the left):
 - **Host Name (or IP address)**: Type the private IP of the target server (e.g., 10.0.10.15).
 - **Port**: 22 .
- . In the **Connection** -> **SSH** -> **Auth** -> **Credentials** category:
 - Browse and select your id_ed25519.ppk private key.
- . In the **Connection** -> **Proxy** category:
 - **Proxy type**: Select **Local**.
 - **Proxy hostname**: Type the Jumhost's Elastic IP (e.g., 54.x.y.z).
 - **Port**: 22 .
 - **Telnet command, or local proxy command**: Type the following (replace paths with your actual PuTTY installation and key paths):

```
plink -batch -v ubuntu@%proxyhost -i
"C:\Users\YOUR_USERNAME\.ssh\id_ed25519.ppk" -nc %host:%port
```

- . Return to the **Session** category:
 - Type a name under "Saved Sessions" (e.g., AWS-Private-10.0.10.15).
 - Click **Save**.
- . Click **Open**. PuTTY will connect to the Jumhost first, proxy the connection, and open the private server terminal transparently!

Method C: Alternatively, using Port Forwarding (SSH Tunnel) via PuTTY

If you prefer standard SSH Port Forwarding:

- . Open **PuTTY** and configure a connection to the Jumphost:
 - **Host Name (or IP address):** Jumphost Elastic IP (54.x.y.z).
 - **User:** ubuntu .
- . Go to **Connection -> SSH -> Tunnels**.
- . In "Add new forwarded port":
 - **Source port:** Type a local random port (e.g., 9022).
 - **Destination:** Type the private IP of your staging server and SSH port (e.g., 10.0.10.15:22).
 - Click **Add**.
- . Go back to the **Session** category, save it as AWS-Jumphost-Tunnel, and click **Open** (enter key passphrase when prompted).
- . Open a **second** PuTTY window:
 - **Host Name (or IP address):** 127.0.0.1 .
 - **Port:** 9022 .
 - Select your private key under **Connection -> SSH -> Auth -> Credentials** and open!

Hybrid Cloud Integration & Costing Guide

#aws

#cloud

#architecture

#vpc

#alb

#waf

#disaster-recovery

#postgresql

#costing

Hybrid Cloud Integration Guide: AWS & On-Premises

1. Executive Summary & Context

To support high-performance operations while maintaining strict cost-efficiency, organizations must determine the most practical strategy for connecting on-premises infrastructure (such as local datacenters, legacy compute, or private office environments) with an AWS-native deployment.


While AWS provides several enterprise-grade physical and virtual private network connectivity options, their upfront, monthly baseline, and data processing charges can be prohibitive for budget-conscious initiatives. Consequently, this guide presents a dual-layered architectural paradigm:

- . **Low-Cost, High-Flexibility Pathways (API and MCP):** The primary recommended choices for agile integrations, utilizing secure HTTP RESTful architectures and the newly released **Model Context Protocol (MCP)** for AI agentic workflows.
- . **Official AWS Hybrid Options (VPN, Direct Connect, and Transit Gateway):** Enterprise-standard private networking pipelines offered for comparative purposes to facilitate corporate security reviews and management approval.

All cost models are calibrated for the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)** and estimated in both **United States Dollars (USD)** and **Malaysian Ringgit (MYR)**, utilizing a baseline conversion exchange rate of **1 USD $\approx$ 4.50 MYR**.

2. Low-Cost Option 1: Secure API-Based Connections

API-based hybrid integration leverages the public internet to transmit request/response payloads between AWS compute nodes (such as our Auto Scaling Groups or Standalone EC2 instances) and on-premises service endpoints. Rather than establishing a persistent, protocol-level network tunnel, communication is transactional, synchronous, or event-driven.



Architectural Implementation & Security

- **Authentication & Authorization:** Connections are secured via OAuth 2.0 Bearer tokens, Amazon Cognito, API Keys, or **Mutual TLS (mTLS)** supported natively by Amazon API Gateway.
- **Network Security & Whitelisting:** Downstream on-premises endpoints are configured to whitelist only the Elastic IP addresses of the AWS NAT Gateway. Inbound AWS traffic to the Application Load Balancer (ALB) is whitelisted at the AWS WAFv2 level or restricted to on-premises gateway IPs.
- **Protocol Support:** Standard REST (HTTPS/JSON) for point-to-point RPCs, or WebSockets (WSS) via Amazon API Gateway for persistent, low-latency bi-directional event streaming.

Financial Breakdown


API-based connections have **no provisioned hourly costs** for physical line terminations. The billing is entirely pay-as-you-go based on data processing and network transit:

- **Amazon API Gateway REST APIs:** \$3.50 per million requests (up to 333 million, with sliding volume discounts).
 - **AWS NAT Gateway Charges:** \$0.045 per hour base charge (\$32.85/month) + \$0.045 per GB of outbound data processed.
 - **AWS Outbound Data Transfer (Internet Egress):** The first 100 GB per month is free. Beyond that, data transferred from AWS to the internet in Malaysia (ap-southeast-5) is charged at **\$0.09 per GB**. On-premises data transfer into AWS is entirely free.
-

3. Low-Cost Option 2: AI-Native MCP-Based Connections

The **Model Context Protocol (MCP)** is an open-standard protocol designed to link Large Language Models (LLMs) securely with external data sources, content repositories, and enterprise applications. Instead of building bespoke API integrations for AI tools, MCP standardizes how AI agents discover and execute tools or fetch contextual files.

In December 2025, AWS introduced native **MCP Proxy Support in Amazon API Gateway**, integrated directly with **Amazon Bedrock AgentCore**. This allows organizations to transform on-premises applications, custom monitoring systems, or local databases into MCP-compatible endpoints without changing existing codebase patterns.



Architectural Implementation & Security

- **Protocol Translation:** On-premises systems host a lightweight “MCP Server” (e.g., exposing local postgres, CSV data, or SSH scripts). Amazon API Gateway’s MCP Proxy translates incoming JSON-RPC 2.0 protocol payloads from AWS Bedrock Agents into standard secure HTTP requests that your local server processes.
- **Enterprise Security & Gateway Services:** Dual authentication is applied:
 1. **Inbound Verification:** Bedrock AgentCore validates the AI agent’s credentials and authorization scopes before initiating requests.
 2. **Outbound Verification:** API Gateway manages secure connections, client certificates (mTLS), and API keys to call the on-premises REST endpoints.
- **Semantic Discovery:** Bedrock AgentCore Gateway allows AI agents to dynamically search and select the most relevant REST APIs that best match the context of the user’s prompt.

Financial Breakdown


- **Amazon Bedrock AgentCore Pricing:** Pay-as-you-go. (Refer to the official Amazon Bedrock AgentCore pricing page for updated

regional rates).

- **API Gateway MCP Proxy Charges:** Treated as standard API Gateway REST API calls (\$3.50 per million requests) with no additional platform premium.
- **Network Egress:** Standard AWS Data Transfer Out (\$0.09/GB for ap-southeast-5 beyond the 100 GB free tier).

4. Official AWS Enterprise Hybrid Connections

For organizations requiring dedicated, low-latency, or highly secure private networking that completely bypasses the public internet, AWS offers three primary structural options.



Option A: AWS Site-to-Site VPN

- **Mechanism:** Establishes encrypted IPsec tunnels over the public internet between an AWS Virtual Private Gateway (VGW) / Transit Gateway and an on-premises firewall/router (Customer Gateway).
- **Bandwidth:** Up to 1.25 Gbps per active tunnel (with ECMP support to group multiple tunnels for higher throughput).
- **Costing:**

- **AWS Connection Fee: \$0.05 per hour** per active VPN connection (\$36.50 per month per site).
- **On-Premises Hardware:** Requires a physical VPN appliance (Cisco, Juniper, Fortinet) or open-source software (StrongSwan, pfSense) on-premises.
- **Data Transfer Out:** Charged at standard AWS Data Transfer Out rates (\$0.09 per GB in ap-southeast-5).

Option B: AWS Direct Connect (DX)

- **Mechanism:** A physical, dedicated, high-speed fiber-optic connection from an on-premises datacenter to an AWS Direct Connect Location. Complete bypass of the public internet.
- **Bandwidth:** Fixed port speeds of 1 Gbps, 10 Gbps, or 100 Gbps (Sub-1G connections are available via AWS Direct Connect Partners).
- **Costing:**
 - **AWS Port Hour Charge: \$0.03 per hour** for 1 Gbps ports (~\$21.90/mo) or **\$2.25 per hour** for 100 Gbps ports (~\$1,642.50/mo).
 - **Data Transfer Out:** Significantly lower than standard internet egress. Direct Connect Data Transfer Out is charged at a highly discounted rate of **\$0.021 to \$0.041 per GB** (region-dependent).
 - **Partner Telecommunication Costs:** Third-party telco providers charge monthly circuit fees to run fiber from your office to the AWS Direct Connect Point of Presence (PoP) (ranging from \$500 to \$5,000+ per month).

Option C: AWS Transit Gateway (TGW)

- **Mechanism:** Acts as a centralized cloud router that simplifies hybrid network topology by interconnecting multiple AWS VPCs, Site-to-Site VPNs, and Direct Connect links through a single logical gateway.
- **Costing:**
 - **AWS Attachment Charge: \$0.05 per hour** per VPC, VPN, or DX attachment (~\$36.50 per month per attachment).

- **Data Processing Fee: \$0.02 per GB** for all data passing through the Transit Gateway.

5. Side-by-Side Hybrid Technology Matrix

Evaluation Criteria	API-Based Connections	MCP-Based Connections	AWS Site-to-Site VPN	AWS Direct Connections
Primary Use Case	Lightweight microservices, Webhooks, database queries	AI Agents, tool calling, Bedrock LLM retrieval	General corporate network extension, secure admin	High-performance bulk transfers, DBs
Underlying Network	Public Internet (Secure TLS)	Public Internet (Secure TLS / mTLS)	Encrypted IPsec Tunnel over Public Internet	Dedicated private physical optic
Max Throughput	Limited only by internet connection	Limited only by internet connection	1.25 Gbps per active tunnel	1 Gbps Gbps
Latency Profile	Variable (reliant on internet ISP)	Variable (reliant on internet ISP)	Stable, moderate (internet routing)	Ultra-deterministic and S
Network Security	Application Layer (TLS, JWT, Mutual TLS)	Application Layer (mTLS, IAM, API Key, Bedrock Core)	Network Layer (IPsec VPN encryption)	Physical isolation (Direct MACsec)
Integration Complexity	Low (Standard HTTP development)	Low-to-Medium (REST config, Bedrock alignment)	Medium (Network routing, IPsec configuration)	High telco physical provisioning

Evaluation Criteria	API-Based Connections	MCP-Based Connections	AWS Site-to-Site VPN	AWS Direct Connections
Implementation Lead Time	Instant	Instant	Hours	30 to 60 days (Telco pull)
AWS Billing Model	Pay-as-you-go (\$3.50/M requests + DTO)	Pay-as-you-go + standard API gateway rates	Hourly (\$0.05/hr) + Data Transfer Out	Hourly (\$0.05/hr) + Parallel DTO
Upfront Capital Cost	\$0.00	\$0.00	\$0.00 (using existing router)	\$1,000 (Hardware cross-connect)

6. Granular Monthly Hybrid Cost Estimations

To assist leadership in budget selection, we model monthly costs based on two enterprise traffic volumes:

- **Low-Volume Integration (10 GB Outbound Data Transfer + 1 Million Requests / Month)**
- **High-Volume Integration (500 GB Outbound Data Transfer + 10 Million Requests / Month)**

All estimates assume deployment in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)** with standard regional rates.

Scenario A: Low-Volume Integration Cost Modeling (10 GB Outbound / 1M Requests)

This represents an agile staging environment, internal dev tools, or an experimental AI agentic system integrating on-premises files.

Low-Volume Scenario Total Monthly Cost (USD)		
API-Based		\$3.50
MCP-Based		\$3.50
S2S VPN		\$36.50
Direct Conn.		\$521.90

1. API-Based Connections

- **API Gateway REST Processing:** 1,000,000 requests * \$3.50/M = \$3.50 USD
- **AWS Data Transfer Out (Egress):** 10 GB (Within free tier threshold of 100 GB) = \$0.00 USD
- **NAT Gateway Hourly Fee (Existing):** Shared with existing infrastructure = \$0.00 USD
- **TOTAL MONTHLY COST: \$3.50 USD / ~RM 15.75 MYR**

2. MCP-Based Connections

- **API Gateway with MCP Proxy:** 1,000,000 requests * \$3.50/M = \$3.50 USD
- **AWS Data Transfer Out (Egress):** 10 GB (Within free tier threshold of 100 GB) = \$0.00 USD
- **Bedrock AgentCore Platform Fee:** Pay-as-you-go nominal charges = \$0.00 USD (Approx)
- **TOTAL MONTHLY COST: \$3.50 USD / ~RM 15.75 MYR**

3. AWS Site-to-Site VPN

- **AWS VPN Connection Charge:** 1 tunnel * 730 hours/month * \$0.05/hr = \$36.50 USD
- **AWS Data Transfer Out (Egress):** 10 GB (Within free tier threshold of 100 GB) = \$0.00 USD

- **On-Premises Router Power/Compute:** Nominal standard overhead = \$0.00 USD
- **TOTAL MONTHLY COST:** \$36.50 USD / ~RM 164.25 MYR

4. AWS Direct Connect (DX)

- **AWS 1 Gbps Port Charge:** 730 hours/month * \$0.03/hr = \$21.90 USD
- **Discounted Data Transfer Out (Egress):** 10 GB * \$0.041/GB = \$0.41 USD
- **Partner Local Loop Circuit:** Sub-1G Partner line (Est. Baseline Partner Fee) = \$500.00 USD
- **TOTAL MONTHLY COST:** \$522.31 USD / ~RM 2,350.40 MYR

Scenario B: High-Volume Integration Cost Modeling (500 GB Outbound / 10M Requests)

This represents active production workloads, automated ETL syncs, enterprise-wide chat agents, or heavy file replication.

High-Volume Scenario Total Monthly Cost (USD)	
API-Based	\$71.00
MCP-Based	\$71.00
S2S VPN	\$72.50
Direct Conn.	\$542.40

1. API-Based Connections

- **API Gateway REST Processing:** 10,000,000 requests * \$3.50/M = \$35.00 USD
- **AWS Data Transfer Out (Egress):** (500 GB - 100 GB Free) = 400 GB * \$0.09/GB = \$36.00 USD
- **TOTAL MONTHLY COST:** \$71.00 USD / ~RM 319.50 MYR

2. MCP-Based Connections

- **API Gateway with MCP Proxy:** $10,000,000 \text{ requests} * \$3.50/\text{M} = \$35.00 \text{ USD}$
- **AWS Data Transfer Out (Egress):** $(500 \text{ GB} - 100 \text{ GB Free}) = 400 \text{ GB} * \$0.09/\text{GB} = \$36.00 \text{ USD}$
- **Bedrock AgentCore Platform Fee:** Pay-as-you-go nominal charges = \$0.00 USD (Approx)
- **TOTAL MONTHLY COST: \$71.00 USD / ~RM 319.50 MYR**

3. AWS Site-to-Site VPN


- **AWS VPN Connection Charge:** $1 \text{ tunnel} * 730 \text{ hours/month} * \$0.05/\text{hr} = \$36.50 \text{ USD}$
- **AWS Data Transfer Out (Egress):** $(500 \text{ GB} - 100 \text{ GB Free}) = 400 \text{ GB} * \$0.09/\text{GB} = \$36.00 \text{ USD}$
- **TOTAL MONTHLY COST: \$72.50 USD / ~RM 326.25 MYR**

4. AWS Direct Connect (DX)

- **AWS 1 Gbps Port Charge:** $730 \text{ hours/month} * \$0.03/\text{hr} = \$21.90 \text{ USD}$
- **Discounted Data Transfer Out (Egress):** $500 \text{ GB} * \$0.041/\text{GB} = \20.50 USD
- **Partner Local Loop Circuit:** Sub-1G Partner line (Est. Baseline Partner Fee) = \$500.00 USD
- **TOTAL MONTHLY COST: \$542.40 USD / ~RM 2,440.80 MYR**

7. Strategic Recommendation & Management Decision Tree

To guide management through the selection process, we present a logical decision matrix to streamline technical approval:



Actionable Roadmap for Approvals

. Phase 1: Zero-Upfront Trial (API / MCP)

- **Recommendation:** Begin with **API-based or MCP-based connections**. It requires \$0 in capital expenditure, zero hardware setups on-premises, and allows the AI team to immediately utilize Bedrock AgentCore and local datasets.
- **Security Action:** Restrict access to the on-premise MCP server by whitelisting the AWS NAT Gateway Elastic IP and enforcing API key rotation.

. Phase 2: Hybrid Network Scaling (Site-to-Site VPN)

- **Trigger:** If corporate security mandates network-level segmentation (e.g., preventing any endpoint from exposing port

443 to the public internet, even with whitelists), transition to **AWS Site-to-Site VPN**.

- **Cost Impact:** Adds a modest **\$36.50/month (RM 164.25)** flat connection charge plus standard outbound data transfer.

. **Phase 3: Ultimate Performance (Direct Connect)**

- **Trigger:** If synchronization volumes exceed **10 TB / month** or latency constraints require single-digit millisecond response times.
- **Cost Impact:** Heavy capital and partner operating expenditures (RM 2,400+ / mo). Highly discounted AWS egress rates make Direct Connect financially viable only at extreme data transfer volumes.

Disaster Recovery Options & Sovereignty Guide

#aws #cloud #architecture #vpc #alb #asg #rds #waf #elasticache #valkey
#dns #disaster-recovery #efs #postgresql #sovereignty #compliance

Disaster Recovery (DR) Options & National Sovereignty Guide

1. Executive Summary & Context

Maintaining system resilience is critical for enterprise deployment. This guide establishes a comprehensive, production-ready Disaster Recovery (DR) playbook modeled directly after the official AWS whitepaper ["Disaster Recovery of Workloads on AWS: Recovery in the Cloud"](#).

This guide addresses three core dimensions:

- . **Existing Architecture Alignment:** Adapting DR practices directly to our Multi-AZ, 3-Tier layout (WAFv2, ALB, private Auto Scaling Groups for Frontend, Backend, and AI tiers, Standalone EC2 staging environments, Multi-AZ RDS PostgreSQL 16 database, Amazon ElastiCache for Valkey, and Amazon EFS shared storage).
- . **Malaysia National Sovereignty:** A strict regulatory compliance review under the **Malaysian Personal Data Protection Act (PDPA) 2010** (including the **Personal Data Protection (Amendment) Act 2024** and the **2025 Cross-Border Personal Data Transfer Guidelines (CBPDT)**).
- . **Financial Transparency (USD / MYR):** Granular cost breakdowns comparing **In-Region Multi-AZ DR** (sovereign, 100% data residency

in Malaysia ap-southeast-5) and **Cross-Region DR** (replicated to Singapore ap-southeast-1 or Jakarta ap-southeast-3).

All currency estimations use a baseline baseline conversion rate of **1 USD ≈ 4.50 MYR**.

2. Malaysia National Sovereignty & Regulatory Compliance

When designing DR topologies for applications processing citizen data in Malaysia, **data sovereignty** and **jurisdictional boundaries** are primary architectural drivers.

Data Residency & Transfer Compliance Pathways	
In-Region DR (ap-southeast-5)	Cross-Region DR (Out of MY)
• 100% Malaysian Sovereignty • No Cross-Border transfer concerns • Fully compliant with 2025 CBPDT • Complete immunity to foreign laws	• Subject to PDPA Section 129 • Requires Transfer Impact (TIA) • Explicit Data Subject Consent • Requires Contractual Clauses

2.1 The Personal Data Protection Act (PDPA) & 2024/2025 Frameworks

- **Section 129 Principal Prohibition:** Under the Malaysian PDPA 2010, transferring personal data outside of Malaysia is prohibited unless specifically exempted by the Minister or falling under specific statutory exceptions.
- **The 2024 Amendments:** The Personal Data Protection (Amendment) Act 2024 introduced mandatory **Data Breach Notifications (DBN)** and the compulsory appointment of a **Data Protection Officer (DPO)**, escalating the legal risk of data exposure.

- **The 2025 CBPDT Guidelines:** Published in April 2025 by the Personal Data Protection Commissioner, the **Cross Border Personal Data Transfer (CBPDT) Guidelines** clarify that data controllers transferring data outside Malaysia must conduct a **Transfer Impact Assessment (TIA)** to ensure the destination provides “adequate protection” substantially similar to the PDPA.

2.2 Sovereign In-Region Multi-AZ DR vs. Cross-Region DR

Sovereign In-Region Multi-AZ DR (Absolute Residency)

- **Architectural Flow:** All primary workloads and disaster recovery assets reside entirely within the three physical Availability Zones of the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)**.
- **PDPA Standing:** Fully compliant out of the box. Because data never crosses Malaysia’s geographical borders, it is **exempt** from Section 129 restrictions, TIAs, or extra consent requirements.
- **Jurisdictional Immunity:** Data remains protected exclusively under Malaysian courts. It is immune to extraterritorial access requests (such as the US CLOUD Act or Singaporean administrative warrants) that might apply if the data was replicated overseas.

Cross-Region DR (Jurisdictional Border Crossing)

- **Architectural Flow:** Backups, read-replicas, or standby environments are copied from ap-southeast-5 to **Singapore (ap-southeast-1)** or **Jakarta (ap-southeast-3)**.
- **PDPA Section 129 Compliance Checklist:** To legally implement Cross-Region replication, the organization must fulfill these strict legal requirements:
 1. **Explicit Data Subject Consent:** Users must actively opt-in to their data being transferred and stored in Singapore/Jakarta, detailed within a clear, updated Privacy Notice.
 2. **Transfer Impact Assessment (TIA):** The DPO must perform and log a formal TIA assessing the recipient country’s data laws (e.g., Singapore’s PDPA 2012 or Indonesia’s UU PDP).

3. **Data Processor Agreements:** Binding contracts with AWS and any third-party processors enforcing technical and organizational safety standards equivalent to Malaysian standard standards.

2.3 Architectural & Engineering Approach to Sovereignty in Cross-Region DR

To mitigate regulatory risks and address national data sovereignty concerns when using Cross-Region DR (e.g. replicating data from Malaysia ap-southeast-5 to Singapore ap-southeast-1 or Jakarta ap-southeast-3), the architectural layout implements several strict technical safeguards:

1. Field-Level Encryption & Tokenization

- **Strategy:** Before any database replication payload, block device copy, or backup snapshot is transmitted across jurisdictional boundaries, all highly sensitive Personally Identifiable Information (PII) – such as national identification numbers, phone numbers, and raw user credentials – is subjected to field-level encryption or tokenization inside the primary Malaysia VPC.
- **Result:** Only fully encrypted, non-readable cyphertext or synthetic tokens are transmitted across borders, while the lookup table and decryption logic remain strictly localized inside ap-southeast-5. Even if foreign authorities subpoena the storage drives in Singapore or Jakarta, they cannot access the actual plaintext citizen identities.

2. Sovereign Key Management & Cryptographic Isolation

- **Strategy:** All AWS KMS (Key Management Service) Customer Managed Keys (CMKs) or Hardware Security Modules (HSMs) used to encrypt cross-region backups and databases are hosted and controlled exclusively inside the Malaysia (ap-southeast-5) region.
- **Result:** The decryption keys are never replicated to Singapore or Jakarta. To boot recovery EC2 instances or decrypt RDS backups during a disaster, foreign hypervisors must dynamically call the Malaysia KMS API. In a geopolitical crisis, the security team can

immediately revoke KMS key access from the primary region, completely neutralizing the foreign environment and rendering the replicated data permanently unreadable.

3. Sovereign Failover Guardrails & Circuit Breakers

- **Strategy:** The Route 53 routing policies and ASG launch parameters are locked behind multi-signature IAM permission controls and active “Sovereignty Circuit Breakers.” Automated failover to the cross-region standby only triggers if the entire Malaysian AWS region is physically unavailable (e.g. undersea cable cuts or natural disasters), not during minor localized software glitches.
- **Result:** This prevents accidental data migration or “regulatory leakage” where application traffic is unnecessarily routed to overseas instances during routine operating conditions.

4. Legal & Compliance Alignment: Transfer Impact Assessments (TIA)

- **Strategy:** Under Section 129 of the PDPA 2010 and the 2025 CBPDT Guidelines, any cross-region flow requires a logged TIA detailing:
 - Technical measures (e.g. KMS, Field-Level Encryption) ensuring identical security.
 - Recipient-region legal framework analysis (e.g. comparing Malaysia’s PDPA with Singapore’s PDPA 2012 or Indonesia’s UU PDP).
 - Explicit Standard Contractual Clauses (SCCs) embedded into the Cloud Service Provider agreement to legally bind AWS to Malaysian-standard data preservation.

3. Disaster Recovery Spectrum Analysis


We analyze the four standard AWS Disaster Recovery options, alongside AWS Elastic Disaster Recovery (AWS DRS) as a modern, cost-optimized, and near-zero RPO hybrid strategy. Each strategy presents a distinct balance of Recovery Point Objective (RPO),

Recovery Time Objective (RTO), implementation complexity, and ongoing operating cost.

Low Cost		
High RPO/RTO		
Backup & Restore	Pilot Light	Warm Standby

Strategy A: Backup and Restore

A passive approach where virtual machine images, database snapshots, and configuration files are regularly backed up and restored to a target environment only after an outage is declared.



1. Architectural Setup & Component Layout

- **Database (RDS):** Automated daily snapshots and transaction logs stored via AWS Backup. In-Region copies are kept in a separate secure AWS account; Cross-Region copies sync to Singapore (ap-southeast-1) S3 buckets.
- **Compute (ASG):** Launch Templates are backed up. AMIs of the Frontend, Backend, and AI Standalone/ASG nodes are pre-baked using our Packer/Ansible pipeline and stored in local or remote AMI catalogs.
- **Shared Storage (EFS):** Daily EFS backups handled via AWS Backup. EFS filesystems are not pre-provisioned in the target environment to save costs; they are recreated during recovery.

- **DNS & Routing (Route 53):** Route 53 pointing to the primary ALB. No active health check routing is configured. Manual DNS update is required during a failover event.
- **Valkey Cache:** No replication. Caching nodes are left unprovisioned in the target area and cold-booted during restoration.

2. Recovery Objectives

- **Recovery Point Objective (RPO): 24 Hours** (determined by backup execution intervals).
- **Recovery Time Objective (RTO): 4 to 8 Hours** (time needed to provision network infrastructure, restore RDS database storage, spin up EC2 instances from AMIs, and update DNS).

3. Trade-offs & Analysis

- **Cost:** Extremely low. No active compute, load balancer, or database instances run in the DR zone.
- **Complexity:** Low configuration complexity, but high manual recovery complexity during an incident.
- **Sovereignty:** If utilizing Cross-Region Backup and Restore, encrypted database snapshots reside in Singapore, which mandates standard PDPA Section 129 consent and TIAs. In-Region backups are completely sovereign.

Strategy B: Pilot Light

A core database and persistent storage layer are kept actively running and synchronized, while the compute/application server layer remains unprovisioned or “turned off” until a disaster occurs.



1. Architectural Setup & Component Layout

- **Database (RDS):** A minimal, single-AZ RDS PostgreSQL read-replica is kept online and continuously synchronized via asynchronous replication.
- **Compute (ASG):** Launch Templates are active in the target VPC, but the ASG **Desired Capacity is set to 0**. Standalone staging environments are not provisioned.
- **Shared Storage (EFS):** Continuous replication to a target EFS filesystem using AWS EFS Replication.
- **DNS & Routing:** Route 53 DNS is configured with active health checks. However, failover requires changing the ASG desired capacity to active levels and waiting for compute nodes to bootstrap.
- **Valkey Cache:** Unprovisioned. Provisioned only during failover.

2. Recovery Objectives

- **Recovery Point Objective (RPO):** **< 5 Minutes** (determined by RDS replica lag and EFS replication latency).
- **Recovery Time Objective (RTO):** **30 to 45 Minutes** (time to promote RDS replica to primary, scale the ASG up from 0 to target size, wait for application health checks to pass, and update Route 53).


3. Trade-offs & Analysis

- **Cost:** Moderate. You pay for continuous RDS replica compute and persistent EFS/S3 storage, but avoid active EC2 compute costs.
- **Complexity:** Medium. Requires automating the promotion of the database and the auto-scaling group scale-up.

- **Sovereignty:** Cross-Region replica databases contain active, unencrypted-in-memory data in Singapore/Jakarta. This triggers strict PDPA compliance, requiring a comprehensive TIA and explicit user consent. In-Region Pilot Light (using a separate VPC in the same region) is fully sovereign.

Strategy C: Warm Standby

A scaled-down but fully functional copy of the entire 3-tier architecture is kept running continuously in the target zone. It handles zero (or minimal) active traffic under normal operations but can instantly scale up to handle the production load if the primary site fails.



1. Architectural Setup & Component Layout

- **Presentation Tier:** A standby ALB is active. AWS WAFv2 is associated and running.
- **Compute Tier (ASG):** The standby ASG runs continuously with a minimal capacity (e.g., 1x t4g.medium instead of the production 2x t4g.xlarge). Standalone staging instances are unprovisioned to optimize cost.
- **Database (RDS):** An active RDS PostgreSQL read-replica runs continuously in the standby zone.
- **Valkey Cache:** A minimal single-node ElastiCache Valkey (cache.t4g.micro) runs to maintain cache readiness.
- **Shared Storage (EFS):** Continuous EFS Replication keeps model weights and configurations in sync.
- **DNS & Routing:** Route 53 Active-Passive failover routing is configured. Health checks automatically redirect traffic to the

standby ALB if the primary fails.

2. Recovery Objectives


- **Recovery Point Objective (RPO): < 1 Minute** (near-real-time database and storage replication).
- **Recovery Time Objective (RTO): 5 to 10 Minutes** (traffic is routed instantly; auto-scaling takes a few minutes to scale compute instances to full production sizing).

3. Trade-offs & Analysis

- **Cost:** High. Requires running a second ALB, active EC2 instances, replica RDS instances, EFS replication, and an ElastiCache node.
- **Complexity:** Medium-High. Requires robust automation for target scaling policies and failover signaling.
- **Sovereignty:** Storing live, unencrypted client data continuously on active databases in foreign regions requires strict PDPA Section 129 audit trails and TIAs.

Strategy D: Multi-Site Active-Active

Traffic is split dynamically across two identical, full-scale production environments running simultaneously in two separate zones. Both sites actively serve client requests.



1. Architectural Setup & Component Layout

- **Presentation Tier:** Active ALBs and WAFv2 are deployed in both sites.
- **Compute Tier (ASG):** Full-scale ASGs (2x t4g.xlarge nodes) run actively in both regions. All Standalone AMI-baking instances are kept in sync.
- **Database (RDS):** Configured as a cross-region active-active database cluster (utilizing Aurora Global Database or custom bi-directional replication engines), or active-passive database with write-routing to the primary site.
- **Valkey Cache:** Independent active Valkey clusters in each region.
- **Shared Storage (EFS):** Continuous bi-directional synchronization or dual EFS staging setups.
- **DNS & Routing:** Route 53 is configured with Geolocation or Latency-based routing to dynamically distribute user requests.

2. Recovery Objectives


- **Recovery Point Objective (RPO):** Near Zero (Real-time).
- **Recovery Time Objective (RTO):** Near Zero (instantaneous failover as both environments are already serving live traffic).

3. Trade-offs & Analysis

- **Cost:** Extremely High. Double the baseline production cost plus significant cross-region data transfer replication fees.
- **Complexity:** Extremely High. Requires solving split-brain scenarios, database write collisions, and high data consistency synchronization challenges.
- **Sovereignty:** Deeply critical. 50% of Malaysian citizen data is processed and stored outside local jurisdiction in real time. This triggers comprehensive PDPA scrutiny, making TIAs, explicit consent notices, and regulatory registration mandatory.

Strategy E: AWS Elastic Disaster Recovery (AWS DRS)

AWS Elastic Disaster Recovery (AWS DRS) minimizes downtime and data loss with fast, reliable recovery of on-premises and cloud-based applications using affordable storage, minimal compute, and point-in-time recovery. It provides sub-minute RPO and minutes-level RTO.



1. Architectural Setup & Component Layout

- **Source Servers:** Install the AWS Elastic Disaster Recovery Replication Agent on target source servers (on-premises or remote cloud nodes) to initiate secure, block-level data replication.
- **Staging Area Subnet:** Data is continuously replicated to a dedicated, low-cost staging area subnet in the designated AWS account and Region (e.g., ap-southeast-5).
- **Compute & Storage Optimization:** Staging costs are highly minimized. Lightweight, automatically managed EC2 replication instances (e.g., t3.small nodes) handle incoming blocks, while affordable EBS volumes (gp3 or low-cost sc1) act as the replication target disks.
- **Non-Disruptive Testing:** Allows running seamless, non-disruptive disaster recovery drills and tests at any time without impacting active replication or source servers.
- **Drill/Recovery Launches:** When a failover or drill is initiated, AWS DRS automatically launches target EC2 instances based on the most up-to-date server state or a specified historical point-in-time state.
- **Failback Replication:** After the primary site issue is resolved, data replication is initiated in reverse back to the primary

site, ensuring seamless failback.

2. Recovery Objectives

- **Recovery Point Objective (RPO): Seconds to Sub-Minute** (continuous, real-time block-level asynchronous replication).
- **Recovery Time Objective (RTO): Minutes** (automated launch, conversion, and orchestration of recovery instances).

3. Trade-offs & Analysis

- **Cost:** Highly Cost-Effective. Replicating servers cost only a nominal hourly software license fee (\$0.028 per server per hour) plus low-cost staging compute (t3.small) and staging EBS volumes, instead of running full warm standby or active-active duplicate environments.
- **Complexity:** Medium. Requires installing replication agents on source servers and configuring launch templates, but failover and failback orchestration are highly automated.
- **Sovereignty:** 100% Sovereign when staging and recovery are targeted inside the **AWS Malaysia Region (ap-southeast-5)**. Because the continuous block replication targets local sovereign physical boundaries, it completely satisfies PDPA and 2025 CBPDT requirements for localized residency, avoiding complex cross-border Transfer Impact Assessments (TIAs).

4. Comprehensive Cost Estimation Model (USD & MYR)

To provide clear financial visibility, we present granular monthly costs for all four DR strategies, comparing:

- . **Sovereign In-Region Multi-AZ DR** (All standby resources deployed inside ap-southeast-5 using a separate DR VPC).
- . **Cross-Region DR** (Standby resources deployed in Singapore ap-southeast-1).

All prices are calibrated for On-Demand rates (~730 hours/month) for our **High-Performance Production Stack**.

4.1 Granular Monthly DR Cost Matrix (USD / Month)

Component	Strategy A: Backup & Restore	Strategy B: Pilot Light	Strategy C: Warm Standby	Strategy D: Active-Active
Networking & NAT	\$0.00	\$0.00 (No outbound needed)	\$32.85 (1x NAT Gateway)	\$32.85 (1x Gateway)
Public IPv4 Addr	\$0.00	\$0.00	\$7.30 (1x NAT IP, 1x ALB IP)	\$10.95 (1x 2x ALB IPs)
ALB & WAFv2	\$0.00	\$0.00	\$25.03 (ALB + WAF ACL)	\$41.46 (ALB + Rules)
Compute (ASG)	\$0.00 (Desired = 0)	\$0.00 (Desired = 0)	\$49.06 (1x t4g.medium warm)	\$196.22 (2x t4g.xlarge active)
Database (RDS)	\$0.00	\$221.92 (1x db.m6g.large replica)	\$443.84 (Multi-AZ replica)	\$443.84 (Multi-AZ active)
Valkey Cache	\$0.00	\$0.00	\$9.34 (cache.t4g.micro)	\$39.71 (cache.t4g.micro)
Storage (EFS)	\$0.00	\$15.00 (50GB replica)	\$15.00 (50GB replica)	\$15.00 (50GB active)
Storage (S3)	\$2.30 (100GB backup)	\$2.30 (100GB backup)	\$2.30 (100GB backup)	\$2.30 (100GB active)
DRS Service Fee	\$0.00	\$0.00	\$0.00	\$0.00

Component	Strategy A: Backup & Restore	Strategy B: Pilot Light	Strategy C: Warm Standby	Strategy D: Active-Active
Data Replication	\$5.00 (AWS Backup fees)	\$15.00 (RDS & EFS sync)	\$30.00 (Live RDS & EFS sync)	\$120.00 (Hadoop sync)
TOTAL (In-Region)	\$7.30 USD	\$254.22 USD	\$614.42 USD	\$902.33 USD
TOTAL (Cross-Region)	\$11.50 USD	\$289.50 USD	\$664.80 USD	\$982.50 USD

Note: Cross-Region costs are slightly higher due to AWS Cross-Region Data Transfer Egress fees (\$0.09/GB from ap-southeast-5) and marginally higher base pricing in ap-southeast-1. AWS DRS cost modeling assumes replicating 3 active EC2 compute nodes (Frontend, Backend, AI Tier) into a secure staging subnet.

4.2 Local Currency Equivalent Comparison (MYR / Month)

Using the exchange baseline of **1 USD ≈ 4.50 MYR**, the monthly DR infrastructure spend equates to:

Monthly DR Cost Comparison (MYR equivalent)		
In-Region (ap-southeast-5)		
└─ Backup & Restore	⇒ RM 32.85	
└─ AWS DRS	⇒ RM 505.35	(Highly Recommended)
└─ Pilot Light	⇒ RM 1,143.99	
└─ Warm Standby	⇒ RM 2,764.89	
└─ Active-Active	⇒ RM 4,060.49	
Cross-Region (Singapore/Jakarta)		
└─ Backup & Restore	⇒ RM 51.75	
└─ AWS DRS	⇒ RM 609.75	
└─ Pilot Light	⇒ RM 1,302.75	
└─ Warm Standby	⇒ RM 2,991.60	
└─ Active-Active	⇒ RM 4,421.25	

5. Decision Matrix & Strategy Trade-offs


To guide executive management through the selection process, we present a technical evaluation matrix:

Evaluation Dimension	Strategy A: Backup & Restore	Strategy B: Pilot Light	Strategy C: Warm Standby	Strategy D: Active-Active
Target RPO	24 Hours	< 5 Minutes	< 1 Minute	Near Real-time
Target RTO	4 to 8 Hours	30 to 45 Minutes	5 to 10 Minutes	Near Zero
Relative Cost	Extremely Low (1x)	Moderate (35x)	High (85x)	Very High (120x)
Engineering Overhead	Low (Standard backups)	Medium (Failover scripts)	High (Automated Route 53)	Extreme (High)
Regulatory Risk	Minimal	Medium (Replica active)	High (Continuous Sync)	Extreme (High)
PDPA Compliance Path	Simple	Requires TIA & Consent	Requires TIA & Consent	Requires Full Audit
Sovereignty Standing (In-Region)	100% Secure (All data, backups, and compute remain strictly in ap-southeast-5. Exempt from Section 129.)	100% Secure (All sync replica DBs and compute templates remain strictly in ap-southeast-5.)	100% Secure (Active standby stack resides 100% within ap-southeast-5 borders.)	100% Secure (Both sites active within Malaysia ap-southeast-5 AZs.)

Evaluation Dimension	Strategy A: Backup & Restore	Strategy B: Pilot Light	Strategy C: Warm Standby	Strategy D: Active-Active
Sovereignty Standing (Cross-Region)	Conditional / Low Risk (Static encrypted backups cross borders. Requires standard TIA and user consent notice.)	Medium Risk (Live RDS replica runs in Singapore/Jakarta. Active memory footprint in foreign jurisdiction; requires robust TIA.)	High Risk (Continuous sync of live application state, DB, cache, and EFS to foreign region. Requires strict SCCs and DPO audit.)	Very High Risk (5 real-time writes queries process foreign soil. High exposure to foreign enforcement require explicit consent audits.)

6. Strategic Recommendations & Roadmap

To balance **business continuity, cost-efficiency, and strict compliance with Malaysia National Sovereignty (PDPA)**, the following phased roadmap is recommended:



Phase 1: Deploy Sovereign In-Region Backup & Restore (Default Baseline)

- **Action:** Configure AWS Backup to take daily snapshots of our Multi-AZ RDS PostgreSQL database, the EFS model volume, and the gp3 EBS volumes of our compute nodes. Store these backups securely inside ap-southeast-5.

- **Sovereignty Advantage:** 100% compliant with the Malaysian PDPA. Zero data crosses national borders, completely bypassing Section 129 regulatory hurdles, saving legal/DPO consultation fees.
- **Cost Impact:** Negligible baseline charge (~\$7.30 USD / RM 32.85 per month).

Phase 2: Implement AWS Elastic Disaster Recovery (AWS DRS) (Near-Zero RPO / RT0)

- **Action:** Install the AWS DRS Replication Agent on active servers to continuously replicate block-level changes asynchronously into a dedicated staging area subnet inside `ap-southeast-5`. Configure DRS launch templates to spin up fully-provisioned target instances upon a drill or failover event.
- **Sovereignty Advantage:** Keeps staging and target areas completely localized within the AWS Malaysia region. Absolute data sovereignty is maintained, satisfying the 2025 CBPDT Guidelines natively while achieving sub-minute RPO.
- **Cost Impact:** Exceptionally cost-effective (~\$112.30 USD / RM 505.35 per month), avoiding active duplicate server overhead until disaster declaration.

Phase 3: Transition to Sovereign In-Region Pilot Light (As Real-Time Hot Read Needs Arise)

- **Action:** If the application scales further and business workflows require an active, continuously queryable read-replica for real-time reads/Bi reports in addition to DR, deploy an RDS PostgreSQL Read Replica in a separate VPC/subnet configuration within the **Malaysia region (ap-southeast-5)**, keeping the ASG desired capacity at 0.
- **Sovereignty Advantage:** Maintains absolute local residency. Enables fast database promotion and compute failover without violating data sovereignty policies.
- **Cost Impact:** Highly cost-optimized (~\$254.22 USD / RM 1,143.99 per month), representing a fraction of the cost of running a full Warm Standby or Cross-Region Active-Active cluster.

RDS PostgreSQL 17 vs. Percona Server for PostgreSQL 17

#aws

#cloud

#architecture

#rds

#dns

#disaster-recovery

#postgresql

#compliance

Deep Technical Comparison: AWS RDS PostgreSQL 17 vs. Self-Installed Percona Server for PostgreSQL 17

Evaluating database platforms for a highly secure, high-performance 3-tier architecture in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)** requires a careful analysis of managed service convenience versus raw open-source control.


This document provides a comprehensive technical comparison between **AWS Relational Database Service (RDS) PostgreSQL 17 (Multi-AZ)** and a **self-installed Percona Server for PostgreSQL 17** deployed on Amazon EC2 Graviton instances. We analyze performance traits, telemetry, extension availability, query planner optimization, cost models, and operational architecture designs.

1. Architectural Designs

To achieve highly available, production-grade resiliency, both options must be designed to withstand Availability Zone (AZ) outages and ensure automatic failover.

Design A: AWS RDS PostgreSQL 17 (Multi-AZ)

AWS RDS simplifies high availability through synchronous replication at the block storage level. A primary DB instance is provisioned in AZ-A, and a hot standby replica is synchronously maintained in AZ-B.





RDS Failover Mechanism

If the primary instance fails, RDS automatically triggers a failover by changing the canonical name record (CNAME) of the DB endpoint to point to the standby instance. This typically completes in **60 to 120 seconds** depending on transaction volumes and crash recovery processes.

Design B: Percona Server for PostgreSQL 17 on EC2 (Patroni, PgBouncer & etcd)

Deploying Percona Server on EC2 requires a self-managed High Availability (HA) stack. This architecture uses **Patroni** as the template for cluster management, **etcd** as the Distributed Consensus Store (DCS), **PgBouncer** for connection pooling, and **pg_backrest** for robust, compressed backup streaming to S3.





Percona Failover Mechanism

Patroni continuously monitors cluster health and registers state changes in the **etcd** Distributed Consensus Store (DCS). If the primary node fails or loses its leader key in etcd, the surviving nodes hold an election. The node with the most up-to-date Write-Ahead Log (WAL) is promoted to primary. PgBouncer is automatically updated with the new write destination, leading to a highly responsive failover time of **10 to 30 seconds**.

2. Performance Comparison & Telemetry

Managed services restrict low-level performance tuning in exchange for operating system abstraction. Percona Server for PostgreSQL unlocks full infrastructure optimization and integrates enterprise features out of the box.

Telemetry and Observability

AWS RDS PostgreSQL

- **Tools:** AWS CloudWatch Metrics, Enhanced Monitoring (OS metrics via an agent), and **RDS Performance Insights**.
- **Capabilities:** Performance Insights is a robust tool that measures database load using **Average Active Sessions (AAS)**. It visualizes

waiting threads, query text, and wait events (e.g., locking, I/O, CPU bottle-necks).

- **Constraints:** CloudWatch collects metrics in discrete blocks (minimum 1-second intervals for Enhanced Monitoring; 1-minute default for CloudWatch). Accessing database internal states is limited, and exporting granular query telemetry to third-party dashboards requires complex agent setups or costly exporters.

Percona Server for PostgreSQL

- **Tools: Percona Monitoring and Management (PMM)** (fully integrated, free, and open-source).
- **Capabilities:** PMM provides deep observability into both the Linux OS kernel and the PostgreSQL database. It captures query-level execution times, index usage, connection behavior, and locks. PMM uses Prometheus for time-series metrics and Grafana for visual dashboards.
- **Query Analytics:** PMM reads directly from the enhanced `pg_stat_monitor` extension, offering chronological query plans, visual query performance over time, and client-metadata tracking.

Extensions & Advanced Observability

The difference in extensibility is a critical factor when choosing between AWS RDS and Percona Server.

Feature / Extension	AWS RDS PostgreSQL 17	Percona Server for PostgreSQL 17	Technical Impact
pg_stat_monitor	Not Supported (Only generic pg_stat_statements is available)	Built-in	pg_stat_monitor tracks query performance using bucket-based intervals capturing exact query plan metrics, timing histograms, client addresses, execution planning times, and query parameters safely without exposing sensitive constants.
Custom Dynamic Libraries	Blocked	Allowed	RDS only supports a predefined list of approved extensions. Custom or proprietary dynamic-link libraries (shared_preload_library) are restricted due to lack of OS-level superuser rights.
pg_repack	Supported (Pre-installed)	Supported	Allows rebuilding tables/indexes online to reclaim bloated disk space without exclusive locks.
pgaudit	Supported (Pre-installed)	Supported	Provides detailed session and object logging for enterprise compliance.

Architectural Tuning Capability

Because AWS RDS abstracts the underlying operating system, users cannot perform system-level resource optimization. On Amazon EC2, Percona Server allows you to fine-tune the OS kernel and hardware integration.

File System Optimization

- **AWS RDS:** Uses the default AWS Linux filesystem structure (typically ext4-based) with standard mount parameters.
- **Percona Server (EC2):** Allows using **XFS** or **ZFS** optimized for PostgreSQL. For example, aligning the filesystem block size (typically 4KB) with PostgreSQL's page size (8KB) prevents **torn-write pages** and reduces WAL write amplification. Mount parameters such as `noatime` can be enabled to bypass write overheads during read-heavy workloads.

Kernel Tuning & Huge Pages

- **AWS RDS:** Configures standard memory allocations dynamically. While it supports standard Large Pages, fine-grained adjustments to kernel-level memory overcommit options or socket queue sizes are inaccessible.
- **Percona Server (EC2):** Enables precise optimization of system memory and swap behaviors:

```
# Optimize kernel memory retention to avoid aggressive swapping
sysctl -w vm.swappiness=1
sysctl -w vm.overcommit_memory=2
sysctl -w vm.overcommit_ratio=80

# Configure and lock Huge Pages to guarantee block buffer efficiency
sysctl -w vm.nr_hugepages=4096
```

Connection Pooling

- **AWS RDS:** Relies on AWS RDS Proxy (an extra paid resource starting at \$0.015 per LCU-hour) or application-side poolers.
 - **Percona Server (EC2):** Allows deploying **PgBouncer** directly on the same EC2 database hosts or dedicated scaling nodes. Configuring PgBouncer in **Transaction Mode** allows the database to easily support tens of thousands of active connections with minimal memory overhead.
-

Query Planner Optimization

Both platforms utilize the core PostgreSQL 17 Query Planner engine, but Percona Server provides more freedom to tune behavior based on exact hardware characteristics.

- **Cost Parameters:** Developers can adjust parameters such as `random_page_cost` and `seq_page_cost`. On RDS, these are modified via dynamic DB Parameter Groups. On Percona Server, they can be configured directly in `postgresql.conf` or set per-session/database to match specific storage benchmarks.
 - **Query Hinting:** Both support `pg_hint_plan` to override planner decisions. However, Percona's tight integration with PMM makes identifying sub-optimal plans and injecting plans significantly easier during active troubleshooting.
-

3. Cost Analysis & Sizing in Malaysia (ap-southeast-5)

To provide an accurate cost comparison, we model both architectures inside the **AWS Malaysia region (ap-southeast-5)** utilizing current On-Demand pricing.

We evaluate two standard sizing tiers:

- **Baseline Tier (Small Business / Staging):** ~8GB RAM, 2 vCPU, 50GB storage.

- **High-Performance Tier (Enterprise Production):** ~16GB RAM, 4 vCPU, 100GB storage.

Note: Calculations assume 730 run-hours per month and an exchange rate baseline of 1 USD $\approx$ 4.50 MYR.

Baseline Cost Model

Sizing Profiles:

- **AWS RDS:** Multi-AZ db.m6g.large (2 vCPU, 8GB RAM) + 50GB gp3 Multi-AZ storage.
- **Percona on EC2:** 2x t4g.large EC2 instances (2 vCPU, 8GB RAM, one Primary, one Replica) + 1x t4g.micro EC2 instance in a private subnet running etcd/DCS for cluster consensus. 50GB gp3 storage per database host.

Sizing & Cost Component	Option A: AWS RDS (Multi-AZ)	Option B: Percona Server on EC2 (HA)	Cost & Technical Analysis
Compute / Host Cost	\$221.92 / month db.m6g.large Multi-AZ hourly rate (\$0.304/hr)	\$104.25 / month 2x t4g.large (\$49.06/mo each) 1x t4g.micro etcd (\$6.13/mo)	EC2 is significantly cheaper as RDS charges a managed database premium for licensing, automatic configuration, and Multi-AZ replication.
Storage (gp3)	\$11.50 / month 50GB Multi-AZ storage (\$0.23/GB-mo)	\$8.00 / month 2x 50GB gp3 storage volumes (\$0.08/GB-mo each)	RDS Multi-AZ storage rates are higher than standard gp3 volumes on EC2 because the synchronous block mirroring process is managed by AWS.
Backup Costs	\$4.75 / month RDS automated snapshot storage (~50GB)	\$1.15 / month pg_backrest compressed backup to S3 (~50GB standard bucket)	pg_backrest performs block-level deduplication and compression, significantly reducing the storage footprint on S3.

Sizing & Cost Component	Option A: AWS RDS (Multi-AZ)	Option B: Percona Server on EC2 (HA)	Cost & Technical Analysis
Monitoring Tooling	\$0.00 (Included basic Performance Insights)	\$0.00 (PMM is open-source and self-hosted)	Performance Insights is included, while PMM can be hosted on a shared utility instance at no extra cost.
Day-2 Operations / Human Maintenance	\$150.00 / month • Estimated 2 engineering hours/mo of DBA administration	\$1,500.00 / month • Estimated 15 engineering hours/mo (Patroni upgrades, OS patching, backups verification)	Crucial Difference: While Percona on EC2 saves on raw AWS resources, it introduces significant human maintenance costs.
TOTAL MONTHLY COST (USD)	\$388.17 USD	\$1,613.40 USD (Including labor) \$113.40 USD (Raw infrastructure only)	For baseline setups, the engineering labor required to manage a high-availability cluster makes RDS the more cost-effective choice.
TOTAL MONTHLY COST (MYR)	~RM 1,746.77 MYR	~RM 7,260.30 MYR (Including labor) ~RM 510.30 MYR (Raw infrastructure)	<i>(Calculated at 1 USD ≈ 4.50 MYR)</i>

High-Performance Cost Model

Sizing Profiles:

- **AWS RDS:** Multi-AZ db.m6g.xlarge (4 vCPU, 16GB RAM) + 100GB gp3 Multi-AZ storage.
- **Percona on EC2:** 2x t4g.xlarge EC2 instances (4 vCPU, 16GB RAM) + 1x t4g.micro EC2 instance for DCS quorum. 100GB gp3 storage per database host.

Sizing & Cost Component	Option A: AWS RDS (Multi-AZ)	Option B: Percona Server on EC2 (HA)	Cost & Technical Analysis
Compute / Host Cost	\$443.84 / month db.m6g.xlarge Multi-AZ hourly rate (\$0.608/hr)	\$202.35 / month 2x t4g.xlarge (\$98.11/mo each) 1x t4g.micro etcd (\$6.13/mo)	The compute cost savings on EC2 increase with larger instance families, making self-managed setups attractive for raw compute efficiency.
Storage (gp3)	\$23.00 / month 100GB Multi-AZ storage (\$0.23/GB-mo)	\$16.00 / month 2x 100GB gp3 storage volumes (\$0.08/GB-mo each)	Raw EBS volumes on EC2 remain significantly more cost-effective than managed Multi-AZ RDS storage.
Backup Costs	\$9.50 / month RDS automated snapshot storage (~100GB)	\$2.30 / month pg_backrest compressed backup to S3 (~100GB standard bucket)	Compression and differential backup strategies on S3 reduce long-term snapshot storage costs.
Monitoring Tooling	\$0.00 (Performance Insights)	\$0.00 (Self-hosted PMM)	Basic Performance Insights is included, while PMM provides advanced query performance tracking.

Sizing & Cost Component	Option A: AWS RDS (Multi-AZ)	Option B: Percona Server on EC2 (HA)	Cost & Technical Analysis
Day-2 Operations / Human Maintenance	\$150.00 / month Estimated 2 engineering hours/mo of DBA administration	\$1,500.00 / month Estimated 15 engineering hours/mo (Patroni maintenance, clustering, state checks)	Human labor remains the dominant cost factor for self-managed architectures.
TOTAL MONTHLY COST (USD)	\$626.34 USD	\$1,720.65 USD (Including labor) \$220.65 USD (Raw infrastructure only)	Even at larger scales, the cost of specialized engineering labor to maintain a reliable HA cluster makes RDS highly competitive.
TOTAL MONTHLY COST (MYR)	~RM 2,818.53 MYR	~RM 7,742.93 MYR (Including labor) ~RM 992.93 MYR (Raw infrastructure)	<i>(Calculated at 1 USD ≈ 4.50 MYR)</i>

4. Operational and Strategic Comparison

Operational Area	Option A: AWS RDS (Multi-AZ)	Option B: Percona Server on EC2	Strategic Decision Guidance
Setup & Provisioning	Instant (Terraform/Console) Provisioned inside secure private subnets within minutes.	Complex Requires configuring OS layers, security groups, etcd clusters, Patroni playbooks, and pg_backrest backups.	Use RDS for faster time-to-market and to avoid operational overhead.
Scalability (Storage & Compute)	Storage Autoscaling & Direct Instance Upgrades Increases volume size dynamically and changes instance classes with minimal downtime.	Manual / Complex Requires resizing EBS volumes, executing OS-level file system growth commands, and scaling instances manually.	RDS provides a highly flexible platform for unpredictable or rapidly scaling workloads.
Minor & Major Upgrades	Fully Automated Upgrades are handled via maintenance windows or simple Terraform version adjustments.	Self-Managed Requires carefully planning node-by-node updates to prevent cluster state mismatches.	RDS is ideal for teams without a dedicated Database Reliability Engineer (DBRE).

Operational Area	Option A: AWS RDS (Multi-AZ)	Option B: Percona Server on EC2	Strategic Decision Guidance
Vendor Lock-in	Moderate Uses open-source PostgreSQL APIs but relies on AWS-specific infrastructure and APIs.	None Completely open-source. The entire stack can be moved to on-premises or other cloud providers without modifications.	Choose Percona on EC2 if you require hybrid portability or multi-cloud compliance.

5. Summary & Recommendation

Choose AWS RDS PostgreSQL 17 if:

- . **You want to minimize operational overhead:** Your engineering team is lean and does not have dedicated DBA resources to manage backups, clustering, or operating system updates.
- . **You need fast deployment:** You need to launch a production-grade database in private subnets within minutes.
- . **You want automated maintenance:** You require automatic patch application, effortless scaling, and managed snapshots.

Choose Percona Server for PostgreSQL 17 on EC2 if:

- . **You need absolute performance optimization:** Your application requires custom filesystems (like ZFS/XFS), kernel-level memory tuning (Huge Pages), or connection pooling at scale.
- . **You want advanced query visibility:** You want to leverage the `pg_stat_monitor` extension and `PMM` without being constrained by the limits of Performance Insights.

- . **You want to avoid vendor lock-in:** You require a platform-agnostic architecture that can run identically on-premises, on bare metal, or on other cloud providers.
- . **You have the engineering bandwidth:** You have a dedicated platform team capable of managing High Availability clusters (Patroni, etcd, pg_backrest).

RAGFlow + Langfuse: AWS vs. On-Premises GPU Integration Guide

#aws #cloud #architecture #vpc #alb #asg #rds #waf #elasticache #valkey
#ssl #disaster-recovery #postgresql #gpu #ragflow #langfuse #sovereignty
#compliance

Deploying RAGFlow & Langfuse: AWS vs. On-Premises GPU Integration Guide

1. Executive Summary

This guide provides a comprehensive architectural and economic analysis of deploying **RAGFlow** (an advanced, layout-aware Retrieval-Augmented Generation engine) paired with **Langfuse** (an open-source LLM engineering and observability platform).


Given that RAGFlow relies heavily on deep document understanding (DeepDoc), optical character recognition (OCR), layout-aware parsing, and embedding vectorization, the choice of underlying compute architecture—specifically **GPU acceleration**—directly impacts ingestion latency, retrieval precision, and query response times.

Furthermore, this guide details how organizations can achieve a highly cost-effective, high-performance hybrid model. Under this paradigm, secure transactional systems run within **AWS Asia Pacific (Malaysia) Region (ap-southeast-5)**, while heavy GPU workloads are processed on-premises or vice versa, interconnected seamlessly via secure **REST APIs** and the **Model Context Protocol (MCP)**.

2. Technology Overview

RAGFlow Architecture

RAGFlow is an open-source RAG engine based on **deep document understanding**. Unlike naive RAG systems that split text strictly by character count, RAGFlow uses visual and layout-aware machine learning models to extract structures, tables, and figures from complex documents (PDFs, Word docs, spreadsheets).



The core backend services are highly decoupled:

- . **RAGFlow API & Task Executor:** Python services managing workflows and document parsing.
- . **DeepDoc Parser:** Implements a multi-stage machine learning pipeline (Layout recognition, Table Structure Recognition [TSR], and OCR).

- . **Storage Tier:** Object storage (MinIO or S3) for original files, and a database (MySQL or PostgreSQL) for transactional metadata.
- . **Search/Vector Tier:** Elasticsearch or Infinity (an AI-native database) for full-text search and dense vector recall.
- . **Caching Tier:** Valkey or Redis for task-broker message queuing and session caching.

Langfuse Architecture


Langfuse is a purpose-built LLM engineering platform that provides:

- **Observability & Tracing:** Near real-time tracing of complex multi-step chains, tool calls, vector retrievals, and generation steps.
- **Prompt Management:** Version-controlled centralized prompt repository.
- **Evaluation & Analytics:** Quantifying LLM latencies, token consumption, and retrieval relevance (e.g., faithfulness, answer relevance).

RAGFlow natively integrates with Langfuse via simple API configurations, allowing operators to trace every document chunk retrieved, prompt assembled, and token generated by downstream LLMs.

3. The Critical Role of GPUs in RAGFlow: Why It Matters

Deploying RAGFlow without a GPU represents a significant architectural bottleneck. The table below outlines how specific sub-systems within RAGFlow utilize hardware and why CPU-only fallbacks are highly inefficient in production.



A. DeepDoc Layout & Visual Parsing (ONNX Runtime)

RAGFlow uses a combination of visual computer vision models (such as YOLOv8 for layout detection, TSR for table structure analysis, and PP-OCrv4/v5 for text extraction).

- **GPU Utility:** These models run via ONNX Runtime using the `CUDAExecutionProvider`.
- **CPU Bottleneck:** Running visual detection models on standard vCPUs causes severe thread saturation, pushing CPU usage to 100% and extending document chunking speeds from **seconds per page** to **minutes per page**.

B. Embedding Generation & Re-ranking (PyTorch / CUDA)

Vector search requires converting text chunks into multi-dimensional dense vectors (using models like BGE-M3 or local HuggingFace embeddings) and re-evaluating candidate matches using cross-encoder re-rankers.

- **GPU Utility:** Vectorizing large batches of documents is an intensely parallelizable matrix-multiplication operation suited

perfectly for GPU Tensor Cores.

- **CPU Bottleneck:** Bulk document ingestion (e.g., a 10,000-page knowledge base) will take days on general-purpose CPUs compared to **minutes or hours on an enterprise GPU**.


C. Local LLM Execution (Ollama / vLLM)

If the organization runs local, self-hosted LLMs (e.g., Qwen-2.5-Instruct, Llama-3-8B) to preserve absolute data sovereignty and avoid third-party API costs.

- **GPU Utility:** Modern LLMs require low-latency KV-caching and rapid auto-regressive decoding, which can only be achieved with high GPU memory bandwidth (VRAM).
- **CPU Bottleneck:** Running an 8B parameters model on a high-spec CPU results in a generation rate of 1 to 3 tokens per second, which is completely unusable for interactive production applications (whereas a GPU easily outputs 50+ tokens per second).

4. AWS-Native Deployment Architecture

Deploying RAGFlow and Langfuse natively in AWS maximizes durability, automatic scaling, and reduces administrative overhead by leveraging managed services.




Components

- . **Network Entrypoint:** AWS WAFv2 filters web-attacks before routing requests through a Multi-AZ **Application Load Balancer (ALB)**.
- . **Compute Tier (Stateless Web):** Langfuse (Node.js/Next.js) and RAGFlow front-ends are deployed on standard, cost-efficient AWS Auto Scaling Groups (ASGs).
- . **Compute Tier (GPU Processing):** Dedicated EC2 instances (specifically **G5** or **G6** families) are deployed in private subnets to run the RAGFlow task executor, ONNX OCR execution, and local LLM/embedding inference.
- . **Database Tier:** Managed **Amazon RDS PostgreSQL (Multi-AZ)** acts as the central transactional database for both Langfuse metadata and RAGFlow document status.

- . **Caching & Queue:** Managed **Amazon ElastiCache for Valkey** provides high-throughput, low-latency key-value queuing.
- . **Storage Tier:** Durability-guaranteed **Amazon S3** replaces local MinIO for holding parsed raw documents and assets.

5. On-Premises Deployment Architecture

When running on-premises, teams rely on a consolidated, cost-efficient hardware stack utilizing standard container orchestration tools.



Components

- . **Hardware Host:** A high-end workstation or rack-mounted server running Ubuntu 24.04/26.04 LTS equipped with at least one NVIDIA GPU (minimum 16GB VRAM, recommended 24GB+ VRAM).
 - . **GPU Driver & Runtime Layers:** Native NVIDIA Driver (e.g., v550+), **CUDA Toolkit (12.1 or 12.4)**, and **NVIDIA Container Toolkit** to allow Docker containers to access GPU hardware directly.
 - . **Orchestration Layer:** **Docker Compose** or a single-node **Kubernetes (K3s)** cluster.
 - . **Storage Engine:** Local NVMe SSD arrays configured in RAID-1 or RAID-10 for low-latency indexing read/writes.
 - . **Services Consolidation:** PostgreSQL, Valkey/Redis, MinIO, Elasticsearch, Langfuse, and RAGFlow running as interconnected container services within a isolated Docker bridge network.
-

6. Head-to-Head Comparison: AWS vs. On-Premises

To support technical leadership in making informed architectural choices, this section analyzes the capital and operating realities of both deployment environments, specifically calibrated for **AWS Asia Pacific (Malaysia) Region (ap-southeast-5)** in both **USD** and **MYR** (assuming 1 USD = 4.50 MYR).

Side-by-Side Matrix

Evaluation Dimension	AWS-Native (ap-southeast-5)	On-Premises GPU Server
Financial Classification	OpEx (Pay-as-you-go, monthly billing)	CapEx (Upfront hardware purchase + power/cooling)
GPU Instance Scarcity	High. GPU instances (g5, p4) often require AWS limit increases or suffer from low regional capacity in Malaysia.	None. Hardware is fully owned and dedicated to the organization.
Ingress/Egress Costs	Egress from AWS to internet is \$0.09 per GB (after 100 GB/mo free). Bulk PDF downloading incurs costs.	None. Data transfer within the local area network is entirely free and unmetered.
Scalability	Near-instantaneous scaling of compute node ASGs. Elastic storage (S3) scales infinitely.	Hardware-constrained. Scaling requires purchasing additional physical GPUs or servers.
Maintenance Overhead	Extremely low. AWS handles RDS Multi-AZ backups, hardware failures, and network routing.	High. Internals teams must handle hardware warranties, disk failures, cooling, and UPS backup.
Sovereignty & Air-Gap	Shared Responsibility model. Subject to cloud compliance policies.	Absolute. Can be run completely air-gapped without any public internet connectivity.

Financial Model

We present a comparative financial projection over a standard **12-month operating cycle**.

The AWS model runs a high-performance, resilient RAGFlow + Langfuse deployment with a single GPU node (g5.xlarge containing an NVIDIA A10G 24GB VRAM). The On-Premises model assumes the purchase of a professional, enterprise-grade AI workstation equipped with an NVIDIA RTX 4090 (24GB VRAM).

12-Month Total Cost Comparison (USD)	
AWS Cloud (OpEx)	\$15,313.20
On-Premises (CapEx)	\$6,400.00

AWS 12-Month Total Cost Breakdown

- . GPU Compute Instance (g5.xlarge - 1 Node):**
 - Hourly on-demand rate: \$1.15 / hr
 - Monthly cost: 730 hours * \$1.15 = \$839.50 USD
 - **12-Month Total: \$10,074.00 USD (RM 45,333.00 MYR)**
- . Auxiliary Compute (ASGs for Web & Langfuse):**
 - 2 x t4g.medium (ARM64, 2 vCPU, 4GB RAM) = 2 * \$0.0384/hr * 730 hrs = \$56.06 USD / mo
 - **12-Month Total: \$672.72 USD (RM 3,027.24 MYR)**
- . Managed Storage (Amazon S3 + RDS Multi-AZ PostgreSQL):**
 - S3 Storage (200 GB) + RDS Postgres (Multi-AZ db.t4g.medium, 100GB gp3) ≈ \$180.00 USD / mo
 - **12-Month Total: \$2,160.00 USD (RM 9,720.00 MYR)**
- . Network Fees (WAF, NAT Gateway, Egress, ALB):**
 - Shared ALB, WAF rules, NAT base, and 200 GB egress ≈ \$200.54 USD / mo
 - **12-Month Total: \$2,406.48 USD (RM 10,829.16 MYR)**
- AWS 12-MONTH COMBINED TOTAL: \$15,313.20 USD / ~RM 68,909.40 MYR**

On-Premises 12-Month Total Cost Breakdown

. Hardware Purchase (One-Time CapEx):

- AMD Threadripper or high-end Intel CPU, 128GB RAM, 2TB NVMe Gen4 SSD, Dual-chamber cooling tower.
- GPU: NVIDIA RTX 4090 (24GB VRAM, Tensor Cores).
- **Upfront cost: \$5,200.00 USD (RM 23,400.00 MYR)**

. Datacenter Rack Power & Cooling (OpEx):

- Avg draw: ~650 Watts under active load.
- Electricity cost in Malaysia (commercial rates of approx. RM 0.39 per kWh) + standard cooling overhead ≈ \$60.00 USD / mo
- **12-Month Total: \$720.00 USD (RM 3,240.00 MYR)**

. Internal Systems Engineer Maintenance (Allocated Labor):


- Fractional engineering time for OS updates, hardware checks, backups ≈ \$40.00 USD / mo
- **12-Month Total: \$480.00 USD (RM 2,160.00 MYR)**

- **ON-PREMISES 12-MONTH COMBINED TOTAL: \$6,400.00 USD / ~RM 28,800.00 MYR**

7. Hybrid Integration: API and MCP Architectures

To minimize cloud costs while maintaining enterprise-grade resilience, organizations can adopt a **Hybrid Deployment Model**.

This approach isolates standard web frontends, authentication portals, and relational metadata within secure, scalable **AWS private subnets**, while routing compute-heavy OCR, layout analysis, and local LLM/embedding inference requests to highly cost-efficient **On-Premises GPU Workstations** situated in corporate datacenters or Cyberjaya offices.



The hybrid model supports two primary secure communication patterns:

Pattern A: Secure REST API Integration

In this pattern, RAGFlow's core engine and databases run in AWS, but the heavy task execution worker (RAGFlow Worker daemon) is deployed on-premises.

- **Execution Flow:**

1. A user uploads a 500-page PDF to the AWS Web Portal. The file is saved directly to **Amazon S3**.
 2. The database logs the parsing task in the **AWS RDS PostgreSQL** queue.
 3. The local on-premises GPU server runs a secure RAGFlow Worker client which continuously polls the AWS database queue (via secure outbound SSL connections or private Site-to-Site VPN).
 4. The on-premises worker pulls the PDF from the secure Amazon S3 pre-signed URL, processes it using local **NVIDIA Tensor Cores** (running YOLOv8 layout and PP-OCR), and uploads the structured vector embeddings back to AWS Elasticsearch/Infinity.
 5. The final trace data is pushed from RAGFlow to the AWS-hosted **Langfuse** server to log execution timelines and parsing latencies.
- **Security Control:** The on-premises workstation needs **no inbound open ports**. It establishes outbound-only HTTPS connections to Amazon S3, RDS, and Langfuse, bypassing corporate firewall ingress limitations.

Pattern B: AI-Native Model Context Protocol (MCP) Integration

The **Model Context Protocol (MCP)** allows AI Agents running on AWS (such as **Amazon Bedrock AgentCore** or custom LangChain/Langgraph agents) to dynamically call on-premises resources, search internal databases, or request local GPU parsing.

- **Execution Flow:**

1. An enterprise AI Agent running in the AWS Cloud receives a query: *"Compare our quarterly financial PDF with the structural charts inside our local Cyberjaya archives."*
2. The agent identifies that the Cyberjaya archives reside on-premises and requires local OCR table-extraction capabilities.
3. The agent initiates an outgoing call via **Amazon API Gateway configured with MCP Proxy Support**.

4. API Gateway translates the request into a standardized JSON-RPC 2.0 payload, routing it over an mTLS-secured connection to the on-premises **Local MCP Server**.
 5. The on-premises MCP Server accesses the local RTX 4090 GPU to execute RAGFlow's DeepDoc visual parser on the requested archival charts.
 6. The structured JSON tables and charts are returned back through API Gateway to the AWS Bedrock Agent, which synthesizes the final response.
- **Tracing Compliance:** Because RAGFlow and Langfuse are fully integrated, the complete round-trip trace—including the AWS Bedrock invocation, the API Gateway transit, the on-premises MCP tool execution, and the final GPU-accelerated generation—is visualised inside the **Langfuse Dashboard** as a unified trace graph.
-

8. Practical Integration & Setup Guide

1. Collect Langfuse Credentials

Before configuring RAGFlow, set up a Langfuse project and acquire the API credentials:

- Sign in to your Langfuse dashboard.
- Navigate to **Settings** **Projects** and select or create a project.
- Generate a new set of API keys. Copy the following variables:
 - LANGFUSE_PUBLIC_KEY
 - LANGFUSE_SECRET_KEY
 - LANGFUSE_HOST (e.g., `https://cloud.langfuse.com` or `http://<your-aws-langfuse-ip>:3000`)

2. Configure RAGFlow to Emit Traces

RAGFlow includes a native Langfuse integration. Key credentials are set per tenant in the RAGFlow ecosystem. This can be configured dynamically through the Web UI:

- . Log in to your RAGFlow web console.
- . Click on your profile avatar in the top-right corner and select **API**.
- . Scroll to the bottom of the API screen to locate the **Langfuse Configuration** section.
- . Fill in the following details:
 - **Host:** Your Langfuse API host URL.
 - **Public Key:** Your LANGFUSE_PUBLIC_KEY .
 - **Secret Key:** Your LANGFUSE_SECRET_KEY .
- . Click **Save**.

Once saved, RAGFlow automatically wraps its internal pipeline chains, retrievals, and generation steps, emitting telemetry directly to your Langfuse workspace on every query or document retrieval.

3. Setting GPU Memory Limits for OCR & Parsers

When self-hosting RAGFlow on-premises or on AWS GPU instances, you can configure how visual parsing tasks distribute memory across GPUs. Add the following environment variables to your RAGFlow .env configuration file:

```
# Allow RAGFlow to access all available local CUDA devices (e.g. 0,1
for dual-GPU)
```

```
PARALLEL_DEVICES="0"
```

```
# Limit maximum GPU memory dedicated to the PP-OCR visual engine to
prevent out-of-memory (OOM) crashes
```

```
OCR_GPU_MEM_LIMIT_MB=4096
```

```
# Define model structure caching strategy (extend vs. shrink)
```

```
OCR_ARENA_EXTEND_STRATEGY="extend"
```

Restart your Docker Compose service to apply limits:

```
docker compose -f docker/docker-compose.yml up -d
```

4. Viewing Unified Traces in Langfuse

After executing a query in RAGFlow (e.g., asking a question in the chat portal), open your Langfuse Project Dashboard:

- Navigate to **Traces**.
- You will see incoming streams prefixed with `ragflow-*`.
- Expanding a trace reveals a nested tree of spans:
 - `ragflow-retrieval`: Logs how long Elasticsearch/Infinity took to execute hybrid keyword/dense-vector queries, accompanied by exact parsed document chunk payloads.
 - `ragflow-rerank`: Quantifies the cross-encoder execution latency.
 - `ragflow-generation`: Visualizes prompt template construction and logs downstream LLM token-generation speeds.

Google Antigravity Skills Integration Guide

#aws

#cloud

#architecture

#rds

#elasticsearch

#valkey

#disaster-recovery

#antigravity

#skills

Google Antigravity Skills Integration Guide

Welcome to the **Google Antigravity Skills Integration Guide** for the AWS Secure 3-Tier Web & AI Infrastructure project. This document outlines how to integrate and share agent knowledge, capabilities, and workflows between different AI assistants—specifically **Google Jules** and **Google Antigravity**—using the open **Agent Skills** ecosystem standard.

By consolidating project-specific operating instructions and architectural boundaries into standard skill specifications, both Google Jules and Google Antigravity can operate with identical high-fidelity context, rulesets, and deployment tools.

1. What are Agent Skills?

Agent Skills represent an open, decentralized standard for extending the capabilities and contextual boundaries of AI agents. Instead of relying purely on generalized system prompts, agents can dynamically discover and read specialized instruction packages tailored for specific repositories, codebases, or workflows.

An Agent Skill is packaged as a standard directory containing a `SKILL.md` file. This markdown file begins with structured YAML front-matter identifying its metadata:

```
---
name: my-skill-name
description: Clear, high-level description of when to apply this skill.
---
# Detailed Instructions
Step-by-step procedures, CLI commands, constraints, and guidelines go
here.
```

2. Where Skills Live: Directory Scopes

Google Antigravity supports both workspace-specific and global scopes for discoverability:

Scope	Location Path	Use Case
Workspace / Project	<workspace-root>/. agents /skills/	Project-specific rulesets, database ports, architectural limits, deployment guides, and custom CLI wrappers.
Global User Scope	~/. gemini /config/skills/	Personal utility commands, code formatter standards, and general-purpose tooling.

Note: While Antigravity natively prioritizes `.agents/skills`, it maintains full backward-compatibility for `.agent/skills` as well.

3. The jules-knowledge Skill

We have defined a dedicated workspace skill for our repository to bridge the knowledge bases of Google Jules and Google Antigravity. This skill is located at:

```
.agents/skills/jules-knowledge/SKILL.md
```

Purpose of jules-knowledge

This skill ensures that whenever you invoke Google Antigravity (using the CLI tool `agy`) or interact with Google Jules in this workspace, the AI assistant automatically:

- . Targets **AWS Asia Pacific (Malaysia)** (`ap-southeast-5`) by default.
 - . Uses **ARM64 Graviton instances** (`t4g.micro` for compute, `db.t4g.micro` for RDS, and `cache.t4g.micro` for Valkey).
 - . Enforces **Ubuntu 26.04 LTS** hardened under the ASIMP framework.
 - . Protects database ingress via port-isolation (blocking direct public routing to port 5432).
 - . Deploys secure session layers with **Amazon ElastiCache for Valkey**.
 - . Follows the mandatory validation pipeline (`./scripts/deploy.sh` and `python scripts/prepare_docs.py`).
-

4. How to Install and Use Skills with npx skills

`npx skills` is a command-line tool developed by Vercel Labs acting as an open package manager for AI agents (including Antigravity, Claude Code, GitHub Copilot, Cursor, and Cline).

Downloading and Syncing Skills

Running `npx skills` packages can place skills in your local user's global configuration (`~/.agents/skills`). To ensure Google Antigravity picks them up properly:

- . **For Workspace Scopes:** Copy the desired skill directory to your project root under `.agents/skills/`:

```
cp -r ~/.agents/skills/some-skill .agents/skills/
```

- . **For Global Scopes (Google Antigravity):** Copy the skill directory to:

```
cp -r ~/.agents/skills/some-skill ~/.gemini/antigravity-cli/skills/
```

5. Seamless Bridge Between Jules and Antigravity

By maintaining our core engineering principles in `.agents/skills/jules-knowledge/SKILL.md` and keeping `AGENTS.md` aligned, we establish a robust bi-directional capability:

- **Google Jules (This Agent):** Uses `AGENTS.md` and this guide to understand and write code conforming to enterprise standards, verifying actions using read-only verification commands.
- **Google Antigravity (The CLI / Local Agent):** Reads the compiled skills within `.agents/skills/` at startup. Typing `/jules-knowledge` or initiating an action matching our database/infrastructure setup will trigger the detailed instructions embedded in the skill.

This unified knowledge strategy guarantees zero-configuration context sync, eliminating environment drift and ensuring both assistants deliver perfectly aligned, region-compliant code.

SOP: Knowledge-First Discovery & Context Preservation Protocol

#okf

#discovery

#context-management

#brain

#dsom

#SOP

#aws

#cloud

SOP: Local Knowledge-First Discovery & OKF Context Protocol

1. Executive Intent

To prevent unnecessary API or CLI queries, remote SSH or Session Manager probes, token window exhaustion, and context loss during agentic sessions, all AI agents must strictly adhere to the **Local Knowledge-First Protocol**. All project facts, architecture details, subnet IPs, security rules, and costing options are documented locally via **OKF v0.1 YAML Frontmatter** inside `.agents/brain/` and `docs/`.

2. Standard Operating Procedure (4-Step Discovery Flow)

[Step 1: User Request / Task]

|



[Step 2: Local OKF Search] — search on .agents/brain/ & docs/ (using 1

|



[Step 3: Local Context Inspection] — read_file target lines on matched

|



[Step 4: AWS Runtime Execution Gate] (ONLY if live runtime state or dep:

Step 1: Local Frontmatter & Metadata Search

Before issuing any AWS CLI command, checking AWS SSM sessions, querying live OpenTofu state, or executing remote scripts against standalone EC2 or active ASG nodes:

- . Search local OKF frontmatter for relevant topics: or description: keywords:
 - Example: Search for “valkey” or “asg” or “postgres” or “dr-options” inside the local files to identify existing design guides and architectural layouts.
- . Read .agents/brain/active_context_manifest.md to see the current session checkpoints.

Step 2: Targeted File Viewing

Once the relevant document is located via OKF frontmatter or titles:

- Read specific segments of the matched local file using read_file to keep the context window lean and avoid loading entire huge files unnecessarily.

Step 3: Local Context Synthesis

Map the local specifications (e.g. ap-southeast-5 regional defaults, instance sizes like t4g.micro, security groups ingress definitions)

to the task at hand. The local repository is the **Single Source of Truth (SSOT)** for configuration intent.

Step 4: AWS Runtime Execution Gate

Probing live AWS resources, executing OpenTofu plans, querying remote database ports, or starting SSM Session Manager connections are authorized **ONLY** when:

- Applying actual infrastructure or configuration updates (e.g., via `./scripts/deploy.sh`).
- Fetching actual live runtime telemetry (e.g., live database statistics or running target group health status) that cannot be modeled statically in documentation.

3. Mandatory Rules Reference

- **Rule 6 (OKF Topics):** All `.md` files must open on line 1 with `---` and contain `topics: [3-5 keywords]` and metadata.
- **Rule 10 (Metadata-First Discovery):** Read targeted file segments to preserve token efficiency.
- **Rule 29 (Local Knowledge-First Mandate):** Always exhaustively search local `.agents/brain/` and `docs/` before digging into remote servers or external web search.

Wazuh Standalone Cloud Installation & Costing Guide

#aws

#cloud

#architecture

#wazuh

#siem

#xdr

#security

#costing

#ec2

Wazuh Standalone Cloud Installation & Costing Guide

Wazuh is an enterprise-grade, open-source security platform providing unified **XDR (Extended Detection and Response)** and **SIEM (Security Information and Event Management)** capabilities. It monitors endpoints, cloud services, and containers, analyzing security events in real-time.

To maintain strict cost isolation and modular design, this guide details how to deploy Wazuh on a **standalone EC2 instance** completely separate from any Auto Scaling Group (ASG), and outlines the **best and cheapest cloud-deployment models** with a dedicated financial breakdown.

1. Wazuh Central Architecture & Sizing

The Wazuh central components consist of three primary elements:

- . **Wazuh Indexer:** A highly scalable, full-text search and analysis engine that indexes and stores alerts generated by the Wazuh server.
- . **Wazuh Server:** The management node that receives and analyzes data from agents, applying rules and raising alerts.

- . **Wazuh Dashboard:** The web user interface for mining, visualizing, and configuring security alerts and platform features.

Sizing and Hardware Guidelines

Wazuh's resource utilization is directly proportional to the event ingestion rate (Events Per Second, or EPS) and the number of active agents.

According to the [Wazuh Quickstart Documentation](#):

- **Official Recommended Sizing (1-100 agents):** 4 vCPUs, 8 GiB RAM, and 50-200 GB SSD storage.
- **Cheapest Standalone Feasibility (1-25 agents):** For testing or low-traffic small environments, Wazuh central components can run successfully on **2 vCPUs and 4 GiB RAM** (e.g., using JVM heap size tuning for OpenSearch/Indexer to prevent out-of-memory states).

2. Separated Costing Models (USD & MYR)

To keep budget planning fully transparent, we analyze four standalone deployment pathways in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)** and compare them with the cheapest global baseline in **US East (N. Virginia, us-east-1)**.

Exchange rate baseline: 1 USD ≈ 4.50 MYR.

Configuration Sizing Matrix

Metric / Resource	Option A: Ultra-Cheap / Dev (Malaysia ap-southeast-5)	Option B: Recommended Cheap Production (Malaysia ap-southeast-5)	Option C: Global Cheap Baseline (US East us-east-1)	Option D: Official AMI Sizing (Malaysia ap-southeast-5)
Instance Type	t4g.medium (ARM64)	t4g.large (ARM64)	t3.medium (x86_64)	c5a.xlarge (x86_64)
vCPU / Memory	2 vCPUs / 4 GiB RAM	2 vCPUs / 8 GiB RAM	2 vCPUs / 4 GiB RAM	4 vCPUs / 8 GiB RAM
Target Workload	1-15 Agents (Dev/Test)	1-35 Agents (Light Prod)	1-15 Agents (Dev/Test)	1-100 Agents (Full Prod)
EBS Storage (gp3)	50 GiB gp3	100 GiB gp3	50 GiB gp3	100 GiB gp3
Network Setup	1 Elastic IP (Static Public IP)	1 Elastic IP (Static Public IP)	1 Elastic IP (Static Public IP)	1 Elastic IP (Static Public IP)

Detailed Cost Breakdown Table (Monthly)

All estimates assume On-Demand usage over a standard month (~730 hours).

Cost Category	Option A (USD / MYR)	Option B (USD / MYR)	Option C (USD Only)	Option D (USD / MYR)
EC2 Compute	\$24.53 / RM 110.39	\$49.06 / RM 220.77	\$30.37	\$136.51 / RM 614.30
gp3 EBS Storage	\$4.00 / RM 18.00	\$8.00 / RM 36.00	\$4.00	\$8.00 / RM 36.00
Public IPv4 Fee	\$3.65 / RM 16.43	\$3.65 / RM 16.43	\$3.65	\$3.65 / RM 16.43
AWS Backup Snapshot	\$2.50 / RM 11.25	\$5.00 / RM 22.50	\$2.50	\$5.00 / RM 22.50
Data Ingress/Egress	Free / Free	Free / Free	Free	Free / Free
TOTAL (Monthly)	\$34.68 / RM 156.07	\$65.71 / RM 295.70	\$40.52	\$153.16 / RM 689.23

Key Cost Efficiency Discoveries

- . **Graviton Efficiency (Option A & B):** Using AWS Graviton (ARM64) t4g instances yields a **20% savings** over x86 counterparts (t3) while providing superior CPU performance per dollar.
- . **Official Marketplace vs. Custom Script:** Launching the pre-built x86 marketplace image restricts sizing to c5a.xlarge (Option D, ~\$153.16/mo). Deploying on a clean Ubuntu 24.04 Graviton server and using the Wazuh installation assistant (Option B, ~\$65.71/mo) **reduces monthly cloud spend by over 57%!**

3. Installation Strategy 1: Clean OS + Assistant (Best & Cheapest)

Deploying a clean, minimal **Ubuntu 24.04 LTS (Noble Numbat)** or **Ubuntu 22.04 LTS (Jammy Jellyfish)** on a Graviton instance (t4g.medium or t4g.large) and executing the Wazuh installation assistant script is the most cost-effective architecture.

Step 1: Provision the Standalone EC2 Instance

Create a standalone EC2 instance using your OpenTofu templates or the AWS Console:

- **Region:** ap-southeast-5 (Malaysia)
- **AMI:** Ubuntu 24.04 LTS (ARM64 / AARCH64)
- **Instance Type:** t4g.large (2 vCPU, 8 GiB RAM - Highly Recommended)
- **Storage:** 100 GB gp3 volume (3,000 IOPS, 125 MB/s baseline throughput)
- **Elastic IP:** Allocate and associate an Elastic IP with the instance to prevent IP changes on reboot.

Step 2: Configure the Firewall (Security Group)

Create a dedicated Security Group for your standalone Wazuh instance with the following ingress rules:

Port	Protocol	Source	Purpose
22	TCP	Admin Office IP / Jumphost	Secure SSH Administration
443	TCP	Admin Office IP / VPN Subnet	Wazuh Web Dashboard (HTTPS)
1514	TCP	Monitored Servers / Agent Subnets	Wazuh Agent secure communication
1515	TCP	Monitored Servers / Agent Subnets	Wazuh Agent enrollment service
55000	TCP	Admin API client / Integrations	Wazuh RESTful API access

Step 3: Run the All-in-One Installation Assistant

SSH into your instance and run the automated script to install Wazuh indexer, server, and dashboard on the same host:

```
# SSH to the instance
ssh -i "your-key.pem" ubuntu@your-wazuh-elastic-ip

# Switch to root
sudo su -

# Download the installation assistant script
curl -sO https://packages.wazuh.com/4.x/wazuh-install.sh

# Execute the All-in-One installation
bash wazuh-install.sh -a
```

Note: If you are running on t4g.medium (4 GiB RAM), the installer might throw a warning that the hardware requirements are not fully met. You can bypass this check safely for small testing environments by passing the -i flag:


```
bash wazuh-install.sh -a -i
```

Step 4: Record Initial Passwords

At the end of a successful installation, the script will output the default credentials and security configuration. Write them down securely:

```
INFO: --- Installation finished ---
```

```
INFO: Admin password: <GENERATED_PASSWORD>
```

```
INFO: You can access the web interface at https://<YOUR_IP>
```

4. Installation Strategy 2: Official AWS Marketplace AMI

Wazuh provides a pre-built Amazon Machine Image (AMI) pre-configured with Amazon Linux 2023, Wazuh manager, Filebeat-OSS, Wazuh indexer, and Wazuh dashboard out of the box.

According to the official [Wazuh AMI Documentation](#):

Step 1: Subscribe in AWS Marketplace

- . Navigate to [Wazuh All-In-One Deployment in the AWS Marketplace](#).
- . Click **View purchase options**, review the software terms, and click **Subscribe** to accept.
- . Once accepted, select **Launch your software** and choose your region (e.g. ap-southeast-5).

Step 2: Launch the Instance via EC2 Console

- . Navigate to the **EC2 Launch Instance** interface.
- . Under AMIs, go to **Browse more AMIs**, select **AWS Marketplace AMIs**, and search for **Wazuh All-In-One Deployment**.

- . Select **c5a.xlarge** (4 vCPU, 8 GiB RAM) or a comparable instance class.
 - . Set storage capacity to at least **100 GiB gp3**.
 - . Select or create a security group ensuring the standard ports (443, 1514, 1515, 22) are whitelisted as shown in Strategy 1.
 - . Select your SSH key pair and launch the instance.
-

5. Critical Post-Installation & Security Hardening

To ensure secure enterprise operations, perform the following verification and hardening procedures immediately after launching:

A. Accessing the Dashboard

Open your web browser and navigate to:

`https://<YOUR_WAZUH_STATIC_IP>`

- **Username:** admin
- **Password (Strategy 1 - Assistant):** The password generated at the end of the script.
- **Password (Strategy 2 - AMI):** The EC2 **Instance ID** with the first letter capitalised (e.g. if the instance ID is `i-0123456789abcdef0`, the password is `I-0123456789abcdef0`). This safety mechanism prevents initial brute-force attempts.

B. Change Default Passwords

To replace the default master passwords with secure credentials:

```
# Navigate to the tool directory
cd /usr/share/wazuh-indexer/plugins/opensearch-security/tools/

# Run the security administrator script to re-configure passwords
# Refer to Wazuh's Password Management guidelines for specific hashes
```

C. Restrict SSH and Dashboard Ingress

Do not leave port 22 (SSH) and port 443 (Dashboard) open to the public internet (0.0.0.0/0). Whitelist access strictly to your organizational CIDR blocks (such as the Cyberjaya development office IP ranges) or configure access via a secure Bastion/SSM Session Manager.

6. Enrolling Wazuh Agents

Once the standalone Wazuh server is fully active, security agents can be registered on target hosts.

For Linux Targets (Debian/Ubuntu):

```
# Download and install target agent pointing to your standalone Wazuh
server IP
wget https://packages.wazuh.com/4.x/apt/pool/main/w/wazuh-agent/wazuh-
agent_4.7.2-1_amd64.deb
sudo WAZUH_MANAGER='<YOUR_WAZUH_STATIC_IP>' dpkg -i wazuh-agent_4.7.2-
1_amd64.deb

# Start the agent service
sudo systemctl daemon-reload
sudo systemctl enable wazuh-agent
sudo systemctl start wazuh-agent
```


For Windows Targets (PowerShell):

```
Invoke-WebRequest -Uri https://packages.wazuh.com/4.x/windows/wazuh-agent-4.7.2-1.msi -OutFile ${env:tmp}\wazuh-agent.msi
msiexec.exe /i ${env:tmp}\wazuh-agent.msi /q WAZUH_MANAGER="
<YOUR_WAZUH_STATIC_IP>"
```

```
Start-Service -Name "Wazuh"
```

Deep State of Mind (DSOM) For My AI Protocol	Harisfazillah Jamel (LinuxMalaysia)	2026-07-26 Standard: UK English	DBP-standard Bahasa Melayu Malaysia (Piawai)	GNU General Public License v3.0
---	-------------------------------------	---------------------------------	--	----------------------------------

Technology Stack Comparison Guide



Technology Stack Comparison & AWS Mapping Guide

This guide provides a comprehensive technical and architectural comparison between the developer's proposed local/on-premises containerised technology stack and our enterprise-grade, highly-available **AWS Secure 3-Tier Architecture**.

We analyse every layer of the developer's architecture—from Frontend React 19 to AI Integration, Messaging, and Databases—mapping them directly to their corresponding AWS-native managed services or secure hosting tiers in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)**.

Additionally, we evaluate **AWS-native cloud alternatives** for external services (such as Twilio, Meta APIs, and OpenAI integrations) to provide a completely integrated cloud solution with corresponding cost structures.

This guide is integrated with our official [Software Licensing & Technology Risk Register \(TS/MC Series\)](#) to track accepted risk and licensing compliance profiles.

1. High-Level Technology Mapping Matrix

The table below outlines how the developer's stack is mapped to our secure AWS architecture, including recommended AWS-native alternatives for third-party components, and maps them to their respective **TS/MC series risk codes**.

Category	Proposed Developer Technology	Our AWS Secure 3-Tier Target	AWS-Native Cloud Alternative
AI Integration	LangChain4j, OpenAI-compatible Provider (Chat, Embeddings, Function Calling)	Hosted on ASG Compute Tier (Graviton t4g.xlarge or GPU-backed instances), calling secure APIs or local inference.	Amazon Bedrock (Serverless access to Anthropic Claude, Meta Llama 3, Cohere, etc with local data privacy/sovereignty compliance).
Backend	Java 21, Spring Boot 3.5.12, Spring Framework 6.2, Spring Security 6.x, Spring Data JPA, Hibernate 6.x	Deployed within Private App Subnets under an Auto Scaling Group (ASG) of Graviton ARM64 instances, managed via AWS Systems Manager (SSM) .	<i>Same application code, but highly scalable and secure.</i>

Category	Proposed Developer Technology	Our AWS Secure 3-Tier Target	AWS-Native Cloud Alternative
Frontend	React 19, Vite, TypeScript, TanStack Router, Ky, shadcn/ui	S3 Bucket + Amazon CloudFront (Global CDN) for static asset distribution OR served via the private ASG behind the Application Load Balancer (ALB) .	Amazon S3 + Amazon CloudFront (Saves compute costs and improves global latency via edge caching).
Mobile Application	React Native, Expo, HeroUI	Connects securely to the Application Load Balancer (ALB) endpoint protected by AWS WAFv2 .	<i>Client-side runtime, standard HTTPS API endpoint mapping.</i>
Authentication & Authorization	Spring Security, JWT Authentication, OAuth 2.0, RBAC	Managed within the Spring Boot backend layer with token caches stored in Amazon ElastiCache Redis .	Amazon Cognito Pools (Serverless user directory, OAuth 2.0 flow generation, MFA, and native security auditing).

Category	Proposed Developer Technology	Our AWS Secure 3-Tier Target	AWS-Native Cloud Alternative
Database	PostgreSQL 16, Docker, Flyway	Fully managed Amazon RDS PostgreSQL (Multi-AZ) in isolated private subnets, with schema migrations coordinated via standalone nodes or CI/CD pipelines.	Amazon RDS for PostgreSQL (Multi-AZ synchronous replication, automated backups, automated patching, zero-SPOF).
Cache / Performance	Redis (Rate limiting, usage counters, session/cache management)	Amazon ElastiCache for Redis (Secure, multi-AZ, and license-compliant caching layer offering 20% lower pricing than legacy Redis OSS).	Amazon ElastiCache for Redis (Dedicated managed caching cluster)
Build Tools	Backend: Maven; Frontend: npm + Vite	Integrated with GitHub Actions and GitLab CI/CD with shared EFS mounts for build caching.	AWS CodeBuild (Fully managed build service)

Category	Proposed Developer Technology	Our AWS Secure 3-Tier Target	AWS-Native Cloud Alternative
Deployment & Infrastructure	Docker, Docker Compose, PostgreSQL, Redis, Backend, Frontend, AI services	Multi-AZ modular infrastructure provisioned via OpenTofu (ALB, ASG, VPC, S3, ElastiCache, IAM, WAFv2, and Standalone compute).	AWS ECS (Elastic Container Service) on Fargate or EKS for managed container orchestration.
Enterprise Standards	Jakarta EE 10, Servlet 6.0, JPA 3.1, Bean Validation 3.0	Run inside the standard Spring Boot JVM running on hardened Ubuntu 26.04 LTS (ASIMP framework) .	<i>Same enterprise compliance parameters.</i>
AI Knowledge Base / RAG	RAGFlow (Document ingestion, parsing, chunking, embedding generation, vector search, retrieval-augmented generation)	Dedicated AI Compute Tier (ASG / Standalone) inside private subnets connected to Amazon EFS and OpenSearch Service (or local vector database).	Amazon Bedrock Knowledge Base (Serverless RAG pipeline utilizing OpenSearch Serverless and Bedrock embeddings).

Category	Proposed Developer Technology	Our AWS Secure 3-Tier Target	AWS-Native Cloud Alternative
Messaging / Omnichannel	Twilio for WhatsApp (WhatsApp Business messaging, inbound/outbound, webhooks, media, callbacks)	Integrated in the backend code via Twilio Java SDK, with webhook receivers protected by AWS WAFv2 .	AWS End User Messaging (Social Channels) (Direct AWS-native WhatsApp Business API integration, eliminating Twilio as an intermediary).
Social Media Messaging	Meta for Developers / Meta Graph API (Facebook Messenger, Instagram Messaging API, webhook subscription, page tokens)	Outbound requests routed via NAT Gateway ; inbound webhooks handled via the ALB or a serverless proxy.	AWS Lambda + Amazon API Gateway (Serverless, auto scaling webhook receivers that scale to zero or inactive to minimise idle costs).
External API Integration	REST API, Webhook, OAuth 2.0, API token management, retry mechanism, request validation	Outbound requests secure-routed through AWS NAT Gateway ; secrets managed inside AWS Secrets Manager .	Amazon API Gateway with custom authorizers and integration throttling.

Category	Proposed Developer Technology	Our AWS Secure 3-Tier Target	AWS-Native Cloud Alternative
Security & SIEM Operations	Traditional/Vendor SIEM platform	Hardened standalone Wazuh SIEM EC2 instance whitelisted strictly for Cyberjaya developer CIDRs.	AWS Security Hub Amazon GuardDuty (Consolidated multi-account cloud-native threat detection).

2. Technical Mapping Details & Architectural Transitions

A. Frontend Web Tier (React 19 + Vite)

- **Proposed Stack:** React 19, Vite, TypeScript, TanStack Router, Ky, shadcn/ui.
- **On-Premises / VM Deployment:** Typically served via Nginx inside a Docker container running on a public-facing virtual machine.
- **AWS Enterprise Architecture [TS-05]:**
 - **Static Site Hosting (CloudFront + S3):** We compile the Vite build artifacts and push them to a secure **Amazon S3** bucket. An **Amazon CloudFront** distribution sits in front of the bucket, serving assets globally through edge locations. This eliminates the need to run EC2 compute for the frontend, reduces latency, and protects against direct S3 bucket scraping.
 - **Secure Dynamic Routing:** API requests (such as `/api/*`) are routed by CloudFront directly to the **Application Load Balancer (ALB)**, while static routes are handled at the edge, creating a unified origin configuration.

B. Backend App Tier (Spring Boot 3.5.12 + Java 21)

- **Proposed Stack:** Spring Boot 3.5.12, Spring Framework 6.2, Spring Security 6.x, Spring Data JPA, Hibernate 6.x, Jakarta EE 10, Maven.
- **On-Premises / VM Deployment:** Executed as a systemd service or Docker container on Server 02 (4 vCPU, 16GB RAM) exposed to the web.
- **AWS Enterprise Architecture [TS-05]:**
 - **Stateless Auto-Scaling Compute:** Deployed across an Auto Scaling Group (ASG) of **AWS Graviton ARM64** instances (e.g., t4g.xlarge to match the 4 vCPU/16GB RAM spec) in **Private App Subnets**.
 - **SSM-Only Access:** Direct SSH is disabled. All staging, auditing, and maintenance is handled passwordlessly via **AWS Systems Manager (SSM)**.
 - **ASIMP Hardening:** Operating systems are locked down to **Ubuntu 26.04 LTS** and audited via the **ASIMP (Ansible System Integrity Management Platform)** framework against CIS Level 2 benchmarks.

C. Caching & Performance (Redis)

- **Proposed Stack:** Redis (used for rate limiting, usage counters, session/cache management).
- **On-Premises / VM Deployment:** A local Docker Redis container running on the same host, with no high availability.
- **AWS Enterprise Architecture [TS-06]:**
 - **Amazon ElastiCache for Valkey:** We map Redis directly to **ElastiCache Valkey** (cache.t4g.medium). Valkey is a fully open-source, license-compliant key-value cache engine that provides full Redis API compatibility but at **20% lower hourly rates** on AWS.
 - **Transit and At-Rest Encryption:** Subnet-isolated caching clusters with forced TLS and strict ingress restricted strictly to the compute security group.

D. Relational Database (PostgreSQL 16 + Flyway)

- **Proposed Stack:** PostgreSQL 16, Docker, Flyway.
- **On-Premises / VM Deployment:** PostgreSQL running inside Docker on Server 04 (4 vCPU, 16GB RAM), risking single-point-of-failure database corruption.
- **AWS Enterprise Architecture [TS-06]:**
 - **Amazon RDS PostgreSQL 16 (Multi-AZ):** We migrate local Postgres to a fully managed **Amazon RDS PostgreSQL** cluster (Multi-AZ) sized at db.m6g.xlarge (4 vCPU, 16GB RAM). RDS handles automatic Multi-AZ synchronous replication, daily automated snapshots, point-in-time recovery (PITR), and security patching.
 - **Network Isolation:** Isolated inside **Private Database Subnets** with zero public internet connectivity. Schema migrations are automatically executed at build-time via GitHub/GitLab runners or a secure Standalone staging EC2 instance.


E. AI Knowledge Base / RAG Platform (RAGFlow)

- **Proposed Stack:** RAGFlow (DeepDoc parser, OCR, semantic chunking, vector search).
- **On-Premises / VM Deployment:** Heavy Docker Compose setup relying on local GPUs or CPU fallbacks, causing processing delays and high latency.
- **AWS Enterprise Architecture [TS-06]:**
 - **GPU Compute & Shared Storage:** RAGFlow executors run on dedicated **G5 instances (g5.xlarge with NVIDIA A10G 24GB VRAM)**. We mount an **Amazon EFS (Elastic File System)** volume across the instances to store shared AI model weights and cached document parsing structures.
 - **Dense Vector Search:** High-performance indexing and dense vector queries are managed natively via **Amazon OpenSearch Service** (with k-NN search enabled) or an isolated PostgreSQL vector database (pgvector extension) on RDS.

3. High-Fidelity AWS-Native Alternatives for “Extra” Items

The developer’s technology stack contains several “extra” third-party SaaS integrations (such as Twilio for WhatsApp, Meta Graph APIs, and OpenAI endpoints).

In enterprise deployments, relying on external SaaS platforms introduces data sovereignty issues, external billing overhead, and architectural latency. Below are the architectural transitions and benefits of choosing **AWS-native replacements**.



Alternative A: Amazon Bedrock (Replaces OpenAI/LangChain4j Endpoints - MC-01 / MC-02 / TS-02)

Instead of routing enterprise chat and sensitive document embeddings to third-party OpenAI endpoints over the public internet, utilise **Amazon Bedrock** loaded with state-of-the-art **Qwen3** or Claude models.

- **AWS Architecture [MC-01]:**
 - Bedrock provides high-speed, serverless API access to leading foundation models (such as **Qwen3 LLM Inference**) entirely within your AWS network.
 - Since traffic never leaves the AWS boundary, this aligns perfectly with the **Malaysian Personal Data Protection Act (PDPA) 2010** data residency and sovereignty requirements.

- Bedrock integrates natively with LangChain4j via standard AWS SDK dependencies, mitigating the commercial SLA lack of LangChain4j by using enterprise-secured network pipelines.
- **Embedding & Query Optimisation [MC-02]:**
 - For static document databases, perform a one-time high-efficiency indexing sweep using Qwen3 Embedding models and store vectors locally in RDS PostgreSQL (pgvector), minimizing ongoing runtime model consumption fees.

Alternative B: AWS End User Messaging (Replaces Twilio for WhatsApp)

Instead of routing WhatsApp Business chats, media files, and webhooks through Twilio as a middleware, utilise **AWS End User Messaging (Social Channels)**.

- **AWS Architecture:**
 - AWS End User Messaging connects your applications directly to the Meta WhatsApp Business API.
 - This eliminates Twilio's per-message platform markup fees, consolidates WhatsApp bills directly into your single AWS invoice, and allows secure, automated webhook handling directly via AWS Lambda.

Alternative C: Amazon Cognito User Pools (Replaces Self-Managed JWT/Spring Security)

Instead of managing sensitive password hashing, JWT token signing, rotation, MFA, and database user tables inside the custom Java Spring Boot code, utilise **Amazon Cognito**.

- **AWS Architecture:**
 - Spring Security delegates authentication to **Amazon Cognito User Pools** via standard OAuth 2.0 / OIDC protocols.
 - Cognito acts as a secure, serverless identity provider that automatically handles multi-factor authentication (MFA), user registration, password reset flows, and brute-force protection.

- Eliminates the security risk of storing credentials directly in RDS and reduces backend JVM memory footprint.

Alternative D: Amazon API Gateway & AWS Lambda (Replaces Custom Webhook Receivers)

The Meta Graph API (Facebook Messenger/Instagram) uses highly dynamic webhooks that send intense bursts of traffic during high-volume communication spikes, saturating Spring Boot JVM threads.

- **AWS Architecture:**

- Public webhooks are routed to **Amazon API Gateway** and processed by serverless **AWS Lambda** functions.
- API Gateway validates incoming request signatures, and AWS Lambda scales instantly to handle tens of thousands of concurrent messages, before routing the parsed messages into an **Amazon SQS** queue for the backend ASG to process asynchronously.
- Minimises idle compute cost (scales to \$0.00 when there are no active messages).

4. Architectural Comparison Summary

The matrix below compares the security, scalability, and operational realities of the developer's proposed local VM stack with our AWS Secure 3-Tier mapping (including the AWS-native alternatives).

Architectural Dimension	Proposed Developer VM / Docker	Our AWS S
High Availability	Single VM per service. Host failure causes total outage.	Multi-AZ i layers. Au
Network Isolation	Compute & DB run on public IPs. High SSH and port-scan risk.	Compute & Ingress re
Security Hardening	Manual OS patching and updates. Unaudited configurations.	ASIMP-harc Continuous
Scalability	Static server resources. Manual resizing is required.	Dynamic ho CPU/memor
API Webhook Resilience	Custom Spring Boot web endpoints susceptible to JVM OOM.	Serverless Bursty Wha
Identity Protection	Self-managed JWT token store. Risks with key rotation & storage.	Amazon Co automated
AI Compliance & Residency	OpenAI requests route overseas. Subject to foreign data transit.	Amazon Be (ap-southe

Software Licensing & Technology Risk Register (TS/MC Series)

#aws

#cloud

#architecture

#licensing

#compliance

#siem

#wazuh

#qwen3

#bedrock

#langchain4j

#ragflow

#langfuse

#costing

Software Licensing & Technology Risk Register (TS/MC Series)

This guide establishes the official **Software Licensing & Technology Risk Register (TS/MC Series)** for our AWS Secure 3-Tier Architecture. It documents the core architectural decisions, commercial/open-source licensing assessments, and operational risks associated with our enterprise-grade web and AI infrastructure.

By classifying key software and AI model components under the **TS (Technology Stack)** and **MC (Model Consumption)** series, we ensure complete traceability, clear ownership of technical risk, and alignment with regulatory standards such as the **Malaysian Personal Data Protection Act (PDPA) 2010**.

1. Technology Stack (TS Series) Registry

The TS Series identifies, tracks, and manages the licensing risks and operational models of software components deployed across our 3-tier architecture.

TS SERIES RISK REGISTER			
ID	Component	Risk Level	Action Taken
TS-02	LangChain4j Framework	Medium-Low	Accepted Risk
TS-04	Wazuh SIEM Self-Hosting	Low	Budget Savings
TS-05	Software Licensing Framework	Low	Audit Standard
TS-06	Self-hosted DB & AI Ops	Medium	Internal Option

TS-02: LangChain4j Integration & SLA Risk

- **Classification:** AI Integration Layer
- **Licensing Profile:** Apache License 2.0 (Highly permissive open-source)
- **Risk Context:** LangChain4j lacks an official commercial Service Level Agreement (SLA) or vendor-backed corporate support structure.
- **Risk Evaluation & Accepted Risk:**
 - **The Risk:** In the event of critical framework bugs, security vulnerabilities (CVEs), or breaking API changes by downstream LLM providers, there is no contractually bound vendor to provide emergency patches or guarantee response times.
 - **Why the Risk is Accepted:** LangChain4j is the de facto standard for building AI/LLM applications in the Java ecosystem (supporting Java 21, Spring Boot 3.5.12, and Spring Security 6.x natively). It has a highly active, robust open-source community, wide enterprise adoption, and regular release cycles.
 - **Mitigation Strategy:**
 1. **Forking Strategy:** Maintain an internal git fork of the specific LangChain4j version in use to allow local patch application independent of upstream releases if needed.
 2. **Dependency Pinning:** Pin dependency versions in Maven pom.xml to prevent automated, untested minor/patch version upgrades.

3. **Abstraction Layer:** Use clean interface abstractions in our Java code so that the underlying LLM orchestration framework can be swapped or modified with minimal business logic disruption.

TS-04: Standalone Wazuh SIEM Hosting

- **Classification:** Enterprise Security & Compliance
- **Licensing Profile:** GPL v3 (GNU General Public License)
- **Risk Context:** Replaces traditional, high-cost third-party SaaS commercial SIEM subscriptions with a standalone self-hosted instance in AWS.
- **Cost & Budget Impact:**
 - By deploying Wazuh on a dedicated, hardened AWS Graviton ARM64 instance (t4g.large), we eliminate the vendor markup fees and per-GB ingestion surcharges common in SaaS SIEM models.
 - This replaces the “vendor SIEM” budget line item under Security, **reducing monthly operational costs by over 57%** (operating at approx. **\$65.71 USD / RM 295.70 MYR per month** on Graviton, compared to \$153+ USD/month on legacy x86 setups).
- **Mitigation Strategy:**
 - Standardise security configurations, restrict public internet access (restrict port 22 and 443 strictly to Cyberjaya developer CIDRs), and backup OpenSearch/Wazuh indexer directories using automated **AWS Backup** schedules.

TS-05: Technology Stack Software & Licensing Framework

- **Classification:** Corporate Compliance & Licensing Audit
- **Licensing Profile:** Governance Framework
- **Context:** Manages open-source and commercial software compliance across the entire 3-tier system to prevent licensing contamination (e.g. copyleft license leakage into proprietary backend code).
- **Compliance Actions:**

- Restrict production runtimes to permissive licenses such as **Apache 2.0, MIT, and BSD**.
- Restrict the introduction of AGPL or strong copyleft licenses in the backend API layer.
- Maintain a dynamic **Software Bill of Materials (SBOM)** generated during GitHub Actions and GitLab CI/CD build phases to automatically flag compliance anomalies.

TS-06: Self-hosted Database Operations (Internal Option)

- **Classification:** Storage & In-House Data Management
- **Licensing Profile:** Open-Source Permissive / Dual-Licensing (PostgreSQL License, BSD-3, Apache 2.0)
- **Risk Context:** Running databases (PostgreSQL, Valkey/Redis) and AI orchestrators (Langfuse, RAGFlow) in-house as self-hosted containers/instances rather than relying on commercial SaaS database vendors.
- **Trade-Off Analysis:**
 - **Benefits:** Complete operational control, zero third-party data transit risk, and elimination of premium SaaS database pricing tiers.
 - **Sovereignty Aligned:** Keeps sensitive AI session telemetry, prompt logs, and document vectors strictly localized within our private subnets in the AWS Malaysia (ap-southeast-5) region, satisfying strict local PDPA data residency mandates.
- **Mitigation Strategy:**
 - Map self-hosted designs directly to managed **Amazon RDS PostgreSQL (Multi-AZ with pgvector)** and **Amazon ElastiCache for Valkey** for production workloads, while retaining containerized self-hosted setups as a cost-optimized “Internal Option” for development and local testing.

2. Model Consumption (MC Series) Registry

The MC Series governs how foundational AI models (LLMs, embedding engines) are consumed, focusing on cost efficiency, data residency, and performance optimisation.

MC SERIES RISK REGISTER			
ID	Component	Optimization	Sovereign Gate
MC-01	Qwen3 LLM Inference	Serverless	Bedrock Private
MC-02	Qwen3 Embed, Index & Query	One-time/Low	Batch/Cached

MC-01: Qwen3 LLM Inference via Amazon Bedrock

- **Classification:** Large Language Model (LLM) Inference
- **Consumption Model:** Serverless API / Provisioned Throughput
- **Architectural Flow:**
 - All conversational queries, reasoning loops, and prompt chains are processed using the state-of-the-art **Qwen3 LLM** (or equivalent highly-optimised multilingual models) running natively within **Amazon Bedrock**.
- **Sovereignty & Compliance Mapping:**
 - By routing Qwen3 model inference directly through Amazon Bedrock, all prompt payloads, context windows, and generated tokens are processed entirely within secure, private AWS regional boundaries (ap-southeast-5).
 - This eliminates data leakage risks associated with routing sensitive corporate information over the public internet to external, third-party AI startups. It provides a clean, audit-ready compliance path for Transfer Impact Assessments (TIAs) under PDPA guidelines.

MC-02: Qwen3 Embedding, Indexing, and Query Optimisation

- **Classification:** Dense Vector Search & Retrieval-Augmented Generation (RAG)
 - **Consumption Model:** Static Ingestion + Low Ongoing Query Costs
 - **Operational Profile:**
 - **Static Documents (One-time Sweep):** Ingest static corporate documents, technical manuals, and knowledge bases using a high-density, one-time embedding sweep using Qwen3 Embedding models.
 - **Storage:** Store the generated high-dimensional dense vectors in **Amazon RDS PostgreSQL (pgvector)** or a secure vector index.
 - **Ongoing Querying (Low Cost):** Since the foundational document corpus is static, ongoing daily operational costs are limited to the tiny vector queries processed during user searches, keeping the average API charge extremely low.
 - **Optimisation Lever:**
 - Utilize batch processing for bulk document ingestion during off-peak hours to take advantage of lower pricing tiers and prevent compute thread saturation in the main Spring Boot application ASG.
-

3. Enterprise Governance & Action Items

To maintain a secure compliance posture, the engineering team must execute the following action items regularly:

- . **Conduct Licensing Sweeps (Quarterly):** Run automated Maven and npm dependency license audits to ensure no copyleft components have bypassed the TS-05 licensing framework.
 - . **Review LangChain4j Versions (Monthly):** Audit the active LangChain4j fork for security patches and ensure any local overrides are documented under the TS-02 risk mitigation log.
 - . **Optimise Bedrock API Allocations (Daily):** Monitor Amazon CloudWatch metrics for Bedrock API invocation counts and token usage to identify opportunities for prompt caching and further lower MC-01/MC-02 ongoing costs.
-

Deep State of Mind (DSOM) For My AI Protocol	Harisfazillah Jamel (LinuxMalaysia)	2026-07-26 Standard: UK English	DBP- standard Bahasa Melayu Malaysia (Piawai)	GNU General Public License v3.0
--	---	--	--	--

Strategic Comparative Review: AWS-Native Managed Platform vs. Self-Hosted Custom Stack

#aws

#cloud

#architecture

#costing

#licensing

#compliance

#postgresql

#valkey

#bedrock

#cognito

#wazuh

#sovereignty

Strategic Comparative Review: AWS-Native Managed Platform vs. Self-Hosted Custom Stack

This review provides a comprehensive, high-fidelity strategic analysis comparing an **AWS-Native Managed Solution** against a **Self-Hosted / On-Premises Custom Stack** inside the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)**.

When building an enterprise-grade, highly available 3-tier architecture for modern web and AI workloads, organisations face a critical architectural decision: leverage AWS managed services to outsource operational risk, or deploy and run an equivalent open-source stack in-house to maximise raw hardware control and avoid vendor lock-in.

This guide looks at the **whole picture**, comparing the entire application lifecycle—including presentation, compute, database, caching, generative AI, messaging, identity, security, and disaster recovery. It integrates our [Software Licensing & Technology Risk Register \(TS/MC Series\)](#), [RDS vs. Percona PostgreSQL Comparison](#), and [Disaster Recovery Playbook](#) into a single, cohesive decision framework.

1. Executive Summary & The Big Picture

The fundamental trade-off between AWS-Native Managed Services and Self-Hosted/Custom Stacks is **Operational Leverage vs. Raw Hardware Control**:

- **AWS-Native Managed Services** abstract infrastructure complexity. AWS guarantees high availability, synchronous multi-AZ replication, automatic failovers, and compliance boundaries. This dramatically reduces the engineering headcount (OpEx) required for Day-2 maintenance, allowing lean teams to focus entirely on product innovation.
- **Self-Hosted Custom Stacks** (whether deployed on raw EC2 instances or local on-premises hardware) eliminate managed service markups and vendor APIs. However, they introduce immense operational complexity. High Availability (HA) must be engineered manually using clustering tools (such as Patroni, etcd, PgBouncer, and custom keepalived scripts). This significantly increases specialized engineering labor costs.

High-Level Architectural Mapping Matrix

The side-by-side matrix below illustrates how the two solutions map across every layer of the architecture, along with their respective **TS/MC Series risk and compliance codes**.

Architectural Layer	AWS-Native Managed Solution	Self-Hosted / Custom Stack	Strategic Trade-Off	Compliance & Risk Code
Presentation Web Layer	Amazon S3 + CloudFront (CDN)	Nginx inside VM / Container	CloudFront improves edge latency and protects against direct scraping; Nginx on VMs requires manual patching.	<u>TS-05</u>
Application Compute	Auto Scaling Groups (ASG) on ARM64 Graviton	Dedicated EC2 / Standalone VMs	ASG offers dynamic scaling and elasticity; dedicated VMs run idle at high cost.	<u>TS-05</u>
Database Tier	Amazon RDS PostgreSQL (Multi-AZ)	Percona Server for PostgreSQL with Patroni & etcd on EC2	RDS automates storage replication and failovers; Percona on EC2 requires dedicated DBRE labor.	<u>TS-06</u>

Architectural Layer	AWS-Native Managed Solution	Self-Hosted / Custom Stack	Strategic Trade-Off	Compliance & Risk Code
Session Cache Layer	Amazon ElastiCache for Valkey	Self-hosted Redis OSS Docker Container	ElastiCache Valkey is fully managed, license-compliant, and 20% cheaper than Redis; Redis OSS runs on-host without HA.	<u>TS-06</u>
Generative AI & RAG	Amazon Bedrock (Serverless Qwen3 / Claude)	RAGFlow + Langfuse on GPU instances (g5.xlarge) + EFS	Bedrock is serverless, private, and out-of-the-box; RAGFlow offers deep document parsing (DeepDoc) but requires GPU compute.	<u>TS-02, MC-01, MC-02</u>

Architectural Layer	AWS-Native Managed Solution	Self-Hosted / Custom Stack	Strategic Trade-Off	Compliance & Risk Code
Identity & Auth	Amazon Cognito User Pools	Custom JWT + Spring Security with database tables	Cognito handles MFA, token rotation, and password flows serverless; custom JWT requires database storage and encryption.	<u>TS-05</u>
Messaging & Webhooks	AWS End User Messaging & API Gateway + Lambda	Twilio WhatsApp API & Spring Boot dynamic endpoints	AWS-native eliminates Twilio per-message markup; serverless API Gateway absorbs dynamic bursts safely.	<i>Third-Party Integration</i>

Architectural Layer	AWS-Native Managed Solution	Self-Hosted / Custom Stack	Strategic Trade-Off	Compliance & Risk Code
Security & SIEM	AWS Security Hub + GuardDuty	Hardened Wazuh SIEM on Graviton EC2	GuardDuty offers cloud-native detection; Wazuh SIEM provides affordable host-level intrusion detection and auditing.	<u>TS-04</u>
Disaster Recovery (DR)	Multi-AZ with automated backups	AWS DRS (Strategy E) continuous replication	Managed Multi-AZ RDS ensures sub-minute recovery; DRS replicates block volumes into low-cost staging subnets.	<u>TS-06</u>


2. Technical Comparison Per Layer

2.1 Compute and Presentation Layer

- **AWS-Native Solution:** App code runs inside stateless Auto Scaling Groups (ASG) using cost-optimized **ARM64 Graviton instances** (t4g.xlarge or c7g.xlarge). Static frontend assets are compiled and hosted in **Amazon S3**, distributed globally via **Amazon CloudFront**.
 - *Benefits:* Scales horizontally based on CPU or memory saturation profiles. There are no idle VM compute costs for static web files. CloudFront absorbs distributed denial of service (DDoS) attacks at the edge, integrated with **AWS WAFv2** for active rate-limiting and OWASP protection.
- **Self-Hosted Solution:** A monolithic virtual machine (such as a single Ubuntu VM running Nginx and Spring Boot).
 - *Drawbacks:* Scaling is vertical, requiring manual VM sizing upgrades and scheduled downtimes. Ingress traffic directly hits the virtual machine, exposing ports (SSH/HTTP) to scanning and exploits. The VM operates at full cost even during periods of zero traffic.

2.2 Database Tier (RDS PostgreSQL vs. Percona on EC2)

- **AWS-Native Solution: Amazon RDS PostgreSQL (Multi-AZ).** High availability is managed synchronously at the block storage level. If AZ-A fails, RDS automatically points the CNAME DNS endpoint to the standby node in AZ-B within **60 to 120 seconds**. Backups are automated, continuous, and integrated with Point-in-Time Recovery (PITR).
- **Self-Hosted Solution: Percona Server for PostgreSQL on EC2 (Patroni, etcd, PgBouncer).** High availability requires a multi-node cluster (two database nodes and a third lightweight consensus node running **etcd**). **Patroni** orchestrates failovers, reducing database redirection latency to **10 to 30 seconds**. Backups are managed using **pg_backrest** streaming to an S3 bucket.
 - *The Catch:* While saving on raw infrastructure, the self-hosted cluster requires expert database engineering (DBRE) to configure, patch, and maintain Patroni playbooks, etcd consensus states, and pg_backrest recovery pipelines.



2.3 Caching Tier (ElastiCache Valkey vs. Self-hosted Redis)

- **AWS-Native Solution: Amazon ElastiCache for Valkey.** Fully managed key-value caching on Graviton (cache.t4g.micro or cache.t4g.medium). Valkey is a modern, open-source replacement for Redis OSS, offering full API compatibility but **20% lower pricing** on AWS.
- **Self-Hosted Solution: Redis OSS on EC2/Docker.** Runs as a local Docker container alongside compute services.
 - *Drawbacks:* The cache layer is single-point-of-failure. If the container or host crashes, all session memory, user rate limits, and web tokens are wiped, causing massive application downtime.

2.4 Generative AI & RAG Tier (Amazon Bedrock vs. GPU EC2)

- **AWS-Native Solution: Amazon Bedrock (Qwen3).** Serverless API calling. All prompt payloads, context windows, and embeddings are processed securely within the AWS regional boundary.
 - *Benefits:* Complete network isolation. Pay-per-token model ensures costs scale only with usage, avoiding any idle GPU hosting fees.
- **Self-Hosted Solution: RAGFlow + Langfuse on dedicated GPU compute.** Deployed on **G5 instances (g5.xlarge with NVIDIA A10G 24GB VRAM)**, mounting **Amazon EFS** for shared AI models and document parsing.
 - *Benefits:* Superior, customized parsing (DeepDoc parser) and OCR. Complete control over prompt pipelines and logging via local Langfuse containers.

2.5 Security and Threat Detection (Cloud-Native vs. Wazuh SIEM)

- **AWS-Native Solution:** AWS GuardDuty, Security Hub, and AWS Config.
 - *Benefits:* Zero-agent deployment. Deeply integrated into cloud control planes, scanning VPC Flow Logs, DNS queries, and IAM roles dynamically.
- **Self-Hosted Solution: Standalone Wazuh SIEM on EC2.**
 - *Benefits:* Highly cost-effective host-level intrusion detection system (HIDS) and compliance auditor. Running Wazuh on a Graviton t4g.large instance **saves over 57% of license and platform fees** compared to traditional enterprise SIEM platforms.

3. Sovereign Compliance & Legal Audits

Operating inside the AWS Malaysia region (ap-southeast-5) mandates strict compliance with the **Malaysian Personal Data Protection Act (PDPA) 2010** and the **2025 CBPDT Guidelines**.

SOVEREIGN COMPLIANCE GATE	
AWS-Native Managed Solution	Self-Hosted / Custom Stack
- Data resides locally inside ap-southeast-5 private subnets.	- Local host-level control matches residency mandates perfectly.
- Satisfies Section 129 PDPA.	- Requires complex local OS rules and audits for compliance.
- Bedrock API traffic is isolated from foreign transit networks.	- Backup streaming uses S3.

- . **Data Residency (Section 129 PDPA):** Both solutions comply with local residency mandates by deploying resources natively inside the Kuala Lumpur region (ap-southeast-5). Payloads are held locally.
- . **Transfer Impact Assessments (TIAs):** Sending sensitive corporate data or customer personally identifiable information (PII) to external, third-party US-based AI models (such as OpenAI endpoints over the public internet) violates strict data transit boundaries. Amazon Bedrock processes Qwen3 models locally in Malaysia, satisfying the requirements of the TIA.
- . **Cryptographic Isolation:** Both models can use AWS Key Management Service (KMS) with customer-managed keys (CMK) to implement envelope encryption on EBS, RDS, and S3 volumes, ensuring total data privacy.

4. Comprehensive Financial Blueprint (TCO Comparison)

To compare the Total Cost of Ownership (TCO) in the Malaysia region, we evaluate the **Enterprise Production Multi-AZ Environment** over 1-year and 3-year timelines.

We assume an exchange rate of **1 USD $\approx$ 4.45 MYR** and calculate both the raw infrastructure charges and the specialized engineering labor (DBRE/SecOps) required to run and maintain the custom stack.

4.1 1-Year Financial Model (USD & MYR)

Sizing & Cost Component	Option A: AWS-Native Managed Platform	Option B: Self-Hosted Custom Stack (EC2)	Financial Analysis
Ingress Load Balancing	\$45.63 / mo (ALB + WAF + 5M reqs)	\$16.43 / mo (Single VM Nginx)	S3 + CloudFront + ALB + WAF offers superior DDoS protection and edge performance.
Compute / Host Cost	\$635.10 / mo (3x c7g.2xlarge ASG)	\$404.70 / mo (2x c7g.2xlarge database hosts + 1x t4g.micro etcd)	Self-hosted compute saves raw instance charges but lacks dynamic auto-scaling elastic margins.
Database & Caching	\$748.08 / mo - RDS db.m7g.xlarge Multi-AZ (\$492.02) 250GB gp3 Multi-AZ storage (\$57.50) - Valkey cache.r7g.large Multi-AZ (\$198.56)	\$216.00 / mo 2x 100GB gp3 database volumes (\$16.00) - Valkey / Redis self-installed on compute (\$0.00) - pg_backrest streaming S3 (~\$200.00)	Managed RDS Multi-AZ storage costs are higher due to synchronous mirroring, but eliminate data loss risk.
Storage & Backups	\$37.50 / mo (EFS + AWS Backup)	\$30.00 / mo (EFS + self-managed scripts)	AWS Backup coordinates automated snapshot rotations.

Sizing & Cost Component	Option A: AWS-Native Managed Platform	Option B: Self-Hosted Custom Stack (EC2)	Financial Analysis
Engineering Labor (OpEx)	\$150.00 / mo Estimated 2 DBA/Ops monitoring hours	\$1,500.00 / mo Estimated 15 expert DBRE/SecOps hours (Patroni/etcd checks, Wazuh upgrades)	The Real Difference: Self-hosting introduces significant human maintenance overhead.
TOTAL MONTHLY COST (USD)	\$1,616.31 USD / month	\$2,167.13 USD / month	When engineering labor is factored in, the AWS-Native platform is more cost-effective.
TOTAL MONTHLY COST (MYR)	~RM 7,192.58 MYR / month	~RM 9,643.73 MYR / month	<i>Calculated at 1 USD ≈ 4.45 MYR</i>
TOTAL 1-YEAR TCO (USD)	\$19,395.72 USD	\$26,005.56 USD	AWS-Native saves \$6,609.84 USD (RM 29,413.79 MYR) in Year 1.

4.2 3-Year Total Cost of Ownership (TCO) Comparison

Over a 3-year lifecycle, organisations can apply **AWS Compute Savings Plans** and **RDS Reserved Instances** to achieve massive discounts (up to 34% off compute and 30% off DB storage).

3-YEAR TCO COMPARISON SUMMARY	
Option A: AWS-Native (Optimised)	Option B: Self-Hosted (EC2)
Infrastructure (3-Yr SP): \$34,030.80	Infrastructure (Raw): \$24,016.80
Labor (Optimised): \$5,400.00	Labor (Ops): \$54,000.00
Total: \$39,430.80 USD (RM 175,467.06)	Total: \$78,016.80 USD (RM 347,174.76)

- **Option A: AWS-Native Managed Solution (3-Year Optimised):**
 - **Infrastructure Cost:** \$945.30 / month × 36 = **\$34,030.80 USD** (MYR 151,437.06)
 - **Engineering Labor:** \$150.00 / month × 36 = **\$5,400.00 USD** (MYR 24,030.00)
 - **Total 3-Year TCO:** **\$39,430.80 USD (RM 175,467.06 MYR)**
- **Option B: Self-Hosted Custom Stack (3-Year Operations):**
 - **Infrastructure Cost:** \$667.13 / month × 36 = **\$24,016.80 USD** (MYR 106,874.76)
 - **Engineering Labor:** \$1,500.00 / month × 36 = **\$54,000.00 USD** (MYR 240,300.00)
 - **Total 3-Year TCO:** **\$78,016.80 USD (RM 347,174.76 MYR)**

Financial Impact:

By opting for the AWS-Native Managed Solution, organisations save **\$38,586.00 USD (~RM 171,707.70 MYR)** over 36 months, representing a **49.4% cost reduction**. This proves that managed database premiums are vastly outweighed by the heavy labor burden of running a custom, high-availability architecture.

5. Strategic Recommendation Matrix

The matrix below serves as a final guide for selecting the optimal architecture based on organizational priorities, talent availability, and performance needs.

Organisational Driver	Choose AWS-Native Managed Solution	Choose Self-Hosted / Custom Stack
Engineering Team Sizing	Lean teams (1-5 engineers) without a dedicated database reliability engineer (DBRE).	Large, specialized infrastructure teams with dedicated DBA, SecOps, and virtualization experts.
Time-to-Market (TTM)	Extremely fast. Production-ready multi-AZ environments launch in minutes.	Slow. Requires manual configuration of cluster consensus, replication lag rules, and monitoring hooks.
Performance Fine-Tuning	Standard. Abstracts the operating system filesystem and memory parameters.	Advanced. Allows custom tuning of the kernel (Huge Pages, filesystem blocks, swap behaviours).
Disaster Recovery (DR)	Standardized. Automated failovers and Multi-AZ replication natively managed by AWS.	Highly complex. Requires Patroni triggers, etcd elections, and manual DNS adjustments.
Vendor Portability	Low. Standardizes on AWS-specific APIs, consoles, and network templates.	High. The entire stack (Percona, Patroni, etcd, Redis) can run identically on-premises or on other clouds.

Deep State of Mind (DSOM) For My AI Protocol	Harisfazillah Jamel (LinuxMalaysia)	2026-07-26 Standard: UK English	DBP- standard Bahasa Melayu Malaysia (Piawai)	GNU General Public License v3.0
--	---	--	--	--

Load Testing Assumptions & Sizing Guide

#aws

#cloud

#architecture

#costing

#performance

#valkey

#rds

#postgresql

#waf

#asg

Load Testing Assumptions & Sizing Guide

This document establishes the official workload assumptions, concurrent user definitions, architectural profiles, and target performance objectives governing the load testing, scalability audits, and capacity planning for our secure **AWS 3-Tier Web and AI Infrastructure**.

These assumptions align directly with the performance testing findings and multi-VU scale-up roadmaps, detailing the dual impact of traffic loads on both system performance and financial costing in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)**.

1. Core Load Testing Objectives & SLA Targets

The primary objective of our load testing is to validate that the multi-tier application stack can scale seamlessly up to 10,000 concurrent Virtual Users (VUs) while adhering to strict Service Level Agreements (SLAs) for security, availability, response times, and cost efficiency.

Service Level Agreement (SLA) Matrix

Metric Category	Target SLA / Threshold	Critical Alert Limit	Diagnostic Source
API/Web Latency (P95)	< 250 ms (Static & Cached)	> 500 ms	Application Load Balancer (ALB)
Transaction Latency (P99)	< 800 ms (DB Dynamic Writes)	> 1,500 ms	RDS Performance Insights
HTTP Error Rate	< 0.01% (Zero Gateway Failures)	> 0.1%	ALB Access Logs & CloudWatch
ASG Compute CPU	< 45% average under peak load	> 75% peak	EC2 CloudWatch Metrics
Valkey Cache Hit Rate	> 99.0% for session lookups	< 95.0%	ElastiCache Redis/Valkey Engine
Database Load (AAS)	< Active vCPU count of DB instance	> DB vCPU limit	RDS Performance Insights (AAS)

2. Multi-VU Sizing Tier Definitions

To structure our scale-up roadmap, we define five distinct workload profiles simulating realistic growth phases. All financial estimations are calculated using a stable reference exchange rate of **1 USD = 4.50 MYR**.

MULTI-VU PERFORMANCE ROADMAP		
Target Load	Deployment Model	Key Scaling Bottleneck Target
100 VU	Baseline Dev/Staging	Idle capacity / Base costs
500 VU	Cost-Optimized	Single points of failure (SPOF)
2,500 VU	High-Performance Prod	Database full scans & disk I/O
5,000 VU	Heavy Concurrency	Valkey session write durability
10,000 VU	Extreme Concurrency	Socket pools & WAF L7 latency

A. Sizing & Resource Allocation Matrix

Target Load (VU)	Deployment Phase / Model	Key Compute Sizing (ASG)	Database Engine Spec	Valkey Cache Sizing	Monitoring & Alerts
100 VU	Baseline Dev / Staging	2x t4g.micro	db.t4g.micro (Multi-AZ)	1x cache.t4g.micro	CloudWatch Metrics, SPOF Alerts
500 VU	Cost-Optimized Staged	2x t4g.medium	db.m6g.large (Multi-AZ)	1x cache.t4g.micro	CloudWatch Metrics, SPOF Alerts
2,500 VU	High-Performance Prod	Avg. 4x t4g.xlarge	db.m6g.xlarge (Multi-AZ)	2x cache.t4g.medium (HA)	CloudWatch Metrics, SPOF Alerts, Latency Monitoring
5,000 VU	Heavy Concurrency Prod	Min. 4x t4g.xlarge	db.m7g.2xlarge (Multi-AZ)	2x cache.t4g.medium (HA)	CloudWatch Metrics, SPOF Alerts, Latency Monitoring, Session Write Durability
10,000 VU	Extreme Concurrency Prod	Min. 8x t4g.xlarge	db.m7g.4xlarge (Multi-AZ)	4x cache.m7g.large (Cluster)	CloudWatch Metrics, SPOF Alerts, Latency Monitoring, Session Write Durability, Socket Pool Monitoring

3. Core Architectural Workload Assumptions

To execute high-fidelity simulated tests, the following application-level and user behavior assumptions are defined:

- . **Read-to-Write Ratio:** The standard web traffic mixture is assumed to be **80% Reads and 20% Writes**.
 - Read operations represent session retrievals, reporting queries, database lookups, and asset downloads.
 - Write operations represent active user logins, database updates (e.g., transactional updates), and reporting uploads.
- . **Session Persistence:** In-memory sessions are handled via ElastiCache for Valkey, with fallback persistence to the relational database layer.
- . **Layer-7 Security Filter Overhead:** All public ingress must transit through AWS WAFv2 with OWASP Top 10 Rules, rate-limiting, and deep-packet payload inspections.
- . **Network Boundaries:** Direct Internet access is prohibited for application servers. Outbound egress is routed exclusively through NAT Gateways.
- . **Auto Scaling Responsiveness:** ASG cooldown periods are set to 180 seconds, triggering scale-up events when average compute CPU exceeds 50% for two consecutive 1-minute evaluation periods.

4. Key Performance Bottlenecks & Critical Remediation

Based on empirical performance testing audits under heavy loads, several system-level bottlenecks were identified and resolved. These findings serve as design rules for future system scale-ups.

A. Database Query Indexing (2,500 VU Bottleneck)

- **Problem:** Under 2,500 concurrent users, RDS MariaDB/PostgreSQL active database load surged to a peak of 8.0 Average Active Sessions (AAS), blocking transactions on database table handlings (wait/io/table/sql/handler).
- **Root Cause:** A complete absence of composite indexes on aggregate and transactional tables, forcing expensive disk filesort operations and full table scans.
- **Remediation:**
 - Create composite index on MariaDB summary tables:
idx_summary_agg (status, tenant_id, total) .
 - Create composite index on reconciliation lookup tables:
idx_recons_lookup (reference_id, is_reconciled) .
 - Create composite index on transaction tables: idx_parking_lookup (refno, user_id) .

B. Storage Disk Write Sync & IOPS (2,500 VU Bottleneck)

- **Problem:** High-concurrency update transactions caused PostgreSQL backend processes to stall, creating an I/O:walSync disk write wait bottleneck.
- **Root Cause:** The default GP3 general-purpose storage volume ran out of write capabilities under sustained, concurrent transactional traffic.
- **Remediation:** Upgrade PostgreSQL storage from GP3 to **Provisioned IOPS (io2)** with at least 5,000 Custom IOPS and 250 MB/s throughput, reducing sync latency from 150ms to under 8ms.

C. Web Application Firewall regional ACL Overhead (10,000 VU Bottleneck)

- **Problem:** Layer-7 WAF inspection rules introduced an additional latency overhead of up to 20ms under heavy loads.
- **Root Cause:** Deep packet inspection evaluating hundreds of complex regex patterns for every request.
- **Remediation:** Nest and optimize Web ACL rules, utilize precise regex match patterns, and deploy explicit scope-down statements

using FieldToMatch settings (e.g., limiting expensive body/payload inspection rules specifically to JSON API requests via UriPath and Method constraints) to prevent inspection overhead on safe static assets while retaining robust baseline protections.

D. Socket Exhaustion & Process Sizing (10,000 VU Bottleneck)

- **Problem:** Standard Nginx connections and PHP-FPM worker pools were exhausted, generating minor Gateway Timeout (504) errors.
- **Remediation:**
 - Increase worker_connections to 20,480 in Nginx configurations.
 - Tune the keep-alive timeout to 15 seconds to recycle socket connections rapidly.
 - Configure the PHP-FPM process manager to use static allocation (pm = static) for dedicated production environments, scaling pm.max_children to 120+ based on physical vCPU capacity and measured per-process memory consumption (assuming ~45MB per PHP-FPM process) to prevent physical RAM exhaustion.

E. In-Memory Session Durability (5,000 VU Warning)

- **Problem:** Under network partition events or failovers, session data stored in Valkey could be lost due to asynchronous replication.
 - **Remediation:**
 - For critical transactional data, route writes directly to RDS.
 - If using Valkey for durable session management:
 - **AWS ElastiCache managed clusters:** Configure the Durability parameter and verify EffectiveDurability=sync is active.
 - **Self-hosted Valkey instances:** Start the database engine explicitly with --durability sync to guarantee synchronous replication logging to replica nodes before acknowledging client writes.
-

5. Load Test Sizing Financial Alignments (FinOps)

Operating high-concurrency systems requires robust resources, which increases the monthly AWS spend. To achieve maximum cost efficiency under scale, we define the following day-2 FinOps levers:

- . **Reserved Instances (RIs):** Commit to a 1-year or 3-year standard contract for the primary Multi-AZ db.m7g database engines to unlock up to **30% - 35% savings** compared to on-demand rates.
- . **Compute Savings Plans:** Secure a 1-year or 3-year Compute Savings Plan to achieve **25% - 43% discounts** on the dynamic Auto Scaling Group EC2 application servers (t4g and c7g families).
- . **S3 Intelligent-Tiering and Lifecycle Policies:** Implement automated rules moving historical logs and backup archives to S3 Glacier Flexible Archive, yielding up to **80% - 84% savings** on long-term storage.
- . **VPC S3 Gateway Endpoints:** Configure free S3 Gateway Endpoints in private subnets, allowing application servers to bypass the NAT Gateway and eliminate NAT Gateway data-processing fees (\$0.045/GB in ap-southeast-5) for log and file transfers.

VPC Module

#aws

#cloud

#architecture

#vpc

#alb

#asg

#rds

#dns

#disaster-recovery

VPC Module

The VPC Module deploys the foundational multi-AZ networking infrastructure for the 3-tier layout. It partitions the network into public subnets, private application subnets, and isolated private database subnets, ensuring high availability and physical logical boundaries.

Architectural Details

- **VPC (`aws_vpc`):** Created with a custom CIDR (default `10.0.0.0/16`) and has `enable_dns_support` and `enable_dns_hostnames` set to `true` to facilitate intra-VPC DNS resolution.
- **Internet Gateway (`aws_internet_gateway`):** Attached to the VPC to route external traffic to and from the public subnets.
- **Subnets (`aws_subnet`):**
 - **Public Subnets:** Deployed across multiple AZs. Maps public IP addresses automatically on launch. Hosts ALBs and NAT Gateways.
 - **Private App Subnets:** Deployed across multiple AZs. Holds ASG instances. Routes outbound internet traffic via NAT Gateways.
 - **Private DB Subnets:** Dedicated subnets for RDS DB placement. Completely lacks routes to the NAT Gateways or Internet Gateways to minimize access.
- **NAT Gateways (`aws_nat_gateway`):** Provisioned in public subnets with a dedicated Elastic IP (`aws_eip`) in each AZ to provide high-availability outbound routes for instances in private subnets.

- **Route Tables and Associations:** Configured to manage separate traffic routes for public, private application, and database subnets.
-

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module's inline documentation at `terraform/modules/vpc/README.md`.

Security Groups Module

#aws

#cloud

#architecture

#alb

#asg

#rds

#disaster-recovery

#postgresql

Security Groups Module

This module defines stateful firewall rules (Security Groups) that isolate each layer of the 3-Tier topology, strictly adhering to the **Zero-Trust Network Principle**.

Firewall Rulesets

Presentation Layer: ALB Security Group (alb_sg)

- **Ingress:** Accepts inbound public traffic on:
 - Port 80 (HTTP) from any IPv4 or IPv6 address (0.0.0.0/0 , ::/0).
 - Port 443 (HTTPS) from any IPv4 or IPv6 address (0.0.0.0/0 , ::/0).
- **Egress:** Allows outbound connections to any destination (0.0.0.0/0) to route filtered requests.

Application Layer: ASG/EC2 Security Group (asg_sg)

- **Ingress:** Accepts traffic **strictly** from the ALB Security Group on the application HTTP port (80). No direct internet access is permitted.
- **Egress:** Allows outbound connections to any destination (0.0.0.0/0) to allow instances to perform updates or retrieve application packages.

Database Layer: RDS Security Group (db_sg)

- **Ingress:** Accepts traffic **strictly** from the ASG/EC2 Security Group on the database port (5432 for PostgreSQL / 3306 for MySQL). No other traffic is allowed.
 - **Egress:** Blocked to restrict exfiltration.
-

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module's inline documentation at [terraform/modules/security_groups/README.md](#) .

Web Application Firewall (WAF) Module

#aws

#cloud

#architecture

#alb

#waf

#disaster-recovery

Web Application Firewall (WAF) Module

The WAF Module provisions an AWS WAFv2 Web ACL (Web Application Firewall) configured with modern security rules to defend against automated layer-7 attacks, SQL Injection (SQLi) attempts, and brute-force traffic floods.

Active Protection Rule Sets

- . **Common Rule Set (OWASP Top 10):**
 - Implements AWS Managed Common Rules (`AWSManagedRulesCommonRuleSet`) which protect against standard vulnerabilities, unauthorized administrative path access, and exploitation techniques.
- . **SQLi Defense (SQL Injection Protection):**
 - Implements AWS Managed SQLi Rules (`AWSManagedRulesSQLiRuleSet`) to identify and block malicious SQL injections in query parameters, headers, or request bodies.
- . **IP Rate Limiting:**
 - Adds a custom rate-based rule to restrict the maximum number of requests a single client IP address can send within a sliding 5-minute window (default: 2000 requests).

Edge Association

- **ALB Association (`aws_wafv2_web_acl_association`):** Directly attaches the regional Web ACL to the Application Load Balancer ARN to filter traffic at the network edge.
-

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module's inline documentation at `terraform/modules/waf/README.md` .

Application Load Balancer (ALB) Module

#aws

#cloud

#architecture

#alb

Application Load Balancer (ALB) Module

This module deploys a highly available, public-facing Application Load Balancer (ALB) that acts as the entry point for all incoming HTTP web traffic. It handles request distribution and active health checking.

Technical Components

- **Load Balancer (`aws_lb`):** An external, internet-facing application load balancer mapped across multiple public subnets to provide high availability.
- **Target Group (`aws_lb_target_group`):** A target group with active health checks to monitor the health of the underlying EC2 instances managed by the Auto Scaling Group.
 - **Path:** /
 - **Protocol:** HTTP
 - **Healthy/Unhealthy Threshold:** 3
 - **Interval:** 30 seconds
 - **Timeout:** 5 seconds
- **Listener (`aws_lb_listener`):** Configured on port 80 (HTTP) to forward traffic to the target group.

Inputs and Outputs

For a detailed list of all input parameters and output values,
please refer to the module's inline documentation at
`terraform/modules/alb/README.md` .

Auto Scaling Group (ASG) Module

#aws

#cloud

#architecture

#asg

#bastion

Auto Scaling Group (ASG) Module

This module deploys an Auto Scaling Group (ASG) of EC2 instances spanning multiple Availability Zones. It automates high-availability application scaling and supports dynamic resource scaling based on real-time CPU consumption.

Technical Details

- **Multi-AZ Resilience:** ASG instances are launched strictly inside private subnets, ensuring they are shielded from direct internet exposure.
- **Dynamic Graviton / ARM64 Auto-detection:**
 - The module looks up the official Amazon Linux 2023 AMI dynamically.
 - If the `instance_type` belongs to the AWS Graviton family (e.g., starts with `t4g`, `m6g`, `c6g` etc.), the module automatically switches to fetch the **ARM64** Amazon Linux 2023 AMI. Otherwise, it defaults to the standard **x86_64** AMI.
- **SSM Integration:** Installs and configures IAM instance profiles with `AmazonSSMManagedInstanceCore` policies, allowing administrators to establish secure terminal shell sessions with instances without deploying bastion hosts or opening public SSH ports.
- **Nginx Bootstrapping:** Provisions instances with an automated bootstrapping `user_data` script that updates system packages,

installs Nginx, starts the service, and creates a customized region-specific welcome index page.

- **Dynamic CPU Scaling:**

- **Scale Out:** Triggers when average CPU utilization matches or exceeds **70%** for two consecutive evaluation periods.
- **Scale In:** Triggers when average CPU utilization falls below or matches **30%** for two consecutive evaluation periods.

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module's inline documentation at `terraform/modules/asg/README.md`.

RDS Module

#aws

#cloud

#architecture

#rds

#ssl

#postgresql

RDS Module

The RDS Module deploys a secure, fully-managed RDS Database instance running in Multi-AZ configuration to ensure maximum availability and seamless automatic failover in the event of an availability zone outage.

Technical Details

- **Multi-AZ Availability:** Automatically provisions a primary database instance in one AZ and a synchronous hot standby instance in a secondary AZ.
 - **Deep Private Isolation:** Placed strictly inside dedicated isolated database subnets without any routes to the internet or NAT gateways.
 - **Dynamic Parameter Group Tuning:**
 - Creates a dedicated parameter group with custom database parameters.
 - Automatically adapts the DB parameter group family (e.g., `postgres16`, `postgres15`, or `mysql8.0`) based on the selected database engine and major version.
 - **Storage Encryption:** Enforces EBS storage encryption (`storage_encrypted = true`) using AWS managed KMS keys.
 - **Storage Autoscaling:** Supports storage scale thresholds (from an initial 20 GB up to 100 GB max) to seamlessly accommodate business data growth without manual intervention.
-

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module's inline documentation at `terraform/modules/rds/README.md`.

Standalone EC2 Instance Module

#aws

#cloud

#architecture

#vpc

#alb

#bastion

#disaster-recovery

Standalone EC2 Instance Module

This module provisions secure Standalone EC2 Instances running **Ubuntu 26.04 LTS** (Noble Numbat successor) to support custom application build-ups, staging tools, developer sandboxes, or other application resources that are not ready or suited for Auto Scaling Group deployment.

Technical Details

- **Secure Network Isolation:**
 - Standalone instances are launched strictly inside the **VPC Private Application Subnets**.
 - They do not have public IP addresses and are completely hidden from direct public internet ingress.
 - Outbound traffic (e.g., package updates via `apt-get`, package validation via `debsums`, and security auditing scans via `ASIMP`) is securely routed through the managed NAT Gateways.
- **SSM Integration:**
 - Deploys dedicated IAM Roles and Instance Profiles with `AmazonSSMManagedInstanceCore` policies attached.
 - Enables developers to establish secure terminal shell sessions with standalone instances using AWS Systems Manager (SSM) without setting up bastion hosts or opening public SSH ports.
- **Chained Security Grouping:**

- Deploys a dedicated security group (standalone-ec2-sg) for standalone instances.
 - Restricts inbound traffic on port 80/443 strictly to the Application Load Balancer (ALB) security group, enabling developers to run web-facing test services or APIs behind the secure ALB.
 - **Bootstrapping user_data:**
 - Standardizes initial packages, applies upgrades, and installs Nginx as a placeholder server.
 - Ready to run custom ASIMP (Ansible System Integrity Management Platform) hardening playbooks.
-

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module's inline documentation at `terraform/modules/standalone_ec2/README.md` .

ElastiCache Valkey Module

#aws

#cloud

#architecture

#vpc

#asg

#elasticache

#valkey

#disaster-recovery

ElastiCache Valkey Module

The ElastiCache Valkey Module deploys a secure, fully-managed **Amazon ElastiCache for Valkey** cache cluster inside the private database subnets of your VPC.

Valkey is a high-performance, fully open-source key-value database stewarded by the Linux Foundation. It acts as a drop-in replacement for Redis OSS, offering identical client protocol and command compatibility while delivering **20% lower on-demand pricing** for self-designed (node-based) clusters.

Technical Details

- **Deep Private Isolation:** Placed strictly inside dedicated database subnets without internet egress.
- **Valkey Engine Selection:** Uses valkey engine with version 7.2 as the default modern cache cluster, avoiding Redis OSS license premium.
- **Transit and At-Rest Encryption:**
 - Enforces TLS in-transit encryption (`transit_encryption_enabled = true`).
 - Enforces AES-256 at-rest encryption (`at_rest_encryption_enabled = true`).
- **Dynamic Ingress Firewalling:**
 - Inbound traffic on port 6379 is allowed exclusively from the application Auto Scaling Group (ASG) Security Group.

- Optionally allows inbound port 6379 connections from standalone staging/developer instances to support 1:1 environment parity for integration testing.
-

Inputs and Outputs

For a detailed list of all input parameters and output values, please refer to the module's inline documentation at `terraform/modules/elasticache/README.md`.

SSH Jumphost Module

#aws

#cloud

#architecture

#vpc

#asg

#jumphost

#bastion

#disaster-recovery

AWS SSH Jumphost (Bastion) Module

This module provisions a highly secure, cost-optimized SSH Jumphost (Bastion) running inside the VPC's public subnet. It is whitelisted exclusively for incoming connections from your developer office network (e.g., Cyberjaya, Selangor, Malaysia) on port 22.

The module also injects ingress SSH security group rules directly into the Auto Scaling Group (ASG) and Standalone EC2 instances' security groups, allowing developers to establish secure SSH proxying/tunnels into the private environment.

Core Features

- **Static Public Routing:** Allocates and associates a dedicated AWS Elastic IP (EIP) with the Jumphost, giving developers a stable public IP to configure in their SSH client configurations.
- **Strict Ingress Whitelisting:** Keying security groups to restrict SSH (TCP port 22) access to the specified office CIDR only, preventing brute-force and port-scanning attempts from other parts of the internet.
- **SSH Chaining to Private Resources:** Attaches managed security group ingress rules to the private subnet ASG and standalone compute groups, enabling passwordless SSH key proxying.
- **Operating System Versatility:** Allows switching between canonical **Ubuntu 24.04/26.04 LTS** (fully compatible with **ASIMP** OS-level hardening) and cloud-optimized **Amazon Linux 2023** via input parameters.

- **SSM-Managed Backup:** Connects to AWS Systems Manager (SSM) by default via an IAM instance profile, providing a safe, passwordless shell backup route for system administrators.
-

Inputs

Name	Description	Type	Default
environment	Environment name for tagging (e.g., production, dev)	string	n/a
vpc_id	The ID of the VPC	string	n/a
public_subnet_ids	List of public subnet IDs in the VPC	list(string)	n/a
instance_type	EC2 compute instance size (Graviton-compatible default)	string	"t4g.micro"
allowed_ssh_cidr	IP CIDR allowed to connect to the Jumphost via SSH	string	"103.188.0.0/16"
jumphost_os	Operating System pattern ('ubuntu' or 'amazon-linux-2023')	string	"ubuntu"

Name	Description	Type	Default
ami_id	Specific AMI ID override (if left empty, AMI is auto-fetched)	string	""
ubuntu_ami_filter_name	The search pattern name for the Ubuntu AMI	string	"ubuntu/images/hvm-ssd-gp3/ubuntu-noble-24.04-*server-*"
asg_sg_id	Security Group ID associated with the Auto Scaling Group	string	n/a
standalone_sg_id	Security Group ID associated with Standalone EC2 instances	string	""

Outputs

Name	Description
jumphost_public_ip	The static Elastic IP address assigned to the Jumphost
jumphost_private_ip	The private IP address of the Jumphost within the VPC
security_group_id	The ID of the security group assigned to the Jumphost
instance_id	The ID of the Jumphost EC2 instance

Automation Scripts

#aws

#cloud

#architecture

#vpc

#waf

Automation Scripts

The project includes CLI helper scripts and bootstrapping scripts under the `scripts/` directory to automate common tasks, local testing, and instance provisioning.

1. Local Deployment Script (`scripts/deploy.sh`)

This script automates the full lifecycle of a local OpenTofu deployment. It performs validation and planning checks before prompting the user for confirmation to apply.

Steps Executed:

- . **OpenTofu Installation Check:** Verifies that the `tofu` CLI is installed and available in the execution path.
- . **Directory Context:** Changes the working directory to `terraform/`.
- . **Environment Configuration Check:** Checks if `terraform.tfvars` exists. If not, it copies it from `terraform.tfvars.example` and prompts the user to review.
- . **Initialization (`tofu init`):** Downloads required providers and module configurations.
- . **Format Verification (`tofu fmt`):** Formats all configuration files recursively to match canonical HCL syntax.
- . **Validation (`tofu validate`):** Verifies syntax correctness and internal variable consistency.

- . **Execution Planning (tofu plan):** Generates a secure plan file `tfplan` .
 - . **User Confirmation Prompt:** Asks `Do you want to apply this deployment? (y/n)` . If `y` , it runs `tofu apply tfplan` .
-

2. Infrastructure Teardown Script (`scripts/destroy.sh`)

This script provides a clean and automated way to destroy all resources managed by OpenTofu, preventing accidental dangling resources or billing charges.

Steps Executed:

- . **OpenTofu Installation Check:** Verifies that the `tofu` CLI is installed and available.
 - . **Initialization:** Confirms that OpenTofu has been initialized and that `.terraform` config exists.
 - . **Resource Destruction (tofu destroy):** Prompts the user to confirm the termination of all resources. Once confirmed, it cleanly removes all VPC components, load balancers, database clusters, auto-scaling groups, and WAF rules.
-

3. Instance Bootstrapping Script (`scripts/user_data.sh`)

This is a standalone bootstrapping script for EC2 instances. It can be used for custom instance images or standalone EC2 deployment tests.

Features:

- Updates the operating system packages.
- Installs and configures an Apache HTTP web server.

- Fetches Instance Metadata Service v2 (IMDSv2) security tokens dynamically.
- Retrieves instance metadata, including the Instance ID and host Availability Zone, displaying this dynamic diagnostic information on a beautifully styled HTML index page.

CI/CD Pipeline Documentation

#aws

#cloud

#architecture

#rds

CI/CD Pipeline Documentation

This project includes a fully automated CI/CD pipeline configured via GitHub Actions in `.github/workflows/opentofu.yml`. It ensures that every code submission is vetted for quality, syntax correctness, and security boundaries.

Pipeline Workflow Trigger Conditions

The pipeline is triggered automatically on:

- **Pull Requests** targeting the main branch.
- **Direct Pushes** or Merges to the main branch.

Job Definitions

The workflow consists of four major jobs designed with security and safety gates:

1. OpenTofu Format and Validate (`opentofu-lint-and-validate`)

- **Environment:** `ubuntu-latest`
- **Steps:**
 - **Checkout Code:** Checks out the repository files.

- **Setup OpenTofu:** Installs OpenTofu version 1.8.2 on the runner using `opentofu/setup-opentofu@v1`.
- **Format Check:** Verifies formatting recursively via `tofu fmt -check -recursive`.
- **Initialize:** Runs `tofu init -backend=false` to configure modules locally.
- **Validate:** Evaluates configurations for semantic validity via `tofu validate`.

2. Check Secrets Availability (`check-secrets`)

- **Environment:** `ubuntu-latest`
- **Purpose:** Checks whether the sensitive OIDC IAM Role `AWS_ROLE_TO_ASSUME` secret is populated without causing parser failures.
- **Steps:**
 - **Evaluate Secret Presence:** Evaluates `secrets.AWS_ROLE_TO_ASSUME` on the runner at step-level, dynamically exporting a boolean output variable `has-aws-role` (set to `true` if populated, or `false` if empty, such as in fork pull requests). This avoids static workflow validation errors.

3. OpenTofu Planning (`opentofu-plan`)

- **Trigger:** Runs only on **Pull Request** events.
- **Dependency:** Requires both `opentofu-lint-and-validate` and `check-secrets` jobs to pass, with `check-secrets.outputs.has-aws-role == 'true'`.
- **Steps:**
 - **Configure AWS OIDC Credentials:** Authenticates securely with AWS using OpenID Connect (OIDC) through `aws-actions/configure-aws-credentials@v4` with `secrets.AWS_ROLE_TO_ASSUME` and `secrets.AWS_REGION`.
 - **OpenTofu Init & Plan:** Performs complete backend initialisation and generates an execution plan showing what changes will be applied.

4. OpenTofu Deployment (opentofu-apply)

- **Trigger:** Runs only on **Pushes** or Merges to the main branch.
 - **Dependency:** Requires both opentofu-lint-and-validate and check-secrets jobs to pass, with check-secrets.outputs.has-aws-role == 'true' .
 - **Steps:**
 - **Configure AWS OIDC Credentials:** Authenticates securely with AWS using OIDC.
 - **OpenTofu Init & Apply:** Runs tofu apply -auto-approve to deploy the target infrastructure.
-

Secret Management Best Practices

The pipeline utilizes modern security standards:

- **No Hardcoded Credentials:** Leverages short-lived, dynamically-exchanged AWS security tokens via OIDC instead of storing long-lived AWS_ACCESS_KEY_ID and AWS_SECRET_ACCESS_KEY secrets.
- **Configurable Region:** Adapts automatically based on secrets.AWS_REGION or defaults gracefully to us-east-1 (or local regional default config).

GitLab CI/CD & AWS EFS Code Deployment

#aws

#cloud

#architecture

#vpc

#alb

#asg

#rds

#bastion

#dns

#disaster-recovery

#gitlab

#efs


GitLab CI/CD & AWS EFS Code Deployment Architecture

This guide provides a comprehensive design and deployment blueprint for automating application delivery. It outlines a secure, scalable workflow where developers use **GitLab** for source control, triggering automated **GitLab CI/CD pipelines** that deploy code directly onto an **Amazon Elastic File System (EFS)**.

By mounting this shared network filesystem (NFSv4) concurrently on both **Auto Scaling Groups (ASGs)** and **Standalone EC2 instances**, application updates become instantly active across all nodes without rebuilding AMIs or performing slow rolling redeployments.

1. High-Level Architecture Overview

The system bridges an external GitLab repository with a secure, private AWS 3-tier VPC. Since EFS endpoints reside deep within private subnets, we implement a highly secure staging-and-sync mechanism.



Core Flow Steps:

- . **Commit & Push:** A developer pushes code or creates a git tag in GitLab.
 - . **GitLab CI/CD Trigger:** The GitLab runner starts, authenticates dynamically to AWS using **OIDC IAM Role Assumption**, runs tests, compiles build artifacts (e.g., node modules, frontend assets), and tars the build.
 - . **Artifact Staging:** The runner uploads the tarball artifact to a secure, private **AWS S3 Deployment Bucket**.
 - . **Automated EFS Syncing:** The runner triggers an **AWS Systems Manager (SSM) Run Command** targeting a dedicated internal Utility/Sync instance (or the Bastion host) that has the EFS volume mounted.
 - . **Atomic Unpacking:** The Utility instance downloads the tarball from S3, unpacks it into the appropriate EFS sub-directory, and atomically swaps symlinks to complete the deploy in milliseconds.
 - . **Instant Access:** Because the EFS volume is mounted across all active ASG instances and standalone staging nodes, the code is immediately available to the running applications without any instance terminations or server reboots.
-

2. Shared EFS Directory & Mounting Strategy

To maintain complete Isolation of Concerns and prevent staging/production resource conflicts, EFS must be structured logically.

2.1 EFS Directory Layout

We partition EFS into two top-level environments (production and staging) with application sub-directories:

```

/mnt/efs/
├─ staging/                                # Used exclusively by Standalone EC2
|   ├─ frontend/                          # Frontend HTML/JS assets
|   └─ backend/                            # Backend API server code
|
└─ production/                            # Used by Auto Scaling Groups (ASGs)
    ├─ frontend/                          # Production Frontend assets
    └─ backend/                            # Production Backend API code

```

2.2 Secure Mounting on EC2 Instances (Ubuntu 26.04 LTS & AL2023)

To ensure reliable, encrypted-in-transit connections to EFS, we use `amazon-efs-utils` over NFSv4 with TLS (Port 2049).

A. Interactive or Bootstrap Script Mounting:

```

# Install EFS Utilities
# For Amazon Linux 2023 (ASG Nodes):
sudo dnf install -y amazon-efs-utils

# For Ubuntu 26.04 LTS (Standalone Staging Nodes):
sudo apt-get update && sudo apt-get install -y binutils cpp
git clone https://github.com/aws/efs-utils.git /tmp/efs-utils
cd /tmp/efs-utils && ./build-deb.sh
sudo apt-get install -y ./build/amazon-efs-utils*.deb

# Create Mount Directories
sudo mkdir -p /mnt/efs

# Mount with TLS enabled (Encapsulated over Port 2049)
# Replace fs-xxxxxx with your actual EFS File System ID
sudo mount -t efs -o tls,iam fs-xxxxxx:/ /mnt/efs

```

B. Persistent Mount Configuration (/etc/fstab):

To guarantee EFS is automatically remounted if an ASG instance launches or a standalone node is rebooted, append the following entry to `/etc/fstab`:

```
fs-xxxxxx:/ /mnt/efs efs _netdev,tls,iam,defaults 0 0
```

(The `_netdev` option prevents the system from attempting to mount EFS before network interfaces are fully initialized).

3. Nginx Configuration & Path Design

To serve code residing on network-attached storage efficiently, Nginx must be configured to minimize file system traversal overhead and handle the static/dynamic paths correctly.

3.1 Nginx Path Allocation

Instance Role	Local Mount Path	App Root Directory (Served by Nginx)
Production Frontend ASG	/mnt/efs	/mnt/efs/production/frontend/
Production Backend ASG	/mnt/efs	/mnt/efs/production/backend/ (App running via PM2/systemd)
Standalone Staging EC2	/mnt/efs	/mnt/efs/staging/frontend/ & /mnt/efs/staging/backend/

3.2 Performance Tuning for Network Filesystems (EFS)

Because EFS is a distributed network filesystem, standard POSIX metadata lookups (like `stat()`, which checks if a file exists or has changed) introduce millisecond network round-trips. Under high web traffic, this can severely throttle Nginx performance.

Mitigation: Enforce Nginx-level file descriptor caching. Add the following directives inside your global `/etc/nginx/nginx.conf` (within the `http` block):

```
# Cache open file descriptors, their sizes, and modification times
open_file_cache          max=2000 inactive=20s;
open_file_cache_valid    30s;
open_file_cache_min_uses 2;
open_file_cache_errors   on;
```

3.3 Production Frontend Nginx Configuration (ASG Node)

Apply this configuration to the instances running in your **Frontend Web ASG** (`/etc/nginx/sites-available/default`):

```

server {
    listen 80;
    server_name app.linuxmalaysia.com;

    # Root pointing to the production frontend EFS path
    root /mnt/efs/production/frontend;
    index index.html index.htm;

    # Performance optimizations for EFS
    sendfile on;
    tcp_nopush on;
    tcp_nodelay on;

    # Enable Gzip Compression to save bandwidth
    gzip on;
    gzip_types text/plain text/css application/json
application/javascript text/xml;

    # Standard SPA Routing or Static File Server
    location / {
        try_files $uri $uri/ /index.html;
    }

    # Route backend API requests to the Backend ALB
    location /api/ {
        # Core VPC DNS Resolver with short TTL to handle Backend ALB
scaling
        resolver 169.254.169.254 valid=10s;
        set $backend_upstream "http://internal-production-backend-alb-
123456.ap-southeast-5.elb.amazonaws.com";

        proxy_pass $backend_upstream;
        proxy_set_header Host $host;
        proxy_set_header X-Real-IP $remote_addr;
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
        proxy_set_header X-Forwarded-Proto $scheme;
    }
}

```

```
# Static Assets Cache Control
location ~* \.
(?:css|js|jpg|jpeg|gif|png|ico|svg|woff|woff2|ttf|otf)$ {
    expires 7d;
    add_header Cache-Control "public, no-transform";
    access_log off;
}

# Custom Error Pages
error_page 500 502 503 504 /50x.html;
location = /50x.html {
    root /usr/share/nginx/html;
}
}
```

3.4 Standalone Staging Nginx Configuration (EC2 Staging Node)

Staging nodes host both frontend and backend APIs on a single standalone instance for cost efficiency. Apply this layout (/etc/nginx/sites-available/default):

```

server {
    listen 80;
    server_name staging.linuxmalaysia.com;

    # Staging Frontend Root Directory
    root /mnt/efs/staging/frontend;
    index index.html;

    sendfile on;

    # Staging Frontend SPA Routing
    location / {
        try_files $uri $uri/ /index.html;
    }

    # Staging API reverse proxy pointing to a local Node.js/Python
service
    location /api/ {
        # Staging backend app usually runs locally on port 5000
        proxy_pass http://127.0.0.1:5000;
        proxy_set_header Host $host;
        proxy_set_header X-Real-IP $remote_addr;
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
        proxy_set_header X-Forwarded-Proto $scheme;

        # Disable buffering for real-time staging logs/WebSocket
support
        proxy_buffering off;
    }
}

```

4. Automated GitLab CI/CD Pipeline Configuration

To trigger automatic deployments based on code versioning, we configure the pipeline to recognize:

- . **Pushes to main branch:** Deploys instantly to the **Staging / Standalone** environment.
- . **Git Release Tags (v*.*.*):** Deploys to the **Production / ASG** environment.

Production-Ready .gitlab-ci.yml Manifest

Place this file in the root of your code repository:

```

stages:
  - lint
  - build
  - deploy

variables:
  AWS_DEFAULT_REGION: "ap-southeast-5" # Malaysia Region
  S3_STAGING_BUCKET: "s3://company-deployment-staging-bucket-ap-southeast-5"
  S3_PROD_BUCKET: "s3://company-deployment-prod-bucket-ap-southeast-5"
  SYNC_INSTANCE_STAGING_ID: "i-0123456789staging" # Standalone EC2 Staging Node ID
  SYNC_INSTANCE_PROD_ID: "i-0123456789utility" # Prod EFS Admin Utility Instance ID

# AWS Authentication via OIDC (No persistent IAM credentials stored in GitLab!)
.aws-auth:
  id_tokens:
    MY_OIDC_TOKEN:
      aud: https://gitlab.com
  before_script:
    - mkdir -p ~/.aws
    - echo "AssumeRoleWithWebIdentity"
    - >
      export $(printf "AWS_ACCESS_KEY_ID=%s AWS_SECRET_ACCESS_KEY=%s AWS_SESSION_TOKEN=%s"
        $(aws sts assume-role-with-web-identity
          --role-arn "arn:aws:iam::112233445566:role/GitLabCI-AWS-OIDC-Role"
          --role-session-name "GitLabPipeline-${CI_PIPELINE_ID}"
          --web-identity-token "${MY_OIDC_TOKEN}"
          --query "Credentials.[AccessKeyId,SecretAccessKey,SessionToken]"
          --output text))

# ===== LINT STAGE =====
code-lint:
  stage: lint

```

```

image: node:20-alpine
script:
  - npm ci
  - npm run lint
rules:
  - if: $CI_COMMIT_BRANCH == "main"
  - if: $CI_COMMIT_TAG =~ /^v\d+\.\d+\.\d+$/

# ===== BUILD STAGE =====
build-assets:
  stage: build
  image: node:20-alpine
  artifacts:
    paths:
      - dist/
      - package.json
      - package-lock.json
    expire_in: 1 week
  script:
    - npm ci
    - npm run build
  rules:
    - if: $CI_COMMIT_BRANCH == "main"
    - if: $CI_COMMIT_TAG =~ /^v\d+\.\d+\.\d+$/

# ===== DEPLOY STAGES =====

# Deploy Staging: Automatically triggered on commits/merges to main
branch
deploy-staging:
  extends: .aws-auth
  stage: deploy
  image: amazon/aws-cli:latest
  dependencies:
    - build-assets
  script:
    # 1. Package the build files
    - tar -czf staging-build.tar.gz dist/ package.json package-

```

lock.json

```
# 2. Upload tarball to Staging S3 Bucket
- aws s3 cp staging-build.tar.gz ${S3_STAGING_BUCKET}/staging-
build.tar.gz

# 3. Trigger AWS SSM Run Command on Staging Instance to pull and
unpack
- >
  aws ssm send-command
  --instance-ids "${SYNC_INSTANCE_STAGING_ID}"
  --document-name "AWS-RunShellScript"
  --comment "Deploy staging files from GitLab CI"
  --parameters 'commands=[
    "aws s3 cp '"${S3_STAGING_BUCKET}"'/staging-build.tar.gz
/tmp/staging-build.tar.gz",
    "mkdir -p /mnt/efs/staging/frontend_new",
    "tar -xzf /tmp/staging-build.tar.gz -C
/mnt/efs/staging/frontend_new",
    "rm -rf /mnt/efs/staging/frontend_old",
    "[ -d /mnt/efs/staging/frontend ] && mv
/mnt/efs/staging/frontend /mnt/efs/staging/frontend_old",
    "mv /mnt/efs/staging/frontend_new /mnt/efs/staging/frontend",
    "rm -rf /mnt/efs/staging/frontend_old /tmp/staging-
build.tar.gz",
    "echo Deployment Successful"
  ]'
rules:
  - if: $CI_COMMIT_BRANCH == "main"

# Deploy Production: Triggered exclusively by creating a Tag (e.g.,
v1.0.0)
deploy-production:
  extends: .aws-auth
  stage: deploy
  image: amazon/aws-cli:latest
  dependencies:
    - build-assets
```



```

script:
  # 1. Package the production files
  - tar -czf prod-build.tar.gz dist/ package.json package-lock.json

  # 2. Upload to Production S3 Bucket
  - aws s3 cp prod-build.tar.gz ${S3_PROD_BUCKET}/prod-
build-${CI_COMMIT_TAG}.tar.gz

  # 3. Trigger SSM Run Command on Utility Admin instance to unpack
onto the shared EFS
  - >
    aws ssm send-command
    --instance-ids "${SYNC_INSTANCE_PROD_ID}"
    --document-name "AWS-RunShellScript"
    --comment "Deploy production files to EFS from GitLab CI"
    --parameters 'commands=[
      "aws s3 cp "'${S3_PROD_BUCKET}'"/prod-build-
"${CI_COMMIT_TAG}"/tar.gz /tmp/prod-build.tar.gz",
      "mkdir -p /mnt/efs/production/frontend_new",
      "tar -xzf /tmp/prod-build.tar.gz -C
/mnt/efs/production/frontend_new",
      "rm -rf /mnt/efs/production/frontend_old",
      "[ -d /mnt/efs/production/frontend ] && mv
/mnt/efs/production/frontend /mnt/efs/production/frontend_old",
      "mv /mnt/efs/production/frontend_new
/mnt/efs/production/frontend",
      "rm -rf /mnt/efs/production/frontend_old /tmp/prod-
build.tar.gz",
      "echo Production Shared EFS Deployment Successful"
    ]'
rules:
  - if: $CI_COMMIT_TAG =~ /^v\d+\.\d+\.\d+$/

```

5. Architectural Evaluation & Better Alternatives

While deploying application code over EFS NFS mounts provides rapid deployment cycles and immediate propagation, it represents a **classic hybrid design pattern with notable operational tradeoffs**.

Below is an engineering analysis comparing EFS deployment to modern cloud-native standards.


5.1 Tradeoffs of the EFS NFS Code-Serving Approach

The Drawbacks:

- . **Network Latency Bottleneck:** Filesystem operations (`stat`, `open`, `read`) over network NFSv4 run significantly slower than on local NVMe/SSD EBS volumes. If an application (e.g., Python/Node.js) imports hundreds of source files at startup, load times are substantially degraded.
- . **Atomic Swap Synchronization Risks:** If an instance executes active code while EFS is being overwritten, it can load partial files, causing runtime exceptions. (Our symlink atomic swap mitigated this, but node-level system file-locks can still lock operations).
- . **Write Concurrency and POSIX File Locking:** Multiple instances running backend APIs attempting to write logs, cache files, or local SQLite data on the same mounted EFS directory will experience lock contention, causing thread blockages.
- . **No Version Immutability (Risk of Corruption):** A single bad build or a malicious script executing on an ASG node can overwrite the shared code folder, instantaneously corrupting/compromising *all* ASG nodes and standalone instances simultaneously.
- . **No Easy Atomic Rollback:** Rollbacks require manual filesystem restores or triggering an old S3 unpack command.

5.2 Proposed Solution 1: Containerization (Docker + Amazon ECS/EKS) – *Highly Recommended*


This is the standard modern cloud-native deployment model.



- **How it works:** GitLab CI/CD builds a standalone Docker Image containing Nginx and the code. It pushes this image to **Amazon Elastic Container Registry (ECR)**. It then triggers an ECS Service rolling update.
- **Why it is superior:**
 - **Complete Isolation & Immutability:** Each container is an exact, unchangeable replica of the tested image.
 - **No Network Storage Bottleneck:** The container runs files off local virtual memory/EBS storage inside the host, resulting in sub-millisecond execution loops.
 - **Zero-Downtime Deployments:** ECS manages blue/green or rolling upgrades automatically, starting healthy new containers before terminating old ones.
 - **Instant rollback:** Rollback is as simple as reverting the active image tag to the previous version in ECR.

5.3 Proposed Solution 2: Stateless S3-Pull Bootstrapping with ASG Instance Refresh

If keeping EC2 instances is mandatory, decouple them from the EFS file system using local EBS drives hydrated from S3.



- **How it works:**

1. GitLab CI/CD compiles the code, zips it, and uploads the immutable artifact to S3 (e.g., `s3://bucket/deploy-v1.0.0.zip`).
2. The pipeline updates an SSM Parameter or updates the **Auto Scaling Group Launch Template** User Data to reference the new version.
3. The pipeline triggers an **ASG Instance Refresh** (a controlled, automated rolling upgrade where AWS replaces old instances with new ones).
4. During bootstrap, the EC2 instance pulls the zip file down from S3, decompresses it locally onto its high-speed EBS volume (`/var/www/html/`), and starts Nginx.

- **Why it is superior:**

- **Performance:** Nginx serves files directly off local high-speed SSDs (EBS) without network NFS lookups.
- **Reliability:** If one instance's file system is corrupted, the other instances in the ASG remain perfectly healthy.
- **Immutability:** Each instance has its own isolated, stable copy of the code. If an auto-scaling event occurs, the newly launched instance pulls the identical versioned zip file.

Costing Estimate



Estimated Costing

We provide two distinct monthly cost estimates for the 3-Tier AWS Architecture in the **AWS Asia Pacific (Malaysia) region (ap-southeast-5)** utilizing On-Demand rates (~730 hours/month). All figures are aligned with our [Software Licensing & Technology Risk Register \(TS/MC Series\)](#) to ensure complete commercial and licensing auditability.

- . **Baseline Cost-Optimized Plan:** Ideal for initial development, testing, staging environments, and low-traffic applications.
- . **High-Performance Developer-Aligned Plan:** Spec'd specifically to fulfill the resource requirements of the **Developer's First Design (Nginx, Backend, RAGFlow, LangFuse)** with production-grade performance.

Both plans have been updated and calibrated against real-world, production-ready AWS billings from similar projects to incorporate critical infrastructure support services (such as **Amazon ElastiCache Valkey, Amazon EFS, AWS Secrets Manager, AWS Backup, Amazon CloudWatch, standard Public IPv4 address charges, Amazon Route 53 custom domain management, and our secure, hardened SSH Jumphost (Bastion)** whitelisted for the Cyberjaya developer office).

Additionally, this guide includes detailed cost models for **AWS-Native Alternatives** to external "extra" integrations (such as Amazon Bedrock instead of OpenAI, Amazon Cognito instead of self-hosted/SaaS Auth, AWS End User Messaging for WhatsApp, and

serverless API Gateway/Lambda webhook routing) so that stakeholders have a complete financial blueprint of a 100% cloud-native architecture.

1. Baseline Cost-Optimized Plan (USD)

This configuration targets baseline usage with smaller resource profiles (t4g.medium compute, db.m6g.large database) to minimize initial expenditures.

Tier / AWS Component	Configuration & Resource Sizing	Hourly/Unit Rate	Est. Monthly Cost (USD)
Networking & NAT	1 NAT Gateway (Outbound API/WhatsApp Traffic)	\$0.045 / hr	\$32.85
	• Base provisioned charge	\$0.045 / GB	\$2.25
	• Estimated ~50 GB Data Processing		
Public IPv4 Addresses	3 Public IPv4 addresses (1 for NAT Gateway, 2 for ALB) • Required public IP address allocation	\$0.005 / hr / IP	\$10.95
Presentation Tier	1 Application Load Balancer (ALB)	\$0.0225 / hr	\$16.43
	• Base LCF/hour charge	\$0.008 / LCU-hr	\$11.68
	• Estimated ~2 LCU average usage		
Security Tier	AWS WAFv2 (Attached to ALB)	\$5.00 / month	\$5.00
	• 1 Regional Web ACL		
	• 3 Rules (OWASP Core, SQLi, Rate Limit)	\$1.00 / rule/mo	\$3.00
	• ~1 Million HTTP/HTTPS Requests	\$0.60 / M-req	\$0.60

Tier / AWS Component	Configuration & Resource Sizing	Hourly/Unit Rate	Est. Monthly Cost (USD)
Compute Tier (ASG)	2x t4g.medium EC2 Instances (ARM Graviton) • 2 vCPU, 4GB RAM each • 2x 30GB gp3 EBS Root Volumes	\$0.0336 / hr / inst \$0.08 / GB-month	\$49.06 \$4.80
Database Tier (RDS)	Multi-AZ db.m6g.large PostgreSQL [TS-06] • 2 vCPU, 8GB RAM (High Availability) • 50GB gp3 Storage	\$0.304 / hr \$0.23 / GB-month	\$221.92 \$11.50
Caching Tier	Amazon ElastiCache Valkey (cache.t4g.micro) [TS-06, TS-05] • 1 Node, 0.5 GB RAM (Session & metadata caching) • <i>Valkey pricing is 20% lower than legacy Redis OSS</i>	\$0.0128 / hr	\$9.34
Storage Tier (S3)	Amazon S3 (Encrypted Uploads & Media) • ~100 GB Standard Storage • ~50,000 PUT/GET API requests	\$0.023 / GB-month Nominal rates	\$2.30 \$0.50

Tier / AWS Component	Configuration & Resource Sizing	Hourly/Unit Rate	Est. Monthly Cost (USD)
Storage Tier (EFS)	Amazon EFS (Elastic File System) • ~10 GB standard shared persistent storage	\$0.30 / GB-month	\$3.00
Secrets Management	AWS Secrets Manager • 2 Secrets (one for RDS DB, one for external API keys)	\$0.40 / secret/mo	\$0.80
Monitoring Tier	Amazon CloudWatch • 3 alarms, custom CPU/Memory dashboard	Nominal rates	\$1.50
Disaster Recovery	AWS Backup • Automated Multi-AZ RDS snapshots & EBS backups (~100 GB)	\$0.05 / GB-month	\$5.00

Tier / AWS Component	Configuration & Resource Sizing	Hourly/Unit Rate	Est. Monthly Cost (USD)
Standalone EC2 Tier	Standalone EC2 Instances (AMI Staging & Baking) • 3x Standalone EC2 Instances (one per ASG group: Frontend, Backend, AI Tier) connected directly to RDS, S3, or EFS to ensure 1:1 environment parity. • Sized as 3x t4g.micro (1 vCPU, 1GB RAM each) • 3x 15GB gp3 EBS Root Volumes (45GB total)	\$0.0084 / hr / inst \$0.08 / GB-month	\$18.39 \$3.60
	Secure SSH Jumphost (Bastion) • 1x t4g.micro EC2 Instance in the Public Subnet (accessible only from Cyberjaya office whitelisted IP) • 15GB gp3 EBS Root Volume • 1 Static Elastic IP allocation	\$0.0084 / hr \$0.08 / GB-mo \$0.005 / hr	\$6.13 \$1.20 \$3.65

Tier / AWS Component	Configuration & Resource Sizing	Hourly/Unit Rate	Est. Monthly Cost (USD)
Data Transfer	Outbound Internet Data Transfer • ~100 GB Outbound (First 100GB Free/mo)	Free Tier	\$0.00
Route 53 Domain	Amazon Route 53 Custom Domain Routing • 1 Public Hosted Zone • Estimated ~2 Million standard queries	\$0.50 / zone / mo \$0.40 / M-req	\$0.50 \$0.80
TOTAL ESTIMATED MONTHLY COST			~\$426.75 USD / month

- **Local Currency Equivalent (MYR): ~RM 1,920 MYR / month**
(calculated at an exchange rate baseline of 1 USD ≈ 4.50 MYR).

2. High-Performance Developer-Aligned Plan (USD)

This configuration scales up computing instances to precisely match the developer’s server specifications (**Server 02, Server 03, and Server 04**) with high performance, dedicated throughput, and larger memory margins.

Tier / AWS Component	Configuration & Resource Sizing	Hourly/Unit Rate	Est. Monthly Cost (USD)
Networking & NAT	1 NAT Gateway (Outbound API/WhatsApp Traffic)	\$0.045 / hr	\$32.85
	• Base provisioned charge	\$0.045 / GB	\$2.25
	• Estimated ~50 GB Data Processing		
Public IPv4 Addresses	3 Public IPv4 addresses (1 for NAT Gateway, 2 for ALB) • Required public IP address allocation	\$0.005 / hr / IP	\$10.95
Presentation Tier	1 Application Load Balancer (ALB)	\$0.0225 / hr	\$16.43
	• Base LCF/hour charge	\$0.008 / LCU-hr	\$11.68
	• Estimated ~2 LCU average usage		
Security Tier	AWS WAFv2 (Attached to ALB)	\$5.00 / month	\$5.00
	• 1 Regional Web ACL		
	• 3 Rules (OWASP Core, SQLi, Rate Limit)	\$1.00 / rule/mo	\$3.00
	• ~1 Million HTTP/HTTPS Requests	\$0.60 / M-req	\$0.60

Tier / AWS Component	Configuration & Resource Sizing	Hourly/Unit Rate	Est. Monthly Cost (USD)
Compute Tier (ASG)	2x t4g.xlarge EC2 Instances (ARM Graviton) • 4 vCPU, 16GB RAM each (Supports Backend, DMS, RAGFlow, LangFuse) • 2x 30GB gp3 EBS Root Volumes	\$0.1344 / hr / inst \$0.08 / GB-month	\$196.22 \$4.80
Database Tier (RDS)	Multi-AZ db.m6g.xlarge PostgreSQL [TS-06] • 4 vCPU, 16GB RAM (Matches Server 04 Data Tier) • 50GB gp3 Multi-AZ Storage	\$0.608 / hr \$0.46 / GB-month	\$443.84 \$23.00
Caching Tier	Amazon ElastiCache Valkey (cache.t4g.medium) [TS-06, TS-05] • 1 Node, 3.09 GB RAM (Production cache & task broker) • <i>Valkey pricing is 20% lower than legacy Redis OSS</i>	\$0.0544 / hr	\$39.71

Tier / AWS Component	Configuration & Resource Sizing	Hourly/Unit Rate	Est. Monthly Cost (USD)
Storage Tier (S3)	Amazon S3 (Encrypted Uploads & Media) • ~100 GB Standard Storage • ~50,000 PUT/GET API requests	\$0.023 / GB-month Nominal rates	\$2.30 \$0.50
Storage Tier (EFS)	Amazon EFS (Elastic File System) • ~50 GB shared network storage for AI model weights / caches	\$0.30 / GB-month	\$15.00
Secrets Management	AWS Secrets Manager • 5 Secrets (RDS, LLM API keys, external integrations, LangFuse, WAF keys)	\$0.40 / secret/mo	\$2.00
Monitoring Tier	Amazon CloudWatch • Logs ingestion (~5 GB), dashboards, custom metric triggers	Nominal rates	\$5.00
Disaster Recovery	AWS Backup • Centralized backup for RDS, EFS, and ASG EBS volumes (~150 GB)	\$0.05 / GB-month	\$7.50

Tier / AWS Component	Configuration & Resource Sizing	Hourly/Unit Rate	Est. Monthly Cost (USD)
Standalone EC2 Tier	Standalone EC2 Instances (AMI Staging & Baking) • 3x Standalone EC2 Instances (one per ASG group: Frontend, Backend, AI Tier) connected directly to RDS, S3, or EFS to ensure 1:1 environment parity. • 1x Frontend: t4g.medium (2 vCPU, 4GB RAM) + 30GB gp3 • 1x Backend: t4g.xlarge (4 vCPU, 16GB RAM) + 30GB gp3 • 1x AI Tier: t4g.xlarge (4 vCPU, 16GB RAM) + 50GB gp3 (110GB gp3 total)	Frontend: \$0.0336 / hr Backend/AI: \$0.1344 / hr Storage: \$0.08 / GB-mo	\$24.53 \$196.22 \$8.80

Tier / AWS Component	Configuration & Resource Sizing	Hourly/Unit Rate	Est. Monthly Cost (USD)
SSH Jumphost Tier	Secure SSH Jumphost (Bastion)		
	• 1x t4g.micro EC2 Instance in the Public Subnet (accessible only from Cyberjaya office whitelisted IP)	\$0.0084 / hr	\$6.13
	• 15GB gp3 EBS Root Volume	\$0.08 / GB-mo	\$1.20
	• 1 Static Elastic IP allocation	\$0.005 / hr	\$3.65
Data Transfer	Outbound Internet Data Transfer • ~100 GB Outbound (First 100GB Free/mo)	Free Tier	\$0.00
Route 53 Domain	Amazon Route 53 Custom Domain Routing	\$0.50 / zone / mo	\$0.50
	• 1 Public Hosted Zone • Estimated ~2 Million standard queries	\$0.40 / M-req	\$0.80
TOTAL ESTIMATED MONTHLY COST			~\$1,064.46 USD / month

- **Local Currency Equivalent (MYR): ~RM 4,790 MYR / month**
(calculated at an exchange rate baseline of 1 USD ≈ 4.50 MYR).

3. Real-World Cost Calibration & Analysis

A comparison with real-world billings from a highly similar production deployment (with a stable run-rate of **\$659.10** for a representative billing cycle, e.g., May 2026) validates our cost modeling and exposes key areas where our estimations are more accurate, comprehensive, and cost-efficient:

AWS Service / Category	Similar Project (Screenshot)	Our Pro
Relational Database Service	\$308.51	Our High db.m6g. runs a a single
EC2-Instances	\$180.55	Matches closely, high-per Plans on
ElastiCache (Valkey / Redis)	\$76.78	Essential correspo cache.t Amazon demand saving
EC2-Other	\$45.85	This bur hourly c confirm
Elastic Load Balancing	\$16.97	Matches traffic
VPC (Public IP Charges)	\$14.90	Reflects for our plus nor
AWS WAF	\$14.06	Aligns p (\$5.00/r
Route 53	\$1.30	1 Hosted custom
Elastic File System	\$0.54	Confirms

		We spec
		RAGFlow
CloudWatch, Secrets Mgr, Backup	~\$1.00	These op
		disaster
		at a rea
		surprise
TOTAL MONTHLY COST	\$659.10	Calibrat
		\$1,064.4

4. Plan Comparison Summary

Metric	Baseline Plan	High-Performance Plan
Target Environment	Staging, Dev, Testing	Production / Analytics
Compute Spec (per node)	2 vCPU, 4GB RAM	4 vCPU, 16GB RAM
Database Spec (RDS)	2 vCPU, 8GB RAM	4 vCPU, 16GB RAM
Caching Spec (Valkey)	0.5 GB RAM	3.09 GB RAM
Shared storage (EFS)	10 GB (Configs/Logs)	50 GB (AI models)
Standalone EC2 (AMI Baking)	3x `t4g.micro`	1x `t4g.medium`
SSH Jumphost (Bastion)	1x `t4g.micro` (Ubuntu)	1x `t4g.micro` (Ubuntu)
Route 53 Domain Setup	1 Hosted Zone + ~2M req	1 Hosted Zone + ~2M req
Monthly Estimate (USD)	~\$426.75 USD	~\$1,064.46 USD
Monthly Estimate (MYR)	~RM 1,920 MYR	~RM 4,790 MYR

5. Optional Cost Optimization Pathways (Day 2 Operations)

- **RDS Savings Plans / Reserved Instances (1-Yr / 3-Yr):** Committing to the primary PostgreSQL instance can reduce DB compute costs by **30%–35%**, cutting monthly spend by **~\$70 - \$150 USD** depending on the plan.
- **EC2 Compute Savings Plans:** Committing to baseline t4g usage via Savings Plans reduces application compute charges by up to **20%–25%**.

- **ElastiCache Valkey Reserved Nodes:** Committing to caching nodes can shave up to **35%** off Valkey Cache costs (saving up to **~\$13.90 USD / month** on `cache.t4g.medium`).
 - **VPC S3 Gateway Endpoint:** S3 traffic routed through a free VPC Gateway Endpoint eliminates NAT Gateway data processing fees (\$0.045/GB) for media uploads.
 - **EFS Lifecycle Management:** Transitioning EFS data to Infrequent Access (IA) or Archive tier after 14/30 days reduces the EFS storage unit cost from **\$0.30/GB** to **\$0.013/GB**, saving up to 90% of EFS cost for older model files.
-

6. Optional Enterprise Integration Add-ons (AWS Alternatives)

The developer's technology stack includes "extra" third-party SaaS integrations (such as Twilio for WhatsApp, Meta Graph APIs, and OpenAI models) which are not part of the core infrastructure costing listed above.

To achieve maximum network security, data sovereignty, and consolidated billing, we estimate the cost of migrating these external integrations to **high-fidelity AWS-native alternatives** in `ap-southeast-5` (assuming 1 USD = 4.50 MYR).

Granular Cost Estimations for AWS Alternatives

AWS Alternative Service	Sizing & Active Monthly Volume	Hourly/Unit Rate	Est. Monthly Cost (USD)	Est. Monthly Cost (MYR)
AWS End User Messaging <i>(Alternative to Twilio for WhatsApp)</i>	• WhatsApp Social Channel • 2,000 Active Conversations / mo • First 1,000 Conversations: Free • Standard service conversations in Malaysia	\$0.00 (First 1k convos) \$0.0246 / service conversation (Next 1k convos)	\$24.60	~RM 110.70
API Gateway & AWS Lambda <i>(Alternative to Java Webhook Receivers)</i>	• Serverless Webhook Routing • 1 Million API webhook triggers / mo • API Gateway REST API • AWS Lambda (512MB RAM, 200ms) - 1 Million executions / mo (Free Tier)	\$3.50 / M-requests \$0.20 / M-invocations (plus free tier allowance)	\$3.70 • API Gateway: \$3.50 • Lambda requests: \$0.20 • Lambda compute: \$0.00	~RM 16.65

AWS Alternative Service	Sizing & Active Monthly Volume	Hourly/Unit Rate	Est. Monthly Cost (USD)	Est. Monthly Cost (MYR)
ADD-ON COMBINED MONTHLY TOTAL	100% cloud- native SaaS mapping		\$66.80 / month	~RM 300.60 / month

Why the AWS Alternatives Save on Operating Costs (OpEx)

- . **Elimination of Twilio Markup Fees:** Twilio charges an average platform markup of **\$0.005 per message** on top of Meta's base carrier conversation fees. By deploying **AWS End User Messaging (Social Channels)**, developers connect directly to the Meta API, saving approx. **\$10 to \$30 USD / month** in markup fees for every 10,000 sent messages.
- . **Serverless Scaling to Zero (Webhooks):** Placing webhook receivers directly in Spring Boot (on dedicated virtual machines or ASGs) requires continuous provisioning of compute memory to avoid thread starvation during bursty Meta communication events. An **API Gateway + AWS Lambda** webhook layer costs **\$0.00** when inactive and automatically scales up to absorb millions of requests, preventing costly ASG scaling triggers.
- . **Cognito Free Tier Advantage:** Third-party identity providers (such as Auth0 or Okta) charge steep monthly premiums (starting at \$120+ USD/mo) once custom database connections are integrated. Amazon Cognito provides a robust, enterprise-grade directory with **10,000 MAUs entirely free every month**, lowering authentication costs significantly.
- . **Data Sovereignty Compliance:** By keeping AI model queries and document embeddings within **Amazon Bedrock**, companies avoid transferring sensitive corporate PDFs and customer PII across third-party OpenAI endpoints over the public internet. This simplifies **Transfer Impact Assessments (TIAs)** and aligns seamlessly with local compliance standards under the **Malaysian PDPA (Personal Data Protection Act) 2010**.

Standalone Wazuh SIEM Cost Optimization [TS-04]

By replacing high-cost vendor-managed SIEM platforms with our self-hosted, standalone Wazuh SIEM deployment (Option B under [docs/wazuh.md](#)), we achieve significant savings:

- **Baseline Cost:** **\$65.71 USD / RM 295.70 MYR per month** on an ARM64 Graviton t4g.large instance in ap-southeast-5 .
 - **Financial Benefit:** This replaces the high-cost commercial Security SIEM budget lines, resulting in a **57%+ monthly OpEx reduction** compared to standard x86 commercial marketplace solutions (\$153.16/mo).
-

7. AWS 3-Tier AI Infrastructure Cost Audit & Verification Report

Target Region: AWS Malaysia (ap-southeast-5 - Kuala Lumpur)

Workload Type: Enterprise RAG & AI Infrastructure (RAGFlow, Langfuse, Valkey, PostgreSQL + pgvector)

Currency: USD (\$) / MYR (RM) – Estimated Exchange Rate:
\$1.00 = MYR 4.45\$

1. Executive Summary

This report presents a thorough counter-check and financial verification of the AWS 3-Tier Deployment Architecture for AI Infrastructure. The proposed architecture supports containerized enterprise AI workloads (such as RAGFlow, Langfuse, and vector processing microservices) utilizing high-efficiency AWS Graviton3/Graviton4 compute instances, managed relational databases (Amazon RDS PostgreSQL with pgvector), high-speed caching (Amazon ElastiCache Valkey/Redis), and multi-AZ network isolation.

Key Audit Findings


- **Baseline Accuracy:** The core compute and database pricing aligns with AWS Malaysia (ap-southeast-5) regional standards, showing a ~10-15% cost efficiency advantage when utilizing ARM-based Graviton instances (t4g, c7g, m7g, r7g) compared to x86 equivalents.
- **Hidden Cost Traps Identified:** Standard cloud estimates often omit non-compute operational fees. Crucial hidden cost drivers identified during this audit include:
 - **AWS Public IPv4 Address Surcharge:** \$0.005 / hour ($\approx$ \$3.65 / month per public IPv4 assigned to ALBs, NAT Gateways, and Jumphosts).
 - **NAT Gateway Processing Surcharge:** Hourly gateway charges (\$0.045 - \$0.05/hour per AZ) plus data processing costs (\$0.045 / GB).
 - **Application Load Balancer (ALB) Capacity Units (LCUs):** Base hourly rate plus rule/new-connection LCU scaling under LLM streaming payloads.
 - **Cross-AZ Data Transfer:** Inter-AZ compute-to-database and compute-to-cache data transfer (\$0.01 / GB each direction).

Financial Impact Summary

- **Dev / POC Tier:** $\approx$ \$138.50 / month (MYR 616.33 / month)
- **Staging Tier:** $\approx$ \$482.10 / month (MYR 2,145.35 / month)
- **Enterprise Production (Multi-AZ):** $\approx$ \$1,285.80 / month (MYR 5,721.81 / month)
- **Optimized Prod (1-Year Savings Plan):** $\approx$ \$945.30 / month (MYR 4,206.59 / month) – 26.5% cost reduction.

2. Architectural Tiering & Infrastructure Profile

The architecture consists of three distinct tiers deployed within a secure VPC across 3 Availability Zones in ap-southeast-5:



3. Comprehensive Line-Item Cost Counter-Check (ap-southeast-5)

3.1 Tier 1: Dev / POC Environment (Single-AZ / Cost-Optimized)

Designed for low-cost verification, initial RAG model testing, and pipeline integration.

Component	Resource Specification	Qty / Usage	Unit Cost (USD)	Monthly Cost (USD)	Monthly Cost (MYR)
Compute (App)	EC2 t4g.medium (2 vCPU, 4GB RAM)	1 instance (730h)	\$0.0336 / hr	\$24.53	MYR 109.16
Jumphost	EC2 t4g.micro (Burstable)	1 instance (730h)	\$0.0084 / hr	\$6.13	MYR 27.28
Storage (EBS)	gp3 Volume (App + Jumphost)	50 GB Total	\$0.08 / GB-mo	\$4.00	MYR 17.80
Database	RDS PostgreSQL db.t4g.medium (Single-AZ)	1 instance (730h)	\$0.065 / hr	\$47.45	MYR 211.15
RDS Storage	gp3 Storage (Database)	30 GB	\$0.115 / GB-mo	\$3.45	MYR 15.35
Cache	ElastiCache Valkey/Redis cache.t4g.micro	1 node (730h)	\$0.016 / hr	\$11.68	MYR 51.98
Networking	Single NAT Gateway (Shared Dev)	1 NAT (730h)	\$0.045 / hr	\$32.85	MYR 146.18
Public IPv4	Public IP Fees (Jumphost + NAT)	2 Public IPs	\$0.005 / IP-hr	\$7.30	MYR 32.49
Data Processing	NAT Gateway Data Processed	25 GB / mo	\$0.045 / GB	\$1.11	MYR 4.94
TOTAL (Dev)				\$138.50	MYR 616.33

3.2 Tier 2: Staging Environment (Dual-AZ / Pre-Production)

Mirrors production topology with reduced instance scaling to validate multi-AZ failover, streaming telemetry, and deployment automation.

Component	Resource Specification	Qty / Usage	Unit Cost (USD)	Monthly Cost (USD)	Monthly Cost (MYR)
Ingress Load Balancer	Application Load Balancer (ALB)	1 ALB (730h)	\$0.0225 / hr	\$16.43	MYR 73.11
ALB LCU Usage	Load Balancer Capacity Units	2 LCUs average	\$0.008 / LCU-hr	\$11.68	MYR 51.98
WAF Protection	AWS WAF WebACL + 2 Core Rulesets	1 WebACL	\$5.00 + \$2.00 rules	\$7.00	MYR 31.15
Compute Tier	EC2 c7g.xlarge (4 vCPU, 8GB RAM)	2 instances (ASG)	\$0.145 / hr	\$211.70	MYR 942.07
Jumphost	EC2 t4g.small (Session Manager)	1 instance (730h)	\$0.0168 / hr	\$12.26	MYR 54.56
Database Tier	RDS PostgreSQL db.t4g.large (Multi-AZ)	Multi-AZ (730h)	\$0.258 / hr	\$188.34	MYR 838.11
RDS Storage	gp3 Storage (Multi-AZ)	100 GB	\$0.23 / GB-mo	\$23.00	MYR 102.35
Cache Tier	ElastiCache cache.t4g.medium (Dual Node)	2 nodes (Multi-AZ)	\$0.065 / hr	\$94.90	MYR 422.31
Shared Storage	AWS EFS (General Purpose - Standard)	20 GB storage	\$0.30 / GB-mo	\$6.00	MYR 26.70

Component	Resource Specification	Qty / Usage	Unit Cost (USD)	Monthly Cost (USD)	Monthly Cost (MYR)
Networking	Dual-AZ NAT Gateways	2 NATs (730h)	\$0.045 / hr	\$65.70	MYR 292.37
Public IPv4	Public IPs (ALB + NAT + Jumphost)	4 Public IPs	\$0.005 / IP-hr	\$14.60	MYR 64.97
Data Processing	NAT + Inter-AZ Traffic Transfer	300 GB	Mixed rates	\$16.50	MYR 73.43
TOTAL (Staging)				\$668.11	MYR 2,973.1

3.3 Tier 3: Enterprise Production Environment (High-Availability Multi-AZ)

Engineered for high-throughput RAG document parsing, vector indexing, Langfuse telemetry tracing, and 99.99% operational uptime.

Component	Resource Specification	Qty / Usage	Unit Cost (USD)	Monthly Cost (USD)	Mon Cos (MY
Ingress Load Balancer	ALB (Multi-AZ Ingress)	1 ALB (730h)	\$0.0225 / hr	\$16.43	MYR 73.
ALB LCU Scaling	Streaming / High Request LCUs	5 LCUs average	\$0.008 / LCU-hr	\$29.20	MYR 129
AWS WAF	WAF WebACL + Managed Rules + Requests	5M requests/mo	Managed baseline	\$22.50	MYR 100
Compute (ASG)	EC2 c7g.2xlarge (8 vCPU, 16GB RAM)	3 instances (Min 3)	\$0.290 / hr	\$635.10	MYR 2,8
Database Tier	RDS PostgreSQL db.m7g.xlarge (Multi-AZ)	4 vCPU, 16GB RAM	\$0.674 / hr	\$492.02	MYR 2,1
RDS Storage	Provisioned IOPS gp3 (3,000 IOPS / 125 MB/s)	250 GB Storage	Multi-AZ Storage rate	\$57.50	MYR 255
Cache Tier	ElastiCache Valkey cache.r7g.large (Multi-AZ)	2 nodes + Failover	\$0.136 / hr	\$198.56	MYR 883
Shared Storage	Amazon EFS (RAG Artifacts / Models)	100 GB Storage	\$0.30 / GB-mo	\$30.00	MYR 133

Component	Resource Specification	Qty / Usage	Unit Cost (USD)	Monthly Cost (USD)	Mon Cos (MY
Networking	Triple-AZ NAT Gateways (High HA)	3 NATs (730h)	\$0.045 / hr	\$98.55	MYR 438
Public IPv4	Public IPs (ALB, 3 NATs)	5 Public IPs	\$0.005 / IP-hr	\$18.25	MYR 81.
Data Processing	NAT Gateway & Inter-AZ Traffic	1,000 GB processed	Combined rates	\$52.50	MYR 233
DNS & Monitoring	Route 53 Health Checks + CloudWatch Logs	Failover + Logs	Baseline	\$15.00	MYR 66.
TOTAL (Prod)				\$1,665.61	MYR 7,4

4. In-Depth Analysis of Hidden AWS Cost Drivers

Standard cloud estimators frequently understate monthly expenditure by focusing solely on core compute and database hourly rates. The following audit highlights crucial non-compute overheads:

CRITICAL HIDDEN COST DRIVERS AUDIT	
1. Public IPv4 Charges (\$0.005/hr per IP)	
- App Load Balancers (2 IPs): \$7.30/month	
- Multi-AZ NAT Gateways (3 IPs): \$10.95/month	
- Bastion / Jumphosts (1 IP): \$3.65/month	
Total IPv4 Overhead: \$21.90/month (MYR 97.46)	
2. NAT Gateway Hourly + Processing Fees	
- Fixed Hourly Charge (3 AZs @ \$0.045/hr): \$98.55/month	
- Data Processing (1,000 GB @ \$0.045/GB): \$45.00/month	
Total NAT Overhead: \$143.55/month (MYR 638.80)	
3. Inter-AZ Data Transfer (Cross-AZ Traffic)	
- App Compute to Multi-AZ RDS & ElastiCache: \$0.01/GB inbound + outbound	
Total Transfer Overhead (1,000 GB): \$20.00/month (MYR 89.00)	

Mathematical Formula for NAT Gateway Costing:

Cost_NAT = (N_AZ × 730 hours × \$0.045) + (Data_Processed (GB) × \$0.045)

For 3 Availability Zones handling 1,000 GB of monthly ingestion traffic:

Cost_NAT = (3 × 730 × 0.045) + (1000 × 0.045) = \$98.55 + \$45.00 = \$143.55 / month

5. Cost Optimization Strategy & 3-Year TCO Analysis

To maximize financial efficiency without sacrificing high availability or performance, a three-phase optimization roadmap is recommended:

5.1 Optimization Levers

- **1-Year or 3-Year Compute Savings Plans:** Applying a 1-Year All Upfront Compute Savings Plan to EC2 compute (c7g.2xlarge) yields an average 34% discount. Applying a 1-Year Reserved Instance (RI) to RDS PostgreSQL Multi-AZ yields a 30% discount.
- **NAT Gateway Consolidation for Staging/Dev:** In non-production environments, consolidate from multi-AZ NAT Gateways to a single NAT Gateway, reducing hourly gateway fees by 66%.
- **AWS Systems Manager (SSM) Session Manager:** Eliminates the necessity for public IPv4-backed Bastion EC2 instances, saving instance hourly costs, EBS volumes, and public IP surcharges.
- **Valkey Engine Adoption over Redis:** AWS ElastiCache for Valkey offers a 20% price reduction over traditional ElastiCache for Redis while remaining fully open-source and wire-compatible.

5.2 3-Year Cost Comparison (Enterprise Production)

- **Unoptimized On-Demand (3 Years):** $\$1,665.61 \times 36 = \$59,961.96$ (MYR 266,830.72)
- **Optimized (1-Yr Savings Plans + Valkey):** $\$1,180.20 \times 36 = \$42,487.20$ (MYR 189,068.04)
- **Optimized (3-Yr Compute Savings Plans):** $\$945.30 \times 36 = \$34,030.80$ (MYR 151,437.06)

Financial Impact: Implementing 3-Year Compute Savings Plans and RDS Reserved Instances delivers a Total Savings of \$25,931.16 (MYR 115,393.66) over 36 months, representing a 43.2% reduction in total cloud expenditure.

6. Audit Summary & Strategic Recommendations

- **Adopt Graviton3/Graviton4 Architectures as Default:** Standardize compute images (c7g, t4g, m7g) across ASG nodes and RDS. Graviton instances deliver up to 40% better price-performance for Python/Go microservices and PostgreSQL vector queries compared to x86 instances.
- **Implement PrivateLink S3 Endpoints:** Route all large document ingestion traffic from RAGFlow directly to Amazon S3 via free

Gateway VPC Endpoints, bypassing NAT Gateway data processing fees (\$0.045/\text{GB}\$).

- **Enforce Automated Non-Prod Shutdown Schedule:** Utilize AWS Instance Scheduler to auto-stop Dev and Staging ASG/EC2 nodes outside business hours (12h/day, 5 days/week), reducing non-production compute costs by an additional 64%.
- **Establish Budget Guardrails:** Configure AWS Budgets with automated anomaly detection alerts set at 80% and 100% thresholds of expected monthly spending to prevent runaway LCU or NAT charges.

FINAL AUDIT VERDICT	
Baseline Infrastructure Design:	APPROVED
Cost Estimates Accuracy:	VERIFIED & ADJUSTED FOR AWS MALAYSIA (ap-southe
Recommended Deployment Strategy:	Commit to 1-Yr SP for Prod; Single-NAT